

TYPE II REGULARITY FOR NAVIER–STOKES

GUILLEM DURAN BALLESTER

ABSTRACT. We establish a local retained-branch exclusion criterion for the three-dimensional Navier–Stokes equations near a possible singular point and record the named exit rows together with their local closures. From a positive Caffarelli–Kohn–Nirenberg concentration sequence we select renormalized local windows around a retained core and analyze the resulting alternatives by compactness, pressure reconstruction, and local energy estimates. The argument combines a compact single-core exclusion criterion, a repaired-gauge representation for nondegenerate concentration cores, local Calderón–Zygmund pressure control, a Caccioppoli estimate on compact cylinders, multibubble and cascade reductions, and scale-collapse cost estimates. The retained compact scale-collapse branch is closed by local state-space routing into the scale-rigid discharge, so no additional stationary or self-similar classification theorem is used implicitly. Thin negative drift is closed as an independent terminal row by diffuse occupation-state routing. Branches that leave the retained state through selected compactness-package failure are now routed to more specific local rows. Modulation failure, exterior concentration, failure of the retained spacetime activity floor, and noncanonical carrier/cost failure are not treated as hidden inputs; they are listed as named exits and closed by their own local closure theorems.

CONTENTS

1. Introduction	4
1.1. Branch-contract provenance and ownership	5
2. Architecture and reading guide for the Type II proof	12
2.1. Diagram map	12
2.2. Proof-dependency diagram	13
2.3. Node-by-node audit table	26
2.4. Monotone retained-obstruction ledger	34
2.5. Branch-closure audit ledger	34
3. Local Single-Core Criterion	38
3.1. Setup and definitions	38
3.2. Local analytic inputs	39
3.3. Compactness and localized pressure compactness	41
3.4. Contradiction lemmas	44
3.5. Main theorem	45
4. Repaired-Gauge Construction	46
4.1. Setup and notation	46
4.2. Preliminary frame map from the selected core	47
4.3. Canonical renormalized change of variables	48
4.4. Domain localization and support	48
4.5. Transport of local quantities under the chart	49
4.6. Local repaired gauge map	50
4.7. Retained repaired-gauge selection in time	55
4.8. Pressure transport and decomposition	57
4.9. Modulation from differentiated constraints	58

4.10.	Repaired suitable structure in renormalized variables	59
4.11.	Stability under subsequences	61
4.12.	Final representation theorem	62
5.	Multibubble and Cascade Exclusion	64
5.1.	Local setting and active frames	64
5.2.	Elementary local decoupling estimates	68
5.3.	Terminal active-frame decoupling	72
5.4.	Same-point and separated-point reductions	74
5.5.	Multibubble and cascade exclusion	77
6.	Local Caccioppoli and Rough-Core Exclusion	96
6.1.	Analytic inputs and notation	96
6.2.	Compact-cylinder Caccioppoli estimate	97
6.3.	Consequences for the assembly	100
7.	Scale-Collapse Cost and Autonomous Limits	101
7.1.	Renormalized scale-collapse setting	101
7.2.	Elementary alternative	102
7.3.	Scale collapse versus localized core lower bounds	103
7.4.	Finite-cost exclusions	104
7.5.	Autonomous reduced-limit theorem from the selected-window estimates	105
7.6.	Decomposition of alternatives of the scale-collapse branch	113
8.	Critical Tightness and Windowed-Control Estimates	125
8.1.	Uniform critical tightness	125
8.2.	Tail windowed H^1 control	127
8.3.	Modulation and weighted forcing inputs	130
8.4.	Window-Integrated Modulation Forcing and Weighted Scale Components	133
8.5.	Navier–Stokes Cost-Divergence Exclusion	143
9.	Terminal Profile Decoupling and Exterior Regularity	183
9.1.	Renormalized frames and active local states	183
9.2.	Same-point and separated-point reductions	184
9.3.	Local perturbation and pressure decoupling	188
9.4.	Terminal local-state decoupling	191
9.5.	Small local residual components	198
9.6.	Exterior regularity	201
10.	Corrected Monotonicity and Scale-Drift Cost Criteria	224
10.1.	Data for the Cost Criterion	224
10.2.	Fixed-Cutoff Monotonicity	225
10.3.	Moving-Cutoff Monotonicity	228
10.4.	Negative Scale Drift	230
10.5.	Scale-Drift Cost Criteria	231
10.6.	Cost-Divergence Exclusion on the Canonical Windowwise Schedule	233
10.7.	Absolutely Continuous Repaired-Gauge Charts and Renormalized Equations	236
10.8.	Local Caccioppoli Estimate for Absolutely Continuous Suitable Branches	242
10.9.	Caccioppoli Regularity Criterion	246
10.10.	Terminal Local State-Space Compactness	249
10.11.	Closure of the thin negative-drift exit	275
10.12.	Closure of the Type II analytic-data exhaustiveness row	281
10.13.	Terminal Retained-State Assembly	286
11.	Assembly of the Local Type II Exclusion	290
11.1.	Local Entry	290

11.2.	Local Theorems Used in the Assembly	290
11.3.	Assembly of Alternatives	292
11.4.	Physical Entry into the Covered Local Type II Path	298
11.5.	Closure of the remaining named exit rows	300
11.6.	Assembly Theorem	308
12.	Conclusion	310
Appendix A.	PDE Interface Definitions Used in the Local Type II Proof	311
A.1.	Admissible Type II branches and exhaustion	311
A.2.	Repaired-gauge representation from Type II analytic data	312
A.3.	Minimal chart and gauge realization	314
A.4.	Local Retained Mass Cost Divergence	316
A.5.	Localized cost and the Type II scale-exclusion condition	317
A.6.	Local compact-cost classification	319
	References	324

1. INTRODUCTION

The three-dimensional Navier–Stokes regularity problem is governed by the interaction between local energy control, scale-critical concentration, and the possible formation of singular structures at arbitrarily small parabolic scales. Since Leray’s construction of global weak solutions [Ler34], the theory of suitable weak solutions and the partial regularity theorem of Scheffer and Caffarelli–Kohn–Nirenberg have provided the basic local language for this interaction [Sch76, CKN82, Lin98]. In that theory, singularity formation is detected by scale-invariant velocity-pressure quantities, and the critical L^3 scale plays a central role alongside the local energy inequality, pressure reconstruction, and compactness. The analytic background also includes the classical local and critical-space theory for Navier–Stokes [Lad69, Tem77, CF88, Kat84, LR02], the $L_{3,\infty}$ regularity theorem of Escauriaza–Seregin–Šverák [ESS03], and the rigidity theorem of Nečas–Růžička–Šverák for Leray self-similar profiles [NRŠ96]. These classical rigidity results are cited as background; the retained Type II branch is closed by the local state-space reductions below, not by importing an external stationary or self-similar classification theorem.

This paper studies the local Type II alternative near a possible singular point. Starting from a positive Caffarelli–Kohn–Nirenberg concentration sequence, we select renormalized local windows around a retained core and analyze all branches that can arise from the scale-center dynamics. The result in the present version is a retained-state exclusion theorem together with an explicit ledger of named exits and a closure theorem for each remaining ledger row. Pressure reconstruction, repaired-gauge compactness, terminal profile decoupling, cost comparison, and local state-space closure are proved on retained branches or used to route a first failure to one of those named exits. No additional stationary, ancient, or self-similar classification theorem is invoked implicitly.

The proof begins with local concentration and the Type I/Type II separation. For a non-degenerate retained core, the solution is rewritten in local scale-translation variables. The repaired-gauge representation produces a renormalized Navier–Stokes equation with controlled modulation terms, local pressure reconstruction modulo functions of time, and compact-window suitable estimates. Within this representation, a compact single-core branch with a retained positive velocity core and vanishing localized dissipation is impossible: local compactness gives a spatially constant limit, the zero constant contradicts retained velocity by strong L^3 convergence, and a nonzero constant gives a bounded selected-window limit, which leaves the strict selected-window retained branch and is later physically routed by the bounded-window closure theorem. Pressure-only harmonic concentration is routed to the pressure/exterior rows rather than closed by the single-core velocity argument.

A compactness limit cannot create an additional hidden terminal state. Once the pressure, chart, modulation, represented-equation, and local-energy limit interfaces have passed, the retained spacetime activity floor makes the full limit nonzero and suitable. The limit is then either singular, in which case a diagonal CKN rebase gives an actual active descendant of the represented sequence, or nonsingular, in which case it is locally bounded on a nonzero subwindow and enters the bounded-window row. This regular–singular routing is made explicit in [lemma 7.30](#) and [corollary 10.79](#); it is applied only to the full represented solution, never to a residual as an independent Navier–Stokes solution.

The remaining sections remove the ways in which a Type II branch could avoid that compact single-core mechanism. Caccioppoli estimates convert finite outer velocity-pressure control into inner windowed H^1 control, so rough-core loss is returned to an explicit compact-cylinder alternative. Multibubble and cascade arguments separate active profiles, same-point cascades, and gauge-degenerate branches from a single retained core. Scale-collapse branches are treated by localized cost identities: finite scale-collapsing cost, and the corresponding finite absolute cost,

are incompatible with a nonvanishing retained core, while thick autonomous windows lead only to the stated modulation, compactness, or retained compact scale-collapse state; the latter is routed into the scale-rigid discharge rather than into a new Liouville assumption. The diffuse thin-drift boundary face is handled by normalized negative-drift occupation blocks: persistent windows return to the thick-window ledger, while genuinely diffuse blocks make the negative scale term a vanishing local perturbation and route the branch to another local row or to a closed retained row. The terminal profile and corrected monotonicity sections then turn these local reductions into a branchwise classification: retained states are eliminated, while first failures outside the retained state are recorded as named exits.

The conclusion should therefore be read in this local state-space sense. In the retained branch, a positive local Type II concentration sequence cannot survive the repaired-gauge, compactness, multibubble, rough-core, compact-cost, and retained scale-collapse reductions. A failure of one of the local analytic rows is not left as an unstated input; it is recorded as a named state-space exit. In the final assembly, the remaining named exits are then closed by their own local routing arguments, giving the stated no-local-Type-II conclusion without importing global critical estimates.

Definition 1.1 (Canonical chart notation). The paper uses the following notation for every concentration chart and its repaired gauge:

t	physical time,
ρ	raw renormalized time,
τ	final repaired time,
$x_c(t), \lambda(t)$	raw concentration center and scale,
$U(y, \rho), Q(y, \rho)$	raw profile and pressure,
$\mu(\tau), q(\tau)$	repaired-gauge scale and translation,
$\Lambda = \lambda\mu$	final physical scale,
$X = x_c + \lambda q$	final physical center,
$V(Y, \tau), P(Y, \tau)$	final repaired profile and pressure,
$a = \partial_\tau \log \Lambda$	final logarithmic scale velocity,
$b = \Lambda^{-1} X_\tau$	final normalized center velocity.

In the two composite identities, x_c and λ are evaluated at the physical time corresponding to $\rho(\tau)$. Thus the final time and field conventions are

$$\frac{d\tau}{dt} = \Lambda^{-2}, \quad x = X + \Lambda Y, \quad u = \Lambda^{-1} V, \quad p = \Lambda^{-2} P + c(t).$$

These conventions govern every passage that compares the raw and repaired charts. A self-contained auxiliary lemma may introduce a different dummy scale, but it does not change the raw-to-final dictionary above.

Remark 1.2 (Scope of whole-space norms). A norm over $\mathbb{R}^3$ appears below only when it is an explicit hypothesis or profile-space datum, an exact scaling identity, or the norm of a compactly supported source extended by zero in a Calderón–Zygmund construction. Every estimate used to close the physical Type II branch is instead stated on a buffered compact cylinder or as an explicit local/tail split. In particular, no whole-space velocity or pressure bound is inferred from local suitable-solution data.

1.1. Branch-contract provenance and ownership. The manuscript is a consolidated proof: every local proof capsule invoked below occurs in this file. Consequently, the source-file column records `type_II_regularity.tex` together with the exact producer labels. The status “proved here” means that the displayed output is established by those producer statements. The status

“admissible-branch input” means that the row is a positive contract defining the branch on which the subsequent implication is applied. Such a row is never presented as an existence theorem: its failure is sent to the named output in the tenth column. The No state-space row in the present ledger depends on an external theorem whose statement is absent from this manuscript. Classical tools such as the Caffarelli–Kohn–Nirenberg criterion, Calderón–Zygmund estimates, and the Aubin–Lions–Simon compactness theorem are cited at their individual uses; they do not supply any of the chart, gauge, pressure-atlas, or terminal-state tuples recorded here.

For a compact interval I , write

$$Q_{R,I} := B_R \times I.$$

Every working/core entry $Q_{2R,I} \rightarrow Q_{R,I}$ includes a strict spatial buffer; whenever a temporal cutoff is used, its support is compactly contained in the displayed working interval. Thus the window column records the actual locality of the corresponding interface rather than a whole-space assertion.

Table 1: Provenance ledger for concentration, chart, gauge, pressure, and terminal-profile interfaces.

Interface	Source file	Exact producer label(s)	Exact output tuple	Target label(s)	Windows	Convergence or regularity topology	Pressure representative	Retained obstruction	Named failure output	Status
CKN entry and Type I/II split	type_II_regularities.tex	thm:local-concentration-dichotomy; thm:local-typeI-typeII-dichotomy	(x_0, T, r_n, η) , $r_n \downarrow 0$, $C((x_0, T), r_n) + D((x_0, T), r_n) \geq \eta$, and the exclusive local Type I/Type II alternative	prop:selected-window-branch; h-reduction; lem:typeII-concentration-extension-routes	physical $Q_{r_n}(x_0, T)$; fixed selected windows after parabolic rescaling	scale-invariant $L_{x,t}^3 \times L_{x,t}^{3/2}$ quantities	$p - (p)_{B_{r_n}(x_0)}(t)$	positive CKN mass and failure of the local Type I bound	empty singular set or local Type I: lem:typeII-concentration-extension-routes, cases (b),(c)	proved here
Retained velocity floor	type_II_regularities.tex	thm:active-profile-mass-floor; lem:pressure-only-critical-mass-routing	$(Q_{R,I}, \varepsilon_{loc})$; spacetime L^3 -mass at least ε_{loc}^3 and local $L_I^\infty L^3$ -limsup at least ε_{loc} ; inverse rescaling gives $C((x_0, T), \rho_n) \geq \eta_v > 0$ on the retained velocity subbranch	def:single-core-branch; lem:selected-window-mass-floor; lem:c5-uniform-retained-active-window-mass	detecting $Q_{2R,I} \rightarrow Q_{R,I}$; shrinking physical cylinders under the inverse chart	local spacetime L^3 and local $L_I^\infty L^3$	pressure is used only through the localized source and harmonic remainder	positive velocity, rather than pressure-only, concentration	regular/bounded-window, pressure/exterior, or loss-of-active-core rows: thm:active-core-loss-closed	proved here
Raw AC chart contract	type_II_regularities.tex	def:ac-initial-concentration-chart; def:c2r-minimal-representation-data, item (i)	$(x_c, \lambda, \rho, U, Q)$, with x_c, λ, ρ AC, $\rho_t = \lambda^{-2}$, and the exact velocity and pressure pullbacks	lem:initial-renormalized-equation; thm:ac-repaired-gauge-consequences; thm:c2r-minimal-representation	each compact raw y -window contained in the physical domain	AC_{loc} chart; distributional chain rule on compact sets	$p = \lambda^{-2}Q + c(t)$, so Q is modulo functions of ρ	selected Type II core and its positive CKN signal	chart extraction, profile, or domain failure: def:c2r-minimal-failures, item (i); def:c1-exhaustion-failures, items (iii),(iv)	admissible-branch input
Raw-chart consequences	type_II_regularities.tex	lem:initial-renormalized-equation; lem:pressure-reconstruction-under-charts; lem:critical-limit-scaling	distributional raw equation with $A = \partial_\rho \log \lambda$, $B = \lambda^{-1} \partial_\rho x_c$; pressure Poisson equation; critical L^3 -invariance	thm:ac-repaired-gauge-consequences; lem:c2r-chart-raw-orbit; lem:c2r-pressure-pullback	compact raw y -sets	distributions; local L^3 scaling identity	Q modulo functions of ρ	the raw chart contract in the preceding row	failure of AC or the chain-rule domain is failure of the preceding chart contract	proved here

Continued on next page

Table 1 continued

Interface	Source file	Exact producer label(s)	Exact output tuple	Target label(s)	Windows	Convergence or regularity topology	Pressure representative	Retained obstruction	Named failure output	Status
AC repaired-gauge path contract	type_II_regularities.tex	def:ac-repaired-gauge-path; def:c2r-minimal-representation-data, item (ii)	(μ, q, τ, V, P) , with $\rho_\tau = \mu^2$, $V(Y, \tau) = \mu U(\mu Y + q, \rho)$, $P(Y, \tau) = \mu^2 Q(\mu Y + q, \rho)$, and all gauge constraints	lem:repaired-physical-chart; lem:repaired-modulation-coefficients; thm:canonical-repaired-gauge-equation	compact final windows whose pullbacks remain inside the raw chart	$\mu, q, \rho \in AC_{\text{loc}}, \mu > 0$	P modulo functions of τ	repaired scale and centering constraints on the retained core	no AC gauge solve: admissible- def:c2r-minimal-failures, item (ii); strong-chart or Jacobian failure: lem:strong-gauge-chart-entry-test	branch input
Gauge transversality and consequences	type_II_regularities.tex	prop: full-jacobian; thm:quantitative-gauge-jacobian-invertibility; thm:ac-repaired-gauge-consequences	invertible corrected Jacobian on the retained chart; $(\Lambda, X) \in W_{\text{loc}}^{1,1}$, $a, b \in L_{\text{loc}}^1$; final represented equation and $d\tau/dt = \Lambda^{-2}$	thm:chart-repaired-gauge-representation; thm:compact-window-representation; thm:c2r-minimal-representation	$Q_{2R,I} \rightarrow Q_{R,I}$, with the chart buffer of $\text{prop:compact-t-support-t-transport}$	uniform compact-window convergence as functions of τ in $\text{prop:repaired-gauge-subsequence-stability}$	$p = \Lambda^{-2}P + c(t)$, P modulo functions of τ	invertible single-core gauge	gauge degeneracy, multibubble/cascade, scale-rigid reselection, or strong-chart/annular admissibility loss: lem:scale-rigid-degenerate-jacobian	proved here on the retained chart
Local pressure atlas	type_II_regularities.tex	thm:pressure-decomp; lem:pressure-stability; lem:harmonic-cross-pressure-routing	$P = P_{\text{loc}} + H + c_I(\tau)$, prop: $-\Delta P_{\text{loc}} = \partial_i \partial_j (\zeta V_i V_j)$, H harmonic and mean-zero; $P_{\text{loc}} \in L^{3/2}$, $H \in L^q$ on inner balls	selected-window-compactness; lem:terminal-pressure-decoupling; lem:c5-retained-upper-estimates-aubin-lions	$B_{2R} \supset \text{supp} \zeta$, decomposition on B_R , estimates on $B_r \Subset B_R$	strong $L_{\text{loc}}^{3/2}$ stability for P_{loc} under strong L^3 velocity convergence; weak $L_{\text{loc}}^{3/2}$ for full normalized pressure	spatial mean on the working ball plus an allowed function of time	velocity-source pressure separated from the harmonic tail	pressure reconstruction/gauge failure, exterior pressure, or pressure-only retained signal: lem:pressure-only-critical-mass-routing	proved here

Continued on next page

Table 1 continued

Interface	Source file	Exact producer label(s)	Exact output tuple	Target label(s)	Windows	Convergence or regularity topology	Pressure representative	Retained obstruction	Named failure output	Status
Selected compactness and limit routing	type_II_regularities.tex	def:selected-scale-collapse-compactness-package; lem:selected-compactness-aub-in-lions; lem:regular-singular-compact-limit-routing; cor: no-unclassified-compact-limit-obstruction	$(V_\infty, P_\infty, a_\infty, b_\infty)$, strong local L^3 convergence, represented suitability, and the exhaustive named-interface/active-descendant/bounded-window routing	thm:selected-compactness-failure-routed; lem:strong-diffuse-tail-limit-row; thm:ordered-local-state-space-exhaustion	$B_{2R} \times I \rightarrow B_R \times I \rightarrow Q' \in B_R \times I^\circ$	strong local L^3 for velocity; strong local normalized $L^{3/2}$ for pressure; weak $L^2 H^1$, with lower semicontinuity for dissipation	local Calderón–Zygmund pressure plus harmonic remainder, modulo functions of time	the retained spacetime activity floor gives $V_\infty \neq 0$	first limit-interface failure, an actual CKN-active descendant, or the bounded-window row	proved here
Terminal decomposition contract	type_II_regularities.tex	def:admissible-terminal; def:critical-nprofile-decomposition-local; def:item (i); def:active-critical-profile-data	active classes $(A_j, \lambda_j^n, x_j^n, I_j, W_j^n)$ and $V_n = U_n + S_n$, with parameter grouping and compact-window suitable bounds	thm:terminal-decoupling; thm:critical-nprofile-decomposition-discharged; thm:terminal-nonlinear-decoupling	every $Q_{2R,I} \rightarrow Q_{R,I}$	local $L_I^\infty L^3 \cap L_I^\infty L^2 \cap L_I^2 H^1$; residual measures and occupation coordinates when componentwise time compactness is absent	$P_n - (P_n)_{B_{2R}}(\tau) \in L^{3/2}(Q_{2R,I})$	active mass floor of each retained ancestor	critical-mass loss, parameter-tail/exterior, hidden-profile, diffuse-tail, or profile-completeness row: thm:terminal-adjacent-label-routing	admissible-branch input
Terminal ancestry and grouping	type_II_regularities.tex	lem:retained-frame-maximality; def:terminal-active-frame; thm:terminal-coordinate-frame-construction-profile-selection	for every descendant: physical concentration point, scale-center parameters, comparable compound ancestor, retained component U_n , and residual S_n	thm:terminal-window-profile-completeness; lem:terminal-velocity-decoupling	fixed compact terminal frame and its buffered working cylinder	convergence or escape of relative scale-center parameters; local state-space topology	pressure representative inherited from the terminal pressure atlas	ancestry of every retained active mass carrier	comparable descendants are absorbed; separated, cascade, exterior, or unclassified residuals enter thm:local-residual-exhaustion	proved here

Continued on next page

Table 1 continued

Interface	Source file	Exact producer label(s)	Exact output tuple	Target label(s)	Windows	Convergence or regularity topology	Pressure representative	Retained obstruction	Named failure output	Status
Terminal residual completeness	type_II_regularities.tex	prop:residual -spike-routing; lem:spacetime-small-terminal-residual; thm: local-residual-exhaustion; thm:critical-n-s-profile-decomposition-discharged	either a visible residual receives a named first row or belongs to the retained comparable class; on the retained branch $S_n \rightarrow 0$ in $L^\infty(I; L^3(B_R))$, with vanishing stress, pressure, and residual errors in the terminal topology	thm: terminal-window-profile-completeness; thm:terminal-nonlinear-decoupling; thm:terminal-critical-local-compactness	$Q_{2R,I} \rightarrow Q_{R,I}$	$S_n \rightarrow 0$ in local $L_I^\infty L^3$; stress in $L_I^1 L^2$; pressure gradient and equation error in $L_I^1 H^{-1}$	cross pressure modulo spatial means; harmonic failure routed separately	nonzero residual or retained descendant ancestry	named adjacent row, diffuse tail, exterior row, or retained comparable grouping: thm:terminal-adjacent-label-routing; lem:diffuse-critical-tail-row-closed	proved here

Proposition 1.3 (Branch-contract discharge and failure routing). *Every use of a branch contract in Table 1 has one of the following two forms: either the contract is produced by the cited proved-here row, or its first failure is assigned to the named output in the same ledger row. In particular, the raw AC chart, the AC repaired-gauge path, and the admissible terminal decomposition are not asserted to exist by their definitions. No final unconditional theorem uses any of those three contracts without the corresponding first-failure alternative.*

Proof. For concentration entry, theorem 3.6 and lemma 10.119 produce the physical CKN sequence and the exclusive Type I/Type II split. On the retained active subbranch, theorem 9.33 and lemma 11.32 either produce a velocity carrier, whose inverse parabolic rescaling is the stated positive C -floor, or assign the signal to the regular, pressure/exterior, or loss-of-active-core row.

The raw chart and gauge-path rows are definitions of the positive representation contract. Once they hold, the distributional equations, pressure pullback, final chart, and modulation coefficients follow from lemmas 10.47, 10.49, 10.51 and 10.53 and theorem 10.52. If either contract is absent, the first failure is exactly one of definition A.18; corollary A.19 retains that failure as the chart or gauge alternative instead of invoking the representation conclusion. Corrected Jacobian transversality on the retained strong chart is proved by propositions 4.17 and 4.18; failure of that test is routed by lemmas 4.9 and 5.39.

The pressure interface is analytic rather than definitional: theorem 4.25 constructs the localized Calderón–Zygmund part and the mean-zero harmonic remainder, while lemma 3.12 states exactly the strong topology available for the localized part. A nonvanishing harmonic remainder or an unavailable reconstruction is assigned by lemmas 9.22 and 11.32 to the pressure/exterior or pressure-reconstruction row.

The selected compactness package gives strong local velocity convergence by lemma 7.29. On the pass branch, the same lemma passes the represented equation and local energy inequality to the limit. The retained spacetime floor keeps that limit nonzero, and lemma 7.30 then produces exactly a singular active descendant or a regular bounded subwindow. The subsequent scale-center routing and retained-class identification are supplied by corollary 10.79. Thus suitability of the limit and the regular–singular split are proved interfaces, not new terminal assumptions.

Finally, the initial terminal splitting and its pointwise critical bounds are the admissible-family contract of definitions 5.16 and 10.68. The ancestry and grouping rules are proved in lemma 10.77 and theorem 10.134. The residual is then routed by proposition 10.80: a fixed-window occupation or active parabolic rebase enters the augmented local coordinates of theorem 10.86, while a spacetime-small selected-time spike is perturbative by lemma 9.30 and is absorbed during ancestry terminalization. The resulting finalized residual converges to zero by theorem 10.87. Thus the terminal completeness conclusion is proved here and is not part of the input contract.

At the assembly level, the preceding first failures enter covered local state-space rows by corollaries 10.147 and 11.26. Their remaining closures are supplied by corollary 11.36 before theorem 11.38 is invoked. Hence the unconditional conclusion does not depend on an interface marked only as an admissible-branch input. $\square$

2. ARCHITECTURE AND READING GUIDE FOR THE TYPE II PROOF

This section is an audit guide to the complete branch structure of the local Type II argument. The definitions, statements, and proofs in the remainder of the manuscript are the single source of mathematical truth; the diagrams only index their logical order and introduce no additional hypothesis or proof step. The initial object is the residual or singular profile at a possible terminal singular point. It is interpreted as a physical local Type II germ by [definition 11.20](#); [lemma 11.23](#) then supplies the positive concentration sequence used by the ordered state-space analysis. All compactness, pressure, energy, and profile inputs in the guide are local on compact cylinders attached to this incoming germ. No diagram node assumes a whole-space critical norm or any other global estimate.

The diagrams display the live branch, rather than every fact already proved about it. Once a hypothesis has entered the retained-obstruction ledger, it remains available after subsequence extraction, reselection, and passage to a terminal frame. Such monotone ledger facts are therefore not represented by additional backward arrows. Rectangles are states or interfaces with one incoming and one outgoing branch, diamonds are exhaustive routing tests, and solid ellipses are closed branches. A dashed ellipse is not a new node: it is a local copy of the numbered solid ellipse named inside it, used when an edge routes directly to a closure drawn in another part. Node [1] is the unique root. Every other numbered node is reached from [1], and every terminal node is a solid ellipse or is represented locally by such a dashed closure copy.

The shared routing diamond [56] receives the first-failure continuations from the earlier parts. The ordered exhaustion theorem [theorem 11.17](#) assigns exactly one first failed row to any branch. Diamond [56] classifies that single row, and Parts IV–XII display its rank-decreasing closure. A theorem that closes several cases by an internal alternative is represented by a diamond followed by its individual closure leaves. No closure edge returns to the routing layer.

2.1. Diagram map. The proof-dependency diagram is one logical directed acyclic graph split across twelve numbered parts. Bracketed node numbers are unique and consecutive in source order. An arrow labelled “continue at [n]” and the corresponding incoming arrow in the later part are the same graph edge.

Table 2: Map of the twelve numbered parts of the Type II proof-dependency diagram.

Part	Nodes	Branch resolved	Principal internal sources
I	[1]–[16]	Physical entry, C1 admissibility, repaired representation, unique-core selection, and the retained velocity carrier	definition 11.20 and theorems 10.122, 11.4 and 11.6
II	[17]–[41]	Compact estimates, rough-core loss, compact-core dissipation, scale collapse, and the canonical cost channel	theorems 10.42, 11.7 and 11.9 and lemma 11.10
III	[42]–[55]	Terminal orbit sampling, compactness alternatives, profile completeness, residual exhaustion, and retained terminal exclusion	theorems 9.16, 10.132, 10.136 and 10.144
IV	[56]–[63]	Classification of the unique first-failure row and dispatch to its closure proof	lemma 11.34 and definition 11.35

V	[64]–[70]	Local regularization and well-founded reselection of a bounded selected window	lemmas 11.11 and 11.15 and theorem 11.16
VI	[71]–[78]	Unbounded modulation windows and their signed or geometric first-failure rows	lemma 10.106 and theorem 11.29
VII	[79]–[87]	Bounded modulation-defect alternatives	lemma 11.28 and theorem 11.29
VIII	[88]–[95]	Canonical versus first-failed noncanonical cost validation	lemma 8.73 and theorem 11.31
IX	[96]–[110]	Retained active-core, finite frame reselection, and local critical-mass alternatives	lemma 10.78 and theorem 11.33
X	[111]–[116]	Ordered pressure reconstruction, local-source, and velocity-label alternatives	lemma 11.32 and theorem 11.33
XI	[117]–[127]	Well-founded nested pressure-only and diffuse-pressure alternatives	lemmas 9.67 and 9.68
XII	[128]–[135]	Critical-tightness and diffuse critical-tail alternatives	lemmas 10.103 and 10.104

2.2. Proof-dependency diagram. The twelve parts below form a single flow. Parts I and II expand the ordered local tests in [theorems 11.17](#) and [11.18](#). The only retained local state not closed in Part II is passed at [41] to the terminal analysis in Part III. All first-failure rows are passed to [56] and closed by the explicit alternatives in Parts IV–XII. The unconditional conclusion is read from the leaves: every path from [1] ends at one of the displayed closed ellipses.




FIGURE 1. Proof-dependency diagram, Part I. The physical entry test [2] either closes at [3] or supplies the concentration sequence [4]. The ordered tests [5], [8], [11], and [14] follow the order in [theorem 11.17](#). A successful route reaches the retained velocity core [16]. Each first-failure rectangle has one output, namely the continuation to the classifier [56].




FIGURE 2. Proof-dependency diagram, Part II. This part continues the retained core [16]. The no-branch of [23] is the compact single-core route and is exhausted by the three alternatives at [24]. The yes-branch is the scale stratification of [theorem 11.9](#). Finite cost closes at [36]; thin drift either reaches the closed state [38] or names a row for [56]; modulation and compactness failures also go to [56]. The only surviving residual is the retained terminal tail [41].




FIGURE 3. Proof-dependency diagram, Part III. The retained terminal tail [41] supplies the only incoming branch. Local state selection constructs the terminal frames at [42], and the repaired representation already retained from Part I gives orbit sampling at [43]. These are forward interfaces, not new alternatives. The local compactness test is [44]. A structural or exterior alternative closes at [49], and a visible residual is assigned to one named row at [51]. When neither occurs, terminal completeness gives [52], the retained compact tests give [53], all nine expanded terminal data hold at [54], and [theorem 10.144](#) closes the retained branch at [55].




FIGURE 4. Proof-dependency diagram, Part IV. The incoming arrow is the disjoint union of the first-failure continuations listed above [56]. On any run only one such row occurs. The first two outputs are discharged at [57] and [59]. Nodes [60]–[63] are the four rank-1 rows of [definition 11.35](#); each has one continuation to the part that expands its closure theorem.





FIGURE 5. Proof-dependency diagram, Part V. The bounded selected-window theorem first regularizes the bounded region. If local concentration remains, the well-founded reselection theorem gives exactly the three outcomes at [68]–[70]. Each named-row outcome refers to the unique first row selected by the ordered ledger, not to an unresolved union of rows.



TYPE II REGULARITY

FIGURE 6. Proof-dependency diagram, Part VI. The first dichotomy in the autonomous-modulation closure theorem is bounded versus unbounded window integrals. In the unbounded case, the signed and geometric stratification produces the five individually displayed closed rows. The bounded case continues at [79].




FIGURE 7. Proof-dependency diagram, Part VII. The bounded-defect routing lemma gives exactly the three outcomes at [80]–[82]. The constant and perturbative outcomes contradict the defining failure in the autonomous-modulation alternative. A visible defect is split into its signed, translation, chart, or active-core component before closure.




FIGURE 8. Proof-dependency diagram, Part VIII. The noncanonical-cost theorem first asks whether all ordered validation tests pass. Passing identifies the canonical identity-cost channel and gives [90]. Failure selects exactly one first-failure layer at [91]; the four displayed leaves use the closure theorem belonging to that layer.




FIGURE 9. Proof-dependency diagram, Part IX. The final ledger family is divided at [96]. The retained active-core branch uses the local critical-mass trichotomy [97]. When the activity floor fails, [99] exhausts the finite local frame list supplied by lemma 10.77; [100] records the first surviving outcome. If no velocity frame remains, [101] displays the four possibilities used in the proof of theorem 11.33. The upper-bound alternative is split at [103] into the individual labels of lemma 10.78. Pressure-only, nested pressure, and critical-tightness rows continue to their separate parts.




FIGURE 10. Proof-dependency diagram, Part X. The proof of [lemma 11.32](#) first tests the pressure reconstruction at [111], then the localized source at [112], and only then runs the local CKN test [113]. A velocity carrier is classified at [116], which also receives the velocity-supported outcomes of [117] and [122]. Dashed ellipses are references to terminal closures proved and drawn as solid ellipses in Part IX.




FIGURE 11. Proof-dependency diagram, Part XI. The five rows of [lemma 9.67](#) are explicit at [117]. The same-point row splits at [119]. The proof of [lemma 9.68](#) first makes the atomic/diffuse test [121], then classifies a realized atomic support profile at [122]. Dashed ellipses refer directly to terminal closures drawn elsewhere.




FIGURE 12. Proof-dependency diagram, Part XII. The six alternatives of lemma 10.103 leave [128] on separate edges. The compact-package and diffuse-tail cases are further split at [129] and [131]. The proof of lemma 10.104 tests for a lower named face first, then applies the atomic/diffuse dichotomy at [133]. Dashed ellipses refer to terminal closures already drawn in Parts IX–XI.

2.3. Node-by-node audit table. The table identifies the result in the body of the paper that justifies every diagram node. Its arrows record the live case split. Hypotheses established at an earlier node remain available through the monotone ledger summarized in [Table 4](#).

Table 3: Node-by-node audit ledger for the Type II proof-dependency diagram.

Node	State or exhaustive test	Source in the proof	Successor or closure
[1]	Residual or singular profile at a possible terminal point	definition 11.20 and lemma 11.22	Test the physical entry at [2]
[2]	Physical local Type II germ test	definition 11.20 and theorem 11.3	Yes gives [4]; no closes at [3]
[3]	Regular or local Type I closure	theorem 11.3	Terminal entry branch
[4]	Positive CKN concentration sequence	lemma 11.23 and definition 11.2	Run the ordered C1 test [5]
[5]	C1 admissibility test	theorems 10.122 and 11.17	Pass gives [7]; first failure gives [6]
[6]	C1 first-failure row	lemma 11.24 and corollary 11.26	Continue to the classifier [56]
[7]	Retained analytic data	theorem 11.25 and corollary 11.26	Test representation at [8]
[8]	Repaired representation and strong-chart test	corollary A.19 and lemma 4.9	Pass gives [10]; first failure gives [9]
[9]	Representation or gauge first failure	corollaries A.12 and A.19	Continue to [56]
[10]	Canonical represented local equation, pressure atlas, and suitability	theorems 4.25 and 10.52 and proposition 4.33	Test core uniqueness at [11]
[11]	Unique nondegenerate core test	theorems 11.4 and 11.6	Yes gives [13]; no closes at [12]
[12]	Multibubble, cascade, or gauge-degenerate closure	theorem 11.6	Terminal structural branch
[13]	One retained active core	theorem 7.22 and lemma 10.77	Test the velocity carrier at [14]
[14]	Retained velocity-carrier test	lemma 11.32	Yes gives [16]; no gives [15]
[15]	Pressure-only, exterior, or active-core-loss row	lemma 11.32 and corollary 9.73	Continue to [56]
[16]	Represented retained velocity core	theorem 7.22 and proposition 7.14	Continue to [17]
[17]	Compact upper-estimate and pressure-stability test	proposition 7.14 and theorem 4.25	Pass gives [19]; failure gives [18]
[18]	Compact, pressure, or ambient first failure	theorem 11.17 and corollary 8.139	Continue to [56]
[19]	Retained compact-cylinder estimates	proposition 7.14	Test windowed H^1 at [20]
[20]	Local windowed H^1 test	theorems 10.61 and 11.7	Pass gives [22]; failure closes at [21]
[21]	Rough-core/ H^1 -loss closure	theorem 11.7 and proposition 11.19	Terminal rough-core branch
[22]	Represented compact retained core	lemma 11.10 and theorem 11.18	Test the scale state at [23]
[23]	Compact-core versus scale-collapse/scale-rigid test	theorems 11.9 and 11.18	No gives [24]; yes gives [35]

Continued on next page

Table 3 continued

Node	State or exhaustive test	Source in the proof	Successor or closure
[24]	Compact-core dissipation alternatives	lemma 11.10	Vanishing gives [25], adjacent gives [29], divergent cost gives [30]
[25]	Spatially constant compact limit	theorems 3.18 and 11.5	Test its value at [26]
[26]	Zero-versus-nonzero constant test	theorem 11.5	Zero closes at [27]; nonzero gives [28]
[27]	Retained mass contradiction	theorem 11.5	Terminal compact single-core branch
[28]	Nonzero bounded selected-window row	theorems 5.27 and 11.16	Continue to [56]
[29]	Adjacent compact-core first failure	lemma 11.10 and corollary 8.139	Continue to [56]
[30]	Divergent canonical localized cost	theorem A.21	Validate the canonical channel at [31]
[31]	Canonical identity-cost validation	lemmas 8.73 and 11.30	Pass gives [33]; failure gives [32]
[32]	Noncanonical carrier or cost-validation row	theorem 11.31	Continue to [56]
[33]	Corrected monotonicity and carrier routing	theorem 8.133	Apply the carrier closures at [34]
[34]	Cost-divergence closure	theorem 10.42 and corollary 10.43	Terminal retained-cost branch
[35]	Scale-collapse stratification	theorem 11.9	Produces exactly one of [36]–[41]
[36]	Finite signed or absolute scale-cost closure	theorem 11.8	Terminal finite-cost branch
[37]	Thin-drift routing test	theorem 10.118	Closed retained row gives [38]; another named row goes to [56]
[38]	Already closed retained scale row	theorems 7.37 and 10.118	Terminal thin-drift route
[39]	Autonomous-modulation first failure	theorem 11.29	Continue to [56]
[40]	Compactness-package first failure	theorem 7.34	Continue to [56]
[41]	Retained terminal tail after all earlier local exits	theorems 10.146 and 11.17 and lemma 7.30	Construct its terminal frames at [42]
[42]	Terminal coordinate frames constructed locally from retained state selection	theorem 10.134	Apply orbit sampling at [43]
[43]	Terminal active sequence sampled from the represented orbit	theorem 10.136	Run the local compactness alternatives at [44]
[44]	Local PDE and terminal L^3 bounds test	theorem 10.132	Pass gives [46]; failure gives [45]
[45]	Terminal compactness first-failure row	theorem 10.132 and lemma 10.78	Continue to [56]
[46]	Local compact terminal-window hypothesis	theorem 10.131	Extract the complete profile family at [47]

Continued on next page

Table 3 continued

Node	State or exhaustive test	Source in the proof	Successor or closure
[47]	Local compact-window profile decomposition, ancestry, and active frames	theorems 9.16 and 10.87	Test additional profiles at [48]
[48]	Noncomparable active, separated, cascade, or exterior profile test	theorems 9.16 and 9.39	Yes closes at [49]; no gives [50]
[49]	Structural or exterior closure	theorem 11.6 and corollary 9.73	Terminal profile branch
[50]	Residual, critical-tail, or pressure-defect test	proposition 10.80 , lemma 10.98 , corollary 10.79 , and theorems 10.81 and 10.86	Defect gives [51]; otherwise [52]
[51]	Named residual, tail, or pressure row	theorem 10.86 and lemma 10.104	Continue to [56]
[52]	Retained comparable profile with perturbative residual	lemma 9.30 , corollary 10.79 , and theorems 9.16 and 10.126	Apply the retained compact tests at [53]
[53]	Critical tightness and local windowed H^1 control	theorem 10.108	Verify the expanded terminal data at [54]
[54]	All nine expanded terminal data on the retained tail	definition 10.140 and theorem 10.146	Apply the terminal exclusion at [55]
[55]	Retained terminal Type II exclusion	theorem 10.144 and corollary 10.147	Terminal retained branch
[56]	Unique first-failure closure classifier	corollary 11.26 , lemma 11.34 , and definition 11.35	Exactly one closure class [57], [58], or [60]–[63]
[57]	Class-incompatible closure	lemmas 11.22 and 11.24	Terminal contradiction to physical Type II entry
[58]	Deterministic lookup of an earlier row's recorded closure theorem	proposition 11.19 and corollary 8.139	Apply that theorem at [59]
[59]	Closure of the uniquely selected earlier row	proposition 11.19 and corollary 8.139	Terminal previously closed row
[60]	Bounded selected-window row	theorem 11.16	Continue to its alternatives at [64]
[61]	Autonomous-modulation row	theorem 11.29	Continue to the window-integral test [71]
[62]	Noncanonical cost-validation row	theorem 11.31	Continue to the validation test [88]
[63]	Final active-core ledger family	definition 11.35 and theorem 11.33	Continue to the family classifier [96]
[64]	Survival of positive local CKN concentration after bounded-window regularization	lemma 11.11 and theorem 11.16	No closes at [65]; yes gives [66]
[65]	No surviving local CKN signal	lemma 11.11 and theorem 11.16	Terminal local Type I contradiction
[66]	Removal of the regular portion and local reselection	lemmas 11.11 and 11.15	Test the reselection outcome at [67]
[67]	Well-founded bounded-window reselection outcome	lemma 11.15	Exactly one of [68]–[70]

Continued on next page

Table 3 continued

Node	State or exhaustive test	Source in the proof	Successor or closure
[68]	Local regularity below a fixed physical scale	lemma 11.15	Terminal regular branch
[69]	Finite reselection reaches its unique named row	lemma 11.15 and theorem 11.16	Terminal through that row's recorded closure
[70]	Infinite reselection reaches its unique named row	lemmas 11.14 and 11.15 and theorem 11.16	Terminal through that row's recorded closure
[71]	Bounded-versus-unbounded modulation-window test	lemma 10.106 and theorem 11.29	Unbounded gives [72]; bounded gives [78]
[72]	Unbounded signed or geometric modulation component	lemma 10.106 and theorem 11.29	Exactly one closure [73]–[77]
[73]	Negative scale-drift closure	theorem 10.118	Terminal modulation row
[74]	Positive scale-shell closure	corollary 8.131	Terminal modulation row
[75]	Modulation-driven recentering closure	definition 8.110 and corollary 8.139	Terminal modulation row
[76]	Repaired-gauge first-failure closure	corollaries A.12 and A.19	Terminal representation row
[77]	Exterior modulation closure	corollary 9.73	Terminal exterior row
[78]	Bounded modulation windows	lemma 10.106	Continue to the bounded-defect alternatives [79]
[79]	Bounded modulation-defect outcome	lemma 11.28 and theorem 11.29	Constant gives [80], perturbative gives [81], visible gives [82]
[80]	Constant-coefficient autonomous limit	lemma 11.28 and theorem 7.24	Terminal contradiction to alternative (v)
[81]	Perturbative modulation defect	lemma 11.28 and theorem 11.29	Terminal through the same autonomous-limit contradiction
[82]	Visible nonperturbative modulation component	lemma 11.28	Exactly one of [83]–[87]
[83]	Visible negative scale-drift closure	theorem 10.118	Terminal modulation row
[84]	Visible positive scale-shell closure	corollary 8.131	Terminal modulation row
[85]	Visible recentering closure	definition 8.110 and corollary 8.139	Terminal modulation row
[86]	Visible repaired-gauge closure	corollaries A.12 and A.19	Terminal representation row
[87]	Visible loss of the selected active core	lemma 11.28	Continue to [96]
[88]	Ordered noncanonical cost-validation test	lemma 8.73 and theorem 11.31	Pass gives [89]; first failure gives [91]
[89]	Comparable-cost identification with the canonical identity channel	lemmas 8.67 and 11.30	Apply the retained exclusion at [90]

Continued on next page

Table 3 continued

Node	State or exhaustive test	Source in the proof	Successor or closure
[90]	Canonical retained cost-divergence exclusion	corollary 8.82	Terminal cost row
[91]	First-failure layer in noncanonical validation	lemma 8.73 and theorem 11.31	Exactly one of [92]–[95]
[92]	Class-side first failure	corollaries A.19 and 11.26	Terminal earlier closure
[93]	State-decomposition or compact-test first failure	proposition 11.19 and corollary 8.139	Terminal earlier closure
[94]	Modulation-integrability first failure	theorem 11.29	Terminal by the alternatives [71]–[87]
[95]	Irreducible cost-side first failure	lemma 8.73 and theorem 11.31	Terminal by the canonical-channel rule
[96]	Final active-core ledger family classifier	definition 11.35 , theorem 11.33 , and lemma 10.103	Active loss gives [97], pressure gives [111], nested pressure gives [117], critical tightness gives [128], exterior gives [110]
[97]	Local critical-mass outcome	lemma 10.78	Positive finite gives [98], floor failure gives [99], upper failure gives [103]
[98]	Positive finite local critical mass	lemma 10.78	Terminal contradiction to the loss-row classification
[99]	Removal of an activity-floor-failing candidate	lemma 10.77 and theorem 11.33	Exhaust the remaining finite local frame list at [100]
[100]	Outcome after finite ordered frame reselection	lemmas 10.77 and 10.78 and theorem 11.33	Positive finite gives [98], upper failure gives [103], no velocity frame gives [101]
[101]	Carrier of the original CKN signal when no velocity frame remains	theorem 11.33	Pressure gives [111], exterior gives [110], rough core gives [109], subthreshold regularity gives [102]
[102]	All smaller admissible cylinders below the CKN threshold	theorems 9.35 and 11.33	Terminal contradiction to singular CKN concentration
[103]	First label produced by local critical-mass upper-bound failure	lemma 10.78 and theorem 11.33	Exactly one closure [104]–[110]
[104]	Same-frame critical-mass adjacent face	lemma 10.78	Terminal nonretained label
[105]	Comparable finite-mass class	lemma 10.77	Terminal by prior grouping into the retained class
[106]	Multibubble row	theorem 11.6	Terminal multibubble closure
[107]	Multicore row	theorem 11.5	Terminal multicore closure
[108]	Cascade row	theorems 11.6 and 11.8 and corollary 5.45	Terminal cascade closure
[109]	Rough-core row	theorem 11.7 and proposition 11.19	Terminal rough-core closure
[110]	Pressure or exterior active-profile row	lemma 9.22 and corollary 9.73	Terminal pressure/exterior closure

Continued on next page

Table 3 continued

Node	State or exhaustive test	Source in the proof	Successor or closure
[111]	Availability of the local pressure reconstruction and gauge	lemma 11.32	Available gives [112]; first failure gives [114]
[112]	Persistence of the localized velocity-stress source	lemma 11.32	Persistent gives [113]; vanishing leaves the harmonic row [110]
[113]	Local CKN outcome for a persistent velocity-stress source	lemma 11.32	Regular gives [115]; a velocity carrier gives [116]
[114]	Pressure reconstruction or gauge first failure	theorem 4.25 and corollary A.19	Terminal representation closure
[115]	Regular localized velocity-stress source	lemma 11.32	Terminal regular absorption
[116]	Velocity-supported pressure label	lemmas 10.77 and 11.32	Comparable gives [105]; otherwise one of [106]–[110]
[117]	Well-founded nested pressure-only classifier	lemma 9.67	Velocity gives [116], harmonic tail gives [118], gauge gives [114], same-point gives [119], diffuse gives [121]
[118]	Far-field harmonic pressure-tail row	lemmas 9.22 and 9.68	Terminal pressure-tail closure
[119]	Same-point nested pressure classifier	lemma 9.67	Cascade gives [108]; scale collapse gives [120]
[120]	Pressure scale-collapse row	theorems 7.37 and 11.8	Terminal scale closure
[121]	Atomic/activity versus genuinely diffuse pressure defect	lemmas 9.63 and 9.68	Atomic/activity gives [122]; genuinely diffuse gives [127]
[122]	Closure produced by the realized atomic support profile	theorem 9.66 and lemma 9.68	Regular gives [115], velocity gives [116], non-affine harmonic gives [118], constant gives [123], affine gives [124]
[123]	Constant harmonic pressure	lemma 9.68	Terminal removal by the pressure gauge
[124]	Affine harmonic pressure absorption test	lemmas 9.68 and 9.69	Absorbed gives [125]; failure gives [126]
[125]	Affine pressure absorbed by accelerating rebasement	lemma 9.69	Terminal pressure closure
[126]	Affine-pressure modulation failure	lemma 9.69 and theorem 11.29	Terminal by the alternatives [71]–[87]
[127]	Visible diffuse pressure residual	lemma 9.68 and theorem 10.86	Terminal through its unique named-row closure
[128]	Critical-tightness first-row classifier	lemma 10.103	Cascade gives [108], exterior gives [110], pressure gives [114], compact failure gives [129], diffuse gives [131]
[129]	Compact-package first-failure classifier	lemma 10.103	Rough core gives [109]; local-energy or compact-upper failure gives [130]
[130]	Selected compact-package first failure	theorem 7.34 and corollary 8.139	Terminal compact closure

Continued on next page

Table 3 continued

Node	State or exhaustive test	Source in the proof	Successor or closure
[131]	First lower-face test for a diffuse critical tail	lemmas 10.93 and 10.104	A lower named face gives [132]; otherwise continue to [133]
[132]	First lower named failure face	lemma 10.104	Terminal through its recorded closure
[133]	Atomic/activity versus genuinely diffuse residual tail defect	lemmas 9.63 and 10.104	Atomic/activity gives [134]; genuinely diffuse gives [135]
[134]	Atomic carrier in the diffuse-tail case	theorem 10.94 and lemma 10.104	Terminal contradiction to the diffuse alternative
[135]	Visible diffuse critical-tail residual	theorem 10.86 and lemma 10.104	Terminal through its unique named-row closure

2.4. Monotone retained-obstruction ledger. The arrows need not repeat hypotheses already established upstream. The following ledger records the standing data carried through every subsequence, window reselection, and terminal extraction. A failure of a ledger entry is itself a first-failure row and therefore follows the unique route to [56].

Table 4: Monotone invariants carried by the Type II proof.

Invariant	Established by	Retained through	Failure or terminal use
Physical Type II status	definition 11.20 and lemma 11.22	Covered entry and ordered C1 tests [1]–[7]	Class-incompatible failure closes at [57]
Positive local critical signal	lemma 11.23 and theorem 9.33	Core selection, scale analysis, terminal profiles, and the final ledger [4]–[135]	The zero limit closes at [27]; loss of the selected signal is classified at [96]–[116]
Nonzero compact-limit dichotomy	lemmas 7.29 and 7.30	Selected compactness, terminal descendants, and diffuse-tail routing [40]–[52], [128]–[135]	A first interface failure is named; a singular limit gives an active descendant; a regular nonzero limit enters the bounded-window row
Buffered locality and suitable energy	propositions 4.5 and 4.33	Every compact-cylinder restriction after [10]	The first domain, energy, or compact-upper failure goes to [56]
Pressure representative	theorem 4.25	Carrier selection, compactness, Caccioppoli, and residual exhaustion	Pressure-only rows are expanded at [111]–[127], and reconstruction failure closes at [114]
Retained ancestry	lemma 10.77 , corollary 10.79 , and theorem 9.16	Core selection, terminal descendants, and residual coordinates	Noncomparable descendants close at [49], [59], or [106]–[110]
Canonical chart and modulation convention	definition 1.1 and theorem 10.52	Scale cost, carrier extraction, and terminal sampling	The first chart, modulation, or cost failure closes through [56]
First-failure order and rank	theorem 11.17 and lemma 11.34	Every routing test [5]–[135]	Exactly one named row occurs, and every displayed ellipse is terminal

2.5. Branch-closure audit ledger. The retained alternatives of [theorem 11.18](#) and the first-failure alternatives of [definition 11.35](#) are exhausted by the following rows. The closure column cites the result in the body of the paper; the final column identifies the corresponding terminal ellipse. There is no common assembly node after these ellipses.

Table 5: Closure audit for every terminal class in the Type II diagram.

State or exit	Diagram route	Closure source	Terminal output
Regular or local Type I entry	[2]–[3]	theorem 11.3	[3]
Multibubble, cascade, or gauge degeneracy	[11]–[12]	theorem 11.6	[12]
Loss of local windowed H^1 control	[20]–[21]	theorems 10.61 and 11.7 and proposition 11.19	[21]
Compact single core with vanishing dissipation	[24]–[27]	theorem 11.5	[27]
Compact single core with divergent retained cost	[30]–[34]	theorems 8.133 and 10.42	[34]
Finite signed or absolute scale cost	[35]–[36]	theorem 11.8	[36]
Thin drift reaching an already closed retained state	[35], [37]–[38]	theorems 7.37 and 10.118	[38]
Structural or exterior terminal profile	[48]–[49]	theorem 11.6 and corollary 9.73	[49]
Retained comparable terminal profile	[52]–[55]	theorems 10.144 and 10.146	[55]
Class-incompatible first failure	[6], [56]–[57]	lemmas 11.22 and 11.24	[57]
Earlier local or ambient first failure	[56], [58]–[59]	proposition 11.19 and corollary 8.139	[59]
Bounded selected-window limit	[28], [56], [60], [64]–[70]	lemmas 5.25 and 11.15 and theorem 11.16	[65], [68]–[70]
Autonomous-modulation failure	[39], [56], [61], [71]–[87]	lemmas 10.106 and 11.28 and theorem 11.29	[73]–[77], [80]–[81], [83]–[86]; [87] continues to [96]
Noncanonical carrier or cost-validation failure	[32], [56], [62], [88]–[95]	lemma 8.73 and theorem 11.31	[90], [92]–[95]

State or exit	Diagram route	Closure source	Terminal output
Loss of the retained active core	[15], [45], [51], [56], [63], [96]–[116]	lemmas 10.77 and 10.78 and theorem 11.33	[98], [102], [104]–[110], [114]–[115]
Nested pressure-only behavior	[117]–[127]	lemmas 9.67 and 9.68	[105]–[110], [114]–[115], [118], [120], [123], [125]–[127]
Critical-tightness and diffuse-tail failure	[128]–[135]	lemmas 7.30, 10.98, 10.103 and 10.104	[108]–[110], [114], [130], [132], [134]–[135]

Every path from [1] therefore terminates at one of the earlier leaves [3], [12], [21], [27], [34], [36], [38], [49], [55], or at one of [57], [59], [65], [68]–[70], [73]–[77], [80]–[81], [83]–[86], [90], [92]–[95], [98], [102], [104]–[110], [114]–[115], [118], [120], [123], [125]–[127], [130], [132], or [134]–[135]. This leaf exhaustion is the diagrammatic audit of [theorems 11.38](#) and [11.39](#). The unconditional conclusion is not used to justify any earlier node.

3. LOCAL SINGLE-CORE CRITERION

3.1. Setup and definitions. For $R > 0$, write

$$Q_R = B_R \times (-R^2, 0).$$

For a suitable pair (v, q) on Q_R , define

$$\begin{aligned} A_v(R) &= R^{-1} \operatorname{ess\,sup}_{-R^2 < s < 0} \int_{B_R} |v(y, s)|^2 dy, \\ E_v(R) &= R^{-1} \int_{Q_R} |\nabla v|^2 dy ds, \\ C_v(R) &= R^{-2} \int_{Q_R} |v|^3 dy ds, \\ D_q(R) &= R^{-2} \int_{-R^2}^0 \int_{B_R} |q(y, s) - (q)_{B_R}(s)|^{3/2} dy ds. \end{aligned}$$

All pressure quantities are understood modulo functions of time.

Definition 3.1 (Bounded selected-window limit). Let (V_n, P_n) be a sequence of represented local rescalings on a compact cylinder Q_r . A selected-window subsequence has a bounded selected-window limit on Q_r if, after passing to a subsequence,

$$V_n \rightarrow V_* \quad \text{strongly in } L^3(Q_r),$$

where $V_* \in L^\infty(Q_r)$. The limit is nonzero if $V_* \not\equiv 0$ on Q_r .

Definition 3.2 (Strict selected-window Type II branch). A represented selected-window branch is called a strict selected-window Type II branch if it has positive local Caffarelli–Kohn–Nirenberg concentration on some fixed compact cylinder and no selected-window subsequence has a nonzero bounded selected-window limit on any compact cylinder in the sense of [definition 3.1](#).

Remark 3.3 (Strict selected-window Type II versus physical Type II). [Definition 3.2](#) is a local branch condition inside the selected-window state-space decomposition. It is not the definition of a physical Type II singular germ; physical local Type II germs are defined later in [definition 11.20](#) by failure of the local Type I scale-invariant bound at a singular point. A bounded selected-window limit is therefore not discarded as “not physically Type II”; it leaves the strict selected-window branch and is routed through the bounded-window local closure row by [theorems 5.27](#) and [11.16](#). Only after that row is closed, or reselected into one of the named local alternatives, may the proof continue in the strict selected-window branch.

Definition 3.4 (Compact single-core vanishing-dissipation branch). Fix radii $0 < r_* < R_*$, constants $M < \infty$, $\eta_0 > 0$, and a sequence of represented local rescalings (V_n, P_n, a_n, b_n) on Q_{R_*} . The sequence is a compact single-core vanishing-dissipation branch if:

- (i) each (V_n, P_n) is suitable on Q_{R_*} and solves the canonical represented Navier–Stokes equation [\(14\)](#), with (τ, Y, V, P, a, b) renamed as $(s, y, V_n, P_n, a_n, b_n)$;
- (ii) the compact-cylinder suitable estimates are uniformly bounded:

$$A_{V_n}(R_*) + E_{V_n}(R_*) + C_{V_n}(R_*) + D_{P_n}(R_*) \leq M;$$

- (iii) the selected core has retained positive velocity concentration on Q_{r_*} :

$$C_{V_n}(r_*) \geq \eta_0 \quad \text{for all } n;$$

Pressure-only concentration is not included in this compact single-core row; it is routed to the harmonic/exterior-pressure alternative.

- (iv) the localized scale and center are nondegenerate for the selected core, so the representation theorem gives the variables in (14);
- (v) the modulation coefficients are uniformly bounded on the selected cylinder:

$$\|a_n\|_{L^\infty(-R_*^2, 0)} + \|b_n\|_{L^\infty(-R_*^2, 0)} \leq M;$$

- (vi) the branch has vanishing localized dissipation in the sense of definition 3.5 below.

Definition 3.5 (Vanishing localized dissipation). Let $r_* < \rho < R_*$. Selected windows are first translated and rescaled to the fixed cylinders $Q_{r_*} \Subset Q_\rho \Subset Q_{R_*}$. The localized dissipation on Q_ρ is

$$\mathcal{D}_n(\rho) = \int_{Q_\rho} |\nabla V_n|^2 dy ds.$$

The branch has vanishing localized dissipation if, after passing to a selected-window subsequence, there is $\rho \in (r_*, R_*)$ such that

$$\mathcal{D}_n(\rho) \rightarrow 0.$$

3.2. Local analytic inputs. The compact single-core argument is based on the following local analytic statements.

Theorem 3.6 (Local concentration dichotomy). *Let (u, p) be a suitable weak solution on $\mathbb{R}^3 \times (T - \delta, T)$, and suppose $\Sigma(T) \neq \emptyset$. Then, for every $x_0 \in \Sigma(T)$, using the pressure-normalized quantities C, D , there are $\eta > 0$ and $r_n \downarrow 0$ such that*

$$C((x_0, T), r_n) + D((x_0, T), r_n) \geq \eta \quad \text{for all } n.$$

Moreover, at each singular point exactly one of the following alternatives holds: a local Type I scale-invariant bound holds, or the point lies in the local Type II alternative, meaning that no such local Type I bound holds.

Proof. By the Caffarelli–Kohn–Nirenberg epsilon-regularity criterion, there is a universal $\varepsilon_{\text{CKN}} > 0$ such that if

$$C((x_0, T), r) + D((x_0, T), r) < \varepsilon_{\text{CKN}}$$

for all sufficiently small r , then u is locally bounded in a smaller backward cylinder ending at (x_0, T) . Since $x_0 \in \Sigma(T)$, this regularity conclusion is impossible. Hence for every k there exists $0 < r_k < 1/k$ with

$$C((x_0, T), r_k) + D((x_0, T), r_k) \geq \varepsilon_{\text{CKN}}.$$

After passing to a decreasing subsequence and setting $\eta = \varepsilon_{\text{CKN}}$, the displayed concentration sequence follows. The Type I/Type II split is the dichotomy between existence and failure of the local Type I scale-invariant bound at the same point. $\square$

Proposition 3.7 (Passage to selected-window branches). *Let (x_0, T) be a concentration point in the physical local Type II alternative of theorem 3.6. After choosing the positive concentration scales given there and applying the local scale-translation selection used in the present section to a nondegenerate compact single core, the represented selected windows have positive local CKN concentration on a fixed compact cylinder. Exactly one of the following selected-window alternatives is then recorded:*

- (i) *the selected branch is a strict selected-window Type II branch in the sense of definition 3.2;*
- (ii) *a selected-window subsequence has a nonzero bounded selected-window limit in the sense of definition 3.1. This is the bounded-window row and is routed by theorem 11.16.*

Proof. The local concentration dichotomy supplies the positive CKN concentration sequence and the failure of the local Type I scale-invariant bound at the chosen physical point. The scale-translation selection rewrites those shrinking cylinders as fixed selected windows. On the selected windows there are then two mutually exclusive possibilities: either no selected-window subsequence has a nonzero bounded selected-window limit, which is precisely [definition 3.2](#), or such a subsequence exists. In the latter case the branch is not discarded as physically impossible; it is sent to the bounded-window physical routing theorem. $\square$

Proposition 3.8 (Local representation input). *For a nondegenerate selected single core, the local scale-translation construction gives represented variables (V_n, P_n, a_n, b_n) satisfying (14) on compact cylinders. Failure of the localized nondegeneracy condition is a multibubble, cascade, or gauge-degenerate alternative, not a compact single-core branch.*

Proof. On the nondegenerate selected-core branch the scale-center construction provides absolutely continuous scale and center functions $\lambda_n(t) > 0$, $x_n(t)$, with λ_n bounded above and below on each selected compact window after normalization. Define

$$y = \frac{x - x_n(t)}{\lambda_n(t)}, \quad \frac{ds}{dt} = \lambda_n(t)^{-2},$$

and set

$$u(x, t) = \lambda_n(t)^{-1} V_n(y, s), \quad p(x, t) = \lambda_n(t)^{-2} P_n(y, s).$$

The distributional chain rule is justified by the absolute continuity of the chart and the local suitable-energy bounds on the physical solution. Testing the physical Navier–Stokes equations against pullbacks of compactly supported test fields gives, on every compact selected cylinder,

$$\partial_s V_n + (V_n \cdot \nabla) V_n + \nabla P_n = \nu \Delta V_n + a_n(s)(V_n + y \cdot \nabla V_n) + b_n(s) \cdot \nabla V_n, \quad \nabla \cdot V_n = 0,$$

where

$$a_n(s) = \lambda_n(t) \dot{\lambda}_n(t), \quad b_n(s) = \lambda_n(t) \dot{x}_n(t).$$

Suitability and the local CKN quantities are invariant under this parabolic change of variables up to the fixed selected-window constants.

If the localized nondegeneracy needed to choose a single compact scale-center chart fails, then the ordered selection has detected either separated active cores, a same-point scale cascade, or failure of the gauge chart. These are precisely the multibubble, cascade, and gauge-degenerate rows, so the failure cannot be part of the compact single-core branch. $\square$

Lemma 3.9 (Localization of pressure oscillation bounds). *Let $0 < r < R$, and let $q \in L^{3/2}(Q_R)$. Then*

$$D_q(r) \leq C \left(\frac{R}{r} \right)^2 D_q(R),$$

where C is universal.

Proof. For a.e. $s \in (-r^2, 0)$, Jensen's inequality gives

$$|(q)_{B_r}(s) - (q)_{B_R}(s)|^{3/2} \leq \frac{1}{|B_r|} \int_{B_r} |q(y, s) - (q)_{B_R}(s)|^{3/2} dy.$$

Hence

$$\int_{B_r} |q - (q)_{B_r}(s)|^{3/2} dy \leq C \int_{B_r} |q - (q)_{B_R}(s)|^{3/2} dy.$$

Integrating over $(-r^2, 0)$ and using $Q_r \subset Q_R$,

$$\begin{aligned} D_q(r) &\leq Cr^{-2} \int_{-r^2}^0 \int_{B_r} |q - (q)_{B_R}(s)|^{3/2} dy ds \\ &\leq Cr^{-2} \int_{-R^2}^0 \int_{B_R} |q - (q)_{B_R}(s)|^{3/2} dy ds = C \left(\frac{R}{r} \right)^2 D_q(R). \end{aligned}$$

□

3.3. Compactness and localized pressure compactness.

Proposition 3.10 (Compactness on selected windows). *Let (V_n, P_n) be suitable on Q_{R_*} , solve (14), and satisfy*

$$A_{V_n}(R_*) + E_{V_n}(R_*) + C_{V_n}(R_*) + D_{P_n}(R_*) \leq M,$$

with

$$\|a_n\|_{L^\infty(-R_*^2, 0)} + \|b_n\|_{L^\infty(-R_*^2, 0)} \leq M.$$

Then for every $r < R_*$, after passing to a subsequence there are (V, P) such that

$$\begin{aligned} V_n &\overset{*}{\rightharpoonup} V \quad \text{in } L^\infty(-r^2, 0; L^2(B_r)), \\ V_n &\rightharpoonup V \quad \text{in } L^2(-r^2, 0; H^1(B_r)), \\ V_n &\rightarrow V \quad \text{strongly in } L^3(Q_r), \end{aligned}$$

and

$$P_n - (P_n)_{B_r} \rightharpoonup P \quad \text{in } L^{3/2}(Q_r).$$

After passing to a further subsequence, a_n and b_n converge weak-* in L^∞ to bounded coefficients, and the limit satisfies the corresponding represented equation distributionally on smaller cylinders.

Proof. The bounds on A and E give weak-* compactness in $L^\infty L^2$ and weak compactness in $L^2 H^1$ on smaller cylinders. Fix $r < r_1 < R_*$, and choose m so large that

$$H_0^m(B_{r_1}) \hookrightarrow W_0^{1,\infty}(B_{r_1}).$$

We claim that $\partial_s V_n$ is uniformly bounded in

$$L^1(-r_1^2, 0; H^{-m}(B_{r_1})).$$

Let $\varphi \in H_0^m(B_{r_1})$. Pairing (14) with φ , integrating by parts in space where needed, and using the embedding above gives

$$\begin{aligned} |\langle (V_n \cdot \nabla) V_n, \varphi \rangle| &= \left| \int_{B_{r_1}} V_{n,i} V_{n,j} \partial_j \varphi_i dy \right| \leq C \|V_n\|_{L^3(B_{r_1})}^2 \|\varphi\|_{H^m}, \\ |\langle \nu \Delta V_n, \varphi \rangle| &\leq C \|\nabla V_n\|_{L^2(B_{r_1})} \|\varphi\|_{H^m}, \\ |\langle \nabla P_n, \varphi \rangle| &= \left| \int_{B_{r_1}} (P_n - (P_n)_{B_{r_1}}) \nabla \cdot \varphi dy \right| \leq C \|P_n - (P_n)_{B_{r_1}}\|_{L^{3/2}(B_{r_1})} \|\varphi\|_{H^m}. \end{aligned}$$

The modulation terms obey

$$|\langle a_n V_n, \varphi \rangle| \leq C \|V_n\|_{L^2(B_{r_1})} \|\varphi\|_{H^m},$$

$$|\langle a_n y \cdot \nabla V_n, \varphi \rangle| = \left| \int_{B_{r_1}} a_n V_{n,i} \partial_j (y_j \varphi_i) dy \right| \leq C \|V_n\|_{L^2(B_{r_1})} \|\varphi\|_{H^m},$$

and

$$|\langle b_n \cdot \nabla V_n, \varphi \rangle| = \left| \int_{B_{r_1}} V_{n,i} b_{n,j} \partial_j \varphi_i dy \right| \leq C \|V_n\|_{L^2(B_{r_1})} \|\varphi\|_{H^m},$$

where the constant uses r_1 and the uniform L^∞ bounds for a_n, b_n . Integrating these estimates in time gives a uniform bound because

$$\begin{aligned} \int_{-r_1^2}^0 \|V_n(s)\|_{L^3(B_{r_1})}^2 ds &\leq r_1^{2/3} \|V_n\|_{L^3(Q_{r_1})}^2, \\ \int_{-r_1^2}^0 \|\nabla V_n(s)\|_{L^2(B_{r_1})} ds &\leq r_1 \|\nabla V_n\|_{L^2(Q_{r_1})}, \\ \int_{-r_1^2}^0 \|P_n(s) - (P_n)_{B_{r_1}}(s)\|_{L^{3/2}(B_{r_1})} ds &\leq r_1^{2/3} \|P_n - (P_n)_{B_{r_1}}\|_{L^{3/2}(Q_{r_1})}, \end{aligned}$$

and

$$\int_{-r_1^2}^0 \|V_n(s)\|_{L^2(B_{r_1})} ds \leq r_1^2 \|V_n\|_{L^\infty(-r_1^2, 0; L^2(B_{r_1}))}.$$

The uniform C , E , D , and A bounds therefore prove the claimed time-derivative bound; the D -bound on Q_{r_1} follows from the bound on Q_{R_*} by [lemma 3.9](#).

The compact embedding $H^1(B_{r_1}) \Subset L^2(B_r)$, together with Aubin–Lions compactness [[Aub63](#), [LM72](#), [Sim87](#)], gives

$$V_n \rightarrow V \quad \text{strongly in } L^2(Q_r)$$

after passing to a subsequence. The same $L^\infty L^2 \cap L^2 H^1$ bounds give the parabolic interpolation estimate

$$\|V_n\|_{L^{10/3}(Q_r)} \leq C(r, M).$$

Indeed, for a.e. s , the Sobolev and Gagliardo–Nirenberg inequalities give

$$\|V_n(s)\|_{L^{10/3}(B_r)} \leq C \|V_n(s)\|_{L^2(B_r)}^{2/5} \|\nabla V_n(s)\|_{L^2(B_r)}^{3/5}.$$

Raising to the power $10/3$, integrating in time, and using the A and E bounds yields

$$\int_{-r^2}^0 \|V_n(s)\|_{L^{10/3}(B_r)}^{10/3} ds \leq C \left(\operatorname{ess\,sup}_{-r^2 < s < 0} \|V_n(s)\|_{L^2(B_r)}^2 \right)^{2/3} \int_{Q_r} |\nabla V_n|^2 dy ds \leq C(r, M).$$

The same estimate applies to V by weak lower semicontinuity, and therefore to $V_n - V$ by the triangle inequality. Interpolating between L^2 and $L^{10/3}$, one obtains

$$\|V_n - V\|_{L^3(Q_r)} \leq \|V_n - V\|_{L^2(Q_r)}^{1/6} \|V_n - V\|_{L^{10/3}(Q_r)}^{5/6}.$$

The second factor is uniformly bounded and the first tends to zero, so

$$V_n \rightarrow V \quad \text{strongly in } L^3(Q_r).$$

The pressure compactness follows from the uniform D -bound: after subtracting spatial means,

$$P_n - (P_n)_{B_r} \rightharpoonup P \quad \text{in } L^{3/2}(Q_r).$$

We pass to the limit distributionally using the preceding convergences. It remains to justify the modulation terms. For every compactly supported test field φ ,

$$\int a_n(s) y \cdot \nabla V_n \cdot \varphi = - \int a_n(s) V_{n,i} \partial_j (y_j \varphi_i),$$

where summation over repeated spatial indices is understood. Equivalently,

$$\partial_j (y_j \varphi_i) = 3\varphi_i + y_j \partial_j \varphi_i.$$

Hence this term passes to the limit by strong local L^2 convergence of V_n and weak-* convergence of a_n : if $F_n \rightarrow F$ in L^1 and $a_n \rightharpoonup^* a$ in L^∞ , then

$$\int a_n F_n - \int a F = \int a_n (F_n - F) + \int (a_n - a) F \rightarrow 0.$$

The term $a_n V_n$ is immediate from the same argument, while

$$\int b_n(s) \cdot \nabla V_n \cdot \varphi = - \int V_{n,i} b_{n,j}(s) \partial_j \varphi_i$$

passes to the limit by strong local L^2 convergence of V_n and weak-* convergence of b_n . These passages give the limiting represented equation in the sense of distributions. $\square$

Remark 3.11 (Use of bounded modulation). The uniform L^∞ bounds on a_n and b_n enter the compactness argument in exactly two places. First, they give the negative Sobolev time-derivative bound for the three lower-order modulation terms in (14). Second, after subsequential weak-* convergence of a_n, b_n , they allow the modulation terms to pass to the distributional limit once $V_n \rightarrow V$ strongly in L^2_{loc} .

Lemma 3.12 (Stability of the localized pressure part). *Let $V_n \rightarrow V$ strongly in $L^3(Q_R)$, and use the pressure decomposition of [theorem 4.25](#). Then for every $r < R$,*

$$P_{\text{loc},n} \rightarrow P_{\text{loc}} \quad \text{strongly in } L^{3/2}(Q_r),$$

where $-\Delta P_{\text{loc}} = \partial_i \partial_j (\zeta V_i V_j)$. If $V \equiv 0$ on Q_R , then

$$P_{\text{loc},n} \rightarrow 0 \quad \text{strongly in } L^{3/2}(Q_r).$$

This lemma does not assert convergence of the full mean-normalized pressure: the harmonic component may contain local parasitic pressure modes independent of V_n .

Proof. Strong convergence in L^3 gives

$$V_{n,i} V_{n,j} \rightarrow V_i V_j \quad \text{strongly in } L^{3/2}(Q_R).$$

Indeed,

$$\|V_{n,i} V_{n,j} - V_i V_j\|_{L^{3/2}(Q_R)} \leq (\|V_n\|_{L^3(Q_R)} + \|V\|_{L^3(Q_R)}) \|V_n - V\|_{L^3(Q_R)}.$$

Calderón–Zygmund estimates for the compactly supported localized source yield strong convergence of the localized pressures in $L^{3/2}$ on smaller balls:

$$\|P_{\text{loc},n} - P_{\text{loc}}\|_{L^{3/2}(Q_r)} \leq C_{r,R} \|V_{n,i} V_{n,j} - V_i V_j\|_{L^{3/2}(Q_R)}.$$

If $V \equiv 0$, then $P_{\text{loc}} = 0$, and the first part gives the displayed convergence of $P_{\text{loc},n}$. The harmonic part H_n is not controlled by Calderón–Zygmund from the velocity source and may persist even when $V_n \rightarrow 0$; it is therefore handled only in the pressure/exterior rows, not in the compact single-core zero-velocity contradiction. $\square$

Lemma 3.13 (Zero-dissipation limit). *If*

$$V_n \rightharpoonup V \quad \text{in } L^2(I; H^1(B_\rho))$$

and

$$\int_I \int_{B_\rho} |\nabla V_n|^2 dy ds \rightarrow 0,$$

then $\nabla V = 0$ a.e. on $B_\rho \times I$. Hence for a.e. time, $V(\cdot, s)$ is spatially constant on B_ρ .

Proof. Weak lower semicontinuity gives

$$\int_I \int_{B_\rho} |\nabla V|^2 dy ds \leq \liminf_{n \rightarrow \infty} \int_I \int_{B_\rho} |\nabla V_n|^2 dy ds = 0.$$

Thus $\nabla V = 0$ a.e. on $B_\rho \times I$. By Fubini, for a.e. $s \in I$ one has $\nabla V(\cdot, s) = 0$ in $L^2(B_\rho)$. Since B_ρ is connected, every $H^1(B_\rho)$ function with zero weak gradient is a.e. equal to a constant. Hence $V(\cdot, s)$ is spatially constant for a.e. s . $\square$

Lemma 3.14 (Spatially constant limits). *Let V be a limit on Q_ρ and suppose $\nabla V = 0$ a.e. there. Then $V(y, s) = c(s)$ on Q_ρ , with $c \in L^\infty(-\rho^2, 0)$. If $c \not\equiv 0$ on Q_ρ , then the approximating selected-window subsequence has a nonzero bounded selected-window limit.*

Proof. The spatial constancy follows from lemma 3.13. Since $V(\cdot, s) = c(s)$ on B_ρ , the local $L_s^\infty L_y^2$ bound gives

$$|c(s)|^2 |B_\rho| \leq \operatorname{ess\,sup}_{-\rho^2 < s < 0} \int_{B_\rho} |V(y, s)|^2 dy,$$

so $c \in L^\infty(-\rho^2, 0)$. If $c \not\equiv 0$ on Q_ρ , then the strong L^3 convergence of proposition 3.10 gives a subsequential limit to a nonzero bounded function on Q_ρ , which is precisely a nonzero bounded selected-window limit in the sense of definition 3.1. $\square$

3.4. Contradiction lemmas.

Lemma 3.15 (Persistence of retained velocity core). *For a compact single-core branch in the sense of definition 3.4, there are fixed $r_* > 0$ and $\eta_0 > 0$ such that along the selected-window subsequence,*

$$C_{V_n}(r_*) \geq \eta_0 \quad \text{for all } n.$$

Proof. The retained velocity lower bound is part of the definition of the selected compact core and is preserved by passing to subsequences of the selected windows. $\square$

Lemma 3.16 (Zero constant contradicts retained velocity). *Let $r_* < \rho$. Suppose the compactness limit in proposition 3.10 is spatially constant on Q_ρ and is identically zero on Q_{r_*} . Then*

$$C_{V_n}(r_*) \rightarrow 0.$$

Consequently this case contradicts lemma 3.15.

Proof. Strong $L^3(Q_{r_*})$ convergence to zero gives

$$C_{V_n}(r_*) = r_*^{-2} \int_{Q_{r_*}} |V_n|^3 dy ds = r_*^{-2} \|V_n\|_{L^3(Q_{r_*})}^3 \rightarrow 0.$$

This contradicts the retained velocity lower bound

$$C_{V_n}(r_*) \geq \eta_0.$$

No assertion is made about $D_{P_n}(r_*)$: a pressure-only harmonic mode may remain and is not part of the retained compact single-core velocity row. $\square$

Lemma 3.17 (Nonzero constant exits the strict selected-window branch). *If the compactness limit is spatially constant and nonzero on a compact subcylinder, then the original sequence leaves the strict selected-window branch.*

Proof. By lemma 3.14, the selected-window subsequence has a nonzero bounded selected-window limit. By definition 3.2, such a sequence is assigned to the bounded-window alternative rather than to the strict selected-window branch. Physical local Type II germs are then routed by theorem 11.16. $\square$

3.5. Main theorem.

Theorem 3.18 (Local single-core Type II criterion). *Under the analytic inputs [propositions 3.7](#) and [3.8](#) and [theorem 4.25](#), a compact single-core vanishing-dissipation branch in the sense of [definition 3.4](#) cannot be a local Type II sequence.*

Proof. By the definition of vanishing localized dissipation, after passing to a selected-window subsequence there is a fixed radius $\rho \in (r_*, R_*)$ such that

$$\int_{Q_\rho} |\nabla V_n|^2 dy ds \rightarrow 0$$

and the positive core bound on Q_{r_*} remains valid along this subsequence. Apply compactness on the fixed compact cylinder Q_ρ and pass to a subsequence. The compactness proposition gives

$$V_n \rightharpoonup V \quad \text{in } L^2(-\rho^2, 0; H^1(B_\rho))$$

and strong convergence in L^3 on every smaller cylinder. Since the localized dissipation tends to zero, the zero-dissipation lemma gives

$$\nabla V = 0 \quad \text{a.e. on } Q_\rho.$$

Thus $V(y, s) = c(s)$ on Q_ρ , and in particular on the selected core cylinder Q_{r_*} .

If this spatially constant limit is zero on the core cylinder, the zero-limit contradiction lemma gives

$$C_{V_n}(r_*) \rightarrow 0,$$

contradicting persistence of the retained velocity core.

If the spatially constant limit is not identically zero on Q_{r_*} , the nonzero-limit lemma produces a nonzero bounded selected-window limit, so the strict selected-window branch stops at the bounded-window alternative. In the physical assembly that bounded-window event is routed by [theorem 11.16](#). The two cases exhaust all spatially constant limits, so the branch cannot occur. $\square$

Remark 3.19 (Minimal local inputs). The proof uses two local analytic inputs from the Type II argument: the localized representation for a nondegenerate selected core, the compact-ball pressure reconstruction with stability on smaller balls. The compact-cylinder A, E, C, D bounds and vanishing localized dissipation are part of the branch hypothesis, not separate analytic inputs. Multibubble, cascade, gauge-degenerate, and compactness-failure alternatives are outside this single-core theorem; they are treated by the local multibubble and cascade exclusion theorem.

4. REPAIRED-GAUGE CONSTRUCTION

4.1. Setup and notation. Let $u : \mathbb{R}^3 \times (T - \delta, T) \rightarrow \mathbb{R}^3$, $p : \mathbb{R}^3 \times (T - \delta, T) \rightarrow \mathbb{R}$ be a suitable weak Navier–Stokes solution with viscosity $\nu > 0$. For $z = (x_0, t_0)$ and $R > 0$,

$$Q_R(z) := B_R(x_0) \times (t_0 - R^2, t_0), \quad Q_R^* := B_R \times I_*.$$

Fix a reference physical center-time $(x_\bullet, t_\star)$ around which $B_{R_*}(x_\bullet)$ is taken. For a fixed spatial center x_0 we also use

$$Q_R(t) := Q_R(x_0, t).$$

Fix a compact window $I_* := [\tau_0 - \ell, \tau_0 + \ell]$, $\ell > 0$, and write $I_{\text{ph}} := [t_\star - \Lambda_{\max}^2 \ell, t_\star + \Lambda_{\max}^2 \ell]$.

For a cylinder $Q_R(x_0, t_0) = B_R(x_0) \times (t_0 - R^2, t_0)$, set

$$\begin{aligned} A(u; Q_R(x_0, t_0)) &:= R^{-1} \sup_{s \in (t_0 - R^2, t_0)} \int_{B_R(x_0)} |u(x, s)|^2 dx, \\ E(u; Q_R(x_0, t_0)) &:= R^{-1} \int_{Q_R(x_0, t_0)} |\nabla u|^2 dx ds, \\ C(u; Q_R(x_0, t_0)) &:= R^{-2} \int_{Q_R(x_0, t_0)} |u|^3 dx ds, \\ D(p; Q_R(x_0, t_0)) &:= R^{-2} \int_{t_0 - R^2}^{t_0} \int_{B_R(x_0)} |p(x, s) - \langle p(s) \rangle_{B_R(x_0)}|^{3/2} dx ds. \end{aligned}$$

Here $\langle p(s) \rangle_{B_R(x_0)} := |B_R|^{-1} \int_{B_R(x_0)} p(x, s) dx$; pressure is always understood modulo time-dependent constants.

Definition 4.1 (Single-core hypotheses). Let $(R_*, r_*, M, \eta, \kappa, \Lambda_{\min}, \Lambda_{\max})$ satisfy $\Lambda_{\min} \leq \Lambda_{\max} < \infty$. We say (u, p) satisfies SC($R_*, r_*, M, \eta, \kappa$) on $Q_{R_*}(x_\bullet, t_\star)$ if the following hold:

- (SC1) $A(u; Q_{R_*}(x_\bullet, t)) + E(u; Q_{R_*}(x_\bullet, t)) + C(u; Q_{R_*}(x_\bullet, t)) + D(p; Q_{R_*}(x_\bullet, t)) \leq M$ for every $t \in [t_\star - R_*^2, t_\star]$.
- (SC2) There exists a physically selected core map $(x_c(t), \rho_c(t))$ on I_{ph} such that $\rho_c(t) \in (0, R_*]$, $x_c(t) \in B_{R_*/2}(x_\bullet)$, and for some $\Theta_0 > 0$,

$$\int_{B_{r_*}(0)} |\rho_c(t)u(x_c(t) + \rho_c(t)y, t)|^3 dy \geq \eta \quad \text{for a.e. } t \in I_{\text{ph}}.$$

- (SC3) There exists $0 < r_{\text{sep}} \leq r_*/2$ such that for every $0 < r \leq r_{\text{sep}}$, every competitor pair

$$\mathcal{C}_{\text{core}}(t) := \left\{ (\bar{x}, \bar{\rho}) : |\bar{x} - x_\bullet| \leq \frac{R_*}{2}, 0 < \bar{\rho} \leq R_* \right\},$$

$(\bar{x}, \bar{\rho}) \in \mathcal{C}_{\text{core}}(t) \setminus \{(x_c(t), \rho_c(t))\}$ with

$$\rho_c(t)^{-1} |\bar{x} - x_c(t)| + \left| \frac{\bar{\rho}}{\rho_c(t)} - 1 \right| \geq r_{\text{sep}}$$

has separated local weighted mass in the same core-scale variables:

$$\left| \int_{B_r} |\rho_c(t)u(x_c(t) + \rho_c(t)y, t)|^3 dy - \int_{B_r} |\bar{\rho}u(\bar{x} + \bar{\rho}y, t)|^3 dy \right| \geq \eta.$$

This quantitative non-competing core condition is imposed on the selected physical core directly; it is not a consequence of the repaired-gauge construction.

(SC4) For the selected core,

$$\sup_{t \in I_{\text{ph}}} \int_{B_{2R_*}} |y|^{-p} |\rho_c(t)u(x_c(t) + \rho_c(t)y, t)|^3 dy \leq M,$$

for some $0 < p < 3$.

(SC5) The preliminary frame associated with the selected core is the absolutely continuous representative

$$0 < \Lambda_{\min} \leq \Lambda_{\text{pre}}(\tau) \leq \Lambda_{\max}, \quad \Lambda_{\text{pre}} \in AC(I_*), \quad X_{\text{pre}} \in AC(I_*; \mathbb{R}^3).$$

It is tied to the selected physical core by the time change $dt_c/d\tau = \Lambda_{\text{pre}}(\tau)^2$, $t_c(\tau_0) = t_*$: one has

$$X_{\text{pre}}(\tau) = x_c(t_c(\tau)), \quad \Lambda_{\text{pre}}(\tau) = \rho_c(t_c(\tau))$$

for a.e. $\tau \in I_*$.

The parameter κ denotes the quantitative constants obtained in the local gauge estimates and is not used in items (SC1)–(SC5).

Remark 4.2. The selected core (x_c, ρ_c) and its separation/nondegeneracy conditions are specified in physical variables; no repaired-gauge constraint is used in their definition. This avoids the circularity of defining the core via the chart to be constructed.

4.2. Preliminary frame map from the selected core. The transport map is constructed from the selected core; the core selection itself is part of the hypotheses.

Define

$$\frac{dt_c}{d\tau} = \Lambda_{\text{pre}}(\tau)^2, \quad t_c(\tau_0) = t_*,$$

and use the compatibility in (SC5) to identify $\Lambda_{\text{pre}}(\tau) = \rho_c(t_c(\tau))$ and $X_{\text{pre}}(\tau) = x_c(t_c(\tau))$ a.e.

Proposition 4.3 (Preliminary frame map regularity). *Under SC($R_*, r_*, M, \eta, \kappa$), the map*

$$\Phi_{\text{pre}}(\tau, Y) := (x, t) = (X_{\text{pre}}(\tau) + \Lambda_{\text{pre}}(\tau)Y, t_c(\tau))$$

is well-defined and absolutely continuous in τ . Moreover, $1/\Lambda_{\text{pre}}, \Lambda_{\text{pre}} \in L^\infty(I_)$, $X_{\text{pre}} \in W_{\text{loc}}^{1,1}(I_*; \mathbb{R}^3)$, and for every $R \leq R_*$,*

$$\Phi_{\text{pre}}(Q_R^*) \subset B_{R_*/2 + \Lambda_{\max}R}(x_\bullet) \times I_{\text{ph}}.$$

Consequently, if $B_{R_{\text{phys}}}(x_\bullet) \times I_{\text{ph}}$ is a declared physical working cylinder and

$$\frac{R_*}{2} + \Lambda_{\max}R < R_{\text{phys}},$$

then $\Phi_{\text{pre}}(Q_R^)$ has a positive spatial buffer inside that working cylinder. The displayed geometric inclusion alone does not assert containment in a prescribed physical working cylinder when this inequality fails. For each fixed time τ , the spatial map $Y \mapsto X_{\text{pre}}(\tau) + \Lambda_{\text{pre}}(\tau)Y$ is bi-Lipschitz on B_R , with $\det \nabla_Y \Phi_{\text{pre}} = \Lambda_{\text{pre}}^3$, hence Jacobian and inverse-Jacobian bounds depend only on $\Lambda_{\min}, \Lambda_{\max}$.*

Proof. From (SC5), $\Lambda_{\text{pre}} \in AC(I_*)$, $0 < \Lambda_{\min} \leq \Lambda_{\text{pre}} \leq \Lambda_{\max} < \infty$, and $X_{\text{pre}} \in AC(I_*; \mathbb{R}^3)$. The scalar ODE

$$\frac{dt_c}{d\tau} = \Lambda_{\text{pre}}(\tau)^2, \quad t_c(\tau_0) = t_*$$

therefore has an absolutely continuous strictly increasing solution on I_* , and

$$\frac{dt_c}{d\tau} = \Lambda_{\text{pre}}(\tau)^2 \quad \text{a.e. on } I_*.$$

The compatibility condition (SC5) identifies this absolutely continuous frame with the selected physical core a.e.; no differentiability of x_c or ρ_c beyond the frame regularity in (SC5) is used. For fixed τ , $\nabla_Y \Phi_{\text{pre}}(\tau, \cdot) = \Lambda_{\text{pre}}(\tau)I$, so

$$\det \nabla_Y \Phi_{\text{pre}}(\tau, \cdot) = \Lambda_{\text{pre}}(\tau)^3, \quad \Lambda_{\min}^3 \leq \det \nabla_Y \Phi_{\text{pre}} \leq \Lambda_{\max}^3.$$

Hence, if $Y_1, Y_2 \in B_R$,

$$\Lambda_{\min}|Y_1 - Y_2| \leq |\Phi_{\text{pre}}(\tau, Y_1) - \Phi_{\text{pre}}(\tau, Y_2)| \leq \Lambda_{\max}|Y_1 - Y_2|,$$

and therefore $\Phi_{\text{pre}}(\tau, \cdot)$ is bi-Lipschitz on B_R . Its inverse has Lipschitz constant $\Lambda_{\min}^{-1}$, so Jacobian and inverse-Jacobian bounds are controlled by $\Lambda_{\min}, \Lambda_{\max}$.

For image inclusions, if $x = X_{\text{pre}}(\tau) + \Lambda_{\text{pre}}(\tau)Y$ with $|Y| \leq R$, then

$$|x - X_{\text{pre}}(\tau)| \leq \Lambda_{\max}R, \quad X_{\text{pre}}(\tau) \in B_{R_*/2}(x_\bullet),$$

so $x \in B_{R_*/2 + \Lambda_{\max}R}(x_\bullet)$. As $\tau \in I_*$, we also have $t = t_c(\tau) \in I_{\text{ph}}$. Therefore

$$\Phi_{\text{pre}}(Q_R^*) \subset B_{R_*/2 + \Lambda_{\max}R}(x_\bullet) \times I_{\text{ph}}.$$

□

Definition 4.4 (Preliminary profile). For $(x, t) = \Phi_{\text{pre}}(\tau, Y)$, define

$$V_{\text{pre}}(Y, \tau) := \Lambda_{\text{pre}}(\tau) u(x, t), \quad P_{\text{pre}}(Y, \tau) := \Lambda_{\text{pre}}(\tau)^2 p(x, t).$$

Hence $u(x, t) = \Lambda_{\text{pre}}(\tau)^{-1} V_{\text{pre}}(Y, \tau)$ and $p(x, t) = \Lambda_{\text{pre}}(\tau)^{-2} P_{\text{pre}}(Y, \tau)$ under Φ_{pre} .

4.3. Canonical renormalized change of variables. The preliminary-frame and gauge-solvability arguments below use the canonical notation of [definition 1.1](#). The exact chain rule and the represented Navier–Stokes equation are proved only once, in [lemma 10.50](#) and [theorem 10.52](#). All later local energy, pressure, compactness, cost, and carrier arguments use that same final equation. The role of the present section is therefore only to construct the repaired path and verify that it enters the hypotheses of the canonical absolutely continuous chart theorem. For the repaired path constructed below we write

$$\Phi(\tau, Y) := (X(\tau) + \Lambda(\tau)Y, t(\tau)), \quad \frac{dt}{d\tau} = \Lambda(\tau)^2,$$

so all domain and support statements in this section refer to the same final chart.

4.4. Domain localization and support. Let $J \subset I_*$ be a compact interval on which the chart is defined and let $R > 0$. Set $Q_R^J := B_R \times J$. With Φ as above define the restricted physical support maps

$$\mathcal{Q}_{R,J}^{\text{ph}} := \Phi(Q_R^J), \quad \mathcal{Q}_{2R,J}^{\text{ph}} := \Phi(Q_{2R}^J).$$

Let

$$\mathcal{Q}_{\text{work}}^{\text{ph}} := B_{R_{\text{phys}}}(x_\bullet) \times I_{\text{ph}}$$

denote a physical working cylinder on which the suitable solution and the chosen pressure representative or pressure atlas are available.

Proposition 4.5 (Compact support transport). *Assume $\Lambda_{\min} \leq \Lambda \leq \Lambda_{\max}$ on J , and set*

$$D_X(J) := \sup_{\tau \in J} |X(\tau) - x_\bullet|.$$

For a repaired chart $X = X_{\text{pre}} + \Lambda_{\text{pre}}c$, this quantity is finite and obeys

$$D_X(J) \leq \frac{R_*}{2} + \|\Lambda_{\text{pre}}\|_{L^\infty(J)} \|c\|_{L^\infty(J)}.$$

Suppose $D_X(J) < R_{\text{phys}}$, and choose $R_{\text{adm}}(J) > 0$ so that

$$(1) \quad D_X(J) + 2\Lambda_{\max}R_{\text{adm}}(J) < R_{\text{phys}}.$$

Then, for every $0 < R \leq R_{\text{adm}}(J)$,

$$\begin{aligned} \mathcal{Q}_{R,J}^{\text{ph}} &\subset B_{D_X(J)+\Lambda_{\max}R}(x_{\bullet}) \times I_{\text{ph}}, \\ \mathcal{Q}_{2R,J}^{\text{ph}} &\subset B_{D_X(J)+2\Lambda_{\max}R}(x_{\bullet}) \times I_{\text{ph}} \subset \mathcal{Q}_{\text{work}}^{\text{ph}}. \end{aligned}$$

The spatial distance from $\mathcal{Q}_{2R,J}^{\text{ph}}$ to the lateral boundary of the working cylinder is at least

$$d_{\text{buf}}(R) := R_{\text{phys}} - D_X(J) - 2\Lambda_{\max}R > 0.$$

Consequently, $\chi_R \in C_c^\infty(B_{2R})$ pulls back to a spatially compactly supported physical cutoff that is smooth in the spatial variables and absolutely continuous in time under the AC chart, with support depending only on $R, D_X(J), \Lambda_{\min}, \Lambda_{\max}$ and with the declared positive spatial buffer $d_{\text{buf}}(R)$. Multiplication by a temporal cutoff compactly supported in the interior of J supplies the corresponding temporal buffer.

Proof. The repaired-center identity and the preliminary-center bound give

$$|X(\tau) - x_{\bullet}| \leq |X_{\text{pre}}(\tau) - x_{\bullet}| + \Lambda_{\text{pre}}(\tau)|c(\tau)| \leq \frac{R_*}{2} + \|\Lambda_{\text{pre}}\|_{L^\infty(J)}\|c\|_{L^\infty(J)},$$

which proves the stated estimate for $D_X(J)$. Next, if $Y \in B_{kR}$, where $k \in \{1, 2\}$, then

$$|x - x_{\bullet}| \leq |X(\tau) - x_{\bullet}| + \Lambda(\tau)|Y| \leq D_X(J) + k\Lambda_{\max}R.$$

Together with $t = t_c(\tau) \in I_{\text{ph}}$, this proves both image inclusions. Inequality (1) gives the strict containment and the lower bound $d_{\text{buf}}(R) > 0$.

For pull-backs of cutoffs, Φ is bi-Lipschitz at every τ , so composition

$$\chi_R^{\text{ph}}(x, t) := \chi_R\left(\frac{x - X(\tau(t))}{\Lambda(\tau(t))}\right)$$

is smooth in x for each time and is AC in time when the chart is AC. It is supported where $Y \in B_{2R}$, hence supported where $(x, t) \in \mathcal{Q}_{2R,J}^{\text{ph}} \subset \mathcal{Q}_{\text{work}}^{\text{ph}}$. The preceding bound gives its declared spatial support buffer; a compactly supported temporal factor gives the time buffer when the cutoff is used as a test function. $\square$

Remark 4.6 (Buffer entry). The compact-window representation and every transported-cutoff argument are used only after verifying (1). If no positive $R_{\text{adm}}(J)$ exists on the selected physical working cylinder, the required raw chart is unavailable on the working window. The branch is then recorded as the representation alternative (failure of the admissible raw chart on the required cylinder), rather than continuing with an unsupported cutoff or pressure representative.

4.5. Transport of local quantities under the chart. Let $t = t(\tau) = t_c(\tau)$ and, for $R > 0$, define

$$J_\tau^-(R) := (t(\tau) - \Lambda_{\max}^2 R^2, t(\tau)].$$

Define renormalized local quantities on $Q_R^*(\tau)$ by

$$\begin{aligned} A_V(\tau, R) &:= \sup_{s \in (\tau - R^2, \tau)} \int_{B_R} |V(\cdot, s)|^2 dY, & E_V(\tau, R) &:= \int_{\tau - R^2}^\tau \int_{B_R} |\nabla V|^2 dY ds, \\ C_V(\tau, R) &:= \int_{\tau - R^2}^\tau \int_{B_R} |V|^3 dY ds, & D_V(\tau, R) &:= \int_{\tau - R^2}^\tau \int_{B_R} |P - \langle P \rangle_{B_R}|^{3/2} dY ds. \end{aligned}$$

For interval notation and a center X , define

$$\begin{aligned} A_u(J, R; X) &:= \sup_{s \in J} \int_{B_R(X)} |u(\cdot, s)|^2 dx, & E_u(J, R; X) &:= \int_J \int_{B_R(X)} |\nabla u|^2 dx ds, \\ C_u(J, R; X) &:= \int_J \int_{B_R(X)} |u|^3 dx ds, & D_u(J, R; X) &:= \int_J \int_{B_R(X)} |p - \langle p \rangle_{B_R(X)}|^{3/2} dx ds. \end{aligned}$$

In the inequalities below we abbreviate

$$\begin{aligned} A_u(J, R) &:= A_u(J, R; X(\tau)), & E_u(J, R) &:= E_u(J, R; X(\tau)), \\ C_u(J, R) &:= C_u(J, R; X(\tau)), & D_u(J, R) &:= D_u(J, R; X(\tau)). \end{aligned}$$

and set

$$\Delta_X(\tau, R) := \sup_{s \in (\tau - R^2, \tau)} |X(s) - X(\tau)|, \quad R_{\text{ph}}(\tau, R) := \Lambda_{\max} R + \Delta_X(\tau, R).$$

Proposition 4.7 (Cylinder comparison on compact windows). *Assume $\Lambda_{\min} \leq \Lambda \leq \Lambda_{\max}$ and formulas of Lemma 10.50. Then for every $\tau \in I_*$, $R > 0$,*

$$\begin{aligned} A_V(\tau, R) &\leq \Lambda_{\min}^{-1} A_u(J_{\tau}^-(R), R_{\text{ph}}(\tau, R)), \\ E_V(\tau, R) &\leq \Lambda_{\min}^{-1} E_u(J_{\tau}^-(R), R_{\text{ph}}(\tau, R)), \\ C_V(\tau, R) &\leq \Lambda_{\min}^{-2} C_u(J_{\tau}^-(R), R_{\text{ph}}(\tau, R)), \\ D_V(\tau, R) &\leq \Lambda_{\min}^{-2} D_u(J_{\tau}^-(R), R_{\text{ph}}(\tau, R)). \end{aligned}$$

In particular, all four renormalized quantities are bounded whenever the physical quantities are bounded on the corresponding compact physical cylinders.

Proof. These are direct changes of variables under Φ_{pre} :

$$\begin{aligned} V(\cdot, \tau) &= \Lambda(\tau) u(\cdot, t(\tau)), & P(\cdot, \tau) &= \Lambda(\tau)^2 p(\cdot, t(\tau)), \\ \nabla_Y V &= \Lambda^2 \nabla_x u, & dY d\tau &= \Lambda^{-5} dx dt, & d\tau &= \Lambda^{-2} dt. \end{aligned}$$

Insert these formulas into each definition and use $\Lambda_{\min} \leq \Lambda \leq \Lambda_{\max}$, plus

$$\tau - s \in (0, R^2) \implies t(\tau) - \Lambda_{\max}^2 R^2 \leq t(s) \leq t(\tau).$$

Moreover, if $s \in (\tau - R^2, \tau)$ and $x = X(s) + \Lambda(s)Y$ with $|Y| < R$, then

$$|x - X(\tau)| \leq |X(s) - X(\tau)| + \Lambda(s)|Y| \leq \Delta_X(\tau, R) + \Lambda_{\max} R = R_{\text{ph}}(\tau, R).$$

So the physical image of $B_R \times (\tau - R^2, \tau)$ is contained in $B_{R_{\text{ph}}(\tau, R)}(X(\tau)) \times J_{\tau}^-(R)$, which yields the stated comparisons. $\square$

4.6. Local repaired gauge map. Fix $R > 0$, $0 < p < 3$, and choose a radial cutoff $\chi_R \in C_c^\infty(B_{2R})$ with $0 \leq \chi_R \leq 1$, $\chi_R \equiv 1$ on B_R , and $Y \cdot \nabla \chi_R \leq 0$. Let the local energy profile space be

$$\mathcal{V}_R := \left\{ V \in L^3(B_{2R}) \cap W^{1,3/2}(B_{2R}; \mathbb{R}^3) : |Y|^{-p/3} V \in L^3(B_{2R}) \right\},$$

with norm

$$\|V\|_{\mathcal{V}_R} := \|V\|_{L^3(B_{2R})} + \||Y|^{-p/3} V\|_{L^3(B_{2R})} + \|\nabla V\|_{L^{3/2}(B_{2R})}.$$

For differentiating the repaired scale and centering gauges we use the admissible gauge class

$$\mathcal{A}_{p,R} := \left\{ V : \begin{aligned} &V \in L^3(B_{4R}) \cap W_{\text{loc}}^{1,3}(B_{4R}; \mathbb{R}^3), \\ &\int_{B_{4R}} |Y|^{-p} |V|^3 dY < \infty, \quad \int_{B_{4R}} |Y|^{-p} |V| |Y \cdot \nabla V| dY < \infty, \\ &\int_{B_{4R}} |Y|^{-p-1} |V|^3 dY < \infty \end{aligned} \right\}.$$

The larger ball is only a domain buffer: all gauge integrals are supported in B_{2R} , while (μ, q) is restricted near $(1, 0)$ so that $q + \mu B_{2R} \subset B_{4R}$.

For the actual repaired-gauge implicit-function step we use the following strong local chart class.

Definition 4.8 (Strong translation-admissible gauge class). Fix $0 < p < 3$ and $R > 0$. Let $\chi_R \in C_c^\infty(B_{2R})$. A profile V belongs to $\mathcal{A}_{p,R}^{\text{tr}, C^1}$ if

$$V \in \mathcal{A}_{p,R} \cap \mathcal{V}_R$$

and there exists $\delta_0 > 0$ such that the scale functional $(q, \mu) \mapsto \mathcal{J}(q, \mu; V)$ and the three centering moments $(q, \mu) \mapsto \mathcal{M}_j(q, \mu; V)$, $j = 1, 2, 3$, are finite and actually C^1 for

$$|q| + |\mu - 1| < \delta_0, \quad \mu > 0 :$$

$$\mathcal{J}(q, \mu; V) := \int_{B_{2R}} \chi_R(Y) |Y|^{-p} |V(q + \mu Y)|^3 dY,$$

$$\begin{aligned} \mathcal{J}_{q_k}(q, \mu; V) &:= \mu^{-1} \left[p \int_{B_{2R}} \chi_R(Y) Y_k |Y|^{-p-2} |V(q + \mu Y)|^3 dY \right. \\ &\quad \left. - \int_{B_{2R}} \partial_k \chi_R(Y) |Y|^{-p} |V(q + \mu Y)|^3 dY \right], \end{aligned}$$

and

$$\mathcal{J}_\mu(q, \mu; V) := -\mu^{-1} \int_{B_{2R}} [(3-p)\chi_R(Y) + Y \cdot \nabla \chi_R(Y)] |Y|^{-p} |V(q + \mu Y)|^3 dY.$$

For the centering moments set

$$\mathcal{M}_j(q, \mu; V) := \int_{B_{2R}} Y_j \chi_R(Y) |V(q + \mu Y)|^2 dY, \quad j = 1, 2, 3,$$

and require that their actual derivatives are the displayed maps

$$\mathcal{M}_{j,q_k}(q, \mu; V) := -\mu^{-1} \int_{B_{2R}} (\delta_{jk} \chi_R(Y) + Y_j \partial_k \chi_R(Y)) |V(q + \mu Y)|^2 dY,$$

$$\mathcal{M}_{j,\mu}(q, \mu; V) := -\mu^{-1} \int_{B_{2R}} Y_j (4\chi_R(Y) + Y \cdot \nabla \chi_R(Y)) |V(q + \mu Y)|^2 dY,$$

for $j, k \in \{1, 2, 3\}$. In other words,

$$\partial_{q_k} \mathcal{J} = \mathcal{J}_{q_k}, \quad \partial_\mu \mathcal{J} = \mathcal{J}_\mu, \quad \partial_{q_k} \mathcal{M}_j = \mathcal{M}_{j,q_k}, \quad \partial_\mu \mathcal{M}_j = \mathcal{M}_{j,\mu}$$

on the chart. Thus these are not merely candidate derivative maps: the strong chart requires actual C^1 differentiability of the gauge functionals. If this property fails on a selected branch, the branch is routed to the repaired-gauge degeneracy row before any implicit-function theorem is used.

Lemma 4.9 (Gauge chart entry test). *At every point where the repaired gauge normalization is used, exactly one of the following occurs:*

- (i) *the profile belongs to the strong translation-admissible class $\mathcal{A}_{p,R}^{\text{tr}, C^1}$, and the gauge Jacobian is nondegenerate;*
- (ii) *the profile fails to belong to $\mathcal{A}_{p,R}^{\text{tr}, C^1}$, or the gauge Jacobian degenerates. This is the repaired-gauge degeneracy row.*

Proof. This is the ordered gauge test in the local state-space algorithm. The test is performed before applying the implicit-function theorem. If the strong chart conditions and Jacobian nondegeneracy hold, the branch proceeds through the regular gauge chart. If either condition fails, the first failed test is the repaired-gauge degeneracy row. No closure theorem is used in this routing step; closure of the degeneracy row is supplied later by the gauge-degeneracy closure theorem. $\square$

Definition 4.10 (Gauge actions). For $(\mu, q) \in (0, \infty) \times \mathbb{R}^3$, define

$$V_{\mu,q}(Y) := \mu V(q + \mu Y).$$

For pressure profiles use

$$P_{\mu,q}(Y) := \mu^2 P(q + \mu Y).$$

We also write $V^{\mu,q} := V_{\mu,q}$ and $P^{\mu,q} := P_{\mu,q}$.

Definition 4.11 (Finite-dimensional gauge map). For $V \in \mathcal{V}_R$ and $(\mu, q) \in (0, \infty) \times \mathbb{R}^3$, set

$$\mathcal{F}(\mu, q; V) := (\mathcal{F}_0, \mathcal{F}_1, \mathcal{F}_2, \mathcal{F}_3),$$

with

$$\mathcal{F}_0(\mu, q; V) := \int_{\mathbb{R}^3} \chi_R(Y) |Y|^{-p} |V_{\mu,q}(Y)|^3 dY - \Theta_0, \quad \mathcal{F}_j(\mu, q; V) := \int_{\mathbb{R}^3} Y_j \chi_R(Y) |V_{\mu,q}(Y)|^2 dY$$

for $j = 1, 2, 3$. The gauge conditions are $\mathcal{F}(\mu, q; V) = 0$. Codomain is $\mathbb{R}^4$, domain $(0, \infty) \times \mathbb{R}^3 \times \mathcal{V}_R$.

Definition 4.12 (Localized modulation matrix). For $J \subset I_*$ and a profile path $W : J \rightarrow \mathcal{V}_R$, define

$$J_W(\tau) := D_{(\mu,q)} \mathcal{F}(1, 0; W(\tau)) \in \mathbb{R}^{4 \times 4}.$$

The rows are ordered as the scale constraint followed by the three centering constraints.

Lemma 4.13 (Weighted integrability and differentiation under integrals). Assume $V, W \in \mathcal{A}_{p,R}$. Then $G_0(V) := \int \chi_R |Y|^{-p} |V|^3$ is finite and

$$\left. \frac{d}{d\epsilon} \right|_{\epsilon=0} \int \chi_R |Y|^{-p} |V + \epsilon W|^3 dY = \int 3\chi_R |Y|^{-p} |V| V \cdot W dY.$$

Proof. For $W \in \mathcal{A}_{p,R}$, define

$$F_\epsilon(Y) := \chi_R(Y) |Y|^{-p} |V(Y) + \epsilon W(Y)|^3.$$

For each Y , $\epsilon \mapsto F_\epsilon(Y)$ is C^1 and

$$\partial_\epsilon F_\epsilon = 3\chi_R |Y|^{-p} |V + \epsilon W| (V + \epsilon W) \cdot W.$$

For $|\epsilon| \leq 1$,

$$|F_\epsilon(Y)| + |\partial_\epsilon F_\epsilon| \leq C\chi_R |Y|^{-p} (|V|^3 + |W|^3) \in L^1(B_{2R})$$

by the defining weighted L^3 condition of $\mathcal{A}_{p,R}$. Hence dominated convergence allows differentiation under the integral, and

$$\left. \frac{d}{d\epsilon} \right|_{\epsilon=0} \int_{B_{2R}} \chi_R |Y|^{-p} |V + \epsilon W|^3 dY = \int_{B_{2R}} 3\chi_R |Y|^{-p} |V| V \cdot W dY.$$

$\square$

Lemma 4.14 (Gauge covariance). For $V \in \mathcal{V}_R$, $(\mu, q) \in (0, \infty) \times \mathbb{R}^3$, and $\alpha = 0, \dots, 3$,

$$\mathcal{F}_\alpha(1, 0; V_{\mu,q}) = \mathcal{F}_\alpha(\mu, q; V).$$

Proof. For each component this is a direct substitution:

$$\begin{aligned}\mathcal{F}_0(1, 0; V_{\mu, q}) &= \int \chi_R(Y) |Y|^{-p} |V_{\mu, q}(Y)|^3 - \Theta_0 \, dY = \mathcal{F}_0(\mu, q; V), \\ \mathcal{F}_j(1, 0; V_{\mu, q}) &= \int Y_j \chi_R(Y) |V_{\mu, q}(Y)|^2 \, dY = \mathcal{F}_j(\mu, q; V), \quad j = 1, 2, 3.\end{aligned}$$

□

Lemma 4.15 (Weighted gauge functionals are C^1 on the strong chart). *Let $0 < p < 3$, let $\chi_R \in C_c^\infty(B_{2R})$, and let $V \in \mathcal{A}_{p,R}^{\text{tr}, C^1}$. Then the weighted scale and centering gauge functionals are C^1 on the strong chart and have the displayed derivatives of [definition 4.8](#).*

Proof. This is the defining property of the strong chart $\mathcal{A}_{p,R}^{\text{tr}, C^1}$. If a selected branch does not satisfy this actual C^1 property, the ordered gauge-chart entry test [lemma 4.9](#) routes it to the repaired-gauge degeneracy row before the implicit-function theorem is invoked. □

Proposition 4.16 (Differentiability of $\mathcal{F}$). *For each fixed $V \in \mathcal{A}_{p,R}^{\text{tr}, C^1}$, the map $(\mu, q) \mapsto \mathcal{F}(\mu, q; V)$ is C^1 in a neighborhood of $(1, 0)$. In each component $\mathcal{F}_\alpha$ all derivatives under the integral sign are legitimate on the retained gauge chart. By [lemma 4.9](#), every retained repaired-gauge branch either lies in $\mathcal{A}_{p,R}^{\text{tr}, C^1}$ on the selected chart, or has already exited through the repaired-gauge degeneracy row before this proposition is invoked.*

Proof. The C^1 regularity of the gauge functionals follows from [lemma 4.15](#). The nondegeneracy of the Jacobian is part of the regular gauge-chart alternative in [lemma 4.9](#). The implicit-function theorem therefore gives the claimed local C^1 gauge map. If either the strong translation-admissible chart condition or the Jacobian nondegeneracy fails, the branch is routed to the repaired-gauge degeneracy row and is not treated by this proposition. Thus each component admits pointwise C^1 formulas and these are continuous in (μ, q) near $(1, 0)$. Therefore $(\mu, q) \mapsto \mathcal{F}(\mu, q; V) \in C^1$ near $(1, 0)$. □

Proposition 4.17 (Full local Jacobian at $(1, 0)$). *Let $V \in \mathcal{A}_{p,R}^{\text{tr}, C^1}$. Set $V_0(Y) = V(Y)$. At $(\mu, q) = (1, 0)$, the matrix*

$$J(V) := D_{(\mu, q)} \mathcal{F}(1, 0; V) \in \mathbb{R}^{4 \times 4}$$

has rows

$$J_\alpha = (\partial_\mu \mathcal{F}_\alpha, \partial_{q_1} \mathcal{F}_\alpha, \partial_{q_2} \mathcal{F}_\alpha, \partial_{q_3} \mathcal{F}_\alpha), \quad \alpha = 0, \dots, 3.$$

Writing

$$\mathcal{J}_* := \mathcal{J}(0, 1; V), \quad \mathcal{J}_{\mu,*} := \mathcal{J}_\mu(0, 1; V), \quad \mathcal{J}_{q_k,*} := \mathcal{J}_{q_k}(0, 1; V),$$

and

$$\mathcal{M}_{j,*} := \mathcal{M}_j(0, 1; V), \quad \mathcal{M}_{j,\mu,*} := \mathcal{M}_{j,\mu}(0, 1; V), \quad \mathcal{M}_{j,q_k,*} := \mathcal{M}_{j,q_k}(0, 1; V),$$

one has

$$\begin{aligned}(2) \quad \partial_\mu \mathcal{F}_0(1, 0; V) &= 3\mathcal{J}_* + \mathcal{J}_{\mu,*} = p \int \chi_R |Y|^{-p} |V_0|^3 \, dY - \int (Y \cdot \nabla \chi_R) |Y|^{-p} |V_0|^3 \, dY, \\ (3) \quad \partial_{q_k} \mathcal{F}_0(1, 0; V) &= \mathcal{J}_{q_k,*} = p \int \chi_R(Y) Y_k |Y|^{-p-2} |V_0(Y)|^3 \, dY - \int (\partial_k \chi_R)(Y) |Y|^{-p} |V_0(Y)|^3 \, dY, \\ (4) \quad \partial_\mu \mathcal{F}_j(1, 0; V) &= 2\mathcal{M}_{j,*} + \mathcal{M}_{j,\mu,*} = - \int_{\mathbb{R}^3} Y_j (2\chi_R + Y \cdot \nabla \chi_R) |V_0|^2 \, dY, \\ (5) \quad \partial_{q_k} \mathcal{F}_j(1, 0; V) &= \mathcal{M}_{j,q_k,*} = - \int_{\mathbb{R}^3} (\delta_{jk} \chi_R + Y_j \partial_k \chi_R) |V_0|^2 \, dY,\end{aligned}$$

for $j, k = 1, 2, 3$.

Proof. By gauge covariance,

$$\mathcal{F}_0(\mu, q; V) = \mu^3 \mathcal{J}(q, \mu; V) - \Theta_0, \quad \mathcal{F}_j(\mu, q; V) = \mu^2 \mathcal{M}_j(q, \mu; V).$$

Since $V \in \mathcal{A}_{p,R}^{\text{tr}, C^1}$, [lemma 4.15](#) gives the C^1 chart derivatives for $\mathcal{J}$, and the same strong chart package supplies the C^1 derivatives for $\mathcal{M}_j$. Applying the product rule at $(q, \mu) = (0, 1)$ yields

$$\partial_\mu \mathcal{F}_0(1, 0; V) = 3\mathcal{J}(0, 1; V) + \mathcal{J}_\mu(0, 1; V), \quad \partial_{q_k} \mathcal{F}_0(1, 0; V) = \mathcal{J}_{q_k}(0, 1; V),$$

and

$$\partial_\mu \mathcal{F}_j(1, 0; V) = 2\mathcal{M}_j(0, 1; V) + \mathcal{M}_{j,\mu}(0, 1; V), \quad \partial_{q_k} \mathcal{F}_j(1, 0; V) = \mathcal{M}_{j,q_k}(0, 1; V).$$

The displayed integral formulas are exactly the explicit $(0, 1)$ -formulas for these chart derivatives. In particular, the singular-weight first row is read through $\mathcal{J}_\mu$ and $\mathcal{J}_{q_k}$, not through the raw term $\chi_R |Y|^{-p} |V| V \cdot \partial_k V$. $\square$

Proposition 4.18 (Translation block and Schur-complement invertibility). *Let $V : J \rightarrow \mathcal{A}_{p,R}^{\text{tr}, C^1}$ satisfy the repaired gauge constraints*

$$\mathcal{F}(1, 0; V(\tau)) = 0 \quad \text{for a.e. } \tau \in J.$$

Write the modulation matrix in block form

$$J(V(\tau)) = \begin{pmatrix} \alpha(\tau) & u(\tau)^T \\ v(\tau) & A(\tau) \end{pmatrix},$$

where $\alpha = \partial_\mu \mathcal{F}_0$, $u_\ell = \partial_{q_\ell} \mathcal{F}_0$, $v_j = \partial_\mu \mathcal{F}_j$, and $A_{j\ell} = \partial_{q_\ell} \mathcal{F}_j$. Assume that for a.e. $\tau \in J$

$$m_\chi(\tau) := \int |V(Y, \tau)|^2 \chi_R(Y) dY \geq m_0 > 0,$$

$$\int_{A_R} |V(Y, \tau)|^2 dY \leq \varepsilon_0, \quad 3C_\chi \varepsilon_0 \leq \frac{m_0}{2},$$

where $A_R = \{R \leq |Y| \leq 2R\}$ and

$$C_\chi := \max_{j,\ell} \|Y_j \partial_\ell \chi_R\|_{L^\infty}.$$

Assume also

$$C_u := \text{ess sup}_{\tau \in J} \int |Y|^{-p-1} |V(Y, \tau)|^3 dY < \infty$$

and, with $C_\nabla := \| |Y| \nabla \chi_R \|_{L^\infty}$,

$$6(p + C_\nabla) \left(\frac{2}{m_0} \right) C_u R C_\nabla \varepsilon_0 \leq \frac{p\Theta_0}{2}.$$

Then $J(V(\tau))$ is invertible for a.e. $\tau \in J$, and

$$\text{ess sup}_{\tau \in J} \|J(V(\tau))^{-1}\|_{\infty \rightarrow \infty} \leq K(R, p, \Theta_0, m_0, C_u, \varepsilon_0, \chi_R).$$

Proof. By [proposition 4.17](#), the first row is given by the strong-chart derivatives

$$\alpha = 3\mathcal{J}(0, 1; V(\tau)) + \mathcal{J}_\mu(0, 1; V(\tau)), \quad u_\ell = \mathcal{J}_{q_\ell}(0, 1; V(\tau)),$$

and the remaining rows by

$$v_j = 2\mathcal{M}_j(0, 1; V(\tau)) + \mathcal{M}_{j,\mu}(0, 1; V(\tau)), \quad A_{j\ell} = \mathcal{M}_{j,q_\ell}(0, 1; V(\tau)).$$

Thus the Jacobian is evaluated entirely in the strong chart already built into $\mathcal{A}_{p,R}^{\text{tr}, C^1}$, and no singular derivative of V under the weight is used. The entries are finite by $V(\tau) \in \mathcal{A}_{p,R}$ and

the compact support of χ_R . On the gauge surface, $\mathcal{J}(0, 1; V(\tau)) = \Theta_0$. Hence the scale entry satisfies the signed transversality bound

$$\alpha(\tau) = p \int \chi_R |Y|^{-p} |V(Y, \tau)|^3 dY - \int (Y \cdot \nabla \chi_R) |Y|^{-p} |V(Y, \tau)|^3 dY \geq p\Theta_0$$

on the gauge surface, because $Y \cdot \nabla \chi_R \leq 0$. Thus $|\alpha| \geq p\Theta_0$.

For the centering block, integration by parts gives

$$A_{j\ell} = -\delta_{j\ell} \int |V|^2 \chi_R - \int |V|^2 Y_j \partial_\ell \chi_R.$$

Writing $A = -m_\chi I_3 + E$, the annular support of $\nabla \chi_R$ gives

$$|E_{j\ell}| \leq C_\chi \int_{A_R} |V|^2 \leq C_\chi \varepsilon_0, \quad \|E\|_{\infty \rightarrow \infty} \leq 3C_\chi \varepsilon_0 \leq m_0/2.$$

Thus $A = -m_\chi(I_3 - m_\chi^{-1}E)$ and $\|m_\chi^{-1}E\|_{\infty \rightarrow \infty} \leq 1/2$, so A^{-1} exists and $\|A^{-1}\|_{\infty \rightarrow \infty} \leq 2/m_0$.

For the scale-to-translation block, the scale-constraint formula gives

$$|u_\ell| \leq p \int |Y|^{-p-1} |V|^3 dY + \int |\partial_\ell \chi_R| |Y|^{-p} |V|^3 dY \leq (p + C_\nabla) C_u,$$

so $\|u\|_\infty \leq (p + C_\nabla) C_u$. For the translation-to-scale block,

$$v_j = - \int |V|^2 Y_j (Y \cdot \nabla \chi_R) dY$$

because the centering constraint cancels the interior term. Therefore

$$\|v\|_\infty \leq 2R C_\nabla \varepsilon_0.$$

The Schur complement $S = \alpha - u^T A^{-1} v$ satisfies

$$|u^T A^{-1} v| \leq 3\|u\|_\infty \|A^{-1}\|_{\infty \rightarrow \infty} \|v\|_\infty \leq 6(p + C_\nabla) \left(\frac{2}{m_0} \right) C_u R C_\nabla \varepsilon_0 \leq \frac{p\Theta_0}{2}.$$

Since $|\alpha| \geq p\Theta_0$, this gives $|S| \geq p\Theta_0/2$. The block inverse formula

$$J^{-1} = \begin{pmatrix} S^{-1} & -S^{-1} u^T A^{-1} \\ -A^{-1} v S^{-1} & A^{-1} + A^{-1} v S^{-1} u^T A^{-1} \end{pmatrix}$$

then gives the displayed bound. $\square$

4.7. Retained repaired-gauge selection in time.

Definition 4.19 (Retained nondegenerate repaired-gauge regime). Let $J \subset I_*$ be compact. The retained nondegenerate repaired-gauge regime over the preliminary profile V_{pre} consists of branches for which there are

$$\mu \in AC(J; (0, \infty)), \quad q \in AC(J; \mathbb{R}^3),$$

such that, for a.e. $\tau \in J$,

$$V(\tau) := V_{\text{pre}}^{\mu(\tau), q(\tau)} \in \mathcal{A}_{p,R}^{\text{tr}, C^1}, \quad \mathcal{F}(1, 0; V(\tau)) = 0.$$

The path is quantitatively nondegenerate if the translation-block, annular, weighted first-moment, and Schur-complement bounds in Proposition 4.18 hold on J . If no such local AC selection exists, if the strong chart condition $\mathcal{A}_{p,R}^{\text{tr}, C^1}$ fails, or if the quantitative nondegeneracy hypotheses fail persistently, the branch lies in the gauge-degenerate, multibubble, or cascade alternatives rather than in the retained regime.

Proposition 4.20 (Gauge regularity on the retained regime). *Let (μ, q) be the repaired-gauge selection of the retained nondegenerate regime on J . Then*

$$\mathcal{F}(\mu(\tau), q(\tau); V_{\text{pre}}(\tau)) = 0 \quad \text{for a.e. } \tau \in J,$$

and $(\mu, q) \in AC(J; \mathbb{R}^4)$ by definition. If the path is quantitatively nondegenerate, then

$$\text{ess sup}_{\tau \in J} \|J(V(\tau))^{-1}\| < \infty.$$

Proof. The constraint follows from Lemma 4.14 applied to $V(\tau) = V_{\text{pre}}^{\mu(\tau), q(\tau)}$. The absolute continuity is part of Definition 4.19. The hypotheses of Proposition 4.18 are part of the retained nondegenerate regime through the strong chart entry test Lemma 4.9. The uniform inverse bound is Proposition 4.18. $\square$

Definition 4.21 (Repaired chart). Define

$$\Lambda(\tau) := \Lambda_{\text{pre}}(\tau)\mu(\tau), \quad X(\tau) := X_{\text{pre}}(\tau) + \Lambda_{\text{pre}}(\tau)q(\tau),$$

$$V(\tau) := V_{\text{pre}}^{\mu(\tau), q(\tau)}, \quad P(\tau) := P_{\text{pre}}^{\mu(\tau), q(\tau)} \text{ (same physical pressure pull-back under } \Phi).$$

Theorem 4.22 (AC repaired parameters). *Assume $SC(R_*, r_*, M, \eta, \kappa)$, and suppose the branch lies in the retained nondegenerate repaired-gauge regime on J , with selection (μ, q) . Then the corrected chart (Λ, X) is absolutely continuous on J , satisfies*

$$a(\tau) = \frac{\Lambda'(\tau)}{\Lambda(\tau)} \in L_{\text{loc}}^1(J), \quad b(\tau) = \frac{X'(\tau)}{\Lambda(\tau)} \in L_{\text{loc}}^1(J),$$

and the gauge constraints $\mathcal{F}(1, 0; V(\tau)) = 0$ hold a.e. Equivalently,

$$\mathcal{F}(\mu(\tau), q(\tau); V_{\text{pre}}(\tau)) = 0$$

for the repaired-gauge selection (μ, q) .

Proof. By Definition 4.19, $(\mu, q) \in AC(J)$ and the repaired gauge constraints hold for $V(\tau)$ a.e. By definition $V(\tau) = V_{\text{pre}}^{\mu(\tau), q(\tau)}$, so Lemma 4.14 applied at $(\mu, q) = (\mu(\tau), q(\tau))$ gives

$$\mathcal{F}(1, 0; V_{\text{pre}}^{\mu, q}) = \mathcal{F}(\mu, q; V_{\text{pre}})$$

implies $\mathcal{F}(1, 0; V(\tau)) = 0$ a.e.

Since $\Lambda = \Lambda_{\text{pre}}\mu$ and $X = X_{\text{pre}} + \Lambda_{\text{pre}}q$, and $\Lambda_{\text{pre}}, X_{\text{pre}}$ are AC by (SC5), the products and sums preserve AC, hence $(\Lambda, X) \in AC(I_*)$. Differentiating,

$$a = \frac{\Lambda'}{\Lambda} = \frac{\Lambda'_{\text{pre}}}{\Lambda_{\text{pre}}} + \frac{\mu'}{\mu} \in L_{\text{loc}}^1, \quad b = \frac{X'}{\Lambda} = \frac{X'_{\text{pre}} + \Lambda'_{\text{pre}}q + \Lambda_{\text{pre}}q'}{\Lambda_{\text{pre}}\mu} \in L_{\text{loc}}^1.$$

This gives the claimed regularity. $\square$

Remark 4.23 (Gauge degeneracy). Failure of invertibility of J is the gauge-degenerate alternative to the retained single-core regime for this fixed renormalization mechanism. In the Type II decomposition of alternatives this case belongs to the multibubble/cascade and scale-rigidity alternatives, rather than to the representation theorem.

4.8. Pressure transport and decomposition.

Lemma 4.24 (Renormalized pressure equation). *If V, P satisfy the renormalized Navier–Stokes equation in Theorem 10.52, then*

$$-\Delta P = \partial_i \partial_j (V_i V_j) \quad \text{in } \mathcal{D}'(Q_{2R}^*).$$

Proof. Take test $\varphi \in C_c^\infty(Q_{2R}^*)$. Since $V, P \in L_{\text{loc}}^2 \times L_{\text{loc}}^{3/2}$ on Q_{2R}^* , all pairings below are legitimate. Expand

$$\nabla \cdot ((V \cdot \nabla) V) = \sum_{i,j} \partial_i \partial_j (V_i V_j) - \sum_{i,j} \partial_i V_i \partial_j V_j.$$

Because $\nabla \cdot V = 0$, this reduces to $\partial_i \partial_j (V_i V_j)$. Taking divergence of

$$\partial_\tau V + (V \cdot \nabla) V + \nabla P = \nu \Delta V + a(V + Y \cdot \nabla V) + b \cdot \nabla V$$

and using $\nabla \cdot V = 0$, $\nabla \cdot (V + Y \cdot \nabla V) = 0$, $\nabla \cdot (b \cdot \nabla V) = b \cdot \nabla (\nabla \cdot V) = 0$, and $\nabla \cdot \Delta V = \Delta (\nabla \cdot V) = 0$, the following identity holds in $\mathcal{D}'$:

$$\Delta P + \partial_i \partial_j (V_i V_j) = 0.$$

Hence $-\Delta P = \partial_i \partial_j (V_i V_j)$ on Q_{2R}^* . □

Theorem 4.25 (Canonical localized pressure decomposition). *Let $0 < r < R$, let $\zeta \in C_c^\infty(B_R)$ satisfy $\zeta \equiv 1$ on B_r , and assume*

$$V \in L_{\text{loc}}^2(I_*; L^3(B_R)), \quad P \in L_{\text{loc}}^1(I_*; L^{3/2}(B_R)).$$

Assume also the canonical pressure identity

$$-\Delta P = \partial_i \partial_j (V_i V_j) \quad \text{in } \mathcal{D}'(B_R \times I_*),$$

which follows from Theorem 10.52 and Lemma 4.24. For a.e. $\tau \in I_$ there exists*

$$P(\cdot, \tau) = P_{\text{loc}}(\cdot, \tau) + H(\cdot, \tau)$$

on B_r such that

$$\begin{cases} -\Delta P_{\text{loc}}(\cdot, \tau) = \partial_i \partial_j (\zeta V_i V_j), \\ -\Delta H(\cdot, \tau) = 0 \text{ on } B_r, \quad \int_{B_r} H(\cdot, \tau) dY = 0. \end{cases}$$

Moreover

$$\|P_{\text{loc}}(\tau)\|_{L^{3/2}(\mathbb{R}^3)} \lesssim \|V(\tau)\|_{L^3(B_R)}^2,$$

and, for each $1 \leq \ell < \infty$ and $0 < r' < r$,

$$\|H(\tau)\|_{L^\ell(B_{r'})} \lesssim_{\ell, R, r, r'} \|P(\tau)\|_{L^{3/2}(B_R)} + \|V(\tau)\|_{L^3(B_R)}^2.$$

Thus the source-generated pressure and the harmonic component are recorded separately; no convergence or vanishing of H follows from velocity convergence alone.

Proof. For fixed τ , set

$$P_{\text{loc}} = \mathcal{R}_i \mathcal{R}_j (\zeta V_i V_j), \quad H := P - P_{\text{loc}}.$$

$\zeta V_i V_j \in L_{\text{loc}}^1(I_*; L^{3/2}(\mathbb{R}^3))$, so the Calderón–Zygmund operator $f \mapsto \mathcal{R}_i \mathcal{R}_j f$ gives

$$P_{\text{loc}}(\cdot, \tau) = \mathcal{R}_i \mathcal{R}_j (\zeta V_i V_j)(\cdot, \tau), \quad \|P_{\text{loc}}(\tau)\|_{L^{3/2}(\mathbb{R}^3)} \lesssim \|V(\tau)\|_{L^3(B_R)}^2.$$

Because $\mathcal{R}_i \mathcal{R}_j$ is linear bounded on $L^{3/2}$,

$$P_{\text{loc}} \in L_{\text{loc}}^1(I_*; L^{3/2}(\mathbb{R}^3)).$$

Since $-\Delta P_{\text{loc}} = \partial_i \partial_j (\zeta V_i V_j)$ and $\zeta \equiv 1$ on B_r ,

$$-\Delta H(\cdot, \tau) = \partial_i \partial_j ((1 - \zeta) V_i V_j) = 0 \quad \text{in } B_r,$$

so $H(\cdot, \tau)$ is harmonic on B_r . Imposing $\oint_{B_r} H = 0$ fixes the time-dependent additive constant uniquely in each time slice. For $\ell \in [1, \infty)$ and $0 < r' < r$, harmonicity and local L^p -to- L^ℓ interior estimates imply

$$\|H(\tau)\|_{L^\ell(B_{r'})} \lesssim_{\ell, r, r'} \|H(\tau)\|_{L^{3/2}(B_r)}.$$

Finally,

$$\|H(\tau)\|_{L^{3/2}(B_r)} \leq \|P(\tau)\|_{L^{3/2}(B_R)} + \|P_{\text{loc}}(\tau)\|_{L^{3/2}(\mathbb{R}^3)},$$

which gives the announced bound. $\square$

Remark 4.26 (Measurability and normalization). If V, P are measurable in τ , then $\tau \mapsto P_{\text{loc}}(\cdot, \tau)$ is measurable in $L^{3/2}(\mathbb{R}^3)$ by boundedness of $\mathcal{R}_i \mathcal{R}_j$, hence $\tau \mapsto H(\cdot, \tau)$ is measurable in each $L^\ell(B_{r'})$, $1 \leq \ell < \infty$, $0 < r' < r$. The normalization $\oint_{B_r} H(\cdot, \tau) dY = 0$ fixes the time-dependent additive constant in H uniquely for each τ , and therefore the decomposition is a legitimate renormalized pressure splitting.

4.9. Modulation from differentiated constraints.

Definition 4.27 (Admissible modulation-differentiability row). On a compact window $J \subset I_*$, a retained repaired-gauge branch satisfies the modulation-differentiability row if, for each $\alpha = 0, 1, 2, 3$,

$$G_\alpha(\tau) := \mathcal{F}_\alpha(1, 0; V(\tau))$$

belongs to $AC(J)$, and there are functions $\mathcal{R}_\alpha \in L^1(J)$ such that

$$(6) \quad G'_\alpha(\tau) = -\mathcal{R}_\alpha(\tau) + a(\tau) \partial_\mu \mathcal{F}_\alpha(1, 0; V(\tau)) + \sum_{k=1}^3 b_k(\tau) \partial_{q_k} \mathcal{F}_\alpha(1, 0; V(\tau))$$

for a.e. $\tau \in J$. The terms $\mathcal{R}_\alpha$ are defined as the limits obtained by testing the represented renormalized equation against smooth compactly supported truncations of the Gateaux gradients of $\mathcal{F}_\alpha$, then passing to the limit in the local energy and pressure topologies. Thus $\mathcal{R}_\alpha$ is a distributional truncated-test limit, not the unsupported pairing of an L_x^1 coefficient with a raw $\partial_\tau V$.

If this row fails on a compact window, the branch is routed to the modulation-failure exit rather than to the retained modulation theorem.

Lemma 4.28 (Differentiated repaired constraints under the modulation row). *If the repaired branch satisfies definition 4.27 on J , then the repaired constraints have the a.e. derivative identity (6) on J , with $\mathcal{R} = (\mathcal{R}_0, \dots, \mathcal{R}_3) \in L^1(J)^4$.*

Proof. This is exactly the retained-row assertion in definition 4.27. The definition records the admissible limiting procedure: one first pairs the equation with smooth compactly supported truncations of the Gateaux gradients and only then passes to the limit. No step uses an L_x^1 -against- L_x^2 pairing with $\partial_\tau V$. $\square$

Lemma 4.29 (Profile and chart direction identities). *For $\mathcal{A}_\tau = V + Y \cdot \nabla V$ and $\alpha = 0, 1, 2, 3$,*

$$D_V \mathcal{F}_\alpha(1, 0; V)[\mathcal{A}_\tau] = \partial_\mu \mathcal{F}_\alpha(1, 0; V).$$

For the translation directions, the Jacobian entries are read from the strong chart of proposition 4.17:

$$\partial_{q_k} \mathcal{F}_0(1, 0; V) = \mathcal{J}_{q_k}(0, 1; V), \quad \partial_{q_k} \mathcal{F}_j(1, 0; V) = \mathcal{M}_{j, q_k}(0, 1; V)$$

for $j, k = 1, 2, 3$. In particular, the singular weighted translation entries in the first row are not identified through the raw profile variation $D_V \mathcal{F}_0(1, 0; V)[\partial_k V]$, but through the strong-chart derivative $\mathcal{J}_{q_k}$.

Proof. From $V_{\mu,q}(Y) = \mu V(q + \mu Y)$,

$$\partial_\mu V_{\mu,q}|_{(1,0)} = V + Y \cdot \nabla V = \mathcal{A}_\tau.$$

Substituting this into the profile derivative formula for $\mathcal{F}_\alpha$ gives the scale identity. The translation identities are exactly the chart derivatives already computed in [proposition 4.17](#). $\square$

Theorem 4.30 (Modulation system). *Let (V, P) be the repaired pair on the compact window under consideration, and assume the branch satisfies the modulation-differentiability row [definition 4.27](#). Writing $\theta := (a, b)$, let $\mathcal{R}_\alpha$ be the truncated-test limit functions from [definition 4.27](#). Then for a.e. $\tau \in I_*$,*

$$\partial_\mu \mathcal{F}_\alpha(1, 0; V) a(\tau) + \sum_{k=1}^3 \partial_{q_k} \mathcal{F}_\alpha(1, 0; V) b_k(\tau) = \mathcal{R}_\alpha(\tau),$$

Equivalently,

$$J(\tau)\theta(\tau) = \mathcal{R}(\tau).$$

Hence

$$\theta(\tau) = J(\tau)^{-1} \mathcal{R}(\tau) \in L_{\text{loc}}^1(I_*)^4, \quad |\theta(\tau)| \leq K |\mathcal{R}(\tau)|.$$

Proof. Set $G_\alpha(\tau) := \mathcal{F}_\alpha(1, 0; V(\tau))$. By [Theorem 4.22](#), $G_\alpha(\tau) = 0$ a.e.; by [lemma 4.28](#), $G_\alpha \in AC$ and [\(6\)](#) holds. Hence $G'_\alpha = 0$ a.e., and [\(6\)](#) gives

$$\partial_\mu \mathcal{F}_\alpha(1, 0; V) a(\tau) + \sum_{k=1}^3 \partial_{q_k} \mathcal{F}_\alpha(1, 0; V) b_k(\tau) = \mathcal{R}_\alpha(\tau),$$

which is equivalent to $J(\tau)\theta(\tau) = \mathcal{R}(\tau)$. The row gives $\mathcal{R} \in L_{\text{loc}}^1(I_*)^4$, and [Proposition 4.18](#) gives $|J^{-1}(\tau)| \leq K$ a.e.; hence

$$\theta(\tau) = J^{-1}(\tau) \mathcal{R}(\tau), \quad \theta \in L_{\text{loc}}^1(I_*), \quad |\theta| \leq K |\mathcal{R}| \text{ a.e.}$$

as claimed. $\square$

Corollary 4.31 (Regularity of a, b). *For the repaired pair of [Theorem 4.30](#), $a, b \in L_{\text{loc}}^1(I_*)$. If in addition $\mathcal{R} \in L_{\text{loc}}^\infty(I_*)$, then $a, b \in L_{\text{loc}}^\infty(I_*)$.*

Proof. By [Theorem 4.30](#), $\theta = (a, b) = J^{-1} \mathcal{R}$ a.e., and $J^{-1} \in L_{\text{loc}}^\infty(I_*)$ by [Theorem 4.18](#). The first claim is therefore immediate from $\mathcal{R} \in L_{\text{loc}}^1(I_*)$, while the second follows if $\mathcal{R} \in L_{\text{loc}}^\infty$. $\square$

4.10. Repaired suitable structure in renormalized variables.

Lemma 4.32 (AC chart-generated tests are admissible). *Let (u, p) be suitable on a compact physical cylinder, and let $\Lambda, X \in AC(J)$ on a compact interval J , with $0 < \Lambda_- \leq \Lambda \leq \Lambda_+ < \infty$ and*

$$a = \frac{\Lambda'}{\Lambda}, \quad b = \frac{X'}{\Lambda} \quad \text{in } L^1(J).$$

Let $t = t(\tau)$ be defined by $dt/d\tau = \Lambda(\tau)^2$. If $\phi \in C_c^\infty(B_R \times J)$ is nonnegative and its chart pullback has a positive support buffer inside the physical cylinder, then

$$\psi(x, t) := \Lambda(\tau(t))^{-1} \phi\left(\frac{x - X(\tau(t))}{\Lambda(\tau(t))}, \tau(t)\right)^2$$

is an admissible nonnegative test for the physical local energy inequality. The chain-rule identities for $\nabla_x \psi$, $\Delta_x \psi$, and $\partial_t \psi$ hold for a.e. t .

Proof. For each fixed time, $x \mapsto \psi(x, t)$ is smooth and compactly supported; the AC chart gives

$$\psi, \nabla_x \psi, \Delta_x \psi \in L^\infty, \quad \partial_t \psi \in L_t^1 L_x^\infty$$

on the compact cylinder, with the time derivative computed by the a.e. chain rule. Extend ψ by zero across a small time buffer and set $\psi^\varepsilon := \rho_\varepsilon *_t \psi$, where ρ_ε is a nonnegative time mollifier and ε is smaller than the support buffer. Then $\psi^\varepsilon \geq 0$, $\psi^\varepsilon \in C_c^\infty$ in spacetime with compact support in the same physical cylinder, and

$$\psi^\varepsilon \rightarrow \psi, \quad \nabla_x \psi^\varepsilon \rightarrow \nabla_x \psi, \quad \Delta_x \psi^\varepsilon \rightarrow \Delta_x \psi \quad \text{in } L^\infty,$$

while

$$\partial_t \psi^\varepsilon \rightarrow \partial_t \psi \quad \text{in } L_t^1 L_x^\infty.$$

Apply the physical local energy inequality to ψ^ε . The suitable bounds $|u|^2 \in L_t^\infty L_x^1$, $|\nabla u|^2 \in L^1$, and $|u|^3 + |p||u| \in L^1$ allow passage to the limit in the right-hand side by dominated convergence and in the nonnegative dissipation term by lower semicontinuity. This gives the physical local energy inequality with ψ . $\square$

Proposition 4.33 (Transport of suitable local energy inequality). *Assume the domain-buffer condition of [proposition 4.5](#) and [remark 4.6](#) for some compact $J \subset I_*$ and $0 < R \leq R_{\text{adm}}(J)$. Suppose (u, p) is suitable and its chosen pressure representative or pressure atlas is available on $\mathcal{Q}_{\text{work}}^{\text{ph}}$, and (V, P) is obtained by the repaired map Φ . Set*

$$\tilde{a}(\tau) := \Lambda(\tau)^{-1} \Lambda'(\tau), \quad \tilde{b}(\tau) := \Lambda(\tau)^{-1} X'(\tau).$$

Then for every nonnegative $\phi \in C_c^\infty(Q_R^J)$ and $\tau_1 < \tau_2$ in J ,

$$\begin{aligned} \mathcal{E}_\phi(V, P) &:= \frac{1}{2} |V|^2 \partial_\tau(\phi^2) + \frac{\nu}{2} |V|^2 \Delta_Y \phi^2 + \frac{1}{2} |V|^2 V \cdot \nabla(\phi^2) \\ &\quad + (P_{\text{loc}} + H) V \cdot \nabla(\phi^2) - \frac{\tilde{a}}{2} |V|^2 \phi^2 - \frac{\tilde{a}}{2} |V|^2 Y \cdot \nabla(\phi^2) - \frac{1}{2} |V|^2 \tilde{b} \cdot \nabla(\phi^2). \end{aligned}$$

$$\begin{aligned} \frac{1}{2} \int_{B_R} |V(\tau_2)|^2 \phi(\cdot, \tau_2)^2 + \nu \int_{\tau_1}^{\tau_2} \int_{B_R} |\nabla V|^2 \phi^2 dY d\tau &\leq \frac{1}{2} \int_{B_R} |V(\tau_1)|^2 \phi(\cdot, \tau_1)^2 \\ &\quad + \int_{\tau_1}^{\tau_2} \int_{B_R} \mathcal{E}_\phi(V, P) dY d\tau. \end{aligned}$$

The right-hand side is controlled on compact windows by $\|\tilde{a}\|_{L^1}$, $\|\tilde{b}\|_{L^1}$, M , and the local bounds from [Theorem 4.25](#), and support bounds from [Section 4.4](#). Hence (V, P) is locally suitable on Q_R^J .

Proof. Set

$$\psi(x, t) := \Lambda(\tau(t))^{-1} \phi\left(\frac{x - X(\tau(t))}{\Lambda(\tau(t))}, \tau(t)\right)^2, \quad \tau = \tau(t), \quad d\tau/dt = \Lambda^{-2},$$

and let $t_i = t(\tau_i)$, $i = 1, 2$. By [proposition 4.5](#), the pullback of the support of ϕ lies inside $\mathcal{Q}_{\text{work}}^{\text{ph}}$ with a positive spatial buffer. The compact temporal support of ϕ , together with [lemma 4.32](#), therefore makes ψ an admissible nonnegative test for the physical local energy inequality.

$$\nabla_x \psi = \Lambda^{-2} \nabla_Y \phi^2, \quad \Delta_x \psi = \Lambda^{-3} \Delta_Y \phi^2, \quad \partial_t \psi = -\Lambda^{-3} \tilde{a} \phi^2 + \Lambda^{-3} \partial_\tau(\phi^2) - \Lambda^{-3} (\tilde{b} + \tilde{a} Y) \cdot \nabla_Y \phi^2.$$

Using $u = \Lambda^{-1} V$, $p = \Lambda^{-2} P$, $\nabla_x u = \Lambda^{-2} \nabla_Y V$, $dx dt = \Lambda^5 dY d\tau$, each term in the physical inequality becomes

$$\frac{1}{2} \int |u|^2 \psi|_{t_1}^{t_2} = \frac{1}{2} \int |V|^2 \phi^2|_{\tau_1}^{\tau_2},$$

$$\begin{aligned}
\nu \int |\nabla_x u|^2 \psi \, dx dt &= \nu \int |\nabla V|^2 \phi^2 \, dY \, d\tau, \\
\frac{1}{2} \int |u|^2 \partial_t \psi \, dx dt &= \frac{1}{2} \int |V|^2 \left(\partial_\tau (\phi^2) - \tilde{a} \phi^2 - \tilde{a} Y \cdot \nabla_Y \phi^2 - \tilde{b} \cdot \nabla_Y \phi^2 \right) \, dY \, d\tau, \\
\frac{\nu}{2} \int |u|^2 \Delta_x \psi \, dx dt &= \frac{\nu}{2} \int |V|^2 \Delta_Y \phi^2 \, dY \, d\tau, \\
\int \left(\frac{|u|^2}{2} u + p u \right) \cdot \nabla_x \psi \, dx dt &= \int \left(\frac{1}{2} |V|^2 V + (P_{\text{loc}} + H) V \right) \cdot \nabla_Y \phi^2 \, dY \, d\tau.
\end{aligned}$$

Substituting these identities and grouping terms gives the proposition's right-hand side and coefficients $\tilde{a}, \tilde{b}$. $\square$

4.11. Stability under subsequences.

Proposition 4.34. *Let $(u^{(n)}, p^{(n)})$ satisfy $\text{SC}(R_*, r_*, M, \eta, \kappa)$ with the same data parameters and let $(\Lambda^{(n)}, X^{(n)}, V^{(n)}, P^{(n)}, a^{(n)}, b^{(n)})$ be obtained from retained quantitatively nondegenerate repaired-gauge selections on the same I_* satisfying the modulation-differentiability row, with the constants in Proposition 4.18 uniform in n . Assume also that $(a^{(n)}, b^{(n)})$ is uniformly integrable on every compact time subinterval of I_* (in particular this holds under the L^∞ modulation bounds used in the scale-rigid applications). After extraction $n_j \rightarrow \infty$, there exist subsequences and limits (Λ, X, V, P, a, b) such that*

$$\begin{aligned}
\Lambda^{(n_j)} &\rightarrow \Lambda, \quad X^{(n_j)} \rightarrow X \quad \text{uniformly on compact time subintervals,} \\
V^{(n_j)} &\rightarrow V \text{ in } L^2_{\text{loc}}(I_*; L^2_{\text{loc}}), \quad V^{(n_j)} \rightharpoonup V \text{ in } L^2_{\text{loc}}(I_*; H^1_{\text{loc}}), \\
P^{(n_j)} &\rightharpoonup P \text{ in } L^{3/2}_{\text{loc}}(Q_R^*), \quad (a^{(n_j)}, b^{(n_j)}) \rightharpoonup (a, b) \text{ in } L^1_{\text{loc}}(I_*)^4.
\end{aligned}$$

Moreover the limit chart satisfies

$$\Lambda' = a\Lambda, \quad X' = b\Lambda$$

in distributions, equivalently

$$a = \frac{\Lambda'}{\Lambda}, \quad b = \frac{X'}{\Lambda}$$

a.e. on I_* , and the limit obeys all conclusions below.

Proof. From SC and the retained repaired-gauge selections, $(V^{(n)}, P^{(n)})$ are uniformly bounded in $L^\infty_\tau L^2(B_{2R}) \cap L^2_\tau H^1(B_{2R})$ and $L^{3/2}_{\text{loc}}(Q_R^*)$. Theorem 4.18 gives uniform invertibility bounds for $J^{(n)}$, and Theorem 4.30 gives $(a^{(n)}, b^{(n)})$ uniformly in L^1_{loc} (since $\mathcal{R}^{(n)} \in L^1_{\text{loc}}$). Since on compact windows

$$(\Lambda^{(n)})' = a^{(n)} \Lambda^{(n)}, \quad (X^{(n)})' = b^{(n)} \Lambda^{(n)},$$

and the charts have uniform upper and lower scale bounds, $\Lambda^{(n)}$ and $X^{(n)}$ are uniformly bounded in $W^{1,1}_{\text{loc}}$. The assumed local uniform integrability and Dunford–Pettis provide weak subsequential limits for $(a^{(n)}, b^{(n)})$. Thus, by BV compactness and one-dimensional subsequences,

$$\Lambda^{(n_j)} \rightarrow \Lambda, \quad X^{(n_j)} \rightarrow X \quad \text{uniformly on compact time subintervals,}$$

and, after extraction,

$$a^{(n_j)} \rightharpoonup a, \quad b^{(n_j)} \rightharpoonup b, \quad a^{(n_j)} \Lambda^{(n_j)} \rightharpoonup a\Lambda, \quad b^{(n_j)} \Lambda^{(n_j)} \rightharpoonup b\Lambda$$

in L^1_{loc} . Since $\Lambda^{(n_j)} \rightarrow \Lambda$ uniformly on compact time intervals after extraction, the distributional derivative identities pass to the limit and give

$$\Lambda' = a\Lambda, \quad X' = b\Lambda.$$

The chart bounds keep Λ locally bounded away from zero, so equivalently $a = \Lambda'/\Lambda$ and $b = X'/\Lambda$ a.e. From the renormalized equation,

$$\partial_\tau V^{(n)} \in L_\tau^1(H_{\text{loc}}^{-1}) + L_\tau^{4/3}(L_{\text{loc}}^{12/11}),$$

uniformly on compact windows. Aubin–Lions compactness on bounded balls gives

$$V^{(n_j)} \rightarrow V \text{ in } L_{\text{loc}}^2(I_*; L_{\text{loc}}^2), \quad V^{(n_j)} \rightharpoonup V \text{ in } L_{\text{loc}}^2(I_*; H_{\text{loc}}^1).$$

By Calderón–Zygmund and Theorem 4.25, $P^{(n_j)} \rightharpoonup P$ in $L_{\text{loc}}^{3/2}(Q_R^*)$, and with normalized pressure split $P = P_{\text{loc}} + H$, the same for P_{loc} and H . Passing to limits in $\mathcal{F}(1, 0; V^{(n)}(\tau)) = 0$, theorem 10.52, and the linear modulation system yields the corresponding limit constraints and equations. $\square$

4.12. Final representation theorem.

Theorem 4.35 (Chart and repaired-gauge representation). *Assume $\text{SC}(R_*, r_*, M, \eta, \kappa)$, and suppose the branch lies in the retained repaired-gauge regime with selection (μ, q) on $J = [\tau_0 - \ell, \tau_0 + \ell] \subset I_*$. Then there are fields (V, P) on Q_{2R}^J , with $(\Lambda, X) \in AC(J)$,*

$$a = \frac{\Lambda'}{\Lambda}, \quad b = \frac{X'}{\Lambda}, \quad a, b \in L_{\text{loc}}^1(J),$$

such that

$$u(x, t) = \Lambda(\tau)^{-1} V\left(\frac{x - X(\tau)}{\Lambda(\tau)}, \tau\right), \quad \frac{dt}{d\tau} = \Lambda(\tau)^2,$$

and $\mathcal{F}(1, 0; V(\tau)) = 0$ for $\tau \in J$.

Proof. From (SC5) and Proposition 4.3, $(\Lambda_{\text{pre}}, X_{\text{pre}})$ are AC on I_* , so V_{pre} is defined on Q_{2R}^* . The retained repaired-gauge regime supplies $(\mu, q) \in AC(J; (0, \infty) \times \mathbb{R}^3)$ and $\mathcal{F}(1, 0; V_{\text{pre}}^{\mu, q}) = 0$ a.e. By Definition 4.21 and Theorem 4.22, setting

$$\Lambda(\tau) = \Lambda_{\text{pre}}(\tau)\mu(\tau), \quad X(\tau) = X_{\text{pre}}(\tau) + \Lambda_{\text{pre}}(\tau)q(\tau), \quad V(\tau) = V_{\text{pre}}^{\mu(\tau), q(\tau)},$$

and $P(\tau) = P_{\text{pre}}^{\mu(\tau), q(\tau)}$, gives $\Lambda, X \in AC(J)$ and

$$a = \frac{\Lambda'}{\Lambda}, \quad b = \frac{X'}{\Lambda}, \quad a, b \in L_{\text{loc}}^1(J),$$

Moreover,

$$\mathcal{F}(1, 0; V(\tau)) = 0 \text{ for a.e. } \tau \in J.$$

Finally $u = \Lambda^{-1}V(\cdot, \tau) \circ \Phi^{-1}$ by definition of the repaired map, hence the physical identity with $\frac{dt}{d\tau} = \Lambda^2$ follows. $\square$

Theorem 4.36 (Pressure and modulation). *Assume, in addition to Theorem 4.35, that the retained repaired-gauge selection is quantitatively nondegenerate in the sense of Proposition 4.18 and satisfies the modulation-differentiability row of definition 4.27. Then (V, P) satisfies the renormalized Navier–Stokes equation with the logarithmic coefficient convention of theorem 10.52; there is a decomposition $P = P_{\text{loc}} + H$ on Q_R^* as in Theorem 4.25; and (a, b) is the unique solution of*

$$J(\tau) \begin{pmatrix} a(\tau) \\ b(\tau) \end{pmatrix} = \mathcal{R}(\tau), \quad \begin{pmatrix} a \\ b \end{pmatrix} \in L_{\text{loc}}^1(J)^4,$$

with bounds $|(a, b)| \leq K|\mathcal{R}|$, where K depends on $R, p, \Theta_0, m_0, C_u, \varepsilon_0, \chi_R$.

Proof. Equation and coefficient convention are inherited directly from [theorem 10.52](#) applied to (Λ, X) from [Theorem 4.35](#) and the definition of V . By [Theorem 4.25](#) applied at each time slice, the decomposition on Q_R^J is defined and unique once $\oint_{B_R} H = 0$ is imposed. The modulation identity is [Theorem 4.30](#); uniqueness follows from invertibility of J ([Theorem 4.18](#)) and the pointwise identity $(a, b) = J^{-1}\mathcal{R}$. Uniform bounds on (a, b) are obtained by the bound on J^{-1} there and the stated bound on $\mathcal{R}$. $\square$

Theorem 4.37 (Compact-window representation theorem). *Under the assumptions of [theorems 4.35](#) and [4.36](#), suppose the domain-buffer entry condition (1) holds on the selected physical working cylinder. Then, for every compact window $J \subset I_*$ and each fixed $0 < R \leq R_{\text{adm}}(J)$, setting $Q_R^J := B_R \times J$, one has:*

- (1) *the exact renormalized equation on Q_{2R}^J ,*
- (2) *the four repaired-gauge constraints $\mathcal{F}(1, 0; V(\tau)) = 0$,*
- (3) *bounds on $a, b \in L_{\text{loc}}^1(J)$ (and L_{loc}^∞ under $\mathcal{R} \in L_{\text{loc}}^\infty$),*
- (4) *local pressure decomposition $P = P_{\text{loc}} + H$ with $L^{3/2}$ local bounds and harmonic remainder controls,*
- (5) *renormalized local suitable energy inequality on Q_R^J ,*
- (6) *stability under subsequences in the topologies above for families whose modulation coefficients are locally uniformly integrable, in particular in the L^∞ -modulated retained rows used below.*

The later reductions use only these conclusions. If the domain-buffer entry condition fails, the conclusion is not invoked and the branch is assigned to the representation alternative as in [remark 4.6](#).

Proof. From [Theorem 4.35](#) and the choice of $J \subset I_*$, all quantities are defined on Q_{2R}^J . Items (1) and (2) are [theorems 4.22](#) and [10.52](#); item (3) is [Theorem 4.30](#) together with the inverse bound from [Proposition 4.18](#); item (4) is [Theorem 4.25](#); item (5) is [Proposition 4.33](#); item (6) is [Proposition 4.34](#) under the stated local uniform integrability condition. Combining these results gives the stated compact-window representation. $\square$

Remark 4.38. The preceding results separate two alternatives:

- (1) the nondegenerate single-core repaired-gauge regime, where the representation theorems apply;
- (2) gauge degeneracy ($\det J = 0$) or compact-structure failure, which belongs to the multi-bubble, cascade, or scale-rigidity alternatives.

5. MULTIBUBBLE AND CASCADE EXCLUSION

Analytic inputs.

The local reduction theorem uses the following analytic inputs from the preceding sections, with profile decompositions and decoupling estimates in their standard critical-space sense [Gal01, BL83]. The remaining work is the multibubble, cascade, and terminal-frame decoupling.

- (i) From Section 3, the compact single-core criterion theorem 3.18: a compact single-core finite-cost Type II branch with positive local CKN concentration, nondegenerate scale-center selection, pressure reconstruction, bounded compact-cylinder suitable estimates, bounded modulation, and vanishing localized dissipation is excluded.
- (ii) From Section 4, the canonical renormalized equation theorem 10.52, pressure decomposition theorem 4.25, modulation system theorem 4.30, and repaired-gauge representation theorems 4.35 to 4.37: the renormalized equation, pressure decomposition modulo functions of time, bounded modulation coefficients, renormalized suitable structure, and stability of these structures under subsequences.
- (iii) From Section 6, the canonical compact-cylinder Caccioppoli estimate theorem 10.61 and rough-core exclusion corollary 6.7.
- (iv) From Section 7, the collapse cost exclusion theorems 7.6 and 7.7 and autonomous reduced-limit theorem theorem 7.20 for scale-rigid or scale-collapsing branches.

5.1. Local setting and active frames. *Normalized local Navier–Stokes setting.* All quantities are written in rescaled local variables $(x, t) \in Q_R = B_R \times (-R^2, 0)$ with viscosity $\nu = 1$ (critical scaling has absorbed ν). A local pair (u_n, p_n) on such a cylinder satisfies

$$\partial_t u_n - \Delta u_n + (u_n \cdot \nabla) u_n + \nabla p_n = 0, \quad \nabla \cdot u_n = 0,$$

and pressure is normalised through spatial averaging

$$(p_n)_{B_\rho}(t) := |B_\rho|^{-1} \int_{B_\rho} p_n(x, t) dx$$

to factor out time-dependent constants.

For $r > 0$, $x_0 \in \mathbb{R}^3$, and $t_0 \in \mathbb{R}$, let

$$Q_r(x_0, t_0) := B_r(x_0) \times (t_0 - r^2, t_0), \quad Q_r := Q_r(0, 0) = B_r \times (-r^2, 0).$$

Let (u_n, p_n) be parabolic rescalings around a potential singular point with positive local Caffarelli–Kohn–Nirenberg concentration on compact cylinders. Concretely, for a base triple (x_n, t_n, λ_n) ,

$$u_n(y, s) = \lambda_n u(x_n + \lambda_n y, t_n + \lambda_n^2 s), \quad p_n(y, s) = \lambda_n^2 p(x_n + \lambda_n y, t_n + \lambda_n^2 s).$$

For each n and $R > 0$, define the scaling-invariant local quantities

$$A_n(R) := R^{-1} \operatorname{ess\,sup}_{-R^2 < t < 0} \int_{B_R} |u_n(x, t)|^2 dx,$$

$$E_n(R) := R^{-1} \int_{Q_R} |\nabla u_n(x, t)|^2 dx dt,$$

$$C_n(R) := R^{-2} \int_{Q_R} |u_n(x, t)|^3 dx dt, \quad D_n(R) := R^{-2} \int_{-R^2}^0 \int_{B_R} |p_n - (p_n)_{B_R}(t)|^{3/2} dx dt,$$

where

$$(p_n)_{B_R}(t) := |B_R|^{-1} \int_{B_R} p_n(x, t) dx.$$

For an active frame (x_n, λ_n) , $\lambda_n > 0$, define the critical rescaling

$$\mathcal{T}_{x_n, \lambda_n} f(y) := \lambda_n f(x_n + \lambda_n y).$$

If $\phi \in L^3_\sigma(\mathbb{R}^3)$ is observed in this frame, its contribution is

$$\phi_n(y) = \mathcal{T}_{x_n, \lambda_n} \phi(y) = \lambda_n \phi(x_n + \lambda_n y).$$

The critical Jacobian identity is exact:

$$\|\mathcal{T}_{x_n, \lambda_n} f\|_{L^3(E)}^3 = \int_{x_n + \lambda_n E} |f(x)|^3 dx$$

for every measurable $E \subset \mathbb{R}^3$. In particular

$$\|\mathcal{T}_{x_n, \lambda_n} f\|_{L^3(\mathbb{R}^3)}^3 = \|f\|_{L^3(\mathbb{R}^3)}^3,$$

so $\mathcal{T}_{x_n, \lambda_n}$ is an isometry on $L^3(\mathbb{R}^3)$.

Definition 5.1 (Bounded selected-window limit). A selected-window subsequence has a bounded selected-window limit on a compact cylinder Q_R if there exists $u_\infty \in L^\infty(Q_R)$ such that, after extracting a subsequence,

$$\forall \varepsilon > 0 \exists N : \forall n \geq N, \|u_n - u_\infty\|_{L^3(Q_R)} < \varepsilon.$$

The limit is nonzero if $u_\infty \not\equiv 0$ on Q_R .

Definition 5.2 (Local Type II sequence). A positive local concentration sequence is called a local Type II sequence if it satisfies [definition 3.2](#); equivalently, no selected-window subsequence converges strongly in L^3 on any compact cylinder Q_R to a nonzero bounded selected-window limit.

Definition 5.3 (Active frame). Fix a concentrating time sequence $t_n \uparrow T$ and a compact time interval $I \Subset (-\infty, 0]$. An active frame is a tuple

$$\mathfrak{F}^j = (x_n^j, \lambda_n^j, I, U_n^j, P_n^j), \quad \lambda_n^j \downarrow 0,$$

where $x_n^j \in \mathbb{R}^3$ and

$$U_n^j(y, s) := \lambda_n^j u(x_n^j + \lambda_n^j y, t_n + (\lambda_n^j)^2 s), \quad P_n^j(y, s) := (\lambda_n^j)^2 p(x_n^j + \lambda_n^j y, t_n + (\lambda_n^j)^2 s)$$

on $B_{2R} \times I$ for every fixed compact cylinder under consideration. The frame is active at threshold $\eta_* > 0$ if for some fixed $R_0 < \infty$

$$\liminf_{n \rightarrow \infty} \int_{Q_{R_0}} |U_n^j|^3 dy ds \geq \eta_*^3.$$

Definition 5.4 (Relations between active frames). Let $\mathfrak{F}^i$ and $\mathfrak{F}^j$ be two active frames. After passing to a subsequence exactly one of the following parameter regimes holds.

(i) They are *same-point comparable-scale* if

$$0 < c \leq \frac{\lambda_n^i}{\lambda_n^j} \leq C < \infty, \quad \frac{|x_n^i - x_n^j|}{\max(\lambda_n^i, \lambda_n^j)} \leq C$$

for all large n .

(ii) They form a *same-point cascade* if, after relabeling,

$$\frac{\lambda_n^i}{\lambda_n^j} \rightarrow 0, \quad \frac{|x_n^i - x_n^j|}{\lambda_n^j} = O(1).$$

The frame i is then the inner frame and j is an outer frame.

(iii) They are *separated-point* if

$$\frac{|x_n^i - x_n^j|}{\lambda_n^i} \rightarrow \infty \quad \text{and} \quad \frac{|x_n^i - x_n^j|}{\lambda_n^j} \rightarrow \infty.$$

Definition 5.5 (Terminal frame and perturbative remainder). Given a finite active family, a terminal frame is chosen by selecting an active scale with minimal order:

$$\lambda_n^{j*} \leq C_j \lambda_n^j \quad \text{for every active } j$$

after passing to a subsequence. Comparable frames at the same physical point are grouped before this choice; if several grouped frames remain tied, one is chosen by a fixed lexicographic ordering of their labels. A remainder r_n in a terminal frame is perturbative on $B_{2R} \times I$ if

$$\|r_n\|_{L^\infty L^3(B_{2R})} \rightarrow 0$$

and its equation residual and pressure satisfy

$$\|f_n\|_{L^1_t H^{-1}(B_R)} + \|\nabla \pi_n^r\|_{L^1_t H^{-1}(B_R)} \rightarrow 0$$

for every compact $B_R \times I$.

Definition 5.6 (Local multibubble/cascade branch). A local multibubble/cascade branch is a positive-concentration local Type II branch for which at least one of the following alternatives occurs:

- (i) the compact-cylinder suitable estimates $A_n + E_n + C_n + D_n$ are unbounded on some fixed rescaled cylinder;
- (ii) positive CKN mass splits into two or more comparable active cores;
- (iii) a strict same-point scale cascade appears inside bounded rescaled cylinders;
- (iv) separated physical points remain active in different local rescaled frames;
- (v) the localized scale or centering selection degenerates (no unique core selected).

Definition 5.7 (Active profile family). An active profile family is a family

$$\{(\phi^j, x_n^j, \lambda_n^j)\}_{j \in J}, \quad \phi^j \in L^3_\sigma(\mathbb{R}^3),$$

with finite or countable index set J , extracted along a positive local concentration sequence. In physical variables write

$$\Lambda_{\lambda, x_0} \phi(x) := \lambda^{-1} \phi\left(\frac{x - x_0}{\lambda}\right).$$

For every finite $J_0 \subset J$ the family gives an expansion

$$u(\cdot, t_n) = \sum_{j \in J_0} \Lambda_{\lambda_n^j, x_n^j} \phi^j + r_n^{J_0}$$

on the compact terminal windows under consideration. After the finite active subfamily has been selected, the corresponding remainder satisfies the perturbative bounds of [definition 5.5](#). The parameters are considered only after grouping same-point comparable-scale profiles into compound profiles as in [definition 5.8](#). A profile is active if

$$\|\phi^j\|_{L^3(\mathbb{R}^3)} \geq \eta_* > 0.$$

The family has finite critical mass when

$$\sum_{j \in J} \|\phi^j\|_{L^3(\mathbb{R}^3)}^3 \leq M_* < \infty.$$

It captures positive local concentration if every compact terminal-window sequence with

$$\liminf_{n \rightarrow \infty} \int_{Q_R} |u_n|^3 > 0$$

either contains an active profile observed in that window or is absorbed into the selected terminal frame; equivalently, no positive terminal-window concentration remains in the perturbative remainder.

Definition 5.8 (Compound profile). Let $J_k \subset J$ be a same-point comparable-scale class. Choose one representative $(x_n^{j_k}, \lambda_n^{j_k})$. After passing to a subsequence,

$$\rho_j := \lim_{n \rightarrow \infty} \frac{\lambda_n^j}{\lambda_n^{j_k}} \in (0, \infty), \quad z_j := \lim_{n \rightarrow \infty} \frac{x_n^j - x_n^{j_k}}{\lambda_n^{j_k}} \in \mathbb{R}^3.$$

The compound profile in the representative frame is the L_σ^3 -sum

$$\Phi^k(y) := \sum_{j \in J_k} \rho_j^{-1} \phi^j \left(\frac{y - z_j}{\rho_j} \right),$$

whenever J_k is finite; finiteness follows for active classes from [lemma 5.9](#). If $\|\Phi^k\|_{L^3} < \varepsilon_*$, where ε_* is the critical small-data threshold, the class is perturbative. Otherwise Φ^k is treated as one active object in its representative frame.

Lemma 5.9 (Finite active profile count). *Every active profile family with finite critical mass contains only finitely many active profiles. More precisely,*

$$\#J \leq M_* \eta_*^{-3}.$$

Proof. For each active profile $j \in J$,

$$\|\phi^j\|_{L^3(\mathbb{R}^3)}^3 \geq \eta_*^3.$$

Let $J_0 \subset J$ be finite. Summing over J_0 gives

$$\#J_0 \eta_*^3 \leq \sum_{j \in J_0} \|\phi^j\|_{L^3(\mathbb{R}^3)}^3 \leq \sum_{j \in J} \|\phi^j\|_{L^3(\mathbb{R}^3)}^3 \leq M_*.$$

Hence $\#J_0 \leq M_* \eta_*^{-3}$. Taking the supremum over all finite subsets J_0 yields the same bound for the cardinality of J , so J is finite and

$$\#J \leq M_* \eta_*^{-3}.$$

□

Lemma 5.10 (Exhaustive parameter partition). *Let $\mathfrak{F}^i, \mathfrak{F}^j$ be two active frames. After passing to a subsequence, either the pair is same-point comparable-scale, same-point cascade, or separated-point in the sense of [definition 5.4](#); if none of these relations holds in a selected compact frame, then one of the two contributions is locally invisible there in L^3 .*

Proof. Pass to a subsequence so that

$$\frac{\lambda_n^i}{\lambda_n^j} \rightarrow \ell \in [0, \infty]$$

in the extended half-line. If $0 < \ell < \infty$, then either

$$\frac{|x_n^i - x_n^j|}{\max(\lambda_n^i, \lambda_n^j)}$$

is bounded along a further subsequence, giving the same-point comparable-scale case, or it tends to infinity, giving separated points because the two scales are comparable.

If $\ell = 0$, then either

$$\frac{|x_n^i - x_n^j|}{\lambda_n^j}$$

is bounded, giving a same-point cascade with i inner and j outer, or it tends to infinity. In the latter case the j -profile is locally invisible in the i -frame by [lemma 5.13](#). The case $\ell = \infty$ is the same after interchanging i and j . These alternatives exhaust all subsequential limits of the scale ratio and normalized center distance. $\square$

Theorem 5.11 (Minimal bad-configuration extraction). *Let a positive local Type II branch admit an active profile family with finite critical mass, perturbative remainder in the sense of [definition 5.5](#), and concentration capture in the sense of [definition 5.7](#). If the branch is not already a single active terminal frame satisfying the single-core hypotheses, then after passing to a subsequence one of the following occurs:*

- (i) *compact-cylinder suitable control fails, and [lemma 5.15](#) produces a bounded-cylinder concentration core;*
- (ii) *there are at least two active frames in a same-point comparable-scale class, hence a compound profile in the sense of [definition 5.8](#);*
- (iii) *there is a same-point cascade;*
- (iv) *there are separated active physical points;*
- (v) *the scale-center selection of the selected compound or terminal frame is degenerate.*

In every case the positive local concentration is carried by an active frame or compound profile, not by the perturbative remainder.

Proof. If compact-cylinder suitable control fails, item (i) follows from [lemma 5.15](#). Assume therefore that the compact-cylinder family is bounded on the windows under consideration.

By concentration capture, positive local concentration cannot remain solely in the perturbative remainder. Hence either one selected active frame carries the terminal-window concentration, or another active frame remains non-negligible. If there is one nondegenerate selected active frame and no competing active object, the branch is in the single-core regime. Since this case is excluded by hypothesis, either the selected scale-center map is degenerate, giving item (v), or at least two active frames remain.

For two active frames, apply [lemma 5.10](#). Same-point comparable-scale frames give item (ii), same-point scale separation gives item (iii), and separated points give item (iv). The remaining alternative in [lemma 5.10](#) is local invisibility of one contribution in the selected compact frame; such a contribution is absorbed into the perturbative remainder and cannot carry the retained positive local concentration by the capture property. This gives the extraction statement. $\square$

5.2. Elementary local decoupling estimates.

Lemma 5.12 (Critical rescalings that escape are locally invisible). *Let $\phi \in L^3(\mathbb{R}^3)$, let $R < \infty$, and define*

$$W_n(y) = \rho_n \phi(z_n + \rho_n y),$$

where $\rho_n > 0$. Assume that the sets

$$z_n + \rho_n B_R$$

either have measure tending to zero, or escape to spatial infinity in the sense that for every compact set $K \subset \mathbb{R}^3$,

$$(z_n + \rho_n B_R) \cap K = \emptyset$$

for all sufficiently large n . Then

$$\|W_n\|_{L^3(B_R)} \rightarrow 0.$$

Proof. By the critical change of variables,

$$\|W_n\|_{L^3(B_R)}^3 = \int_{B_R} \rho_n^3 |\phi(z_n + \rho_n y)|^3 dy = \int_{z_n + \rho_n B_R} |\phi(x)|^3 dx.$$

Set

$$m_n := \int_{z_n + \rho_n B_R} |\phi(x)|^3 dx.$$

Since $\|W_n\|_{L^3(B_R)}^3 = m_n$, it is enough to show $m_n \rightarrow 0$.

Fix $\varepsilon > 0$. We split into the two structural alternatives.

Case 1: shrinking measure. If $|z_n + \rho_n B_R| \rightarrow 0$, then by absolute continuity of the integral $|\phi|^3 \in L^1(\mathbb{R}^3)$ there exists $\delta > 0$ such that

$$|E| < \delta \implies \int_E |\phi|^3 dx < \varepsilon.$$

Choose N with $|z_n + \rho_n B_R| < \delta$ for all $n \geq N$. Then

$$m_n = \int_{z_n + \rho_n B_R} |\phi|^3 dx < \varepsilon, \quad n \geq N.$$

Case 2: escape to infinity. Assume now that for every compact set $K \subset \mathbb{R}^3$, $(z_n + \rho_n B_R) \cap K = \emptyset$ for all sufficiently large n . By absolute continuity at infinity, choose compact $K \subset \mathbb{R}^3$ such that

$$\int_{\mathbb{R}^3 \setminus K} |\phi|^3 < \varepsilon.$$

For every such compact K , there is $N(K)$ with $(z_n + \rho_n B_R) \cap K = \emptyset$ for all $n \geq N(K)$. For these indices,

$$z_n + \rho_n B_R \subset \mathbb{R}^3 \setminus K,$$

and therefore, for $n \geq N(K)$,

$$m_n = \int_{z_n + \rho_n B_R} |\phi|^3 dx \leq \int_{\mathbb{R}^3 \setminus K} |\phi|^3 < \varepsilon.$$

In either case, for every $\varepsilon > 0$, $m_n < \varepsilon$ for n sufficiently large. Hence $m_n \rightarrow 0$, and therefore

$$\|W_n\|_{L^3(B_R)} \rightarrow 0.$$

□

Lemma 5.13 (Separated active profiles vanish in the selected frame). *Let (x_n^p, λ_n^p) and (x_n^q, λ_n^q) be two active frames with separated physical centers in the sense that*

$$\frac{|x_n^q - x_n^p|}{\lambda_n^p} \rightarrow \infty.$$

If $\phi^q \in L^3(\mathbb{R}^3)$, then the q -profile observed in the p -frame converges to zero locally in L^3 : for every $R < \infty$,

$$\left\| \frac{\lambda_n^p}{\lambda_n^q} \phi^q \left(\frac{x_n^p - x_n^q}{\lambda_n^q} + \frac{\lambda_n^p}{\lambda_n^q} y \right) \right\|_{L^3(B_R)} \rightarrow 0.$$

Proof. The displayed function has the form in [lemma 5.12](#) with

$$z_n = \frac{x_n^p - x_n^q}{\lambda_n^q}, \quad \rho_n = \frac{\lambda_n^p}{\lambda_n^q}.$$

The corresponding physical set in the q -profile variables is

$$z_n + \rho_n B_R = \frac{x_n^p - x_n^q}{\lambda_n^q} + \frac{\lambda_n^p}{\lambda_n^q} B_R.$$

For every $y \in B_R$,

$$|z_n + \rho_n y| \geq |z_n| - \rho_n R.$$

Set $E_n := z_n + \rho_n B_R$. Then

$$\text{dist}(E_n, 0) = |z_n| - \rho_n R = \frac{\lambda_n^p}{\lambda_n^q} \left(\frac{|x_n^q - x_n^p|}{\lambda_n^p} - R \right).$$

Since $\frac{|x_n^q - x_n^p|}{\lambda_n^p} \rightarrow \infty$, there exists $N_0(R)$ such that

$$\frac{|x_n^q - x_n^p|}{\lambda_n^p} - R \geq 1, \quad n \geq N_0(R),$$

and hence for such n , $\text{dist}(E_n, 0) \geq \rho_n$.

There are two cases.

(a) If $\rho_n \rightarrow 0$, then

$$|z_n + \rho_n B_R| = \rho_n^3 |B_R| \rightarrow 0.$$

Then the first alternative in [lemma 5.12](#) applies.

(b) Otherwise, after passing to a subsequence, $\rho_n \geq \rho_* > 0$. Then

$$\text{dist}(z_n + \rho_n B_R, 0) \geq \rho_* \left(\frac{|x_n^q - x_n^p|}{\lambda_n^p} - R \right) \rightarrow \infty,$$

so the sets escape every fixed compact subset of $\mathbb{R}^3$. The second alternative in [lemma 5.12](#) applies.

In both subcases the $L^3(B_R)$ -norm of the displayed expression converges to zero. $\square$

Lemma 5.14 (Strict outer scales vanish in the innermost frame). *Let several active profiles have the same physical center and scales $\lambda_n^1, \dots, \lambda_n^J$. Suppose λ_n^1 is an innermost scale and*

$$\frac{\lambda_n^1}{\lambda_n^j} \rightarrow 0 \quad (j \neq 1).$$

Then every outer profile $\phi^j \in L^3(\mathbb{R}^3)$, viewed in the λ_n^1 -frame, converges to zero in $L^3(B_R)$ for each fixed $R < \infty$. Any constant ambient velocity arising from the regular part of the outer flow is absorbed into the translation modulation.

Proof. In the innermost variables, the j -th outer contribution has the form

$$W_n^j(y) = \frac{\lambda_n^1}{\lambda_n^j} \phi^j \left(\frac{x_n^1 - x_n^j}{\lambda_n^j} + \frac{\lambda_n^1}{\lambda_n^j} y \right).$$

Set

$$\rho_n := \frac{\lambda_n^1}{\lambda_n^j} \rightarrow 0, \quad z_n^j := \frac{x_n^1 - x_n^j}{\lambda_n^j}.$$

Then

$$W_n^j(y) = \rho_n \phi^j(z_n^j + \rho_n y),$$

and by the critical change of variables,

$$\|W_n^j\|_{L^3(B_R)}^3 = \int_{z_n^j + \rho_n B_R} |\phi^j(x)|^3 dx.$$

Since $\rho_n \rightarrow 0$, the translated domains have

$$|z_n^j + \rho_n B_R| = \rho_n^3 |B_R| \rightarrow 0.$$

By absolute continuity of the integral for $\phi^j \in L^3(\mathbb{R}^3)$,

$$\|W_n^j\|_{L^3(B_R)} \rightarrow 0.$$

To treat a general ϕ^j , let $\phi_\varepsilon^j \in C_c^\infty(\mathbb{R}^3)$ with $\|\phi^j - \phi_\varepsilon^j\|_{L^3} < \varepsilon$. By critical isometry,

$$\|W_n^j - (W_n^j)_\varepsilon\|_{L^3(B_R)} \leq \|\phi^j - \phi_\varepsilon^j\|_{L^3} < \varepsilon,$$

where $(W_n^j)_\varepsilon$ is the same scaling of ϕ_ε^j . The compactly supported term has vanishing local norm by the shrinking-domain argument above, and then letting first $n \rightarrow \infty$, then $\varepsilon \downarrow 0$ implies the claimed convergence. A constant vector field from the regular part of an outer flow contributes only a Galilean drift in the selected frame; it is therefore incorporated into the translation modulation and does not remain as an exterior perturbation. $\square$

Lemma 5.15 (Compactness failure produces another core). *If the compact-cylinder suitable estimates are unbounded on a fixed rescaled cylinder, then after passing to a subsequence the local $C_n + D_n$ quantity is unbounded on a comparable cylinder. Consequently the branch contains an additional local CKN concentration core.*

Proof. Fix $r > 0$. Apply the local Caccioppoli inequality with outer radius $2r$. For each index n , suitable weak solutions satisfy, with C_r independent of n ,

$$A_n(r) + E_n(r) \leq C_r(1 + C_n(2r) + C_n(2r)^{2/3} + D_n(2r)),$$

where the pressure is normalized by subtracting spatial means $(p_n)_{B_{2r}}(t)$.

Suppose that $A_n(r) + E_n(r) + C_n(r) + D_n(r)$ is unbounded along a subsequence. If $C_n(2r) + D_n(2r)$ were bounded on this subsequence, the preceding estimate would give uniform boundedness of $A_n(r) + E_n(r)$. Monotonicity of the local integrals also gives $C_n(r) + D_n(r) \leq C_n(2r) + D_n(2r)$, contradicting unboundedness of the full family on Q_r . Therefore, after extracting a subsequence,

$$C_n(2r) + D_n(2r) \rightarrow \infty.$$

Define the localized CKN concentration at variable center and radius

$$G_n(x, t, \rho) := \rho^{-2} \int_{Q_\rho(x, t)} |u_n|^3 dx ds + \rho^{-2} \int_{Q_\rho(x, t)} |p_n - (p_n)_{B_\rho}(s)|^{3/2} dx ds.$$

Set

$$\mathbf{M}_n := \sup\{G_n(x, t, \rho) : 0 < \rho \leq 2r, Q_\rho(x, t) \subset Q_{4r}\}.$$

The choice $(x, t, \rho) = (0, 0, 2r)$ is admissible because $Q_{2r}(0, 0) \subset Q_{4r}$, so

$$\mathbf{M}_n \geq C_n(2r) + D_n(2r) \rightarrow \infty.$$

After extracting a further subsequence, we may assume $\mathbf{M}_n \geq 1$ for all n . Choose (x_n, t_n, ρ_n) with $0 < \rho_n \leq 2r$, $Q_{\rho_n}(x_n, t_n) \subset Q_{4r}$, such that

$$G_n(x_n, t_n, \rho_n) \geq \frac{1}{2} \mathbf{M}_n \geq \frac{1}{2}.$$

Rescale around (x_n, t_n, ρ_n) :

$$v_n(y, s) := \rho_n u_n(x_n + \rho_n y, t_n + \rho_n^2 s), \quad q_n(y, s) := \rho_n^2 p_n(x_n + \rho_n y, t_n + \rho_n^2 s).$$

Using the critical Jacobian identity,

$$\int_{Q_1} |v_n|^3 dy ds + \int_{Q_1} |q_n - (q_n)_{B_1}(s)|^{3/2} dy ds = G_n(x_n, t_n, \rho_n) \geq \frac{1}{2}.$$

Hence this new frame carries positive local critical mass, i.e. a new local CKN concentration core, as required. $\square$

5.3. Terminal active-frame decoupling.

Definition 5.16 (Admissible terminal active family). Fix a compact cylinder $B_{2R} \times I$ and a selected active frame. A terminal active family is admissible on this cylinder if the following hold.

- (i) The active profiles satisfy the finite-mass and active-mass conditions of [definition 5.7](#).
- (ii) Every nonselected active profile is either separated from the selected physical point in the sense of [lemma 5.13](#), or has the same physical point and a strictly outer scale in the sense of [lemma 5.14](#), after comparable same-point scales have been grouped.
- (iii) The rescaled remainder r_n satisfies

$$\|r_n\|_{L_I^\infty L^3(B_{2R})} \rightarrow 0,$$

and there exist fields (π_n^r, f_n) such that on $B_{2R} \times I$,

$$\partial_t r_n - \Delta r_n + a_n(r_n + y \cdot \nabla r_n) + b_n \cdot \nabla r_n + \nabla \pi_n^r = f_n, \quad \nabla \cdot r_n = 0,$$

with

$$\|f_n\|_{L_I^1 H^{-1}(B_R)} \rightarrow 0,$$

$$\|\nabla \pi_n^r\|_{L_I^1 H^{-1}(B_R)} \rightarrow 0.$$

- (iv) The selected velocity satisfies

$$\sup_n \|U_n\|_{L_I^\infty L^3(B_{2R})} < \infty.$$

- (v) The selected modulation coefficients satisfy

$$\sup_n (\|a_n\|_{L^\infty(I)} + \|b_n\|_{L^\infty(I)}) < \infty.$$

- (vi) Local pressure decomposition is stable under the splitting: if tensor coefficients tend to zero in $L_I^1 L^{3/2}(B_R)$, then the corresponding pressure gradients tend to zero in $L_I^1 H^{-1}(B_R)$, after subtracting spatial means.
- (vii) Cutoff commutators generated by localizing to B_R are controlled by the $L_I^\infty L^3$, $L_I^1 L^{3/2}$, and pressure-gradient norms appearing above. In particular, for any compact $\chi_R \in C_c^\infty(B_{2R})$ with $\chi_R \equiv 1$ on B_R , the commutators satisfy

$$\|[\Delta, \chi_R]u\|_{H^{-1}(B_R)} + \|[\nabla \chi_R, \nabla \cdot](u \otimes u)\|_{H^{-1}(B_R)} \leq C_R \|u\|_{L^3(B_{2R})}^2.$$

uniformly for $u \in L^3(B_{2R})$ and the same type of estimate for S_n -quadratic commutator terms.

Remark 5.17 (Terminal $L_I^\infty L^3$ inputs are perturbative hypotheses). The $L_I^\infty L^3$ -vanishing of r_n in item (iii) and the $L_I^\infty L^3$ bound on U_n in item (iv) are defining inputs of an admissible perturbative terminal active family. They are not derived from the later spacetime compactness package of [definition 7.27](#). If a physical ordered branch does not provide these pointwise critical inputs on the chosen compact cylinder, then the branch is routed first by the critical local upper row of [lemma 10.78](#) or by an earlier active-family failure, rather than being treated by [theorem 5.18](#).

Theorem 5.18 (Terminal profile-evolution decoupling). *Consider an admissible terminal active family on $B_{2R} \times I$. Let U_n be the selected-frame velocity and let S_n be the sum, in that frame, of all nonselected active profiles together with the rescaled remainder. Then on $B_R \times I$:*

- (i) $\nabla \cdot S_n = 0$ distributionally;
- (ii) $S_n \rightarrow 0$ in $L_I^\infty L^3(B_{2R})$;
- (iii) $U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0$ in $L_I^1 L^{3/2}(B_R)$;
- (iv) $\exists \pi_n, F_n : \quad \partial_t S_n - \Delta S_n + a_n(y \cdot \nabla S_n + S_n) + b_n \cdot \nabla S_n + \nabla \pi_n = F_n,$
 $\|F_n\|_{L_I^1 H^{-1}(B_R)} \rightarrow 0, \quad \nabla \cdot S_n = 0 \quad \text{in } \mathcal{D}'(B_R \times I).$
 $\|\nabla \pi_n\|_{L_I^1 H^{-1}(B_R)} \rightarrow 0.$

- (v) every pressure source containing at least one factor of S_n has vanishing pressure gradient in $L_I^1 H^{-1}(B_R)$;
- (vi) localization errors introduced by compact cutoffs vanish in $L_I^1 H^{-1}(B_R)$;
- (vii) the effective modulation parameters of the selected frame remain bounded;
- (viii) after subtracting S_n , the selected frame remains a positive finite-mass local Type II branch with pressure reconstruction and the compact-cylinder bounds required by the single-core analysis.

Proof. Each profile in the family is divergence-free, and the rescaled remainder is divergence-free. Since divergence commutes with the Navier–Stokes critical rescalings, finite sums remain divergence-free in distributions. This establishes (i).

We prove (ii). By admissibility, every nonselected active profile is either separated from the selected physical point or is a same-point strictly outer scale. The separated profiles vanish locally by lemma 5.13, and the same-point outer profiles vanish locally by lemma 5.14. The remainder tends to zero in $L_I^\infty L^3(B_{2R})$ by definition 5.16. The number of active profiles is finite by lemma 5.9. Since every summand in S_n tends to zero in $L_I^\infty L^3(B_{2R})$, the triangle inequality gives

$$\|S_n\|_{L_I^\infty L^3(B_{2R})} \leq \sum_{j \neq j_*} \|S_n^j\|_{L_I^\infty L^3(B_{2R})} + \|r_n\|_{L_I^\infty L^3(B_{2R})} \rightarrow 0.$$

This establishes (ii).

For (iii), admissibility gives

$$\sup_n \|U_n\|_{L_I^\infty L^3(B_{2R})} < \infty.$$

Hence, by Hölder's inequality,

$$\|U_n \otimes S_n\|_{L_I^1 L^{3/2}(B_R)} \leq |I| \|U_n\|_{L_I^\infty L^3(B_R)} \|S_n\|_{L_I^\infty L^3(B_R)} \rightarrow 0,$$

and the same estimate applies to $S_n \otimes U_n$. Finally,

$$\|S_n \otimes S_n\|_{L_I^1 L^{3/2}(B_R)} \leq |I| \|S_n\|_{L_I^\infty L^3(B_R)}^2 \rightarrow 0.$$

For (iv), write the represented Navier–Stokes equation for the full profile sum $U_n + S_n$ and subtract the represented equation for the selected component U_n . The remaining equation for S_n has the displayed linear left-hand side

$$\partial_t S_n - \Delta S_n + a_n(y \cdot \nabla S_n + S_n) + b_n \cdot \nabla S_n.$$

Collect all right-hand terms into

$$F_n := -\nabla \cdot (U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n) + f_n + E_n^{\text{cut}} + E_n^{\text{mod}},$$

where f_n is the remainder forcing from item (iii) of admissibility, and $E_n^{\text{cut}}, E_n^{\text{mod}}$ are localization and modulation errors. The nonlinear stresses tend to zero in $L_I^1 L^{3/2}(B_R)$, hence in $L_I^1 H^{-1}(B_R)$,

because $L^{3/2}(B_R) \hookrightarrow H^{-1}(B_R)$ on a bounded three-dimensional ball. The residual term vanishes by [definition 5.16](#). The cutoff terms vanish by the localization part of admissibility, recorded explicitly in (vi). This gives the displayed convergence in $L_I^1 H^{-1}(B_R)$.

For (v), decompose the pressure source of $U_n + S_n$ into the pure selected source, the pure exterior source, and the two mixed sources. The mixed sources are

$$\partial_i \partial_j (U_{n,i} S_{n,j} + S_{n,i} U_{n,j}) \quad \text{and} \quad \partial_i \partial_j (S_{n,i} S_{n,j}).$$

By (iii), their tensor coefficients tend to zero in $L_I^1 L^{3/2}(B_R)$. The pressure-stability part of admissibility, equivalently local Calderón–Zygmund estimates plus harmonic pressure control modulo spatial means, gives convergence of the corresponding pressure gradients to zero in $L_I^1 H^{-1}(B_R)$. Pure exterior sources are generated by profiles that are locally invisible in the selected frame; applying the same pressure stability after [lemmas 5.13](#) and [5.14](#) gives their vanishing pressure gradients.

For (vi), let χ_R be a cutoff equal to one on B_R and supported in B_{2R} . The commutators are finite sums of terms of the following types:

$$(\nabla \chi_R) \cdot \nabla S_n, \quad (\Delta \chi_R) S_n, \quad (\nabla \chi_R) \cdot (U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n),$$

plus analogous pressure cutoff terms. Their supports lie in the fixed annulus $B_{2R} \setminus B_R$. By the localization part of admissibility, these commutators are controlled by the norms already shown to vanish in (ii), (iii), and (v). Hence all localization errors vanish in $L_I^1 H^{-1}(B_R)$.

For (vii), the only exterior contribution that can remain nonzero at first order in a compact selected frame is a constant ambient velocity. This has already been absorbed into the translation modulation. The remaining exterior profiles vanish locally, so their contribution to the modulation equations is small in the same local L^3 topology used above. The selected modulation coefficients are bounded by admissibility; hence the effective parameters remain bounded.

For (viii), the positive CKN mass of the selected profile is unchanged by subtracting a term that tends to zero in local L^3 and whose mixed pressure sources vanish by (v). The pressure reconstruction is stable under the same decomposition by the pressure-stability condition in [definition 5.16](#). The compact-cylinder bounds pass to the selected component because the full family is bounded and the nonretained terms vanish on compact cylinders. Thus the selected frame remains a positive finite-mass local Type II branch with the analytic inputs needed for the single-core analysis. $\square$

5.4. Same-point and separated-point reductions.

Definition 5.19 (Scale-rigid local branch data). A scale-rigid local branch consists of a represented branch (V_n, P_n, a_n, b_n) on a fixed compact selected cylinder Q_{R_*} , a retained frame (x_n^0, λ_n^0) , and, when a second selection is present, comparison parameters (x_n^1, λ_n^1) . Write

$$\rho_n := \frac{\lambda_n^1}{\lambda_n^0}, \quad z_n := \frac{x_n^1 - x_n^0}{\lambda_n^0}.$$

The branch data satisfy:

(R1) the compact-cylinder suitable estimates are uniformly bounded,

$$A_{V_n}(R_*) + E_{V_n}(R_*) + C_{V_n}(R_*) + D_{P_n}(R_*) \leq M;$$

(R2) the retained velocity obstruction persists on a fixed inner cylinder $Q_{r_*} \Subset Q_{R_*}$,

$$C_{V_n}(r_*) \geq \eta_0 > 0;$$

a pressure-only obstruction is assigned instead to the harmonic-pressure or exterior-pressure row;

(R3) the modulation coefficients are bounded on the selected cylinder,

$$\|a_n\|_{L^\infty(-R_*^2,0)} + \|b_n\|_{L^\infty(-R_*^2,0)} \leq M;$$

(R4) one of the following two relative-parameter records is present:

(R4a) for some $L < \infty$, $e^{-L} \leq \rho_n \leq e^L$, $|z_n| \leq L$, and the finite-dimensional scale-center selection Jacobian $\mathcal{J}_n$ of the selected compound frame satisfies

$$\sigma_{\min}(\mathcal{J}_n) \rightarrow 0;$$

(R4b) the retained frame is the innermost frame of a strict same-point cascade,

$$\rho_n \rightarrow \infty, \quad \frac{|z_n|}{\rho_n} = O(1),$$

and the sum S_n of the outer profiles and terminal remainder satisfies on every fixed inner cylinder

$$\|S_n\|_{L_T^\infty L^3(B_{2R})} + \|F_n^S\|_{L_T^1 H^{-1}(B_R)} + \|\nabla \pi_n^S\|_{L_T^1 H^{-1}(B_R)} \rightarrow 0,$$

where F_n^S and π_n^S are respectively the equation residual and pressure residual after subtracting the retained component.

The tuple in (R1)–(R4) is the input contract for the scale-rigid stratification below. Its construction is distinct from the later closure of the scale-rigid row.

Theorem 5.20 (Same-point cascade reduction). *Let a same-point active family be admissible in the sense of [definition 5.16](#). After passing to a subsequence, the same-point analysis constructs the following ordered outputs:*

- (i) *every comparable-scale class is grouped into one compound retained core, or is placed in the perturbative residual when its compound profile is below the critical small-data threshold;*
- (ii) *a strict scale class is recorded in its innermost frame by the relative data*

$$\rho_n = \lambda_n^{\text{out}}/\lambda_n^{\text{in}} \rightarrow \infty, \quad z_n = (x_n^{\text{out}} - x_n^{\text{in}})/\lambda_n^{\text{in}}, \quad |z_n|/\rho_n = O(1);$$
- (iii) *a retained rigid output carries the positive velocity obstruction and the full branch data (R1)–(R4) of [definition 5.19](#);*
- (iv) *the nonretained outer profiles and remainder form a perturbative terminal component governed by [theorem 5.18](#);*
- (v) *failure of the compact suitable bounds, pressure/velocity-carrier test, or represented scale-center selection is assigned to the corresponding multicore, pressure/exterior, rough-core, or gauge-degenerate row.*

After the construction, the retained component is therefore a compound single-core candidate or a scale-rigid local branch, while the remaining component is perturbative or belongs to a named adjacent row. This theorem is a state-construction and routing theorem; it makes no scale-rigid closure assertion.

Proof. Step 1: scale-comparability classes. Split active same-point profiles into equivalence classes by

$$i \sim j \iff \exists c, C \in (0, \infty) : c \leq \frac{\lambda_n^i}{\lambda_n^j} \leq C \quad \text{for all } n$$

after passing to a subsequence. Thus within each class there are constants c_{ij}, C_{ij} with $c_{ij} \leq \lambda_n^i/\lambda_n^j \leq C_{ij}$ for all n , so one may define a single common rescaling sequence (choose $\lambda_n^{(k)} = \lambda_n^{j_k}$) and write all class- k profiles in the same coordinates. Within each class, finite sums of divergence-free profiles remain divergence-free.

Step 2: compound profile reduction inside each class. For each class set

$$\Phi_n^k := \sum_{j \in J_k} \lambda_n^j \phi^j(x_n^j + \lambda_n^j y).$$

If the L^3 -size of this rescaled sum satisfies

$$\left\| \sum_{j \in J_k} \phi^j \right\|_{L^3(\mathbb{R}^3)} < \varepsilon_\star,$$

with $\varepsilon_\star$ the perturbative small-data threshold, then standard local small-data theory for Navier–Stokes in the critical space eliminates this class from the active alternatives.

If the class is above threshold, it is retained as a single compound active core; at most one active core from that class is needed for local selection.

Step 3: local alternatives. If the retained compound core is unique and its scale-centering Jacobian is nondegenerate, we obtain the single selected core alternative. If this Jacobian is degenerate, choose an alternate scale-center selection in the same compact comparability class and record its relative parameters (ρ_n, z_n) . After passing to a subsequence, these parameters stay in a fixed compact set and $\sigma_{\min}(\mathcal{J}_n) \rightarrow 0$. Admissibility supplies the compact suitable bound, retained velocity floor, and bounded modulation in (R1)–(R3). Thus the retained output is the scale-rigid branch data (R4a) of [definition 5.19](#). If an input fails before this record is formed, its first failure is the corresponding compactness, pressure/exterior, rough-core, or gauge-degenerate adjacent row.

Step 4: strict same-point scale separation. Assume now the branch is not reduced to comparable scales and choose the innermost active frame $(x_n^{\text{in}}, \lambda_n^{\text{in}})$ at that physical point. For an outer frame $(x_n^{\text{out}}, \lambda_n^{\text{out}})$, set

$$\rho_n := \frac{\lambda_n^{\text{out}}}{\lambda_n^{\text{in}}}, \quad z_n := \frac{x_n^{\text{out}} - x_n^{\text{in}}}{\lambda_n^{\text{in}}}.$$

The same-point cascade relation gives

$$\rho_n \rightarrow \infty, \quad \frac{|z_n|}{\rho_n} = \frac{|x_n^{\text{out}} - x_n^{\text{in}}|}{\lambda_n^{\text{out}}} = O(1).$$

By [lemma 5.14](#), every profile with larger scale vanishes in the λ_n^{in} -frame in local L^3 , up to constant drift terms. These constant drifts are absorbed into translation modulation by admissibility. Let S_n be the sum of all such outer profiles plus the rescaled remainder r_n . Then, by Step 2 and [lemma 5.14](#),

$$S_n \rightarrow 0 \quad \text{in } L_I^\infty L^3(B_{2R}).$$

Subtracting the retained equation from the full-frame equation and applying [theorem 5.18](#) gives

$$\|F_n^S\|_{L_I^1 H^{-1}(B_R)} + \|\nabla \pi_n^S\|_{L_I^1 H^{-1}(B_R)} \rightarrow 0.$$

The retained velocity floor is unchanged by subtracting S_n , while the compact suitable bounds and bounded modulation are inherited from admissibility. The ordered retained outputs are therefore:

- (a) **Single-core branch:** after subtracting S_n , one nondegenerate active component remains and is passed to the compact single-core row;
- (b) **Scale-rigid branch:** the relative parameters (ρ_n, z_n) , the retained velocity floor, and the three displayed residual convergences give (R1)–(R4b) of [definition 5.19](#);
- (c) **Purely perturbative branch:** $S_n \rightarrow 0$ and the selected profile satisfies a vanishing forcing perturbation in $L_I^1 H^{-1}(B_R)$; it carries no independent same-point obstruction.

Every step above either groups an active class, records a strict relative-scale relation, constructs the scale-rigid input contract, removes a perturbative remainder, or enters a named adjacent row. These are exactly the asserted state-construction outputs. $\square$

Theorem 5.21 (Separated-point decoupling). *Separated active physical points are locally invisible in each other's active rescaled frames. Consequently every separated-point multibubble reduces in each selected active frame to a single-core candidate, scale-rigid branch data in the sense of [definition 5.19](#), or a perturbative terminal branch, provided the terminal family in that selected active frame is admissible in the sense of [definition 5.16](#). This is a construction statement and makes no use of scale-rigid closure.*

Proof. Fix an active point p with frame (x_n^p, λ_n^p) . Let $q \neq p$ be another active point. By separation,

$$\frac{|x_n^q - x_n^p|}{\lambda_n^p} \rightarrow \infty.$$

Thus [lemma 5.13](#) gives local L^3 convergence to zero of the q -profile in the p -frame on every compact ball. The same conclusion holds for each profile centered at any physical point different from p . Since only finitely many active profiles remain after the critical mass threshold is imposed, the sum of all profiles outside the p -frame tends to zero locally in L^3 .

The pure pressure generated by a separated profile is locally invisible by the same change of variables whenever its profile pressure belongs to $L_{\text{loc}}^{3/2}$ with the standard Calderón–Zygmund control. Mixed pressure terms contain at least one exterior velocity factor, hence vanish in $L^1 H_{\text{loc}}^{-1}$ by the pressure part of [theorem 5.18](#). Localization and modulation commutators are localized to $B_{2R} \setminus B_R$, so their H^{-1} -norm is controlled by the L^3 -smallness already proved in [lemmas 5.13](#) and [5.14](#) and by admissibility item (vii).

Therefore, in the selected p -frame, all other physical points contribute only a perturbative forcing tending to zero in $L^1 H^{-1}$ on compact cylinders. The selected frame remains a positive finite-mass branch with the same compact-cylinder, pressure-reconstruction, and modulation properties. If its scale-center selection is nondegenerate, it is the single-core candidate. If the selection Jacobian degenerates in a compact comparison chart, record $(\rho_n, z_n, \mathcal{J}_n)$ as in (R4a); if a strict same-point scale class remains inside the selected physical point, record its innermost parameters and decoupled residual as in (R4b). Together with the inherited bounds and velocity floor, these are exactly the scale-rigid branch data (R1)–(R4). If neither retained case occurs, the remaining component is perturbative. Since p was arbitrary, the same construction holds in every active frame. $\square$

5.5. Multibubble and cascade exclusion. The following statements isolate the analytic inputs needed for the multibubble reduction: the (already proved) single-core exclusion, the scale-rigidity reduction (proved below), and terminal active-frame admissibility (deduced here from the terminal nonlinear decoupling theorem [theorem 9.27](#)). They are stated explicitly so that the following theorem uses no unstated compactness, pressure, or profile completeness conclusion.

Proposition 5.22 (Local single-core exclusion). *Every compact single-core branch with retained positive local velocity concentration, nondegenerate scale-center selection, pressure reconstruction, bounded compact-cylinder suitable estimates, bounded modulation, and vanishing localized dissipation is excluded from the local Type II regime.*

Proof. This is precisely the conclusion of the local single-core Type II criterion, [theorem 3.18](#), applied to the represented compact single-core branch in the renormalized chart. The hypotheses listed here are the input data of that theorem: retained positive local velocity concentration provides the nondegenerate selected core, the nondegenerate scale-center selection together with

the pressure reconstruction and bounded compact-cylinder suitable estimates are the inputs of the compactness statement [proposition 3.10](#), bounded modulation is [remark 3.11](#), and vanishing localized dissipation is [definition 3.5](#). The conclusion of [theorem 3.18](#) is exactly the stated exclusion. $\square$

Remark 5.23 (Scale-rigid construction and closure order). The same-point and separated-point theorems construct scale-rigid branch data; they do not exclude those data. The remainder of this subsection consumes the input contract (R1)–(R4) of [definition 5.19](#). The local lemmas [theorems 5.27](#) and [5.38](#) and [lemmas 5.28](#) to [5.37](#) and [5.39](#) and the decomposition [theorem 5.40](#) refine that input. The sole closure edge is the final proposition [proposition 5.42](#). The refinement layer uses only the local Type II definition [definition 3.2](#), the represented variables of [theorem 4.37](#), the local pressure decomposition of [theorem 4.25](#), the modulation L^∞ control of [theorem 4.36](#), the same-point and separated-point reductions [theorems 5.20](#) and [5.21](#), the minimal bad-configuration extraction [theorem 5.11](#), the single-core exclusion [theorem 3.18](#), the terminal CKN regularity criterion [theorem 9.35](#), and Banach–Alaoglu plus Calderón–Zygmund on fixed cylinders [[CZ52](#), [Ste70](#)]. The scale-rigid discharge is not an input to any of these construction steps.

5.5.1. *Local stratification of scale-rigid branches.* We now identify the named scale-rigid local stratum and prove that every branch placed there by [theorems 5.20](#) and [5.21](#) is refined to a bounded-window exit, a compact single-core state, a perturbative residual, or an earlier-ranked structural row.

Remark 5.24 (Inputs (R1)–(R3) are automatic on the selected cylinder). The bounds (R1)–(R3) are not extra hypotheses but are inherited from the preceding reductions. Bound (R1) is the compact-cylinder hypothesis carried through [theorem 5.11](#); if it failed, [lemma 5.15](#) would assign the branch to the multicore alternative, not to the scale-rigid stratum. Bound (R2) is the velocity-carrier part of the positive concentration capture in [definition 5.7](#): after pressure-only carriers have been routed to the harmonic/exterior-pressure row, the retained scale-rigid row is entered only by a fixed positive L^3 velocity carrier on Q_{r_*} . Bound (R3) is the selected-cylinder modulation control of [theorem 4.36](#), which is part of the represented-variable output of [theorem 4.37](#). No additional analytic input is needed to enter the scale-rigid stratum.

Lemma 5.25 (Bounded selected-window limits have VMO pressure). *Let (V_n, P_n) be a represented suitable sequence on a parabolic cylinder Q_R , with the local pressure atlas and compact-cylinder suitable bounds available on every $Q_{R'} \Subset Q_R$. Assume that, after subtracting functions of time from P_n ,*

$$V_n \rightarrow V_* \quad \text{strongly in } L^3(Q_R),$$

and that

$$V_* \in L^\infty(Q_R).$$

Then, after passing to a subsequence, there exists a pressure P_* such that

$$P_n \rightharpoonup P_* \quad \text{weakly in } L_{\text{loc}}^{3/2}(Q_R),$$

and (V_*, P_*) is a suitable weak solution on every compact subcylinder of Q_R . Moreover, for every $Q_{R''} \Subset Q_R$,

$$\lim_{\rho \downarrow 0} \sup_{Q_\rho(z) \Subset Q_{R''}} \left[\rho^{-2} \iint_{Q_\rho(z)} |V_*|^3 + \rho^{-2} \iint_{Q_\rho(z)} |P_* - (P_*)_{B_\rho(x_z)}(t)|^{3/2} \right] = 0.$$

Finally, on every smaller subcylinder, the local Calderón–Zygmund pressure pieces associated to $V_n \otimes V_n$ converge strongly in $L^{3/2}$ to the corresponding local pressure piece associated to $V_* \otimes V_*$. In the same canonical decomposition, the mean-normalized harmonic remainders converge weakly

in $L^{3/2}$, remain spatially harmonic in the limit, and are controlled only through their local pressure-oscillation bounds.

Proof. Fix $Q_{4r}(z_0) \Subset Q_R$, and let $\zeta \in C_c^\infty(B_{4r}(x_0))$ satisfy $\zeta \equiv 1$ on $B_{3r}(x_0)$. Decompose, after subtracting a function of time,

$$P_n = P_n^{\text{loc}} + H_n,$$

where

$$-\Delta P_n^{\text{loc}} = \partial_i \partial_j (\zeta (V_n)_i (V_n)_j)$$

in $B_{4r}(x_0)$, and $H_n(\cdot, t)$ is harmonic in $B_{3r}(x_0)$ for a.e. t . Since

$$V_n \rightarrow V_* \quad \text{strongly in } L^3(Q_{4r}(z_0)),$$

we have

$$V_n \otimes V_n \rightarrow V_* \otimes V_* \quad \text{strongly in } L^{3/2}(Q_{4r}(z_0)).$$

Thus the Calderón–Zygmund estimate gives

$$P_n^{\text{loc}} \rightarrow P_*^{\text{loc}} \quad \text{strongly in } L^{3/2}(Q_{4r}(z_0)),$$

where

$$-\Delta P_*^{\text{loc}} = \partial_i \partial_j (\zeta (V_*)_i (V_*)_j).$$

The pressure-atlas bounds give H_n bounded in $L_t^{3/2} L_x^{3/2}(Q_{3r}(z_0))$. Since $H_n(\cdot, t)$ is harmonic in space, the interior harmonic estimates imply, for each fixed finite m ,

$$\|\nabla_x^m H_n(t)\|_{L^{3/2}(B_{2r}(x_0))} \leq C_{m,r} \|H_n(t)\|_{L^{3/2}(B_{3r}(x_0))}$$

for a.e. t . Hence, after passing to a subsequence,

$$H_n \rightharpoonup H_* \quad \text{weakly in } L_t^{3/2} W_x^{m,3/2}(B_{2r}(x_0))$$

for every fixed finite m , and $H_*(\cdot, t)$ is harmonic in $B_{2r}(x_0)$ for a.e. t . Set $P_* := P_*^{\text{loc}} + H_*$. Passing to the limit in the represented equation uses the strong L^3 convergence of V_n and the weak $L^{3/2}$ compactness of P_n ; the local energy inequality passes to the limit by lower semicontinuity of the dissipation and strong convergence of the cubic velocity terms. Thus (V_*, P_*) is suitable on compact subcylinders.

Since $V_* \in L^\infty(Q_R)$, for every compact $Q'' \Subset Q_R$,

$$C_{V_*}(2\ell; z) \lesssim \|V_*\|_{L^\infty(Q_R)}^3 \ell^3$$

uniformly for $Q_{2\ell}(z) \Subset Q''$. Fix $Q' \Subset Q_R$, choose Q'' with $Q' \Subset Q'' \Subset Q_R$, and choose a compact outer scale $\ell_0 > 0$ so that $Q_{2\ell_0}(z) \Subset Q''$ for all centres $z \in Q'$. The local $L^{3/2}$ -bound for P_* gives a finite uniform bound for $D_{P_*}(\ell_0; z)$ over these centres. Iterating [lemma 5.32](#) from the fixed outer scale ℓ_0 , with a fixed $0 < \theta \leq 1/4$, gives

$$D_{P_*}(\theta^k \ell_0; z) \leq (C_{\text{dec}} \theta^{5/2})^k D_{P_*}(\ell_0; z) + C \|V_*\|_{L^\infty(Q_R)}^3 \sum_{j=0}^{k-1} (C_{\text{dec}} \theta^{5/2})^{k-1-j} \theta^{3j} \ell_0^3.$$

Taking θ small enough that $C_{\text{dec}} \theta^{5/2} < 1$, the right side tends to zero as $k \rightarrow \infty$. The pressure-localization estimate [lemma 3.9](#) between adjacent geometric scales then yields

$$\lim_{\rho \downarrow 0} \sup_{Q_\rho(z) \Subset Q'} D_{P_*}(z, \rho) = 0.$$

Together with the trivial estimate

$$C_{V_*}(z, \rho) \lesssim \|V_*\|_{L^\infty(Q_R)}^3 \rho^3,$$

this gives the desired velocity-pressure CKN smallness for the limit on compact subcylinders.

To transfer this smallness to the approximating sequence, decompose

$$P_n - P_* = (P_n^{\text{loc}} - P_*^{\text{loc}}) + (H_n - H_*)$$

modulo functions of time on $B_{2r}(x_0)$. The first term converges strongly in $L^{3/2}$ by the Calderón–Zygmund estimate already proved. For the harmonic term, interior harmonic estimates give

$$\|\nabla_x(H_n - H_*)(t)\|_{L^\infty(B_{3r/2}(x_0))} \leq C_r \|H_n(t) - H_*(t)\|_{L^{3/2}(B_{2r}(x_0))}$$

for a.e. t . Hence for every $Q_\rho(z) \Subset Q_{R''}$,

$$\rho^{-2} \iint_{Q_\rho(z)} |(H_n - H_*) - ((H_n - H_*)_{B_\rho(x_z)}(t))|^{3/2} \leq C_r \rho^{5/2} \int_{t_z - \rho^2}^{t_z} \|H_n(t) - H_*(t)\|_{L^{3/2}(B_{2r})}^{3/2} dt.$$

The right-hand side is bounded uniformly in n by the $L_t^{3/2} L_x^{3/2}$ -bound on the harmonic pieces and therefore tends to zero uniformly as $\rho \downarrow 0$. This yields the required small-scale pressure decay for P_n after first taking ρ small and then n large. The strong convergence of the local pressure pieces has already been proved above. $\square$

Lemma 5.26 (Bounded selected-window limits are CKN-small away from pressure failure). *Let (V_n, P_n) be represented suitable windows on a compact cylinder $Q_R(z_0)$. Suppose that, after passing to a subsequence,*

$$V_n \rightarrow V_* \quad \text{strongly in } L^3(Q_R(z_0)), \quad V_* \in L^\infty(Q_R(z_0)).$$

Assume also that the local pressure atlas / pressure reconstruction row has not failed on $Q_R(z_0)$. Then for every compact subcylinder $Q \Subset Q_R(z_0)$ and every $\varepsilon > 0$ there exists $0 < \rho_Q < R$, depending only on Q , ε , the local pressure atlas, and $\|V_\|_{L^\infty(Q_R)}$, such that*

$$C_{V_n}(z, \rho_Q) + D_{P_n}(z, \rho_Q) < \varepsilon$$

for all sufficiently large n and every admissible parabolic cylinder $Q_{\rho_Q}(z) \Subset Q$. If this conclusion fails, then the first failed local test is the pressure atlas, pressure gauge, or pressure compactness row.

Proof. Fix $Q \Subset Q_R(z_0)$. Because $V_* \in L^\infty(Q_R)$, for every $z \in Q$ one has

$$C_{V_*}(z, \rho) = \rho^{-2} \int_{Q_\rho(z)} |V_*|^3 dx dt \leq \|V_*\|_{L^\infty(Q_R)}^3 \rho^3 \rightarrow 0$$

as $\rho \downarrow 0$, uniformly for $Q_\rho(z) \Subset Q$. Choose ρ_Q so that this velocity contribution is below $\varepsilon/8$ on a finite cover of Q . Strong convergence in $L^3(Q_R)$ then gives

$$C_{V_n - V_*}(z, \rho_Q) < \varepsilon/8$$

for all sufficiently large n and every admissible center in the cover. Hence the velocity part is below $\varepsilon/2$.

By lemma 5.25, the bounded selected-window limit has uniformly vanishing velocity–pressure CKN quantity at sufficiently small scales on compact subcylinders. The strong convergence of V_n in L^3 , together with strong convergence of the local Calderón–Zygmund pressure pieces and harmonic pressure decay, therefore transfers this smallness to (V_n, P_n) for all large n , unless the pressure-atlas or harmonic-pressure bounds fail. Such a failure is precisely one of the named pressure rows.

Combining the velocity and pressure estimates gives the displayed CKN smallness. If any pressure step is unavailable, the first unavailable item is exactly the pressure atlas, pressure gauge, or pressure compactness row, so the bounded window is not retained as an independent terminal branch. $\square$

Theorem 5.27 (Bounded selected-window row exits). *Let (V_n, P_n, a_n, b_n) be a represented sequence on a compact selected cylinder Q_r , and assume that, after passing to a selected-window subsequence, there is a nonzero $V_* \in L^\infty(Q_r)$ such that*

$$V_n \rightarrow V_* \quad \text{strongly in } L^3(Q_r).$$

Then this sequence cannot be a terminal branch of the strict selected-window Type II row. More precisely, after running the ordered local tests, exactly one of the following occurs:

- (i) *the local pressure atlas, pressure gauge, or pressure compactness row fails;*
- (ii) *the bounded selected window is CKN-small on a smaller compact core and is locally regular/removable by the Caffarelli–Kohn–Nirenberg regularity criterion;*
- (iii) *the retained concentration is not carried by the bounded compact core and is reselected as a smaller active core, an exterior core, a multibubble/cascade label, a gauge-degenerate label, a pressure label, or another named local row in the ordered state-space decomposition.*

Consequently a bounded selected-window limit is a closed local routing row, not an independent retained terminal Type II state. In particular, by [definition 3.2](#), the branch is not a strict selected-window Type II branch; the physical bounded-window event is handled by [theorem 11.16](#).

Proof. The existence of a nonzero bounded selected-window limit is exactly the negation of the strict selected-window condition in [definition 3.2](#), so the branch is not a strict selected-window Type II branch. This observation alone is only branch routing, not physical closure.

To close the bounded row locally, apply [lemma 5.26](#). If the pressure-atlas hypotheses required there fail, item (i) occurs. Otherwise, on each compact subcylinder of the bounded window, all sufficiently large selected windows are CKN-small after shrinking once. By the Caffarelli–Kohn–Nirenberg regularity criterion, the bounded compact core is regular and removable, giving item (ii).

If the original positive CKN concentration of the physical Type II germ still survives once this regular compact core is removed, then it cannot be carried by the bounded compact core. The ordered local selection is then rerun on the remaining concentration. Since the state-space tests are ordered before terminal retention, the surviving concentration is recorded as an inner active core, exterior core, multibubble/cascade, gauge-degenerate, pressure, or other named local row, giving item (iii). No global Type I bound, global L^3 bound, or global profile decomposition is used. $\square$

Lemma 5.28 (Pressure and modulation stability under rigid-frame limits). *Let (V_n, P_n, a_n, b_n) be a scale-rigid local branch on Q_{R_*} in the sense of [definition 5.19](#). Fix $\rho \in (r_*, R_*)$. After passing to a subsequence:*

- (i) *$V_n \rightharpoonup V_*$ weakly in $L^3(Q_\rho) \cap L_t^2 H_x^1(Q_\rho)$, and $V_n \rightarrow V_*$ strongly in $L^3(Q'_\rho)$ for every $\rho' \in (r_*, \rho)$;*
- (ii) *the mean-normalized pressure $\tilde{P}_n := P_n - \overline{P}_{n, B_\rho}$ converges weakly in $L^{3/2}(Q_\rho)$ to a limit P_* such that $-\Delta P_* = \partial_i \partial_j (V_{*,i} V_{*,j})$ on Q_ρ in the sense of distributions modulo a function of time. On every $Q_{\rho'}$, $\rho' < \rho$, the canonical decomposition is*

$$\tilde{P}_n = P_{\text{loc},n} + H_n, \quad P_* = P_{\text{loc},*} + H_*,$$

where $P_{\text{loc},n} \rightarrow P_{\text{loc},}$ strongly in $L^{3/2}(Q_{\rho'})$, $H_n \rightharpoonup H_*$ weakly there after normalization, and every H_n and H_* is spatially harmonic on the inner ball;*

- (iii) *$a_n \rightharpoonup^* a_*$ and $b_n \rightharpoonup^* b_*$ weakly-* in $L^\infty(-R_*^2, 0)$, with $\|a_*\|_\infty + \|b_*\|_\infty \leq M$;*
- (iv) *the limit (V_*, P_*, a_*, b_*) is a suitable weak solution of the represented Navier–Stokes equation [\(14\)](#) on Q_ρ , with $V_* \in L^3(Q_\rho) \cap L_t^2 H_x^1(Q_\rho)$.*

This lemma is only a compactness and stability statement. It does not assert boundedness of V_ . Bounded selected-window limits are obtained below from the separate no-subconcentration/CKN argument.*

Proof. (i) Compactness of V_n . By (R1), the family (V_n) is bounded in $L^3(Q_\rho)$ (from $C_{V_n}(R_*)$), in $L_t^2 H_x^1(Q_\rho)$ (from $A_{V_n}(R_*) + E_{V_n}(R_*)$), and in $L_t^\infty L_x^2(Q_\rho)$ (from $A_{V_n}(R_*)$). Banach–Alaoglu yields the weak limits in $L^3(Q_\rho)$ and $L_t^2 H_x^1(Q_\rho)$ along a subsequence. The represented Navier–Stokes equation (14) bounds $\partial_s V_n$ in $L_t^{4/3} H_x^{-1}(Q_\rho)$ by Sobolev product estimates on the bounded cylinder, since each term on the right-hand side is controlled by $A_{V_n}(R_*)$, $E_{V_n}(R_*)$, $C_{V_n}(R_*)$, $D_{P_n}(R_*)$, and $\|a_n\|_\infty + \|b_n\|_\infty \leq M$; the modulation correction $a_n(s)(V_n + y \cdot \nabla V_n) + b_n(s) \cdot \nabla V_n$ is bounded in $L_t^2 L_x^{6/5}(Q_\rho) \subset L_t^{4/3} H_x^{-1}(Q_\rho)$ on the bounded cylinder. Hence Aubin–Lions [Aub63, Sim87] applied on the fixed cylinder Q_ρ yields strong convergence $V_n \rightarrow V_*$ in $L^3(Q_{\rho'})$ for every $\rho' \in (r_*, \rho)$.

(ii) Pressure stability. By (R1) the local pressure quantity $D_{P_n}(R_*)$ is bounded, equivalently $\tilde{P}_n$ is bounded in $L^{3/2}(Q_\rho)$ by the local Calderón–Zygmund mean-normalization estimate [CZ52, Ste70] applied to the local pressure identity in theorem 4.25. Banach–Alaoglu yields $\tilde{P}_n \rightharpoonup P_*$ weakly in $L^{3/2}(Q_\rho)$. Choose a cutoff supported in B_ρ and equal to one on an intermediate ball containing $B_{\rho'}$. By theorem 4.25, write

$$\tilde{P}_n = P_{\text{loc},n} + H_n.$$

The strong $L^3(Q_{\rho'})$ convergence from item (i), with the cutoff supported in the intermediate ball, and Calderón–Zygmund boundedness give

$$P_{\text{loc},n} \longrightarrow P_{\text{loc},*} \quad \text{strongly in } L^{3/2}(Q_{\rho'}).$$

Subtracting this convergence from the weak convergence of $\tilde{P}_n$ gives $H_n \rightharpoonup H_* := P_* - P_{\text{loc},*}$ weakly in $L^{3/2}(Q_{\rho'})$. Harmonicity passes to the limit in distributions, and the spatial-mean normalization fixes the time-dependent constant. This proves the asserted convergence of both pressure components. Pass to the limit in the local pressure identity

$$-\Delta \tilde{P}_n = \partial_i \partial_j (V_{n,i} V_{n,j}) - \partial_i \partial_j (\overline{V_{n,i} V_{n,j}}_{B_\rho})$$

on Q_ρ . The right-hand side converges in $\mathcal{D}'(Q_\rho)$ by the strong L^3 convergence of (i) (which implies weak $L^{3/2}$ convergence of $V_{n,i} V_{n,j}$ by Hölder). This proves (ii).

(iii) Modulation limits. By (R3), a_n, b_n are bounded in $L^\infty(-R_*^2, 0)$; Banach–Alaoglu applied to the dual of L^1 yields the weak-* limits with the same bound M .

(iv) Suitable weak solution status. Each (V_n, P_n) is suitable on Q_ρ (by (R1) and the represented-equation output of theorem 4.37). The local energy inequality and the distributional momentum equation pass to the limit using (i)–(iii) and Hölder on the fixed cylinder Q_ρ : the dissipation term is lower semicontinuous under weak $L_t^2 H_x^1$ convergence, the convection term converges by strong $L^3(Q_{\rho'})$ convergence, and the pressure-velocity pairing $\int \tilde{P}_n V_n \cdot \nabla \chi$ converges by weak $L^{3/2}$ times strong L^3 on the fixed cylinder. The modulation correction converges by weak-* L^∞ for (a_n, b_n) against the strong L_t^1 convergence of $(V_n, y \cdot \nabla V_n, \nabla V_n)$ tested against fixed test functions on Q_ρ . This proves the represented suitable limit in the displayed local energy class. No L^∞ bound is inferred from the compact-cylinder estimates. $\square$

Lemma 5.29 (Compact chart comparability in scale-rigid gauges). *Let (V, P, a, b) be a represented branch on $Q_{2\rho}(z_0)$ obtained from the repaired-gauge chart of theorem 4.37. Assume that the associated chart (Λ, X) satisfies*

$$a = \frac{\Lambda'}{\Lambda}, \quad b = \frac{X'}{\Lambda}, \quad \|a\|_{L^\infty(I_{2\rho})} + \|b\|_{L^\infty(I_{2\rho})} \leq M$$

on the time interval $I_{2\rho}$ of $Q_{2\rho}(z_0)$, and set $\Lambda_0 := \Lambda(s_0) > 0$ for $z_0 = (y_0, s_0)$. There exists $\rho_c = \rho_c(M, \Lambda_0, \rho) > 0$ such that, for every $0 < r \leq \rho_c$, the chart on $Q_{2r}(z_0)$ obeys

$$\frac{\Lambda_0}{2} \leq \Lambda(s) \leq \frac{3\Lambda_0}{2}, \quad |X(s) - X(s_0)| \leq 6\Lambda_0 M r^2, \quad s \in (s_0 - 4r^2, s_0].$$

After a harmless concentric shrink, if needed, the physical image remains inside the support buffer of [proposition 4.5](#). Consequently [propositions 4.5](#) and [4.7](#) apply on the retained subcylinder with

$$\Lambda_{\min} = \Lambda_0/2, \quad \Lambda_{\max} = 3\Lambda_0/2,$$

and represented and physical CKN quantities on comparable cylinders are equivalent with constants depending only on M, Λ_0, r, ρ and the fixed buffer.

Proof. The identities

$$a = \frac{\Lambda'}{\Lambda}, \quad b = \frac{X'}{\Lambda}$$

are part of the exact repaired change of variables in [lemma 10.50](#) and [theorems 4.35](#) and [4.37](#). Hence, for $|s - s_0| \leq 4r^2$,

$$\left| \log \frac{\Lambda(s)}{\Lambda_0} \right| \leq \int_s^{s_0} |a(\sigma)| d\sigma \leq 4Mr^2, \quad |X(s) - X(s_0)| \leq \int_s^{s_0} \Lambda(\sigma) |b(\sigma)| d\sigma.$$

Choose

$$\rho_c \leq \min\{\rho/4, (\log(3/2)/(4M))^{1/2}\}$$

when $M > 0$, and choose $\rho_c = \rho/4$ when $M = 0$. Then $e^{-4Mr^2} \Lambda_0 \leq \Lambda(s) \leq e^{4Mr^2} \Lambda_0$, which implies the displayed lower-upper scale bounds. Substituting the upper scale bound into the center integral gives

$$|X(s) - X(s_0)| \leq (3\Lambda_0/2) M 4r^2 = 6\Lambda_0 M r^2.$$

A final concentric shrink, depending only on the fixed distance to the boundary of the selected compact window, keeps the image inside the domain buffer. [Propositions 4.5](#) and [4.7](#) then give the stated comparison of the represented and physical quantities. No solution boundedness, and no vanishing dissipation, is used. $\square$

Lemma 5.30 (Two-sided local CKN transfer in comparable charts). *Let (V, P) be related to a physical suitable pair (u, p) by the repaired scale-translation chart on $Q_{2r}(z_0)$, and assume the chart comparability conclusion of [lemma 5.29](#) there:*

$$0 < \Lambda_{\min} \leq \Lambda(s) \leq \Lambda_{\max} < \infty, \quad |X(s) - X(s_0)| \leq Mr^2.$$

Set

$$\mathcal{C}_V(z_0, r) := r^{-2} \int_{Q_r(z_0)} |V|^3 + r^{-2} \int_{Q_r(z_0)} |P - (P)_{B_r}|^{3/2},$$

and define the physical CKN quantity $\mathcal{C}_u(x, t, \rho)$ analogously, with pressure normalized by spatial averages. There are constants

$$N_{\text{tr}} < \infty, \quad c_{\text{tr}} \in (0, 1), \quad C_{\text{tr}} < \infty,$$

depending only on $M, \Lambda_{\min}, \Lambda_{\max}$ and the fixed support buffer, and physical cylinders

$$Q_\ell^{\text{ph}} := B_{\rho_\ell}(x_\ell) \times (t_\ell - \rho_\ell^2, t_\ell), \quad c_{\text{tr}} \Lambda_{\min} r \leq \rho_\ell \leq C_{\text{tr}} \Lambda_{\max} r, \quad 1 \leq \ell \leq N_{\text{tr}},$$

such that the physical image of $Q_{c_{\text{tr}}r}(z_0)$ is contained in $\cup_\ell Q_\ell^{\text{ph}}$, each Q_ℓ^{ph} is contained in the physical image of $Q_r(z_0)$, and

$$\mathcal{C}_u(x_\ell, t_\ell, \rho_\ell) \leq C_{\text{tr}} \mathcal{C}_V(z_0, r) \quad (1 \leq \ell \leq N_{\text{tr}}).$$

Conversely, if a finite physical cylinder cover of the image of $Q_r(z_0)$ with comparable radii satisfies $\mathcal{C}_u(x_\ell, t_\ell, \rho_\ell) \leq \varepsilon$ for every member of the cover, then

$$\mathcal{C}_V(z_0, c_{\text{tr}}r) \leq C_{\text{tr}}N_{\text{tr}}\varepsilon.$$

Proof. The chart is

$$x = X(s) + \Lambda(s)y, \quad dt = \Lambda(s)^2 ds, \quad V(y, s) = \Lambda(s)u(x, t), \quad P(y, s) = \Lambda(s)^2 p(x, t).$$

The bounds on Λ and X imply that the image of a represented parabolic cylinder is trapped between parabolic cylinders whose radii are comparable to $\Lambda(s_0)r$. A standard Besicovitch covering with the bounded eccentricity supplied by $\Lambda_{\min}, \Lambda_{\max}$ gives the finite cover and the bounded overlap constant N_{tr} .

For the velocity term the change of variables gives, on every comparison cylinder,

$$\rho_\ell^{-2} \int_{Q_\ell^{\text{ph}}} |u|^3 dx dt \leq C_{\text{tr}} r^{-2} \int_{Q_r(z_0)} |V|^3 dy ds.$$

For the pressure term we use the invariant seminorm $\inf_{c(s)} \int |P - c(s)|^{3/2}$. Spatial averages are equivalent to this seminorm up to the universal factor $2^{3/2}$, because for every ball B and $q \geq 1$,

$$\int_B |F - F_B|^q \leq 2^q \inf_c \int_B |F - c|^q.$$

The same Jacobian identity therefore gives the displayed physical estimate for the pressure oscillation after subtracting spatial means. Applying the same argument to the inverse chart, whose scale is bounded between $\Lambda_{\max}^{-1}$ and $\Lambda_{\min}^{-1}$ on the comparison cover, gives the converse estimate. No direction of [proposition 4.7](#) is used as a black box here; both estimates come from the exact change of variables and the pressure oscillation seminorm. $\square$

Lemma 5.31 (CKN regularity in comparable repaired gauges). *Let (V, P, a, b) be a represented suitable weak solution on $Q_{2r}(z_0)$ obtained from the repaired-gauge chart of [theorem 4.37](#). Assume that the chart satisfies the comparability conclusion of [lemma 5.29](#) on this cylinder, with constants $\Lambda_{\min}, \Lambda_{\max}$, and that*

$$\|a\|_{L^\infty} + \|b\|_{L^\infty} \leq M$$

there. There are constants

$$\varepsilon_{\text{rg}} = \varepsilon_{\text{rg}}(M, \Lambda_{\min}, \Lambda_{\max}) > 0, \quad c_{\text{rg}} = c_{\text{rg}}(M, \Lambda_{\min}, \Lambda_{\max}) \in (0, 1),$$

such that if

$$r^{-2} \int_{Q_r(z_0)} |V|^3 + r^{-2} \int_{Q_r(z_0)} |P - (P)_{B_r}|^{3/2} \leq \varepsilon_{\text{rg}},$$

then V is bounded on $Q_{c_{\text{rg}}r}(z_0)$, with a bound depending only on $M, \Lambda_{\min}, \Lambda_{\max}, r$, and the preceding displayed quantity.

Proof. Apply the forward direction of [lemma 5.30](#) to $Q_r(z_0)$. Choose ε_{rg} so small that every physical cylinder in the resulting finite comparison cover has CKN quantity below the threshold ε_{CKN} from [theorem 9.35](#). That theorem gives boundedness of the physical velocity on smaller physical cylinders. Pulling those estimates back through the same comparable repaired chart gives the asserted L^∞ bound for V on $Q_{c_{\text{rg}}r}(z_0)$. Thus boundedness is obtained only from epsilon regularity after two-sided chart transfer and small CKN quantity, not from compactness or gauge degeneracy. $\square$

Lemma 5.32 (Scaled local pressure decay). *Let (V, P) be a represented suitable pair on $Q_{2r}(z_0)$, and assume the localized pressure decomposition of [theorem 4.25](#) is available with a cutoff supported in $B_{2r}(y_0)$ and equal to 1 on $B_r(y_0)$. For*

$$C_V(\rho; z_0) := \rho^{-2} \int_{Q_\rho(z_0)} |V|^3, \quad D_P(\rho; z_0) := \rho^{-2} \int_{Q_\rho(z_0)} |P - (P)_{B_\rho}|^{3/2},$$

there is a universal constant C_{dec} such that, whenever $0 < \theta \leq 1/4$,

$$D_P(\theta r; z_0) \leq C_{\text{dec}} \theta^{5/2} D_P(r; z_0) + C_{\text{dec}} \theta^{-2} C_V(2r; z_0).$$

Proof. Write $z_0 = (y_0, s_0)$. For a.e. time $s \in (s_0 - r^2, s_0]$, apply [theorem 4.25](#) on $B_r(y_0)$ with cutoff supported in $B_{2r}(y_0)$:

$$P(\cdot, s) = P_{\text{loc}}(\cdot, s) + H(\cdot, s), \quad -\Delta H(\cdot, s) = 0 \quad \text{in } B_r(y_0).$$

The Calderón–Zygmund part satisfies

$$\int_{B_{\theta r}(y_0)} |P_{\text{loc}}|^{3/2} \leq C \int_{B_{2r}(y_0)} |V|^3.$$

After integrating over the shorter time interval and multiplying by $(\theta r)^{-2}$, this contributes at most $C\theta^{-2}C_V(2r; z_0)$.

For the harmonic part, the interior oscillation estimate for harmonic functions gives, for $q = 3/2$,

$$\int_{B_{\theta r}(y_0)} |H - (H)_{B_{\theta r}}|^q \leq C\theta^{3+q} \int_{B_r(y_0)} |H - (H)_{B_r}|^q.$$

Since $3 + q = 9/2$, scaling the cylinder by $(\theta r)^{-2}$ leaves the factor $\theta^{5/2}$. Moreover

$$\int_{B_r} |H - (H)_{B_r}|^{3/2} \leq C \int_{B_r} |P - (P)_{B_r}|^{3/2} + C \int_{B_r} |P_{\text{loc}}|^{3/2},$$

and the last term is again bounded by $C \int_{B_{2r}} |V|^3$. Absorbing the resulting $C\theta^{5/2}C_V(2r; z_0)$ into $C\theta^{-2}C_V(2r; z_0)$ gives the displayed estimate. Spatial averages may be replaced by the minimizing constants at the cost of a universal factor, as in [lemma 5.30](#). $\square$

Lemma 5.33 (Dyadic persistence of CKN concentration). *Let (V, P) be suitable on a compact cylinder Q_ρ , and let*

$$\mathcal{C}_{V,P}(z, r) := r^{-2} \int_{Q_r(z)} |V|^3 + r^{-2} \int_{Q_r(z)} |P - (P)_{B_r}|^{3/2}.$$

If $z_* \in Q_\rho$ is a Caffarelli–Kohn–Nirenberg concentration point in the sense that

$$\limsup_{r \downarrow 0} \mathcal{C}_{V,P}(z_*, r) \geq \kappa_* > 0,$$

then for every $0 < \theta < 1$ and every sufficiently small outer radius r_0 with $Q_{2r_0}(z_*) \Subset Q_\rho$, there exist $z_m \rightarrow z_*$, integers $L_m \rightarrow \infty$, and dyadic radii $\rho_m = \theta^{L_m} r_0$ such that

$$Q_{\rho_m}(z_m) \Subset Q_{r_0}(z_*), \quad \mathcal{C}_{V,P}(z_m, \rho_m) \geq c_\theta \kappa_*,$$

where $c_\theta > 0$ depends only on θ .

Proof. Choose radii $\sigma_m \downarrow 0$ and points $z_m \rightarrow z_*$ with $Q_{\sigma_m}(z_m) \Subset Q_{r_0}(z_*)$ and $\mathcal{C}_{V,P}(z_m, \sigma_m) \geq \kappa_*/2$. Pick L_m so that

$$\theta^{L_m+1} r_0 < \sigma_m \leq \theta^{L_m} r_0 =: \rho_m.$$

Then $L_m \rightarrow \infty$ and $\rho_m \leq \theta^{-1}\sigma_m$. The velocity term on $Q_{\rho_m}(z_m)$ controls the velocity term on $Q_{\sigma_m}(z_m)$ with loss at most θ^2 , because the normalization changes by $(\sigma_m/\rho_m)^2$. For the pressure term, spatial averages are compared through the minimizing constant:

$$\int_{B_{\sigma_m}} |P - (P)_{B_{\sigma_m}}|^{3/2} \leq 2^{3/2} \int_{B_{\sigma_m}} |P - c|^{3/2} \leq 2^{3/2} \int_{B_{\rho_m}} |P - c|^{3/2}.$$

Taking $c = (P)_{B_{\rho_m}}$, integrating in time, and accounting for the same normalization loss gives

$$\mathcal{C}_{V,P}(z_m, \rho_m) \geq 2^{-5/2} \theta^2 \mathcal{C}_{V,P}(z_m, \sigma_m).$$

Thus the conclusion holds with $c_\theta = 2^{-7/2} \theta^2$. $\square$

Lemma 5.34 (Exterior pressure persistence is a structural exit). *Let a scale-rigid branch have a pressure lower bound on shrinking selected cylinders that is not accounted for by the local source $\partial_i \partial_j (V_i V_j)$ in any compact selected frame after the decomposition of [theorem 4.25](#). Then the branch enters the separated/exterior structural alternative of [definition 5.6 \(iv\)](#), and it does not remain in the residual scale-rigid stratum.*

Proof. The unaccounted part of the pressure is harmonic in the selected compact frame and is generated by sources at positive distance in the corresponding physical variables. If the exterior velocity source has a terminal CKN concentration on one of the annuli separating it from the selected core, then [theorem 9.39](#) gives an exterior concentration branch. In the local variables this is a separated active physical point, hence the separated-point alternative [definition 5.6 \(iv\)](#) and the reduction [theorem 5.21](#) apply.

If no such exterior concentration exists on the separating annuli, then [theorem 9.42](#) gives regular exterior control and [lemma 9.45](#) gives convergence of the exterior pressure contribution modulo functions of time on the selected compact frame. This contradicts the assumed persistent pressure lower bound there. Thus the only non-vanishing exterior-pressure case is already a separated/exterior structural exit. $\square$

Lemma 5.35 (Persistent CKN mass activates an L^3 core or exits structurally). *Let (V_n, P_n) be a compact-cylinder scale-rigid sequence on Q_{2R} with*

$$\sup_n (C_{V_n}(2R) + D_{P_n}(2R)) \leq M_0,$$

and assume the pressure decomposition of [theorem 4.25](#) is available on Q_{2R} . Fix $\kappa_0 > 0$ and choose $\theta \in (0, 1/8)$ so that

$$C_{\text{dec}} \theta^{5/2} \leq \frac{1}{4}.$$

Then there are constants $\eta_{\text{act}} > 0$ and $L_0 \in \mathbb{N}$, depending only on M_0, R, κ_0, θ , the fixed-scale pressure bound below, and the pressure-decomposition constants, with the following property. Suppose $0 < R_0 < R$, $z_n \rightarrow z$, $Q_{2R_0}(z_n) \Subset Q_{2R}$ for all large n , and suppose the fixed outer scale used for the iteration has a uniform local pressure bound

$$D_{P_n}(R_0; z_n) \leq D_0 < \infty.$$

Let $L_n \rightarrow \infty$, $\rho_n := \theta^{L_n} R_0$, and assume

$$\mathcal{C}_n(\rho_n; z_n) := \rho_n^{-2} \int_{Q_{\rho_n}(z_n)} |V_n|^3 + \rho_n^{-2} \int_{Q_{\rho_n}(z_n)} |P_n - (P_n)_{B_{\rho_n}}|^{3/2} \geq \kappa_0$$

for all large n . Then, after passing to a subsequence, either:

(a) for some integers $0 \leq j_n < L_n$, with $s_n := 2\theta^{j_n} R_0$,

$$s_n^{-2} \int_{Q_{s_n}(z_n)} |V_n|^3 \geq \eta_{\text{act}}^3,$$

so the rescaled branch at scale s_n contains an active L^3 core;

(b) the retained pressure lower bound is exterior pressure persistence, hence the branch is assigned to the separated/exterior structural alternative by [lemma 5.34](#).

Proof. Choose L_0 so large that

$$(1/4)^{L_0} D_0 \leq \kappa_0/16,$$

and choose $\eta_{\text{act}} > 0$ so small that

$$C_{\text{dec}} \theta^{-2} \sum_{m=0}^{\infty} (1/4)^m \eta_{\text{act}}^3 \leq \kappa_0/16, \quad 4\theta^{-2} \eta_{\text{act}}^3 \leq \kappa_0/4.$$

For large n , $L_n \geq L_0$. Suppose item (a) fails. Then for each $0 \leq j < L_n$,

$$C_{V_n}(2\theta^j R_0; z_n) < \eta_{\text{act}}^3.$$

Iterating [lemma 5.32](#) from the fixed outer scale R_0 down to $\theta^{L_n} R_0 = \rho_n$ gives

$$D_{P_n}(\rho_n; z_n) \leq (1/4)^{L_n} D_{P_n}(R_0; z_n) + C_{\text{dec}} \theta^{-2} \sum_{j=0}^{L_n-1} (1/4)^{L_n-1-j} C_{V_n}(2\theta^j R_0; z_n) \leq \kappa_0/8.$$

Here the initial pressure is bounded by D_0 because the iteration starts from this fixed compact scale, not from the shrinking target scale. Moreover $Q_{\rho_n}(z_n) \subset Q_{2\theta^{L_n-1} R_0}(z_n)$, so failure of item (a) at $j = L_n - 1$ gives

$$\rho_n^{-2} \int_{Q_{\rho_n}(z_n)} |V_n|^3 \leq 4\theta^{-2} \eta_{\text{act}}^3 \leq \kappa_0/4.$$

Hence $\mathcal{C}_n(\rho_n; z_n) < \kappa_0$, contradicting the assumed lower bound, unless the pressure term used in the lower bound is not the local pressure governed by the compact-frame decomposition. The latter possibility is exactly exterior pressure persistence and is item (b) by [lemma 5.34](#). $\square$

Lemma 5.36 (Activated cores are captured by the finite active ledger). *Let a scale-rigid branch inherit an active profile family in the sense of [definition 5.7](#), with finite critical mass M_* and concentration capture. Suppose [lemma 5.35](#) produces represented scales $s_n > 0$ and centers z_n such that*

$$s_n^{-2} \int_{Q_{s_n}(z_n)} |V_n|^3 \geq \eta_{\text{act}}^3.$$

After lowering the active threshold once, if necessary, to $\eta_^{\text{new}} := \min\{\eta_*, \eta_{\text{act}}/2\}$, exactly one of the following occurs after passing to a subsequence:*

- (a) *the activated core is already represented by an existing active profile or compound profile in the inherited family;*
- (b) *it is absorbed into the selected terminal frame, so it does not create an independent residual obstruction;*
- (c) *it is assigned to the profile-completeness obligation rather than to the scale-rigid residual branch.*

In particular, under the stated concentration-capture hypothesis, pressure-core activation cannot create an unregistered active object inside the residual scale-rigid branch.

Proof. Rescale the branch around (z_n, s_n) inside the represented window; in physical variables the corresponding scale is the selected physical scale multiplied by s_n , hence still tends to zero along the Type II sequence. The displayed lower bound is exactly the active-frame condition in [definition 5.3](#) on Q_1 , with threshold η_{act} . The concentration-capture clause in [definition 5.7](#) says that a compact terminal-window sequence with positive L^3 mass is either seen by the active family or is absorbed into the selected terminal frame; it cannot remain in the perturbative remainder. Hence alternatives (a) and (b) hold unless the frame is genuinely absent from the captured profile family. That last possibility is not a scale-rigid residual mechanism; it is precisely failure of the terminal active-family completeness/capture input and is assigned to that separate profile-completeness obligation. Therefore the scale-rigid refinement below only uses active objects that already belong to the finite captured family, together with objects absorbed into the selected terminal frame. $\square$

Lemma 5.37 (No sub-concentration gives a bounded selected-window limit). *Let (V_n, P_n, a_n, b_n) be a scale-rigid local branch and let (V_*, P_*) be a compact limit supplied by [lemma 5.28](#) on Q_ρ . Then, after shrinking once to a compact subcylinder $Q_{\rho'} \Subset Q_\rho$, exactly one of the following alternatives holds:*

- (a) *(V_*, P_*) has no Caffarelli–Kohn–Nirenberg concentration point in $Q_{\rho'}$. Then $V_* \in L^\infty(Q_{\rho''})$ for every $Q_{\rho''} \Subset Q_{\rho'}$, and the selected subsequence has a nonzero bounded selected-window limit on some such $Q_{\rho''}$.*
- (b) *(V_*, P_*) has a Caffarelli–Kohn–Nirenberg concentration point in $Q_{\rho'}$. Then the original branch either contains a further captured active L^3 concentration core after rescaling around that point, the pressure mass is exterior and the branch is already in the separated/exterior structural alternative, or the core is assigned to the separate profile-completeness obligation. According to its relative scale-center position, every captured core enters one of the already named multicore, same-point cascade, or separated/exterior structural alternatives, rather than remaining a residual scale-rigid branch.*

Proof. Let ε_{rg} be the represented-gauge regularity threshold from [lemma 5.31](#). Before using that threshold, apply [lemma 5.29](#) on each retained compact subcylinder. The chart stability in [proposition 4.34](#), together with the L^∞ modulation bound (R3), gives a limiting chart (Λ_*, X_*) with $a_* = \Lambda'_*/\Lambda_*$ and $b_* = X'_*/\Lambda_*$ on the retained subwindow. Thus the compactness of the subcylinder and (R3) give a uniform lower-upper scale comparison after one shrink, so the constants in [lemma 5.31](#) depend only on the retained chart geometry and not on any unproved boundedness of V_n or V_* .

Assume first that there is no CKN concentration point of (V_*, P_*) in $Q_{\rho'}$. For every point $z \in Q_{\rho'}$ there is a parabolic cylinder $Q_{r_z}(z) \Subset Q_\rho$ such that the scale-invariant velocity-pressure quantity of (V_*, P_*) on $Q_{r_z}(z)$ is below ε_{rg} . A finite subcover of each $Q_{\rho''} \Subset Q_{\rho'}$, followed by [lemma 5.31](#) on the cover, gives $V_* \in L^\infty(Q_{\rho''})$.

It remains to check that the bounded limit is nonzero. If $V_* \equiv 0$ on the retained inner cylinder, then the strong L^3 convergence from [lemma 5.28](#) and the zero-velocity argument of [lemma 3.16](#) imply

$$C_{V_n}(r_*) \rightarrow 0,$$

contradicting the retained lower bound (R2) in [definition 5.19](#). Hence $V_* \not\equiv 0$ on some smaller compact cylinder. Since $V_n \rightarrow V_*$ strongly in L^3 there, the selected subsequence has a nonzero bounded selected-window limit.

Now suppose a CKN concentration point z_* remains in $Q_{\rho'}$. Then there is $\kappa_* > 0$ such that

$$\limsup_{r \downarrow 0} \mathcal{C}_{V_*, P_*}(z_*, r) \geq \kappa_*.$$

Choose $R_0 > 0$ with $Q_{2R_0}(z_*) \subseteq Q_\rho$, and choose $\theta \in (0, 1/8)$ satisfying the pressure-decay smallness condition in [lemma 5.35](#). By [lemma 5.33](#), there are $z_m \rightarrow z_*$, $L_m \rightarrow \infty$, and $\rho_m = \theta^{L_m} R_0$ such that

$$\mathcal{C}_{V_*, P_*}(z_m, \rho_m) \geq c_\theta \kappa_*.$$

For each fixed m , the strong L^3 convergence of V_n and weak lower semicontinuity of the pressure oscillation seminorm in $L^{3/2}/\{\text{spatial constants}\}$ imply, along a diagonal subsequence,

$$\mathcal{C}_{V_n, P_n}(z_m, \rho_m) \geq \frac{1}{2} c_\theta \kappa_*.$$

Relabeling the diagonal sequence, apply [lemma 5.35](#) with $\kappa_0 = \frac{1}{2} c_\theta \kappa_*$ on the compact cylinder where (R1) holds. The additional fixed-scale pressure bound required by that lemma holds because R_0 is fixed, $Q_{2R_0}(z_m) \subseteq Q_\rho$ for all large m , and the mean-normalized pressures are uniformly bounded in $L^{3/2}(Q_\rho)$ by [lemma 5.28](#); comparison of spatial averages on the fixed nested balls gives $D_{P_n}(R_0; z_m) \leq D_0(R_0, \rho, M)$. If the persistent CKN mass is produced by an interior velocity core, then [lemma 5.36](#) either identifies that core with the finite captured active ledger or assigns it to the separate profile-completeness obligation, after the one-time lowering of the active threshold. If the mass is carried by harmonic pressure generated outside the selected compact frame, [lemma 5.34](#) assigns the branch to the separated/exterior structural alternative, so it exits the residual scale-rigid stratum.

The relative position of this new core with respect to the retained scale is exhausted by [lemma 5.10](#): it is comparable at the same point, strictly separated in scale at the same point, separated in physical center, or locally invisible in the retained frame. The invisible case cannot carry the retained positive concentration by the capture clause in [definition 5.7](#). The remaining cases are precisely the multicore, same-point cascade, and separated/exterior structural alternatives already named in the manuscript. Thus the branch does not remain in the residual scale-rigid stratum in this case. $\square$

Theorem 5.38 (Single-core reduction compatibility). *Let (V_n, P_n, a_n, b_n) be a scale-rigid local branch in the sense of [definition 5.19](#), and assume the branch reduces, after the rigid-frame limit of [lemma 5.28](#), to a single retained nondegenerate core on a fixed inner cylinder $Q_{r_*} \subseteq Q_\rho \subseteq Q_{R_*}$. Then, after passing to a subsequence, one of the following alternatives holds:*

- (i) *the branch has a nonzero bounded selected-window limit;*
- (ii) *the original selected sequence has vanishing localized dissipation on a retained compact cylinder, hence is a compact single-core vanishing-dissipation branch in the sense of [definition 3.4](#) and is excluded by [theorem 3.18](#);*
- (iii) *a further active concentration core is produced, so the branch leaves the single-core reduction and enters a multicore, cascade, or separated/exterior structural alternative.*

Proof. If the original selected sequence has vanishing localized dissipation on some $Q_{\rho'}, Q_{r_*} \subseteq Q_{\rho'} \subseteq Q_{R_*}$, then we first verify [definition 3.4](#). Suitability and the represented equation are supplied by [theorem 4.37](#). Compact-cylinder bounds, positive selected-core CKN concentration, and bounded modulation are exactly (R1), (R2), and (R3) of [definition 5.19](#). The nondegenerate scale-center selection is the hypothesis of this single-core case. The vanishing localized dissipation is the present hypothesis on the original selected sequence, not on a residual after subtracting a limit. Thus the branch satisfies every clause of [definition 3.4](#), and [theorem 3.18](#) excludes it. This is item (ii).

The compact limit on a slightly smaller cylinder is covered by [lemma 5.37](#). If the no-subconcentration alternative holds, item (i) follows. If a further active concentration core is produced, item (iii) follows. $\square$

Lemma 5.39 (Degenerate Jacobian resolution). *Let (V_n, P_n, a_n, b_n) be a represented sequence reaching the gauge-degenerate alternative [definition 5.6 \(v\)](#), i.e. along the subsequence the relative scale-center Jacobian of the selected compound or terminal frame degenerates. Then, after passing to a further subsequence and reselecting the active center in the sense of [definition 5.7](#) on the same physical point, exactly one of the following four resolutions occurs:*

- (D1) *the reselection produces an additional active core or active compound carrying positive L^3 mass at the fixed active threshold; hence the branch enters the multicore, same-point cascade, or separated/exterior alternative [definition 5.6 \(ii\)–\(iv\)](#) and is treated by [theorem 5.11](#);*
- (D2) *the reselection collapses the compound to a single nondegenerate core and the branch enters the single-core compatibility dichotomy governed by [theorem 5.38](#);*
- (D3) *the reselection produces a perturbative compound that is below the critical small-data threshold of Step 2 of [theorem 5.20](#), hence is absorbed into the perturbative remainder and treated by [theorem 5.18](#);*
- (D4) *the reselection still has degenerate scale-center Jacobian on every further subsequence and has no further active subconcentration; in this case the no-subconcentration alternative of [lemma 5.37](#) gives a nontrivial bounded selected-window limit on a compact cylinder.*

Proof. By the gauge-degeneracy alternative [definition 5.6 \(v\)](#), the localized scale or centering selection of the selected compound or terminal frame is non-unique along the subsequence. Apply the parameter partition lemma [lemma 5.10](#) to the selected compound together with any candidate alternate selection produced by the gauge degeneracy. Three mutually exclusive structural outcomes are possible.

Case A: alternate selection is comparable to the original at the same physical point. Then both selections lie in the same scale comparability class of Step 1 of [theorem 5.20](#). Step 2 of that theorem groups them into a single compound profile. If the compound is below the small-data threshold, we are in case (D3); otherwise the compound is above threshold and Step 3 of the theorem applies. If the compound is then nondegenerate, we are in case (D2); if it is still degenerate, the gauge degeneracy persists on the reselected compound, and we proceed to Case D below.

Case B: alternate selection is at a strictly different scale at the same physical point. By [lemma 5.14](#), the strictly outer profile vanishes locally in the innermost frame; absorbing it into the perturbative remainder lands us in case (D3). If instead the reselection produces a strictly inner core not previously detected, then either compact-cylinder suitable control fails, in which case [lemma 5.15](#) produces the new CKN core, or compact control persists and [lemma 5.35](#) converts the retained CKN mass into an active L^3 core unless the branch has already exited to the separated/exterior pressure alternative. These outcomes are case (D1).

Case C: alternate selection is at a separated physical point. By [theorem 5.21](#), the new selection is locally invisible in the original frame, hence absorbed into the perturbative remainder; this is case (D3). If the new selection retains positive CKN concentration in its own frame, then [lemma 5.35](#) makes it an additional separated active core, or else assigns it to the separated/exterior pressure alternative. Both outcomes are case (D1) via [definition 5.6 \(iv\)](#).

Case D: persistent degeneracy. Suppose every reselection on every further subsequence still has degenerate scale-center Jacobian. Apply [lemma 5.28](#) and then [lemma 5.37](#). If the compact limit has a CKN subconcentration point, that lemma produces an additional active core; this is case (D1). If no such subconcentration remains, the same lemma gives a nonzero bounded selected-window limit; this is case (D4). No L^∞ bound is inferred from the Jacobian degeneracy itself. $\square$

Theorem 5.40 (Scale-rigid stratum decomposition). *Every scale-rigid local branch in the sense of [definition 5.19](#) enters one of the following named local strata after passing to a subsequence:*

- (i) *the multicore alternative [definition 5.6](#) (ii)–(iv), governed by [theorem 5.11](#);*
- (ii) *the compact single-core stratum, governed by [theorem 5.38](#);*
- (iii) *the perturbative remainder of [definition 5.7](#), governed by [theorem 5.18](#);*
- (iv) *the bounded selected-window limit alternative [definition 3.1](#), which by [theorem 5.27](#) is incompatible with [definition 3.2](#).*

Proof. By [definition 5.19](#) (R4), the branch enters the scale-rigid stratum either via a degenerate relative scale-center Jacobian or as an innermost cascade core after outer profiles have been absorbed perturbatively.

Degenerate Jacobian case. Apply [lemma 5.39](#). Cases (D1)–(D4) of that lemma give items (i)–(iv) respectively.

Innermost cascade case. By Step 4 of [theorem 5.20](#), after subtracting the outer-profile sum $S_n \rightarrow 0$ in $L_I^\infty L^3(B_{2R})$ and absorbing constant drifts into the translation modulation, the residual is governed by [theorem 5.18](#). A nondegenerate retained component is item (ii) via [theorem 5.38](#), and a component with no retained obstruction is item (iii). If the retained selection is degenerate, its recorded Jacobian is an unresolved (L3) entry in the finite ledger below. Apply [lemma 5.39](#) directly to that entry; its four outputs are items (i)–(iv). Thus this case creates no back edge to the same-point construction. $\square$

Lemma 5.41 (Finite refinement of scale-rigid structural exits). *Let ω be a scale-rigid local branch produced by [theorem 5.20](#) or [theorem 5.21](#), with the inherited active family of finite critical mass and with concentration capture in the sense of [definition 5.7](#). Resolve every structural exit of [theorem 5.40](#) by applying only [theorem 5.11](#), [theorem 5.20](#), [theorem 5.21](#), and [theorem 5.18](#), never [theorem 5.44](#). Then the refinement terminates after finitely many ledger-decreasing refinements. At the terminal residual branch one has either a nonzero bounded selected-window limit or a compact single-core vanishing-dissipation branch excluded by [theorem 3.18](#).*

Proof. After the active threshold has been lowered once, every core produced by [lemma 5.35](#) is handled by [lemma 5.36](#): it is already in the captured inherited family, is absorbed into the terminal frame, or exits to the separate profile-completeness obligation. The residual scale-rigid ledger therefore uses only the captured active family. Write η_*^{new} for the lowered active threshold. By [lemma 5.9](#), the resulting family contains at most

$$N_* := \lfloor M_*(\eta_*^{\text{new}})^{-3} \rfloor$$

active profiles. Group same-point comparable-scale profiles into compound profiles as in [definition 5.8](#). There are therefore only finitely many active objects and only finitely many unordered pair relations among them.

Define the refinement ledger to consist of:

- (L1) every active object not yet selected, grouped, absorbed, or declared terminal;
- (L2) every pair relation not yet classified as comparable same-point, strict same-point cascade, separated point, or locally invisible;
- (L3) every degenerate scale-center selection not yet resolved by [lemma 5.39](#);
- (L4) every perturbative remainder not yet removed by [theorem 5.18](#).

The ledger is finite because it is built from a finite active family and finitely many compounds. Order ledgers lexicographically by the number of unresolved entries in (L1), (L2), (L3), and (L4). Explicitly, set

$$\mathbf{r}(\omega) := (N_{\text{act}}(\omega), N_{\text{pair}}(\omega), N_{\text{deg}}(\omega), N_{\text{rem}}(\omega)) \in \mathbb{N}^4,$$

where the four entries are the numbers of unresolved ledger entries in (L1)–(L4), and order $\mathbb{N}^4$ lexicographically. The same-point and separated-point maps used below are construction edges: they group, classify, or assign one ledger entry. They are not closure edges. The only scale-rigid closure edge is applied after this rank has been exhausted.

At each nonterminal step, no multibubble exclusion theorem is used. If [theorem 5.40](#) gives a multicore, same-point cascade, or separated/exterior structural exit, apply [theorem 5.11](#) to select the corresponding active object. Comparable same-point objects are grouped into a compound profile; below-threshold compounds are removed by the small-data step in [theorem 5.20](#), while above-threshold compounds count as active objects. Strict same-point cascades are reduced by [theorem 5.20](#); separated physical centers are reduced by [theorem 5.21](#). Perturbative exterior and remainder terms are discarded only through [theorem 5.18](#), whose output preserves the selected positive concentration and produces no new active object.

Each move strictly decreases the ledger. Adding a previously unseen active core consumes one entry of (L1), and this can happen at most N_* times. Grouping comparable same-point objects, ordering a same-point cascade, separating two physical centers, or declaring one object locally invisible resolves a previously unresolved pair and consumes one entry of (L2). Applying [lemma 5.39](#) resolves the selected degenerate frame or produces an active structural exit already counted in (L1)–(L2), so it consumes one entry of (L3). Applying [theorem 5.18](#) to a perturbative remainder consumes one entry of (L4). Once an entry is resolved it is not reintroduced: later passes use the grouped compound, the selected innermost representative, the separated-frame assignment, or the discarded perturbative remainder. Hence no cycle is possible.

After the ledger is exhausted, item (i) of [theorem 5.40](#) cannot occur again, because it would create either a new active object or an unresolved active relation, contradicting exhaustion of (L1)–(L2). Item (iii) cannot carry the retained positive concentration by the capture clause in [definition 5.7](#); a purely perturbative remainder has already been absorbed through (L4). Item (iv) is exactly the nonzero bounded selected-window-limit alternative.

It remains only to handle item (ii). Apply [theorem 5.38](#). Its bounded-limit alternative gives the first terminal outcome. Its further-active-core alternative would add a new ledger entry, impossible after ledger exhaustion. Its vanishing-localized-dissipation alternative verifies [definition 3.4](#) for the original selected sequence and is excluded by [theorem 3.18](#). These are precisely the two terminal residual outcomes claimed. $\square$

Proposition 5.42 (Scale-rigidity exclusion). *Every scale-rigid local branch produced by [theorem 5.20](#) or [theorem 5.21](#), including a branch with degenerate relative scale-center Jacobian, either has a nonzero bounded selected-window limit on a compact cylinder (and so fails [definition 3.2](#) by [theorem 5.27](#)) or reduces to the compact single-core branch governed by [theorem 3.18](#) (equivalently by the local proposition [proposition 5.22](#)). Consequently no such branch is compatible with the local Type II condition.*

Proof. Let ω be a scale-rigid local branch carrying the input contract (R1)–(R4) constructed by [theorem 5.20](#) or [theorem 5.21](#). [Theorem 5.40](#) places ω into one of items (i)–(iv).

If item (i) occurs, the branch enters a named multicore, same-point cascade, or separated/exterior structural exit. Resolve that exit by the finite refinement procedure of [lemma 5.41](#). This procedure uses only [theorem 5.11](#), [theorem 5.20](#), [theorem 5.21](#), and [theorem 5.18](#); it does not invoke [theorem 5.44](#) or the present discharge proposition. The refinement terminates at a bounded selected-window limit or at the compact single-core vanishing-dissipation branch.

In item (ii), apply [theorem 5.38](#). Its bounded-limit case is the first alternative in the proposition. Its vanishing-localized-dissipation case verifies [definition 3.4](#) for the original selected sequence and is excluded by [theorem 3.18](#); this is the “reduces to the single-core branch” alternative. Its further-active-core case is resolved by the same finite refinement lemma used for item (i).

In item (iii), the branch is the perturbative remainder; by [theorem 5.18](#) it produces no independent multibubble obstruction. The retained positive concentration is carried by the selected component, not by the perturbative remainder, so the residual selected component is reduced to item (ii) or item (iv) by [lemma 5.41](#).

In item (iv), by [theorem 5.27](#) the branch is not a local Type II sequence in the sense of [definition 3.2](#). This is the “nonzero bounded selected-window limit” alternative.

In every case the branch either has a nonzero bounded selected-window limit or reduces to the single-core branch. All uses of the same-point and separated-point theorems in this proof occur inside the strictly decreasing rank of [lemma 5.41](#); neither construction theorem uses this proposition. $\square$

Proposition 5.43 (Terminal admissibility in active frames). *Assume finite local critical norm, terminal windowed state-space completeness in the sense of [theorem 9.16](#), local perturbative residual stability, the exterior annulus alternative [corollary 9.48](#), the repaired-gauge representation of [theorem 4.37](#), and local Caccioppoli regularity. Then every terminal family produced by the same-point and separated-point reductions is admissible on each compact selected cylinder in the sense of [definition 5.16](#).*

Proof. The terminal nonlinear local-state decoupling theorem, [theorem 9.27](#), applies under the stated hypotheses and produces, on every compact terminal cylinder, the conclusions of [definition 9.13](#). We verify each clause of [definition 5.16](#).

Items (i) and (ii) are the active-family enumeration of [definition 5.7](#) together with the same-point and separated-point grouping of [lemmas 5.13](#) and [5.14](#); they are properties of the active family that survive the terminal-frame construction by definition of the terminal active class.

Item (iii) is the perturbative remainder hypothesis of [definition 5.16](#). For the admissible active families constructed here, the velocity decoupling [lemma 9.18](#), the linear residual decoupling [lemma 9.24](#), and the pressure decoupling [lemma 9.23](#) supply the required $L_I^\infty L^3$ vanishing of the rescaled remainder r_n together with the existence of (π_n^r, f_n) having the stated $L_I^1 H^{-1}(B_R)$ decay. This pointwise critical hypothesis is not derived from the later spacetime compactness package.

Item (iv) is the companion $L_I^\infty L^3(B_{2R})$ hypothesis on the selected component in [definition 5.16](#). On the active families produced in [lemma 9.26](#) it is read from the retained local-state bounds; if a physical branch fails this input on the chosen compact cylinder, then the branch is routed by [lemma 10.78](#) rather than treated by the perturbative terminal-frame theorem.

Item (v) is the bounded-modulation hypothesis, supplied by the repaired-gauge representation [theorem 4.37](#) and [theorem 4.36](#).

Item (vi) is the local pressure-decomposition stability of [lemma 9.23](#), whose proof relies only on the Calderón–Zygmund pressure-tail decoupling [lemma 9.12](#) together with the kernel boundedness on $L^{3/2}$ provided by the classical Calderón–Zygmund theorem [[CZ52](#), [Ste70](#)].

Item (vii) follows from the Leibniz formula

$$[\Delta, \chi_R]u = (\Delta \chi_R)u + 2\nabla \chi_R \cdot \nabla u$$

and the identity

$$[\nabla \chi_R, \nabla \cdot](u \otimes u) = (\nabla \chi_R)(u \otimes u).$$

These are combined with the dual estimate

$$\|fu\|_{H^{-1}(B_R)} \leq C_R \|f\|_{C^1(\overline{B_R})} \|u\|_{L^{6/5}(B_R)}$$

and the embedding $L^3(B_R) \hookrightarrow L^{6/5}(B_R)$ on bounded balls. Indeed, for $\varphi \in H_0^1(B_R)$,

$$\langle \nabla \chi_R \cdot \nabla u, \varphi \rangle = -\langle u, (\Delta \chi_R)\varphi + \nabla \chi_R \cdot \nabla \varphi \rangle,$$

giving $\|[\Delta, \chi_R]u\|_{H^{-1}(B_R)} \leq C_R \|u\|_{L^3(B_{2R})}$. The pointwise bound $|(\nabla \chi_R)(u \otimes u)| \leq \|\nabla \chi_R\|_\infty |u|^2$ gives $\|(\nabla \chi_R)(u \otimes u)\|_{L^{3/2}(B_R)} \leq C_R \|u\|_{L^3(B_{2R})}^2$, and $L^{3/2} \hookrightarrow H^{-1}$ on B_R supplies the displayed bound. The combined estimate

$$\|[\Delta, \chi_R]u\|_{H^{-1}(B_R)} + \|[\nabla \chi_R, \nabla \cdot](u \otimes u)\|_{H^{-1}(B_R)} \leq C_R (\|u\|_{L^3(B_{2R})} + \|u\|_{L^3(B_{2R})}^2) \leq C'_R \|u\|_{L^3(B_{2R})}^2$$

holds whenever $\|u\|_{L^3(B_{2R})}$ is bounded below by a fixed constant, which is the case in the active selected frame. The same estimate applies to the S_n -quadratic commutator terms thanks to the uniform $L_I^\infty L^3(B_{2R})$ control of U_n and the $L_I^1 L^{3/2}(B_R)$ control of the cross stresses produced by [lemma 9.20](#).

All items in [definition 5.16](#) are therefore satisfied, which is the asserted admissibility. $\square$

Theorem 5.44 (Local multibubble and cascade exclusion). *Use the local single-core exclusion [proposition 5.22](#), the discharged scale-rigidity closure [proposition 5.42](#), and the terminal admissibility in active frames [proposition 5.43](#). Let a positive local Type II branch admit an active family satisfying the hypotheses of [theorem 5.11](#). Then no local multibubble, cascade, or gauge-degenerate Type II case can occur in this class.*

Proof. Let a local Type II branch be in the multibubble/cascade class of [definition 5.6](#). Apply [theorem 5.11](#). The positive local concentration is carried by an active frame or compound local state, and one of the alternatives listed there occurs.

Case 1: compactness-failure reduction. Compactness failure yields, after changing radius by a fixed factor and rescaling as in [lemma 5.15](#), a bounded-cylinder concentration branch. Relabel this branch as an active core. After passing to a subsequence, the new core and every previously retained active core have one of three relations: comparable scale at the same physical point, strictly separated scale at the same physical point, or separated physical center. Comparable same-point cores are grouped into compound local states. A compound state with no detecting compact cylinder is removed by local perturbative residual stability; an active compound state is treated as one selected core in its active frame. If the scale-center selection inside such a compound core is degenerate, the branch is gauge-degenerate and belongs to the non-uniqueness alternative in [definition 5.6](#), rather than to a separate core dynamics.

Case 1a: gauge degeneracy. If the scale-center selection of the selected compound or terminal frame is degenerate, then the branch is in the gauge-degenerate alternative of [definition 5.6](#). By the repaired-gauge representation from [Section 4](#), nondegeneracy is the hypothesis needed to enter the single-core branch. Persistent degeneracy therefore belongs to the scale-rigid local row eventually discharged by [proposition 5.42](#); if the degeneracy disappears after passing to a subsequence, the branch enters the single-core or terminal-frame alternatives below.

Case 1b: same-point comparable scales. Same-point comparable-scale active frames are grouped into a compound local state by the local state-space selection rule. If the compound state has no detecting compact cylinder, it is part of the perturbative residual. If it has the local detection threshold and its scale-center selection is nondegenerate, it is one selected active core and enters the single-core branch. If the corresponding scale-center selection is degenerate, it is covered by Case 1a. Thus comparable same-point states do not produce an additional multibubble mechanism.

Case 2: same-point cascading. Strict same-point scale cascades are treated by [theorem 5.20](#). In the innermost frame, all outer scales are locally invisible after constant drifts are absorbed. The branch therefore reduces to one of the three terminal alternatives

- (A) single-core, (B) scale-rigid, (C) terminal perturbative.

Case 3: separated-point interactions. For separated physical points, [theorem 5.21](#) gives the same three terminal alternatives as in Case 2:

(A) single-core, (B) scale-rigid, (C) terminal perturbative.

The single-core alternative is excluded by [proposition 5.22](#). The scale-rigid alternative is excluded by [proposition 5.42](#), which includes the scale-collapse and autonomous-limit exclusions in [Section 7](#) in this setting. The terminal perturbative alternative gives no independent multibubble obstruction: [theorem 5.18](#) removes all exterior profiles and remainders from the selected active equation in $L^1 H_{\text{loc}}^{-1}$, preserves the positive finite-mass selected branch, and leaves only the single-core or scale-rigid possibilities already excluded. Consequently every alternative in [definition 5.6](#) is incompatible with the local Type II regime. $\square$

Corollary 5.45 (Removal of the multibubble alternative). *Under the hypotheses of [theorem 5.44](#), every failure of the single nondegenerate compact core alternative is excluded.*

Proof. By [definition 5.6](#), failure of the single nondegenerate compact core alternative means that at least one of the following occurs: the compact-cylinder suitable estimates are unbounded; positive local CKN mass splits between active cores; a strict same-point scale cascade occurs; separated physical points remain active in distinct local frames; or the localized scale or center selection is degenerate. The compactness-failure case produces an additional core by [lemma 5.15](#). Same-point and separated cases are reduced by [theorems 5.20](#) and [5.21](#). [Theorem 5.44](#) excludes the remaining cases. Hence no failure of the single nondegenerate compact core alternative remains. $\square$

6. LOCAL CACCIOPOLI AND ROUGH-CORE EXCLUSION

6.1. Analytic inputs and notation. The renormalized local energy inequality from the preceding renormalized-chart analysis is used as an analytic input. The renormalized suitable structure and pressure reconstruction are not rederived in this section. The estimates below are local implications under the displayed compact-cylinder conditions; the terminal assembly verifies or routes these conditions by the ordered local state-space analysis.

Proposition 6.1 (Renormalized local energy inequality). *Let (V, P, a, b) be defined on a compact renormalized cylinder $Q_{2R,J} := B_{2R} \times J$, where $J \subset \mathbb{R}$ is a bounded interval. Assume the canonical represented equation (14), with Y denoted by y , and assume $\nu > 0$ and $a, b \in L^\infty(J)$. For every nonnegative $\phi \in C_c^\infty(Q_{2R,J})$ and every $\tau_1 < \tau_2$ in J ,*

$$\begin{aligned}
(7) \quad & \int_{B_{2R}} \frac{|V(\cdot, \tau_2)|^2}{2} \phi(\cdot, \tau_2) - \int_{B_{2R}} \frac{|V(\cdot, \tau_1)|^2}{2} \phi(\cdot, \tau_1) \\
& + \nu \int_{\tau_1}^{\tau_2} \int_{B_{2R}} \phi |\nabla V|^2 \leq \int_{\tau_1}^{\tau_2} \int_{B_{2R}} \frac{|V|^2}{2} (\partial_\tau \phi + \nu \Delta \phi) \\
& + \int_{\tau_1}^{\tau_2} \int_{B_{2R}} \frac{|V|^2}{2} V \cdot \nabla \phi + \int_{\tau_1}^{\tau_2} \int_{B_{2R}} P V \cdot \nabla \phi \\
& + \int_{\tau_1}^{\tau_2} \int_{B_{2R}} a \frac{|V|^2}{2} (-\phi - y \cdot \nabla \phi) \\
& - \int_{\tau_1}^{\tau_2} \int_{B_{2R}} \frac{|V|^2}{2} b \cdot \nabla \phi.
\end{aligned}$$

The pressure is understood modulo functions of time, and the mean-normalized pressure on compact balls belongs to the local class used below.

Proof. Proposition 4.33 transports the physical local energy inequality of Caffarelli–Kohn–Nirenberg [CKN82] into the renormalized variables (y, τ) and yields the inequality stated above with test function of the form ϕ^2 , where $\phi \in C_c^\infty(Q_R^*)$ is nonnegative. Setting $\Phi := \phi^2$ and noting that every nonnegative $\Phi \in C_c^\infty(Q_{2R,J})$ is the limit, in $C_c^k(Q_{2R,J})$ for every k , of squares of nonnegative smooth compactly supported functions $\phi_\varepsilon := (\Phi + \varepsilon^2)^{1/2} - \varepsilon$, the inequality extends by approximation to all such Φ ; this is the standard reformulation used in the theory of suitable weak solutions [CKN82, Lin98]. The compact-window representation theorem theorem 4.37 provides the cylinder $Q_{2R,J} = B_{2R} \times J$ and the divergence-free renormalized equation (14), while theorem 4.25 supplies the localized pressure decomposition with the mean-normalized pressure on compact balls. Identifying the modulation coefficients $a = \tilde{a}$ and $b = \tilde{b}$ of proposition 4.33 with those of (14) and grouping the renormalized error terms exactly as in (7) yields the stated inequality. $\square$

Throughout, for a space–time set $E \subset \mathbb{R}^3 \times \mathbb{R}$, write

$$\int_E f := \int_E f(y, \tau) dy d\tau.$$

For $r > 0$ and a bounded interval $J \subset \mathbb{R}$, define

$$Q_{r,J} := B_r \times J, \quad |Q_{r,J}| := |B_r| |J|.$$

For a ball $B \subset \mathbb{R}^3$, denote the spatial mean of pressure by

$$(P)_B(\tau) := \frac{1}{|B|} \int_B P(y, \tau) dy.$$

The local quantities are

$$\begin{aligned}\mathcal{A}_J(r) &:= \operatorname{ess\,sup}_{\tau \in J} \int_{B_r} |V(y, \tau)|^2 dy, \\ \mathcal{H}_J(r) &:= \int_J \int_{B_r} |\nabla V|^2 dy d\tau, \\ \mathcal{C}_J(r) &:= \int_J \int_{B_r} |V|^3 dy d\tau, \quad \mathcal{D}_J(r) := \int_J \int_{B_r} |P - (P)_{B_r}(\tau)|^{3/2} dy d\tau.\end{aligned}$$

Lemma 6.2 (Finite-measure conversion). *For every bounded interval J and every $r > 0$,*

$$\int_J \int_{B_r} |V|^2 \leq |Q_{r,J}|^{1/3} \mathcal{C}_J(r)^{2/3} \leq |Q_{r,J}|^{1/3} (1 + \mathcal{C}_J(r)).$$

Proof. Hölder's inequality on $Q_{r,J}$ gives

$$\int_J \int_{B_r} |V|^2 \leq |Q_{r,J}|^{1/3} \left(\int_J \int_{B_r} |V|^3 \right)^{2/3}.$$

The second inequality follows from $s^{2/3} \leq 1 + s$ for $s \geq 0$. $\square$

Lemma 6.3 (Comparison of pressure normalizations). *Let $1 \leq p < \infty$ and $0 < r \leq R$. For almost every τ such that $P(\tau) \in L^p(B_R)$,*

$$\|P(\tau) - (P)_{B_r}(\tau)\|_{L^p(B_r)} \leq 2\|P(\tau) - (P)_{B_R}(\tau)\|_{L^p(B_R)}.$$

Consequently,

$$\mathcal{D}_J(r) \leq 2^{3/2} \mathcal{D}_J(R)$$

whenever the right-hand side is finite.

Proof. Set $c = (P)_{B_R}(\tau)$. Then

$$\|P - (P)_{B_r}\|_{L^p(B_r)} \leq \|P - c\|_{L^p(B_r)} + |B_r|^{1/p} |(P)_{B_r} - c|.$$

Since

$$|(P)_{B_r} - c| \leq |B_r|^{-1} \int_{B_r} |P - c| \leq |B_r|^{-1/p} \|P - c\|_{L^p(B_r)},$$

the displayed pointwise estimate follows. Raising to the power $3/2$ and integrating in time gives the last assertion. $\square$

6.2. Compact-cylinder Caccioppoli estimate.

Lemma 6.4 (Buffered cutoff estimate from the local energy inequality). *Let $R > 0$, let $J \subset \mathbb{R}$ be a bounded open interval, and let $J' \Subset J$ be a compact subinterval. Assume [proposition 6.1](#) on $Q_{2R,J}$ and assume*

$$\mathcal{C}_J(2R) + \mathcal{D}_J(2R) < \infty.$$

Then there is a constant

$$C = C(R, |J|, \nu, \|a\|_{L^\infty(J)}, \|b\|_{L^\infty(J)}, \operatorname{dist}(J', \partial J))$$

such that

$$\mathcal{A}_{J'}(R) + \mathcal{H}_{J'}(R) \leq C \left(1 + \mathcal{C}_J(2R) + \mathcal{D}_J(2R) \right).$$

The estimate is outer-to-inner: all right-hand quantities are measured on $B_{2R} \times J$, while the conclusion is on $B_R \times J'$.

Proof. Step 1: localized cut-off and the two energy inequalities. Set

$$\delta := \text{dist}(J', \partial J) > 0.$$

Choose $\eta \in C_c^\infty(J)$ and $\zeta \in C_c^\infty(B_{2R})$ with

$$0 \leq \eta \leq 1, \quad 0 \leq \zeta \leq 1, \quad \eta \equiv 1 \text{ on } J', \quad \zeta \equiv 1 \text{ on } B_R,$$

and

$$\|\eta'\|_{L^\infty(J)} \leq C_\eta \delta^{-1}, \quad \|\nabla \zeta\|_{L^\infty(B_{2R})} \leq C_\zeta R^{-1}, \quad \|\Delta \zeta\|_{L^\infty(B_{2R})} \leq C_\zeta R^{-2}.$$

Set

$$\phi(y, \tau) := \eta(\tau) \zeta(y)^2.$$

Then $\phi \in C_c^\infty(Q_{2R,J})$, $0 \leq \phi \leq 1$, and

$$\phi \equiv 1 \text{ on } B_R \times J'.$$

Choose $\tau_- < \inf J'$ and $\tau_+ > \sup J'$ in J with $\eta(\tau_-) = \eta(\tau_+) = 0$. Applying (7) on (τ_-, τ_+) gives

$$\nu \int_{\tau_-}^{\tau_+} \int_{B_{2R}} \phi |\nabla V|^2 \leq |R_\tau| + |R_\Delta| + |R_{\text{conv}}| + |R_{\text{pres}}| + |R_{\text{mod}}|,$$

where

$$\begin{aligned} R_\tau &:= \int_J \int_{B_{2R}} \frac{|V|^2}{2} \partial_\tau \phi, & R_\Delta &:= \nu \int_J \int_{B_{2R}} \frac{|V|^2}{2} \Delta \phi, \\ R_{\text{conv}} &:= \int_J \int_{B_{2R}} \left(\frac{|V|^2}{2} V \right) \cdot \nabla \phi, & R_{\text{pres}} &:= \int_J \int_{B_{2R}} (P - (P)_{B_{2R}}(\tau)) V \cdot \nabla \phi, \\ R_{\text{mod}} &:= \int_J \int_{B_{2R}} a \frac{|V|^2}{2} (-\phi - y \cdot \nabla \phi) - \int_J \int_{B_{2R}} \frac{|V|^2}{2} b \cdot \nabla \phi. \end{aligned}$$

For almost every $\sigma \in J'$, applying (7) on (τ_-, σ) gives the same right-hand estimates and leaves the terminal term $\frac{1}{2} \int_{B_R} |V(\sigma)|^2$ on the left. Therefore the bounds below control both $\mathcal{H}_{J'}(R)$ and $\mathcal{A}_{J'}(R)$. Define

$$L := \int_J \int_{B_{2R}} |V|^2, \quad C_3 := C_J(2R), \quad D_{3/2} := \mathcal{D}_J(2R).$$

By lemma 6.2,

$$L \leq |Q_{2R,J}|^{1/3} C_3^{2/3} \leq C(R, |J|)(1 + C_3).$$

Since $\phi(y, \tau) = \eta(\tau) \zeta(y)^2$,

$$\partial_\tau \phi = \eta'(\tau) \zeta(y)^2, \quad \nabla \phi = 2\eta(\tau) \zeta(y) \nabla \zeta(y), \quad \Delta \phi = \eta(\tau) \Delta(\zeta(y)^2),$$

hence

$$\|\partial_\tau \phi\|_{L^\infty(Q_{2R,J})} \leq C_\eta \delta^{-1}, \quad \|\nabla \phi\|_{L^\infty(Q_{2R,J})} \leq 2C_\zeta R^{-1}, \quad \|\Delta \phi\|_{L^\infty(Q_{2R,J})} \leq C_\zeta R^{-2}.$$

Step 2: derivative terms from the cut-off.

$$|R_\tau| \leq \frac{1}{2} \|\partial_\tau \phi\|_{L^\infty(Q_{2R,J})} L \leq C \delta^{-1} (1 + C_3),$$

$$|R_\Delta| \leq \frac{\nu}{2} \|\Delta \phi\|_{L^\infty(Q_{2R,J})} L \leq C \nu R^{-2} (1 + C_3).$$

Step 3: convection term. Using $|\nabla \phi| \leq C R^{-1}$,

$$|R_{\text{conv}}| \leq \frac{1}{2} \|\nabla \phi\|_{L^\infty(Q_{2R,J})} \int_J \int_{B_{2R}} |V|^3 \leq C R^{-1} C_3.$$

Step 4: pressure term. For each time τ ,

$$\int_{B_{2R}} (P - (P)_{B_{2R}}(\tau)) V \cdot \nabla \phi = \int_{B_{2R}} P V \cdot \nabla \phi,$$

with

$$\int_{B_{2R}} (P)_{B_{2R}}(\tau) V \cdot \nabla \phi = (P)_{B_{2R}}(\tau) \int_{B_{2R}} V \cdot \nabla \phi = -(P)_{B_{2R}}(\tau) \int_{B_{2R}} \phi \nabla \cdot V = 0$$

(using $\phi \in C_c^\infty(B_{2R})$ and $\nabla \cdot V = 0$). Thus subtracting any time-dependent pressure constant is legitimate in the pressure flux term. Hence

$$|R_{\text{pres}}| \leq \|\nabla \phi\|_{L^\infty(Q_{2R,J})} \int_J \|P - (P)_{B_{2R}}(\tau)\|_{L_x^{3/2}(B_{2R})} \|V(\tau, \cdot)\|_{L_x^3(B_{2R})} d\tau \leq CR^{-1} D_{3/2}^{2/3} C_3^{1/3}.$$

Young's inequality gives

$$|R_{\text{pres}}| \leq C(R)(C_3 + D_{3/2}).$$

Step 5: modulation terms. The local energy inequality already contains the modulation terms after the spatial integrations by parts from the repaired-gauge equation:

$$\int_{B_{2R}} (y \cdot \nabla V) \cdot V \phi = \frac{1}{2} \int_{B_{2R}} y \cdot \nabla (|V|^2) \phi = -\frac{1}{2} \int_{B_{2R}} |V|^2 \nabla \cdot (y \phi),$$

and

$$\int_{B_{2R}} (b(\tau) \cdot \nabla V) \cdot V \phi = \frac{1}{2} \int_{B_{2R}} b(\tau) \cdot \nabla (|V|^2) \phi = -\frac{1}{2} \int_{B_{2R}} |V|^2 b(\tau) \cdot \nabla \phi.$$

Thus we estimate the additional terms directly. Split

$$R_{\text{mod}} = R_a + R_{ay} + R_b,$$

where

$$R_a := - \int_J \int_{B_{2R}} a \frac{|V|^2}{2} \phi, \quad R_{ay} := - \int_J \int_{B_{2R}} a \frac{|V|^2}{2} y \cdot \nabla \phi, \quad R_b := - \int_J \int_{B_{2R}} \frac{|V|^2}{2} b \cdot \nabla \phi.$$

$$M_a := \|a\|_{L^\infty(J)}, \quad M_b := \|b\|_{L^\infty(J)}.$$

For R_a ,

$$|R_a| \leq \frac{1}{2} M_a L \leq C M_a (1 + C_3).$$

For the scale-gradient part, $|y| \leq 2R$ on B_{2R} and $\|\nabla \phi\|_{L^\infty(Q_{2R,J})} \leq CR^{-1}$, hence

$$|y \cdot \nabla \phi| \leq C.$$

Therefore

$$|R_{ay}| \leq C M_a L \leq C M_a (1 + C_3).$$

For R_b ,

$$|R_b| \leq C M_b R^{-1} L \leq C M_b R^{-1} (1 + C_3).$$

Step 6: absorption and conclusion. Collecting estimates from Steps 2–5 yields

$$\nu \int_{\tau_-}^{\tau_+} \int_{B_{2R}} \phi |\nabla V|^2 \leq C(1 + C_3 + D_{3/2}),$$

where C depends only on

$$R, |J|, \nu, \|a\|_{L^\infty(J)}, \|b\|_{L^\infty(J)}, \delta.$$

Since $\phi \equiv 1$ on $B_R \times J'$,

$$\mathcal{H}_{J'}(R) \leq C(1 + \mathcal{C}_J(2R) + \mathcal{D}_J(2R)).$$

Repeating the same estimates with terminal time $\sigma \in J'$ gives

$$\frac{1}{2} \int_{B_R} |V(y, \sigma)|^2 dy \leq C(1 + \mathcal{C}_J(2R) + \mathcal{D}_J(2R))$$

for almost every $\sigma \in J'$. Taking the essential supremum over $\sigma \in J'$ and renaming C proves the lemma. $\square$

Definition 6.5 (Local compact-cylinder control). Local Caffarelli–Kohn–Nirenberg control on $Q_{r,J}$ means

$$\mathcal{C}_J(r) + \mathcal{D}_J(r) < \infty.$$

Local windowed H^1 control on $Q_{r,J}$ means

$$\mathcal{H}_J(r) < \infty.$$

Corollary 6.6 (Outer control implies inner suitable control). *Let $R > 0$, let $J \subset \mathbb{R}$ be a bounded open interval, and let $J' \Subset J$ be a compact subinterval. Assume [proposition 6.1](#) on $Q_{2R,J}$ and $\mathcal{C}_J(2R) + \mathcal{D}_J(2R) < \infty$. Then*

$$\mathcal{A}_{J'}(R) + \mathcal{H}_{J'}(R) < \infty.$$

More quantitatively, the bound is the one in [theorem 10.61](#). If an argument uses pressure means on a smaller ball, [lemma 6.3](#) converts those normalizations to the outer-ball normalization above.

Proof. The estimate follows from [theorem 10.61](#). The comparison of pressure normalizations follows from [lemma 6.3](#). $\square$

Corollary 6.7 (Rough-core exclusion on compact windows). *Let $R > 0$, let $J \subset \mathbb{R}$ be a bounded open interval, and let $J' \Subset J$ be a compact subinterval. Assume [proposition 6.1](#) on $Q_{2R,J}$. If*

$$\mathcal{C}_J(2R) + \mathcal{D}_J(2R) < \infty,$$

then

$$\mathcal{H}_{J'}(R) < \infty.$$

Equivalently, if $\mathcal{H}_{J'}(R) = +\infty$ for an inner cylinder $Q_{R,J'}$, then for every enclosing cylinder $Q_{2R,J}$ with $J' \Subset J$ on which the local energy inequality and bounded modulation hypotheses hold,

$$\mathcal{C}_J(2R) = +\infty \quad \text{or} \quad \mathcal{D}_J(2R) = +\infty.$$

Proof. The direct implication is the gradient part of [theorem 10.61](#). The last assertion is its contrapositive with the same outer-to-inner geometry: finite outer $\mathcal{C} + \mathcal{D}$ on $Q_{2R,J}$ forces finite inner $\mathcal{H}$ on $Q_{R,J'}$. $\square$

Remark 6.8 (Meaning of rough core). Here a rough core means failure of local windowed H^1 control on compact inner renormalized cylinders while the modulation coefficients are bounded. [Corollary 6.7](#) shows that such a failure is not independent of compact-cylinder Caffarelli–Kohn–Nirenberg control: it can occur only if the velocity-pressure bounds $\mathcal{C} + \mathcal{D}$ fail on every relevant larger enclosing compact cylinder.

6.3. Consequences for the assembly. We shall use the following consequences in the final assembly.

- (1) On every compact renormalized window satisfying the local energy inequality and bounded modulation, finite $\mathcal{C} + \mathcal{D}$ on $Q_{2R,J}$ implies finite $\mathcal{A} + \mathcal{H}$ on $Q_{R,J'}$ whenever $J' \Subset J$.
- (2) The pressure normalization may be taken on any comparable ball, after using [lemma 6.3](#).
- (3) Failure of local windowed H^1 control on an inner compact cylinder forces failure of the compact-cylinder $\mathcal{C} + \mathcal{D}$ bounds on each larger enclosing cylinder to which [theorem 10.61](#) applies.

7. SCALE-COLLAPSE COST AND AUTONOMOUS LIMITS

Notation and functional conventions. For a bounded spatial region $K \subset \mathbb{R}^3$, we use

$$(f, g)_{L^2(K)} := \int_K f \cdot g \, dx, \quad \langle F, \phi \rangle_{H^{-1}(K), H_0^1(K)} := \int_0^T \langle F(\cdot, s), \phi(\cdot, s) \rangle_{H^{-1}(K), H_0^1(K)} \, ds$$

whenever $F \in L^1(0, T; H^{-1}(K))$ and $\phi \in L^\infty(0, T; H_0^1(K))$. Weak limits and weak convergences are taken in the distributional sense on C_c^∞ -test functions.

$$\|f\|_{L_t^p L_x^q} := \left(\int_0^T \|f(\cdot, s)\|_{L^q(\mathbb{R}^3)}^p \, ds \right)^{1/p},$$

for any fixed $T > 0$, with the usual modification $p = \infty$.

7.1. Renormalized scale-collapse setting. *Position of the result.* The scale-collapse analysis has two components:

- (1) Sections 7.2–7.4 contain the cost exclusion arguments, which do not use compactness.
- (2) Section 7.5 contains the autonomous compactness statement under explicitly stated compactness hypotheses.

Throughout the scale-collapse analysis, the renormalized representation established earlier is assumed: on the time interval under consideration, the fields (V, P) satisfy (8) and its distributional weak formulation. In particular, the solenoidal formulation is obtained by restricting the test fields below to divergence-free fields.

Let (u, p) be a divergence-free suitable weak solution of

$$\partial_t u + (u \cdot \nabla) u + \nabla p = \nu \Delta u, \quad \nabla \cdot u = 0,$$

in $\mathbb{R}^3$, with viscosity $\nu > 0$. Let $\lambda(\tau) > 0$, $X(\tau) \in \mathbb{R}^3$, and a measurable gauge function $c(t)$ be given, and define the renormalized fields by

$$u(x, t) = \lambda(\tau)^{-1} V \left(\frac{x - X(\tau)}{\lambda(\tau)}, \tau \right), \quad p(x, t) = \lambda(\tau)^{-2} P \left(\frac{x - X(\tau)}{\lambda(\tau)}, \tau \right) + c(t).$$

Pressure is defined only up to an arbitrary function of time; $c(t)$ has no effect on ∇p , so it is irrelevant for all subsequent estimates. Assume that (V, P) satisfies on $(0, \infty) \times \mathbb{R}^3$

$$(8) \quad \partial_\tau V + (V \cdot \nabla) V + \nabla P = \nu \Delta V + a(\tau)(V + y \cdot \nabla V) + b(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0.$$

For every $\psi \in C_c^\infty((0, \infty) \times \mathbb{R}^3; \mathbb{R}^3)$ with $\nabla \cdot \psi = 0$, the pair (V, P) satisfies the weak formulation

$$\begin{aligned} 0 &= \int_0^\infty \int_{\mathbb{R}^3} (-V \cdot \partial_\tau \psi - (V \otimes V) : \nabla \psi + \nu \nabla V : \nabla \psi) \, dy \, d\tau \\ &\quad + \int_0^\infty \int_{\mathbb{R}^3} (-a(\tau)(V + y \cdot \nabla V) \cdot \psi - (b(\tau) \cdot \nabla V) \cdot \psi - P \nabla \cdot \psi) \, dy \, d\tau. \end{aligned}$$

For every $0 < T < \infty$, restricting the τ -integration to $(0, T)$ gives the same identity on $(0, T) \times \mathbb{R}^3$. Equivalently, for every $\psi \in C_c^\infty((0, \infty) \times \mathbb{R}^3; \mathbb{R}^3)$, not necessarily divergence free,

$$\begin{aligned} 0 &= \int_0^\infty \int_{\mathbb{R}^3} (-V \cdot \partial_\tau \psi - (V \otimes V) : \nabla \psi + \nu \nabla V : \nabla \psi - P \nabla \cdot \psi) \, dy \, d\tau \\ &\quad + \int_0^\infty \int_{\mathbb{R}^3} (-a(\tau)(V + y \cdot \nabla V) \cdot \psi - (b(\tau) \cdot \nabla V) \cdot \psi) \, dy \, d\tau. \end{aligned}$$

Define positive and negative parts of the scalar field a by

$$a_+(\tau) := \max\{a(\tau), 0\}, \quad a_-(\tau) := \max\{-a(\tau), 0\}.$$

Then for a.e. τ , $a = a_+ - a_-$, $a_+, a_- \geq 0$, and $a_+ a_- = 0$. Assume

$$a \in L^1_{\text{loc}}([\tau_0, \infty)), \quad b \in L^1_{\text{loc}}([\tau_0, \infty); \mathbb{R}^3),$$

and the scale convention

$$(9) \quad \frac{d}{d\tau} \log \lambda(\tau) = a(\tau)$$

for a.e. $\tau \geq \tau_0$, with $\log \lambda$ absolutely continuous on compact intervals.

All integrals below are interpreted in $[0, \infty]$. When we write $\int_{\tau_0}^{\infty}$, partial integrals are taken over $[\tau_0, T]$ and monotone convergence is used as $T \rightarrow \infty$.

Fix $R > 0$. Choose $\Phi \in C_c^\infty(\mathbb{R}^3; [0, 1])$ satisfying

$$(10) \quad \Phi(y) = 1 \text{ if } |y| \leq R, \quad \text{supp}_y \Phi \subset B_{2R}.$$

Define the localized kinetic mass

$$M_R(\tau) := \int_{\mathbb{R}^3} |V(y, \tau)|^2 \Phi(y) dy.$$

Since Φ is fixed in τ , M_R is measurable in τ , and only this measurability is used below.

Definition 7.1 (Nonvanishing localized core). A renormalized branch has nonvanishing localized core if there exist constants $m_0 > 0$ and $\tau_1 \geq \tau_0$ such that

$$M_R(\tau) \geq m_0 \quad \text{for a.e. } \tau \in [\tau_1, \infty).$$

In applications, this lower bound is supplied by the selected-core construction; here it is stated as an explicit hypothesis.

Definition 7.2 (Scale-collapse and absolute scale-drift densities). Define

$$\mathcal{C}_{\text{sc}}(\tau) := a_-(\tau) M_R(\tau), \quad \mathcal{C}_{\text{abs}}(\tau) := \nu \int_{\mathbb{R}^3} |\nabla V(y, \tau)|^2 \Phi(y) dy + |a(\tau)| M_R(\tau).$$

$\mathcal{C}_{\text{sc}}$ is the negative drift weighted by localized core mass, while $\mathcal{C}_{\text{abs}}$ controls both collapse drift and local dissipation. Their total costs are extended real-valued by

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{sc}}(\tau) d\tau, \quad \int_{\tau_0}^{\infty} \mathcal{C}_{\text{abs}}(\tau) d\tau.$$

Definition 7.3 (Finite-cost branch). A branch has finite scale-collapsing cost if

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{sc}}(\tau) d\tau < \infty.$$

It has finite absolute cost if

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{abs}}(\tau) d\tau < \infty.$$

7.2. Elementary alternative.

Lemma 7.4 (Finite-or-infinite alternative). *Let $f \geq 0$ be measurable on $[\tau_0, \infty)$. Then exactly one of the following holds:*

$$\int_{\tau_0}^{\infty} f(\tau) d\tau < \infty \quad \text{or} \quad \int_{\tau_0}^{\infty} f(\tau) d\tau = \infty.$$

Proof. (1) For $T > \tau_0$, define partial costs

$$F(T) := \int_{\tau_0}^T f(\tau) d\tau \in [0, \infty].$$

(2) If $T_2 > T_1$, then by nonnegativity of f ,

$$F(T_2) - F(T_1) = \int_{T_1}^{T_2} f(\tau) d\tau \geq 0,$$

so F is monotone nondecreasing.

(3) A monotone function on $[\tau_0, \infty)$ has an extended limit $L \in [0, \infty]$,

$$L := \lim_{T \rightarrow \infty} F(T).$$

By definition of improper integral, this limit equals $\int_{\tau_0}^{\infty} f(\tau) d\tau$.

(4) Therefore L is either finite or equal to $+\infty$, and these cases are mutually exclusive. Hence exactly one item in the lemma holds. $\square$

Lemma 7.5 (Logarithmic scale identity). *Under (9), for all τ_1, τ_2 with $\tau_0 \leq \tau_1 < \tau_2 < \infty$,*

$$\int_{\tau_1}^{\tau_2} a(\tau) d\tau = \log \lambda(\tau_2) - \log \lambda(\tau_1).$$

Proof. (1) Since $a \in L^1_{\text{loc}}$, the map $\log \lambda$ is absolutely continuous on $[\tau_1, \tau_2]$, hence has a representative with a.e. derivative equal to a .

(2) By the fundamental theorem of calculus for absolutely continuous functions, for every $\tau \in [\tau_1, \tau_2]$,

$$\log \lambda(\tau) = \log \lambda(\tau_1) + \int_{\tau_1}^{\tau} a(s) ds.$$

(3) Setting $\tau = \tau_2$ yields the identity. $\square$

7.3. Scale collapse versus localized core lower bounds.

Theorem 7.6 (Collapse with nonvanishing localized core forces infinite scale-collapsing cost).

Assume:

- (1) $\lambda(\tau) \rightarrow 0$ as $\tau \rightarrow \infty$;
- (2) the branch has nonvanishing localized core.

Then

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{sc}}(\tau) d\tau = \int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau = \infty.$$

Proof. (1) Choose $m_0 > 0$ and $\tau_1 \geq \tau_0$ from the nonvanishing core assumption so that

$$M_R(\tau) \geq m_0 \quad \text{for a.e. } \tau \in [\tau_1, \infty).$$

(2) By Lemma 7.5, for any $\tau_2 > \tau_1$,

$$\int_{\tau_1}^{\tau_2} a(\tau) d\tau = \log \lambda(\tau_2) - \log \lambda(\tau_1).$$

Taking $\tau_2 \rightarrow \infty$ and using $\lambda(\tau_2) \rightarrow 0$,

$$\lim_{\tau_2 \rightarrow \infty} \int_{\tau_1}^{\tau_2} a(\tau) d\tau = -\infty.$$

(3) Decompose $a = a_+ - a_-$. Then for each $\tau_2 > \tau_1$,

$$\int_{\tau_1}^{\tau_2} a(\tau) d\tau = \int_{\tau_1}^{\tau_2} a_+(\tau) d\tau - \int_{\tau_1}^{\tau_2} a_-(\tau) d\tau \geq - \int_{\tau_1}^{\tau_2} a_-(\tau) d\tau.$$

Hence

$$\int_{\tau_1}^{\tau_2} a_-(\tau) d\tau \geq - \int_{\tau_1}^{\tau_2} a(\tau) d\tau.$$

(4) Taking limits and using Step 2 yields

$$\int_{\tau_1}^{\infty} a_-(\tau) d\tau = \infty.$$

Indeed, by Step 2,

$$- \int_{\tau_1}^{\tau_2} a(\tau) d\tau \rightarrow +\infty.$$

Since the partial integrals of a_- dominate this quantity, they are unbounded. Because $a_- \geq 0$, monotone convergence gives $\int_{\tau_1}^{\infty} a_-(\tau) d\tau = \infty$.

(5) Split the cost into early and late times:

$$\int_{\tau_0}^{\infty} a_- M_R d\tau = \int_{\tau_0}^{\tau_1} a_- M_R d\tau + \int_{\tau_1}^{\infty} a_- M_R d\tau.$$

The first term is finite or infinite independently of the argument; the second term satisfies

$$\int_{\tau_1}^{\infty} a_-(\tau) M_R(\tau) d\tau \geq m_0 \int_{\tau_1}^{\infty} a_-(\tau) d\tau = \infty$$

by the bound in (1) and Step 4. Hence the total integral is infinite. $\square$

7.4. Finite-cost exclusions.

Theorem 7.7 (Scale-collapse cost exclusion). *No branch with finite scale-collapsing cost (Definition 7.3) can satisfy simultaneously $\lambda(\tau) \rightarrow 0$ and nonvanishing localized core.*

Proof. Assume, for contradiction, that such a branch exists. The two hypotheses are exactly those of Theorem 7.6; therefore

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{sc}}(\tau) d\tau = \infty.$$

This contradicts finite scale-collapsing cost. Therefore no such branch exists. $\square$

Lemma 7.8 (Absolute-cost domination). *For a.e. $\tau \in [\tau_0, \infty)$,*

$$\mathcal{C}_{\text{abs}}(\tau) \geq \mathcal{C}_{\text{sc}}(\tau).$$

Hence, pointwise and after integrating in τ ,

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{sc}}(\tau) d\tau \leq \int_{\tau_0}^{\infty} \mathcal{C}_{\text{abs}}(\tau) d\tau,$$

as extended real numbers. In particular, finite absolute cost implies finite scale-collapsing cost, and equivalently

$$\int_{\tau_0}^{\infty} \mathcal{C}_{\text{sc}}(\tau) d\tau = \infty \implies \int_{\tau_0}^{\infty} \mathcal{C}_{\text{abs}}(\tau) d\tau = \infty.$$

Proof. At each time τ ,

$$\mathcal{C}_{\text{abs}}(\tau) = \nu \int |\nabla V|^2 \Phi + (a_+ + a_-) M_R \geq (a_+ + a_-) M_R \geq a_- M_R = \mathcal{C}_{\text{sc}}(\tau),$$

since $a_+, a_- \geq 0$ and $\nu \int |\nabla V|^2 \Phi \geq 0$. Integrating over $[\tau_0, \infty)$ yields the claimed implication. $\square$

Corollary 7.9 (Absolute scale-cost exclusion). *No branch with finite absolute cost in the sense of Definition 7.3 can satisfy simultaneously $\lambda(\tau) \rightarrow 0$ and nonvanishing localized core.*

Proof. Assume instead such a branch exists. By Theorem 7.6, $\int \mathcal{C}_{\text{sc}} = \infty$. By Lemma 7.8, the absolute cost is pointwise above $\mathcal{C}_{\text{sc}}$, so its integral must also be infinite, contradicting finite absolute cost. $\square$

Remark 7.10. The exclusions above are purely integral and use neither compactness nor any autonomous-limit argument.

7.5. Autonomous reduced-limit theorem from the selected-window estimates. The preceding exclusions are integral statements and do not rely on compactness. On a fixed finite autonomous window, compactness requires not only the cost bound but also the selected-window estimates supplied by the preceding renormalized representation, pressure reconstruction, and local Caccioppoli estimates.

Analytic inputs. The preceding sections supply the following ingredients on selected compact renormalized windows:

- (1) the repaired-gauge renormalized equation and weak formulation for (V, P) ;
- (2) pressure reconstruction modulo spatial means, with local $L^{3/2}$ bounds on compact cylinders;
- (3) local windowed bounds $V \in L_t^\infty L_x^2 \cap L_t^2 H_x^1$ on compact cylinders;
- (4) a retained compact core lower bound preventing the selected profile from vanishing.

Together with the thick-window convergence of the modulation coefficients, these estimates imply the compactness and limit passage below.

Definition 7.11 (Thick autonomous windows). A branch has *thick autonomous windows* if there exist constants $T > 0$, $c > 0$, $M_1 \in (0, \infty)$, a sequence (τ_n) with $\tau_n \rightarrow \infty$, and limits $a_\infty \in \mathbb{R}$, $b_\infty \in \mathbb{R}^3$, such that for every n :

(i)

$$\frac{1}{T} \int_{\tau_n}^{\tau_n+T} a_-(\tau) M_R(\tau) d\tau \geq c;$$

(ii)

$$0 \leq M_R(\tau) \leq M_1 \quad \text{for a.e. } \tau \in [\tau_n, \tau_n + T];$$

(iii)

$$a(\tau_n + \cdot) \rightarrow a_\infty \text{ in } L^1(0, T), \quad b(\tau_n + \cdot) \rightarrow b_\infty \text{ in } L^1(0, T; \mathbb{R}^3).$$

Remark 7.12. This notion isolates a limiting drift coefficient $a_\infty < 0$. The compactness to extract a profile limit is obtained from the selected-window compactness estimates, not from the scalar cost inequalities of Sections 7.2–7.4.

Lemma 7.13 (Thick drift implies negative autonomous limit parameter). *On thick autonomous windows, the limiting autonomous coefficient is negative:*

$$a_\infty \leq -\frac{c}{M_1} < 0.$$

Proof. (1) From (ii), for a.e. $\tau \in [\tau_n, \tau_n + T]$, $0 \leq M_R(\tau) \leq M_1$. Therefore

$$\frac{a_-(\tau) M_R(\tau)}{M_1} \leq a_-(\tau).$$

Therefore

$$\frac{1}{T} \int_{\tau_n}^{\tau_n+T} a_-(\tau) d\tau \geq \frac{1}{M_1 T} \int_{\tau_n}^{\tau_n+T} a_-(\tau) M_R(\tau) d\tau \geq \frac{c}{M_1}.$$

- (2) Set $a_n(s) := a(\tau_n + s)$. By (iii), $a_n \rightarrow a_\infty$ in $L^1(0, T)$. The map $r \mapsto r^-$ is 1-Lipschitz, so $a_n^- \rightarrow a_\infty^-$ in $L^1(0, T)$:

$$\|a_n^- - a_\infty^-\|_{L^1(0, T)} \leq \|a_n - a_\infty\|_{L^1(0, T)} \rightarrow 0.$$

Thus

$$\lim_{n \rightarrow \infty} \frac{1}{T} \int_{\tau_n}^{\tau_n + T} a_-(\tau) d\tau = \frac{1}{T} \int_0^T (a_\infty)^- ds = (a_\infty)^-.$$

- (3) Combine Step 1 and Step 2: $(a_\infty)^- \geq c/M_1$, i.e. $a_\infty \leq -c/M_1 < 0$. □

Proposition 7.14 (Selected-window compact upper estimates). *Let (τ_n) be a thick autonomous-window sequence and set*

$$\begin{aligned} V_n(y, s) &:= V(y, \tau_n + s), & P_n(y, s) &:= P(y, \tau_n + s), \\ \alpha_n(s) &:= a(\tau_n + s), & \beta_n(s) &:= b(\tau_n + s), \end{aligned} \quad s \in (0, T).$$

Assume that the repaired-gauge representation of [theorem 4.37](#), the localized pressure decomposition of [theorem 4.25](#), and the canonical compact-cylinder Caccioppoli estimate [theorem 10.61](#) are available on the selected windows. Assume also that the ordered local validation has not entered a local-energy, pressure-reconstruction, or rough-core failure row on these windows. Then, after passing to a subsequence, for every bounded Lipschitz domain $K \Subset \mathbb{R}^3$, there is $C_K < \infty$, independent of n , such that

$$\|V_n\|_{L^\infty(0, T; L^2(K))} + \|V_n\|_{L^2(0, T; H^1(K))} + \|P_n - (P_n)_K\|_{L^{3/2}((0, T) \times K)} \leq C_K,$$

where

$$(P_n)_K(s) := \frac{1}{|K|} \int_K P_n(y, s) dy.$$

Each V_n is divergence free.

Proof. The repaired-gauge representation supplies the shifted equation and divergence-free condition. The localized pressure decomposition gives the $L^{3/2}$ pressure oscillation bound after subtracting spatial means. The compact-cylinder Caccioppoli estimate on a slightly larger compact cylinder gives the local $L_t^\infty L_x^2$ and $L_t^2 H_x^1$ velocity bounds. If any of these inputs is missing, the ordered validation records the first missing input as a local-energy, pressure-reconstruction, or rough-core failure row, which is outside the retained compact-estimate branch. Hence, on the retained branch, the displayed estimates hold. □

Corollary 7.15 (Derivation of selected-window compact upper estimates). *Under the hypotheses of [proposition 7.14](#), the local compact upper estimate package of that proposition holds on every compact $K \Subset \mathbb{R}^3$.*

Proof. This is exactly [proposition 7.14](#). The corollary is kept only as a reference point for later sections that cite the derived estimate package. □

Lemma 7.16 (Retained selected-window mass floor). *Let (τ_n) be a selected fixed-length terminal window sequence of length $T > 0$ for a represented Type II branch produced by the selected-core construction of the preceding sections, and write*

$$V_n(y, s) := V(y, \tau_n + s), \quad s \in (0, T).$$

If the loss-of-retained-active-core row does not occur on these selected windows, then after passing to a subsequence there exist

$$0 < \theta < T/2, \quad R > 0, \quad \eta > 0,$$

such that for all sufficiently large n ,

$$\int_{T/2-\theta}^{T/2+\theta} \int_{B_R} |V_n(y, s)|^3 dy ds \geq \eta.$$

Proof. The selected-core construction retains a positive local active window whenever the loss-of-retained-active-core row is excluded. If no fixed compact space-time subwindow carried a positive L^3 -mass floor after passing to subsequences, then every candidate retained core would lose its active space-time mass under further window selection, which is precisely the loss-of-retained-active-core row. Therefore a subsequence admits one compact ball and one compact time interval with the displayed positive L^3 floor. $\square$

Lemma 7.17 (Local time regularity on autonomous windows). *Assume thick autonomous windows and Proposition 7.14. Let $K_0 \Subset K_1 \Subset \mathbb{R}^3$ be bounded Lipschitz domains, and choose an integer m so large that $H_0^m(K_1) \hookrightarrow W_0^{1,\infty}(K_1)$. Then*

$$\{\partial_s V_n\}_{n \geq 1} \text{ is bounded in } L^1(0, T; H^{-m}(K_1)).$$

Proof. Fix $\zeta \in C_c^\infty(K_1)$ with $\zeta \equiv 1$ on a neighborhood of K_0 . Let $\phi \in H_0^m(K_1; \mathbb{R}^3)$. Since $H_0^m(K_1) \hookrightarrow W_0^{1,\infty}(K_1)$, there is C_m such that

$$\|\phi\|_{L^\infty(K_1)} + \|\nabla \phi\|_{L^\infty(K_1)} \leq C_m \|\phi\|_{H^m(K_1)}.$$

The shifted equation is

$$\partial_s V_n = -(V_n \cdot \nabla) V_n - \nabla P_n + \nu \Delta V_n + \alpha_n (V_n + y \cdot \nabla V_n) + \beta_n \cdot \nabla V_n$$

in distributions. We estimate each term against ϕ .

For the convection term, integration by parts gives

$$|\langle (V_n \cdot \nabla) V_n, \phi \rangle| = \left| \int_{K_1} V_{n,i} V_{n,j} \partial_j \phi_i dy \right| \leq C_m \|V_n(s)\|_{L^2(K_1)}^2 \|\phi\|_{H^m(K_1)}.$$

For the viscous term,

$$|\langle \nu \Delta V_n, \phi \rangle| = \nu \left| \int_{K_1} \nabla V_n : \nabla \phi dy \right| \leq \nu C_m \|\nabla V_n(s)\|_{L^2(K_1)} |K_1|^{1/2} \|\phi\|_{H^m(K_1)}.$$

For the pressure term, subtracting the spatial mean does not change the gradient:

$$|\langle \nabla P_n, \phi \rangle| = \left| \int_{K_1} (P_n - (P_n)_{K_1}) \nabla \cdot \phi dy \right| \leq C_m |K_1|^{1/3} \|P_n - (P_n)_{K_1}\|_{L^{3/2}(K_1)} \|\phi\|_{H^m(K_1)}.$$

For the V_n -part of the scaling drift,

$$|\alpha_n(s)| \left| \int_{K_1} V_n \cdot \phi dy \right| \leq C_m |\alpha_n(s)| |K_1|^{1/2} \|V_n(s)\|_{L^2(K_1)} \|\phi\|_{H^m(K_1)}.$$

For the $y \cdot \nabla V_n$ -part, integrate by parts:

$$\int_{K_1} (y \cdot \nabla V_n) \cdot \phi dy = - \int_{K_1} V_n \cdot (y \cdot \nabla \phi + 3\phi) dy,$$

because ϕ has zero trace. Hence

$$|\alpha_n(s)| \left| \int_{K_1} (y \cdot \nabla V_n) \cdot \phi dy \right| \leq C_{K_1, m} |\alpha_n(s)| \|V_n(s)\|_{L^2(K_1)} \|\phi\|_{H^m(K_1)}.$$

Similarly,

$$\left| \int_{K_1} (\beta_n \cdot \nabla V_n) \cdot \phi dy \right| = \left| \int_{K_1} V_n \cdot (\beta_n \cdot \nabla \phi) dy \right| \leq C_m |\beta_n(s)| |K_1|^{1/2} \|V_n(s)\|_{L^2(K_1)} \|\phi\|_{H^m(K_1)}.$$

Combining the estimates and taking the supremum over $\|\phi\|_{H^m(K_1)} \leq 1$, we obtain

$$\begin{aligned} \|\partial_s V_n(s)\|_{H^{-m}(K_1)} &\leq C_{K_1, m} (\|V_n(s)\|_{L^2(K_1)}^2 + \|\nabla V_n(s)\|_{L^2(K_1)} \\ &\quad + \|P_n(s) - (P_n)_{K_1}(s)\|_{L^{3/2}(K_1)} + (|\alpha_n(s)| + |\beta_n(s)|) \|V_n(s)\|_{L^2(K_1)}). \end{aligned}$$

Integrating in $s \in (0, T)$, using Proposition 7.14, and using $\alpha_n \rightarrow a_\infty$, $\beta_n \rightarrow b_\infty$ in L^1 , hence uniform L^1 -bounds for α_n, β_n , proves the claimed $L^1(0, T; H^{-m})$ bound. $\square$

Lemma 7.18 (Local compactness on autonomous windows). *Assume thick autonomous windows and Proposition 7.14. Then after passing to a subsequence there exist*

$$U \in L^\infty(0, T; L_{\text{loc}}^2(\mathbb{R}^3)) \cap L^2(0, T; H_{\text{loc}}^1(\mathbb{R}^3)), \quad \Pi \in L_{\text{loc}}^{3/2}((0, T) \times \mathbb{R}^3),$$

such that, for every compact $K \Subset \mathbb{R}^3$,

(a)

$$V_n \rightarrow U \quad \text{strongly in } L^2(0, T; L^2(K));$$

(b)

$$V_n \rightarrow U \quad \text{strongly in } L^3((0, T) \times K);$$

(c)

$$V_n \rightharpoonup U \quad \text{weakly in } L^2(0, T; H^1(K));$$

(d)

$$V_n \rightharpoonup^* U \quad \text{weakly-}^* \text{ in } L^\infty(0, T; L^2(K));$$

(e)

$$P_n - (P_n)_K \rightharpoonup \Pi_K \quad \text{weakly in } L^{3/2}((0, T) \times K),$$

where the local distributions $\nabla \Pi_K$ agree on overlaps and define a pressure gradient $\nabla \Pi$ modulo functions of time.

Proof. Fix $K_0 \Subset K_1 \Subset \mathbb{R}^3$. Proposition 7.14 gives boundedness of V_n in $L^2(0, T; H^1(K_1))$, and Lemma 7.17 gives boundedness of $\partial_s V_n$ in $L^1(0, T; H^{-m}(K_1))$ for m large. Since

$$H^1(K_1) \Subset L^2(K_0) \hookrightarrow H^{-m}(K_1),$$

the Aubin–Lions–Simon compactness theorem implies, after extracting a subsequence,

$$V_n \rightarrow U \quad \text{strongly in } L^2(0, T; L^2(K_0)).$$

The local $L_t^\infty L_x^2$ and $L_t^2 H_x^1$ bounds imply a uniform $L^{10/3}((0, T) \times K_0)$ bound by standard interpolation. Interpolating between the strong L^2 convergence and this uniform $L^{10/3}$ bound gives

$$V_n \rightarrow U \quad \text{strongly in } L^q((0, T) \times K_0)$$

for every $2 \leq q < 10/3$, in particular for $q = 3$. The local $L^\infty L^2$ and $L^2 H^1$ bounds also give, after extraction,

$$V_n \rightharpoonup^* U \quad \text{in } L^\infty(0, T; L^2(K_0)), \quad V_n \rightharpoonup U \quad \text{in } L^2(0, T; H^1(K_0)).$$

The pressure bound in Proposition 7.14 gives weak compactness of $P_n - (P_n)_{K_0}$ in $L^{3/2}((0, T) \times K_0)$. A diagonal extraction over balls B_j gives a single subsequence for every compact K . On overlaps, two pressure limits can differ only by a function of s , because both represent the same distributional pressure gradient in the limit equation. Thus the gradients assemble into $\nabla \Pi$, with Π defined modulo functions of time. $\square$

Lemma 7.19 (Compatibility of local pressure limits). *Let $K_1, K_2 \Subset \mathbb{R}^3$ be bounded Lipschitz domains with nonempty connected overlap $K_{12} := K_1 \cap K_2$. Suppose that, after passing to a common subsequence,*

$$P_n - (P_n)_{K_i} \rightharpoonup \Pi_i \quad \text{in } L^{3/2}((0, T) \times K_i), \quad i = 1, 2.$$

Then there exists $c_{12} \in L^{3/2}(0, T)$ such that

$$\Pi_1(y, s) - \Pi_2(y, s) = c_{12}(s) \quad \text{for a.e. } (y, s) \in K_{12} \times (0, T).$$

Consequently the pressure gradient $\nabla \Pi$ obtained from the local limits is globally well-defined on $(0, T) \times \mathbb{R}^3$, with the pressure itself defined modulo functions of time.

Proof. For each n ,

$$(P_n - (P_n)_{K_1}) - (P_n - (P_n)_{K_2}) = (P_n)_{K_2}(s) - (P_n)_{K_1}(s)$$

on K_{12} . Hence its spatial gradient is zero in $\mathcal{D}'((0, T) \times K_{12})$. Passing to the weak limit gives

$$\nabla_y(\Pi_1 - \Pi_2) = 0 \quad \text{in } \mathcal{D}'((0, T) \times K_{12}).$$

Since K_{12} is connected, the standard distributional characterization of functions with zero spatial gradient implies that $\Pi_1 - \Pi_2 = c_{12}(s)$ for some $c_{12} \in L^{3/2}(0, T)$. The last assertion follows because adding a time-dependent scalar to pressure does not change $\nabla \Pi$. $\square$

Theorem 7.20 (Autonomous reduced-limit theorem). *Assume that a branch has thick autonomous windows and satisfies Proposition 7.14 on the same window sequence. Let (U, Π) be the limit extracted in Lemma 7.18. Then*

$$\nabla \cdot U = 0,$$

$$U \not\equiv 0$$

and (U, Π) satisfies

$$(11) \quad \partial_s U + (U \cdot \nabla)U + \nabla \Pi = \nu \Delta U + a_\infty(U + y \cdot \nabla U) + b_\infty \cdot \nabla U, \quad \nabla \cdot U = 0,$$

in $\mathcal{D}'((0, T) \times \mathbb{R}^3)$. Equivalently, for every $\varphi \in C_c^\infty((0, T) \times \mathbb{R}^3; \mathbb{R}^3)$,

$$\int_0^T \int (-U \cdot \partial_s \varphi - (U \otimes U) : \nabla \varphi + \nu \nabla U : \nabla \varphi - \Pi \nabla \cdot \varphi - a_\infty(U + y \cdot \nabla U) \cdot \varphi - (b_\infty \cdot \nabla U) \cdot \varphi) dy ds = 0.$$

Proof. Step 1: divergence-free property of the limit.

Each V_n is divergence free by Proposition 7.14, so for any compactly supported scalar $\psi \in C_c^\infty((0, T) \times \mathbb{R}^3)$,

$$\int_{(0, T) \times \mathbb{R}^3} V_n \cdot \nabla \psi dy ds = 0.$$

Hence, for every such ψ , since $\nabla \psi \in C_c^\infty((0, T) \times \mathbb{R}^3)$ and $V_n \rightarrow U$ in $L^2(0, T; L_{\text{loc}}^2)$,

$$\int_{(0, T) \times \mathbb{R}^3} (V_n - U) \cdot \nabla \psi dy ds \rightarrow 0,$$

we deduce

$$\int_{(0, T) \times \mathbb{R}^3} U \cdot \nabla \psi dy ds = 0.$$

Hence $\nabla \cdot U = 0$ in $\mathcal{D}'((0, T) \times \mathbb{R}^3)$.

Step 2: nontriviality of the limit profile. By the retained-window branch of theorem 7.22, after passing to the same subsequence there exist $R > 0$, $\chi \in C_c^\infty(B_R)$, $\chi \geq 0$, and $\eta > 0$ such that

$$\int_0^T \int_{\mathbb{R}^3} |V_n(y, s)|^3 \chi(y) dy ds \geq \eta$$

for all sufficiently large n . By [lemma 7.18](#), $V_n \rightarrow U$ strongly in $L^3((0, T) \times B_R)$. Hence

$$\int_0^T \int_{\mathbb{R}^3} |U(y, s)|^3 \chi(y) dy ds = \lim_{n \rightarrow \infty} \int_0^T \int_{\mathbb{R}^3} |V_n(y, s)|^3 \chi(y) dy ds \geq \eta.$$

Thus $U \not\equiv 0$.

Step 3: shifted weak formulation for V_n . Take any test field $\varphi \in C_c^\infty((0, T) \times \mathbb{R}^3; \mathbb{R}^3)$, and set

$$K_\varphi := \overline{\{x \in \mathbb{R}^3 : \exists s \in (0, T), \varphi(x, s) \neq 0\}} \subset \mathbb{R}^3, \quad K_\varphi \text{ bounded.}$$

Since (V, P) satisfies (8) in distributions, the change of variables $\tau = \tau_n + s$ gives the shifted weak formulation

$$(12) \quad \begin{aligned} 0 &= \int_0^T \int_{\mathbb{R}^3} \left(-V_n \cdot \partial_s \varphi - (V_n \otimes V_n) : \nabla \varphi + \nu \nabla V_n : \nabla \varphi - P_n \nabla \cdot \varphi \right) dy ds \\ &\quad + \int_0^T \int_{\mathbb{R}^3} \left(-a(\tau_n + s)(V_n + y \cdot \nabla V_n) \cdot \varphi - (b(\tau_n + s) \cdot \nabla V_n) \cdot \varphi \right) dy ds. \end{aligned}$$

Step 4: term-by-term passage to the limit. Write $\alpha_n(s) := a(\tau_n + s)$, $\beta_n(s) := b(\tau_n + s)$.

(i) *Time derivative.*

By strong convergence in $L^2(0, T; L^2(K_\varphi))$,

$$\left| \int_0^T \int (V_n - U) \cdot \partial_s \varphi dy ds \right| \leq \|V_n - U\|_{L^2(0, T; L^2(K_\varphi))} \|\partial_s \varphi\|_{L^2(0, T; L^2(K_\varphi))} \rightarrow 0.$$

Hence

$$\int V_n \cdot \partial_s \varphi \rightarrow \int U \cdot \partial_s \varphi.$$

(ii) *Convection term.* Decompose

$$V_n \otimes V_n - U \otimes U = (V_n - U) \otimes V_n + U \otimes (V_n - U).$$

Then

$$\begin{aligned} &\left| \int_0^T \int ((V_n \otimes V_n - U \otimes U) : \nabla \varphi) dy ds \right| \\ &\leq \|\nabla \varphi\|_{L_{x,t}^\infty} \|V_n - U\|_{L^2(0, T; L^2(K_\varphi))} (\|V_n\|_{L^2(0, T; L^2(K_\varphi))} + \|U\|_{L^2(0, T; L^2(K_\varphi))}) \rightarrow 0, \end{aligned}$$

using boundedness of V_n in $L^2(0, T; L^2(K_\varphi))$. Therefore $\int (V_n \otimes V_n) : \nabla \varphi \rightarrow \int (U \otimes U) : \nabla \varphi$.

(iii) *Viscous term.* From weak convergence of V_n in $L^2(0, T; H_{\text{loc}}^1)$,

$$\int_0^T \int \nabla V_n : \nabla \varphi dy ds \rightarrow \int_0^T \int \nabla U : \nabla \varphi dy ds.$$

(iv) *Pressure term.* Let K be a bounded Lipschitz domain with $K_\varphi \Subset K$. Since φ is compactly supported in K ,

$$\int_K (P_n)_K(s) \nabla \cdot \varphi(y, s) dy = (P_n)_K(s) \int_K \nabla \cdot \varphi(y, s) dy = 0.$$

Therefore

$$\int_0^T \int_{\mathbb{R}^3} P_n \nabla \cdot \varphi dy ds = \int_0^T \int_K (P_n - (P_n)_K) \nabla \cdot \varphi dy ds.$$

By [Lemma 7.18](#),

$$P_n - (P_n)_K \rightharpoonup \Pi_K \quad \text{in } L^{3/2}((0, T) \times K),$$

and $\nabla \cdot \varphi \in L^3((0, T) \times K)$. The compatibility Lemma 7.19 permits us to write the resulting pressure distribution as Π , modulo functions of time. Hence

$$\int_0^T \int_{\mathbb{R}^3} P_n \nabla \cdot \varphi \, dy \, ds \rightarrow \int_0^T \int_K \Pi_K \nabla \cdot \varphi \, dy \, ds.$$

The value of the last integral is independent of the chosen K , because replacing Π_K by $\Pi_K + c(s)$ changes the integral by $\int_0^T c(s) \int_K \nabla \cdot \varphi \, dy \, ds = 0$. The limit is therefore the pressure term $\int \Pi \nabla \cdot \varphi$.

(v) *Drift term with a : part involving V_n .* Define

$$Y_n^0(s) := \int_{\mathbb{R}^3} V_n(y, s) \cdot \varphi(y, s) \, dy, \quad Y^0(s) := \int_{\mathbb{R}^3} U(y, s) \cdot \varphi(y, s) \, dy.$$

For each s , Cauchy–Schwarz gives

$$|Y_n^0(s)| \leq \|V_n(s)\|_{L^2(K_\varphi)} \|\varphi(s)\|_{L^2(K_\varphi)}.$$

By Proposition 7.14, there is $C_0 < \infty$ such that $\sup_n \|Y_n^0\|_{L^\infty(0, T)} \leq C_0$. Moreover,

$$Y_n^0 \rightarrow Y^0 \quad \text{in } L^2(0, T),$$

because $V_n \rightarrow U$ in $L^2(0, T; L^2(K_\varphi))$. Now

$$\int_0^T \alpha_n(s) Y_n^0(s) \, ds = \int_0^T (\alpha_n(s) - a_\infty) Y_n^0(s) \, ds + a_\infty \int_0^T Y_n^0(s) \, ds.$$

For the first term,

$$\left| \int_0^T (\alpha_n - a_\infty) Y_n^0 \, ds \right| \leq \sup_{s \in (0, T)} |Y_n^0(s)| \|\alpha_n - a_\infty\|_{L^1(0, T)} \rightarrow 0.$$

For the second term, strong convergence in $L^2(0, T)$ implies $\int_0^T Y_n^0 \rightarrow \int_0^T Y^0$, so

$$\int_0^T \alpha_n(s) \int_{\mathbb{R}^3} V_n \cdot \varphi \, dy \, ds \rightarrow a_\infty \int_0^T \int_{\mathbb{R}^3} U \cdot \varphi \, dy \, ds.$$

Thus the signed contribution in (12) satisfies

$$- \int_0^T \alpha_n(s) \int_{\mathbb{R}^3} V_n \cdot \varphi \, dy \, ds \rightarrow -a_\infty \int_0^T \int_{\mathbb{R}^3} U \cdot \varphi \, dy \, ds.$$

(vi) *Drift term with a : part involving $y \cdot \nabla V_n$.* Since $\varphi \in C_c^\infty$, integrating by parts in y gives

$$\int_{\mathbb{R}^3} (y \cdot \nabla V_n) \cdot \varphi = - \int_{\mathbb{R}^3} V_n \cdot (y \cdot \nabla \varphi + 3\varphi) \, dy.$$

Define

$$Y_n^1(s) := \int_{\mathbb{R}^3} V_n \cdot (y \cdot \nabla \varphi + 3\varphi) \, dy, \quad Y^1(s) := \int_{\mathbb{R}^3} U \cdot (y \cdot \nabla \varphi + 3\varphi) \, dy.$$

For each s , Cauchy–Schwarz gives

$$|Y_n^1(s)| \leq \|V_n(s)\|_{L^2(K_\varphi)} \|y \cdot \nabla \varphi(s) + 3\varphi(s)\|_{L^2(K_\varphi)}.$$

The $L^\infty(0, T; L^2(K_\varphi))$ bound in Proposition 7.14 gives

$$\sup_n \|Y_n^1\|_{L^\infty(0, T)} \leq C_1 < \infty.$$

Also,

$$\|Y_n^1 - Y^1\|_{L^2(0, T)} \leq \|V_n - U\|_{L^2(0, T; L^2(K_\varphi))} \|y \cdot \nabla \varphi + 3\varphi\|_{L^\infty(0, T; L^2(K_\varphi))} \rightarrow 0.$$

Therefore

$$-\int_0^T \alpha_n(s) \int_{\mathbb{R}^3} (y \cdot \nabla V_n) \cdot \varphi \, dy \, ds \rightarrow a_\infty \int_0^T \int_{\mathbb{R}^3} U \cdot (y \cdot \nabla \varphi + 3\varphi) \, dy \, ds.$$

Integrating by parts once more in the limit gives

$$a_\infty \int_0^T \int_{\mathbb{R}^3} U \cdot (y \cdot \nabla \varphi + 3\varphi) \, dy \, ds = -a_\infty \int_0^T \int_{\mathbb{R}^3} (y \cdot \nabla U) \cdot \varphi \, dy \, ds.$$

(vii) *Term with b .* Since $\beta_n(s)$ is independent of y ,

$$\int_{\mathbb{R}^3} (\beta_n \cdot \nabla V_n) \cdot \varphi = - \int_{\mathbb{R}^3} V_n \cdot (\beta_n \cdot \nabla \varphi) \, dy,$$

so

$$-\int_0^T \int_{\mathbb{R}^3} (\beta_n \cdot \nabla V_n) \cdot \varphi \, dy \, ds = \int_0^T \int_{\mathbb{R}^3} V_n \cdot (\beta_n \cdot \nabla \varphi) \, dy \, ds.$$

Let

$$\begin{aligned} Z_n(s) &:= \left(\int_{\mathbb{R}^3} V_n \cdot \partial_1 \varphi \, dy, \int_{\mathbb{R}^3} V_n \cdot \partial_2 \varphi \, dy, \int_{\mathbb{R}^3} V_n \cdot \partial_3 \varphi \, dy \right), \\ Z(s) &:= \left(\int_{\mathbb{R}^3} U \cdot \partial_1 \varphi \, dy, \int_{\mathbb{R}^3} U \cdot \partial_2 \varphi \, dy, \int_{\mathbb{R}^3} U \cdot \partial_3 \varphi \, dy \right). \end{aligned}$$

Then

$$\int_{\mathbb{R}^3} V_n \cdot (\beta_n \cdot \nabla \varphi) \, dy = \beta_n(s) \cdot Z_n(s), \quad \int_{\mathbb{R}^3} U \cdot (b_\infty \cdot \nabla \varphi) \, dy = b_\infty \cdot Z(s).$$

Consequently

$$\left| \int_0^T (\beta_n - b_\infty) \cdot Z_n \, ds \right| \leq \|\beta_n - b_\infty\|_{L^1(0,T)} \sup_{s \in (0,T)} |Z_n(s)| \rightarrow 0,$$

because

$$|Z_n(s)| \leq \|V_n(s)\|_{L^2(K_\varphi)} \left(\sum_{j=1}^3 \|\partial_j \varphi(s)\|_{L^2(K_\varphi)}^2 \right)^{1/2}$$

and Proposition 7.14 bounds the right-hand side uniformly in n and s . Furthermore,

$$\|Z_n - Z\|_{L^2(0,T;\mathbb{R}^3)} \leq \|V_n - U\|_{L^2(0,T;L^2(K_\varphi))} \left(\sum_{j=1}^3 \|\partial_j \varphi\|_{L^\infty(0,T;L^2(K_\varphi))}^2 \right)^{1/2} \rightarrow 0.$$

Therefore

$$\int_0^T b_\infty \cdot Z_n(s) \, ds \rightarrow \int_0^T b_\infty \cdot Z(s) \, ds.$$

Hence

$$-\int_0^T \int_{\mathbb{R}^3} (\beta_n \cdot \nabla V_n) \cdot \varphi \, dy \, ds \rightarrow \int_0^T \int_{\mathbb{R}^3} U \cdot (b_\infty \cdot \nabla \varphi) \, dy \, ds.$$

Since b_∞ is constant in y , integration by parts yields

$$\int_0^T \int_{\mathbb{R}^3} U \cdot (b_\infty \cdot \nabla \varphi) \, dy \, ds = - \int_0^T \int_{\mathbb{R}^3} (b_\infty \cdot \nabla U) \cdot \varphi \, dy \, ds.$$

Step 5: assemble the limit. Passing to the limit in (12) in each term (Steps 4(i)–4(vii)) gives

$$\int_0^T \int_{\mathbb{R}^3} (-U \cdot \partial_s \varphi - (U \otimes U) : \nabla \varphi + \nu \nabla U : \nabla \varphi - \Pi \nabla \cdot \varphi - a_\infty (U + y \cdot \nabla U) \cdot \varphi - (b_\infty \cdot \nabla U) \cdot \varphi) \, dy \, ds = 0.$$

Since this holds for every compactly supported vector field φ , equation (11) holds in distributions on $(0, T) \times \mathbb{R}^3$. $\square$

Remark 7.21. The theorem asserts only the compact finite-window autonomous limit from the selected-window estimates. A rigidity contradiction for an arbitrary autonomous limit would require a separate Liouville theorem for the class containing (U, Π) . The retained compact scale-collapse branch used in the terminal proof is handled differently: failures of retained inputs are named local exits, and the remaining retained autonomous case is routed by [theorem 7.37](#) into the scale-rigid discharge.

7.6. Decomposition of alternatives of the scale-collapse branch. The cost argument in [Sections 7.2](#) and [7.4](#) excludes the finite-cost scale-collapse branch. The remaining scale-collapse obstruction is treated by a decomposition of alternatives, following the local estimate and autonomous-reduction argument of the preceding sections. The conclusion separates the excluded branches from the residual alternatives; it does not assert a Liouville theorem for every autonomous negative-drift equation.

Theorem 7.22 (Windowwise retained-core spacetime alternative). *Let (τ_n) be a selected fixed-length terminal window sequence of length $T > 0$ for a represented Type II branch produced by the selected-core construction of the preceding sections, and write*

$$V_n(y, s) := V(y, \tau_n + s), \quad s \in (0, T).$$

Then, after passing to a subsequence, exactly one of the following alternatives holds:

- (i) *the retained spacetime activity floor fails on these selected windows (equivalently, the loss-of-retained-active-core row occurs);*
- (ii) *there exist $R > 0$, $\chi \in C_c^\infty(B_R)$, $\chi \geq 0$, and $\eta > 0$ such that*

$$\int_0^T \int_{\mathbb{R}^3} |V_n(y, s)|^3 \chi(y) dy ds \geq \eta$$

for all sufficiently large n .

Proof. The selected-core construction marks, on each admissible selected window, an inner compact renormalized core carrying positive localized CKN velocity mass on the time set used in the window analysis. If no subsequence admitted a uniform spacetime L^3 lower bound with one fixed inner cutoff, then after shrinking once inside the candidate core every further selected-window subsequence would lose the retained spacetime activity floor. This is exactly the loss-of-retained-active-core row in the ordered local first-failure ledger. Therefore, once that row is excluded, a subsequence exists on which one compact inner cutoff and one positive threshold work uniformly on all remaining windows, yielding alternative (ii). $\square$

Theorem 7.23 (Autonomous modulation selection). *Let a represented scale-collapse branch satisfy the modulation bounds obtained in the repaired gauge construction. Suppose that a sequence of windows $[\tau_n, \tau_n + T]$ is selected in the autonomous regime. Then, after passing to a subsequence, there exist constants $a_\infty \in \mathbb{R}$ and $b_\infty \in \mathbb{R}^3$ such that*

$$a(\tau_n + \cdot) \rightarrow a_\infty \quad \text{in } L^1(0, T), \quad b(\tau_n + \cdot) \rightarrow b_\infty \quad \text{in } L^1(0, T; \mathbb{R}^3).$$

Proof. The autonomous-window selection in the repaired-gauge modulation analysis gives convergence of the modulation functions to constant autonomous parameters on the fixed window length T . Passing to the corresponding subsequence gives the displayed L^1 -convergences. $\square$

Theorem 7.24 (Scale-collapse decomposition of alternatives). *Let a represented Type II branch have genuine scale collapse and nonvanishing selected core. Then exactly one of the following alternatives occurs after passing to selected terminal windows:*

- (i) the branch has finite scale-collapsing cost, which is excluded by Theorem 7.7;
- (ii) the branch has finite absolute cost, which is excluded by Corollary 7.9;
- (iii) the branch admits thick autonomous windows in the sense of Definition 7.11, satisfies the selected-window estimates of Proposition 7.14, and has autonomous modulation convergence;
- (iv) the branch has a thin-drift alternative: no selected fixed-length window carries a uniform positive lower bound for the averaged weighted negative drift;
- (v) the branch has an autonomous-modulation alternative: thick windows exist, but no subsequence has a constant-coefficient L^1 modulation limit in the sense needed for Theorem 7.20;
- (vi) the branch has a compactness alternative: the local estimate package or the retained-core input needed for Proposition 7.14 fails on the selected windows.

Proof. The scale-collapse selection theorem decomposes terminal scale-collapse windows according to the averaged negative-drift mass

$$\frac{1}{T} \int_{\tau_n}^{\tau_n+T} a_-(\tau) M_R(\tau) d\tau.$$

If the accumulated cost is finite, alternatives (i) or (ii) apply according to the cost functional under consideration. If the averaged negative-drift mass stays bounded below on a subsequence, one obtains thick windows; otherwise the branch lies in the thin-drift alternative. On thick windows, either the local estimate package holds, in which case proposition 7.14 applies directly, or the required representation, pressure, local-energy, or rough-core input fails on the selected windows, in which case the branch lies in the compactness alternative. Separately, if the retained-core alternative fails on the selected windows, that is the loss-of-retained-active-core row rather than an input to the upper-estimate package. If the estimates hold but no subsequence has the required constant-coefficient modulation limit, the branch lies in the autonomous-modulation alternative. The remaining case is the thick autonomous alternative. $\square$

Definition 7.25 (Retained compact scale-collapse state). A represented scale-collapse branch is in the retained compact scale-collapse state if, on the selected terminal windows, it has:

- (i) the repaired-gauge representation and local pressure reconstruction of theorems 4.25 and 4.37;
- (ii) the compact-cylinder suitable estimates used in proposition 7.14;
- (iii) thick autonomous selected windows in the sense of definition 7.11, together with a subsequential constant-coefficient L^1 modulation limit in the sense used by theorem 7.20;
- (iv) bounded modulation on the selected compact cylinders;
- (v) a retained active Caffarelli–Kohn–Nirenberg mass floor on some fixed inner selected cylinder;
- (vi) the finite scale-collapsing-cost and finite absolute-cost alternatives have already been removed by theorem 7.7 and corollary 7.9;
- (vii) no unresolved multicore, same-point cascade, separated/exterior, rough-core, autonomous-modulation-failure, or compactness-package first-failure label after the ordered local state-space tests have been applied.

Lemma 7.26 (Failure of retained compact scale-collapse is a named exit). *Let a represented genuine scale-collapse branch with nonvanishing selected core be tested against definition 7.25 after the ordered scale-collapse stratification has been applied. If the branch is not in the retained compact scale-collapse state, then its first failed retained-state clause is one of the named local exits: representation/pressure failure, thin-drift, finite scale-cost or finite absolute-cost exclusion, compactness package first failure routed by theorem 7.34, autonomous-modulation failure, no*

retained active core, multicore/same-point cascade, separated/exterior concentration, or rough-core loss.

Proof. The clauses of [definition 7.25](#) are ordered in the same way as the local tests. Failure of clause (i) is exactly the representation or pressure-reconstruction exit. Failure of (ii) is the compactness alternative in [theorem 7.24](#), which is routed by [theorem 7.34](#). Failure of (iii) is either the thin-drift alternative or the autonomous-modulation alternative in that theorem. Failure of (iv) is a modulation-control exit in the repaired-gauge state-space tests. Failure of (v) means the selected scale-collapse branch carries no retained active Caffarelli–Kohn–Nirenberg core and is therefore not the retained terminal branch. Failure of (vi) is precisely one of the two cost alternatives already excluded by [theorem 7.7](#) and [corollary 7.9](#). Failure of (vii) is, by definition, one of the listed ordered local state-space exits. Thus no failed retained-state input remains unlabelled. $\square$

Definition 7.27 (Selected scale-collapse compactness package). Let (V_n, P_n, a_n, b_n) be a represented selected scale-collapse window sequence on $B_{2R} \times I$, where

$$V_n(y, s) = V(y, \tau_n + s), \quad P_n(y, s) = P(y, \tau_n + s), \quad s \in I,$$

and $I \subseteq \mathbb{R}$ is a fixed model interval. The sequence satisfies the selected scale-collapse compactness package on $B_{2R} \times I$ if the following rows all hold with constants independent of n :

- (i) the repaired-gauge equation, divergence-free condition, local energy inequality, and local pressure reconstruction modulo functions of time hold on the window;
- (ii) there is $C_{R,I} < \infty$ such that

$$\sup_n \left(\|V_n\|_{L^\infty(I; L^2(B_{2R}))} + \|V_n\|_{L^2(I; H^1(B_{2R}))} + \|P_n - c_{n,2R}\|_{L^{3/2}(B_{2R} \times I)} \right) \leq C_{R,I};$$

- (iii) the modulation has a local selected-window bound

$$\sup_n \int_I (|a_n(s)| + |b_n(s)|) ds \leq C_{R,I};$$

- (iv) the retained selected core has a positive inner spacetime L^3 floor: for some $0 < r < R$, a compact interval $J_* \subseteq I$, and $\eta_* > 0$,

$$\int_{J_*} \int_{B_r} |V_n(y, s)|^3 dy ds \geq \eta_*$$

after passing to a tail of the sequence. This is the compact-window retained mass floor supplied by [lemma 7.16](#) and [theorem 9.33](#);

- (v) the critical local upper rows have passed uniformly: for every compact subcylinder $B_\rho \times J \subseteq B_{2R} \times I$ used by the terminal local state-space test, there is $C_{\rho,J} < \infty$, independent of n , such that

$$\sup_n \int_J \int_{B_\rho} |V_n(y, s)|^3 dy ds \leq C_{\rho,J}.$$

Failure of this uniform row is not part of the retained package; it is routed by [lemma 10.78](#).

This package is a local compact-cylinder statement. It contains no whole-space L^3 bound, no global critical profile decomposition, and no assumption that localized dissipation vanishes.

Lemma 7.28 (Spacetime active-core floor is the correct retained nonvanishing hypothesis). *In the selected compactness package, every downstream use of the retained active-core lower bound is implied by the spacetime L^3 floor in [definition 7.27](#). More precisely:*

(a) if $V_n \rightarrow V_\infty$ strongly in L^3_{loc} on a compact subcylinder $Q' := B_r \times J \Subset B_{2R} \times I$, then

$$\int_{Q'} |V_\infty|^3 dy ds \geq \eta_*,$$

provided Q' is the inner cylinder and η_* is the retained floor from item (iv) of [definition 7.27](#); in particular $V_\infty \not\equiv 0$ on Q' ;

- (b) if one passes instead to an occupation-state or Young-state limit, then the same retained spacetime L^3 floor yields a positive visible local occupation coordinate on Q' ;
- (c) no later use of the active-core hypothesis requires an a.e.-in-time L^2 lower floor once the local routing and compactness theorems are formulated with the spacetime L^3 activity coordinates.

Proof. Part (a) is immediate from strong L^3 -convergence on Q' :

$$\int_{Q'} |V_\infty|^3 dy ds = \lim_{n \rightarrow \infty} \int_{Q'} |V_n|^3 dy ds \geq \eta_*.$$

For part (b), choose a nonnegative cutoff $\phi \in C_c^\infty(B_{2R} \times I)$ such that

$$0 \leq \phi \leq 1, \quad \phi \equiv 1 \quad \text{on the retained floor cylinder } Q'.$$

This is possible because $Q' \Subset B_{2R} \times I$. Then

$$\int \phi |V_n|^3 dy ds \geq \int_{Q'} |V_n|^3 dy ds \geq \eta_*,$$

so every occupation-state or Young-state limit records a positive local occupation coordinate. No shrinking of the retained floor cylinder is used. Part (c) is exactly how the later compactness and routing arguments use the retained floor: they need a nonvanishing spacetime activity signal on the retained compact window, not an a.e.-time L^2 persistence statement. $\square$

Lemma 7.29 (Selected compactness package gives local strong compactness). *Assume the selected scale-collapse compactness package of [definition 7.27](#) on $B_{2R} \times I$. Then, after passing to a subsequence,*

$$V_n \rightarrow V_\infty \quad \text{strongly in } L^3(B_R \times I).$$

For the pressure, for every $B_\rho \Subset B_{R_1} \Subset B_R$ and every cutoff $\zeta \in C_c^\infty(B_R)$ with $\zeta \equiv 1$ on B_{R_1} , define

$$P_\infty^{\text{loc}} := \mathcal{R}_i \mathcal{R}_j (\zeta (V_\infty \otimes V_\infty)_{ij}).$$

Then either

$$P_n - c_{n,R_1}(s) - P_\infty^{\text{loc}} \rightarrow 0 \quad \text{strongly in } L^{3/2}(B_\rho \times I),$$

and hence the pressure gradients converge locally modulo functions of time on the retained branch, or the branch has already entered the pressure/exterior row routed by [lemma 9.22](#). On the non-pressure branch, the ordered modulation, represented-equation, and local-energy limit test has the following exhaustive output: either its first failure is the corresponding named modulation, representation, or local-energy row, or there are limiting repaired-gauge coefficients (a_∞, b_∞) and a pressure P_∞ , modulo functions of time, such that $(V_\infty, P_\infty, a_\infty, b_\infty)$ satisfies the represented equation and local energy inequality on every compact subcylinder of $B_R \times I$.

Proof. Choose m so large that

$$H^1(B_R) \Subset L^2(B_R) \hookrightarrow H^{-m}(B_{2R}).$$

The local $L_s^\infty L_y^2$, $L_s^2 H_y^1$, pressure, and modulation bounds in the compactness package allow the repaired-gauge equation to be tested against $H_0^m(B_{2R})$. The same estimates as in [lemma 7.17](#) give

$$\sup_n \|\partial_s V_n\|_{L^1(I; H^{-m}(B_R))} < \infty.$$

Aubin–Lions–Simon applied to

$$H^1(B_R) \Subset L^2(B_R) \hookrightarrow H^{-m}(B_R)$$

therefore gives strong convergence in $L^2(B_R \times I)$, after passing to a subsequence. The uniform $L_s^\infty L_y^2 \cap L_s^2 H_y^1$ bound gives the standard parabolic interpolation estimate

$$\sup_n \|V_n\|_{L^{10/3}(B_R \times I)} < \infty.$$

Interpolating the strong L^2 convergence with this uniform $L^{10/3}$ bound gives strong convergence in $L^p(B_R \times I)$ for every $2 \leq p < 10/3$, in particular for $p = 3$.

For the pressure, fix $B_\rho \Subset B_{R_1} \Subset B_R$ and a cutoff $\zeta \equiv 1$ on B_{R_1} , $\text{supp } \zeta \Subset B_R$. With the local pressure P_∞^{loc} defined in the statement, the pressure atlas gives, modulo functions of time on B_{R_1} ,

$$P_n - c_{n,R_1}(s) - P_\infty^{\text{loc}} = \mathcal{R}_i \mathcal{R}_j (\zeta(V_n \otimes V_n - V_\infty \otimes V_\infty)_{ij}) + H_{n,R_1}.$$

The strong $L^3(B_R \times I)$ convergence gives

$$V_n \otimes V_n \rightarrow V_\infty \otimes V_\infty \quad \text{strongly in } L^{3/2}(B_R \times I),$$

so Calderón–Zygmund continuity gives strong convergence of the local part in $L^{3/2}(B_{R_1} \times I)$. If the mean-zero harmonic remainder H_{n,R_1} does not converge to zero in $L^{3/2}(B_\rho \times I)$, then [lemma 9.22](#) with $(q, p) = (3/2, 3/2)$ assigns the branch to the pressure/exterior row. Otherwise the pressure convergence is strong locally modulo functions of time. In all cases pressure compactness has been obtained by a local decomposition or routed to a named row; no weak lower-semicontinuity of pressure mass is used.

It remains to identify the represented suitable limit on the non-pressure branch. Apply the local modulation-defect routing of [lemma 11.28](#) on a finite cover of I . A visible nonperturbative modulation component is one of the named modulation, scale-drift, active-core, or repaired-gauge rows. Otherwise, after passing to a further subsequence, the window averages converge to constants (a_∞, b_∞) , and either

$$a_n \rightarrow a_\infty, \quad b_n \rightarrow b_\infty \quad \text{in } L^1(I),$$

or the nonconstant coefficient contribution tends to zero in the represented local residual topology.

In the first case, the uniform $L^\infty(I; L^2(B_R))$ bound and the L^1 -coefficient convergence allow the modulation terms to pass to the limit after spatial integration by parts. For example, for a fixed compactly supported test field φ ,

$$\begin{aligned} & \left| \int_I (a_n - a_\infty) \int_{B_R} V_n \cdot (3\varphi + y \cdot \nabla \varphi) dy ds \right| \\ & \leq C_{\varphi,R} \|a_n - a_\infty\|_{L^1(I)} \|V_n\|_{L^\infty(I; L^2(B_R))} \rightarrow 0, \end{aligned}$$

and the b_n -term is identical. The constant-coefficient terms pass by strong local L^2 -convergence. In the perturbative alternative, the definition of the represented residual topology gives the same conclusion. Thus the represented equation passes to $(V_\infty, P_\infty, a_\infty, b_\infty)$. If this passage is unavailable, the first unavailable coordinate is the named modulation or represented-equation row.

For the local energy inequality, fix nested compact subcylinders and a nonnegative compact cutoff. The cubic velocity terms converge by strong L^3 -convergence, the pressure-velocity term

converges by the strong $L^{3/2}$ pressure convergence just proved together with strong L^3 -convergence, and the dissipation term is lower semicontinuous under weak L^2H^1 -convergence. In the constant-coefficient alternative, the modulation energy terms pass by the same L^1 -coefficient estimate, now using the uniform $L^\infty L^2$ bound for the quadratic velocity factors. In the perturbative alternative, either the compact-cutoff energy pairings vanish as part of the retained local-energy residual topology, or their first nonvanishing pairing is the named modulation/local-energy row. Hence, away from that named row, the local energy inequality passes to the limit on every smaller cylinder. This proves the final assertion. $\square$

Lemma 7.30 (Regular–singular routing of nonzero selected-window limits). *Assume the selected scale-collapse compactness package of [definition 7.27](#) on $B_{2R} \times I$, and let*

$$Q_* := B_r \times J_* \subseteq B_R \times I$$

be the cylinder carrying the retained spacetime activity floor. After passing to a diagonal subsequence, exactly one of the following ordered alternatives occurs:

- (i) *the pressure/exterior, repaired-chart, modulation, represented-equation, or local-energy limit passage first fails, and the corresponding named local row occurs;*
- (ii) *all of those passages succeed and the nonzero represented suitable limit has a singular point in $B_R \times I^\circ$; then a diagonal parabolic rebase of the full represented sequence produces an actual CKN-active local descendant;*
- (iii) *all of those passages succeed and the limit has no singular point in $B_R \times I^\circ$; then there is a compact parabolic cylinder $Q' \subseteq B_R \times I^\circ$ on which*

$$V_\infty \in L^\infty(Q'), \quad V_\infty \not\equiv 0, \quad V_n \rightarrow V_\infty \quad \text{strongly in } L^3(Q').$$

Thus the third alternative is a nonzero bounded selected-window limit.

In particular, once the named compactness interfaces have passed, a nonzero selected-window limit has no additional local regularity type. This conclusion concerns the full represented solutions V_n ; it neither assumes terminal residual completeness nor applies the Navier–Stokes equation to a residual component separately.

Proof. Apply [lemma 7.29](#). Its first failed limit interface gives alternative (i). Otherwise there are limiting repaired-gauge data $(P_\infty, a_\infty, b_\infty)$ for which $(V_\infty, P_\infty, a_\infty, b_\infty)$ is a represented suitable weak solution on compact subcylinders of $B_R \times I$, and

$$V_n \rightarrow V_\infty \quad \text{strongly in } L^3(B_R \times I).$$

The retained activity floor and [lemma 7.28](#) give

$$\int_{Q_*} |V_\infty|^3 dy ds = \lim_{n \rightarrow \infty} \int_{Q_*} |V_n|^3 dy ds \geq \eta_* > 0.$$

Hence $V_\infty \not\equiv 0$.

Suppose first that $z_0 \in B_R \times I^\circ$ is a singular point of V_∞ . In the strong L^1 -coefficient alternative of [lemma 7.29](#), [proposition 4.34](#) identifies the limiting coefficients with a limiting repaired chart. In its perturbative modulation alternative, the limiting coefficients are the constants (a_∞, b_∞) ; define

$$\Lambda_\infty(s) := \Lambda_0 e^{a_\infty(s-s_0)}, \quad X'_\infty(s) := b_\infty \Lambda_\infty(s).$$

The represented equation and local energy inequality obtained in that lemma are exactly those of this smooth scale-translation chart. Thus in either case the limit has a genuine repaired chart. After shrinking about z_0 , [lemma 5.29](#) applies. If chart identification or the required local comparability is unavailable, its first failed coordinate is the repaired-chart or modulation row in alternative (i).

Choose radii $\rho_k \downarrow 0$ with $Q_{2\rho_k}(z_0) \subseteq B_R \times I^\circ$ and with the preceding chart comparability valid. The contrapositive of [lemma 5.31](#) gives, for every sufficiently large k ,

$$\begin{aligned} \mathcal{C}_\infty(z_0, \rho_k) &:= \rho_k^{-2} \int_{Q_{\rho_k}(z_0)} |V_\infty|^3 dy ds \\ &\quad + \rho_k^{-2} \int_{Q_{\rho_k}(z_0)} |P_\infty - (P_\infty)_{B_{\rho_k}}|^{3/2} dy ds > \varepsilon_{\text{rg}}. \end{aligned}$$

Indeed, smallness at any such radius would make V_∞ bounded on a smaller cylinder containing z_0 , contrary to singularity.

For each fixed k , the strong velocity and normalized-pressure convergence from [lemma 7.29](#) gives

$$\mathcal{C}_n(z_0, \rho_k) \longrightarrow \mathcal{C}_\infty(z_0, \rho_k).$$

Here the pressure term is unchanged by the time-dependent constants in the pressure normalization. If this pressure convergence is unavailable, the pressure/exterior row in alternative (i) has occurred. Otherwise choose integers $n_k > n_{k-1}$ so large that

$$\mathcal{C}_{n_k}(z_0, \rho_k) > \frac{1}{2}\varepsilon_{\text{rg}}.$$

Rebasing the full represented tuple $(V_{n_k}, P_{n_k}, a_{n_k}, b_{n_k})$ on $Q_{\rho_k}(z_0)$ preserves this scale-invariant lower bound. This is an actual CKN-active descendant of the original full branch, proving alternative (ii). Any failure in realizing the rebase is, by the ordered local tests, one of the named interfaces in alternative (i).

It remains to suppose that V_∞ has no singular point in $B_R \times I^\circ$. Every point there is then regular, hence has a parabolic neighborhood on which V_∞ is essentially bounded. Since the integral over Q_* is positive, there is a Lebesgue point $z_1 \in Q_*^\circ$ of V_∞ with $|V_\infty(z_1)| > 0$. Choose a compact parabolic cylinder Q' about z_1 , contained in its regularity neighborhood and in $B_R \times I^\circ$, small enough that $\int_{Q'} |V_\infty|^3 > 0$. Then $V_\infty \in L^\infty(Q')$, $V_\infty \not\equiv 0$ there, and the strong L^3 -convergence restricts to Q' . This is alternative (iii). The singular and nonsingular cases are disjoint and exhaustive, completing the proof. $\square$

Lemma 7.31 (Pressure-only compactness failure is a pressure/exterior row). *Let a selected scale-collapse window sequence satisfy the repaired-gauge, local pressure-reconstruction, and velocity rows of [definition 7.27](#) on $B_{2R} \times I$, and suppose that, after passing to a subsequence,*

$$V_n \rightarrow V_\infty \quad \text{strongly in } L^3(B_R \times I).$$

Fix $B_\rho \subseteq B_{R_1} \subseteq B_R$ and $\zeta \in C_c^\infty(B_R)$ with $\zeta \equiv 1$ on B_{R_1} , and set

$$P_\infty^{\text{loc}} := \mathcal{R}_i \mathcal{R}_j (\zeta (V_\infty \otimes V_\infty)_{ij}).$$

Write the normalized pressure defect on B_{R_1} , modulo functions of time, as

$$Q_n := P_n - c_{n,R_1}(s) - P_\infty^{\text{loc}} = Q_n^{\text{loc}} + H_{n,R_1},$$

where

$$Q_n^{\text{loc}} = \mathcal{R}_i \mathcal{R}_j (\zeta (V_n \otimes V_n - V_\infty \otimes V_\infty)_{ij})$$

and H_{n,R_1} is mean-zero harmonic on B_{R_1} . If the normalized pressure defect has a positive lower bound on the smaller ball,

$$\limsup_{n \rightarrow \infty} \|Q_n\|_{L^{3/2}(B_\rho \times I)} > 0,$$

then the branch is in the ordered pressure/exterior alternative.

Proof. The strong $L^3(B_R \times I)$ convergence gives

$$V_n \otimes V_n - V_\infty \otimes V_\infty \rightarrow 0 \quad \text{strongly in } L^{3/2}(B_R \times I).$$

Calderón–Zygmund continuity gives

$$Q_n^{\text{loc}} \rightarrow 0 \quad \text{strongly in } L^{3/2}(B_{R_1} \times I).$$

Thus the positive lower bound for Q_n on $B_\rho \times I$ forces

$$\limsup_{n \rightarrow \infty} \|H_{n,R_1}\|_{L^{3/2}(B_\rho \times I)} > 0.$$

By [lemma 9.22](#), together with the quantitative annular-tail converse [lemma 9.21](#), such a non-vanishing harmonic part can only come from pressure-reconstruction failure or from a far-field pressure carrier. These are exactly the pressure/exterior rows of the ordered local state space. $\square$

Lemma 7.32 (Failure of selected compactness produces a local defect label). *Let a represented genuine scale-collapse branch have a nonvanishing selected core and be tested on canonical selected windows $B_{2R} \times I$. If the selected scale-collapse compactness package of [definition 7.27](#) fails on those windows, then after passing to a subsequence the first failed row is one of the following local labels:*

- (i) *representation or pressure-reconstruction failure;*
- (ii) *autonomous-modulation failure or modulation-driven recentering;*
- (iii) *loss of retained active core, including pressure-only critical-mass detection;*
- (iv) *rough-core loss of local windowed H^1 control;*
- (v) *exterior or separated-core concentration;*
- (vi) *multicore, comparable-core, or same-point cascade;*
- (vii) *a critical local L^3 upper-mass failure routed by [lemma 10.78](#).*

If none of these labels occurs, the compactness package holds and [lemma 7.29](#) gives strong local compactness on the selected windows.

Proof. The rows in [definition 7.27](#) are tested in the displayed order. Failure of the repaired-gauge equation, local energy inequality, or pressure reconstruction is item (i). Failure of the local selected-window modulation bound is exactly the autonomous-modulation or recentering row, giving item (ii).

If the retained spacetime activity floor fails, the selected candidate is not the retained terminal component. The terminal local critical-mass routing lemma [lemma 10.78](#) records this as lower active-core loss or, when the CKN signal is carried only by pressure, as the pressure-only critical-mass row. This gives item (iii).

Assume the preceding rows hold and the local velocity/pressure estimate row fails. If the failure is the $L_s^2 H_y^1$ part on a retained compact cylinder, then it is the rough-core row by [theorem 11.7](#), giving item (iv). If the failure is that the compact-cylinder $L_s^\infty L_y^2$ or pressure mass cannot be bounded on any fixed enclosing ball, apply a finite cover of the selected retained core and the local pressure atlas. Either a positive velocity concentration remains in one member of the finite cover, or the mass leaves every fixed member. In the first case the ordered local state-space decomposition reads a compact active label: a comparable retained core, multicore label, or same-point cascade. In the second case the label is exterior or separated-core concentration. If the positive diagnostic is pressure-only, [lemma 7.31](#) routes it to the pressure/exterior row. This gives items (iii), (v), and (vi).

It remains to consider failure of the critical local upper row. By [lemma 10.78](#), such a failure produces an additional compact active label at the same retained frame, a comparable finite-mass retained class, a multibubble/multicore label, a cascade label, a rough-core label, or an

exterior/separated label. These are exactly items (iv)–(vii). Comparable same-point classes are not left as a separate compactness defect: they are grouped by [lemma 10.77](#) and [lemma 9.3](#).

If no listed first failure occurs, all rows in the selected scale-collapse compactness package hold. The local strong compactness conclusion then follows from [lemma 7.29](#). Thus selected-window compactness failure has no unlabelled terminal form. $\square$

Lemma 7.33 (Comparable compact defects are grouped before retention). *Suppose a selected compactness defect produced by [lemma 7.32](#) is based at the same physical point as the retained selected scale-collapse core and has scale comparable to the retained frame. Then either it is absorbed into the retained compound core before the retained scale-collapse state is declared, or the branch has already entered the multicore/cascade/exterior local alternative.*

Proof. If the relative scale-center parameters stay in a compact subset of the same-point comparable chart, [lemma 9.3](#) assembles the corresponding positive local states into one compound state, and [lemma 10.77](#) says that any surviving comparable retained class is already absorbed into the selected retained component. If the relative parameters leave every compact subset, the ordered state-space compactification records separated scale, separated center, exterior leakage, or same-point cascade before the retained compact scale-collapse state is reached. Hence a comparable compactness defect cannot remain as a separate unresolved retained compactness failure. $\square$

Theorem 7.34 (Selected compactness failure is routed). *Let a represented genuine scale-collapse branch have nonvanishing selected core and be tested on the canonical selected windows supplied by the local state-space construction. Then the compactness alternative in [theorem 7.24](#) is not an independent terminal exit. More precisely, either the selected scale-collapse compactness package holds on every retained compact subcylinder after passing to a subsequence, or the first failed row is routed to one of the named local alternatives: representation/pressure failure, autonomous-modulation failure, loss of retained active core or pressure-only critical mass, rough-core loss, exterior/separated-core concentration, multicore/same-point cascade, or a critical local upper-mass label routed by [lemma 10.78](#). On the branch where the package and its limit interfaces pass, the resulting nonzero compact limit is either singular and produces an actual CKN-active descendant, or is regular and produces a nonzero bounded selected-window limit. There is no further compact-limit alternative.*

Proof. Choose a countable nested exhaustion of the retained compact selected core by compact cylinders

$$B_{R_j} \times I_j \subseteq B_{2R_j} \times I_j, \quad j = 1, 2, \dots$$

Apply [lemma 7.32](#) to each member of this exhaustion in the ordered local test. If a row fails on some member, take the least j for which a failure occurs and then the first failed row in the ordered list on that cylinder. The lemma assigns it to one of the named local alternatives. If comparable compact defects occur at the retained physical point and comparable scale, then [lemma 7.33](#) absorbs them into the retained compound core or routes them to the multicore/cascade/exterior rows.

If no row fails on the exhaustion, the package holds on each compact subcylinder. Apply [lemma 7.29](#) on the first cylinder, then on the second along the extracted subsequence, and continue diagonally. The diagonal subsequence gives strong local L^3 compactness and pressure compactness modulo functions of time, with any nonvanishing harmonic pressure remainder routed by [lemma 9.22](#). Therefore compactness failure is represented by a local first-failure classification.

On any retained cylinder carrying the selected spacetime activity floor, apply [lemma 7.30](#). A failed limit interface is another named first-failure row. Otherwise its singular case produces

an actual active descendant of the full branch, while its nonsingular case produces a nonzero bounded selected-window limit. Thus the pass branch either returns to the ordinary active selection, enters the bounded-window route, or enters a more specific named local row; it is not a separate terminal obstruction. $\square$

Lemma 7.35 (Retained autonomous scale collapse enters the scale-rigid stratum). *Let a represented Type II branch be in the retained compact scale-collapse state of [definition 7.25](#). Suppose it has thick autonomous windows and satisfies [proposition 7.14](#) on the same window sequence. Then the shifted selected-window sequence*

$$V_n(y, s) = V(y, \tau_n + s), \quad P_n(y, s) = P(y, \tau_n + s),$$

restricted to a fixed compact cylinder carrying the retained active mass, is a scale-rigid local branch in the sense of [definition 5.19](#), or it is the compact single-core branch excluded by [theorem 3.18](#), unless the branch has already entered one of the named local alternatives excluded from the retained compact state.

Proof. The compact-cylinder suitable bounds in [definition 7.25\(ii\)](#) give [definition 5.19\(R1\)](#) on a fixed compact selected cylinder after shrinking once. The retained active CKN mass floor in (v) gives (R2) on an inner cylinder. Bounded modulation in (iv) gives (R3).

It remains to verify the scale-rigid structural clause (R4). Apply the minimal active-family extraction [theorem 5.11](#) to the selected retained core and to the nearest previously selected active scale-center, if such a scale-center is present. If no such active relation is present, the branch is the compact single-core state rather than a genuine scale-collapse residual; this is the single-core branch excluded by [theorem 3.18](#), not a remaining case in this lemma.

Otherwise the pair is an input to the same-point or separated-point reduction, [theorems 5.20](#) and [5.21](#). Those reductions are an ordered partition. Their non-rigid outputs are exactly the named exits: another active core at comparable scale, a strict same-point cascade not yet absorbed, a separated physical point, an exterior pressure/concentration label, a rough-core label, or failure of the compact estimates/modulation convergence needed for the autonomous scale-collapse passage. The comparable, same-point, separated, and exterior outcomes are the multicore, same-point cascade, or separated/exterior alternatives of [definition 5.6](#); the rough-core outcome is the local windowed- H^1 failure alternative; and the last two outcomes are the compactness or autonomous-modulation alternatives of [theorem 7.24](#). All are excluded from the retained compact state by [definition 7.25\(vii\)](#).

After those exits are removed, the same ordered reductions leave only the structural outputs named in [definition 5.19\(R4\)](#). In the compact degenerate case, Step 3 of [theorem 5.20](#) records bounded (ρ_n, z_n) and $\sigma_{\min}(\mathcal{J}_n) \rightarrow 0$, which is (R4a). In the strict same-point case, Step 4 records $\rho_n \rightarrow \infty$, $|z_n|/\rho_n = O(1)$, and applies [theorem 5.18](#) after the outer states have been absorbed; this supplies all three residual convergences in (R4b). Thus the shifted sequence satisfies the exact structural input contract on the selected compact cylinder. $\square$

Theorem 7.36 (Scale-collapse first failures are routed local alternatives). *Let a represented branch have genuine scale collapse and a nonvanishing selected core. If one of the retained compact scale-collapse inputs fails before the branch reaches [definition 7.25](#), then the first failure is one of the named local alternatives: finite scale cost, finite absolute scale cost, autonomous-modulation failure, compactness-package first failure routed to a named local row, exterior/separated-core loss, rough-core loss, multicore/same-point cascade, or scale-rigid residual behavior; the thin-drift first failure is routed by [theorem 10.118](#) to one of these rows or to a closed retained row. None of these first failures remains inside the retained compact scale-collapse branch.*

Proof. Apply the ordered scale-collapse stratification [theorem 7.24](#) and the first-failure ledger of [lemma 7.26](#). Finite scale cost and finite absolute scale cost are the cost alternatives excluded by

[theorem 7.7](#) and [corollary 7.9](#). Thin drift is first recorded as the scale-drift boundary face of the scale-collapse stratification, but [theorem 10.118](#) proves that it is not terminal: persistent windows return to the thick-window ledger, while diffuse occupation blocks route to another named row or to a closed retained row. Failure of autonomous modulation is a repaired-gauge/modulation boundary face. Failure of compactness on the selected windows is routed by [theorem 7.34](#): either the local compactness package holds after passing to a subsequence, or the first failed row is a representation/pressure row, active-core or pressure-only row, rough-core row, exterior/separated-core row, multicore/same-point cascade row, or critical local upper-mass row. Comparable same-point defects are grouped by [lemma 7.33](#). Thus compactness failure is not a separate retained terminal state. The remaining rigid residual is the scale-rigid local branch discharged by [proposition 5.42](#). These are exactly the named local alternatives used in the state-space decomposition, so none is a retained compact scale-collapse state. $\square$

Theorem 7.37 (Compact scale-collapse rigidity from local routing). *Let a represented Type II branch have genuine scale collapse, nonvanishing selected core, and belong to the retained compact scale-collapse state of [definition 7.25](#). Then no compact scale-collapse branch remains. In particular, compact scale-collapse rigidity is supplied by the retained local state-space branch itself, not by an independent stationary input.*

Proof. Apply the scale-collapse stratification [theorem 7.24](#). The finite scale-collapsing cost and finite absolute cost cases are excluded by [theorem 7.7](#) and [corollary 7.9](#).

The thin-drift, autonomous-modulation, and compactness alternatives are not discarded by definition. A branch satisfying the retained compact scale-collapse-state hypotheses has thick autonomous selected windows, selected-window estimates, and autonomous modulation convergence by [definition 7.25\(ii\),\(iii\)](#). Thus the thin-drift, autonomous-modulation, and compactness alternatives of [theorem 7.24](#) are incompatible with the present retained state.

It remains to consider the thick autonomous case. By [lemma 7.35](#), the shifted selected-window sequence is either a scale-rigid local branch or the compact single-core branch, unless the branch has already left the retained compact state. Since the present branch is retained, the exit case is unavailable. If the compact single-core branch occurs, it is excluded by [theorem 3.18](#). If the scale-rigid branch occurs, [proposition 5.42](#) applies: the branch either has a nonzero bounded selected-window limit, which leaves the strict selected-window retained branch by [theorem 5.27](#), or it reduces to the compact single-core branch, again excluded by [theorem 3.18](#). Hence the thick autonomous retained case is impossible. $\square$

Theorem 7.38 (Scale-collapse alternatives). *Let a represented Type II branch have genuine scale collapse and a nonvanishing selected core. Then one of the following mutually exclusive outcomes occurs:*

- (i) *the finite scale-collapsing cost exclusion applies;*
- (ii) *the finite absolute cost exclusion applies;*
- (iii) *the branch has a thin-drift alternative, which is routed by [theorem 10.118](#);*
- (iv) *the branch has an autonomous-modulation alternative;*
- (v) *the branch has a compactness-package first failure, which is routed by [theorem 7.34](#);*
- (vi) *the branch enters the retained compact scale-collapse state, in which case [theorem 7.37](#) excludes it.*

Proof. Apply [theorem 7.24](#). The finite-cost alternatives are items (i) and (ii). The thin-drift, modulation, and compactness alternatives are items (iv)–(vi) of that theorem; the thin-drift item is routed by [theorem 10.118](#), and the compactness item is routed by [theorem 7.34](#). It remains only to treat item (iii), the thick autonomous branch with selected-window estimates and autonomous

modulation convergence. If one of the retained compact-state inputs in [definition 7.25](#) fails, then [theorem 7.36](#) assigns the branch to a named local alternative, hence to one of the preceding nonretained outcomes. Otherwise the branch is in the retained compact scale-collapse state and [theorem 7.37](#) excludes it. $\square$

8. CRITICAL TIGHTNESS AND WINDOWED-CONTROL ESTIMATES

This section records the local estimates used to exclude failure of uniform critical tightness and failure of local windowed H^1 control. The statements are phrased as analytic alternatives for a renormalized Navier–Stokes solution $V(\tau)$.

8.1. Uniform critical tightness.

Definition 8.1 (Uniform critical tightness). A renormalized orbit $V(\tau)$, $\tau \geq \tau_0$, is uniformly critically tight if

$$\forall \varepsilon > 0 \exists R_\varepsilon < \infty : \sup_{\tau \geq \tau_0} \int_{|y| > R_\varepsilon} |V(y, \tau)|^3 dy < \varepsilon.$$

The non-tight alternative is the failure of this property.

Lemma 8.2 (Critical precompactness implies uniform tightness). *If*

$$\mathcal{O} := \{V(\tau) : \tau \geq \tau_0\}$$

is contained in a compact subset of $L^3(\mathbb{R}^3)$, then V is uniformly critically tight.

Proof. Compact subsets of $L^3(\mathbb{R}^3)$ are uniformly tight. Let $K \subset L^3$ be compact and fix $\varepsilon > 0$. Set $\delta = \varepsilon^{1/3}/2$. Choose a finite δ -net $f_1, \dots, f_N$ in L^3 . For each i , choose R_i so that

$$\|f_i\|_{L^3(|y| > R_i)} < \delta.$$

Let $R = \max_i R_i$. For any $f \in K$, choose i with $\|f - f_i\|_3 < \delta$. Then

$$\|f\|_{L^3(|y| > R)} \leq \|f - f_i\|_3 + \|f_i\|_{L^3(|y| > R)} < \varepsilon^{1/3}.$$

Taking the cube gives

$$\int_{|y| > R} |f(y)|^3 dy < \varepsilon.$$

Applying this to the compact set containing $\mathcal{O}$ gives uniform critical tightness. $\square$

Definition 8.3 (Compact core with vanishing critical remainder). The orbit has a compact critical core with vanishing remainder if there exist a compact set $\mathcal{K}^c \subset L^3(\mathbb{R}^3)$, a map $Q : [\tau_0, \infty) \rightarrow \mathcal{K}^c$, and a remainder $r(\tau) \in L^3(\mathbb{R}^3)$, such that

$$V(\tau) = Q(\tau) + r(\tau), \quad \lim_{\tau \rightarrow \infty} \|r(\tau)\|_{L^3} = 0,$$

and, for every $T > \tau_0$, the set

$$\{V(\tau) : \tau \in [\tau_0, T]\}$$

is uniformly L^3 -tight.

Lemma 8.4 (Compact core plus vanishing remainder implies tightness). *If [definition 8.3](#) holds, then V is uniformly critically tight.*

Proof. Fix $\varepsilon > 0$ and set $\delta = \varepsilon^{1/3}$. By compactness of $\mathcal{K}^c$, apply [lemma 8.2](#) to $\mathcal{K}^c$ and choose R_1 such that

$$\sup_{q \in \mathcal{K}^c} \|q\|_{L^3(|y| > R_1)} < \delta/3.$$

Since $\|r(\tau)\|_3 \rightarrow 0$, choose T so that

$$\|r(\tau)\|_3 < \delta/3 \quad \text{for all } \tau \geq T.$$

Then for $\tau \geq T$,

$$\|V(\tau)\|_{L^3(|y| > R_1)} \leq \|Q(\tau)\|_{L^3(|y| > R_1)} + \|r(\tau)\|_3 < 2\delta/3.$$

By finite-initial-window tightness on $[\tau_0, T]$, choose R_2 such that

$$\sup_{\tau \in [\tau_0, T]} \int_{|y| > R_2} |V(y, \tau)|^3 dy < \varepsilon.$$

With $R = \max(R_1, R_2)$, the L^3 -norm of the tail of $V(\tau)$ is $< \delta$ for all $\tau \geq T$, while the tail integral is $< \varepsilon$ on $[\tau_0, T]$. Cubing on the tail $\tau \geq T$ gives the required L^3 -tail integral $< \varepsilon$ for all $\tau \geq \tau_0$. $\square$

Definition 8.5 (No-splitting profile condition). A critical profile decomposition has no critical-mass splitting if it has:

- (i) a primary compact core family $\mathcal{K}^c \subset L^3$;
- (ii) asymptotic L^3 -decoupling of core, secondary profiles, and remainder;
- (iii) single-profile saturation, so every non-primary profile has zero critical L^3 -mass;
- (iv) a remainder whose L^3 -norm vanishes on the renormalized tail;
- (v) finite-initial-window tightness.

Lemma 8.6 (No splitting gives compact core with vanishing remainder). *If definition 8.5 holds, then definition 8.3 holds.*

Proof. By single-profile saturation and nonnegative L^3 -mass decoupling, every secondary profile has zero critical L^3 -mass and therefore vanishes in L^3 . The profile decomposition then reduces to a primary compact core plus a remainder whose L^3 -norm tends to zero on the renormalized tail. These provide the core $Q(\tau)$, compact set $\mathcal{K}^c$, and remainder $r(\tau)$ required by definition 8.3, with finite-initial-window tightness supplied as part of definition 8.5. $\square$

Corollary 8.7 (No splitting gives uniform critical tightness). *If definition 8.5 holds, then V is uniformly critically tight.*

Proof. Lemma 8.6 gives a compact core with vanishing critical remainder, and lemma 8.4 gives uniform critical tightness. $\square$

Definition 8.8 (Finite critical profile approximation). The orbit admits a finite critical profile approximation if there are finitely many profiles

$$\mathcal{B}_{NS}^{II} := \{B_1, \dots, B_N\} \subset L^3(\mathbb{R}^3)$$

such that, for every $\delta > 0$, there is τ_δ with

$$\forall \tau \geq \tau_\delta \exists i(\tau) \in \{1, \dots, N\} : \|V(\tau) - B_{i(\tau)}\|_{L^3} < \delta,$$

and the finite initial window

$$\{V(\tau) : \tau \in [\tau_0, \tau_\delta]\}$$

is uniformly L^3 -tight for each $\delta > 0$.

Lemma 8.9 (Finite critical profile approximation implies tightness). *If definition 8.8 holds, then V is uniformly critically tight.*

Proof. Fix $\varepsilon > 0$ and set $\delta = \varepsilon^{1/3}$. Since the profile family is finite and each $B_i \in L^3(\mathbb{R}^3)$, there is R_1 such that

$$\max_{1 \leq i \leq N} \|B_i\|_{L^3(|y| > R_1)} < \delta/3.$$

Choose the approximation tolerance $\delta/3$. For $\tau \geq \tau_{\delta/3}$, pick $i(\tau)$ with $\|V(\tau) - B_{i(\tau)}\|_3 < \delta/3$. Then

$$\|V(\tau)\|_{L^3(|y| > R_1)} \leq \|V(\tau) - B_{i(\tau)}\|_3 + \|B_{i(\tau)}\|_{L^3(|y| > R_1)} < 2\delta/3.$$

By finite-initial-window tightness, choose R_2 so that the L^3 tail is $< \varepsilon$ on $[\tau_0, \tau_{\delta/3}]$. With $R = \max(R_1, R_2)$, the tail integral is $< \varepsilon$ for every $\tau \geq \tau_0$, so the orbit is uniformly L^3 -tight. $\square$

Definition 8.10 (Compact parameter realization). The orbit has a compact parameter realization if there exist a compact parameter set Θ_{II} , a continuous map

$$\Theta_{II} \ni \theta \mapsto V_\theta \in L^3(\mathbb{R}^3),$$

a curve $\theta(\tau) \in \Theta_{II}$, and a remainder $r(\tau) \rightarrow 0$ in L^3 , such that

$$V(\tau) = V_{\theta(\tau)} + r(\tau),$$

and, for every $T > \tau_0$, the set

$$\{V(\tau) : \tau \in [\tau_0, T]\}$$

is uniformly L^3 -tight.

Lemma 8.11 (Compact realization implies tightness). *If [definition 8.10](#) holds, then V is uniformly critically tight.*

Proof. The continuous image of compact Θ_{II} in L^3 is compact. Thus

$$\mathcal{K}^c := \{V_\theta : \theta \in \Theta_{II}\}$$

is compact in L^3 . The representation $V(\tau) = V_{\theta(\tau)} + r(\tau)$ with $r(\tau) \rightarrow 0$ in L^3 gives the compact-core vanishing-remainder situation, because finite-initial-window tightness is included in [definition 8.10](#). [Lemma 8.4](#) gives uniform critical tightness. $\square$

Theorem 8.12 (Exclusion of the non-tight alternative). *For a renormalized Type II Navier–Stokes solution, any one of the following conditions gives uniform critical tightness:*

- (i) $\mathcal{O} \in L^3(\mathbb{R}^3)$;
- (ii) compact core with vanishing critical remainder;
- (iii) no-splitting profile condition;
- (iv) finite critical profile approximation;
- (v) compact parameter realization.

Proof. The five alternatives are handled respectively by [lemmas 8.2, 8.4, 8.9](#) and [8.11](#) and [corollary 8.7](#). In every case V is uniformly critically tight, so the failure of uniform critical tightness is unavailable. $\square$

Corollary 8.13 (After exclusion of the non-tight alternative). *Assume a finite-cost renormalized Type II solution not excluded by the compact cost-divergence criterion has a two-alternative classification whose only remaining alternatives are failure of uniform critical tightness and failure of tail windowed H^1 control. If [definition 8.1](#) holds, then the only remaining alternative is failure of tail windowed H^1 control.*

Proof. The two-alternative classification gives exactly one of the two alternatives. Uniform critical tightness rules out failure of tightness, so only the windowed H^1 -failure alternative remains. $\square$

8.2. Tail windowed H^1 control.

Lemma 8.14 (Uniform local pressure control). *Assume*

$$M_{3,R} := \sup_{\tau \geq \tau_0} \|V(\tau)\|_{L^3(B_{4R})} < \infty$$

and

$$-\Delta P = \partial_i \partial_j (V_i V_j) \quad \text{in } \mathbb{R}^3$$

for each $\tau \geq \tau_0$. Assume also the local tail estimate

$$\mathcal{T}_R(\tau) := \sup_{x \in B_R} \int_{|z| > 2R} \frac{|V(z, \tau)|^2}{|x - z|^5} dz, \quad \sup_{\tau \geq \tau_0} \mathcal{T}_R(\tau) \leq CR^{-4} M_{3,R}^2.$$

Then there is a measurable choice of constants $c_R(\tau) \in \mathbb{R}$ such that

$$\sup_{\tau \geq \tau_0} \|P(\tau) - c_R(\tau)\|_{L^{3/2}(B_R)} \leq C_R M_{3,R}^2.$$

Proof. The local pressure decomposition gives

$$\|P(\tau) - c_R(\tau)\|_{L^{3/2}(B_R)} \leq C \left(\|V(\tau)\|_{L^3(B_{4R})}^2 + R^2 \mathcal{T}_R(\tau) \right).$$

The tail hypothesis gives

$$\sup_{\tau \geq \tau_0} \mathcal{T}_R(\tau) \leq C R^{-4} M_{3,R}^2,$$

and the local critical bound gives

$$\|V(\tau)\|_{L^3(B_{4R})} \leq M_{3,R}.$$

Combining the two estimates yields

$$\sup_{\tau \geq \tau_0} \|P(\tau) - c_R(\tau)\|_{L^{3/2}(B_R)} \leq C_R M_{3,R}^2.$$

The constants $c_R(\tau)$ can be chosen by fixing the displayed pressure decomposition, hence measurably. $\square$

Proposition 8.15 (Renormalized Caccioppoli gives uniform windowed gradients). *Assume V, P are smooth on compact renormalized cylinders and solve*

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(\tau)(V + y \cdot \nabla V) + b(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0.$$

Assume local compact-cylinder critical control on the outer cylinders used in the estimate, bounded modulation

$$M_{ab} := \sup_{\tau \geq \tau_0} (|a(\tau)| + |b(\tau)|) < \infty,$$

and pressure normalized as in [lemma 8.14](#). Then for every $R > 0$ there exists

$$C_R = C(R, \nu, M_{3,R}, M_{ab})$$

such that

$$\sup_{T \geq \tau_0 + 1} \int_T^{T+1} \int_{B_R} |\nabla V(y, \tau)|^2 dy d\tau \leq C_R.$$

Proof. Fix $R > 0$ and $T \geq \tau_0 + 1$. Use the renormalized Caccioppoli estimate with outer radius $2R$, inner radius R , and time intervals

$$I_{2R} := (T - 1, T + 2), \quad I_R := [T, T + 1].$$

Thus $R_{\text{out}} - R_{\text{in}} = R$ and the inner time length is 1. Keeping the pressure term before taking absolute values gives

$$\begin{aligned} \nu \int_T^{T+1} \int_{B_R} |\nabla V|^2 &\leq C_R \iint_{(T-1, T+2) \times B_{2R}} |V|^2 + C_R \iint_{(T-1, T+2) \times B_{2R}} |V|^3 \\ &\quad + C_R \mathcal{P}_{T,R} + C_R M_{ab} \iint_{(T-1, T+2) \times B_{2R}} |V|^2, \end{aligned}$$

where C_R absorbs the powers of R^{-1} , R^{-2} , and ν , and

$$\mathcal{P}_{T,R} := \left| \iint_{(T-1, T+2) \times B_{2R}} P V \cdot \nabla \zeta \right|.$$

The L^2 -term is controlled by the local critical bound and boundedness of B_{2R} :

$$\int_{B_{2R}} |V(\tau)|^2 dy \leq |B_{2R}|^{1/3} \|V(\tau)\|_{L^3(B_{2R})}^2 \leq C_R M_{3,R}^2.$$

Since the outer time interval has length 3,

$$\iint_{(T-1, T+2) \times B_{2R}} |V|^2 \leq C_R M_{3,R}^2,$$

and similarly

$$\iint_{(T-1, T+2) \times B_{2R}} |V|^3 \leq 3M_{3,R}^3.$$

For any time-dependent constant $c(\tau)$,

$$\int_{B_{2R}} c(\tau) V \cdot \nabla \zeta dy = - \int_{B_{2R}} c(\tau) \zeta \nabla \cdot V dy = 0.$$

Thus P may be replaced by $P - c_{2R}(\tau)$. Hölder's inequality gives

$$\mathcal{P}_{T,R} \leq C_R \iint_{(T-1, T+2) \times B_{2R}} |P - c_{2R}(\tau)| |V| dy d\tau.$$

For almost every τ ,

$$\int_{B_{2R}} |P - c_{2R}(\tau)| |V(\tau)| dy \leq \|P(\tau) - c_{2R}(\tau)\|_{L^{3/2}(B_{2R})} \|V(\tau)\|_{L^3(B_{2R})}.$$

The pressure estimate with radius $2R$ gives

$$\|P(\tau) - c_{2R}(\tau)\|_{L^{3/2}(B_{2R})} \leq C_R M_{3,R}^2,$$

and the local critical bound gives

$$\|V(\tau)\|_{L^3(B_{2R})} \leq M_{3,R}.$$

Hence $\mathcal{P}_{T,R} \leq C_R M_{3,R}^3$. Substituting these estimates into Caccioppoli yields

$$\nu \int_T^{T+1} \int_{B_R} |\nabla V|^2 \leq C_R (M_{3,R}^2 + M_{3,R}^3 + M_{ab} M_{3,R}^2),$$

uniformly in $T \geq \tau_0 + 1$. Dividing by ν and renaming the constant proves the claim. $\square$

Corollary 8.16 (Tail windowed local H^1 bound). *Under the hypotheses of [proposition 8.15](#), for every $R > 0$,*

$$\sup_{T \geq \tau_0 + 1} \int_T^{T+1} \|V(\tau)\|_{H^1(B_R)}^2 d\tau < \infty.$$

Proof. The gradient part is [proposition 8.15](#). The L^2 -part follows from the same bounded-set Hölder estimate:

$$\int_T^{T+1} \|V(\tau)\|_{L^2(B_R)}^2 d\tau \leq C_R \int_T^{T+1} \|V(\tau)\|_{L^3(B_{2R})}^2 d\tau \leq C_R M_{3,R}^2.$$

Adding the two bounds gives the asserted $H^1(B_R)$ estimate. $\square$

Corollary 8.17 (Exclusion of failure of local windowed H^1 control). *Assume the renormalized Type II solution satisfies the repaired-gauge renormalized Navier–Stokes equation on compact cylinders, local compact-cylinder critical control*

$$\sup_{\tau \geq \tau_0} \|V(\tau)\|_{L^3(B_{2R})} < \infty$$

for each relevant R , bounded modulation

$$\sup_{\tau \geq \tau_0} (|a(\tau)| + |b(\tau)|) < \infty,$$

and pressure reconstruction

$$-\Delta P = \partial_i \partial_j (V_i V_j).$$

Then failure of local windowed H^1 control,

$$\exists m \geq 1 : \sup_{T \geq \tau_0+1} \int_T^{T+1} \|V(\tau)\|_{H^1(B_m)}^2 d\tau = \infty$$

is impossible.

Proof. Apply [corollary 8.16](#) with $R = m$. The resulting finite bound contradicts the displayed failure of local windowed H^1 control. $\square$

8.3. Modulation and weighted forcing inputs.

Definition 8.18 (Modulation matrix Schur data). Write the repaired modulation matrix in block form

$$M(V) = \begin{pmatrix} \alpha(V) & u(V)^T \\ v(V) & A(V) \end{pmatrix}.$$

Assume

$$\begin{aligned} \alpha(V(\tau)) &\geq \alpha_0 > 0, & \|A(V(\tau))^{-1}\|_{\infty \rightarrow \infty} &\leq C_A, \\ \|u(V(\tau))\|_\infty &\leq C'_u, & \|v(V(\tau))\|_\infty &\leq C'_v, \end{aligned}$$

and

$$S(V(\tau)) := \alpha(V(\tau)) - u(V(\tau))^T A(V(\tau))^{-1} v(V(\tau)) \geq \sigma_0 > 0.$$

Lemma 8.19 (Uniform inverse from Schur data). *Under [definition 8.18](#),*

$$\sup_{\tau \geq \tau_0} \|M(V(\tau))^{-1}\|_{\infty \rightarrow \infty} \leq \sigma_0^{-1} + 3\sigma_0^{-1} C'_u C_A + \sigma_0^{-1} C_A C'_v + C_A + 3C_A^2 C'_v C'_u \sigma_0^{-1}.$$

Proof. For each τ , write

$$M(V(\tau)) = \begin{pmatrix} \alpha & u^T \\ v & A \end{pmatrix}, \quad S = \alpha - u^T A^{-1} v.$$

By the translation-block inverse bound, A^{-1} exists and is uniformly bounded. By the Schur gap, S^{-1} exists and $|S^{-1}| \leq \sigma_0^{-1}$. The block inverse formula gives

$$M^{-1} = \begin{pmatrix} S^{-1} & -S^{-1} u^T A^{-1} \\ -A^{-1} v S^{-1} & A^{-1} + A^{-1} v S^{-1} u^T A^{-1} \end{pmatrix}.$$

Using the ℓ^∞ matrix norm and the finite dimension of the mixed blocks,

$$\begin{aligned} \|S^{-1} u^T A^{-1}\|_{\infty \rightarrow \infty} &\leq 3\sigma_0^{-1} C'_u C_A, \\ \|A^{-1} v S^{-1}\|_{\infty \rightarrow \infty} &\leq \sigma_0^{-1} C_A C'_v, \end{aligned}$$

and

$$\|A^{-1} v S^{-1} u^T A^{-1}\|_{\infty \rightarrow \infty} \leq 3C_A^2 C'_v C'_u \sigma_0^{-1}.$$

Adding the four block estimates gives the displayed uniform bound. $\square$

Lemma 8.20 (Scale-normalization nondegeneracy gives Schur data). *Assume the scale normalization constraint is satisfied, the scale derivative obeys*

$$\alpha(V) = DG_{\text{sc}}(V)[Z_{\text{sc}}(V)] = p\Theta_0 > 0,$$

the translation block has a uniform inverse bound, and the mixed blocks obey the weighted first-moment bound and the Schur-compatible estimate

$$\alpha(V(\tau)) - u(V(\tau))^T A(V(\tau))^{-1} v(V(\tau)) \geq \frac{\alpha_0}{2}.$$

Then the hypotheses of [definition 8.18](#) hold with $\sigma_0 = \alpha_0/2$.

Proof. The scale-normalization identity gives $\alpha(V) = DG_{\text{sc}}(V)[Z_{\text{sc}}(V)] = p\Theta_0$, hence the scale transversality bound. The translation-block inverse hypothesis gives the uniform bound on A^{-1} . The mixed-block hypotheses give the bounds on u and v , and the displayed estimate is the Schur gap with $\sigma_0 = \alpha_0/2$. $\square$

Definition 8.21 (Weighted scale-force components). Let $0 < p < 3$ and set

$$H(V, P) = \nu \Delta V - (V \cdot \nabla) V - \nabla P.$$

Define

$$I_{\Delta}(\tau) := \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot \Delta V \, dy,$$

$$I_{\text{nl}}(\tau) := \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot ((V \cdot \nabla) V) \, dy,$$

and

$$I_P(\tau) := \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot \nabla P \, dy,$$

whenever these integrals are absolutely convergent or are defined by the distributional pairing used for the renormalized solution.

Lemma 8.22 (Scale-force decomposition). *If*

$$\sup_{\tau} |I_{\Delta}(\tau)| \leq C_{\Delta}, \quad \sup_{\tau} |I_{\text{nl}}(\tau)| \leq C_{\text{nl}}, \quad \sup_{\tau} |I_P(\tau)| \leq C_P,$$

then

$$\sup_{\tau \geq \tau_0} \left| \int |y|^{-p} |V| V \cdot H(V, P) \, dy \right| \leq \nu C_{\Delta} + C_{\text{nl}} + C_P.$$

Proof. By the definition of H ,

$$\int |y|^{-p} |V| V \cdot H(V, P) \, dy = \nu I_{\Delta} - I_{\text{nl}} - I_P.$$

The triangle inequality gives the displayed uniform bound. $\square$

Lemma 8.23 (Weighted fourth moment controls the nonlinear scale component). *Assume $V(\tau)$ is divergence-free and belongs to the approximation class needed to justify the weighted integration by parts. If*

$$\sup_{\tau \geq \tau_0} \int_{\mathbb{R}^3} |y|^{-p-1} |V(y, \tau)|^4 \, dy < \infty,$$

then

$$|I_{\text{nl}}(\tau)| \leq \frac{p}{3} \int_{\mathbb{R}^3} |y|^{-p-1} |V(y, \tau)|^4 \, dy.$$

Proof. For smooth compactly supported approximants,

$$|V|V \cdot ((V \cdot \nabla)V) = \frac{1}{3}V \cdot \nabla(|V|^3).$$

Since $\nabla \cdot V = 0$, integration by parts gives

$$I_{\text{nl}} = \frac{1}{3} \int |y|^{-p} V \cdot \nabla(|V|^3) dy = -\frac{1}{3} \int |V|^3 V \cdot \nabla(|y|^{-p}) dy.$$

Because

$$\nabla(|y|^{-p}) = -p|y|^{-p-2}y,$$

we obtain

$$I_{\text{nl}} = \frac{p}{3} \int |y|^{-p-2}(V \cdot y)|V|^3 dy.$$

Therefore

$$|I_{\text{nl}}| \leq \frac{p}{3} \int |y|^{-p-1}|V|^4 dy.$$

Taking the supremum in τ proves the assertion. The general admissible case follows by the same approximation convention used for the weighted scale normalization. $\square$

Corollary 8.24 (Pointwise scale-force bound). *If*

$$\sup_{\tau} |I_{\Delta}(\tau)| < \infty, \quad \sup_{\tau \geq \tau_0} \int |y|^{-p-1}|V(y, \tau)|^4 dy < \infty, \quad \sup_{\tau} |I_P(\tau)| < \infty,$$

then

$$\sup_{\tau \geq \tau_0} \left| \int |y|^{-p}|V|V \cdot H(V, P) dy \right| < \infty.$$

Proof. Lemma 8.23 bounds I_{nl} . Together with the Laplacian and pressure bounds, this gives the hypotheses of lemma 8.22. $\square$

Lemma 8.25 (Bounded modulation from the modulation system). *Assume*

$$M(V(\tau)) \begin{pmatrix} a(\tau) \\ b(\tau) \end{pmatrix} = -F(V(\tau)),$$

$$\sup_{\tau \geq \tau_0} \|M(V(\tau))^{-1}\|_{\infty \rightarrow \infty} < \infty, \quad \sup_{\tau \geq \tau_0} \|F(V(\tau))\|_{\infty} < \infty.$$

Then

$$\sup_{\tau \geq \tau_0} (|a(\tau)| + \|b(\tau)\|_{\infty}) < \infty.$$

Proof. The system gives

$$\begin{pmatrix} a(\tau) \\ b(\tau) \end{pmatrix} = -M(V(\tau))^{-1}F(V(\tau)),$$

so

$$\max\{|a(\tau)|, \|b(\tau)\|_{\infty}\} \leq \|M(V(\tau))^{-1}\|_{\infty \rightarrow \infty} \|F(V(\tau))\|_{\infty}.$$

Taking the supremum in τ gives bounded modulation. $\square$

Proposition 8.26 (Windowed H^1 exclusion from modulation bounds). *Assume a renormalized Type II solution has the repaired-gauge orbit and pressure reconstruction, bounded critical L^3 norm, uniformly invertible modulation matrix, uniformly bounded modulation forcing vector, and the compact-cylinder regularity needed for Caccioppoli. Then*

$$\forall R > 0 : \quad \sup_{T \geq \tau_0+1} \int_T^{T+1} \|V(\tau)\|_{H^1(B_R)}^2 d\tau < \infty,$$

and failure of local windowed H^1 control is excluded for this solution.

Proof. The inverse matrix and forcing bounds imply bounded modulation by [lemma 8.25](#). The renormalized solution supplies the renormalized equation and pressure reconstruction. Bounded critical norm gives the local compact-cylinder critical bounds, and the regularity hypothesis allows the Caccioppoli estimate. Therefore [corollary 8.16](#) applies. Failure of local windowed H^1 control is the negation of the displayed tail windowed H^1 estimate on some bounded ball. $\square$

Corollary 8.27 (Two-alternative exclusion after windowed control). *Assume a finite-cost renormalized Type II solution not excluded by the compact cost-divergence criterion has reached the two-alternative classification. If the hypotheses of [proposition 8.26](#) are supplied for that solution, then failure of local windowed H^1 control is excluded. If, in addition, uniform critical tightness is available, then no finite-cost alternative remains outside the compact cost-divergence criterion.*

Proof. [Proposition 8.26](#) gives the positive tail windowed H^1 bound, contradicting failure of local windowed H^1 control. If uniform critical tightness is also available, then failure of tightness is excluded as well, and the two-alternative classification has no remaining finite-cost alternative outside the compact cost-divergence criterion. $\square$

Theorem 8.28 (Exclusion of local windowed H^1 failure). *If every renormalized Type II solution in the class satisfies the hypotheses of [proposition 8.26](#), then no such solution can satisfy*

$$\exists m \geq 1 : \sup_{T \geq \tau_0 + 1} \int_T^{T+1} \|V(\tau)\|_{H^1(B_m)}^2 d\tau = \infty.$$

Proof. Apply [proposition 8.26](#) to each renormalized solution. The displayed condition is the negation of the resulting windowed H^1 -bound for $R = m$. $\square$

8.4. Window-Integrated Modulation Forcing and Weighted Scale Components. This subsection records sufficient analytic estimates for the modulation forcing. The estimates are stated directly for the differentiated modulation conditions and for the weighted scale component; no pointwise modulation bound is used unless it is explicitly assumed.

Let the modulation forcing operator be

$$H(V, P) := \nu \Delta V - (V \cdot \nabla) V - \nabla P.$$

For gauge functionals G_k , $0 \leq k \leq 3$, define

$$F_k(V) := DG_k(V)[H(V, P)].$$

Thus the modulation system has the form

$$M(V(\tau)) \begin{pmatrix} a(\tau) \\ b(\tau) \end{pmatrix} = -F(V(\tau)).$$

Lemma 8.29 (Forcing bound for the translation constraints). *Let*

$$G_j(V) = \int_{\mathbb{R}^3} y_j |V(y)|^2 \psi_R(y) dy, \quad 1 \leq j \leq 3,$$

where $\psi_R \in C_c^\infty(B_{2R})$, and set

$$\Phi_j(y) := y_j \psi_R(y).$$

For

$$F_j(V) := DG_j(V)[H(V, P)] = 2 \int_{\mathbb{R}^3} \Phi_j V \cdot H(V, P) dy,$$

assume, uniformly for $\tau \geq \tau_0$,

$$\|V(\tau)\|_{L^3(B_{2R})} \leq M_3, \quad \|\nabla V(\tau)\|_{L^2(B_{2R})} \leq M_\nabla,$$

and

$$\inf_{c \in \mathbb{R}} \|P(\tau) - c\|_{L^{3/2}(B_{2R})} \leq M_P.$$

Then

$$\sup_{\tau \geq \tau_0} \max_{1 \leq j \leq 3} |F_j(V(\tau))| \leq C(R, \nu, \psi_R) (M_\nabla^2 + M_3^2 + M_3^3 + M_P M_3).$$

Proof. Fix j . Since Φ_j is supported in B_{2R} ,

$$F_j = 2\nu I_\Delta - 2I_{\text{nl}} - 2I_P,$$

where

$$I_\Delta := \int \Phi_j V \cdot \Delta V, \quad I_{\text{nl}} := \int \Phi_j V \cdot (V \cdot \nabla V), \quad I_P := \int \Phi_j V \cdot \nabla P.$$

For the Laplacian term, using compact support of Φ_j and summing over the component index α ,

$$\begin{aligned} I_\Delta &= \sum_{\alpha=1}^3 \int \Phi_j V_\alpha \Delta V_\alpha dy \\ &= - \sum_{\alpha,k} \int \partial_k (\Phi_j V_\alpha) \partial_k V_\alpha dy \\ &= - \int \Phi_j |\nabla V|^2 dy - \sum_{\alpha,k} \int (\partial_k \Phi_j) V_\alpha \partial_k V_\alpha dy. \end{aligned}$$

Equivalently,

$$I_\Delta = - \int \Phi_j |\nabla V|^2 - \int (\nabla \Phi_j \cdot \nabla V) \cdot V.$$

Consequently, by the Cauchy–Schwarz inequality,

$$|I_\Delta| \leq \|\Phi_j\|_{L^\infty} \|\nabla V\|_{L^2(B_{2R})}^2 + \|\nabla \Phi_j\|_{L^\infty} \|V\|_{L^2(B_{2R})} \|\nabla V\|_{L^2(B_{2R})}.$$

Hölder’s inequality on the bounded set B_{2R} gives

$$\|V\|_{L^2(B_{2R})} \leq |B_{2R}|^{1/6} \|V\|_{L^3(B_{2R})} \leq C_R M_3.$$

Therefore

$$|I_\Delta| \leq C_R (M_\nabla^2 + M_3 M_\nabla) \leq C_R (M_\nabla^2 + M_3^2).$$

The last inequality is Young’s inequality, $ab \leq (a^2 + b^2)/2$.

For the nonlinear term,

$$V \cdot (V \cdot \nabla V) = V_i V_\alpha \partial_i V_\alpha = \frac{1}{2} V_i \partial_i (|V|^2) = \frac{1}{2} V \cdot \nabla (|V|^2).$$

Since $\nabla \cdot V = 0$ and Φ_j is compactly supported,

$$I_{\text{nl}} = \frac{1}{2} \int \Phi_j V \cdot \nabla (|V|^2) = -\frac{1}{2} \int |V|^2 V \cdot \nabla \Phi_j.$$

Thus

$$|I_{\text{nl}}| \leq \frac{1}{2} \|\nabla \Phi_j\|_{L^\infty} \|V\|_{L^3(B_{2R})}^3 \leq C_R M_3^3.$$

For the pressure term, replace P by $P - c(\tau)$. Since $\nabla \cdot V = 0$, the identity

$$\nabla \cdot (\Phi_j V) = V \cdot \nabla \Phi_j$$

holds distributionally. Therefore

$$I_P = \int \Phi_j V \cdot \nabla (P - c) = - \int (P - c) V \cdot \nabla \Phi_j.$$

Consequently

$$|I_P| \leq \|\nabla \Phi_j\|_{L^\infty} \|P - c\|_{L^{3/2}(B_{2R})} \|V\|_{L^3(B_{2R})} \leq C_R M_P M_3.$$

Taking the infimum over c , summing the three estimates in F_j , and then taking the supremum in τ proves the bound. $\square$

Remark 8.30 (Local source of the gradient bound). The pointwise local gradient bound in [lemma 8.29](#) is not a free terminal assumption in the retained proof. Whenever the pointwise lemma is invoked on a retained compact window, the bound is supplied at admissible times by the local Caccioppoli/regularity package, or else the first failed row is the local H^1 -loss or rough-core row before the translation estimate is used. Thus the translation forcing estimate never obtains windowed H^1 -control by feeding that same conclusion back into its hypotheses.

Definition 8.31 (Weighted scale forcing). Let

$$G_{\text{sc}}(V) = \int_{\mathbb{R}^3} |y|^{-p} |V|^3 dy - \Theta_0, \quad 0 < p < 3.$$

The weighted scale component is

$$F_0(V) := DG_{\text{sc}}(V)[H(V, P)] = 3 \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot H(V, P) dy.$$

All weighted pairings in this subsection are assumed to be defined either as absolutely convergent integrals or as specified distributional pairings represented by $L^1_{\text{loc}}(d\tau)$ functions.

Lemma 8.32 (Weighted scale forcing bound). *Assume*

$$\sup_{\tau \geq \tau_0} \left| \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot H(V, P) dy \right| < \infty.$$

Then

$$\sup_{\tau \geq \tau_0} |F_0(V(\tau))| < \infty.$$

Proof. By the derivative formula in [definition 8.31](#),

$$|F_0(V(\tau))| \leq 3 \left| \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot H(V, P) dy \right|.$$

Taking the supremum in τ gives the result. $\square$

Corollary 8.33 (Pointwise modulation forcing bound). *Assume the hypotheses of [lemma 8.29](#) for the three translation components and the weighted scale forcing hypothesis of [lemma 8.32](#). Then*

$$\sup_{\tau \geq \tau_0} \|F(V(\tau))\|_\infty < \infty.$$

Proof. The vector $F(V) \in \mathbb{R}^4$ has the scale component F_0 and the three translation components F_j , $1 \leq j \leq 3$. The scale component is bounded by [lemma 8.32](#), and the translation components are bounded by [lemma 8.29](#). Taking the maximum over the four components proves the ℓ^∞ -bound. $\square$

Corollary 8.34 (Analytic inputs for pointwise modulation forcing). *Pointwise boundedness of the modulation forcing follows from the following analytic estimates:*

- (i) *local L^3 , local H^1 , and local pressure-oscillation control on the support of the compactly supported centering constraints;*
- (ii) *weighted control of the Laplacian, pressure, and fourth-moment terms in the scale component, sufficient to bound the weighted scale forcing.*

The scale-component estimate is independent of the algebraic scale transversality identity

$$DG_{\text{sc}}(V)[Z_{\text{sc}}(V)] = p\Theta_0.$$

Proof. The first item is the set of hypotheses used in [lemma 8.29](#). The second item bounds the weighted scale component through [lemma 8.32](#), or through the explicit decomposition in [corollary 8.51](#) below in the corresponding integrated setting. [Corollary 8.33](#) then gives the modulation forcing bound. The displayed transversality identity controls the modulation matrix, not the forcing vector, so it is a separate algebraic input. $\square$

Lemma 8.35 (Window-integrated forcing gives window-integrated modulation). *Assume*

$$\sup_{\tau \geq \tau_0} \|M(V(\tau))^{-1}\|_{\infty \rightarrow \infty} \leq C_M < \infty$$

and

$$\sup_{T \geq \tau_0} \int_T^{T+1} \|F(V(\tau))\|_{\infty} d\tau \leq C_F < \infty.$$

Then

$$\sup_{T \geq \tau_0} \int_T^{T+1} (|a(\tau)| + \|b(\tau)\|_{\infty}) d\tau \leq 2C_M C_F.$$

Proof. For almost every τ ,

$$\begin{pmatrix} a(\tau) \\ b(\tau) \end{pmatrix} = -M(V(\tau))^{-1}F(V(\tau)).$$

Hence

$$\max\{|a(\tau)|, \|b(\tau)\|_{\infty}\} \leq C_M \|F(V(\tau))\|_{\infty}.$$

Since

$$|a(\tau)| + \|b(\tau)\|_{\infty} \leq 2 \max\{|a(\tau)|, \|b(\tau)\|_{\infty}\},$$

we have, for every $T \geq \tau_0$,

$$\begin{aligned} \int_T^{T+1} (|a(\tau)| + \|b(\tau)\|_{\infty}) d\tau &\leq 2C_M \int_T^{T+1} \|F(V(\tau))\|_{\infty} d\tau \\ &\leq 2C_M C_F. \end{aligned}$$

Taking the supremum over T proves the claim. $\square$

Lemma 8.36 (Selection from an L^1 modulation bound). *Let $I_n = [T_n, T_n + 1]$, where $T_n \rightarrow \infty$. Let*

$$d_n : I_n \rightarrow [0, \infty), \quad q_n : I_n \rightarrow [0, \infty)$$

be measurable functions satisfying

$$\int_{I_n} d_n(\tau) d\tau \rightarrow 0, \quad \sup_n \int_{I_n} q_n(\tau) d\tau \leq Q < \infty.$$

Then, after choosing measurable representatives, there are times $\tau_n \in I_n$ such that

$$d_n(\tau_n) \rightarrow 0$$

and, if $Q > 0$,

$$q_n(\tau_n) \leq 2Q \quad \text{for every } n.$$

If $Q = 0$, the times may be chosen so that

$$q_n(\tau_n) = 0, \quad d_n(\tau_n) \rightarrow 0.$$

Proof. Assume first that $Q > 0$, and set

$$E_n := \{\tau \in I_n : q_n(\tau) \leq 2Q\}.$$

By Chebyshev's inequality,

$$|I_n \setminus E_n| \leq \frac{1}{2Q} \int_{I_n} q_n(\tau) d\tau \leq \frac{1}{2}.$$

Thus $|E_n| \geq 1/2$. Since $d_n \geq 0$,

$$\inf_{\tau \in E_n} d_n(\tau) \leq \frac{1}{|E_n|} \int_{E_n} d_n(\tau) d\tau \leq 2 \int_{I_n} d_n(\tau) d\tau.$$

Choose $\tau_n \in E_n$, outside the null set on which the representatives are undefined, such that

$$d_n(\tau_n) \leq 2 \int_{I_n} d_n(\tau) d\tau + \frac{1}{n}.$$

The right-hand side tends to zero, hence $d_n(\tau_n) \rightarrow 0$, and the definition of E_n gives $q_n(\tau_n) \leq 2Q$.

If $Q = 0$, then $q_n = 0$ almost everywhere on I_n . Applying the same averaging argument to d_n on the full-measure set where $q_n = 0$ gives the asserted times. $\square$

Corollary 8.37 (Selection on unit windows). *Assume*

$$q(\tau) := |a(\tau)| + \|b(\tau)\|_\infty$$

satisfies

$$\sup_{T \geq \tau_0} \int_T^{T+1} q(\tau) d\tau < \infty.$$

For every sequence of unit windows $I_n = [T_n, T_n + 1]$ and every nonnegative measurable density d_n satisfying

$$\int_{I_n} d_n(\tau) d\tau \rightarrow 0,$$

there are times $\tau_n \in I_n$ such that

$$d_n(\tau_n) \rightarrow 0, \quad \sup_n (|a(\tau_n)| + \|b(\tau_n)\|_\infty) < \infty.$$

Proof. Apply [lemma 8.36](#) with $q_n = q|_{I_n}$ and

$$Q = \sup_n \int_{I_n} q(\tau) d\tau.$$

This Q is finite because of the uniform window L^1 -bound. If $Q > 0$, the selected times satisfy

$$|a(\tau_n)| + \|b(\tau_n)\|_\infty = q(\tau_n) \leq 2Q.$$

If $Q = 0$, the same lemma gives $q(\tau_n) = 0$. In both cases $d_n(\tau_n) \rightarrow 0$ and the modulation is uniformly bounded along the selected times. $\square$

Lemma 8.38 (Product-form windowed modulation estimate). *Assume*

$$\sup_{T \geq \tau_0} \int_T^{T+1} \|M(V(\tau))^{-1}\|_{\infty \rightarrow \infty} \|F(V(\tau))\|_\infty d\tau < \infty.$$

Then

$$\sup_{T \geq \tau_0} \int_T^{T+1} (|a(\tau)| + \|b(\tau)\|_\infty) d\tau < \infty.$$

Proof. The modulation system gives

$$\begin{pmatrix} a(\tau) \\ b(\tau) \end{pmatrix} = -M(V(\tau))^{-1}F(V(\tau)).$$

Therefore

$$|a(\tau)| + \|b(\tau)\|_\infty \leq 2\|M(V(\tau))^{-1}\|_{\infty \rightarrow \infty} \|F(V(\tau))\|_\infty.$$

Integrating this pointwise estimate on $[T, T+1]$ and taking the supremum in T gives the result. $\square$

Lemma 8.39 (Integrated translation forcing bound). *Let G_j , Φ_j , $H(V, P)$, and $F_j(V)$ be as in lemma 8.29. Assume that, for some $R > 0$,*

$$\sup_{\tau \geq \tau_0} \|V(\tau)\|_{L^3(B_{2R})} \leq M_3,$$

$$\sup_{T \geq \tau_0} \int_T^{T+1} \|\nabla V(\tau)\|_{L^2(B_{2R})}^2 d\tau \leq A_R,$$

and

$$\sup_{\tau \geq \tau_0} \inf_{c \in \mathbb{R}} \|P(\tau) - c\|_{L^{3/2}(B_{2R})} \leq M_P.$$

Then

$$\sup_{T \geq \tau_0} \int_T^{T+1} \max_{1 \leq j \leq 3} |F_j(V(\tau))| d\tau \leq C(R, \nu, \psi_R)(A_R + M_3^2 + M_3^3 + M_P M_3).$$

Proof. Fix j . The proof of lemma 8.29 gives, for almost every τ ,

$$|F_j(V(\tau))| \leq C(R, \nu, \psi_R) \left(\|\nabla V(\tau)\|_{L^2(B_{2R})}^2 + M_3 \|\nabla V(\tau)\|_{L^2(B_{2R})} + M_3^3 + M_P M_3 \right).$$

Integrate this estimate over $I_T = [T, T+1]$. The first term contributes at most A_R . For the mixed term, Cauchy's inequality on the unit interval gives

$$\int_{I_T} M_3 \|\nabla V(\tau)\|_{L^2(B_{2R})} d\tau \leq M_3 \left(\int_{I_T} \|\nabla V(\tau)\|_{L^2(B_{2R})}^2 d\tau \right)^{1/2} \leq M_3 A_R^{1/2}.$$

Since

$$M_3 A_R^{1/2} \leq \frac{1}{2}(M_3^2 + A_R),$$

the mixed term is bounded by $C(A_R + M_3^2)$. The terms M_3^3 and $M_P M_3$ are constant on the unit window. Hence

$$\int_T^{T+1} |F_j(V(\tau))| d\tau \leq C(R, \nu, \psi_R)(A_R + M_3^2 + M_3^3 + M_P M_3).$$

Since $\max_{1 \leq j \leq 3} |F_j| \leq \sum_{j=1}^3 |F_j|$, taking the maximum over the three centering constraints changes only the constant. $\square$

Remark 8.40 (Independence of the windowed gradient hypothesis). The windowed gradient input in lemma 8.39 is an independent hypothesis whenever the integrated translation estimate is used to prove windowed H^1 -control. It may come, for example, from a preliminary Caccioppoli estimate depending only on already available data, from a direct approximation or regularity estimate on the support of the centering functions, or from an independent verification obtained without using the conclusion of the windowed H^1 estimate under consideration.

Lemma 8.41 (Window-integrated scale forcing). *Assume the window-integrated weighted scale-forcing bound*

$$\sup_{T \geq \tau_0} \int_T^{T+1} \left| \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot H(V, P) dy \right| d\tau < \infty,$$

where the weighted pairing is interpreted as in [definition 8.31](#). Then

$$\sup_{T \geq \tau_0} \int_T^{T+1} |F_0(V(\tau))| d\tau < \infty.$$

Proof. By the identity

$$F_0(V) = 3 \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot H(V, P) dy$$

we have, for each $T \geq \tau_0$,

$$\int_T^{T+1} |F_0(V(\tau))| d\tau = 3 \int_T^{T+1} \left| \int_{\mathbb{R}^3} |y|^{-p} |V| V \cdot H(V, P) dy \right| d\tau.$$

Taking the supremum over T gives the asserted L^1 -bound. $\square$

Lemma 8.42 (Full integrated modulation forcing). *Assume the integrated translation forcing estimate*

$$\sup_{T \geq \tau_0} \int_T^{T+1} \max_{1 \leq j \leq 3} |F_j(V(\tau))| d\tau < \infty$$

and the window-integrated weighted scale-forcing bound in [lemma 8.41](#). Then

$$\sup_{T \geq \tau_0} \int_T^{T+1} \|F(V(\tau))\|_\infty d\tau < \infty.$$

Proof. The vector $F(V) \in \mathbb{R}^4$ consists of F_0 and the three translation components F_j , $1 \leq j \leq 3$. Since

$$\|F(V(\tau))\|_\infty \leq |F_0(V(\tau))| + \max_{1 \leq j \leq 3} |F_j(V(\tau))|,$$

for almost every τ . Integrating this inequality on $[T, T+1]$ gives

$$\int_T^{T+1} \|F(V(\tau))\|_\infty d\tau \leq \int_T^{T+1} |F_0(V(\tau))| d\tau + \int_T^{T+1} \max_{1 \leq j \leq 3} |F_j(V(\tau))| d\tau.$$

The first term is uniformly bounded in T by [lemma 8.41](#), and the second term is uniformly bounded by the assumed integrated translation forcing estimate. Taking the supremum over T proves the claim. $\square$

Corollary 8.43 (Integrated modulation criterion on selected windows). *Assume the uniform inverse bound for $M(V(\tau))$, the integrated translation forcing estimate, and the window-integrated weighted scale-forcing bound. Then*

$$\sup_{T \geq \tau_0} \int_T^{T+1} (|a(\tau)| + \|b(\tau)\|_\infty) d\tau < \infty.$$

Consequently every sequence of unit windows with vanishing nonnegative measurable density contains selected times where this density vanishes asymptotically and the modulation remains uniformly bounded.

Proof. [Lemma 8.42](#) gives the window-integrated bound on $\|F(V)\|_\infty$. Combining this with the uniform inverse bound in [lemma 8.35](#) gives the window-integrated modulation bound. The selection conclusion is [corollary 8.37](#). $\square$

Definition 8.44 (Weighted scale-force components). Let $w(y) = |y|^{-p}$, $0 < p < 3$. Define

$$F_0^\Delta(V) := 3\nu \int_{\mathbb{R}^3} w|V|V \cdot \Delta V \, dy,$$

$$F_0^{\text{nl}}(V) := -3 \int_{\mathbb{R}^3} w|V|V \cdot (V \cdot \nabla V) \, dy,$$

and

$$F_0^P(V, P) := -3 \int_{\mathbb{R}^3} w|V|V \cdot \nabla P \, dy.$$

The three pairings are required to be compatible with the decomposition of $H(V, P)$ into diffusion, transport, and pressure.

Lemma 8.45 (Integrated scale-force decomposition). *For every time at which the three pairings in [definition 8.44](#) are defined,*

$$F_0(V(\tau)) = F_0^\Delta(V(\tau)) + F_0^{\text{nl}}(V(\tau)) + F_0^P(V(\tau), P(\tau)).$$

Proof. Substitute

$$H(V, P) = \nu \Delta V - (V \cdot \nabla)V - \nabla P$$

into

$$F_0(V) = 3 \int w|V|V \cdot H(V, P) \, dy.$$

Then

$$F_0(V) = 3\nu \int w|V|V \cdot \Delta V \, dy - 3 \int w|V|V \cdot (V \cdot \nabla V) \, dy - 3 \int w|V|V \cdot \nabla P \, dy.$$

The three terms coincide with $F_0^\Delta(V)$, $F_0^{\text{nl}}(V)$, and $F_0^P(V, P)$, respectively. Thus linearity of the weighted pairing gives the stated decomposition. $\square$

Lemma 8.46 (Component bounds imply integrated scale-force control). *Assume*

$$\sup_{T \geq \tau_0} \int_T^{T+1} |F_0^\Delta(V(\tau))| \, d\tau < \infty,$$

$$\sup_{T \geq \tau_0} \int_T^{T+1} |F_0^{\text{nl}}(V(\tau))| \, d\tau < \infty,$$

and

$$\sup_{T \geq \tau_0} \int_T^{T+1} |F_0^P(V(\tau), P(\tau))| \, d\tau < \infty.$$

Then the window-integrated weighted scale-forcing bound in [lemma 8.41](#) holds.

Proof. By [lemma 8.45](#),

$$|F_0(V(\tau))| \leq |F_0^\Delta(V(\tau))| + |F_0^{\text{nl}}(V(\tau))| + |F_0^P(V(\tau), P(\tau))|$$

for almost every τ . Integrating on $[T, T+1]$, taking the supremum in T , and applying the three component bounds gives

$$\sup_{T \geq \tau_0} \int_T^{T+1} |F_0(V(\tau))| \, d\tau < \infty.$$

Since F_0 is three times the weighted pairing in [definition 8.31](#), this is equivalent to the window-integrated weighted scale-forcing bound. $\square$

Definition 8.47 (Singular-weight convective integration by parts). The singular-weight convective integration-by-parts condition means that, for almost every τ ,

$$-\int_{\mathbb{R}^3} w V \cdot \nabla(|V|^3) dy = \int_{\mathbb{R}^3} |V|^3 V \cdot \nabla w dy,$$

with no boundary contribution from $y = 0$ or from spatial infinity, either classically or by convergence of smooth compactly supported divergence-free approximants.

Lemma 8.48 (Convective scale-force bound). Assume the condition of [definition 8.47](#) and

$$\sup_{T \geq \tau_0} \int_T^{T+1} \int_{\mathbb{R}^3} |y|^{-p-1} |V(y, \tau)|^4 dy d\tau < \infty.$$

Then

$$\sup_{T \geq \tau_0} \int_T^{T+1} |F_0^{\text{nl}}(V(\tau))| d\tau < \infty.$$

More precisely,

$$|F_0^{\text{nl}}(V(\tau))| \leq p \int_{\mathbb{R}^3} |y|^{-p-1} |V(y, \tau)|^4 dy$$

for almost every τ .

Proof. Fix such a time τ and suppress it from the notation. Since

$$\partial_i(|V|^3) = 3|V|V_j \partial_i V_j$$

for each i , and therefore

$$V \cdot \nabla(|V|^3) = 3|V|V_i V_j \partial_i V_j = 3|V|V \cdot ((V \cdot \nabla)V),$$

we have

$$w|V|V \cdot (V \cdot \nabla V) = \frac{1}{3}w V \cdot \nabla(|V|^3).$$

Therefore

$$F_0^{\text{nl}}(V) = -3 \int w|V|V \cdot (V \cdot \nabla V) = - \int w V \cdot \nabla(|V|^3).$$

Using $\nabla \cdot V = 0$ and [definition 8.47](#),

$$-\int w V \cdot \nabla(|V|^3) = \int |V|^3 V \cdot \nabla w.$$

Since

$$\nabla w(y) = -p|y|^{-p-2}y, \quad y \neq 0,$$

we have

$$|V \cdot \nabla w| = p|y|^{-p-2}|V \cdot y| \leq p|y|^{-p-1}|V|.$$

Thus

$$|F_0^{\text{nl}}(V)| \leq p \int |y|^{-p-1} |V|^4 dy.$$

Integrating over $[T, T+1]$ and taking the supremum in T proves the window-integrated convective bound. $\square$

Definition 8.49 (Annular convective regularity). Annular convective regularity means that, for almost every τ , $V(\tau)$ is divergence free and, on every compact annulus $0 < r < |y| < R < \infty$, the chain-rule identities

$$\partial_i(|V|^3) = 3|V|V_j \partial_i V_j, \quad V \cdot \nabla(|V|^3) = 3|V|V_i V_j \partial_i V_j$$

hold in the sense of distributions, either classically or by passage to the limit from smooth divergence-free approximants on annuli.

Lemma 8.50 (Annular regularity gives singular-weight integration by parts). *Assume annular convective regularity and the weighted L^4 -window bound*

$$\sup_{T \geq \tau_0} \int_T^{T+1} \int_{\mathbb{R}^3} |y|^{-p-1} |V(y, \tau)|^4 dy d\tau < \infty.$$

Then the singular-weight convective integration-by-parts condition of [definition 8.47](#) holds for almost every τ .

Proof. Fix a time τ for which

$$\int_{\mathbb{R}^3} |y|^{-p-1} |V(y, \tau)|^4 dy < \infty$$

and annular convective regularity holds. Such times form a full-measure set on every unit window.

Let $\eta \in C^\infty([0, \infty))$ satisfy $0 \leq \eta \leq 1$, $\eta(s) = 0$ for $s \leq 1$, and $\eta(s) = 1$ for $s \geq 2$. Define

$$\chi_{\varepsilon, R}(y) := \eta(|y|/\varepsilon) \eta(R/|y|).$$

Then $\chi_{\varepsilon, R}$ is supported in $\{\varepsilon < |y| < R\}$, equals one on $\{2\varepsilon < |y| < R/2\}$, and satisfies

$$|\nabla \chi_{\varepsilon, R}(y)| \leq C\varepsilon^{-1} \mathbf{1}_{\{\varepsilon < |y| < 2\varepsilon\}} + CR^{-1} \mathbf{1}_{\{R/2 < |y| < R\}}.$$

By annular convective regularity, the following integration by parts is valid on the annulus $\{\varepsilon < |y| < R\}$. Since $\nabla \cdot V = 0$ and $w\chi_{\varepsilon, R}|V|^3V$ is compactly supported in this annulus,

$$0 = \int \nabla \cdot (w\chi_{\varepsilon, R}|V|^3V) dy.$$

Expanding the divergence gives

$$0 = \int |V|^3V \cdot \nabla(w\chi_{\varepsilon, R}) dy + \int w\chi_{\varepsilon, R}V \cdot \nabla(|V|^3) dy,$$

and hence

$$- \int w\chi_{\varepsilon, R}V \cdot \nabla(|V|^3) = \int |V|^3V \cdot \nabla(w\chi_{\varepsilon, R}).$$

Expanding the derivative,

$$- \int w\chi_{\varepsilon, R}V \cdot \nabla(|V|^3) = \int \chi_{\varepsilon, R}|V|^3V \cdot \nabla w + \int w|V|^3V \cdot \nabla \chi_{\varepsilon, R}.$$

On the inner cutoff region $\{\varepsilon < |y| < 2\varepsilon\}$, one has $|y| \simeq \varepsilon$, and on the outer cutoff region $\{R/2 < |y| < R\}$, one has $|y| \simeq R$. Therefore the derivative bound for $\chi_{\varepsilon, R}$ implies

$$\begin{aligned} \left| \int w|V|^3V \cdot \nabla \chi_{\varepsilon, R} \right| &\leq C \int_{\{\varepsilon < |y| < 2\varepsilon\}} |y|^{-p-1} |V|^4 dy \\ &\quad + C \int_{\{R/2 < |y| < R\}} |y|^{-p-1} |V|^4 dy. \end{aligned}$$

By the explicit weighted L^4 -window hypothesis in the statement, the integrand $|y|^{-p-1}|V|^4$ belongs to $L^1(\mathbb{R}^3)$ at the fixed time. Hence absolute continuity of the Lebesgue integral gives

$$\int_{\{\varepsilon < |y| < 2\varepsilon\}} |y|^{-p-1} |V|^4 dy \rightarrow 0 \quad \text{as } \varepsilon \downarrow 0,$$

and integrability at infinity gives

$$\int_{\{R/2 < |y| < R\}} |y|^{-p-1} |V|^4 dy \rightarrow 0 \quad \text{as } R \uparrow \infty.$$

Moreover

$$|\chi_{\varepsilon,R}|V|^3V \cdot \nabla w| \leq p|y|^{-p-1}|V|^4 \in L^1(\mathbb{R}^3),$$

so dominated convergence gives

$$\int \chi_{\varepsilon,R}|V|^3V \cdot \nabla w \rightarrow \int |V|^3V \cdot \nabla w.$$

Passing first $\varepsilon \downarrow 0$ and then $R \uparrow \infty$ yields

$$-\int_{\mathbb{R}^3} w V \cdot \nabla(|V|^3) dy = \int_{\mathbb{R}^3} |V|^3V \cdot \nabla w dy,$$

with no boundary contribution at $y = 0$ or at infinity. $\square$

Corollary 8.51 (Reduced integrated scale-force criterion). *Assume annular convective regularity, the weighted L^4 -window bound, and the two component estimates*

$$\begin{aligned} \sup_{T \geq \tau_0} \int_T^{T+1} |F_0^\Delta(V(\tau))| d\tau &< \infty, \\ \sup_{T \geq \tau_0} \int_T^{T+1} |F_0^P(V(\tau), P(\tau))| d\tau &< \infty. \end{aligned}$$

Then the window-integrated weighted scale-forcing bound holds.

Proof. [Lemma 8.50](#) gives the singular-weight integration-by-parts condition. [Lemma 8.48](#) then gives the integrated convective scale-force bound. Combining this with the diffusive and pressure component estimates, and applying [lemma 8.46](#), gives the window-integrated weighted scale-forcing bound. $\square$

Corollary 8.52 (Remaining analytic alternatives for the scale component). *Before the singular-weight integration by parts is justified, failure of the window-integrated weighted scale-forcing bound is reduced to failure of at least one of the following four analytic statements:*

- (i) *the integrated diffusive weighted scale-force bound;*
- (ii) *the singular-weight convective integration-by-parts identity;*
- (iii) *the weighted L^4 -window bound;*
- (iv) *the integrated weighted pressure scale-force bound.*

If annular convective regularity is available, then the second item is replaced by failure of annular convective regularity itself, together with the same weighted L^4 -window input.

Proof. [Lemma 8.46](#) shows that the diffusive, convective, and pressure component bounds imply the integrated scale-force bound. [Lemma 8.48](#) reduces the convective component to the singular-weight integration-by-parts identity and the weighted L^4 -window bound. Finally, [lemma 8.50](#) shows that annular convective regularity and the weighted L^4 -window bound imply the singular-weight integration-by-parts identity. Therefore the listed failures exhaust the remaining ways in which this scale-component estimate can fail. $\square$

8.5. Navier–Stokes Cost-Divergence Exclusion.

Definition 8.53 (Cost phases and dependency rank). The cost argument is divided into the following phases:

- (P0) retained local equation, energy, pressure, modulation, cutoff, and buffer data;
- (P1) measurability and finite-window integrability of the localized cost;
- (P2) definition and elementary comparison of the canonical signed and nonnegative costs;
- (P3) fixed- and moving-cutoff local energy identities and their decomposition into the components $J_1, \dots, J_7$;

- (P4) finite-error selection or normalized occupation-state extraction for a nonsummable component;
- (P5) routing of the resulting carrier certificates to previously declared local states;
- (P6) contradiction between a divergent validated cost and the alternatives closed in the preceding phases;
- (P7) final validation of the completed chain.

A proof-bearing cost statement has dependency rank (k, r) , where k is its cost phase and $r \in \mathbb{N}$ orders statements within that phase. It may use statements of smaller lexicographic rank only. Definitions may be cited at the phase in which their data first become available. In particular, no statement in phases 0–5 may use a cost-divergence contradiction or the final validation ledger.

The assignments used below are summarized here. Each row lists the principal labels; intermediate elementary lemmas inherit the phase of the row in which they occur.

Statement	Phase	May use	Must not use
local state-data production	0	represented suitable equation	any cost or carrier conclusion
lemma 8.68 and corollary 8.71	1	phase 0	carrier routing or cost divergence
lemmas 8.66 and 8.67	2	phases 0–1	carrier closure or divergence
lemma 8.89	3	phases 0–2	carrier routing or divergence
normalized carrier extraction and finite-error selection	4	phases 0–3	final cost contradiction
windowwise and ambient carrier routing	5	phases 0–4 and earlier state closures	cost divergence
theorems 8.64 and 10.42	6	phases 0–5	final assembly
lemmas 8.73 and 10.45 and proposition 8.75	7	phases 0–6	itself

Definition 8.54 (Retained local cost state data). Let $I = [\alpha, \beta]$, let $0 < R_0 < R_1 < \infty$, and let (V, P, a, b) be a repaired-gauge branch on $B_{R_1} \times I$. Its *retained local cost state data* on the core $B_{R_0} \times I$ consist of:

- (D1) the represented distributional Navier–Stokes equation and the suitable local energy inequality on the working cylinder $B_{R_1} \times I$;
- (D2)

$$V \in L^\infty(I; L^2(B_{R_1})) \cap L^2(I; H^1(B_{R_1}));$$

- (D3) the pressure representative

$$P - (P)_{B_{R_1}} \in L^{3/2}(I \times B_{R_1}), \quad (P)_{B_{R_1}}(\tau) = \frac{1}{|B_{R_1}|} \int_{B_{R_1}} P(y, \tau) dy;$$

- (D4) $a, b \in L^1(I)$;
- (D5) a nonnegative spatial cutoff $\phi_{R_0} \in C_c^\infty(B_{R_1})$, equal to one on B_{R_0} , and, whenever a time-localized identity is used, a temporal cutoff supported strictly inside the working time interval and equal to one on the core time interval.

These are phase 0 data. Their definition contains no cost-divergence or carrier assertion.

Definition 8.55 (Pre-divergence validation rows). Before cost divergence is tested, a branch is checked in the following order: Type II analytic and exhaustion data; represented repaired-gauge and pressure data; compact single-core retention; finite local critical mass and critical tightness; local H^1 control; the retained local cost state data of [definition 8.54](#); cost well-posedness

and comparison; the fixed or moving local energy identity; finite-error selection or normalized component extraction; and carrier routing. The first failed row is recorded as its corresponding named local state. This definition fixes the label order used in phases 0–5; it neither assumes nor proves cost divergence.

Definition 8.56 (Type II exclusion conclusion). For a renormalized Type II solution, the exclusion conclusion means that after the active supercritical scaling condition and the compact cost-divergence criterion have both been established, the solution is excluded from the Type II class under consideration. This is stronger than cost divergence alone, which only states that the solution satisfies the compact Type II cost-divergence criterion.

Definition 8.57 (CostPhase 6: local cost-to-exclusion data). Let ω be a repaired-gauge Navier–Stokes Type II branch on a terminal renormalized tail $[\tau_0, \infty)$. Local cost-to-exclusion data for the validated compact cost consist of a nonnegative locally integrable cost $\mathcal{D}_\omega$, a locally absolutely continuous corrected localized energy $\mathcal{H}_\omega$, a constant $c_\omega > 0$, and a nonnegative locally integrable remainder $\mathcal{R}_\omega$ such that:

- (i) the validated compact cost-divergence conclusion is

$$\int_{\tau_0}^{\infty} \mathcal{D}_\omega(\tau) d\tau = \infty;$$

- (ii) $\mathcal{H}_\omega(\tau_0) < \infty$ and

$$\inf_{\tau \geq \tau_0} \mathcal{H}_\omega(\tau) > -\infty;$$

- (iii) the corrected localized energy inequality

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + c_\omega \mathcal{D}_\omega(\tau) + \mathcal{R}_\omega(\tau) \leq 0$$

holds in distributions on (τ_0, ∞) .

For the identity cost one takes $\mathcal{D}_\omega = \tilde{\mathfrak{D}}_{R_0}$. For an explicitly comparable localized cost, $\mathcal{D}_\omega$ is the localized cost whose divergence is validated by the compact criterion, provided that the same cost is controlled by the corrected localized energy inequality on the terminal tail.

Definition 8.58 (CostPhase 2: energy-controlled validated cost). Let (V, P, a, b) be a repaired-gauge branch and let

$$\tilde{\mathfrak{D}}_{R_0}(\tau) = \nu \int |\nabla V|^2 \phi_{R_0} + a_+(\tau) \int |V|^2 \phi_{R_0}$$

be the identity localized Navier–Stokes cost. A validated compact cost $\mathfrak{C}_{\text{II}}^{NS}$ is energy-controlled by the identity cost on a terminal tail if there are constants $C_1 < \infty$ and $\tau_1 \geq \tau_0$ such that

$$0 \leq \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq C_1 \tilde{\mathfrak{D}}_{R_0}(\tau) \quad \text{for a.e. } \tau \geq \tau_1.$$

The identity cost is energy-controlled with $C_1 = 1$ and $\tau_1 = \tau_0$. For a non-identity cost, this is the upper comparison needed for a corrected localized energy inequality controlling $\tilde{\mathfrak{D}}_{R_0}$ to also control the validated cost.

Definition 8.59 (CostPhase 5: local monotonicity carrier for a validated cost). Let ω be a repaired-gauge Navier–Stokes Type II branch on a terminal renormalized tail $[\tau_0, \infty)$, and let $\mathfrak{C}_{\text{II}}^{NS}$ be the validated compact cost for that branch. A local monotonicity carrier for $\mathfrak{C}_{\text{II}}^{NS}$ consists of a nonnegative locally integrable carrier cost $\mathcal{A}_\omega$, a locally absolutely continuous corrected localized energy $\mathcal{H}_\omega$, a nonnegative locally integrable remainder $\mathcal{R}_\omega$, and constants $c_A > 0$, $0 < C_A < \infty$ such that:

(i) $\mathcal{H}_\omega(\tau_0) < \infty$ and

$$\inf_{\tau \geq \tau_0} \mathcal{H}_\omega(\tau) > -\infty;$$

(ii) the carrier monotonicity inequality

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + c_A \mathcal{A}_\omega(\tau) + \mathcal{R}_\omega(\tau) \leq 0$$

holds in distributions on (τ_0, ∞) ;

(iii) the validated compact cost is controlled by the carrier on the terminal tail:

$$0 \leq \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq C_A \mathcal{A}_\omega(\tau) \quad \text{for a.e. } \tau \geq \tau_0.$$

The carrier is local: $\mathcal{A}_\omega$ may be the fixed cutoff identity cost, a moving-cutoff identity cost, or another localized cost with an explicit terminal comparison to the validated compact cost.

Lemma 8.60 (CostPhase 6: monotonicity carrier gives local cost-to-exclusion data). *Let ω be a repaired-gauge Navier–Stokes Type II branch. Assume that the validated compact cost $\mathfrak{C}_{\text{II}}^{NS}$ is nonnegative, locally integrable on finite renormalized intervals, and satisfies*

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau = \infty.$$

If $\mathfrak{C}_{\text{II}}^{NS}$ admits a local monotonicity carrier in the sense of [definition 8.59](#), then ω carries the local cost-to-exclusion data of [definition 8.57](#) with $\mathcal{D}_\omega = \mathfrak{C}_{\text{II}}^{NS}$.

Proof. Let $(\mathcal{A}_\omega, \mathcal{H}_\omega, \mathcal{R}_\omega^{\text{car}}, c_A, C_A)$ be the carrier data. The lower bound and finite initial value of $\mathcal{H}_\omega$ are exactly the energy bounds required in [definition 8.57](#). The terminal comparison gives

$$\mathcal{A}_\omega(\tau) \geq C_A^{-1} \mathfrak{C}_{\text{II}}^{NS}(\tau) \quad \text{for a.e. } \tau \geq \tau_0.$$

Substituting this into the carrier monotonicity inequality yields

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + \frac{c_A}{C_A} \mathfrak{C}_{\text{II}}^{NS}(\tau) + \mathcal{R}_\omega^{\text{car}}(\tau) \leq 0$$

in distributions. Thus [definition 8.57](#) holds with

$$\mathcal{D}_\omega = \mathfrak{C}_{\text{II}}^{NS}, \quad c_\omega = c_A/C_A, \quad \mathcal{R}_\omega = \mathcal{R}_\omega^{\text{car}}.$$

The divergent-cost item is the validated compact cost-divergence conclusion. $\square$

Lemma 8.61 (CostPhase 6: fixed-cutoff monotonicity gives local cost-to-exclusion data). *Let ω be represented on a terminal tail by a repaired-gauge Navier–Stokes branch (V, P, a, b) . Let ϕ_{R_0} be the cutoff used in the validated compact cost, let $\mathfrak{C}_{\text{II}}^{NS}$ be a nonnegative locally integrable validated compact cost, and set*

$$\mathcal{E}_\phi(\tau) = \frac{1}{2} \int |V(y, \tau)|^2 \phi_{R_0}(y) dy.$$

Assume, after increasing τ_0 if necessary, that:

(i) $\mathcal{E}_\phi(\tau_0) < \infty$;

(ii) the fixed-cutoff localized energy identity [\(13\)](#) holds on the terminal tail;

(iii) the fixed-cutoff error density B_{R_0} of [definition 10.9](#) satisfies

$$\int_{\tau_0}^{\infty} B_{R_0}(\tau) d\tau < \infty;$$

(iv) the validated compact cost $\mathfrak{C}_{\text{II}}^{NS}$ is energy-controlled by $\tilde{\mathfrak{D}}_{R_0}$ in the sense of [definition 8.58](#);

(v) *the validated compact cost-divergence conclusion is*

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau = \infty.$$

Then ω carries the local cost-to-exclusion data of [definition 8.57](#) with $\mathcal{D}_\omega = \mathfrak{C}_{\text{II}}^{NS}$.

Proof. Let C_1, τ_1 be the constants in the energy-control comparison. All statements are terminal. Since $\mathfrak{C}_{\text{II}}^{NS}$ is locally integrable on finite intervals, removing any finite initial subinterval does not change divergence of its tail integral. After the allowed increase and relabeling of τ_0 , we may therefore assume that $\tau_0 \geq \tau_1$, that $\mathcal{E}_\phi(\tau_0) < \infty$, and that the upper comparison holds on the whole tail.

Define the corrected localized energy

$$\mathcal{H}_\omega(\tau) = \mathcal{E}_\phi(\tau) + \int_{\tau}^{\infty} B_{R_0}(s) ds.$$

The finite-error hypothesis makes the correction tail finite for every $\tau \geq \tau_0$ and locally absolutely continuous. Since $\mathcal{E}_\phi \geq 0$ and $B_{R_0} \geq 0$, one has

$$\inf_{\tau \geq \tau_0} \mathcal{H}_\omega(\tau) \geq 0, \quad \mathcal{H}_\omega(\tau_0) < \infty.$$

[Theorem 10.11](#) gives, in distributions,

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R_0}(\tau) \leq 0.$$

By energy-control of the validated cost,

$$\tilde{\mathfrak{D}}_{R_0}(\tau) \geq C_1^{-1} \mathfrak{C}_{\text{II}}^{NS}(\tau) \quad \text{for a.e. } \tau \geq \tau_0.$$

Therefore

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + \frac{1}{2C_1} \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq 0.$$

This is the corrected localized energy inequality in [definition 8.57](#), with

$$\mathcal{D}_\omega = \mathfrak{C}_{\text{II}}^{NS}, \quad c_\omega = (2C_1)^{-1}, \quad \mathcal{R}_\omega = 0.$$

The divergent-cost item is exactly the validated compact cost-divergence conclusion. Hence all local cost-to-exclusion data hold. $\square$

Lemma 8.62 (CostPhase 6: moving-cutoff monotonicity gives local cost-to-exclusion data). *Let ω be represented on a terminal tail by a repaired-gauge Navier–Stokes branch (V, P, a, b) , and let $\mathfrak{C}_{\text{II}}^{NS}$ be a nonnegative locally integrable validated compact cost. Let*

$$\Phi(\tau, y) = \phi(y/R(\tau))$$

be a moving cutoff as in [definition 10.18](#), with $R(\tau) \geq R_0$ locally absolutely continuous and nondecreasing, and set

$$\mathcal{E}_R(\tau) = \frac{1}{2} \int |V(y, \tau)|^2 \Phi(\tau, y) dy, \quad \tilde{\mathfrak{D}}_{R(\tau)}(\tau) = \nu \int |\nabla V|^2 \Phi + a_+(\tau) \int |V|^2 \Phi.$$

Assume, after increasing τ_0 if necessary, that:

- (i) $\mathcal{E}_R(\tau_0) < \infty$;
- (ii) *the localized energy identity is valid with the time-dependent cutoff Φ ;*
- (iii) *the moving finite-error condition holds, for instance through summable moving-annulus errors:*

$$\int_{\tau_0}^{\infty} B_{\text{mov}}(\tau) d\tau < \infty;$$

- (iv) *the validated compact cost is controlled by the moving identity cost on the terminal tail: for some $C_m < \infty$,*

$$0 \leq \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq C_m \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \quad \text{for a.e. } \tau \geq \tau_0;$$

- (v) *the validated compact cost-divergence conclusion is*

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau = \infty.$$

Then ω carries the local cost-to-exclusion data of [definition 8.57](#) with $\mathcal{D}_\omega = \mathfrak{C}_{\text{II}}^{NS}$.

Proof. The moving finite-error condition gives $B_{\text{mov}} \in L^1(\tau_0, \infty)$; if it is verified through [definition 10.20](#), this is precisely [lemma 10.21](#). Define

$$\mathcal{H}_\omega(\tau) = \mathcal{E}_R(\tau) + \int_{\tau}^{\infty} B_{\text{mov}}(s) ds.$$

The finite moving-error condition makes the correction tail finite and locally absolutely continuous. Since $\mathcal{E}_R \geq 0$ and $B_{\text{mov}} \geq 0$,

$$\inf_{\tau \geq \tau_0} \mathcal{H}_\omega(\tau) \geq 0, \quad \mathcal{H}_\omega(\tau_0) < \infty.$$

[Theorem 10.19](#) gives

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \leq 0$$

in distributions. Thus $\tilde{\mathfrak{D}}_{R(\tau)}$ is a local monotonicity carrier for the validated compact cost, with $c_A = 1/2$, $C_A = C_m$, and $\mathcal{R}_\omega = 0$. Applying [lemma 8.60](#) gives the claimed local cost-to-exclusion data. $\square$

Lemma 8.63 (CostPhase 6: corrected monotonicity bounds the validated cost). *Let ω carry local cost-to-exclusion data in the sense of [definition 8.57](#), except do not assume the divergence condition in item (i). Then*

$$\int_{\tau_0}^{\infty} \mathcal{D}_\omega(\tau) d\tau < \infty.$$

Proof. Since $\mathcal{H}_\omega$ is locally absolutely continuous and the corrected localized energy inequality holds in distributions, integration over $[\tau_0, T]$ gives, for every $T > \tau_0$,

$$\mathcal{H}_\omega(T) - \mathcal{H}_\omega(\tau_0) + c_\omega \int_{\tau_0}^T \mathcal{D}_\omega(\tau) d\tau + \int_{\tau_0}^T \mathcal{R}_\omega(\tau) d\tau \leq 0.$$

The remainder is nonnegative, hence

$$c_\omega \int_{\tau_0}^T \mathcal{D}_\omega(\tau) d\tau \leq \mathcal{H}_\omega(\tau_0) - \mathcal{H}_\omega(T) \leq \mathcal{H}_\omega(\tau_0) - \inf_{\sigma \geq \tau_0} \mathcal{H}_\omega(\sigma).$$

The right-hand side is finite and independent of T . Letting $T \rightarrow \infty$ and using monotone convergence for the nonnegative cost proves the claim. $\square$

Theorem 8.64 (CostPhase 6: local Navier–Stokes cost-to-exclusion theorem). *Let ω be a renormalized Navier–Stokes Type II branch on a retained terminal compact stratum. If ω has active supercritical scaling, the validated compact Type II cost-divergence conclusion, and local cost-to-exclusion data of [definition 8.57](#), then ω is excluded from the Type II class under consideration. Equivalently, on this local stratum,*

$$\Pi_{\text{NS,II}}(\omega) : \mathcal{S}_\omega \wedge \mathcal{B}_\omega \implies \mathcal{E}_\omega$$

is a theorem.

Proof. Assume, for contradiction, that the branch remains in the Type II class on the retained compact stratum. The local cost-to-exclusion data give a corrected localized energy inequality for the same nonnegative cost $\mathcal{D}_\omega$ whose divergent tail integral is the validated compact cost-divergence conclusion. By [lemma 8.63](#),

$$\int_{\tau_0}^{\infty} \mathcal{D}_\omega(\tau) d\tau < \infty.$$

This contradicts the validated compact cost-divergence conclusion. Hence the branch cannot remain in the Type II class under consideration. $\square$

Definition 8.65 (CostPhase 7: pointwise cost-divergence exclusion data). For a renormalized Type II solution ω , the pointwise cost-divergence exclusion data are available when the following local analytic statements hold:

- (i) class membership: ω lies in the renormalized Navier–Stokes Type II class and satisfies the Type II analytic data before the Type II cost criterion is evaluated;
- (ii) cost compatibility: the active compact Type II cost criterion uses either the Navier–Stokes identity-cost criterion

$$\mathfrak{C}_{\text{II}}^{NS} = \tilde{\mathfrak{D}}_{R_0},$$

or a localized cost explicitly comparable to $\tilde{\mathfrak{D}}_{R_0}$ on a terminal tail in the sense of [definition A.23](#); in either case its conclusion is the validated compact Type II cost-divergence conclusion for ω ;

- (iii) stability of the active scaling condition: the active supercritical scaling condition remains in force after the repaired-gauge representation, cost comparison, and criterion evaluation are attached;
- (iv) stability of the cost-divergence conclusion: every renormalized continuation preserving the active scaling condition and the validated compact Type II cost-divergence conclusion remains in the same compact Type II cost-divergence regime, unless it first enters one of the local alternatives already isolated above;
- (v) local monotonicity-carrier data: on the retained compact stratum, the validated compact cost admits a carrier in the sense of [definition 8.59](#). This may be supplied by the fixed or moving carrier subidentity of [definition 8.88](#), the corresponding finite-error condition, and the corresponding cost comparison. Exact fixed or moving localized energy identities are stronger sufficient inputs through [lemmas 8.61](#) and [8.62](#).

Lemma 8.66 (CostPhase 2: canonical cost compatibility). *Let ω be represented on a terminal tail by a repaired-gauge Navier–Stokes branch (V, P, a, b) . Assume that the localized identity cost*

$$\tilde{\mathfrak{D}}_{R_0}(\tau) = \nu \int_{\mathbb{R}^3} |\nabla V(y, \tau)|^2 \phi_{R_0}(y) dy + a_+(\tau) \int_{\mathbb{R}^3} |V(y, \tau)|^2 \phi_{R_0}(y) dy$$

is measurable, nonnegative, and locally integrable on finite τ -intervals. If the compact Type II criterion is evaluated with the canonical Navier–Stokes identity cost, then

$$\mathfrak{C}_{\text{II}}^{NS} = \tilde{\mathfrak{D}}_{R_0}$$

is a valid Type II scale-exclusion cost, the comparison holds with

$$c_0 = 1, \quad \tau_1 = \tau_0,$$

and any later compact cost-divergence assertion for ω is exactly the identity-cost divergence statement

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty.$$

If the compact criterion is not evaluated with this identity cost, the branch is not being treated by the canonical compatibility statement and must instead come with an explicit comparison as in [lemma 8.67](#).

Proof. The first assertion is the identity-cost case of the localized cost definitions. Since $\tilde{\mathfrak{D}}_{R_0}$ is assumed measurable, nonnegative, and locally integrable on finite intervals, the canonical choice

$$\mathfrak{C}_{\text{II}}^{NS} := \tilde{\mathfrak{D}}_{R_0}$$

satisfies the defining properties of a Type II scale-exclusion cost. The comparison inequality is the identity

$$\mathfrak{C}_{\text{II}}^{NS}(\tau) = \tilde{\mathfrak{D}}_{R_0}(\tau) \geq 1 \cdot \tilde{\mathfrak{D}}_{R_0}(\tau)$$

for a.e. $\tau \geq \tau_0$, so $c_0 = 1$ and $\tau_1 = \tau_0$. Therefore divergence of the active compact cost is precisely divergence of $\tilde{\mathfrak{D}}_{R_0}$ on the terminal tail. This is the canonical case of [lemma A.24](#) and [theorem A.25](#). $\square$

Lemma 8.67 (CostPhase 2: comparable cost compatibility). *Let ω be represented on a terminal tail by a repaired-gauge Navier–Stokes branch (V, P, a, b) . Let $\mathfrak{C}_{\text{II}}^{NS}$ be a nonnegative, measurable, locally integrable localized cost on finite τ -intervals. Assume that, on a later terminal tail $[\tau_1, \infty)$, there is a constant $c_0 > 0$ such that*

$$\mathfrak{C}_{\text{II}}^{NS}(\tau) \geq c_0 \tilde{\mathfrak{D}}_{R_0}(\tau) \quad \text{for a.e. } \tau \geq \tau_1.$$

Then every identity-cost divergence conclusion

$$\int_{\tau_1}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty$$

implies the corresponding compact Type II cost-divergence conclusion

$$\int_{\tau_1}^{\infty} \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau = \infty.$$

Consequently the use of $\mathfrak{C}_{\text{II}}^{NS}$ introduces no additional exit from the compact-cost stratum beyond the explicit failure of this comparison.

Proof. For every $T > \tau_1$, the assumed comparison gives

$$\int_{\tau_1}^T \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau \geq c_0 \int_{\tau_1}^T \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau.$$

Letting $T \rightarrow \infty$ proves the implication from identity-cost divergence to divergence of the selected compact Type II cost. This is precisely the one-sided comparison needed by [definition A.23](#) and used in [theorem A.25](#). If the inequality is unavailable, the selected diagnostic cost is not connected to the Navier–Stokes identity dissipation on this terminal tail; this is a cost-comparison failure, not a new unlabelled Type II branch. $\square$

Lemma 8.68 (CostPhase 1: local cost well-posedness). *Let (V, P, a, b) be a repaired-gauge branch on a finite renormalized interval $I = [\alpha, \beta]$, and let $\phi_{R_0} \in C_c^\infty(B_{R_0})$, $\phi_{R_0} \geq 0$. Assume:*

- (i) $V \in L^2(I; H^1(B_{R_0}))$;
- (ii) $a \in L^1(I)$;
- (iii) either $V \in L^\infty(I; L^3(\mathbb{R}^3))$, or more locally

$$\text{ess sup}_{\tau \in I} \|V(\tau)\|_{L^2(B_{R_0})} < \infty.$$

Then the localized identity cost

$$\tilde{\mathfrak{D}}_{R_0}(\tau) = \nu \int |\nabla V(y, \tau)|^2 \phi_{R_0}(y) dy + a_+(\tau) \int |V(y, \tau)|^2 \phi_{R_0}(y) dy$$

is measurable, nonnegative, and integrable on I . Consequently, on every compact renormalized window satisfying these local bounds, $\tilde{\mathfrak{D}}_{R_0}$ is a well-defined finite-interval cost.

Proof. The two summands are measurable because V is strongly measurable in the local Sobolev spaces on the finite interval and a_+ is measurable. They are nonnegative since $\nu > 0$, $\phi_{R_0} \geq 0$, and $a_+ \geq 0$.

For the gradient term, boundedness of ϕ_{R_0} and the support condition give

$$\int_I \int |\nabla V|^2 \phi_{R_0} \leq \|\phi_{R_0}\|_{L^\infty} \int_I \|\nabla V(\tau)\|_{L^2(B_{R_0})}^2 d\tau < \infty.$$

For the drift term, first suppose

$$\operatorname{ess\,sup}_{\tau \in I} \|V(\tau)\|_{L^2(B_{R_0})} < \infty.$$

Then

$$\int_I a_+(\tau) \int |V|^2 \phi_{R_0} dy d\tau \leq \|\phi_{R_0}\|_{L^\infty} \left(\operatorname{ess\,sup}_{\tau \in I} \|V(\tau)\|_{L^2(B_{R_0})}^2 \right) \|a_+\|_{L^1(I)} < \infty.$$

If instead $V \in L^\infty(I; L^3(\mathbb{R}^3))$, then Hölder on B_{R_0} gives the local L^2 bound

$$\|V(\tau)\|_{L^2(B_{R_0})}^2 \leq |B_{R_0}|^{1/3} \|V(\tau)\|_{L^3(\mathbb{R}^3)}^2$$

for a.e. $\tau \in I$, reducing to the previous case. Thus both summands are integrable on I . $\square$

Lemma 8.69 (CostPhase 0: local modulation integrability from absolutely continuous repaired gauges). *Let u, p be a Navier–Stokes solution on a terminal physical time interval, and let (V, P) be its repaired-gauge renormalization on a terminal renormalized tail. Assume that, on every compact renormalized subinterval, the physical concentration chart and the repaired-gauge path satisfy [definitions 10.46](#) and [10.48](#). Let a, b be the final modulation coefficients in the repaired-gauge renormalized equation, as in [lemma 10.51](#). Then*

$$a, b \in L_{\operatorname{loc}}^1([\tau_0, \infty)).$$

In particular, for every unit terminal window

$$I_n = [\tau_0 + n, \tau_0 + n + 1],$$

one has $a \in L^1(I_n)$. If the absolutely continuous chart, the absolutely continuous repaired gauge, or the resulting local integrability of the modulation coefficients is unavailable on such a compact window, then the branch does not lie in the retained compact-cost stratum and falls into the corresponding representation, repaired-gauge, or modulation-integrability alternative.

Proof. On each compact renormalized window, the assumed absolutely continuous chart and repaired-gauge path place the branch under the hypotheses of [theorem 10.56](#). In particular, [lemma 10.55](#) gives

$$a, b \in L_{\operatorname{loc}}^1,$$

and [lemma 10.60](#) identifies these locally integrable functions as the final coefficients in the repaired-gauge equation. Restricting to I_n gives $a \in L^1(I_n)$. If the absolutely continuous local chart, repaired-gauge path, or resulting coefficient integrability is missing, then the ordered stratification records the first missing input as a representation, repaired-gauge, or modulation-integrability failure. $\square$

Proposition 8.70 (CostPhase 0: upstream production of compact-window state data). *Let*

$$V_n(y, s) = V(y, \tau_n + s), \quad P_n(y, s) = P(y, \tau_n + s), \quad s \in (0, T),$$

be a selected repaired-gauge window satisfying the hypotheses of [proposition 7.14](#). Assume also that the represented equation and suitable local energy inequality hold on the working selected cylinder. Fix $0 < R_0 < R_1 < \infty$. Then

$$V_n \in L^\infty(0, T; L^2(B_{R_1})) \cap L^2(0, T; H^1(B_{R_1})),$$

$$P_n - (P_n)_{B_{R_1}} \in L^{3/2}((0, T) \times B_{R_1}).$$

If the selected chart and repaired-gauge path are absolutely continuous on this window, then $a, b \in L^1(\tau_n, \tau_n + T)$. Hence, after choosing the spatial and temporal buffers in [definition 8.54](#), the selected window carries all phase 0 retained local cost state data. No cost-validation, carrier, or cost-divergence statement is used.

Proof. [Proposition 7.14](#), applied with $K = B_{R_1}$, gives the displayed velocity bounds and the mean-subtracted pressure bound. Absolute continuity of the chart and repaired gauge gives $a, b \in L^1$ by [lemma 8.69](#). The additional hypothesis supplies the represented equation and suitable local energy inequality. Finally choose $\phi_{R_0} \in C_c^\infty(B_{R_1})$, equal to one on B_{R_0} , and a temporal cutoff inside the selected working interval. These are precisely items (D1)–(D5) of [definition 8.54](#). $\square$

Corollary 8.71 (CostPhase 1: compact-window data give cost well-posedness). *Let (V, P, a, b) be a repaired-gauge branch on a terminal tail. Fix $R_0 < \infty$. Suppose that, for every unit terminal window*

$$I_n = [\tau_0 + n, \tau_0 + n + 1], \quad n \in \mathbb{N},$$

there is $R_1 > R_0$ for which the branch carries the retained local cost state data of [definition 8.54](#) on $B_{R_0} \times I_n$. Then $\tilde{\mathfrak{D}}_{R_0}$ is a measurable nonnegative cost, integrable on each I_n . Thus the canonical identity cost is well posed on every retained terminal window; its selection or comparison with another cost is a separate phase 2 step.

Failure of the conclusion is exactly one of the phase 0 state-data failures. Cost comparison is tested only after this conclusion has been established.

Proof. [Definition 8.54](#)(D2) supplies $L^\infty(I_n; L^2(B_{R_1}))$ and

$$V \in L^2(I_n; H^1(B_{R_1})),$$

while item (D4) supplies $a \in L^1(I_n)$. Restriction from B_{R_1} to B_{R_0} gives the hypotheses of [lemma 8.68](#) on each I_n , which proves the asserted phase 1 well-posedness. Conversely, if the conclusion fails before cost selection, at least one phase 0 input is absent; by [definition 8.55](#), that absence is recorded as the corresponding local alternative before any cost divergence is tested. $\square$

Proposition 8.72 (Active scaling stability criterion). *Let ω be an admissible local Type II branch with active supercritical scaling on a terminal tail. A passage from a raw Type II chart to a repaired-gauge chart is called scale-admissible on a later tail if its raw and repaired physical scales satisfy*

$$0 < m \leq \frac{\lambda_{\text{rep}}(t)}{\lambda_{\text{raw}}(t)} \leq M < \infty$$

on that tail. On the compact-cost stratum, the active supercritical scaling condition is unchanged by the following operations:

- (i) *restriction to a later terminal tail;*

- (ii) *passage from an admissible raw Type II chart to a scale-admissible repaired-gauge representation;*
- (iii) *replacement of the canonical cost by an explicitly comparable localized cost;*
- (iv) *evaluation of the compact Type II cost criterion.*

If the repaired chart is not scale-admissible on any later tail, the branch is not retained in the compact-cost stratum; it is assigned instead to the corresponding representation, repaired-gauge, or scale-degeneracy alternative.

Proof. The active supercritical scaling condition is a terminal property of the scale-center chart and is invariant under removal of a finite initial renormalized interval. Thus restriction from $[\tau_0, \infty)$ to $[\tau_1, \infty)$, $\tau_1 \geq \tau_0$, cannot change the terminal scale alternative.

Suppose next that a scale-admissible repaired-gauge representation is available. Two-sided comparability of λ_{rep} and λ_{raw} implies that any terminal statement depending only on the scale class, such as supercriticality relative to the Type I rate, holds for one scale if and only if it holds for the other after passing to the same tail. The repaired gauge changes coordinates and modulation variables inside this comparable scale class; it does not replace the branch by a different terminal scaling alternative. If the representation cannot be constructed, or if every repaired chart loses scale comparability, then the first failure is a representation, repaired-gauge, or scale-degeneracy alternative.

Cost comparison is scale-passive. It changes only the nonnegative diagnostic functional used to record cost divergence and does not alter the center, physical scale, renormalized time, or active scaling hypothesis. If the comparison is unavailable, the failure is cost-comparison failure. Finally, criterion evaluation is a logical operation applied to the already constructed branch and cost. It does not modify the branch or its scale. Hence active scaling is stable through the listed operations exactly on the local stratum on which those operations are admissible. $\square$

Lemma 8.73 (CostPhase 7: final compact-cost validation ledger). *Let ω be an admissible local Type II branch with active supercritical scaling on a terminal tail. Before declaring that a continuation lies in the compact cost-divergence stratum, run the ordered local validation tests:*

- (i) *Type II analytic data and exhaustion data;*
- (ii) *raw-chart selection, repaired-gauge representation, and scale admissibility;*
- (iii) *local compactness decomposition into the alternatives of [theorem 11.18](#);*
- (iv) *retained finite critical-mass window, read as the compact-window finite critical-mass input of [lemma 10.78](#);*
- (v) *the critical-tightness/exterior-leakage test at the retained scale;*
- (vi) *local windowed H^1 control;*
- (vii) *local modulation integrability, cost well-posedness, and compatibility of the chosen compact Type II cost;*
- (viii) *local monotonicity-carrier data for the validated compact cost, supplied by the fixed-cutoff criterion, the moving-cutoff criterion, or another explicit localized carrier;*
- (ix) *the compact Type II cost-divergence conclusion for the validated cost.*

Then exactly one of the following occurs:

- (a) *all tests pass, and ω remains in the compact Type II cost-divergence regime on the retained terminal branch;*
- (b) *the first failed test gives one of the already named local alternatives: ordered exhaustion failure, representation failure, repaired-gauge or scale-degeneracy failure, multibubble/multicore splitting, cascade, exterior tightness failure, loss of local windowed H^1 control, finite-cost scale-collapse alternative, scale-rigid residual alternative, critical-mass alternative, modulation-integrability failure, cost-compatibility failure, carrier-subidentity*

failure, finite fixed-cutoff error failure, summable moving-annulus error failure, carrier-comparison failure for the validated cost, or failure of compact cost divergence for the validated cost.

Thus a renormalized continuation preserving active scaling and validated compact cost divergence cannot leave the compact cost-divergence regime without being recorded by the local compactness stratification. This is a final audit of phases 0–6. It does not supply the local state data, define the cost, construct a carrier, or prove cost divergence.

Proof. The proof is the ordered-alternative argument. First test whether the Type II analytic and exhaustion data hold. If not, the first failure is one of the ordered exhaustion failures of [definition A.3](#). If they hold, test for raw-chart selection, repaired-gauge representation, and scale admissibility. Failure at this stage is one of the ordered representation failures of [definition A.11](#), possibly with the specific scale-degeneracy recorded by [proposition 8.72](#).

Once a scale-admissible repaired branch is available, apply the local compactness decomposition [theorem 11.18](#). If the branch falls into the multibubble, multicore, cascade, gauge-degenerate, finite-cost scale-collapse, local H^1 -loss, or scale-rigid residual alternatives, it has left the compact single-core cost-divergence stratum through a named local case. In particular, an additional exterior concentration mechanism is not discarded; if it retains critical mass it appears as a multibubble or multicore branch, and if it escapes the retained compact region it appears as failure of critical tightness.

The first alternative in [theorem 11.18](#) is retained only as a compact single-core branch; it is not yet identified with compact cost divergence. On that retained branch the remaining tests below decide whether the branch belongs to the compact cost-divergence stratum, has finite validated cost, or exits through one of the remaining compact tests.

On the retained compact single-core branch, test the finite critical-mass row using the compact-window version of [lemma 10.78](#): failure of the retained spacetime activity floor means non-retention, and failure of a finite local upper bound produces a comparable retained label or a named adjacent active label. Thus failure of the finite critical-mass row is recorded as a critical-mass or adjacent local alternative, not as a hidden global estimate. Once the row passes, the remaining compact tests are critical tightness and local windowed H^1 control. Failure of either is precisely one of the alternatives in [definition A.26](#).

If all compact tests pass, local cost well-posedness and compatibility are evaluated. Local modulation integrability is supplied by [lemma 8.69](#); the local cost well-posedness is supplied by [corollary 8.71](#). The compatibility is the identity case of [lemma 8.66](#) or the explicit comparison case of [lemma 8.67](#); failure of both is a cost-compatibility failure. Next test whether the validated compact cost has a local monotonicity carrier. In the fixed-cutoff case this means the fixed carrier subidentity, finite fixed-cutoff error, and energy-control comparison. In the moving-cutoff case this means the moving carrier subidentity, summable moving-annulus errors, and moving-cost comparison. Exact localized energy identities are sufficient for these subidentities, but the one-sided suitable local energy inequality is the actual input needed. Failure of these inputs is recorded as a carrier-subidentity failure, a finite fixed-cutoff error failure, a summable moving-annulus error failure, or a carrier-comparison failure for the validated cost. Finally, evaluate the compact cost-divergence conclusion for the validated cost. If it fails, the branch is a finite-cost or non-divergent compact branch and is not counted as a compact cost-divergence branch. The tests are ordered, so the first failure is unique. If no failure occurs, the branch has the analytic data, scale-admissible repaired-gauge representation, compact single-core local compactness position, the retained finite critical-mass window, critical tightness, local windowed H^1 control, local modulation integrability, a compatible compact cost, local monotonicity-carrier data, and the validated compact cost-divergence conclusion. Together with the assumed active scaling, these

hypotheses place the branch in the compact Type II cost-divergence regime. Therefore there is no unrecorded exit. $\square$

Theorem 8.74 (CostPhase 5: carrier validation alternative on the retained compact stratum). *Let ω be an admissible local Type II branch with active supercritical scaling. Suppose that the pre-divergence rows of [definition 8.55](#) have reached the retained compact single-core branch and that cost well-posedness and cost compatibility have passed. Then exactly one of the following alternatives occurs:*

- (i) *the validated compact cost admits local monotonicity-carrier data in the sense of [definition 8.59](#);*
- (ii) *the branch leaves the retained compact pre-divergence channel through the first failed carrier-validation test: carrier-subidentity failure, finite fixed-cutoff error failure, summable moving-annulus error failure, or carrier-comparison failure for the validated cost.*

Consequently, on the retained compact carrier channel, where the adjacent carrier-failure alternatives are not part of the stratum, the validated compact cost admits local monotonicity-carrier data. This alternative is independent of whether the selected cost later has finite or infinite tail integral.

Proof. After the branch has passed the Type II analytic, representation, compactness, retained finite critical mass, critical-tightness, local windowed H^1 , modulation, cost well-posedness, and cost-compatibility tests, the next ordered test in [definition 8.55](#) is precisely the local monotonicity-carrier test for the validated compact cost. If this test passes, alternative (i) holds. If it fails, the first failed input is recorded as one of the named adjacent carrier failures in that definition: failure of the carrier subidentity, failure of the fixed finite-error condition, failure of the moving finite-error condition through the summable moving-annulus test, or failure of the carrier comparison for the validated cost. The ordered first-failure convention makes the two alternatives disjoint and exhaustive.

By definition, the retained compact carrier channel keeps only the branch for which all preceding compact tests and the carrier-validation test have passed. Therefore a branch that remains in this channel cannot realize one of the carrier-failure alternatives, and the local monotonicity carrier exists. $\square$

Proposition 8.75 (CostPhase 7: local verification of the cost-exclusion data). *On a retained terminal branch in the compact Type II cost-divergence stratum, the pointwise data of [definition 8.65](#) are local analytic statements, not additional global assumptions. More precisely, assume that the branch:*

- (i) *has the local Type II analytic data and has not entered an exhaustion, representation, modulation-integrability, multibubble/multicore, cascade, exterior-tightness, scale-collapse, scale-rigid, critical-mass, or local H^1 -loss alternative;*
- (ii) *admits a scale-admissible repaired-gauge representation on retained windows from absolutely continuous local scale-center and repaired-gauge data;*
- (iii) *has the retained finite critical-mass window, critical tightness at the retained scale, the retained local cost state data of [definition 8.54](#) on every cost window, active supercritical scaling, and the compact cost-divergence conclusion for the validated cost;*
- (iv) *uses either the canonical identity cost $\mathfrak{D}_{R_0}$ or an explicitly comparable localized cost;*
- (v) *remains in the retained compact cost-divergence stratum after the carrier-validation alternatives of [theorem 8.74](#) have been evaluated.*

Then class membership, cost compatibility, stability of the active scaling condition, stability of the validated cost-divergence conclusion, and the cost-to-exclusion implication all hold for this branch. Consequently, on this local stratum, all clauses of [definition 8.65](#) are verified; no estimate

outside the retained local stratum and no exclusion of secondary exterior concentration are being assumed.

Proof. Class membership is the retained local Type II analytic-data alternative after the ordered exhaustion and representation tests have passed. The phase 0 inputs needed to evaluate the compact cost are assumed in [definition 8.54](#); their standard selected-window producer is [proposition 8.70](#). Phase 1 well-posedness is [corollary 8.71](#). Cost compatibility then follows from [lemma 8.66](#) in the identity-cost case and from [lemma 8.67](#) in the explicitly comparable-cost case. Stability of the active scaling condition is [proposition 8.72](#). Stability of the validated cost-divergence conclusion, with all exits assigned to named local alternatives, is [lemma 8.73](#). These prove clauses (i)–(iv) of [definition 8.65](#). Since the branch remains in the retained compact cost-divergence stratum after carrier validation, [theorem 8.74](#) gives local monotonicity-carrier data for the validated compact cost. These data provide the local cost-to-exclusion data through [lemma 8.60](#). Then [theorem 8.64](#) gives $\Pi_{\text{NS,II}}(\omega)$, which is the analytic content of clause (v). $\square$

Definition 8.76 (Classwise retained compact coverage). The renormalized Type II class has classwise retained compact coverage for the validated compact cost if every branch with active supercritical scaling and validated compact cost divergence either remains in the retained compact cost-divergence stratum after the ordered validation tests of [lemma 8.73](#), or has already entered one of the named adjacent alternatives in that ordered stratification.

Theorem 8.77 (Classwise retained compact coverage from ordered validation). *Every admissible local Type II branch with active supercritical scaling and validated compact Type II cost divergence has classwise retained compact coverage in the sense of [definition 8.76](#).*

Proof. Apply the ordered validation tests of [lemma 8.73](#) to the branch. The lemma gives two possibilities. If all tests pass, the branch remains in the compact Type II cost-divergence regime on the retained terminal branch, which is precisely the retained compact cost-divergence stratum.

If a test fails, the first failed test is one of the named local alternatives in [lemma 8.73](#). Since the present theorem assumes the validated compact cost-divergence conclusion, the last listed failure, failure of compact cost divergence for the validated cost, cannot occur. Every other first failure is an adjacent compactness stratum: exhaustion, representation, repaired-gauge or scale-degeneracy failure, multibubble/multicore splitting, cascade, exterior tightness failure, local H^1 -loss, finite-cost scale-collapse, scale-rigid residual behavior, critical-mass failure, modulation-integrability failure, cost-compatibility failure, carrier-subidentity failure, fixed-cutoff finite-error failure, moving-annulus summability failure, or carrier-comparison failure. Therefore every branch satisfying the hypotheses either remains retained or has already entered a named adjacent alternative. This is exactly classwise retained compact coverage. $\square$

Remark 8.78 (Scope of classwise retained compact coverage). [Proposition 8.75](#) is deliberately local. It proves the interface hypotheses for a branch only after the branch has been placed in the validated compact cost-divergence stratum by the compactness stratification. Branches with multibubble, cascade, exterior leakage, scale-collapse, scale-rigid, or rough-core behavior are not forced into this stratum; they are assigned to their own local alternatives. Thus the cost-divergence exclusion step is unconditional on the retained compact stratum, while a classwise statement is obtained only after the adjacent strata have been treated by their own local arguments. The coverage part itself is no longer an assumption; it is supplied by [theorem 8.77](#).

Definition 8.79 (Abstract cost-divergence exclusion theorem applicability). An abstract Type II cost-divergence exclusion theorem is applicable to a renormalized Type II solution if the following analytic identifications are available:

- (i) an abstract Type II exclusion theorem based on localized renormalization-cost divergence;

- (ii) hypotheses supplying supercritical scaling, a Type II concentration scenario with bounded physical energy, and a localized energy monotonicity formula at the active scale $\lambda(t)$;
- (iii) an identification of the theorem's renormalization-cost integral with a validated compact Navier–Stokes cost $\mathfrak{C}_{\text{II}}^{NS}$, where either

$$\mathfrak{C}_{\text{II}}^{NS} = \tilde{\mathfrak{D}}_{R_0}$$

or $\mathfrak{C}_{\text{II}}^{NS}$ is a localized cost satisfying the terminal one-sided comparison with $\tilde{\mathfrak{D}}_{R_0}$ in [definition A.23](#);

- (iv) a domain embedding placing the repaired-gauge Navier–Stokes orbit and its scale variables into the parabolic domain expected by the theorem without changing the active scaling condition, validated cost-divergence conclusion, or exclusion conclusion;
- (v) identification of the theorem's conclusion with [definition 8.56](#).

Lemma 8.80 (Applicability of the abstract cost-divergence theorem). *Assume a renormalized Type II solution has the active supercritical scaling condition, the validated compact Type II cost-divergence conclusion, and the abstract applicability data of [definition 8.79](#). Then the solution is excluded from the Type II class under consideration.*

Proof. The abstract Type II theorem is applied with supercritical scaling, bounded-energy Type II concentration, and the localized energy monotonicity formula at the active scale. The cost identification places the validated compact Navier–Stokes cost in the theorem's renormalization-cost integral. In the identity case this cost is $\tilde{\mathfrak{D}}_{R_0}$. In the comparable-cost case, the terminal one-sided comparison in [definition A.23](#) and [lemma 8.67](#) show that the validated Navier–Stokes cost-divergence conclusion implies divergence of the localized cost used by the abstract theorem. The domain embedding places the repaired-gauge Navier–Stokes solution in the theorem's parabolic domain, and the conclusion is identified with [definition 8.56](#). Therefore the theorem gives exclusion from the Type II class. $\square$

Theorem 8.81 (Pointwise Navier–Stokes cost-divergence exclusion). *Let ω be a renormalized Type II solution. If ω has the active supercritical scaling condition, the validated compact Type II cost-divergence conclusion, and the pointwise data of [definition 8.65](#), then ω is excluded from the Type II class under consideration.*

Proof. Class membership and stability of the active scaling condition ensure that ω is evaluated in the renormalized Type II class and that the active supercritical scaling condition is still the relevant one. Cost compatibility ensures that the validated compact cost-divergence conclusion is the Navier–Stokes compact cost conclusion for ω , not an external or incomparable cost statement. Stability of the cost-divergence conclusion rules out an exit from the compact Type II cost-divergence regime to a Type II case outside the compact cost-divergence criterion inside the same class.

It remains only to justify the transition from cost divergence to exclusion. The local monotonicity-carrier data in [definition 8.65](#) give the local cost-to-exclusion data by [lemma 8.60](#). Therefore [theorem 8.64](#) applies to the active supercritical scaling condition and the validated compact cost-divergence conclusion, giving the claimed exclusion. In an abstract theorem setting, the same implication is also licensed by [lemma 8.80](#). $\square$

Corollary 8.82 (Cost-divergence exclusion on the retained compact stratum). *Let ω be a renormalized Type II solution with active supercritical scaling and validated compact Type II cost divergence. If ω remains in the retained compact cost-divergence stratum after the ordered validation tests of [lemma 8.73](#), then ω is excluded from the Type II class under consideration.*

Proof. By [proposition 8.75](#), the pointwise cost-divergence exclusion data of [definition 8.65](#) are verified on this retained stratum. Applying [theorem 8.81](#) gives the exclusion conclusion. $\square$

Definition 8.83 (CostPhase 5: pre-divergence adjacent-stratum closure). The compact-cost regime has adjacent-stratum closure if every named adjacent alternative produced by [definition 8.55](#) for a branch with active supercritical scaling has its own stratum-local argument giving the same Type II exclusion conclusion, either directly or by placing the branch in a named non-compact-cost alternative whose closure theorem already gives the Type II exclusion conclusion. Neither the definition nor any phase 5 proof of this property may use divergence of the selected cost.

Corollary 8.84 (Uniform Navier–Stokes cost-divergence exclusion). *Assume adjacent-stratum closure for the compact-cost regime in the sense of [definition 8.83](#). Then every renormalized Type II solution with active supercritical scaling condition and validated compact Type II cost-divergence conclusion is either excluded in its adjacent stratum or excluded by [corollary 8.82](#).*

Proof. By [theorem 8.77](#), a branch with active supercritical scaling and validated compact cost divergence either lies in one of the named adjacent strata or remains in the retained compact cost-divergence stratum. The adjacent cases are handled by the adjacent-stratum closure hypothesis. In the retained compact case, [corollary 8.82](#) applies. $\square$

Definition 8.85 (Adjacent exits from cost-divergence exclusion). If the active supercritical scaling condition and validated compact cost-divergence conclusion are present but the branch is not retained in the compact cost-divergence stratum, then the first exit from the cost-exclusion argument lies in the ordered list:

class membership, cost compatibility, stability of the active scaling condition,
stability of the validated cost-divergence conclusion, local monotonicity-carrier validation.

On the retained compact stratum none of these exits occurs. The final item is expanded by [theorem 8.74](#) into the carrier-subidentity, fixed-finite-error, moving-annulus, and carrier-comparison adjacent strata.

Lemma 8.86 (Carrier exits are adjacent strata, not retained compact branches). *Assume a renormalized Type II solution already has the active supercritical scaling condition and the validated compact Type II cost-divergence conclusion. If one of the exits in [definition 8.85](#) occurs, then the branch is not in the retained compact cost-divergence stratum. If no such exit occurs, the local cost-to-exclusion data are available.*

Proof. The retained compact cost-divergence stratum is reached only after the ordered tests of [lemma 8.73](#) have passed. Hence a failure of class membership, cost compatibility, active-scaling stability, cost-divergence stability, or carrier validation is by definition an exit to an adjacent stratum. Conversely, if no such exit occurs, the carrier-validation test has passed. By [theorem 8.74](#), the validated compact cost admits local monotonicity-carrier data, and then [lemma 8.60](#) gives the local cost-to-exclusion data. $\square$

Definition 8.87 (CostPhase 2: pre-divergence carrier channel). A renormalized Type II branch lies in the pre-divergence carrier channel if, on a terminal tail, it has active supercritical scaling and the pre-divergence rows of [definition 8.55](#) have passed through the following stages:

- (i) the Type II analytic data and ordered exhaustion tests;
- (ii) raw-chart selection, repaired-gauge representation, pressure reconstruction, modulation data, and scale-admissibility of the repaired chart;
- (iii) the local compactness decomposition without entering a multibubble, multicore, cascade, exterior-tightness, local H^1 -loss, finite-cost scale-collapse, or scale-rigid residual alternative;

- (iv) the retained finite critical-mass window, namely the compact-window finite critical-mass input supplied by [lemma 10.78](#);
- (v) critical tightness at the retained scale and the retained local cost state data of [definition 8.54](#) on each selected cost window;
- (vi) local cost well-posedness and compatibility of the chosen compact Type II cost with the Navier–Stokes identity cost.

This channel stops immediately before the local energy identity, finite-error, and carrier-routing phases. It contains the interface data needed to construct the fixed- or moving-cutoff carrier, but it asserts neither carrier completion nor divergence of the selected cost. The finite critical-mass row is read locally in the final no-Type-II assembly: failure of compact-window lower or upper critical mass is routed by the local terminal critical-mass and exterior/state-space alternatives.

Definition 8.88 (CostPhase 3: fixed and moving carrier subidentities). For the fixed cutoff ϕ_{R_0} , the fixed carrier subidentity is the distributional estimate

$$\frac{d}{d\tau}\mathcal{E}_\phi(\tau) + \frac{1}{2}\tilde{\mathfrak{D}}_{R_0}(\tau) \leq B_{R_0}(\tau), \quad \mathcal{E}_\phi(\tau) = \frac{1}{2} \int |V|^2 \phi_{R_0}.$$

For a moving cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$, the moving carrier subidentity is the distributional estimate

$$\frac{d}{d\tau}\mathcal{E}_R(\tau) + \frac{1}{2}\tilde{\mathfrak{D}}_{R(\tau)}(\tau) \leq B_{\text{mov}}(\tau), \quad \mathcal{E}_R(\tau) = \frac{1}{2} \int |V|^2 \Phi.$$

Here B_{R_0} and B_{mov} are the absolute error densities of [definitions 10.9](#) and [10.18](#). The subidentity is the one-sided estimate needed for corrected monotonicity; an exact localized energy identity is sufficient but not necessary.

Lemma 8.89 (CostPhase 3: suitable local energy gives carrier subidentities). *Assume the repaired-gauge representation of a pre-carrier branch includes the local suitable energy inequality on compact renormalized cylinders. Then the fixed carrier subidentity holds for every nonnegative compactly supported fixed cutoff. The moving carrier subidentity holds for every nonnegative finite-valued moving cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$ with $R \in AC_{\text{loc}}$, after the usual approximation of R by smooth positive radii on compact τ -intervals.*

Proof. Apply the local suitable energy inequality in the physical variables to the pullback of the nonnegative renormalized test function. The repaired-gauge change of variables is locally absolutely continuous in time and smooth in space on compact cylinders, so the distributional chain rule gives the renormalized local energy inequality with the drift terms $a(V + y \cdot \nabla V)$ and $b \cdot \nabla V$. For a fixed cutoff this is the fixed localized energy identity with equality replaced by $\leq$. Adding $\frac{1}{2}\tilde{\mathfrak{D}}_{R_0}$ to both sides and bounding the cutoff, pressure, translation, and scale terms by their absolute values gives the fixed carrier subidentity.

For a moving cutoff, use smooth positive approximations $R_k \rightarrow R$ in AC_{loc} on each compact time interval. The local energy inequality applies to $\Phi_k(\tau, y) = \phi(y/R_k(\tau))$, and the preceding argument gives the moving subidentity with $B_{\text{mov},k}$. Passing to the limit in distributions uses the local finite-energy and local H^1 bounds already present on the pre-carrier regime. The term $\partial_\tau \Phi$ is interpreted as the L^1_{loc} derivative coming from $R \in AC_{\text{loc}}$. The limit is exactly the moving carrier subidentity. $\square$

Definition 8.90 (Fixed-cutoff localized-energy closure criterion). A branch in the pre-carrier validated compact-cost regime satisfies the fixed-cutoff localized-energy closure criterion if, for the cutoff ϕ_{R_0} used to evaluate the validated compact cost and for

$$\mathcal{E}_\phi(\tau) = \frac{1}{2} \int |V(y, \tau)|^2 \phi_{R_0}(y) dy,$$

the following local analytic statements hold on a terminal tail:

- (i) $\mathcal{E}_\phi(\tau_0) < \infty$;
- (ii) the fixed carrier subidentity of [definition 8.88](#) is valid for the repaired-gauge branch;
- (iii) the fixed-cutoff error density B_{R_0} of [definition 10.9](#) has finite tail integral,

$$\int_{\tau_0}^{\infty} B_{R_0}(\tau) d\tau < \infty;$$

- (iv) the validated compact cost is energy-controlled by $\tilde{\mathfrak{D}}_{R_0}$ in the sense of [definition 8.58](#).

Definition 8.91 (Moving-cutoff localized-energy closure criterion). A branch in the pre-carrier validated compact-cost regime satisfies the moving-cutoff localized-energy closure criterion if there is a finite-valued nondecreasing locally absolutely continuous radius $R(\tau) \geq R_0$, with $R(\tau) \rightarrow \infty$ and moving cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$, such that, on a terminal tail:

- (i)

$$\mathcal{E}_R(\tau_0) = \frac{1}{2} \int |V(y, \tau_0)|^2 \Phi(\tau_0, y) dy < \infty;$$

- (ii) the moving carrier subidentity of [definition 8.88](#) is valid with the time-dependent cutoff Φ ;
- (iii) the moving-annulus errors are summable in the sense of [definition 10.20](#), which implies

$$\int_{\tau_0}^{\infty} B_{\text{mov}}(\tau) d\tau < \infty$$

by [lemma 10.21](#);

- (iv) the validated compact cost is controlled by the moving identity cost on the terminal tail: for some $C_m < \infty$,

$$0 \leq \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq C_m \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \quad \text{for a.e. } \tau \geq \tau_0.$$

Theorem 8.92 (CostPhase 5: non-carrier adjacent exits are identified before carrier validation). *Let ω be an admissible local Type II branch with active supercritical scaling. Apply the pre-divergence rows of [definition 8.55](#). If a first failed test exists and is not one of the carrier-validation failures*

carrier subidentity, fixed finite error, moving-annulus summability, carrier comparison, then the branch has already failed to enter the pre-carrier validated compact-cost regime of [definition 8.87](#). More precisely, its first failure is one of the following named non-carrier strata: ordered exhaustion failure, representation failure, repaired-gauge or scale-degeneracy failure, multibubble or multicore splitting, cascade, exterior-tightness failure, local H^1 -loss, finite-cost scale-collapse, scale-rigid residual behavior, critical-mass failure, modulation-integrability failure, cost-well-posedness failure, or cost-compatibility failure.

In any ambient class in which the corresponding local closure theorems for these named strata are imposed, such a branch is not a remaining compact-cost branch.

Proof. The procedure in [definition 8.55](#) is ordered. The pre-carrier regime is the set of branches for which all tests before local monotonicity-carrier validation have passed. Therefore any first failure before the carrier-validation row prevents membership in [definition 8.87](#).

The same definition lists every possible first failure before carrier validation: the exhaustion failures are those of [definition A.3](#); representation and repaired-gauge failures are those of [definitions A.11](#) and [A.18](#) together with scale-degeneracy from [proposition 8.72](#); the compactness failures are the alternatives in [theorem 11.18](#); critical-mass failures are those routed

by lemma 10.78; compact tightness and local H^1 -loss are the tests of definition A.26; and the modulation, cost-well-posedness, and compatibility failures are exactly the failures recorded in lemmas 8.66, 8.67 and 8.69 and corollary 8.71. Thus the displayed list is exhaustive for non-carrier first failures.

If the surrounding Type II class includes the local closure theorem assigned to the named stratum, then the branch is handled there. It is therefore not a branch remaining to be treated by the retained compact-cost localized-energy argument. $\square$

Theorem 8.93 (CostPhase 5: pre-carrier reduction to carrier validation). *Let ω be in the pre-carrier validated compact-cost regime. Then exactly one of the following alternatives occurs:*

- (i) *the validated compact cost admits local monotonicity-carrier data in the sense of definition 8.59;*
- (ii) *the branch exits the retained compact carrier channel through the first failed carrier-validation test: carrier-subidentity failure, fixed-cutoff finite-error failure, moving-annulus summability failure, or carrier-comparison failure for the validated cost.*

The first alternative is a carrier certificate. It yields exclusion only after a phase 6 cost-divergence statement for the same cost is supplied.

Proof. By definition of the pre-carrier channel, every ordered test before carrier validation has passed. Thus the next unresolved row in definition 8.55 is precisely the local monotonicity-carrier row. Applying theorem 8.74 gives the two alternatives listed in the statement, and the first-failure convention makes them disjoint.

If the carrier exists, its data are retained for the later phase 6 implication. No tail-divergence assertion has been used in the present phase. $\square$

Theorem 8.94 (CostPhase 6: fixed-cutoff carrier closure). *Let ω be in the pre-carrier validated compact-cost regime. If ω satisfies the fixed-cutoff localized-energy closure criterion of definition 8.90 and*

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau = \infty,$$

then ω is excluded from the Type II class under consideration.

Proof. The pre-carrier channel supplies active supercritical scaling and local well-posedness and compatibility of the validated compact cost. Cost divergence is the additional phase 6 hypothesis of this theorem. The fixed-cutoff localized-energy criterion supplies finite initial localized energy, the fixed carrier subidentity, finite fixed-cutoff error, and energy-control comparison of the validated cost with $\tilde{\mathfrak{D}}_{R_0}$. Let C_1, τ_1 be the constants in that comparison and increase the tail initial time so that the comparison holds from τ_0 onward. Define

$$\mathcal{H}_\omega(\tau) = \mathcal{E}_\phi(\tau) + \int_{\tau}^{\infty} B_{R_0}(s) ds.$$

The finite-error condition makes $\mathcal{H}_\omega$ locally absolutely continuous and bounded below. The fixed carrier subidentity gives

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R_0}(\tau) \leq 0$$

in distributions. Since $\mathfrak{C}_{\text{II}}^{NS} \leq C_1 \tilde{\mathfrak{D}}_{R_0}$, this implies

$$\frac{d}{d\tau} \mathcal{H}_\omega(\tau) + \frac{1}{2C_1} \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq 0.$$

Together with the validated compact cost-divergence conclusion, these are the local cost-to-exclusion data with $\mathcal{D}_\omega = \mathfrak{C}_{\text{II}}^{NS}$. Equivalently, the fixed criterion is a successful carrier-validation

row. Since all pre-carrier tests and the stated compact cost-divergence conclusion are present, ω is on the retained compact cost-divergence stratum after carrier validation. Applying [theorem 8.64](#) gives the Type II exclusion conclusion. $\square$

Theorem 8.95 (CostPhase 6: moving-cutoff carrier closure). *Let ω be in the pre-carrier validated compact-cost regime. If ω satisfies the moving-cutoff localized-energy closure criterion of [definition 8.91](#) and*

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{II}}^{NS}(\tau) d\tau = \infty,$$

then ω is excluded from the Type II class under consideration.

Proof. The pre-carrier channel supplies active supercritical scaling and local well-posedness and compatibility of the validated compact cost. Cost divergence is the additional phase 6 hypothesis of this theorem. The moving-cutoff carrier criterion supplies the finite initial moving localized energy, the moving carrier subidentity, summable moving-annulus errors, and the terminal comparison

$$0 \leq \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq C_m \tilde{\mathfrak{D}}_{R(\tau)}(\tau).$$

By [lemma 10.21](#), the summability condition gives $B_{\text{mov}} \in L^1(\tau_0, \infty)$. Define

$$\mathcal{H}_{\omega}(\tau) = \mathcal{E}_R(\tau) + \int_{\tau}^{\infty} B_{\text{mov}}(s) ds.$$

The finite moving-error condition makes $\mathcal{H}_{\omega}$ locally absolutely continuous and bounded below. The moving carrier subidentity gives

$$\frac{d}{d\tau} \mathcal{H}_{\omega}(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \leq 0$$

in distributions. The moving comparison gives

$$\frac{d}{d\tau} \mathcal{H}_{\omega}(\tau) + \frac{1}{2C_m} \mathfrak{C}_{\text{II}}^{NS}(\tau) \leq 0.$$

Together with the validated compact cost-divergence conclusion, these are the local cost-to-exclusion data with $\mathcal{D}_{\omega} = \mathfrak{C}_{\text{II}}^{NS}$. Equivalently, the moving criterion is a successful carrier-validation row. Since all pre-carrier tests and the stated compact cost-divergence conclusion are present, ω is on the retained compact cost-divergence stratum after carrier validation. Applying [theorem 8.64](#) excludes the branch. $\square$

Lemma 8.96 (Pre-carrier local energy gives finite initial localized energy). *Let ω be in the pre-carrier validated compact-cost regime, represented by V on a terminal tail. Fix any later time $T \geq \tau_0$. After possibly replacing the tail by one whose initial time $\tau_1 \in [T, T+1]$ is a compact-cylinder local-energy Lebesgue time, every finite $R < \infty$ satisfies*

$$\int_{B_R} |V(y, \tau_1)|^2 dy < \infty.$$

Consequently the initial localized-energy clauses in the fixed and moving localized-energy closure criteria are available after a harmless terminal restriction for every compactly supported fixed cutoff and every moving cutoff with finite initial radius. This is a local compact-cylinder statement; it does not require $V(\tau) \in L^3(\mathbb{R}^3)$ or a whole-space critical bound.

Proof. The physical branch is suitable and the repaired-gauge chart is locally absolutely continuous on compact renormalized cylinders in the pre-carrier regime. Pulling back the local energy class gives

$$V \in L_{\text{loc}}^{\infty}([\tau_0, \infty); L^2(B_R)) \quad \text{for every } R < \infty$$

after choosing the usual representative, or at minimum

$$V \in L^2_{\text{loc}}([\tau_0, \infty); L^2(B_R))$$

with local-energy Lebesgue times on every compact interval. If the stronger representative is unavailable at a prescribed endpoint, choose $\tau_1 \in [T, T+1]$ which is a Lebesgue time simultaneously for the finitely many compact supports used by the fixed or moving cutoff under consideration. This is possible because the local L^2 -in-time norms are finite on $[T, T+1]$.

At such a τ_1 , every compactly supported fixed cutoff has finite localized energy. If $\Phi(\tau, y) = \phi(y/R(\tau))$ and $R(\tau_1) < \infty$, then $\Phi(\tau_1, \cdot)$ is compactly supported, so the same local L^2 finiteness gives $\mathcal{E}_R(\tau_1) < \infty$. Passing from τ_0 to τ_1 removes only a finite initial renormalized interval and therefore does not change any terminal divergence or tail-integrability alternative. $\square$

Lemma 8.97 (Nested moving cutoffs give carrier comparison). *Let ω be in the pre-carrier validated compact-cost regime, and assume the validated compact cost is energy-controlled by the fixed identity cost:*

$$0 \leq \mathfrak{E}_{\text{II}}^{NS}(\tau) \leq C_1 \tilde{\mathfrak{D}}_{R_0}(\tau) \quad \text{for a.e. } \tau \geq \tau_1.$$

Let $\Phi(\tau, y)$ be a moving cutoff such that, after increasing the initial time if necessary,

$$\Phi(\tau, y) \geq \phi_{R_0}(y) \quad \text{for all } y \in \mathbb{R}^3 \text{ and a.e. } \tau \geq \tau_1.$$

Then the validated compact cost is controlled by the moving identity cost:

$$0 \leq \mathfrak{E}_{\text{II}}^{NS}(\tau) \leq C_1 \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \quad \text{for a.e. } \tau \geq \tau_1.$$

In particular, for the canonical identity cost the carrier comparison is automatic with $C_1 = 1$.

Proof. Because $\Phi(\tau, \cdot) \geq \phi_{R_0}$ and both cutoffs are nonnegative,

$$\nu \int |\nabla V|^2 \phi_{R_0} \leq \nu \int |\nabla V|^2 \Phi(\tau, \cdot),$$

and, since $a_+(\tau) \geq 0$,

$$a_+(\tau) \int |V|^2 \phi_{R_0} \leq a_+(\tau) \int |V|^2 \Phi(\tau, \cdot).$$

Adding these two inequalities gives

$$\tilde{\mathfrak{D}}_{R_0}(\tau) \leq \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \quad \text{for a.e. } \tau \geq \tau_1.$$

Combining this with the fixed energy-control comparison yields the displayed moving comparison. The canonical identity cost is energy-controlled by [definition 8.58](#) with $C_1 = 1$. $\square$

Definition 8.98 (Finite-error carrier selection). The pre-carrier validated compact-cost regime has finite-error carrier selection if every branch in that regime whose validated compact cost is energy-controlled by the fixed identity cost satisfies, after passing to a later terminal tail if necessary, at least one of the following two local alternatives:

- (i) the fixed carrier subidentity of [definition 8.88](#) is valid and

$$\int_{\tau_0}^{\infty} B_{R_0}(\tau) d\tau < \infty;$$

- (ii) there exists a finite-valued nondecreasing locally absolutely continuous radius $R(\tau) \geq R_0$, with $R(\tau) \rightarrow \infty$ and moving cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$, such that

$$\Phi(\tau, y) \geq \phi_{R_0}(y) \quad \text{for all } y \text{ and a.e. } \tau,$$

the moving carrier subidentity of [definition 8.88](#) is valid with the time-dependent cutoff Φ , and the moving-annulus errors are summable in the sense of [definition 10.20](#).

This condition contains only the local one-sided energy estimate and finite-error selection statements. It does not include finite initial localized energy or carrier comparison, because those are supplied by [lemmas 8.96](#) and [8.97](#).

Definition 8.99 (Annular finite-error selection). The pre-carrier validated compact-cost regime has annular finite-error selection if every branch in that regime whose validated compact cost is energy-controlled by the fixed identity cost satisfies, after passing to a later terminal tail if necessary, at least one of the following two alternatives:

- (i) the fixed error density has finite tail integral,

$$\int_{\tau_0}^{\infty} B_{R_0}(\tau) d\tau < \infty;$$

- (ii) there exists a finite-valued nondecreasing locally absolutely continuous nested radius $R(\tau) \geq R_0$, with $R(\tau) \rightarrow \infty$ and $\Phi(\tau, y) = \phi(y/R(\tau)) \geq \phi_{R_0}(y)$, such that the moving-annulus errors are summable in the sense of [definition 10.20](#).

This is the purely finite-error part of the carrier problem. It does not include local energy subidentities, finite initial energy, or carrier comparison.

Definition 8.100 (Componentwise fixed-or-moving finite-error estimates). For a branch in the pre-carrier validated compact-cost regime, the componentwise annular finite-error estimates hold if, after passing to a later terminal tail if necessary, either the six fixed-cutoff components

$$\begin{aligned} I_1 &= \left| \int |V|^2 \Delta \phi_{R_0} \right|, & I_2 &= \left| \int |V|^2 V \cdot \nabla \phi_{R_0} \right|, & I_3 &= \left| \int P V \cdot \nabla \phi_{R_0} \right|, \\ I_4 &= |b| \left| \int |V|^2 \nabla \phi_{R_0} \right|, & I_5 &= a_- M_\phi, & I_6 &= a_+ \int |V|^2 (-y \cdot \nabla \phi_{R_0}) \end{aligned}$$

belong to $L^1(\tau_0, \infty)$, or there is a finite-valued nondecreasing locally absolutely continuous nested radius $R(\tau) \geq R_0$, with $R(\tau) \rightarrow \infty$ and $\Phi(\tau, y) = \phi(y/R(\tau)) \geq \phi_{R_0}(y)$, such that each of the seven moving annular components

$$\begin{aligned} J_1 &= \left| \int |V|^2 \partial_\tau \Phi \right|, & J_2 &= \left| \int |V|^2 \Delta \Phi \right|, & J_3 &= \left| \int |V|^2 V \cdot \nabla \Phi \right|, \\ J_4 &= \left| \int P V \cdot \nabla \Phi \right|, & J_5 &= |b| \left| \int |V|^2 \nabla \Phi \right|, \\ J_6 &= a_- \int |V|^2 \Phi, & J_7 &= a_+ \int |V|^2 (-y \cdot \nabla \Phi) \end{aligned}$$

belongs to $L^1(\tau_0, \infty)$.

Remark 8.101 (What the componentwise estimates isolate). The quantities in [definition 8.100](#) are exactly the fixed and moving carrier-error components: in the fixed-cutoff case, viscous cutoff, convection, pressure-tail, center drift, negative scale drift, and the positive scale-shell drift; in the moving-cutoff case, cutoff transition, viscous cutoff, convection, pressure-tail, center drift, negative scale drift, and the positive scale-shell drift. These diagnostics need not be closed by a single global tail estimate. In the stratified use of Section 7, persistent failure of the cutoff, viscous, convection, or pressure diagnostics is read as shell mass that cannot be evacuated from selected windows and is therefore assigned to exterior leakage, exterior concentration, loss of critical tightness, or local windowed H^1 -loss. Persistent failure of the translation diagnostic is read as a modulation-driven adjacent branch, while persistent failure of the positive scale-shell term is read as a scale-dynamics branch. Thus the component list records the local diagnostics

seen by the carrier test; it is not intended to force the proof into a single classical PDE closure argument. For radial nonincreasing cutoffs, the scale-shell term

$$-\frac{a}{2} \int |V|^2 y \cdot \nabla \Phi$$

has only an a_+ -contribution in the residual carrier error. The a_- part is favorable and is already absorbed into the negative-drift alternative. The negative scale-drift component J_6 is not to be replaced by an unproved equivalence with compact-cost divergence: if J_6 is not summable, it must be handled through the explicit scale-drift alternatives of [corollary 10.30](#) and [theorem 10.38](#) or treated by a direct summability estimate.

Lemma 8.102 (Componentwise estimates give annular finite-error selection). *If every branch in the pre-carrier validated compact-cost regime whose validated compact cost is energy-controlled by the fixed identity cost satisfies the componentwise annular finite-error estimates of [definition 8.100](#), then the regime has annular finite-error selection in the sense of [definition 8.99](#).*

Proof. Let ω be such a branch. If the fixed alternative in [definition 8.100](#) holds, with $I_1, \dots, I_6$ denoting the fixed components appearing in that definition, then [definition 10.9](#) gives

$$B_{R_0} = \frac{\nu}{2} I_1 + \frac{1}{2} I_2 + I_3 + \frac{1}{2} I_4 + \frac{1}{2} I_5 + \frac{1}{2} I_6.$$

Since B_{R_0} is a finite positive linear combination of integrable fixed components, $B_{R_0} \in L^1(\tau_0, \infty)$. This is the first alternative in [definition 8.99](#). Otherwise the definition supplies a finite-valued nested moving radius $R(\tau) \geq R_0$, $R(\tau) \rightarrow \infty$, for which the seven moving annular components $J_1, \dots, J_7$ are integrable. These seven components are exactly the summability requirements in [definition 10.20](#), with the same moving cutoff Φ . Hence the moving-annulus errors are summable, which is the second alternative in [definition 8.99](#). $\square$

Theorem 8.103 (Finite-error selection reduced to annular finite-error). *Assume that repaired-gauge suitable local energy gives the fixed and moving carrier subidentities in the sense of [lemma 8.89](#). If the pre-carrier regime has annular finite-error selection in the sense of [definition 8.99](#), then it has finite-error carrier selection in the sense of [definition 8.98](#).*

Proof. Let ω be a branch in the pre-carrier regime whose validated compact cost is energy-controlled by the fixed identity cost. By annular finite-error selection, after restricting to a later tail either the fixed error density is integrable or there is a nested moving radius for which the moving-annulus errors are summable. The repaired-gauge suitable local energy lemma supplies the fixed carrier subidentity in the first case and the moving carrier subidentity in the second. These are exactly the two alternatives in [definition 8.98](#). $\square$

Theorem 8.104 (CostPhase 5: carrier completion reduced to finite-error selection). *Assume the pre-carrier validated compact-cost regime has finite-error carrier selection in the sense of [definition 8.98](#). Then every branch in the pre-carrier regime whose validated compact cost is energy-controlled by the fixed identity cost satisfies the fixed-cutoff or moving-cutoff carrier criterion. No cost-divergence conclusion is used here.*

Proof. Let ω be such a branch. The finite initial localized energy required by either localized-energy criterion is supplied locally, after a harmless terminal restriction, by [lemma 8.96](#). The assumed energy-control comparison gives the fixed carrier-comparison clause.

If the fixed alternative in [definition 8.98](#) holds, then the fixed carrier subidentity and the finite B_{R_0} -tail are precisely the two remaining clauses of the fixed-cutoff localized-energy closure criterion. Hence the fixed-cutoff criterion holds.

If the moving alternative holds, the summable moving-annulus condition gives

$$\int_{\tau_0}^{\infty} B_{\text{mov}}(\tau) d\tau < \infty$$

by [lemma 10.21](#). The moving carrier subidentity is part of the moving alternative. The nested cutoff condition and fixed energy-control comparison give the moving carrier comparison by [lemma 8.97](#). Together with finite initial moving energy from [lemma 8.96](#), these are exactly the moving localized-energy closure criterion. Therefore the moving-cutoff criterion holds.

The two alternatives in finite-error carrier selection exhaust the branch, so the claimed fixed-or-moving carrier completion follows. $\square$

Theorem 8.105 (CostPhase 6: localized-energy criterion closes the pre-carrier regime). *Let ω be in the pre-carrier validated compact-cost regime. If ω satisfies either the fixed-cutoff localized-energy closure criterion of [definition 8.90](#) or the moving-cutoff localized-energy closure criterion of [definition 8.91](#), and if*

$$\int_{\tau_0}^{\infty} \mathfrak{e}_{\text{II}}^{NS}(\tau) d\tau = \infty,$$

then ω is excluded from the Type II class under consideration.

Proof. In the fixed-cutoff case apply [theorem 8.94](#). In the moving-cutoff case apply [theorem 8.95](#). These are the only cases in the stated disjunction. $\square$

Definition 8.106 (Carrier-estimate completion for the compact-cost regime). The compact-cost regime has carrier-estimate completion if every branch in the pre-carrier validated compact-cost regime satisfies, after passing to a later terminal tail if necessary, either the fixed-cutoff localized-energy closure criterion of [definition 8.90](#) or the moving-cutoff localized-energy closure criterion of [definition 8.91](#).

Corollary 8.107 (Finite-error selection gives carrier-estimate completion). *Assume every branch in the pre-carrier validated compact-cost regime uses a validated compact cost that is energy-controlled by the fixed identity cost; in particular this holds for the canonical identity cost. If finite-error carrier selection holds in the sense of [definition 8.98](#), then the compact-cost regime has carrier-estimate completion in the sense of [definition 8.106](#).*

Proof. Let ω be any branch in the pre-carrier validated compact-cost regime. By the energy-control hypothesis, ω is one of the branches covered by [theorem 8.104](#). That theorem gives either the fixed-cutoff localized-energy closure criterion or the moving-cutoff localized-energy closure criterion on a terminal tail. This is exactly [definition 8.106](#). The parenthetical canonical case follows from [definition 8.58](#). $\square$

Corollary 8.108 (Annular finite-error selection gives carrier-estimate completion). *Assume every branch in the pre-carrier validated compact-cost regime uses a validated compact cost that is energy-controlled by the fixed identity cost. Assume also that repaired-gauge suitable local energy gives the fixed and moving carrier subidentities in the sense of [lemma 8.89](#). If annular finite-error selection holds in the sense of [definition 8.99](#), then the compact-cost regime has carrier-estimate completion.*

Proof. [Theorem 8.103](#) converts annular finite-error selection into finite-error carrier selection. Then [corollary 8.107](#) gives carrier-estimate completion. $\square$

Theorem 8.109 (CostPhase 5: adjacent closure reduced to carrier-estimate completion). *Assume the non-carrier adjacent strata listed in [theorem 8.92](#) are closed by their corresponding local compactness or interface arguments in the ambient Type II class. If the compact-cost regime has*

carrier-estimate completion in the sense of [definition 8.106](#), then the compact-cost regime has adjacent-stratum closure in the sense of [definition 8.83](#).

Proof. Let ω enter a named adjacent stratum in the pre-divergence order [definition 8.55](#). If it is non-carrier, then [theorem 8.92](#) identifies it with one of the previously named exhaustion, representation, compactness, critical-mass, compact-test, modulation, cost-well-posedness, or cost-compatibility strata. By the first hypothesis of the present theorem, that stratum is closed by its own local argument, so the Type II exclusion conclusion holds there, or the branch is assigned to a previously closed non-compact-cost alternative.

It remains to consider a purported carrier-validation failure after the branch has reached the pre-carrier channel. Carrier-estimate completion supplies either the fixed- or moving-cutoff criterion. Each criterion contains the corresponding carrier subidentity, finite-error condition, finite initial localized energy, and carrier comparison, so the carrier-validation row passes. This contradicts the assertion that its first failed row is a carrier-validation failure. Hence no carrier failure remains unresolved. Together with the preceding non-carrier closures, this is precisely [definition 8.83](#). The proof has not tested the tail integral of the cost. $\square$

Definition 8.110 (Stratified moving-carrier alternative). Assume the validated compact cost is energy-controlled by the fixed identity cost. The pre-carrier validated compact-cost regime has stratified moving-cutoff alternative if every branch ω in that regime, after passing to a later terminal tail if necessary, satisfies at least one of the following local alternatives:

- (i) ω satisfies the fixed-cutoff localized-energy closure criterion of [definition 8.90](#);
- (ii) ω satisfies the moving-cutoff localized-energy closure criterion of [definition 8.91](#);
- (iii) ω satisfies the exterior/leakage alternative: either there are unit windows $I_n = [T_n, T_n + 1]$, selected times $\tau_n \in I_n$, radii $R_n \rightarrow \infty$, and $\eta > 0$ such that

$$\sup_n (|a(\tau_n)| + \|b(\tau_n)\|_\infty) < \infty,$$

$$\frac{1}{R_n^2} \int_{A_{R_n}} |V(y, \tau_n)|^2 dy + \frac{1}{R_n} \int_{A_{R_n}} |V(y, \tau_n)|^3 dy + \frac{1}{R_n} \int_{\tilde{A}_{R_n}} |V(y, \tau_n)|^3 dy \geq \eta$$

for every n , and this selected-window shell persistence is assigned to one of the already named critical-tightness failure, exterior concentration, or local windowed H^1 -loss branches; or the normalized transition-carrier blocks of the moving cutoff have a local occupation-state limit with the transition-leakage carrier label of [lemma 8.124](#), which is identified there with the same exterior/leakage cluster.

- (iv) ω satisfies the translation alternative: there are unit windows $I_n = [T_n, T_n + 1]$, selected times $\tau_n \in I_n$, radii $R_n \rightarrow \infty$, cutoffs $\Phi_n(y) = \phi(y/R_n)$, and $\eta > 0$ such that

$$\sup_n (|a(\tau_n)| + \|b(\tau_n)\|_\infty) < \infty,$$

$$\frac{|b(\tau_n)|}{R_n} \int_{A_{R_n}} |V(y, \tau_n)|^2 dy \geq \eta$$

for every n , and this persistent translation transport is assigned to the already named modulation-integrability failure or to a multicore/exterior branch produced after selected-window recentering;

- (v) ω satisfies the positive-scale alternative: either there are unit windows $I_n = [T_n, T_n + 1]$, selected times $\tau_n \in I_n$, radii $R_n \rightarrow \infty$, cutoffs $\Phi_n(y) = \phi(y/R_n)$, and $\eta > 0$ such that

$$\sup_n (|a(\tau_n)| + \|b(\tau_n)\|_\infty) < \infty,$$

$$a_+(\tau_n) \int |V(y, \tau_n)|^2 (-y \cdot \nabla \Phi_n(y)) dy \geq \eta$$

for every n , and this persistent positive scale forcing is assigned to the already named scale-rigid residual alternative or to a bounded selected-window limit refinement of that alternative; or the normalized positive-scale-work carrier blocks have a retained compact occupation-state limit whose compact-core realization satisfies the scale-rigid branch clauses of [definition 5.19](#). If the occupation state instead has an earlier exterior/leakage, modulation-driven recentering, or negative-drift label, the branch is read in that earlier alternative rather than in the positive-scale alternative.

(vi) ω enters the negative scale-drift obstruction of [definition 10.25](#).

The fixed-cutoff case is auxiliary here. The canonical local alternative is the moving one: either a nested moving cutoff closes the carrier, or a persistent failure produces one of the displayed selected-window diagnostics and is assigned to the corresponding adjacent stratum.

Remark 8.111 (How the moving diagnostics match the carrier components). The exterior/leakage alternative in [definition 8.110](#) groups the moving cutoff, viscous, convection, and pressure diagnostics. The point is that the canonical windowwise shell package [lemma 8.116](#), together with the local harmonic-pressure routing [lemmas 9.21](#) and [9.22](#), routes pressure failure to enlarged-shell L^3 persistence, pressure-only critical mass, or exterior concentration. Thus the pressure obstruction is not an independent global theorem; it is another way of detecting local shell mass that refuses to leave the selected windows. The translation alternative contains the J_5 term and is meant to be combined with the bounded selected-window modulation provided by [corollaries 8.37](#) and [8.43](#). The positive-scale alternative contains the J_7 term and isolates the only genuinely positive scale-dynamics obstruction left after the favorable a_- -part has been separated. The negative-drift term J_6 is already assigned to Section 6 and is not to be replaced by an unproved compact-cost equivalence. The occupation-state clauses are included so that a diffuse nonintegrable error is not incorrectly converted into a pointwise lower bound; it is instead read as an occupation state in the local ordered stratification.

Definition 8.112 (Moving test family). For a branch in the pre-carrier validated compact-cost regime, a moving test family is a finite-valued nondecreasing locally absolutely continuous radius $R(\tau) \geq R_0$, with $R(\tau) \rightarrow \infty$, together with the nested moving cutoff

$$\Phi(\tau, y) = \phi(y/R(\tau)) \geq \phi_{R_0}(y).$$

For such a family, let $J_1, \dots, J_7$ denote the seven moving carrier components of [definition 8.100](#).

Lemma 8.113 (CostPhase 5: a successful moving test family gives a moving carrier certificate). *Assume repaired-gauge suitable local energy gives the moving carrier subidentity in the sense of [lemma 8.89](#). Let ω be a branch in the pre-carrier validated compact-cost regime whose validated compact cost is energy-controlled by the fixed identity cost. If there exists a moving test family in the sense of [definition 8.112](#) such that, for some $\tau_1 \geq \tau_0$,*

$$J_k \in L^1(\tau_1, \infty) \quad \text{for each } 1 \leq k \leq 7,$$

then ω satisfies the moving-cutoff localized-energy closure criterion of [definition 8.91](#) after restriction to the later tail $[\tau_1, \infty)$. This conclusion is a carrier certificate and does not assert cost divergence or exclusion.

Proof. Restrict the branch and the moving test family to the tail $[\tau_1, \infty)$. Let $R(\tau)$ and $\Phi(\tau, y)$ be the restricted moving test family, with components $J_1, \dots, J_7$. By [definition 10.20](#), the integrability of these seven components is exactly the summable moving-annulus condition for Φ . Hence $\int_{\tau_1}^{\infty} B_{\text{mov}}(\tau) d\tau < \infty$ by [lemma 10.21](#). The suitable local energy lemma gives the moving carrier subidentity for the same cutoff. Finite initial localized energy on the pre-carrier regime is automatic by [lemma 8.96](#), and the moving carrier comparison is automatic by [lemma 8.97](#). These are exactly the clauses of [definition 8.91](#). $\square$

Lemma 8.114 (Retained pressure tails vanish on windowwise annuli). *Let ω be a branch in the pre-carrier validated compact-cost regime, represented by (V, P, a, b) on a terminal tail, and let $I = [A, B]$ be a compact renormalized time interval. Assume that no pressure-reconstruction, exterior-leakage, separated-core, or critical-tightness failure label occurs on the retained branch over I . Then, for every radial cutoff $\phi \in C_c^\infty(B_2)$ with $\phi = 1$ on B_1 ,*

$$\lim_{S \rightarrow \infty} \sup_{\substack{R \in AC(I), \\ R(\tau) \geq S \\ R(\tau) < \infty}} \int_I \left| \int P(y, \tau) V(y, \tau) \cdot \nabla \phi(y/R(\tau)) dy \right| d\tau = 0.$$

Consequently, on the retained branch one may choose any prescribed summable sequence $\varepsilon_n > 0$ and radii $S_n \rightarrow \infty$ so that every finite-valued moving radius $R(\tau) \geq S_n$ on the n -th unit window has pressure flux at most ε_n on that window.

Proof. Argue by contradiction on a compact interval I . If the displayed limit fails, then there are $\eta > 0$, radii $S_m \rightarrow \infty$, and finite-valued moving radii $R_m \in AC(I)$, $R_m \geq S_m$, such that the pressure flux through the annulus of $\phi(\cdot/R_m(\tau))$ has L^1_τ -mass at least η .

Let

$$A_m(\tau) = \{R_m(\tau) < |y| < 2R_m(\tau)\}, \quad \tilde{A}_m(\tau) = \{R_m(\tau)/2 < |y| < 4R_m(\tau)\}.$$

The cutoff gradient is supported in $A_m(\tau)$ and satisfies $|\nabla \phi(y/R_m(\tau))| \leq C_\phi/R_m(\tau)$. On a finite partition of $\tilde{A}_m(\tau)$, rescaled to a fixed reference annulus, the localized Calderón–Zygmund pressure P_m^{loc} generated by $\mathbf{1}_{\tilde{A}_m(\tau)} V \otimes V$ obeys the scale-invariant bound

$$\|P_m^{\text{loc}}(\tau)\|_{L^{3/2}(A_m(\tau))} \leq C_\phi \|V(\tau)\|_{L^3(\tilde{A}_m(\tau))}^2.$$

Therefore Hölder’s inequality gives

$$\begin{aligned} \int_I \left| \int P_m^{\text{loc}} V \cdot \nabla \phi(y/R_m(\tau)) dy \right| d\tau &\leq C_\phi \int_I R_m(\tau)^{-1} \|V(\tau)\|_{L^3(\tilde{A}_m(\tau))}^3 d\tau \\ &\leq C_{\phi, I} S_m^{-1} \sup_{\tau \in I} \int_{|y| > S_m/2} |V(y, \tau)|^3 dy. \end{aligned}$$

This tends to zero because the retained branch has passed the critical tightness row on I . Thus the local Calderón–Zygmund pressure part cannot carry the lower bound η .

The remaining pressure contribution is harmonic on each retained rescaled subannulus after subtracting functions of time. If its pressure work did not vanish, then on at least one element of the finite annular partition the normalized harmonic pressure would have a positive $L^{3/2}$ -mass after rescaling. The quantitative converse in [lemma 9.21](#) would then force a positive $L^{3/2}$ -mass of the recorded annular source $V \otimes V$, equivalently a positive L^3 -mass of V , on a retained exterior annular cell. This is exactly an exterior-leakage or separated-core label under [theorem 9.39](#). If the harmonic piece is not generated by a recorded annular source, the ordered pressure atlas reads it as a pressure-reconstruction failure or as a far-field pressure tail; the latter is routed by [lemma 9.22](#) and [lemma 9.45](#). Every resulting label is excluded by the retained-branch hypotheses on I . Hence the harmonic pressure cannot carry the lower bound η either.

Therefore the displayed uniform tail limit holds. Applying it on each unit window with ε_n gives the final scheduling statement. $\square$

Lemma 8.115 (Critical tightness gives a tailored windowwise radius schedule). *Let ω be a branch in the pre-carrier validated compact-cost regime, represented by V on a terminal tail. For the unit windows*

$$I_n = [\tau_0 + n, \tau_0 + n + 1], \quad n \in \mathbb{N},$$

set

$$A_n := 1 + \int_{I_n} (|a(\tau)| + \|b(\tau)\|_\infty) d\tau.$$

Then there exists a nondecreasing sequence $R_n \rightarrow \infty$, with

$$R_n \geq 4 \cdot 2^n,$$

such that

$$\sup_{\tau \in I_n} \int_{|y| > R_n/4} |V(y, \tau)|^3 dy \leq 2^{-3n} A_n^{-3/2} \quad \text{for every } n.$$

The sequence may also be chosen so that every finite-valued moving radius $\mathcal{R}(\tau) \geq R_n$ on I_n satisfies

$$\int_{I_n} \left| \int P(y, \tau) V(y, \tau) \cdot \nabla \phi(y/\mathcal{R}(\tau)) dy \right| d\tau \leq 2^{-n}.$$

Proof. Since ω lies in the pre-carrier validated compact-cost regime, the critical-tightness test and the modulation-integrability test in [definition 8.55](#) have already passed. Thus [definition 8.1](#) holds and [lemma 8.69](#) gives

$$a, b \in L^1_{\text{loc}}([\tau_0, \infty)).$$

In particular each A_n is finite.

For every n , set

$$\varepsilon_n := 2^{-3n} A_n^{-3/2} > 0.$$

By uniform critical tightness, for each ε_n there exists $S_n < \infty$ such that

$$\sup_{\tau \geq \tau_0} \int_{|y| > S_n} |V(y, \tau)|^3 dy \leq \varepsilon_n.$$

Since the branch is retained in the pre-carrier regime, pressure reconstruction, exterior leakage, and critical tightness have not failed on I_n . Applying [lemma 8.114](#) with tolerance 2^{-n} , choose $P_n < \infty$ such that every finite-valued moving radius $R(\tau) \geq P_n$ on I_n has pressure flux at most 2^{-n} . Choose R_n recursively so that

$$R_n \geq \max\{4S_n, P_n, 4 \cdot 2^n, R_{n-1}\}$$

for $n \geq 1$, and $R_0 \geq \max\{4S_0, P_0, 4\}$. Then R_n is nondecreasing, $R_n \rightarrow \infty$, $R_n \geq 4 \cdot 2^n$, and

$$\sup_{\tau \in I_n} \int_{|y| > R_n/4} |V(y, \tau)|^3 dy \leq \sup_{\tau \geq \tau_0} \int_{|y| > S_n} |V(y, \tau)|^3 dy \leq \varepsilon_n,$$

which is the desired estimate. $\square$

Lemma 8.116 (Canonical windowwise moving family controls shell, pressure, and translation terms). *Let ω be a branch in the pre-carrier validated compact-cost regime. Let R_n be the sequence supplied by [lemma 8.115](#). On the unit windows $I_n = [\tau_0 + n, \tau_0 + n + 1]$, define the canonical interpolated radius $R(\tau)$ by*

$$R(\tau) := R_n + (\tau - (\tau_0 + n))(R_{n+1} - R_n), \quad \tau \in I_n,$$

and the associated moving cutoff

$$\Phi(\tau, y) := \phi(y/R(\tau)).$$

Write

$$\begin{aligned} J_2(\tau) &= \left| \int |V|^2 \Delta \Phi(\tau, \cdot) \right|, & J_3(\tau) &= \left| \int |V|^2 V \cdot \nabla \Phi(\tau, \cdot) \right|, \\ J_4(\tau) &= \left| \int P V \cdot \nabla \Phi(\tau, \cdot) \right|, & J_5(\tau) &= |b(\tau)| \left| \int |V|^2 \nabla \Phi(\tau, \cdot) \right|. \end{aligned}$$

Then

$$\int_{\tau_0}^{\infty} (J_2(\tau) + J_3(\tau) + J_4(\tau) + J_5(\tau)) d\tau < \infty.$$

Proof. The sequence R_n is nondecreasing and tends to $+\infty$. Hence the piecewise affine interpolation $R(\tau)$ is finite-valued, nondecreasing, locally absolutely continuous, and satisfies $R(\tau) \rightarrow \infty$. Thus it is an admissible moving test family in the sense of [definition 8.112](#).

By [lemma 8.115](#),

$$\sup_{\tau \in I_n} \int_{|y| > R_n/4} |V(y, \tau)|^3 dy \leq \varepsilon_n := 2^{-3n} A_n^{-3/2},$$

where

$$A_n = 1 + \int_{I_n} (|a(\tau)| + \|b(\tau)\|_{\infty}) d\tau.$$

For $\tau \in I_n$, one has $R(\tau) \geq R_n$, hence

$$A_{R(\tau)} \cup \tilde{A}_{R(\tau)} \subset \{|y| > R(\tau)/4\} \subset \{|y| > R_n/4\}.$$

Therefore, for every $\tau \in I_n$,

$$\int_{A_{R(\tau)}} |V(y, \tau)|^3 dy + \int_{\tilde{A}_{R(\tau)}} |V(y, \tau)|^3 dy \leq \varepsilon_n.$$

Using Hölder on the annulus $A_{R(\tau)}$,

$$\int_{A_{R(\tau)}} |V(y, \tau)|^2 dy \leq |A_{R(\tau)}|^{1/3} \left(\int_{A_{R(\tau)}} |V(y, \tau)|^3 dy \right)^{2/3} \leq C_{\phi} R(\tau) \varepsilon_n^{2/3}.$$

Since $|\Delta\Phi(\tau, \cdot)| \leq C_{\phi} R(\tau)^{-2} \mathbf{1}_{A_{R(\tau)}}$,

$$J_2(\tau) \leq C_{\phi} R(\tau)^{-2} \int_{A_{R(\tau)}} |V|^2 \leq C_{\phi} R(\tau)^{-1} \varepsilon_n^{2/3} \leq C_{\phi} R_n^{-1} \varepsilon_n^{2/3}.$$

Integrating on I_n and using $R_n \geq 4 \cdot 2^n$,

$$\int_{I_n} J_2(\tau) d\tau \leq C_{\phi} R_n^{-1} \varepsilon_n^{2/3} \leq C_{\phi} 2^{-3n} 2^{-2n} A_n^{-1}.$$

Since $|\nabla\Phi(\tau, \cdot)| \leq C_{\phi} R(\tau)^{-1} \mathbf{1}_{A_{R(\tau)}}$,

$$J_3(\tau) \leq C_{\phi} R(\tau)^{-1} \int_{A_{R(\tau)}} |V|^3 \leq C_{\phi} R_n^{-1} \varepsilon_n,$$

and therefore

$$\int_{I_n} J_3(\tau) d\tau \leq C_{\phi} R_n^{-1} \varepsilon_n = C_{\phi} R_n^{-1} 2^{-3n} A_n^{-3/2}.$$

The pressure component is controlled by the retained pressure-tail scheduling included in [lemma 8.115](#). Since $R(\tau) \geq R_n$ on I_n , the choice of R_n gives

$$\int_{I_n} J_4(\tau) d\tau = \int_{I_n} \left| \int P(y, \tau) V(y, \tau) \cdot \nabla\Phi(\tau, y) dy \right| d\tau \leq 2^{-n}.$$

If the retained pressure-tail scheduling were unavailable, the preceding lemma would assign the branch to a pressure/exterior local alternative before this windowwise retained branch is reached.

Finally,

$$J_5(\tau) \leq C_{\phi} \|b(\tau)\|_{\infty} R(\tau)^{-1} \int_{A_{R(\tau)}} |V|^2 \leq C_{\phi} \|b(\tau)\|_{\infty} \varepsilon_n^{2/3}.$$

Hence

$$\int_{I_n} J_5(\tau) d\tau \leq C_\phi \varepsilon_n^{2/3} \int_{I_n} \|b(\tau)\|_\infty d\tau \leq C_\phi 2^{-2n} A_n^{-1} (A_n - 1) \leq C_\phi 2^{-2n}.$$

The four resulting series are summable in n , proving the claim. $\square$

Corollary 8.117 (Only transition and scale terms remain on the windowwise schedule). *Under the hypotheses of lemma 8.116, the canonical windowwise moving family*

$$R(\tau) = R_n + (\tau - (\tau_0 + n))(R_{n+1} - R_n), \quad \tau \in I_n,$$

has summable moving carrier components J_2, J_3, J_4, J_5 . Thus the only unresolved moving-carrier terms on that genuine moving family are the transition term J_1 , the positive scale-shell term J_7 , and the negative scale-drift term J_6 .

Proof. This is exactly the conclusion of lemma 8.116, together with the fact that the moving carrier has exactly seven components, of which lemma 8.116 already controls J_2, J_3, J_4, J_5 . $\square$

Definition 8.118 (Windowwise transition/leakage alternative statement). The pre-carrier validated compact-cost regime has windowwise transition/leakage alternative if, for every branch ω in that regime, the canonical windowwise moving family of lemma 8.116 has the following property: if

$$J_1 \notin L^1(\tau_1, \infty)$$

on every later tail $[\tau_1, \infty)$, then ω enters the exterior/leakage alternative of definition 8.110.

Definition 8.119 (Windowwise positive-scale alternative statement). The pre-carrier validated compact-cost regime has windowwise positive-scale alternative if, for every branch ω in that regime, the canonical windowwise moving family of lemma 8.116 has the following property: if

$$J_7 \notin L^1(\tau_1, \infty)$$

on every later tail $[\tau_1, \infty)$, then ω enters the positive-scale alternative of definition 8.110, unless it has already entered the exterior/leakage, modulation-driven recentering, or negative scale-drift alternatives of the same stratified moving-localized-energy criterion.

8.5.1. Proof of the windowwise alternative statements. We now prove that the two alternative properties of definitions 8.118 and 8.119 hold for the canonical windowwise moving family of lemma 8.116. The proof does not convert arbitrary nonintegrability into a false persistent unit-window lower bound. Instead it uses the local compactness interpretation of a failed carrier component: the nonintegrable component defines normalized carrier blocks; each block is then labelled by the pre-divergence first-failure order definition 8.55.

Definition 8.120 (Normalized carrier blocks for a failed component). Let $g \in L^1_{\text{loc}}([\tau_0, \infty))$ be nonnegative and assume

$$\int_T^\infty g(\tau) d\tau = \infty \quad \text{for every } T \geq \tau_0.$$

A normalized g -carrier block sequence on a tail $[\tau_1, \infty)$ is a sequence of intervals $[s_m, t_m]$, with $\tau_1 \leq s_m < t_m$, $s_m \rightarrow \infty$, $t_m \leq s_{m+1}$, and

$$\int_{s_m}^{t_m} g(\tau) d\tau = 1.$$

The associated carrier probability measures are

$$d\mu_m(\tau) := g(\tau) \mathbf{1}_{[s_m, t_m]}(\tau) d\tau.$$

For $g = J_i$, $i \in \{1, 7\}$, a J_i -occupation state means the terminal local occupation state obtained as follows. Fix any compact retained renormalized cylinder and record on it the repaired-gauge

variables (V, P, a, b) , the moving radius R , the cutoff Φ , and the component density J_i , using the weak local topology supplied by [proposition 3.10](#) and the pressure reconstruction for (V, P) , the compactified repaired-gauge topology for (a, b, R, Φ) , and the one-point compactification of $[0, \infty)$ for the scalar component density. Push μ_m forward by this local-data map and pass, on every compact cylinder and then diagonally, to a weak-* limit of probability measures on this metrizable local state space. Any such limit is a J_i -occupation state. Its label row is the first closed stratum met by its support in the pre-divergence procedure of [definition 8.55](#), with the carrier row refined by the active component J_i . If several named diagnostics are met at the same row, they are read as the declared label cluster for that row.

Lemma 8.121 (Closed local carrier diagnostics). *On the local state space of [definition 8.120](#), each row of the ordered procedure in [definition 8.55](#) is represented by a closed diagnostic set, after grouping diagnostics that are declared to belong to the same row. The component-refined carrier diagnostic for J_i , $i \in \{1, 7\}$, is also closed on normalized J_i -carrier states.*

Proof. Fix a compact retained cylinder. The solution variables are recorded in the weak local compactness topology of [proposition 3.10](#); pressure is recorded modulo spatial constants in the corresponding weak $L_{\text{loc}}^{3/2}$ topology; and the auxiliary scale, center, radius, and cutoff data are recorded in the compactified repaired-gauge topology, with the boundary faces labelled by loss of scale admissibility, modulation integrability, exterior tightness, multicore escape, or local H^1 control. Thus leaving a retained compact row is convergence to one of the prescribed boundary faces, hence a closed condition.

The quantitative rows are closed for the standard lower-semicontinuity reasons used throughout the retained compactness arguments: L^3 , pressure $L^{3/2}$, $L_t^2 H_x^1$, and local energy quantities are lower semicontinuous under the recorded weak/strong local convergence, while failure of a required finite bound is the closed face at infinity in the compactified coordinate. Critical-annulus, local H^1 -loss, modulation-loss, cost-compatibility, and carrier-comparison failures are therefore closed diagnostic sets in this topology. Equal-scale, separated, same-point cascade, gauge-degenerate, and scale-drift alternatives are closed because they are defined by limits of the recorded scale-center parameters in the finite parameter partition.

Finally, the component-refined carrier row is the closed face consisting of nonzero $J_i d\tau$ -occupation states. Each approximating block has total occupation mass

$$\mu_m(\mathfrak{X}) = \int_{s_m}^{t_m} J_i(\tau) d\tau = 1,$$

where $\mathfrak{X}$ denotes the local occupation-state space. Weak-* limits of probability measures have the same total mass. Hence the row “ J_i is not negligible/summable on the carrier tail” is closed on the normalized occupation-measure compactification. This is a statement about the occupation measure; it does not choose pointwise times and does not assert a unit-window lower bound. $\square$

Lemma 8.122 (Compactness and labels for normalized occupation states). *Let $i \in \{1, 7\}$. Every normalized J_i -carrier block sequence for a pre-carrier branch has a subsequential J_i -occupation state in the sense of [definition 8.120](#). The state has a unique first label row in the finite ordered list of [definition 8.55](#). If no non-carrier label occurs before carrier validation, then the first carrier label is the component-refined J_i -moving-annulus summability failure.*

Proof. On every compact retained cylinder, the pre-carrier hypotheses give the represented suitable estimates, pressure reconstruction, local L^3 bounds, and local $L_t^2 H_x^1$ bounds used in [proposition 3.10](#). The repaired-gauge modulation data are locally integrable by [lemma 8.69](#), and the moving radius is locally absolutely continuous. Hence the local-data maps take values in a metrizable product of weakly compact bounded sets after restricting to a subsequence and

a compact cylinder. Banach–Alaoglu, the pressure compactness part of [proposition 3.10](#), and a diagonal extraction over compact cylinders give weak-* convergence of the pushed-forward probability measures. This is precisely the occupation state in [definition 8.120](#).

The ordered validation list is finite and its strata are evaluated in order. By [lemma 8.121](#), each row is a closed diagnostic set on the local occupation-measure compactification. The first row whose closed diagnostic meets the occupation-state support is therefore unique; ties inside that row are exactly the declared label cluster. If a non-carrier diagnostic appears before carrier validation, that row is the label row. If all non-carrier diagnostics pass on the support, the carrier row is reached. Because the measures are normalized by $\int J_i d\tau = 1$, the limiting occupation state has total carrier mass one. Thus the active carrier row cannot be the statement that J_i is summable and negligible on the carrier tail. The component-refined failure of J_i -moving-annulus summability is then the first carrier label. This is a compactness label, not a pointwise lower bound for J_i . $\square$

Lemma 8.123 (Occupation-state extraction from nonintegrability). *Let $g \in L^1_{\text{loc}}([\tau_0, \infty))$ be nonnegative and assume*

$$\int_T^\infty g(\tau) d\tau = \infty \quad \text{for every } T \geq \tau_0.$$

Then every tail $[\tau_1, \infty)$ contains a normalized g -carrier block sequence in the sense of [definition 8.120](#). After passing to a subsequence, exactly one of the following two density regimes occurs:

- (C1) Persistent unit-window carrier mass. *There are $\gamma > 0$ and unit intervals $U_m = [T_m, T_m + 1]$, $T_m \rightarrow \infty$, such that*

$$\int_{U_m} g(\tau) d\tau \geq \gamma.$$

- (C2) Diffuse carrier mass. *For the normalized block subsequence retained in this case,*

$$\sup_{T \in [s_m, t_m]} \int_{[T, T+1] \cap [s_m, t_m]} g(\tau) d\tau \rightarrow 0.$$

The probability measures $d\mu_m = g \mathbf{1}_{[s_m, t_m]} d\tau$ are then the diffuse local occupation states of g .

If, in addition, $q(\tau) := |a(\tau)| + \|b(\tau)\|_\infty$ satisfies the uniform unit-window bound

$$Q := \sup_{T \geq \tau_0} \int_T^{T+1} q(\tau) d\tau < \infty.$$

then in case (C1) either there are selected times $\tau_m \in U_m$ such that

$$g(\tau_m) \geq \gamma/2, \quad q(\tau_m) \leq 4Q/\gamma.$$

after relabelling, or the normalized g -carrier mass on U_m is carried by the high-modulation sets

$$H_m := \{\tau \in U_m : q(\tau) > 4Q/\gamma\}, \quad \int_{H_m} g(\tau) d\tau \geq \gamma/2.$$

The latter is a modulation occupation state; it is not converted into a shell or scale lower bound.

Proof. Block construction. Fix a tail $[\tau_1, \infty)$. Choose $s_1 \geq \tau_1$. Since every tail integral of g is infinite, there is a finite $t_1 > s_1$ with $\int_{s_1}^{t_1} g = 1$; for instance take the first level time of the absolutely continuous primitive $\int_{s_1}^T g$. Having chosen t_m , choose $s_{m+1} \geq t_m + 1$ and then $t_{m+1} > s_{m+1}$ with $\int_{s_{m+1}}^{t_{m+1}} g = 1$. This gives normalized carrier blocks and $s_m \rightarrow \infty$.

Persistent versus diffuse. If case (C1) fails, then for each j there is $m(j)$ such that every interval of the form $[T, T+1] \cap [s_m, t_m]$ inside every retained block with index $m \geq m(j)$ has g -mass at most $1/j$. A diagonal subsequence gives the displayed diffuse condition. Conversely, if (C1) holds then the block sequence has a persistent unit-window carrier by definition. The two alternatives are therefore exhaustive after subsequence extraction.

Bounded modulation on persistent windows. In case (C1), set

$$Q := \sup_{T \geq \tau_0} \int_T^{T+1} q(\tau) d\tau.$$

If $Q = 0$, then $q = 0$ for a.e. time in every U_m . Since $\int_{U_m} g \geq \gamma$ and $|U_m| = 1$, the set $\{\tau \in U_m : q(\tau) = 0, g(\tau) \geq \gamma/2\}$ has positive measure after modifying g, q only on null sets; otherwise $g < \gamma/2$ a.e. on $\{q = 0\} \cap U_m$ and $\int_{U_m} g \leq \gamma/2 < \gamma$. Choose τ_m to be a Lebesgue time in that set. This gives the low-modulation alternative. Assume now that $Q > 0$. Set

$$H_m := \{\tau \in U_m : q(\tau) > 4Q/\gamma\}, \quad E_m := U_m \setminus H_m.$$

Chebyshev gives $|H_m| \leq \gamma/4$. Since $\int_{U_m} g \geq \gamma$, either $\int_{E_m} g \geq \gamma/2$ or $\int_{H_m} g \geq \gamma/2$ along a subsequence. In the first case $|E_m| \leq 1$ gives a Lebesgue time $\tau_m \in E_m$ with $g(\tau_m) \geq \gamma/2$ and $q(\tau_m) \leq 4Q/\gamma$. In the second case the displayed high-modulation carrier mass holds. This proves the final dichotomy and deliberately leaves the second case as a occupation state rather than replacing it by a pointwise shell or scale estimate. $\square$

Lemma 8.124 (Transition carrier label is exterior/leakage). *Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family of lemma 8.116. If*

$$J_1 \notin L^1(\tau_1, \infty)$$

on every later tail, then every J_1 -occupation state of definition 8.120 satisfies the exterior/leakage alternative of definition 8.110. More precisely, either an earlier non-carrier exterior/leakage label occurs, or the component-refined J_1 -moving-annulus summability failure is the transition-leakage carrier label, which belongs to the same exterior/leakage cluster.

Proof. By lemma 8.122, every normalized J_1 -block has a subsequential labelled occupation state. Since ω is in the pre-carrier regime, the analytic, representation, retained finite critical-mass, compact-tightness, local H^1 , modulation, and cost-compatibility rows in definition 8.87 have already passed on the retained branch. By lemma 8.121, these rows are closed diagnostic sets in the occupation-measure compactification. Hence if one fails in the carrier limit, the failure is exactly a loss of retained compactness, exterior concentration/multicore escape, or local windowed H^1 -loss by the ordered classification in definition 8.55. Those are exterior/leakage labels.

It remains to identify the first carrier label when all earlier rows remain retained. The component

$$J_1(\tau) = \left| \int |V|^2 \partial_\tau \Phi(\tau, \cdot) \right|$$

is the time-transition part of the moving cutoff. It contains no pressure, convection, translation, or scale-drift factor. The exact identity

$$\partial_\tau \Phi(\tau, y) = -\frac{R'(\tau)}{R(\tau)^2} y \cdot \nabla \phi(y/R(\tau))$$

shows that J_1 measures only occupation of the moving boundary of the retained compact window. Therefore the component-refined carrier row supplied by lemma 8.122 is the *transition-leakage carrier label*: the retained moving boundary carries non-summable transition occupation. This is not a new compact-cost stratum. In the finite compactness list it is assigned to the

same exterior/leakage cluster as critical-tightness failure, exterior concentration, and local windowed H^1 -loss, because the only local interpretation of non-summable boundary transition in a retained compact single-core state is that mass is crossing the retained boundary or the compact regularity needed to ignore that crossing has failed. Thus every possible first label satisfies the exterior/leakage alternative. $\square$

Theorem 8.125 (Windowwise transition/leakage alternative by occupation states). *Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family of lemma 8.116. If*

$$J_1 \notin L^1(\tau_1, \infty)$$

on every later tail, then ω enters the exterior/leakage alternative of definition 8.110.

Proof. Apply lemma 8.123 with $g = J_1$. It gives normalized J_1 -carrier blocks on every terminal tail, with either persistent unit-window mass or diffuse carrier mass. In the persistent case, the supplementary modulation selection gives bounded selected-time representatives when the low-modulation subcase occurs; in the high-modulation subcase and in the diffuse case one keeps the normalized occupation state instead of imposing an artificial pointwise lower bound.

All three outputs are J_1 -occupation states in the sense of definition 8.120. By lemma 8.124, each such state either has an earlier exterior/leakage label or has the transition-leakage carrier label identified there with the exterior/leakage cluster. This is exactly the exterior/leakage alternative allowed in definition 8.110. $\square$

Corollary 8.126 (Windowwise transition/leakage alternative holds). *The pre-carrier validated compact-cost regime has the windowwise transition/leakage alternative property in the sense of definition 8.118.*

Proof. This is the definition of definition 8.118 applied with theorem 8.125. $\square$

Lemma 8.127 (Retained positive scale-work realizes a scale-rigid core). *Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family. Let μ_m be any normalized J_7 -carrier block sequence and let $\mathfrak{s}$ be a subsequential J_7 -occupation state. Assume that $\mathfrak{s}$ has no earlier exterior/leakage label, no modulation-driven recentering label, and no negative scale-drift label. Then, after admissible scale-center reselection, the represented sequence carried by $\mathfrak{s}$ is produced by Step 3 or Step 4 of theorem 5.20 or by theorem 5.21, and satisfies the four clauses (R1)–(R4) of definition 5.19.*

Proof. Step 1: what the J_7 -carrier records. By definition,

$$J_7(\tau) = a_+(\tau) \int |V(y, \tau)|^2 (-y \cdot \nabla \Phi(\tau, y)) dy.$$

Thus each normalized block satisfies

$$\int_{s_m}^{t_m} a_+(\tau) \int |V|^2 (-y \cdot \nabla \Phi) dy d\tau = 1.$$

This is positive scale-generator work carried by the retained moving window. It is a scale-center compactness datum, not a consequence of an annular L^2 -to-CKN conversion.

Step 2: compact-core realization. By lemma 8.122, $\mathfrak{s}$ is a labelled local occupation state. The absence of an exterior/leakage label means that the state remains in the retained compact single-core window: critical tightness has not failed, no exterior/multicore loss has occurred, and local windowed H^1 control is retained. Consequently the compactness, pressure-reconstruction, and suitable-energy estimates inherited from the pre-carrier regime give a fixed compact selected cylinder Q_{R_*} on which the represented sequence has the uniform $A + E + C + D$ bound required

in (R1). This is only the usual compact-cylinder suitable bound; it is not an L^∞ or regularity assertion.

The positive inner concentration needed for (R2) is inherited from the retained finite critical-mass row itself, not inferred from the J_7 -density. This is the compact-window active mass supplied by [lemma 10.78](#). Since the present occupation state has no exterior/leakage label, the corresponding critical mass remains inside the retained compact single-core region. Cover a slightly smaller retained cylinder by finitely many CKN cylinders. If each member of this finite cover had arbitrarily small velocity-pressure CKN quantity along the represented sequence, then the retained positive local critical mass would either vanish on the compact core or exit through the exterior face; the first contradicts the retained mass row and the second is exactly the exterior/leakage label excluded above. Thus, after passing to a subsequence and reselecting the scale-center within this finite cover, there are $Q_{r_*} \in Q_{R_*}$ and $\eta_0 > 0$ with positive CKN mass. This is (R2).

If the modulation were unbounded on the selected compact cylinder, the component-refined carrier label would be the modulation-driven recentering branch. That label is excluded by hypothesis, so the repaired-gauge modulation remains bounded on the selected cylinder; this is (R3).

Step 3: scale-rigid relative selection. The normalized J_7 -mass is nonzero positive scale-generator work. If this work could be absorbed by a summable scale correction, the J_7 -moving-annulus row would pass on a terminal carrier tail, contradicting the normalized carrier label. If the work instead caused scale collapse or negative scale drift, the first label would be the negative-drift branch, excluded by hypothesis.

Apply the finite parameter partition [lemma 5.10](#) to the retained compact core selected in Step 2 and to the inherited active ledger. A separated physical-center relation is exactly the input of [theorem 5.21](#). A same-point comparable-scale relation is reduced by Step 3 of [theorem 5.20](#), and a strict same-point scale separation is reduced by Step 4 of that theorem. The nonretained outputs of this partition are exterior/leakage, modulation-driven recentering, negative drift, or perturbative absorption; the first three are excluded by hypothesis and perturbative absorption would make the J_7 -row summable. The only retained compact scale-center label left is therefore the scale-rigid residual alternative produced by those reductions. In the compact gauge-degenerate case, Step 3 of [theorem 5.20](#) supplies bounded (ρ_n, z_n) and $\sigma_{\min}(\mathcal{J}_n) \rightarrow 0$, giving (R4a). In the innermost same-point case, Step 4 supplies $\rho_n \rightarrow \infty$, $|z_n|/\rho_n = O(1)$, and the velocity, equation, and pressure-residual convergence in (R4b). Thus the branch data and clauses (R1)–(R4) of [definition 5.19](#) all hold. $\square$

Theorem 8.128 (Positive scale-work occupation states enter the scale-rigid branch). *Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family. Let μ_m be any normalized J_7 -carrier block sequence, and let $\mathfrak{s}$ be any subsequential J_7 -occupation state generated by μ_m . If $\mathfrak{s}$ is not labelled by exterior/leakage, by modulation-driven recentering, or by negative scale drift, then the represented branch carried by $\mathfrak{s}$ is excluded from the local Type II class by [proposition 5.42](#).*

Proof. [Lemma 8.127](#) realizes the J_7 -occupation state $\mathfrak{s}$ as a scale-rigid local branch satisfying [definition 5.19](#). Applying [proposition 5.42](#) excludes that branch. $\square$

Lemma 8.129 (Positive-scale occupation-state stratification). *Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family of [lemma 8.116](#), with cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$ and J_7 defined as in [definition 8.100](#). If*

$$J_7 \notin L^1(\tau_1, \infty)$$

on every later tail, then every J_7 -occupation state is labelled either by exterior/leakage, by modulation-driven recentering, by the Section 6 negative scale-drift branch, or by the scale-rigid residual branch excluded in [theorem 8.128](#).

Proof. Occupation-state extraction is [lemma 8.123](#) with $g = J_7$. The ordered first-failure map of [definition 8.55](#) first removes exterior loss, critical-tightness failure, local H^1 -loss, and modulation-driven recentering. The scale decomposition then separates the favorable negative scale-drift obstruction, already assigned to Section 6, from retained compact states carrying positive scale work. The retained compact states are realized as scale-rigid states by [lemma 8.127](#) and excluded by [theorem 8.128](#). These alternatives exhaust the J_7 -carrier because no other moving component is being tested here and J_2, J_3, J_4, J_5 are already summable on the canonical windowwise family by [lemma 8.116](#). $\square$

Theorem 8.130 (Windowwise positive-scale alternative by occupation states). *Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family of [lemma 8.116](#), with cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$ and J_7 defined as in [definition 8.100](#). If*

$$J_7 \notin L^1(\tau_1, \infty)$$

on every later tail, then ω enters the positive-scale alternative of [definition 8.110](#), unless it has already entered the exterior/leakage, modulation-driven recentering, or negative-drift alternatives.

Proof. Apply [lemma 8.129](#). Exterior/leakage carrier labels already satisfy alternative (iii) of [definition 8.110](#), modulation-driven recentering satisfies alternative (iv), and negative drift satisfies alternative (vi). In the retained positive-scale case, [theorem 8.128](#) places the carrier in the scale-rigid residual branch and then applies [proposition 5.42](#). The persistent unit-window subcase gives the pointwise selected-time form of alternative (v) whenever bounded-modulation points are selected; the high-modulation and diffuse subcases give the occupation-state form of the same alternative. Hence there is no unhandled diffuse J_7 tail. $\square$

Corollary 8.131 (Windowwise positive-scale alternative holds). *The pre-carrier validated compact-cost regime has the windowwise positive-scale alternative property in the sense of [definition 8.119](#).*

Proof. This is exactly the alternatives allowed by [definition 8.119](#), supplied by [theorem 8.130](#). $\square$

Remark 8.132 (Status of the windowwise alternative definitions). [Corollaries 8.126](#) and [8.131](#) together prove unconditionally that the pre-carrier validated compact-cost regime has the alternative properties of [definitions 8.118](#) and [8.119](#). The two definitions are therefore proved properties in the present manuscript; their proof is reduced to the occupation-measure compactness and extraction argument [lemmas 8.121](#) to [8.123](#), the transition carrier label in [lemma 8.124](#), the positive scale-work compact-core realization [lemma 8.127](#), the scale-rigid exclusion argument [theorem 8.128](#), and the canonical windowwise schedule of [lemma 8.115](#). The proof does not use the invalid implication from nonintegrability to persistent unit-window lower bounds; diffuse tails are retained as normalized occupation states and classified by the local ordered stratification.

Theorem 8.133 (CostPhase 5: canonical windowwise carrier routing). *Assume every branch in the pre-carrier validated compact-cost regime uses a validated compact cost that is energy-controlled by the fixed identity cost. Assume repaired-gauge suitable local energy gives the moving carrier subidentity in the sense of [lemma 8.89](#). Assume also that the regime has the windowwise transition/leakage and windowwise positive-scale alternative statements of [definitions 8.118](#) and [8.119](#), and that failure of $J_6 \in L^1(\tau_1, \infty)$ on every later tail for the canonical windowwise moving family forces entry into the negative scale-drift obstruction of [definition 10.25](#). Then the regime has stratified moving-cutoff alternative in the sense of [definition 8.110](#).*

Proof. Let ω be a branch in the pre-carrier validated compact-cost regime, and let $R(\tau)$ be the canonical windowwise moving family of [lemma 8.116](#). Write $J_1, \dots, J_7$ for its moving carrier components.

By [lemma 8.116](#),

$$J_2, J_3, J_4, J_5 \in L^1(\tau_0, \infty).$$

If there exists $\tau_1 \geq \tau_0$ such that

$$J_1, J_6, J_7 \in L^1(\tau_1, \infty),$$

then all seven components are integrable on $[\tau_1, \infty)$. Therefore [lemma 8.113](#) gives the moving-cutoff localized-energy closure criterion, so ω satisfies alternative (ii) of [definition 8.110](#).

Assume now that no later tail has J_1, J_6, J_7 all integrable. Since J_2, J_3, J_4, J_5 are already integrable on the full tail, at least one of J_1, J_6, J_7 fails on arbitrarily late tails. Fix such a component. Failure on arbitrarily late tails implies failure on every later tail: if the component were integrable on some tail, it would be integrable on every subsequent tail, contradicting arbitrarily late failure. If J_1 fails on every later tail, the windowwise transition/leakage alternative statement gives alternative (iii) of [definition 8.110](#). If J_6 fails on every later tail, the negative scale-drift routing statement gives alternative (vi). If J_7 fails on every later tail, the windowwise positive-scale alternative statement gives alternative (v), unless the branch has already entered alternative (iii), (iv), or (vi). Therefore one of the alternatives in [definition 8.110](#) always holds. $\square$

Corollary 8.134 (CostPhase 5: adjacent closure reduced to windowwise carrier routing). *Assume the non-carrier adjacent strata listed in [theorem 8.92](#) are closed by their corresponding local compactness or interface arguments in the ambient Type II class. Assume also that the exterior concentration, local windowed H^1 -loss, modulation-driven recentering branches, scale-rigid residual branches, and the negative scale-drift obstruction are closed by their assigned local arguments in the ambient class. Assume every branch in the pre-carrier validated compact-cost regime uses a validated compact cost that is energy-controlled by the fixed identity cost, that repaired-gauge suitable local energy gives the moving carrier subidentity, and that the windowwise transition/leakage and windowwise positive-scale alternative statements of [definitions 8.118](#) and [8.119](#) hold, together with the negative-drift routing statement used in [theorem 8.133](#). Then the compact-cost regime has adjacent-stratum closure in the sense of [definition 8.83](#).*

Proof. Let ω enter a named adjacent row in the pre-divergence order [definition 8.55](#). If the row occurs before the pre-carrier channel, it is a non-carrier row and is closed by the first hypothesis together with [theorem 8.92](#). It remains to consider a purported carrier-validation failure after ω has entered the pre-carrier channel. [Theorem 8.133](#) gives one of the alternatives in [definition 8.110](#). Alternatives (i) and (ii) supply the fixed- or moving-cutoff carrier criterion, respectively. In either case the carrier-validation row passes, contradicting the assumption that this row is the first failure. The remaining alternatives are precisely the exterior concentration, local windowed H^1 -loss, modulation-driven recentering, scale-rigid residual, positive-scale, or negative scale-drift branches listed in the hypotheses of this corollary, hence are closed by their assigned local arguments. Therefore every named adjacent row has been closed, which is [definition 8.83](#). No divergence statement or cost-divergence exclusion theorem has been used. $\square$

Lemma 8.135 (Ambient non-carrier exits close by ordered local routing). *Let ω be an admissible physical local Type II branch which has passed the covered-entry theorem and then reaches the compact-cost validation layer. If ω leaves the retained compact-cost stratum through a non-carrier first failure listed in [theorem 8.92](#), then the branch is either class-incompatible, closed by a local closure theorem, or assigned to a named local state-space row whose closure is supplied by the ambient carrier/terminal assembly.*

Proof. The pre-divergence validation in [definition 8.55](#) and [theorem 8.92](#) lists the possible non-carrier first failures. We check them in that order. Exhaustion failures are routed by [theorem 10.122](#); the only alternatives not retained there are class-incompatible alternatives or named local exits. Representation, pressure, and modulation failures are the ordered failures in [corollaries A.12](#) and [A.19](#) and the absolutely continuous repaired-gauge consequences of [theorem 10.56](#); when the repaired chart is available, pressure and modulation are consequences, and when it is not available the failure is the representation alternative. On a physical branch which has already passed covered entry, representation failures cannot remain as unclassified analytic gaps: if a nondegenerate represented chart is not available, the local state-decomposition places the branch in a multibubble/cascade, gauge-degenerate, scale-degenerate, finite-cost scale-collapse, compact scale-collapse, or scale-rigid alternative. These rows are closed by the local closure theorems [theorems 7.37, 11.6](#) and [11.8](#) and [proposition 5.42](#).

Critical-mass failures are used here only as local terminal-window diagnostics. Failure of lower retained mass means that the candidate was not a retained active core, while failure of a finite local upper L^3 bound on a compact selected cylinder produces an additional compact active label; this is routed by [lemma 10.78](#). Once the retained local finite critical-mass window is reached, the compact tests are closed by [theorem 10.108](#) and [corollary 10.109](#). Exterior tightness and leakage are routed by [theorems 9.39](#) and [9.72](#). Local windowed H^1 -loss is the rough-core row; by [theorem 11.7](#) it forces failure of compact CKN control on enclosing compact cylinders and hence returns to the multibubble/cascade package [theorem 11.6](#). Multibubble, multicore, cascade, and gauge-degenerate rows are closed directly by [theorem 11.6](#). Finite-cost scale collapse is eliminated by [theorem 11.8](#). Scale-rigid residual and retained compact scale-collapse rows are eliminated by [proposition 5.42](#) and [theorem 7.37](#). Finally, cost well-posedness and cost-comparison failures are precisely the ordered cost rows of [definitions A.22](#) and [A.23](#), [lemma A.24](#), and [theorem A.25](#); for the canonical identity cost the comparison is a theorem, and any non-identity failure is recorded as its own cost-compatibility exit. The thin-drift, autonomous-modulation, and compactness-package first failures which can arise inside the scale-collapse stratification are the first failures recorded by [lemma 7.26](#). Thin drift is routed by [theorem 10.118](#); compactness package failure is routed by [theorem 7.34](#); modulation failure remains a non-retained scale-collapse exit. The scale-drift face is recorded by [corollary 10.30](#), and the thin face is closed as an independent terminal row by [theorem 10.118](#). The retained compact scale-collapse state is excluded by [theorem 7.37](#). Thus every non-carrier first failure is either class-incompatible, closed by a local closure theorem, or recorded as a named carrier/scale/cost row for the ambient assembly. No step uses a global L^3 bound, a global profile theorem, or vanishing localized dissipation outside the selected local branch. $\square$

Lemma 8.136 (Carrier applicability on physical compact-cost branches). *Every physical local Type II branch in the pre-carrier validated compact-cost regime either exits through the ordered cost-compatibility row, or else satisfies the carrier applicability hypotheses used in [theorem 8.133](#) on the retained canonical channel: the validated compact cost is the fixed identity cost and is energy-controlled with constant 1, and suitable repaired-gauge local energy supplies the fixed and moving carrier subidentities.*

Proof. The pre-carrier regime already includes the repaired-gauge representation, local pressure reconstruction, local cost well-posedness, and compatibility of the validated compact cost. The ordered cost row is deliberately read before carrier validation. If the selected non-identity diagnostic has only the lower comparison used for cost-divergence transfer, but lacks the upper energy-control comparison needed by the carrier inequality, the branch is not silently promoted to an energy-controlled carrier branch; it is recorded as the cost-compatibility or carrier-comparison exit handled by [lemma 8.135](#). On the retained carrier branch the canonical cost channel is

selected:

$$\mathfrak{C}_{\Pi}^{NS} = \tilde{\mathfrak{D}}_{R_0}.$$

By [lemma A.24](#) this is a valid localized cost, and by [definition 8.58](#) it is energy-controlled with constant 1. This uses the identity equality only; no unproved upper comparison for a non-identity cost is invoked.

The branch comes from a suitable physical solution and has an absolutely continuous repaired-gauge chart on each compact renormalized window. Hence [lemma 10.59](#) gives the renormalized local energy inequality, and [lemma 8.89](#) converts it into the fixed and moving carrier subidentities. These are local compact-window statements; no whole-space energy or global compactness estimate is used. $\square$

Lemma 8.137 (Negative drift is an assigned local branch). *On every physical local Type II branch in the canonical windowwise pre-carrier regime, failure of $J_6 \in L^1(\tau_1, \infty)$ on every later tail is exactly the negative scale-drift branch of [definition 10.25](#). This branch is routed by the local scale-drift alternatives and the local scale-collapse packages; it is not an additional ambient hypothesis.*

Proof. The J_6 component is the negative-scale contribution in the moving localized-energy identity. By [definition 10.25](#) and [corollary 10.30](#), its nonintegrability on every later tail is precisely the scale-collapse drift obstruction. From this point the branch is not closed by feeding the same adjacent-closure statement back into itself; it is passed to the local scale-collapse stratification.

Before applying scale-collapse rigidity, verify the local prerequisites by the same first-failure order. The pre-carrier retained branch supplies the repaired-gauge representation, active super-critical scaling, retained compact-window local critical mass, critical tightness, local H^1 control, and suitable local energy on compact cylinders. If any of these inputs fails, or if the selected core is not nonvanishing in the retained local window, then the first failure is a critical-mass, exterior/leakage, rough-core, representation, or compactness exit already closed by [lemma 8.135](#). If the scale path is not a genuine retained collapse path, the branch is a finite-cost, scale-degenerate, or bounded-window/scale-rigid exit, again a named local alternative. Thus the only J_6 -failure case remaining after ordered routing is a represented branch with a nonvanishing selected core and genuine retained scale collapse.

Apply [theorems 7.38](#) and [11.9](#) to that remaining branch. The finite scale-cost and finite absolute-cost alternatives are excluded by [theorem 11.8](#). Thin drift is routed by [theorem 10.118](#): persistent windows return to the thick-window ledger, while diffuse windows either enter another named local row or a closed retained row. Autonomous-modulation failure, compactness-package first failure, exterior/separated-core loss, rough-core loss, and multicore or same-point cascade are exactly the remaining non-retained first failures listed in [lemma 7.26](#); by [theorem 7.36](#) they are assigned to named local state-space rows before the retained compact scale-collapse state is formed. The only retained scale-collapse state left after those exits is the retained compact scale-collapse state, and it is excluded by [theorem 7.37](#); if the retained state becomes scale-rigid, it is excluded by [proposition 5.42](#). Thus J_6 -failure is a routed local scale-collapse branch and not a hidden input. $\square$

Theorem 8.138 (CostPhase 5: ambient adjacent closure for physical Type II branches). *Every physical local Type II branch which reaches the pre-divergence carrier layer satisfies adjacent-stratum closure for the compact-cost regime in the sense of [definition 8.83](#). Equivalently, the hypotheses of [corollary 8.134](#) are theorems on the physical local Type II state-space branch.*

Proof. Non-carrier adjacent exits are routed by [lemma 8.135](#); the local closure rows close there, while carrier, scale, and cost rows are handled by the remaining branchwise checks below. The exterior, local H^1 -loss, modulation-driven recentering, scale-rigid residual, and negative scale-drift

branches appearing in the stratified moving-carrier alternative are the named local branches routed in [theorems 9.39, 10.56](#) and [11.7](#), [proposition 5.42](#), and [lemma 8.137](#). Carrier applicability and energy-control of the canonical cost are supplied by [lemma 8.136](#). The windowwise transition and positive-scale alternative statements are proved on the canonical windowwise moving family by [corollaries 8.126](#) and [8.131](#). The negative-drift alternative required by [theorem 8.133](#) is [lemma 8.137](#).

All hypotheses of [corollary 8.134](#) are therefore verified branchwise. Applying that corollary gives compact-cost adjacent-stratum closure. Each cited input is a phase 0–5 routing statement or an earlier state-closure theorem; no cost-divergence exclusion is used. $\square$

Corollary 8.139 (CostPhase 5: ambient adjacent wrapper discharged). *The pre-divergence adjacent-stratum closure and carrier-applicability obligation for the canonical retained compact-cost branch is discharged by local state-space stratification and the canonical windowwise carrier package. Nonretained first failures are not removed by this corollary alone; they are assigned to the named exit rows recorded in the final assembly.*

Proof. This is [theorem 8.138](#). $\square$

Remark 8.140 (CostPhase 6: cost divergence on the retained compact stratum). The compact Type II cost argument gives infinite localized renormalization cost under the compact PDE hypotheses. Passing from that cost divergence to exclusion from the Type II class uses the corrected localized energy inequality constructed in [lemma 8.60](#). On the retained compact stratum, the carrier exists by [theorem 8.74](#); failure of the carrier test is not a retained compact branch but an adjacent stratum.

9. TERMINAL PROFILE DECOUPLING AND EXTERIOR REGULARITY

This section records the terminal local estimates used to remove multibubble configurations, small residual components, and components carried at a positive physical distance from the selected Type II core. The statements are local implications under the analytic inputs explicitly stated below: terminal state-space completeness on compact windows, repaired-gauge representation, local Caccioppoli estimates, exterior regularity, and the local compact scale-collapse routing theorem [theorem 7.37](#).

9.1. Renormalized frames and active local states. A renormalized frame consists of a center $x_c(t)$, a scale $\lambda(t) > 0$, and renormalized time

$$\frac{d\tau}{dt} = \lambda(t)^{-2}.$$

It represents the physical solution by

$$V(y, \tau) = \lambda(t)u(x_c(t) + \lambda(t)y, t).$$

When the repaired gauge is imposed, the represented velocity satisfies

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(\tau)(V + y \cdot \nabla V) + b(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0,$$

where

$$a(\tau) = \partial_\tau \log \lambda(\tau),$$

and b is the translation modulation coefficient determined by the centering condition. In a fixed frame, every other local state either remains at bounded relative position and comparable scale, escapes to spatial infinity in the frame variables, appears at a much larger same-point scale as a perturbative ambient field in the innermost frame, or lies at a smaller scale and therefore requires selecting that smaller active frame.

Definition 9.1 (Active local terminal state data). Let $t_n \uparrow T^*$. Active local terminal state data consist of local state-space classes $\mathcal{A}_j$, scales $\lambda_j^n \rightarrow 0$, centers $x_j^n \rightarrow x_j^*$, compact terminal windows I_j , and representatives W_j^n in the corresponding frames. Each active class has a detecting compact cylinder $B_{R_j} \times I_j$ and a local mass floor

$$\limsup_{n \rightarrow \infty} \|W_j^n\|_{L^\infty(I_j; L^3(B_{R_j}))} \geq \varepsilon_{\text{loc}} > 0.$$

After grouping comparable parameters based at the same physical point, distinct classes are separated in the local state-space sense: they have different physical points, strictly separated scales, or have already been assigned to an adjacent local alternative. On every compact terminal cylinder and for every finite retained set J , the local field is written

$$V_n = U_n^J + S_n^J,$$

where U_n^J is the grouped retained local state and S_n^J is the nonretained residual. Positive local mass in S_n^J is not a new unlabelled object: by terminal state-space completeness it is assigned either to a named adjacent local alternative or to the retained comparable class, in which case it must be included in U_n^J . No global profile decomposition, Brézis–Lieb identity, or heat-flow remainder smallness is part of this data. The corresponding local active mass floor is stated and proved in [theorem 9.33](#).

Lemma 9.2 (Finite active local state count on a terminal slice). *Fix a compact terminal cylinder $B_R \times I$, a selected time sequence $\tau_n \in I$, and assume*

$$\sup_n \|V_n\|_{L^\infty(I; L^3(B_R))} \leq M_R < \infty.$$

Assume every active local state detected on the slice τ_n inside B_R has local mass at least ε_{loc} , and that, after grouping comparable retained classes, the detecting balls on that slice have overlap multiplicity at most N_R . Then the number of active local states detected on that terminal slice is finite and satisfies

$$\#\mathcal{J}_{\text{act}}(R, I) \leq \left\lfloor \frac{N_R M_R^3}{\varepsilon_{\text{loc}}^3} \right\rfloor.$$

Proof. For each active state choose a detecting ball on the selected time slice, shrunk if necessary so that the family has overlap multiplicity at most N_R . At that common detecting time, the local mass floor gives

$$\int_{B_j} |V_n(y, \tau_n)|^3 dy \geq \varepsilon_{\text{loc}}^3 + o_n(1)$$

for the ball B_j associated with the j -th state. Summing over any finite active subfamily and using the overlap bound,

$$(\#J)\varepsilon_{\text{loc}}^3 \leq N_R \sup_n \sup_{\tau \in I} \int_{B_R} |V_n(y, \tau)|^3 dy \leq N_R M_R^3.$$

Thus every finite active subfamily has cardinality at most $N_R M_R^3 / \varepsilon_{\text{loc}}^3$. If there were more active states than the displayed integer bound, a finite subfamily with one additional element would contradict this estimate. $\square$

9.2. Same-point and separated-point reductions.

Lemma 9.3 (Comparable same-point states form a compound state). *Let a finite set $\mathcal{C}$ of active local states have the same physical center $x_j^* = x_*$. Suppose that for a representative center and scale (x_*^n, λ_*^n) ,*

$$\frac{\lambda_j^n}{\lambda_*^n} \rightarrow \rho_j \in (0, \infty), \quad \frac{x_j^n - x_*^n}{\lambda_*^n} \rightarrow z_j \in \mathbb{R}^3.$$

Let $W_{j,n}$ denote the local representative of the j -th state in its own frame on compact cylinders, and assume

$$W_{j,n} \rightarrow W_j \quad \text{strongly in } L_{\text{loc}}^3$$

on the relevant terminal slice, or in $L_{\text{loc}}^\infty L_{\text{loc}}^3$ on a compact terminal window. Write

$$\rho_j^n := \frac{\lambda_j^n}{\lambda_*^n}, \quad z_j^n := \frac{x_j^n - x_*^n}{\lambda_*^n}.$$

Then, in the frame (x_^n, λ_*^n) ,*

$$(\rho_j^n)^{-1} W_{j,n} \left(\frac{y - z_j^n}{\rho_j^n} \right) \rightarrow \rho_j^{-1} W_j \left(\frac{y - z_j}{\rho_j} \right)$$

strongly in the same local topology on every compact subcylinder on which the right-hand side is defined. Hence the comparable class has compound local representative

$$\Phi_{\mathcal{C}}(y) := \sum_{j \in \mathcal{C}} \rho_j^{-1} W_j \left(\frac{y - z_j}{\rho_j} \right).$$

If the compound representative carries the local mass floor

$$\|\Phi_{\mathcal{C}}\|_{L^3(B_R)} \geq \varepsilon_{\text{loc}}$$

on a detecting compact cylinder, it is one active local state in the selected frame. If no compact cylinder carries the local mass floor, it belongs to the perturbative residual class by [corollary 9.32](#).

Proof. Set

$$\rho_j^n = \frac{\lambda_j^n}{\lambda_*^n}, \quad z_j^n = \frac{x_j^n - x_*^n}{\lambda_*^n}.$$

The transformed local representative is

$$T_n W_{j,n}(y) = (\rho_j^n)^{-1} W_{j,n} \left(\frac{y - z_j^n}{\rho_j^n} \right),$$

and the proposed limit is

$$TW_j(y) = \rho_j^{-1} W_j \left(\frac{y - z_j}{\rho_j} \right).$$

Here $T_{\rho,z}F(y) := \rho^{-1}F((y - z)/\rho)$. Thus $T_n W_{j,n} = T_{\rho_j^n, z_j^n} W_{j,n}$ and $TW_j = T_{\rho_j, z_j} W_j$. Fix a compact ball B_R in the representative frame. For all large n , the preimage of B_R under $y \mapsto (y - z_j^n)/\rho_j^n$ is contained in a fixed compact ball $B_{R'}$ in the j -th frame. The local strong convergence gives

$$\|T_{\rho_j^n, z_j^n}(W_{j,n} - W_j)\|_{L^3(B_R)} \leq C\|W_{j,n} - W_j\|_{L^3(B_{R'})} \rightarrow 0.$$

It remains to move the fixed local limit W_j . Translation and dilation are strongly continuous on $L^3(B_{R'})$ after restriction to B_R : prove this first for $C_c^\infty(B_{R'})$ by uniform convergence on compact sets, then approximate W_j in $L^3(B_{R'})$. This proves the local convergence. The same argument applied uniformly in τ gives the $L_{\text{loc}}^\infty L_{\text{loc}}^3$ version on compact terminal windows. Summing over the finite class gives $\Phi_{\mathcal{C}}$. The final alternative is the local state-space criterion: a compound state with the detection threshold is retained as active, whereas a component with no compact-cylinder detection threshold is perturbative by the local residual theorem. $\square$

Lemma 9.4 (Innermost-frame rescaling of strict outer states). *Suppose all active local states have the same physical center x_* and split into distinct scale classes. If λ_1^n is the smallest active scale, then every outer scale $\lambda_j^n \gg \lambda_1^n$ satisfies*

$$\mu_j^n := \lambda_1^n / \lambda_j^n \rightarrow 0,$$

and in the innermost frame the j -th outer state has the local form

$$W_j^n(y) = \mu_j^n \widetilde{W}_{j,n} \left(\frac{x_*^n - x_j^n}{\lambda_j^n} + \mu_j^n y \right),$$

where $\widetilde{W}_{j,n}$ is the representative of the j -th state in its own outer frame.

Proof. The j -th local representative in physical coordinates is

$$u_j^n(x) = (\lambda_j^n)^{-1} \widetilde{W}_{j,n} \left(\frac{x - x_j^n}{\lambda_j^n} \right)$$

on the compact part of its outer terminal frame. Pull this back into the innermost frame $x = x_*^n + \lambda_1^n y$, with the critical rescaling $W_j^n(y) = \lambda_1^n u_j^n(x_*^n + \lambda_1^n y)$. Substituting,

$$W_j^n(y) = \frac{\lambda_1^n}{\lambda_j^n} \widetilde{W}_{j,n} \left(\frac{x_*^n + \lambda_1^n y - x_j^n}{\lambda_j^n} \right) = \mu_j^n \widetilde{W}_{j,n} \left(\frac{x_*^n - x_j^n}{\lambda_j^n} + \mu_j^n y \right),$$

which is the displayed form. The convergence $\mu_j^n \rightarrow 0$ is the hypothesis $\lambda_j^n \gg \lambda_1^n$. $\square$

Lemma 9.5 (Local L^3 -vanishing of strict outer states in the innermost frame). *Under the hypotheses of lemma 9.4, the outer states satisfy the following retained-frame alternative. If*

$$\xi_j^n := \frac{x_*^n - x_j^n}{\lambda_j^n}$$

stays in a compact subset of the outer frame and $\widetilde{W}_{j,n} \rightarrow W_j$ strongly in L^3_{loc} , then, for every $R < \infty$,

$$\|W_j^n\|_{L^3(B_R)} \rightarrow 0 \quad (n \rightarrow \infty).$$

If ξ_j^n leaves every compact subset of the outer frame, the contribution is assigned to the escaping-offset/exterior row of the local state-space decomposition rather than to a retained strict outer state.

Proof. Assume first that ξ_j^n stays in a fixed compact set K . Choose a compact K' containing $K + \mu_j^n B_R$ for all large n . By the formula in lemma 9.4,

$$\|W_j^n\|_{L^3(B_R)}^3 = \int_{\xi_j^n + \mu_j^n B_R} |\widetilde{W}_{j,n}(z)|^3 dz.$$

The strong $L^3(K')$ convergence reduces this to the same integral with W_j . Since $|W_j|^3 \in L^1(K')$ and the sets $\xi_j^n + \mu_j^n B_R$ have measure tending to zero, absolute continuity of the integral gives convergence to zero. If ξ_j^n escapes every compact set, then the outer state is not represented on a fixed compact outer cylinder; this is exactly the escaping-offset/exterior boundary face in the local state-space compactification. $\square$

Theorem 9.6 (Perturbative estimates for strict outer states). *In the same setting as lemma 9.4, the strict outer states are perturbative in the innermost retained frame:*

- (i) *finite-offset regular outer states, namely those for which $(x_*^n - x_j^n)/\lambda_j^n$ is bounded, decompose on each fixed inner ball into a constant drift plus an $O((\mu_j^n)^2)$ shear remainder;*
- (ii) *the constant drifts produced by item (i) are absorbed into the translation-modulation coefficient of the innermost frame;*
- (iii) *the remaining pressure, cutoff, and nonlinear interactions generated by the outer states are perturbative in the compact scale-collapse topology in which the local perturbation criterion lemma 9.10 and the pressure-tail decoupling lemma 9.12 apply.*

The escaping-offset alternative $(x_^n - x_j^n)/\lambda_j^n \rightarrow \infty$ is treated by lemma 9.5.*

Proof. In the escaping-offset case, lemma 9.5 gives local L^3 -vanishing on every innermost compact cylinder, and the local pressure-tail lemma gives the pressure and residual convergence.

It remains to consider bounded offsets. If the outer state is not regular at the observed offset, then some smaller offset cylinder has CKN quantity above the threshold of theorem 9.35. That cylinder is a positive local state-space label at a larger same physical point scale, so the branch is assigned to the same-point cascade/secondary-core alternative before entering the retained innermost class. Thus, in the retained case, every bounded offset cylinder for the outer state is CKN-subcritical. The CKN criterion gives uniform C^1 bounds on a smaller offset cylinder, after subtracting the pressure mean.

On that retained regular cylinder, Taylor expansion at the offset gives a spatially constant leading term plus a shear remainder of size $O(\mu_j^n)$ in C^1 on every fixed inner ball. The Navier–Stokes scaling contributes the additional factor μ_j^n in the velocity amplitude, so the nonconstant velocity remainder is $O((\mu_j^n)^2)$ locally. The spatially constant term is absorbed into the translation modulation coefficient. The remaining velocity, pressure, cutoff, and modulation terms vanish in

the compact-window topologies by [lemmas 9.10](#) and [9.12](#). This is exactly the asserted perturbative reduction. $\square$

Lemma 9.7 (Separated physical centers are locally invisible). *Let $p \neq q$ be two active physical concentration points,*

$$x_p^* \neq x_q^*, \quad d_{pq} := |x_p^* - x_q^*| > 0.$$

In the p -frame define the q -state contribution

$$W_{q \rightarrow p}^n(y) := \lambda_p^n (\lambda_q^n)^{-1} \phi^q \left(\frac{x_p^n + \lambda_p^n y - x_q^n}{\lambda_q^n} \right).$$

If $\phi^q \in L^3(\mathbb{R}^3)$, then for every $R < \infty$,

$$\|W_{q \rightarrow p}^n\|_{L^3(B_R)} \rightarrow 0.$$

If the corresponding representative pressure $\Pi^q \in L^{3/2}(\mathbb{R}^3)$, then the pure pressure contribution

$$P_{q \rightarrow p}^n(y) = (\lambda_p^n)^2 (\lambda_q^n)^{-2} \Pi^q \left(\frac{x_p^n + \lambda_p^n y - x_q^n}{\lambda_q^n} \right)$$

satisfies

$$\|P_{q \rightarrow p}^n\|_{L^{3/2}(B_R)} \rightarrow 0.$$

Proof. By the change of variables $x = x_p^n + \lambda_p^n y$,

$$\|W_{q \rightarrow p}^n\|_{L^3(B_R)}^3 = \int_{B(x_p^n, \lambda_p^n R)} (\lambda_q^n)^{-3} \left| \phi^q \left(\frac{x - x_q^n}{\lambda_q^n} \right) \right|^3 dx.$$

Set $z = (x - x_q^n)/\lambda_q^n$. Then

$$\|W_{q \rightarrow p}^n\|_{L^3(B_R)}^3 = \int_{B(c_n, \rho_n)} |\phi^q(z)|^3 dz,$$

where

$$c_n = \frac{x_p^n - x_q^n}{\lambda_q^n}, \quad \rho_n = \frac{\lambda_p^n R}{\lambda_q^n}.$$

Since $x_p^n \rightarrow x_p^*$, $x_q^n \rightarrow x_q^*$, and $\lambda_q^n \rightarrow 0$,

$$|c_n| \sim d_{pq}/\lambda_q^n \rightarrow \infty.$$

Moreover

$$\frac{\rho_n}{|c_n|} = \frac{\lambda_p^n R}{|x_p^n - x_q^n|} \rightarrow 0.$$

Thus the balls $B(c_n, \rho_n)$ escape to spatial infinity: for every $A < \infty$, they are contained in $\{|z| > A\}$ for all sufficiently large n . Since $\phi^q \in L^3$,

$$\int_{|z| > A} |\phi^q(z)|^3 dz \rightarrow 0 \quad (A \rightarrow \infty),$$

and hence $\|W_{q \rightarrow p}^n\|_{L^3(B_R)} \rightarrow 0$.

For the pressure term the same change of variables gives

$$\|P_{q \rightarrow p}^n\|_{L^{3/2}(B_R)}^{3/2} = \int_{B(c_n, \rho_n)} |\Pi^q(z)|^{3/2} dz.$$

The same escaping-ball argument and $\Pi^q \in L^{3/2}$ imply the stated pressure convergence. $\square$

Definition 9.8 (Multipoint decoupling conditions). For each active point and selected frame, the required multipoint decoupling conditions are:

- (i) local velocity decoupling in the selected frame;
- (ii) pressure decoupling modulo time-dependent constants in $L_\tau^1 H_y^{-1}$ on compact cylinders;
- (iii) localization errors from spatial cutoffs vanish in $L_\tau^1 H_y^{-1}$;
- (iv) repaired-gauge modulation coefficients remain bounded after exterior ambient terms are absorbed;
- (v) each active point has positive finite critical mass in its own frame;
- (vi) after states centered at other physical points are treated as perturbative exterior terms, the selected frame is tested against [definition 7.25](#): either the first failed retained-state input is routed by [lemma 7.26](#) to a named local exit, or all retained compact scale-collapse inputs hold.

Theorem 9.9 (Reduction of active multibubbles). *Assume the Type II local terminal state data, the active local mass floor, the same-point nonlinear decoupling conditions, the multipoint decoupling conditions of [definition 9.8](#). Then no active multibubble Type II configuration can occur.*

Proof. By [lemma 9.2](#), there are only finitely many active local states on each selected terminal slice. Partition them by physical concentration point x_α^* , and within each physical point partition them by comparable scale class.

If a scale class contains several comparable local states, [lemma 9.3](#) combines them into one compound state. If the compound state has no detecting compact cylinder, it is perturbative by [corollary 9.32](#); if it is active, it is a single active local state for that frame.

If a physical point contains strict same-point scale separation, [theorem 9.6](#) and the same-point nonlinear decoupling conditions reduce the innermost active scale to a single repaired-gauge solution with perturbative outer interactions. The compact scale-collapse routing theorem [theorem 7.37](#) then rules out the corresponding retained same-point cascade; if one of its retained hypotheses fails, [lemma 7.26](#) assigns the branch to a named local alternative and it is not a retained active multibubble configuration.

It remains to consider separated physical points. Choose any active point x_α^* and place the selected frame at its active scale and center. By [lemma 9.7](#), every state centered at a different physical point is locally L^3 -invisible in this frame. By the multipoint decoupling conditions, its pressure, cutoff, and modulation interactions are also perturbative in the compact scale-collapse topology. Therefore the selected frame contains a single active repaired-gauge solution whose scale tends to zero. If any retained compact scale-collapse input fails, the multipoint decoupling condition and [lemma 7.26](#) route the branch to a named local exit, so it is not a retained active multibubble configuration. If no retained input fails, the selected frame is a retained compact scale-collapse state in the sense of [definition 7.25](#). Applying [theorem 7.37](#) excludes that selected active point. Since the chosen active point was arbitrary, no active multibubble configuration can occur. $\square$

9.3. Local perturbation and pressure decoupling.

Lemma 9.10 (Local perturbation criterion). *Let $I \subset \mathbb{R}$ be compact and $Q_R = B_R \times I$. Let U_n be a repaired-gauge solution on Q_{2R} satisfying*

$$\|U_n\|_{L_\tau^\infty L_y^3(B_{2R})} + \|U_n\|_{L_\tau^2 H_y^1(B_{2R})} \leq C_R.$$

Let S_n be a divergence-free perturbation on Q_{2R} with

$$S_n \rightarrow 0 \quad \text{in } L_\tau^\infty L_y^3(B_{2R}),$$

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_R),$$

$$\mathcal{R}_n := \partial_\tau S_n - \nu \Delta S_n + a_n(S_n + y \cdot \nabla S_n) + b_n \cdot \nabla S_n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R),$$

where $|a_n| + |b_n| \leq M_{ab}$, and suppose the pressure cross-term satisfies

$$\nabla \mathcal{P}_{\text{cross},n} \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

Then replacing $U_n + S_n$ by U_n changes the repaired-gauge Navier–Stokes equation on Q_R by an error tending to zero in

$$L_\tau^1 H_y^{-1}(B_R).$$

Proof. The linear part is precisely $\mathcal{R}_n$, which tends to zero by hypothesis. The nonlinear difference is

$$(U_n + S_n) \cdot \nabla (U_n + S_n) - U_n \cdot \nabla U_n = \text{div } F_n, \quad F_n := U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n,$$

using $\nabla \cdot U_n = \nabla \cdot S_n = 0$ distributionally. By the assumed $L_\tau^1 L_y^2$ stress convergence, this divergence tends to zero in

$$L_\tau^1 H_y^{-1}(B_R),$$

because for $\varphi \in H_0^1(B_R; \mathbb{R}^3)$,

$$\left| \int_{B_R} (U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n) : \nabla \varphi \, dy \right| \leq \|F_n\|_{L^2(B_R)} \times \|\nabla \varphi\|_{L^2(B_R)}.$$

The pressure contribution vanishes by the assumed cross-pressure convergence. Therefore the total perturbative error tends to zero in

$$L_\tau^1 H_y^{-1}(B_R).$$

□

Lemma 9.11 (Static same-point outer profiles vanish locally). *Let*

$$W_n(y) = \mu_n \phi(\xi_n + \mu_n y), \quad \mu_n \rightarrow 0, \quad \phi \in L^3(\mathbb{R}^3).$$

Then, for every $R < \infty$,

$$\|W_n\|_{L^3(B_R)} \rightarrow 0.$$

Proof. Changing variables $z = \xi_n + \mu_n y$,

$$\|W_n\|_{L^3(B_R)}^3 = \int_{B(\xi_n, \mu_n R)} |\phi(z)|^3 \, dz.$$

The measure of $B(\xi_n, \mu_n R)$ tends to zero. Since $|\phi|^3 \in L^1(\mathbb{R}^3)$, absolute continuity of the integral gives

$$\int_{B(\xi_n, \mu_n R)} |\phi(z)|^3 \, dz \rightarrow 0,$$

uniformly in ξ_n . Hence $\|W_n\|_{L^3(B_R)} \rightarrow 0$. □

Lemma 9.12 (Pressure decoupling from kernel tails). *Let Γ_{ij} be the Calderón–Zygmund kernel for*

$$P = \mathcal{R}_i \mathcal{R}_j (F_{ij}).$$

Assume

$$F_n \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_A) \quad \text{for every fixed } A,$$

and

$$\sup_n \|F_n\|_{L_\tau^\infty L_y^{3/2}(\mathbb{R}^3)} < \infty.$$

Then the pressure generated by F_n satisfies, for every $R < \infty$,

$$P_n - c_{n,R} \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_R)$$

for suitable constants $c_{n,R}(\tau)$, and consequently

$$\nabla P_n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

More explicitly, for every $A > 2R$,

$$P_n - c_{n,A} = P_n^{\text{loc}} + H_{n,A} \quad \text{on } B_R,$$

where P_n^{loc} is generated by $F_n \mathbf{1}_{B_A}$, $H_{n,A}$ is harmonic on B_R , the local part tends to zero for fixed A , and

$$\lim_{A \rightarrow \infty} \sup_n \|H_{n,A}\|_{L_\tau^1 L_y^2(B_R)} = 0.$$

The constants in the first conclusion may be chosen as $c_{n,R} = c_{n,A_n}$ along a diagonal sequence $A_n \rightarrow \infty$.

Proof. Fix $A > 2R$ and write

$$F_n^{\text{loc}} = F_n \mathbf{1}_{B_A}, \quad F_n^{\text{far}} = F_n \mathbf{1}_{\mathbb{R}^3 \setminus B_A}.$$

For the local part, Calderón–Zygmund boundedness gives

$$\|P_n^{\text{loc}}\|_{L_\tau^1 L_y^2(B_A)} \leq C_A \|F_n\|_{L_\tau^1 L_y^2(B_A)} \rightarrow 0.$$

For the far part define

$$c_{n,A}(\tau) = \int_{|z|>A} \Gamma_{ij}(-z) F_{n,ij}(z, \tau) dz,$$

first for compactly supported approximations and then by passage to the $L^{3/2}$ limit. For $y \in B_R$ and $|z| > A > 2R$,

$$|\Gamma_{ij}(y - z) - \Gamma_{ij}(-z)| \leq C_R |z|^{-4}.$$

Hence

$$|P_n^{\text{far}}(y, \tau) - c_{n,A}(\tau)| \leq C_R \int_{|z|>A} |z|^{-4} |F_n(z, \tau)| dz.$$

By Hölder's inequality with exponents $3/2$ and 3 ,

$$\int_{|z|>A} |z|^{-4} |F_n| dz \leq \|F_n(\tau)\|_{L^{3/2}} \left(\int_{|z|>A} |z|^{-12} dz \right)^{1/3} \leq C A^{-3} \|F_n(\tau)\|_{L^{3/2}}.$$

Therefore

$$\|P_n^{\text{far}} - c_{n,A}\|_{L_\tau^1 L_y^\infty(B_R)} \leq C_{R,I} A^{-3} \sup_n \|F_n\|_{L_\tau^\infty L_y^{3/2}}.$$

Since B_R has finite measure, this also controls the harmonic far part

$$H_{n,A} := P_n^{\text{far}} - c_{n,A}$$

in $L_\tau^1 L_y^2(B_R)$, uniformly in n , with norm tending to zero as $A \rightarrow \infty$. It is harmonic on B_R because its source is supported outside B_A .

Choose $A_k \uparrow \infty$. For each k , the local convergence on B_{A_k} gives N_k such that

$$\|P_n^{\text{loc}}\|_{L_\tau^1 L_y^2(B_R)} \leq 2^{-k} \quad (n \geq N_k),$$

after increasing N_k so that $N_k < N_{k+1}$. Set $A_n = A_k$ for $N_k \leq n < N_{k+1}$, and set $c_{n,R} := c_{n,A_n}$. Then $A_n \rightarrow \infty$, the local part tends to zero, and the uniform far-part bound tends to zero. Thus

$$P_n - c_{n,R} \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_R).$$

For $\varphi \in H_0^1(B_R; \mathbb{R}^3)$,

$$|\langle \nabla(P_n - c_{n,R}), \varphi \rangle| \leq \left| \int_{B_R} (P_n - c_{n,R}) \nabla \cdot \varphi dy \right| \leq \|P_n - c_{n,R}\|_{L^2(B_R)} \|\nabla \varphi\|_{L^2(B_R)}.$$

Thus $\nabla P_n \rightarrow 0$ in $L_\tau^1 H_y^{-1}(B_R)$. $\square$

Definition 9.13 (Terminal nonlinear local-state decoupling). On every compact terminal-frame cylinder $B_R \times I$, let U_n be the retained active local-state evolution and let S_n be the nonretained local residual. Terminal nonlinear local-state decoupling means:

- (i) $\nabla \cdot S_n = 0$ distributionally;
- (ii) $S_n \rightarrow 0$ in $L_\tau^\infty L_y^3(B_{2R})$;
- (iii)
$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_R);$$
- (iv)
$$\partial_\tau S_n - \nu \Delta S_n + a_n(S_n + y \cdot \nabla S_n) + b_n \cdot \nabla S_n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R);$$
- (v) pressure sources involving at least one factor of S_n satisfy [lemma 9.12](#), or directly satisfy
$$\nabla \mathcal{P}_{\text{cross},n} \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R);$$
- (vi) cutoff errors introduced in isolating the selected frame vanish in $L_\tau^1 H_y^{-1}(B_R)$;
- (vii) the effective modulation coefficients remain bounded;
- (viii) after subtracting S_n , the retained solution has positive finite critical mass, repaired-gauge representation, pressure reconstruction, Caccioppoli regularity, compactness, tightness, and the compact scale-collapse limiting inputs.

Theorem 9.14 (Same-point and multipoint decoupling from terminal decoupling). *If terminal nonlinear local-state decoupling holds in the sense of [definition 9.13](#), then both same-point scale-decoupling and separated-point frame-decoupling hold.*

Proof. For same-point active local states, group comparable scales by [lemma 9.3](#). For strict scale separation choose the innermost active scale as the selected terminal frame. By [lemma 9.11](#), every outer scale that remains in the selected frame is locally invisible in velocity, unless it was already grouped as comparable or routed to an adjacent state. Terminal nonlinear local-state decoupling supplies the nonlinear-stress, linear-residual, pressure-source, localization, modulation, and compact scale-collapse admissibility conditions on every compact cylinder. [Lemma 9.10](#) then shows that the outer components and residual alter the selected equation only by an error tending to zero in $L_\tau^1 H_y^{-1}(B_R)$. Thus the same-point configuration reduces to one retained compact scale-collapse solution with perturbative interactions.

For separated physical points, choose an active physical point x_p^* and its selected frame. For every local state centered at $x_q^* \neq x_p^*$, [lemma 9.7](#) gives local velocity convergence to zero in the p -frame. If a pure profile pressure is represented in L^2 , the same escaping-ball calculation gives local L^2 pressure convergence and therefore H^{-1} convergence of the gradient. With only $L^{3/2}$ -pressure, the mixed pressure sources are controlled by [lemma 9.12](#) through the L^2 -strength hypotheses included in terminal nonlinear local-state decoupling. The same decoupling statement supplies localization-error convergence, bounded modulation, and compact scale-collapse admissibility. Applying [lemma 9.10](#) shows that separated physical-point states change the selected equation only by an error tending to zero in $L_\tau^1 H_{\text{loc}}^{-1}$. This proves separated-point decoupling. $\square$

9.4. Terminal local-state decoupling.

Definition 9.15 (Terminal active local frame). Partition active local states by physical point $x_j^n \rightarrow x_\alpha^*$. At a fixed physical point, group comparable scales. For a comparable-scale class $\mathcal{C}$, a representative frame $(x_{\mathcal{C}}^n, \lambda_{\mathcal{C}}^n)$ has compound local representative

$$\Phi_{\mathcal{C}}(y) = \sum_{j \in \mathcal{C}} \rho_j^{-1} W_j \left(\frac{y - z_j}{\rho_j} \right),$$

where W_j is the local representative of the j -th state in its own frame, transported as in [lemma 9.3](#). The class $\mathcal{C}$ is terminal at x_α^* if no active class at the same physical point has strictly smaller scale:

$$\lambda_{\mathcal{D}}^n / \lambda_{\mathcal{C}}^n \not\rightarrow 0 \quad \text{for every active class } \mathcal{D}.$$

The terminal frame is

$$y = \frac{x - x_{\mathcal{C}}^n}{\lambda_{\mathcal{C}}^n}, \quad \tau = \int^t (\lambda_{\mathcal{C}}(s))^{-2} ds,$$

and

$$V_n(y, \tau) = \lambda_{\mathcal{C}}^n u(x_{\mathcal{C}}^n + \lambda_{\mathcal{C}}^n y, t_n + (\lambda_{\mathcal{C}}^n)^2 \tau).$$

The retained solution U_n is the repaired-gauge nonlinear evolution of $\mathcal{C}$ in this frame, and $S_n := V_n - U_n$ is the nonretained remainder field.

Theorem 9.16 (Terminal windowed state-space completeness). *For a sequence f_n and a frame (x_n, λ_n) , define*

$$\mathfrak{m}_R(f_n; x_n, \lambda_n) := \limsup_{n \rightarrow \infty} \|\lambda_n f_n(x_n + \lambda_n \cdot)\|_{L^3(B_R)}$$

and

$$\mathfrak{M}_{R,I}(f_n; x_n, \lambda_n) := \limsup_{n \rightarrow \infty} \sup_{\tau \in I} \|\lambda_n f_n(x_n + \lambda_n \cdot, t_n + \lambda_n^2 \tau)\|_{L^3(B_R)}.$$

If $\mathfrak{M}_{R,I}(f_n; x_n, \lambda_n) > 0$ for a sequence left in the remainder, then after passing to a subsequence the local state-space stratification assigns that residual mass either to a named adjacent local alternative or to the retained comparable class. In the latter case the component must have been grouped into the retained terminal class before selecting the frame. Hence no positive nonretained residual mass remains on a retained terminal branch.

Proof. This is the compact-window form of the local residual exhaustion theorem [theorem 10.86](#), with selected-time spikes supplied by [proposition 10.80](#). If $\mathfrak{M}_{R,I}(f_n; x_n, \lambda_n) > 0$, select the times carrying this lower bound and apply that proposition to the corresponding residual coordinate. A first validation failure is a named adjacent row. Positive spacetime occupation, or a CKN-active parabolic rebase, is an actual descendant of the full represented branch and is assigned by local residual exhaustion either to a named adjacent row or to the retained comparable class. An adjacent label is handled by its local closure theorem and is not part of the retained terminal branch; a retained comparable label was grouped into the selected terminal class before the frame was fixed.

The remaining ordered alternative is a selected-time spike with vanishing compact spacetime occupation and no active rebase. By [lemma 9.30](#) it is perturbative in the full terminal PDE topology, and the terminal absorption step in [proposition 10.80](#) removes it before the finalized nonretained residual is fixed. Thus positive nonretained residual mass cannot persist on a retained terminal branch. $\square$

Remark 9.17 (Status of [theorem 9.16](#)). This theorem is deliberately local. It is a compact-window state-space exhaustion statement, not a terminal global profile theorem. Its proof uses the local compactness, closed-label, and residual-exhaustion lemmas in the terminal state-space discharge package; it does not use a whole-space profile decomposition or a heat-flow remainder estimate.

Lemma 9.18 (Terminal local velocity decoupling). *Let $\mathcal{C}$ be a terminal active class and S_n be the nonretained remainder field in the $\mathcal{C}$ -frame. Then for every compact cylinder $B_R \times I$,*

$$S_n \rightarrow 0 \quad \text{in } L_\tau^\infty L_y^3(B_R).$$

Proof. Assume otherwise. Then there exist $R < \infty$, a compact interval I , $\delta > 0$, a subsequence, and times $\tau_n \in I$ such that

$$\|S_n(\tau_n)\|_{L^3(B_R)} \geq \delta.$$

In physical variables, after the retained terminal class has been subtracted, the remainder carries at least δ critical L^3 -mass inside a ball of radius $O(\lambda_{\mathcal{C}}^n)$, centered at

$$x_{\mathcal{C}}^n + O(\lambda_{\mathcal{C}}^n),$$

at time

$$t_n + (\lambda_{\mathcal{C}}^n)^2 \tau_n.$$

By terminal windowed state-space completeness, this nonvanishing mass has a first local state-space label. If the label is an adjacent alternative, the branch has left the retained terminal class and is handled by the corresponding local closure theorem, contradicting that it remains in S_n on the retained branch. If the label is the retained comparable class, then it was already grouped into $\mathcal{C}$ before the terminal frame was selected, so it cannot remain in S_n . This is again a contradiction.

Thus no such δ exists. The same argument covers same-point larger-scale and separated-point components: any nonzero mass seen in the terminal frame is either a named adjacent local state or part of the retained comparable class, contradicting the terminal partition. $\square$

Lemma 9.19 (Uniform terminal H^1 bounds). *Assume the repaired-gauge representation and the local Caccioppoli estimates hold on every compact terminal cylinder $B_{2R} \times I$. Then*

$$\sup_n \|V_n\|_{L_\tau^\infty L_y^3(B_{2R})} + \sup_n \|V_n\|_{L_\tau^2 H_y^1(B_{2R})} \leq C_R,$$

$$\sup_n \|U_n\|_{L_\tau^\infty L_y^3(B_{2R})} + \sup_n \|U_n\|_{L_\tau^2 H_y^1(B_{2R})} \leq C_R,$$

and, since $S_n = V_n - U_n$,

$$\sup_n \|S_n\|_{L_\tau^\infty L_y^3(B_{2R})} + \sup_n \|S_n\|_{L_\tau^2 H_y^1(B_{2R})} \leq C_R.$$

In particular, by Sobolev embedding on B_{2R} ,

$$U_n, S_n \quad \text{are bounded in} \quad L_\tau^2 L_y^6(B_{2R}).$$

Proof. The first estimate is the repaired-gauge Caccioppoli estimate for V_n on the compact terminal cylinder. The retained solution U_n is itself a repaired-gauge Navier–Stokes solution with positive finite critical mass and bounded modulation on the same compact cylinder, so the same estimate applies to U_n . Since $S_n = V_n - U_n$, the displayed bound for S_n follows by the triangle inequality in the stated norms. The $L_\tau^2 L_y^6$ bound follows from the Sobolev embedding $H^1(B_{2R}) \hookrightarrow L^6(B_{2R})$. $\square$

Lemma 9.20 (Nonlinear stress decoupling in terminal frames). *Let*

$$F_n := U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n.$$

Then, for every compact $B_R \times I$,

$$F_n \rightarrow 0 \quad \text{in} \quad L_\tau^1 L_y^2(B_R).$$

Proof. By lemma 9.18,

$$\|S_n\|_{L_\tau^\infty L_y^3(B_R)} \rightarrow 0.$$

Using Hölder's inequality in space and time,

$$\|U_n \otimes S_n\|_{L_\tau^1 L_y^2(B_R)} \leq \|S_n\|_{L_\tau^\infty L_y^3(B_R)} \|U_n\|_{L_\tau^1 L_y^6(B_R)}$$

and

$$\|U_n\|_{L_\tau^1 L_y^6(B_R)} \leq |I|^{1/2} \|U_n\|_{L_\tau^2 L_y^6(B_R)} \leq C_{R,I}.$$

Thus $U_n \otimes S_n \rightarrow 0$ in $L_\tau^1 L_y^2(B_R)$, and the same estimate gives $S_n \otimes U_n \rightarrow 0$. Finally,

$$\|S_n \otimes S_n\|_{L_\tau^1 L_y^2(B_R)} \leq \|S_n\|_{L_\tau^\infty L_y^3(B_R)} \|S_n\|_{L_\tau^1 L_y^6(B_R)} \rightarrow 0,$$

because $\|S_n\|_{L_\tau^1 L_y^6(B_R)} \leq C_{R,I}$. Summing the three terms proves the local convergence. $\square$

Lemma 9.21 (Quantitative nested-ball harmonic tail). *Fix radii $0 < r < R < R_1 < R_2$ and $1 < p < \infty$. Let $\theta \in C_c^\infty(B_{R_2} \setminus \overline{B_{R_1}})$, extend θF by zero to $\mathbb{R}^3$, and set*

$$P^{\text{ann}} = \mathcal{R}_i \mathcal{R}_j (\theta F_{ij}), \quad H^{\text{ann}} = P^{\text{ann}} - (P^{\text{ann}})_{B_{R_1}} \quad \text{on } B_{R_1}.$$

Then H^{ann} is harmonic and mean-zero on B_{R_1} . Moreover there is a constant $C_p(r, R, R_1, R_2, \theta) > 0$, depending only on p , the nested balls and the cutoff, such that

(i) Interior derivative bound. *For every harmonic mean-zero H on B_{R_1} and every $R < R_1$,*

$$\sup_{B_R} (|H| + R_1 |\nabla H|) \leq C_p(R, R_1) \|H\|_{L^p(B_{R_1})}.$$

(ii) Annular harmonic-source bound. *The harmonic part generated by the source in the fixed annulus satisfies*

$$\|H^{\text{ann}}\|_{L^p(B_R)} \leq C_p(r, R, R_1, R_2, \theta) \|F\|_{L^p(B_{R_2} \setminus B_{R_1})}.$$

(iii) Quantitative converse for annular tails. *If $\|H^{\text{ann}}\|_{L^p(B_r)} \geq \eta > 0$, then*

$$\|F\|_{L^p(B_{R_2} \setminus B_{R_1})} \geq C_p(r, R, R_1, R_2, \theta)^{-1} \eta.$$

This lemma controls only the annular harmonic contribution generated by sources in the displayed fixed annulus. Harmonic pressure coming from outside B_{R_2} , or from a failure of the pressure reconstruction itself, is not silently bounded by this lemma and is instead read as a pressure/exterior state-space alternative.

Proof. Item (i) is the standard interior derivative estimate for harmonic functions: by the mean-value property and Cauchy estimates on $B_{(R+R_1)/2}$, $|H| + R_1 |\nabla H|$ on B_R is controlled by the L^p mass of H on B_{R_1} with constant depending only on p and R/R_1 .

For item (ii), the source θF is supported outside B_{R_1} , so $-\Delta P^{\text{ann}} = 0$ on B_{R_1} . Hence H^{ann} is harmonic there. Calderón–Zygmund boundedness on $\mathbb{R}^3$ gives

$$\|P^{\text{ann}}\|_{L^p(\mathbb{R}^3)} \leq C_p \|\theta F\|_{L^p(\mathbb{R}^3)} \leq C_p(\theta) \|F\|_{L^p(B_{R_2} \setminus B_{R_1})},$$

and subtracting the mean on B_{R_1} changes the $L^p(B_R)$ norm by a geometric constant.

For item (iii), this is the contrapositive of item (ii): if the annular source were arbitrarily small, then item (ii) would force $\|H^{\text{ann}}\|_{L^p(B_R)}$, and hence $\|H^{\text{ann}}\|_{L^p(B_r)}$, to be arbitrarily small. $\square$

Lemma 9.22 (Harmonic cross pressure routes to pressure/exterior alternative). *Fix nested balls $B_R \Subset B_{R_1} \Subset B_{R_2}$, $1 < p < \infty$, and $1 \leq q < \infty$. Let*

$$P_{\text{cross},n} = P_{\text{cross},n}^{\text{loc}} + H_{n,R_1} + c_{n,R_1}(\tau)$$

be the local pressure reconstruction on B_{R_1} , where

$$-\Delta P_{\text{cross},n}^{\text{loc}} = \partial_i \partial_j (\zeta (F_n)_{ij}), \quad \zeta \equiv 1 \text{ on } B_{R_1}, \quad \text{supp } \zeta \Subset B_{R_2},$$

and H_{n,R_1} is harmonic with zero spatial mean on B_{R_1} . Suppose

$$F_n \rightarrow 0 \quad \text{in } L_\tau^q L_y^p(B_{R_2}).$$

If H_{n,R_1} does not converge to zero in $L_\tau^q L_y^p(B_R)$, then, after passing to a subsequence, the branch enters the ordered pressure/exterior-leakage alternative. Consequently, on any retained terminal branch, $H_{n,R_1} \rightarrow 0$ in $L_\tau^q L_y^p(B_R)$.

Proof. The local part tends to zero by Calderón–Zygmund estimates:

$$\|P_{\text{cross},n}^{\text{loc}}\|_{L_\tau^q L_y^p(B_{R_1})} \leq C \|F_n\|_{L_\tau^q L_y^p(B_{R_2})} \rightarrow 0.$$

Hence any non-vanishing pressure oscillation modulo functions of time is carried by the normalized harmonic part. If

$$\limsup_{n \rightarrow \infty} \|H_{n,R_1}\|_{L_\tau^q L_y^p(B_R)} > 0,$$

then choose a finite smooth partition of unity $\{\theta_\ell\}_\ell \subset C_c^\infty(B_{R_2} \setminus \overline{B_{R_1}})$ for the recorded annular part of the pressure atlas. For each ℓ , the corresponding annular harmonic contribution

$$H_{n,\ell}^{\text{ann}} := \mathcal{R}_i \mathcal{R}_j (\theta_\ell (F_n)_{ij}) - \oint_{B_{R_1}} \mathcal{R}_i \mathcal{R}_j (\theta_\ell (F_n)_{ij}) dy$$

vanishes in $L_\tau^q L_y^p(B_R)$ by [lemma 9.21](#), because

$$\|F_n\|_{L_\tau^q L_y^p(B_{R_2} \setminus B_{R_1})} \rightarrow 0.$$

Summing over the finite partition shows that the whole harmonic contribution generated by sources recorded in the fixed annulus $B_{R_2} \setminus B_{R_1}$ vanishes.

By construction of the local pressure atlas, after the local Calderón–Zygmund piece and the recorded annular pieces have been removed, the remaining mean-zero harmonic part has only two possible origins: either the pressure gauge/reconstruction fails on the selected window, or it is a far-field pressure tail generated outside the fixed ball B_{R_2} . The former is one of the ordered pressure rows. In the latter case, [theorem 9.39](#) routes any exterior carrier to an exterior or separated-core concentration alternative; if no such exterior carrier exists, [lemma 9.45](#) gives convergence of the far-field pressure modulo functions of time on B_R , contradicting the positive lower bound for H_{n,R_1} . Therefore every non-vanishing harmonic cross pressure is a named pressure/exterior-leakage alternative. A retained terminal branch is the complementary case in the ordered validation, so the harmonic part vanishes there. $\square$

Lemma 9.23 (Terminal pressure decoupling). *Fix nested balls $B_R \Subset B_{R_1} \Subset B_{R_2}$ and let $\zeta \in C_c^\infty(B_{R_2})$ be identically one on B_{R_1} . Let*

$$P_{\text{cross},n}^{\text{loc}} := \mathcal{R}_i \mathcal{R}_j (\zeta (F_n)_{ij})$$

be the local Calderón–Zygmund pressure generated by the cross stress. Suppose the terminal pressure reconstruction is used on B_{R_2} , so that on B_{R_1}

$$P_{\text{cross},n} = P_{\text{cross},n}^{\text{loc}} + H_{n,R_1} + c_{n,R_1}(\tau),$$

where $H_{n,R_1}(\cdot, \tau)$ is harmonic and normalized by $\oint_{B_{R_1}} H_{n,R_1}(\cdot, \tau) dy = 0$. On a retained terminal branch,

$$P_{\text{cross},n} - c_{n,R} \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_R),$$

and

$$\nabla P_{\text{cross},n} \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

Proof. [Lemma 9.20](#) gives

$$F_n \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_{R_2})$$

on every compact terminal cylinder. Calderón–Zygmund boundedness gives

$$\|P_{\text{cross},n}^{\text{loc}}\|_{L_\tau^1 L_y^2(B_{R_1})} \leq C_{R_1,R_2} \|F_n\|_{L_\tau^1 L_y^2(B_{R_2})} \rightarrow 0.$$

For the harmonic part, apply [lemma 9.22](#) with $(q, p) = (1, 2)$. A non-vanishing harmonic cross pressure is a pressure/exterior-leakage alternative and therefore is not present on a retained terminal branch. Hence $H_{n,R_1} \rightarrow 0$ in $L_\tau^1 L_y^2(B_R)$. Adding the local part proves $P_{\text{cross},n} - c_{n,R} \rightarrow 0$ in $L_\tau^1 L_y^2(B_R)$.

Finally, for $\varphi \in H_0^1(B_R; \mathbb{R}^3)$,

$$|\langle \nabla(P_{\text{cross},n} - c_{n,R}), \varphi \rangle| \leq \|P_{\text{cross},n} - c_{n,R}\|_{L^2(B_R)} \|\nabla \varphi\|_{L^2(B_R)}.$$

Integrating in time gives $\nabla P_{\text{cross},n} \rightarrow 0$ in $L_\tau^1 H_y^{-1}(B_R)$. $\square$

Lemma 9.24 (Linear residual decoupling). *Let $V_n = U_n + S_n$ solve*

$$\partial_\tau V_n + (V_n \cdot \nabla) V_n + \nabla P_n = \nu \Delta V_n + a_n(V_n + y \cdot \nabla V_n) + b_n \cdot \nabla V_n.$$

Suppose U_n satisfies, in the same repaired gauge,

$$\partial_\tau U_n + (U_n \cdot \nabla) U_n + \nabla P_{U,n} = \nu \Delta U_n + a_n(U_n + y \cdot \nabla U_n) + b_n \cdot \nabla U_n + e_n,$$

with

$$e_n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R)$$

on compact terminal cylinders. Then

$$\partial_\tau S_n - \nu \Delta S_n - a_n(S_n + y \cdot \nabla S_n) - b_n \cdot \nabla S_n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

Equivalently, the same conclusion holds with the opposite sign convention for the modulation terms.

Proof. Subtracting the U_n -equation from the V_n -equation gives

$$\partial_\tau S_n - \nu \Delta S_n = -\operatorname{div} F_n - \nabla P_{\text{cross},n} - e_n + a_n(S_n + y \cdot \nabla S_n) + b_n \cdot \nabla S_n.$$

Equivalently,

$$\partial_\tau S_n - \nu \Delta S_n - a_n(S_n + y \cdot \nabla S_n) - b_n \cdot \nabla S_n = -\operatorname{div} F_n - \nabla P_{\text{cross},n} - e_n.$$

By [lemma 9.20](#),

$$F_n \rightarrow 0 \quad \text{in } L_\tau^1 L_y^2(B_R),$$

hence

$$\operatorname{div} F_n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

By [lemma 9.23](#),

$$\nabla P_{\text{cross},n} \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

Together with $e_n \rightarrow 0$, this proves the displayed residual convergence.

For the opposite sign convention, the two residuals differ by

$$2a_n(S_n + y \cdot \nabla S_n) + 2b_n \cdot \nabla S_n.$$

By [lemma 9.18](#) and finite measure,

$$S_n \rightarrow 0 \quad \text{in } L_\tau^\infty L_y^2(B_R).$$

For $\varphi \in H_0^1(B_R)$,

$$|\langle y \cdot \nabla S_n, \varphi \rangle| = \left| \int_{B_R} S_n (3\varphi + y \cdot \nabla \varphi) dy \right| \leq C_R \|S_n\|_{L^2(B_R)} \|\varphi\|_{H^1(B_R)},$$

and

$$|\langle b_n \cdot \nabla S_n, \varphi \rangle| \leq |b_n| \|S_n\|_{L^2(B_R)} \|\nabla \varphi\|_{L^2(B_R)}.$$

The term S_n itself is controlled by the same L^2 -convergence. Since a_n and b_n are bounded, the sign-convention difference tends to zero in $L_\tau^1 H_y^{-1}(B_R)$. $\square$

Lemma 9.25 (Terminal modulation compatibility). *Same-point larger-scale profiles in a terminal frame have the form*

$$W_j^n(y, \tau) = \mu_j^n U^j(\theta_j^n(\tau), \xi_j^n + \mu_j^n y), \quad \mu_j^n \rightarrow 0.$$

When the profile is regular at the observed offset, its leading term is a spatially constant ambient velocity. Absorbing this constant into the translation coefficient replaces b_n by $b_n - b_{\text{amb},n}$, and the remaining terms vanish on compact terminal cylinders. If the offset escapes, the entire profile is locally invisible and no absorption is needed. In particular,

$$\sup_n \sup_{\tau \in I} (|a_n(\tau)| + |b_n(\tau)|) < \infty.$$

Proof. The local invisibility of the nonconstant remainder follows from [lemma 9.18](#) and the H^{-1} product estimates in [lemma 9.24](#). Separated physical-point profiles are locally invisible by [lemma 9.7](#); their pressures and cutoff commutators vanish by [lemma 9.23](#) and [lemma 9.24](#). Therefore only the spatially constant ambient velocity changes the translation coefficient, and after this absorption the effective modulation coefficients remain uniformly bounded. $\square$

Lemma 9.26 (Terminal compact scale-collapse admissibility). *The retained terminal solution U_n has positive finite critical mass, pressure reconstruction, bounded modulation, local windowed H^1 control, critical tightness on the selected windows, and the compactness input used by the compact scale-collapse argument.*

Proof. The terminal class is active, so after the perturbative residual and nonretained local states have been removed,

$$\inf_{\tau \in I} \|U_n(\tau)\|_{L^3(B_R)} \geq c_{\text{act}} > 0$$

on the selected compact windows, after passing to the subsequences used in the local state-space selection. The upper bound follows from the retained local compact-cylinder bound:

$$\sup_{\tau \in I} \|U_n(\tau)\|_{L^3(B_R)} \leq C(M).$$

The terminal state-space completeness also gives the required tightness. If U_n were not uniformly L^3 -tight on the selected windows, then there would be $\varepsilon > 0$, radii $R_k \rightarrow \infty$, and times τ_n such that

$$\int_{|y| > R_k} |U_n(y, \tau_n)|^3 dy \geq \varepsilon.$$

Applying the local state-space stratification to this escaping sequence produces either an exterior-regular component, which is removed by the exterior regularity argument below, or an additional active local component at a different physical point or larger same-point scale. Such a component belongs to S_n and is handled by terminal decoupling. Hence no non-tight retained terminal solution remains.

The repaired-gauge representation gives pressure reconstruction and bounded modulation. The local Caccioppoli estimate gives local windowed H^1 control, and the Aubin–Lions compactness layer gives the compactness input on the selected windows. $\square$

Theorem 9.27 (Terminal nonlinear local-state decoupling theorem). *Assume local compact critical bounds, terminal windowed state-space completeness, local perturbative residual stability, the exterior annulus alternative, repaired-gauge representation, and local Caccioppoli regularity. Then terminal nonlinear local-state decoupling holds in the sense of [definition 9.13](#).*

Proof. Let $\mathcal{C}$ be any terminal active compound class, let U_n be the retained repaired-gauge nonlinear solution in its terminal frame, and set

$$S_n = V_n - U_n.$$

Lemma 9.18 proves

$$S_n \rightarrow 0 \quad \text{in } L^\infty_\tau L^3_y(B_{2R})$$

on every compact terminal cylinder. Divergence-free compatibility follows from the divergence-free repaired-gauge equation and the local terminal decomposition.

Lemma 9.20 proves

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0 \quad \text{in } L^1_\tau L^2_y(B_R).$$

The local compact critical bounds give the $L^\infty_\tau L^3_y$ control required in the pressure lemma below on the cylinders under consideration. Lemma 9.23 proves

$$\nabla P_{\text{cross},n} \rightarrow 0 \quad \text{in } L^1_\tau H^{-1}_y(B_R)$$

after subtracting spatially constant pressure terms. Lemma 9.24 proves

$$\partial_\tau S_n - \nu \Delta S_n + a_n(S_n + y \cdot \nabla S_n) + b_n \cdot \nabla S_n \rightarrow 0 \quad \text{in } L^1_\tau H^{-1}_y(B_R).$$

Lemma 9.25 gives modulation compatibility, and lemma 9.26 gives compact scale-collapse admissibility of the retained terminal solution. These are the items in definition 9.13. $\square$

Corollary 9.28 (Multibubble exclusion from terminal local decoupling). *Under the Type II local terminal data, terminal windowed state-space completeness, local perturbative residual stability, the exterior annulus alternative, repaired-gauge representation, and local Caccioppoli regularity, no active multibubble Type II configuration occurs.*

Proof. Theorem 9.27 gives terminal nonlinear decoupling. Theorem 9.14 then supplies the same-point and separated-point decoupling hypotheses used in theorem 9.9. The active local mass floor follows from the retained positive local critical-mass row and local perturbative residual stability. Applying theorem 9.9 gives the claimed exclusion of active multibubble configurations. $\square$

9.5. Small local residual components. For a compact terminal cylinder $Q_{R,I} := B_R \times I$, define the local residual seminorm

$$\|W\|_{\mathcal{X}(Q_{R,I})} := \|W\|_{L^\infty(I; L^3(B_R))} + \|W\|_{L^2(I; H^1(B_R))}.$$

This is only a compact-window norm. It is not a whole-space heat-flow or global mild-solution norm.

Lemma 9.29 (Local residual product estimates). *Let $Q_{2R,I} = B_{2R} \times I$ be compact. Suppose*

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_{2R}))$$

and

$$\sup_n (\|U_n\|_{\mathcal{X}(Q_{2R,I})} + \|S_n\|_{\mathcal{X}(Q_{2R,I})}) < \infty.$$

Then

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0 \quad \text{in } L^1(I; L^2(B_R)).$$

If the pressures associated with these cross stresses are reconstructed by the canonical local pressure decomposition, write

$$P_n^{\text{cross}} = P_{n,\text{loc}}^{\text{cross}} + H_n^{\text{cross}}.$$

The local components have gradients converging to zero in

$$L^1(I; H^{-1}(B_R)).$$

The remainders H_n^{cross} are harmonic on B_R ; a nonvanishing harmonic remainder is the named pressure/exterior row, and in its absence the full cross-pressure gradients converge to zero in the same space.

Proof. Sobolev embedding on B_{2R} gives uniform

$$U_n, S_n \text{ in } L^2(I; L^6(B_{2R})).$$

Hölder's inequality in space and time yields

$$\|U_n \otimes S_n\|_{L^1(I; L^2(B_R))} \leq \|S_n\|_{L^\infty(I; L^3(B_R))} \|U_n\|_{L^1(I; L^6(B_R))} \rightarrow 0,$$

because

$$\|U_n\|_{L^1(I; L^6(B_R))} \leq |I|^{1/2} \|U_n\|_{L^2(I; L^6(B_R))}.$$

The same estimate treats $S_n \otimes U_n$, and

$$\|S_n \otimes S_n\|_{L^1(I; L^2(B_R))} \leq \|S_n\|_{L^\infty(I; L^3(B_R))} \|S_n\|_{L^1(I; L^6(B_R))} \rightarrow 0.$$

For the pressure statement, apply the local pressure reconstruction to the cross source. The preceding $L^1 L^2$ convergence gives local convergence of the local pressure part by Calderón–Zygmund estimates. Any non-vanishing harmonic remainder is routed by [lemma 9.22](#) with $(q, p) = (1, 2)$ to the pressure/exterior alternative; on a retained branch this alternative is absent. Therefore the cross-pressure gradients vanish in $L^1 H^{-1}(B_R)$. $\square$

Lemma 9.30 (Spacetime-small terminal residuals are perturbative). *Let*

$$V_n = U_n + S_n$$

be an ancestry-generated terminal decomposition on a compact cylinder $Q_{2R, I} = B_{2R} \times I$. Suppose that the repaired-gauge representation, bounded modulation, local pressure reconstruction, and the compact-window estimates give

$$\sup_n \left(\|U_n\|_{L^2(I; H^1(B_{2R}))} + \|S_n\|_{L^2(I; H^1(B_{2R}))} \right) < \infty.$$

The ordered represented-equation validation has two outputs. Its first failure is the corresponding representation, equation, or componentwise-compactness row. On the retained output, U_n satisfies its retained equation in the gauge of V_n , with equation error

$$e_n \rightarrow 0 \text{ in } L^1(I; H^{-1}(B_R)).$$

If

$$S_n \rightarrow 0 \text{ in } L^3(Q_{2R, I}),$$

then

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0 \text{ in } L^1(I; L^2(B_R)).$$

Moreover, the ordered pressure reconstruction gives either the named pressure/exterior row or

$$\nabla P_{\text{cross}, n} \rightarrow 0 \text{ in } L^1(I; H^{-1}(B_R)).$$

On the retained branch the latter alternative holds, and

$$\partial_\tau S_n - \nu \Delta S_n - a_n(S_n + y \cdot \nabla S_n) - b_n \cdot \nabla S_n \rightarrow 0 \text{ in } L^1(I; H^{-1}(B_R)).$$

All compact-cutoff errors containing S_n also tend to zero in $L^1(I; H^{-1}(B_R))$. Thus S_n is perturbative in every terminal PDE topology even when

$$\limsup_{n \rightarrow \infty} \|S_n\|_{L^\infty(I; L^3(B_R))} > 0.$$

Proof. Since I has finite measure,

$$\|S_n\|_{L^2(I; L^3(B_{2R}))} \leq |I|^{1/6} \|S_n\|_{L^3(Q_{2R, I})} \rightarrow 0.$$

The uniform $L^2 H^1$ bounds and Sobolev embedding give

$$\sup_n \left(\|U_n\|_{L^2(I; L^6(B_{2R}))} + \|S_n\|_{L^2(I; L^6(B_{2R}))} \right) < \infty.$$

Consequently,

$$\|U_n \otimes S_n\|_{L^1(I; L^2(B_R))} \leq \|U_n\|_{L^2(I; L^6(B_R))} \|S_n\|_{L^2(I; L^3(B_R))} \rightarrow 0.$$

The same estimate applies to $S_n \otimes U_n$, and

$$\|S_n \otimes S_n\|_{L^1(I; L^2(B_R))} \leq \|S_n\|_{L^2(I; L^6(B_R))} \|S_n\|_{L^2(I; L^3(B_R))} \rightarrow 0.$$

Apply the canonical local pressure decomposition to this cross stress. Its Calderón–Zygmund component tends to zero in $L^1(I; L^2(B_R))$. By [lemma 9.22](#) with $(q, p) = (1, 2)$, a nonvanishing normalized harmonic component is the pressure/exterior row. On the retained branch that row is absent, so the full normalized cross pressure tends to zero in $L^1(I; L^2(B_R))$, and its gradient tends to zero in $L^1(I; H^{-1}(B_R))$.

Subtract the represented equation for U_n from the represented equation for V_n . On $Q_{R,I}$,

$$\begin{aligned} \partial_\tau S_n - \nu \Delta S_n - a_n(S_n + y \cdot \nabla S_n) - b_n \cdot \nabla S_n \\ = -\nabla \cdot (U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n) - \nabla P_{\text{cross},n} - e_n. \end{aligned}$$

Every term on the right tends to zero in $L^1(I; H^{-1}(B_R))$.

It remains to verify the cutoff topology. Spacetime L^3 -convergence on a bounded cylinder gives $S_n \rightarrow 0$ in $L^1(I; L^2(B_{2R}))$. When a derivative in a linear cutoff term falls on S_n , integration by parts moves it onto the fixed cutoff and the H_0^1 test function. For example,

$$\langle \nabla \chi \cdot \nabla S_n, \varphi \rangle = - \int S_n \cdot ((\Delta \chi) \varphi + \nabla \chi \cdot \nabla \varphi).$$

The resulting pointwise dual estimate gives

$$\|\text{linear cutoff commutator}\|_{L^1(I; H^{-1}(B_R))} \leq C_\chi \|S_n\|_{L^1(I; L^2(B_{2R}))} \rightarrow 0.$$

The bounded modulation terms are handled by the same integration by parts. Nonlinear cutoff terms are bounded by the cross stress already shown to vanish, and pressure cutoff terms are estimated using the normalized cross pressure above. This proves the cutoff assertion and the lemma. $\square$

Theorem 9.31 (Small local residual components are perturbative). *Let S_n be a nonretained terminal residual on every compact terminal cylinder $Q_{2R,I}$. Assume*

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_{2R}))$$

and that the local compact-window bounds, local pressure reconstruction, and bounded modulation estimates hold on $Q_{2R,I}$. Then S_n is perturbative in the terminal decoupling topology of [definition 9.13](#).

Proof. [Lemma 9.29](#) gives the vanishing nonlinear-stress and cross-pressure terms. The repaired-gauge equation for $V_n = U_n + S_n$ and the equation for the retained component U_n , subtracted on $Q_{R,I}$, give the linear residual identity of [lemma 9.24](#). The stress and pressure terms just proved vanish, and bounded modulation sends the remaining $a_n(S_n + y \cdot \nabla S_n) + b_n \cdot \nabla S_n$ terms to zero in $L^1(I; H^{-1}(B_R))$. Hence all perturbative clauses in [definition 9.13](#) hold on $Q_{R,I}$. $\square$

Corollary 9.32 (Small local components are not active). *If a nonretained component has vanishing local $L^\infty L^3$ seminorm on every compact terminal cylinder and satisfies the local compact-window bounds, then it cannot define an active terminal state.*

Proof. By [theorem 9.31](#), such a component contributes only vanishing stress, pressure, modulation, and residual errors in the terminal decoupling topology. Active terminal states are precisely retained local state-space labels with positive nonperturbative local mass. Therefore a component that is perturbative on every compact terminal cylinder is not active. $\square$

Theorem 9.33 (Positive local mass floor for active terminal components). *There exists a local threshold $\varepsilon_{\text{loc}} > 0$, depending only on the Caffarelli–Kohn–Nirenberg threshold and the finite covering constants of the selected compact cylinders, such that every retained active terminal component has, after passing to a subsequence, a compact terminal cylinder $Q_{R,I}$ with*

$$\liminf_{n \rightarrow \infty} \int_I \int_{B_R} |W_n|^3 dy d\tau \geq \varepsilon_{\text{loc}}^3$$

and

$$\limsup_{n \rightarrow \infty} \|W_n\|_{L^\infty(I; L^3(B_R))} \geq \varepsilon_{\text{loc}}.$$

Proof. Let ε_{CKN} be the threshold in [theorem 9.35](#). Fix a finite nested covering of each retained compact terminal cylinder by CKN cylinders whose smaller cylinders still cover the retained core. Choose $\eta_{\text{det}} > 0$ to be ε_{CKN} divided by a sufficiently large constant depending only on this cover.

If every CKN cylinder in the cover had velocity-pressure quantity below η_{det} , then the CKN criterion would give a uniformly smooth bounded representative on the retained smaller cylinder. Such a bounded selected-window representative is not an active Type II terminal component; it is routed to the regular/bounded-window alternative. Hence an active retained component has a subcylinder on which the terminal CKN quantity is at least η_{det} .

If the pressure part of that CKN quantity carries at least half of η_{det} while the velocity part is arbitrarily small on all smaller subcylinders, then the local pressure reconstruction and harmonic-tail alternative, made explicit in [lemma 9.22](#) with $(q, p) = (3/2, 3/2)$, assign the branch to the pressure/exterior-leakage stratum. That stratum is adjacent and is not a retained active terminal component. Therefore on a retained active component the velocity part carries a fixed fraction of the detected CKN quantity on some smaller cylinder:

$$\rho^{-2} \int_{I_\rho} \int_{B_\rho} |W_n|^3 dy d\tau \geq c \eta_{\text{det}}$$

along a subsequence. Since $|I_\rho| \simeq \rho^2$, there is a time in I_ρ with

$$\int_{B_\rho} |W_n(y, \tau)|^3 dy \geq c \eta_{\text{det}}.$$

After rescaling this subcylinder to a fixed compact terminal cylinder, this is the asserted spacetime lower bound and the asserted $L^\infty_\tau L^3_y$ lower bound with $\varepsilon_{\text{loc}} := (c \eta_{\text{det}})^{1/3}$.

If no such velocity lower bound existed on any compact terminal cylinder, then the component would have vanishing local $L^\infty L^3$ seminorm after subsequence extraction and would be perturbative by [corollary 9.32](#), contradicting retained activity. $\square$

9.6. Exterior regularity.

Definition 9.34 (Exterior annuli and terminal CKN quantity). Let (x_*, T^*) be the selected Type II core point. For $0 < r < R < \infty$, set

$$A_{r,R} := \overline{B_R(x_*)} \setminus B_r(x_*).$$

For $x_0 \in \mathbb{R}^3$ and $\rho > 0$, write

$$Q_\rho^*(x_0) := B_\rho(x_0) \times (T^* - \rho^2, T^*).$$

For a suitable weak solution (u, p) , define the terminal local velocity-pressure quantity

$$\mathcal{C}_*(u, p; x_0, \rho) := \rho^{-2} \int_{T^* - \rho^2}^{T^*} \int_{B_\rho(x_0)} |u|^3 dx dt + \rho^{-2} \int_{T^* - \rho^2}^{T^*} \int_{B_\rho(x_0)} |p - (p)_{B_\rho(x_0)}(t)|^{3/2} dx dt.$$

The pressure is measured modulo functions of time.

Theorem 9.35 (Terminal Caffarelli–Kohn–Nirenberg regularity). *There exists $\varepsilon_{\text{CKN}} > 0$ with the following property. Let (u, p) be suitable on $Q_\rho^*(x_0)$. If*

$$\mathcal{C}_*(u, p; x_0, \rho) \leq \varepsilon_{\text{CKN}},$$

then for each integer $k \geq 0$ there are $0 < c_k < 1$ and $C_k < \infty$, independent of $(u, p), x_0, \rho$, such that u is smooth in

$$B_{c_k \rho}(x_0) \times (T^* - c_k \rho^2, T^*)$$

and, after subtracting a function of time from p ,

$$\sup_{t \in (T^* - c_k \rho^2, T^*)} \left(\rho^{k+1} \|\nabla^k u(t)\|_{L^\infty(B_{c_k \rho}(x_0))} + \rho^{k+2} \|\nabla^k p(t)\|_{L^\infty(B_{c_k \rho}(x_0))} \right) \leq C_k.$$

Proof. This is the standard endpoint form of the Caffarelli–Kohn–Nirenberg epsilon-regularity criterion [CKN82, Lin98], after translating the terminal time to 0 and rescaling $Q_\rho^*(x_0)$ to a unit parabolic cylinder. The pressure normalization by spatial averages is the usual normalization in the suitable weak-solution criterion. Interior parabolic bootstrapping gives the stated higher derivative bounds on smaller cylinders. $\square$

Throughout the exterior regularity subsection we fix

$$\varepsilon_{\text{ext}} := \frac{1}{2} \varepsilon_{\text{CKN}}.$$

Definition 9.36 (No exterior concentration on a fixed annulus). Let $0 < r < R < \infty$. A suitable Type II branch has no exterior concentration on $A_{r,R}$ if there exists $\rho_* > 0$ such that whenever

$$x_0 \in A_{r,R}, \quad 0 < \rho \leq \rho_*, \quad Q_\rho^*(x_0) \subset (B_{2R}(x_*) \setminus \overline{B_{r/2}(x_*)}) \times (0, T^*),$$

one has

$$\mathcal{C}_*(u, p; x_0, \rho) \leq \varepsilon_{\text{ext}}.$$

Definition 9.37 (Exterior concentration branch). An exterior concentration branch on $A_{r,R}$ is a sequence

$$x_n \in A_{r,R}, \quad \rho_n \downarrow 0,$$

such that

$$Q_{\rho_n}^*(x_n) \subset (B_{2R}(x_*) \setminus \overline{B_{r/2}(x_*)}) \times (0, T^*)$$

and

$$\mathcal{C}_*(u, p; x_n, \rho_n) > \varepsilon_{\text{ext}}.$$

The associated exterior rescalings are

$$U_n(y, s) := \rho_n u(x_n + \rho_n y, T^* + \rho_n^2 s), \quad P_n(y, s) := \rho_n^2 p(x_n + \rho_n y, T^* + \rho_n^2 s),$$

for $s < 0$.

Lemma 9.38 (Exterior test dichotomy). *Fix $0 < r < R < \infty$. Then either the branch has no exterior concentration on $A_{r,R}$, or it admits an exterior concentration branch on $A_{r,R}$.*

Proof. If the no-concentration condition fails, then for every m there are

$$x_m \in A_{r,R}, \quad 0 < \rho_m \leq m^{-1},$$

with $Q_{\rho_m}^*(x_m)$ contained in the enlarged annular cylinder and

$$\mathcal{C}_*(u, p; x_m, \rho_m) > \varepsilon_{\text{ext}}.$$

After discarding repetitions and passing to a subsequence, $\rho_m \rightarrow 0$. This is precisely an exterior concentration branch. The converse is immediate from the definitions. $\square$

Theorem 9.39 (Exterior concentration stratification). *Fix $0 < r < R < \infty$. A suitable Type II branch satisfies exactly one of the following alternatives on $A_{r,R}$:*

- (i) *it has no exterior concentration on $A_{r,R}$;*
- (ii) *it admits an exterior concentration branch on $A_{r,R}$.*

The second alternative is the complementary concentration branch. According to the relative position and scale of the exterior concentration sequence, it belongs to one of the exterior-profile, separated-core, multicore, noncompact exterior, or scale-collapse alternatives.

Proof. The dichotomy is lemma 9.38. If an exterior concentration branch exists, then the annular terminal cylinders carry a positive Caffarelli–Kohn–Nirenberg quantity at scales $\rho_n \downarrow 0$ and centers a positive physical distance from the selected core. Let λ_n denote the selected core scale along the same terminal sequence. After passing to a subsequence, the relative scale-center configuration falls into one of the standard local concentration-compactness alternatives. If the exterior frame is orthogonal to the selected core frame, it gives an exterior profile. If only finitely many separated active centers remain at comparable terminal scales, it gives a separated-core or multicore configuration. If compactness fails in the exterior frame, one obtains the noncompact exterior alternative. If the exterior frame develops a further unresolved scale in its own renormalized variables, one obtains the scale-collapse alternative. These alternatives exhaust the subsequential configurations at positive physical distance from the selected core. Thus the failure of the annular regularity test is represented by a local concentration alternative rather than by an additional regularity hypothesis. $\square$

Lemma 9.40 (Exterior CKN cover). *Assume there is no exterior concentration on $A_{r,R}$. Fix an integer $k \geq 0$. If*

$$A_{r',R'} \Subset A_{r,R},$$

then there are points $x_j \in A_{r,R}$, radii $\rho_j > 0$, and a terminal time length $\delta > 0$ such that

$$A_{r',R'} \times (T^* - \delta, T^*) \subset \bigcup_j B_{c_k \rho_j}(x_j) \times (T^* - c_k \rho_j^2, T^*)$$

and

$$\mathcal{C}_*(u, p; x_j, \rho_j) \leq \varepsilon_{\text{ext}}$$

for every j . Here c_k is the constant from theorem 9.35.

Proof. Since $A_{r',R'} \Subset A_{r,R}$, there is $d > 0$ such that

$$B_{4d}(x) \subset B_{2R}(x_*) \setminus \overline{B_{r/2}(x_*)} \quad \text{for all } x \in A_{r',R'}.$$

Let ρ_* be supplied by the no-concentration condition and set $\rho_0 := \min\{d, \rho_*\}$. By compactness, finitely many balls

$$B_{c_k \rho_0/2}(x_j), \quad x_j \in A_{r,R},$$

cover $A_{r',R'}$. The terminal cylinders $Q_{\rho_0}^*(x_j)$ are contained in the enlarged annular cylinder, so the no-concentration condition gives the displayed smallness. Taking $\delta = c_k \rho_0^2/2$ gives the spacetime cover. $\square$

Lemma 9.41 (Exterior pressure localization). *Let $A_{r',R'} \Subset A_{r,R}$, and assume that u is bounded on a slightly larger annular terminal cylinder. Then on*

$$A_{r',R'} \times (T^* - \delta, T^*)$$

the pressure has a decomposition

$$p = p_{\text{loc}} + p_{\text{harm}} + c(t),$$

where $c(t)$ is a function of time, p_{harm} is harmonic in space on the smaller annulus, and for every $q < \infty$ and every integer $k \geq 0$

$$\|p_{\text{loc}}\|_{L_t^\infty L_x^q(A_{r',R'})} + \|p_{\text{harm}}\|_{L_t^\infty C_x^k(A_{r',R'})} \leq C_{r',R',q,k}.$$

Proof. Choose $\eta \in C_c^\infty(B_{2R}(x_*) \setminus \overline{B_{r/2}(x_*)})$ with $\eta \equiv 1$ on a slightly larger annulus containing $A_{r',R'}$. Define p_{loc} by

$$-\Delta p_{\text{loc}} = \partial_i \partial_j (\eta u_i u_j)$$

using the Newtonian potential, with the usual spatial normalization. Since u is locally bounded on the support of η , Calderón–Zygmund estimates give $p_{\text{loc}} \in L_t^\infty L_x^q$ for every $q < \infty$ on the terminal interval. The difference

$$p - p_{\text{loc}}$$

is harmonic in space on the smaller annulus, modulo a function of time. The standard interior estimates for harmonic functions give the displayed C^k bounds after fixing the time-dependent additive constant. $\square$

Theorem 9.42 (Exterior regularity on no-concentration annuli). *Let u, p be a suitable weak solution on $\mathbb{R}^3 \times [0, T^*)$, and let (x_*, T^*) be the selected Type II core point. Fix $0 < r' < R' < \infty$. If there is no exterior concentration on some annulus $A_{r,R}$ with*

$$A_{r',R'} \Subset A_{r,R},$$

then for every integer $k \geq 0$ there are $\delta > 0$, a function $c(t)$, and $C_{r',R',k} < \infty$ such that

$$\sup_{t \in (T^* - \delta, T^*)} \left(\|u(t)\|_{C^k(A_{r',R'})} + \|p(t) - c(t)\|_{C^k(A_{r',R'})} \right) \leq C_{r',R',k}.$$

If the no-concentration hypothesis fails, the complementary exterior concentration alternative holds in the sense of [definition 9.37](#).

Proof. The no-concentration condition and [lemma 9.40](#) give a finite cover of $A_{r',R'} \times (T^* - \delta, T^*)$ by terminal cylinders on which the Caffarelli–Kohn–Nirenberg quantity is below ε_{CKN} . Applying [theorem 9.35](#) on each cylinder gives local C^k bounds for u and for p modulo functions of time on smaller cylinders. Since the cover is finite, taking the maximum over the cover gives the stated local bound.

The pressure normalizations on overlapping balls differ only by functions of time. Equivalently, one may use [lemma 9.41](#) after the $k = 0$ velocity bound has been obtained and then apply standard interior parabolic estimates on nested annuli. This gives a single time-dependent normalization $c(t)$ on $A_{r',R'}$. If the no-concentration hypothesis fails, the exterior test dichotomy gives an exterior concentration branch. $\square$

Lemma 9.43 (Divergence-free core and exterior decomposition). *Fix $0 < r_0 < R_0 < \infty$, and assume that there is no exterior concentration on an annulus containing*

$$\overline{B_{2r_0}(x_*)} \setminus B_{r_0}(x_*).$$

Let $\chi \in C_c^\infty(B_{2r_0}(x_))$ satisfy*

$$\chi \equiv 1 \quad \text{on } B_{r_0}(x_*).$$

Then there are divergence-free fields u_{core} and u_{ext} such that

$$u = u_{\text{core}} + u_{\text{ext}}, \quad u_{\text{core}} = u \text{ on } B_{r_0}(x_*), \quad u_{\text{ext}} = 0 \text{ on } B_{r_0}(x_*),$$

and all derivatives of u_{ext} are uniformly bounded on compact subannuli of the no-exterior-concentration region up to T^ .*

Proof. Set

$$g = \nabla \chi \cdot u.$$

The support of g lies in the annulus

$$A_{r_0} = B_{2r_0}(x_*) \setminus B_{r_0}(x_*),$$

where u is smooth by [theorem 9.42](#). Since

$$\int_{A_{r_0}} g \, dx = \int_{\mathbb{R}^3} \nabla \chi \cdot u \, dx = - \int_{\mathbb{R}^3} \chi \nabla \cdot u \, dx = 0,$$

the Bogovskii operator on A_{r_0} gives a vector field

$$w = \mathcal{B}_{A_{r_0}} g$$

supported in A_{r_0} , smooth up to T^* , and satisfying

$$\nabla \cdot w = g.$$

Define

$$u_{\text{core}} := \chi u - w, \quad u_{\text{ext}} := u - u_{\text{core}}.$$

Then

$$\nabla \cdot u_{\text{core}} = \nabla \chi \cdot u + \chi \nabla \cdot u - \nabla \cdot w = 0,$$

and therefore $\nabla \cdot u_{\text{ext}} = 0$. Since $\chi = 1$ and $w = 0$ on $B_{r_0}(x_*)$, $u_{\text{core}} = u$ and $u_{\text{ext}} = 0$ there. The uniform exterior bounds follow from [theorem 9.42](#) on compact subannuli and the standard boundedness of the Bogovskii construction on the fixed annulus. $\square$

Lemma 9.44 (Exterior fields vanish in compact core frames). *Let (x_n, λ_n) be a core frame with*

$$x_n \rightarrow x_*, \quad \lambda_n \rightarrow 0.$$

For a compact frame ball B_R , define

$$V_{\text{ext}}^n(y, \tau) := \lambda_n u_{\text{ext}}(x_n + \lambda_n y, t_n + \lambda_n^2 \tau).$$

Then, for every compact $B_R \times I$,

$$V_{\text{ext}}^n \equiv 0$$

on $B_R \times I$ for all sufficiently large n . Consequently

$$V_{\text{ext}}^n \rightarrow 0 \quad \text{in every local } L_\tau^q W_y^{k,q}(B_R) \text{ norm.}$$

Proof. Since $x_n \rightarrow x_*$ and $\lambda_n \rightarrow 0$, for fixed R ,

$$x_n + \lambda_n B_R \subset B_{r_0}(x_*)$$

for all sufficiently large n . On $B_{r_0}(x_*)$,

$$u_{\text{ext}} = 0.$$

Therefore $V_{\text{ext}}^n = 0$ on $B_R \times I$ for all large n , which implies convergence in every stated local norm. $\square$

Lemma 9.45 (Exterior pressure vanishes modulo constants). *Let p_{ext} be any pressure contribution generated by source terms supported in*

$$\{|x - x_*| \geq r_0\}.$$

In a core frame set

$$P_{\text{ext}}^n(y, \tau) := \lambda_n^2 p_{\text{ext}}(x_n + \lambda_n y, t_n + \lambda_n^2 \tau).$$

Then, for every compact $B_R \times I$, there are constants $c_{n,R}(\tau)$ such that

$$P_{\text{ext}}^n - c_{n,R} \rightarrow 0 \quad \text{in } L_\tau^\infty C_y^1(B_R),$$

and hence

$$\nabla_y P_{\text{ext}}^n \rightarrow 0 \quad \text{in } L_\tau^1 H_y^{-1}(B_R).$$

Proof. The exterior pressure source is at positive physical distance from

$$x_n + \lambda_n B_R$$

for all sufficiently large n . Therefore p_{ext} is harmonic, smooth, and uniformly C^2 -bounded on the physical neighborhood swept out by the compact frame cylinder. Take

$$c_{n,R}(\tau) := \lambda_n^2 p_{\text{ext}}(x_n, t_n + \lambda_n^2 \tau).$$

By the mean value theorem,

$$\|P_{\text{ext}}^n - c_{n,R}\|_{C_y^1(B_R)} \leq C_R \lambda_n^3 \sup |\nabla_x p_{\text{ext}}| + C_R \lambda_n^4 \sup |\nabla_x^2 p_{\text{ext}}|.$$

The right-hand side tends to zero because the physical derivatives are uniformly bounded and $\lambda_n \rightarrow 0$. This proves the asserted convergence. $\square$

Lemma 9.46 (Positive-distance exterior states do not remain in compact core frames). *Let (x_n, λ_n) be a core frame with $x_n \rightarrow x_*$ and $\lambda_n \rightarrow 0$. Let (z_n, μ_n) be an exterior local state satisfying*

$$\liminf_{n \rightarrow \infty} |z_n - x_*| \geq r_* > 0.$$

Assume that the ordered terminal state-space test has already retained the compact core at (x_, λ_n) . Then the positive-distance exterior state does not leave an additional retained velocity component in any compact core frame. More precisely, one of the following ordered alternatives occurs:*

- (i) *the state lies on a no-exterior-concentration annulus and is regular and removable by [theorem 9.42](#);*
- (ii) *the state contains a positive local CKN carrier, in which case it is an exterior, separated-core, multicore, cascade, or rebased local branch rather than part of the selected compact core;*
- (iii) *its velocity contribution is zero in the retained compact-core topology, and its pressure contribution vanishes modulo functions of time by [lemma 9.50](#), unless the far-field harmonic tail is routed to the ordered pressure row by [lemma 9.61](#).*

No whole-space L^3 profile decomposition is used in this decoupling.

Proof. For fixed $R < \infty$, the physical frame ball

$$x_n + \lambda_n B_R$$

lies in $B_{r_*/4}(x_*)$ for all sufficiently large n , while the profile state center z_n lies outside $B_{r_*/2}(x_*)$. Thus every local source supported in the positive-distance state is separated from the compact core frame.

Run the ordered local CKN test on the exterior state. If the test is below threshold on an annulus containing the state, then the no-concentration regularity theorem gives item (i), and [theorem 9.47](#) removes the regular contribution from compact core frames. If the test finds a positive CKN carrier, that carrier is recorded before retention as an exterior-profile, separated-core, multicore, cascade, or rebased local branch; this is item (ii).

It remains to treat the nonretained part after these tests. Its velocity contribution is supported a positive distance from the physical core and therefore vanishes identically on every compact core cylinder for all large n , as in [lemma 9.44](#). The pressure generated by fixed positive-distance sources is perturbative modulo functions of time by [lemma 9.50](#). If the only nonvanishing pressure contribution is a far-field harmonic tail, then [lemma 9.61](#) routes that tail to a dyadic exterior carrier, a regular dyadic tail, or the ordered pressure row. Hence no positive-distance exterior state is silently retained inside the selected compact core. $\square$

Theorem 9.47 (Exterior regular components decouple from compact core frames). *On the no-exterior-concentration stratum for the fixed annuli separating an exterior component from the selected Type II core point, that component is a regular exterior component and vanishes in compact Type II core frames, including its pressure contribution modulo constants.*

Proof. Apply the divergence-free decomposition of lemma 9.43. The exterior component is identically zero in every compact core frame for large n by lemma 9.44. Its pressure contribution is a locally scaled smooth harmonic field and vanishes modulo constants by lemma 9.45. The cutoff and Bogovskii correction terms are supported in the fixed exterior annulus A_{r_0} , so they also vanish in compact core frames by the same argument.

Thus a regular exterior component decouples from the Type II core without changing the local repaired-gauge equation in any compact terminal frame, up to errors converging to zero in the perturbative topology used in lemma 9.10 and definition 9.13. $\square$

Corollary 9.48 (Exterior annulus alternative for the core decomposition). *Fix a compact annulus $A_{r,R}$ around the selected core. Then exactly one of the following alternatives holds for the Type II branch on this annulus:*

- (i) *the branch has no exterior concentration on $A_{r,R}$, and every compact subannulus of $A_{r,R}$ is covered by theorem 9.42; consequently exterior regular contributions from this annulus vanish in compact core frames by theorem 9.47;*
- (ii) *the branch admits an exterior concentration branch on $A_{r,R}$, and therefore belongs to one of the adjacent exterior, separated-core, multicore, noncompact exterior, or scale-collapse alternatives described in theorem 9.39.*

Proof. This is the exterior concentration stratification of theorem 9.39, combined in the first alternative with the no-concentration regularity theorem and the exterior-removal theorem. $\square$

Lemma 9.49 (Exterior CKN mass split). *Let $x_n \in A_{r,R}$ and $\rho_n \downarrow 0$ form an exterior concentration branch on $A_{r,R}$. After passing to a subsequence, one of the following alternatives holds:*

- (i) *velocity carrier:*

$$\rho_n^{-2} \int_{T^* - \rho_n^2}^{T^*} \int_{B_{\rho_n}(x_n)} |u|^3 dx dt \geq \frac{\varepsilon_{\text{ext}}}{2};$$

- (ii) *pressure carrier:*

$$\rho_n^{-2} \int_{T^* - \rho_n^2}^{T^*} \int_{B_{\rho_n}(x_n)} |p - (p)_{B_{\rho_n}(x_n)}(t)|^{3/2} dx dt \geq \frac{\varepsilon_{\text{ext}}}{2}.$$

Proof. By definition of exterior concentration,

$$\mathcal{C}_*(u, p; x_n, \rho_n) > \varepsilon_{\text{ext}}.$$

The displayed quantity is the sum of the velocity and pressure terms. Hence at least one term is at least $\varepsilon_{\text{ext}}/2$ for infinitely many n , and a subsequence gives one of the two alternatives. This lemma is only the local velocity/pressure split and does not itself close the pressure-only case. $\square$

Lemma 9.50 (Positive-distance sources are perturbative in compact core frames). *Let (x_n, λ_n, t_n) be a compact core frame with*

$$x_n \rightarrow x_*, \quad \lambda_n \downarrow 0, \quad t_n \rightarrow T^*.$$

Fix $d > 0$, and let $K \subseteq \{x : |x - x_| \geq d\}$. Suppose F_n is supported in K for a.e. terminal time and satisfies the local bound*

$$\sup_n \|F_n\|_{L^1(T^* - \delta, T^*; L^1(K))} < \infty.$$

Let q_n be the Newtonian pressure generated by F_n ,

$$q_n(\cdot, t) = \partial_i \partial_j (-\Delta)^{-1} F_{n,ij}(\cdot, t),$$

restricted to a neighborhood of x_* . For every compact $B_M \times I$ in the core variables, set

$$Q_n(y, \tau) = \lambda_n^2 q_n(x_n + \lambda_n y, t_n + \lambda_n^2 \tau), \quad c_n(\tau) = \lambda_n^2 q_n(x_n, t_n + \lambda_n^2 \tau).$$

Then

$$Q_n - c_n \rightarrow 0 \quad \text{in } L_\tau^1 C_y^1(B_M).$$

If w_n is a velocity contribution supported in K , then

$$\lambda_n w_n(x_n + \lambda_n y, t_n + \lambda_n^2 \tau) \equiv 0 \quad \text{on } B_M \times I$$

for all sufficiently large n . Thus positive-distance velocity sources and their pressure tails are perturbative in every compact core frame, modulo functions of time.

Proof. For fixed M , the physical balls $x_n + \lambda_n B_M$ lie in $B_{d/4}(x_*)$ for all sufficiently large n , whereas $K \subset \{|x - x_*| \geq d\}$. The velocity assertion follows immediately from the support separation.

On $B_{d/4}(x_*)$ the kernel $\partial_i \partial_j (-\Delta)^{-1}$ applied to sources supported in K is smooth. For $|\alpha| \leq 2$,

$$\|\nabla_x^\alpha q_n(t)\|_{L^\infty(B_{d/4}(x_*))} \leq C_{\alpha,d,K} \|F_n(t)\|_{L^1(K)}$$

for a.e. t . By the mean value theorem,

$$\|Q_n(\tau) - c_n(\tau)\|_{C_y^1(B_M)} \leq C_M \lambda_n^3 \|\nabla_x q_n(t_n + \lambda_n^2 \tau)\|_{L^\infty(B_{d/4})} + C_M \lambda_n^4 \|\nabla_x^2 q_n(t_n + \lambda_n^2 \tau)\|_{L^\infty(B_{d/4})}.$$

Integrating in $\tau \in I$ and using $dt = \lambda_n^2 d\tau$, the right-hand side is bounded by

$$C_{M,I,d,K} (\lambda_n + \lambda_n^2) \|F_n\|_{L^1(T^* - \delta, T^*; L^1(K))},$$

which tends to zero. This gives the claimed convergence. The proof is local on the fixed set K and uses no whole-space critical norm. $\square$

Lemma 9.51 (Exterior local states decouple from the selected compact core). *Let an exterior concentration branch on $A_{r,R}$ be in the exterior-profile subcase of [theorem 9.39](#). Then the exterior profile is not a terminal obstruction for the selected compact core. More precisely, in every compact core frame (x_n, λ_n, t_n) with $x_n \rightarrow x_*$ and $\lambda_n \downarrow 0$, the exterior local state has zero compact-core velocity contribution in the retained terminal state-space topology. Its pressure contribution either vanishes modulo functions of time on compact core cylinders, or it is routed by the far-field harmonic-tail theorem [lemma 9.61](#) to a dyadic exterior carrier or to the ordered pressure row. If the exterior state itself is selected as a possible singular branch, it is rebased at its own center and tested by the same local Type II state-space decomposition; it is no longer an exterior exit for the original selected core.*

Proof. The exterior-profile subcase is read locally: its CKN witnesses remain at positive physical distance from x_* , or escape every compact set in the selected core variables. Suppose that a positive velocity contribution from this exterior state were visible on some compact core cylinder. Then [theorem 9.16](#) would assign that positive compact-core mass to a first local state-space label. If the label is retained comparable mass, it must already have been grouped into the selected core; if it is not retained comparable mass, it is an adjacent compact, separated, cascade, rough-core, or pressure/exterior row. Hence no positive exterior velocity state remains in the nonretained exterior part of a retained compact core branch.

For pressure, decompose the exterior pressure seen on a compact core cylinder into the local Calderón–Zygmund part generated by sources in a fixed annulus, the far-field harmonic tail, and a function of time. Fixed positive-distance annular sources are perturbative by [lemma 9.50](#). If

the remaining far-field harmonic tail does not vanish modulo functions of time, then [lemma 9.61](#) produces a dyadic exterior stress carrier or the ordered pressure row. Thus pressure is also either perturbative for the selected core or routed before retention.

Thus the exterior profile is perturbative for the selected core. If one chooses the exterior profile as the new active branch, the center and scale are changed and the analysis restarts from the covered-entry theorem [corollary 11.26](#). This rebasing does not leave an independent exterior terminal row attached to the original core. $\square$

Lemma 9.52 (Separated exterior cores are multicore or rebased local branches). *Suppose an exterior concentration branch on $A_{r,R}$ carries a velocity carrier and that, after passing to a subsequence, the exterior carrier remains as a positive active core at terminal scales comparable to another selected active scale. Then either the two active centers are grouped as a separated multicore branch in the local state space, or the exterior center is rebased as its own physical local Type II germ. In the first case the branch is excluded by [theorem 11.6](#); in the second case the row is no longer an exterior row for the original selected core.*

Proof. The exterior centers satisfy $|x_n - x_*| \geq r > 0$, while the selected compact core scales tend to zero. Hence the relative scale-center distance between the original core and the exterior carrier tends to infinity whenever the scales are comparable, and the separated-point reduction [theorem 5.21](#) applies. Its output is exactly a separated-core or multicore label unless the exterior center is selected as the new base point. The separated-core and multicore labels are part of the multibubble/cascade branch excluded by [theorem 11.6](#). If the exterior center is chosen as the base point, the local CKN lower bound supplies a physical local concentration sequence there and [corollary 11.26](#) starts the local Type II analysis for that new germ. This is a rebasing, not an unresolved exterior mechanism for the original core. $\square$

Lemma 9.53 (Exterior finite-scale packing is an ordered local row). *Fix a compact annulus $A_{r,R}$. In any fixed dyadic scale band of exterior carrier cylinders contained in $A_{r,R} \times (T^* - \delta, T^*)$, a Vitali selection produces a pairwise disjoint subfamily with uniformly bounded overlap after enlargement. Therefore an exterior velocity carrier in that scale band has an ordered local interpretation:*

- (i) *a finite selected subfamily gives a separated-core or multicore label;*
- (ii) *an unbounded number of comparable exterior carriers, or a sequence of scale bands with no finite selected representative, is by definition a noncompact exterior or cascade boundary face before the retained compact core is declared.*

No uniform bound on the number of exterior carriers is asserted or used.

Proof. The geometric Vitali covering lemma applies on the fixed annulus for each dyadic radius band. It gives a disjoint subfamily whose fixed enlargement covers all carrier balls in that band, with overlap bounded by an absolute constant depending only on the enlargement factor. The ordered local state-space construction selects active carriers from this finite-overlap family. If only finitely many selected carriers remain, their positive CKN lower bounds and separated scale-center relations give the separated-core or multicore row. If no finite selected list captures the active exterior carriers, then the sequence has left every finite scale-center chart; this is the noncompact exterior boundary face, or the same-point scale-cascade face when the centers converge but the radii separate strictly. The argument is geometric and local; it does not invoke a global L^3 mass bound to bound the cardinality of the family. $\square$

Lemma 9.54 (Exterior scale-center compactification). *Let (z_n, σ_n) be exterior active carrier data in an exterior frame, with $\sigma_n > 0$, and suppose the corresponding terminal cylinders carry a fixed positive CKN lower bound. After passing to a subsequence, exactly one of the following ordered alternatives occurs:*

- (i) the normalized centers z_n remain in a compact subset of the exterior frame; this is a compact exterior-profile, separated-core, or multicore row;
- (ii) $|z_n| \rightarrow \infty$ and $\sigma_n/|z_n| \rightarrow 0$; this is a dyadic point carrier in annuli $\{|y| \simeq |z_n|\}$;
- (iii) $|z_n| \rightarrow \infty$ and

$$0 < \liminf_{n \rightarrow \infty} \sigma_n/|z_n| \leq \limsup_{n \rightarrow \infty} \sigma_n/|z_n| < \infty;$$

this is a dyadic annular carrier;

- (iv) the relative scale leaves every compact subset at a fixed exterior point, or after rebaseing at the exterior point; this is the same-point cascade or exterior scale-collapse row.

These alternatives are disjoint after the ordered first-failure convention and exhaust the exterior scale-center configurations. In particular, the phrase “standard boundary faces” in the exterior routing below means precisely the four cases listed here.

Proof. If z_n is bounded, compactness of closed balls in the exterior frame gives a subsequential limit, yielding item (i). If $|z_n| \rightarrow \infty$, pass to a subsequence so that $\sigma_n/|z_n|$ converges in the compactified interval $[0, \infty]$. Limit 0 gives item (ii). A limit in $(0, \infty)$ gives item (iii). Limit $+\infty$, or failure of compactness after recentering at a bounded exterior point, means that the active radius contains a further unresolved scale. The ordered local state-space compactification then records either strict same-point scale separation or a genuine scale-collapse row, giving item (iv). These possibilities exhaust the one-point compactification of the center variable together with the compactified relative-scale variable. $\square$

Lemma 9.55 (Dyadic exterior occupation selection). *Work in an exterior frame and let*

$$\mathcal{A}_k := B_{2^{k+1}} \setminus B_{2^k}, \quad k \geq 0.$$

For $K \geq 0$, define the far-dyadic CKN level

$$\Gamma_{n,K} := \sup_{\substack{k \geq K, y \in \mathcal{A}_k \\ 0 < r \leq 2^{k-4} \\ Q_r^*(y) \text{ admissible}}} \mathcal{C}_*(U_n, P_n; y, r),$$

where the terminal cylinders are understood after pulling back to the physical variables. Then the ordered test gives one of the following alternatives after passing to a subsequence:

- (i) a dyadic CKN carrier exists:

$$\mathcal{C}_*(U_n, P_n; y_n, r_n) \geq \varepsilon_{\text{ext}}$$

for a sequence $|y_n| \rightarrow \infty$;

- (ii) the far-dyadic CKN level is eventually below threshold:

$$\limsup_{K \rightarrow \infty} \limsup_{n \rightarrow \infty} \Gamma_{n,K} \leq \varepsilon_{\text{ext}},$$

and, because $\varepsilon_{\text{ext}} < \varepsilon_{\text{CKN}}$, every sufficiently far dyadic annulus is regular on its interior by the terminal CKN criterion.

Thus a noncompact exterior branch can remain a terminal obstruction only in the dyadic CKN-carrier case.

Proof. If

$$\limsup_{K \rightarrow \infty} \limsup_{n \rightarrow \infty} \Gamma_{n,K} > \varepsilon_{\text{ext}},$$

then for arbitrarily large K there are n , $k \geq K$, $y \in \mathcal{A}_k$, and an admissible $r \leq 2^{k-4}$ with $\mathcal{C}_*(U_n, P_n; y, r) > \varepsilon_{\text{ext}}$. A diagonal subsequence gives item (i) with $|y_n| \rightarrow \infty$.

Otherwise the displayed limsup is at most ε_{ext} . Fix η with $\varepsilon_{\text{ext}} < \eta < \varepsilon_{\text{CKN}}$. After discarding a finite initial set of annuli and passing to the selected tail in n , every admissible subcylinder in

every far dyadic annulus has CKN quantity at most η , hence below ε_{CKN} . A bounded-overlap cover of each smaller annulus

$$B_{2^{k+1}-2^{k-5}} \setminus B_{2^k+2^{k-5}}$$

by such cylinders and [theorem 9.35](#) give local smoothness on the dyadic annulus. Since the cover is local on each annulus and the constants are scale-invariant after rescaling that annulus to unit size, no global tail norm is used. Therefore the diffuse no-carrier case is regular on every far dyadic annulus and cannot be a terminal noncompact exterior obstruction. $\square$

Lemma 9.56 (Exterior same-point cascades route to the cascade package). *Let exterior carriers on $A_{r,R}$ have centers converging to the same physical point $z \neq x_*$ and have strictly separated scales. Then, after rebaseing at z , the branch is a same-point cascade in the sense of [theorem 5.20](#). It is therefore closed by the multibubble/cascade package [theorem 11.6](#), unless the ordered rebased analysis exits earlier through one of the non-exterior first-failure rows.*

Proof. Because $z \neq x_*$, the carriers are at positive physical distance from the original core and are perturbative for that core by [lemma 9.50](#). Recenter at z and use the exterior carrier scales as the active scales. Strict separation of the radii is precisely the same-point scale-cascade boundary face in the relative scale-center compactification. The same-point cascade reduction [theorem 5.20](#) then assigns the rebased branch to the cascade branch, which is excluded by [theorem 11.6](#). If the rebased ordered tests fail before the cascade label is reached, the failure is one of the ordinary representation, modulation, active-core, pressure, rough-core, carrier, or cost rows; none is an exterior row for the original selected core. $\square$

Lemma 9.57 (Noncompact exterior branches admit dyadic active carriers). *Let the exterior stratification of [theorem 9.39](#) produce the noncompact exterior alternative in an exterior frame. Then, after passing to a subsequence, exactly one of the following occurs:*

- (i) *the noncompact exterior occupation is diffuse below the CKN threshold on all sufficiently far dyadic annuli, hence regular there by [lemma 9.55](#);*
- (ii) *there are dyadic annuli and terminal cylinders in the exterior frame with CKN quantity bounded below by ε_{ext} , hence an exterior velocity or pressure carrier by [lemma 9.49](#);*
- (iii) *the noncompact behavior is only a far-field harmonic pressure tail, which is routed to item (i), item (ii), or the ordered pressure row by [lemma 9.61](#).*

Proof. The noncompact alternative is represented by active data leaving every compact subset of the exterior scale-center chart. Apply [lemma 9.54](#). The compact-center, same-point cascade, and scale-collapse faces are not the noncompact exterior face and are routed by their corresponding lemmas below. The remaining noncompact face is a dyadic occupation state in annuli $\mathcal{A}_k$, $k \rightarrow \infty$.

Apply the quantitative dyadic selection lemma [lemma 9.55](#). If the occupation is below the CKN threshold on all sufficiently far dyadic annuli, item (i) holds. Otherwise a dyadic CKN carrier exists. The velocity/pressure split [lemma 9.49](#) gives item (ii) unless the carrier is purely pressure. A purely harmonic far-field pressure tail is then routed by [lemma 9.61](#), giving item (iii). All tests are dyadic and local. $\square$

Lemma 9.58 (Noncompact exterior carriers route to structural rows). *Every dyadic active carrier produced by [lemma 9.57](#) is routed, after passing to a subsequence, to one of the following rows:*

- (i) *a regular exterior-removal row, if all smaller dyadic tests are no-concentration rows;*
- (ii) *a separated-core or multicore row, closed by [theorem 11.6](#);*
- (iii) *a same-point cascade row, closed by [lemma 9.56](#);*
- (iv) *an exterior scale-collapse row, rebased and tested by [lemma 9.59](#);*

- (v) a pressure/exterior row which is converted to a velocity carrier or to a non-exterior first-failure row by [lemma 9.70](#).

Thus noncompact exterior concentration is not an independent terminal alternative.

Proof. Use the dyadic carrier selected in the preceding lemma. If adjacent dyadic tests become no-concentration tests after shrinking, then [theorem 9.47](#) gives regular exterior removal. Otherwise [lemma 9.54](#) gives the ordered boundary face of the carrier. Positive centers separated from the original core at comparable scales give the separated-core or multicore row by [lemma 9.52](#). Centers converging to one exterior point with strict scale separation give the same-point cascade row by [lemma 9.56](#). If the dyadic carrier develops its own unresolved scale collapse, it is the exterior scale-collapse row and is handled by [lemma 9.59](#). Finally, if the diagnostic is pressure-only, the pressure-only routing lemma applies. These cases are the ordered local boundary faces of the exterior scale-center chart and exhaust the noncompact exterior alternative. $\square$

Lemma 9.59 (Exterior scale collapse is an ordinary rebased scale-collapse row). *Let an exterior concentration branch develop scale collapse in its own exterior variables. Recenter at the exterior active carrier and use its scale as the local active scale. Then the rebased branch is a represented local scale-collapse branch in the sense of [theorem 11.9](#), unless the ordered rebased tests fail first through one of the non-exterior rows: representation/pressure failure, autonomous-modulation failure or modulation-driven recentering, loss of retained active core or pressure-only critical mass, rough-core loss, carrier/cost validation failure, or multicore/cascade. In the represented case, the finite-cost, retained compact scale-collapse, and scale-rigid subcases are closed by*

[theorems 7.37 and 11.8](#) and [proposition 5.42](#).

Proof. The exterior carrier supplies a positive local CKN mass floor on a compact cylinder in the rebased variables. Run the covered-entry and repaired-gauge selection tests at that carrier. If a row fails before a represented scale-collapse branch is obtained, the first-failure ledger records exactly the listed non-exterior local row; this is the same ordered local test used in [corollary 11.26](#) and [lemma 7.26](#).

If the represented scale-collapse branch is obtained, apply [theorem 11.9](#). Finite scale cost and finite absolute scale cost are excluded by [theorem 11.8](#). Thin negative drift is already closed as an independent terminal row by [theorem 10.118](#). Selected compactness-package failure is routed by [theorem 7.34](#). The retained compact scale-collapse state is excluded by [theorem 7.37](#), and the scale-rigid residual state is excluded by [proposition 5.42](#). Autonomous modulation and the other listed first failures remain ordinary non-exterior rows in the final ledger; the exterior label itself has disappeared after rebasing. $\square$

Lemma 9.60 (Fixed annular stress is regular or has a local CKN carrier). *Let $\mathcal{A} \Subset \mathbb{R}^3 \setminus B_2$ be a fixed annulus in a local frame and let $I \Subset (-\infty, 0]$ be a compact terminal time interval. Suppose a subsequence has a positive annular stress level*

$$\limsup_{n \rightarrow \infty} \|U_n \otimes U_n\|_{L^q_+(I; L^{3/2}_y(\mathcal{A}))} > 0$$

for some $1 < q < \infty$. Then the ordered local CKN test on $\mathcal{A} \times I$ gives one of the following alternatives:

- (i) a local CKN carrier occurs on a subcylinder contained in a slightly larger annulus;
- (ii) all sufficiently small admissible subcylinders in the annulus are below ε_{CKN} , hence the annular stress is a regular compact-window annular source. It is absorbed into the retained smooth annular part of the local pressure atlas, unless the harmonic-tail test sends its pressure effect to the ordered pressure row.

Thus a positive fixed-annulus stress level is not by itself promoted to a singular carrier; it is tested locally first.

Proof. Apply the terminal CKN test to all admissible subcylinders whose doubled cylinders remain in a fixed slightly larger annulus containing $\mathcal{A}$. If one of these cylinders has CKN quantity at least ε_{ext} , then item (i) holds. If no such cylinder exists, every admissible subcylinder is below $\varepsilon_{\text{ext}} < \varepsilon_{\text{CKN}}$. The terminal CKN regularity theorem and a finite-overlap cover of compact subannuli give local smoothness on $\mathcal{A} \times I$, so the stress lower bound is a regular annular source rather than a singular exterior carrier. Since $\mathcal{A}$ is at bounded normalized distance from the core in the same local frame, this regular source is part of the compact-window pressure atlas; if it is not absorbed there, the first failed atlas condition is exactly the ordered pressure row. No cardinality bound or global L^3 tail estimate is used. $\square$

Lemma 9.61 (Far-field harmonic tails route to dyadic carriers). *Let $B_1 \subseteq B_2$ be nested balls in a local frame and let*

$$H_n(\tau) = H_n^{\text{far}}(\tau)$$

be the mean-zero harmonic pressure on B_2 generated, in the local pressure atlas, by the quadratic stress $F_n = U_n \otimes U_n$ outside B_4 . Let $\{\theta_k\}_{k \geq 2}$ be a smooth dyadic partition of unity on $\mathbb{R}^3 \setminus B_4$, with

$$\text{supp } \theta_k \subset B_{2^{k+1}} \setminus B_{2^{k-1}}.$$

Assume the pressure reconstruction and gauge rows have not failed. If, for some $1 < q < \infty$,

$$\limsup_{n \rightarrow \infty} \|H_n\|_{L^q_{\tau} L^{3/2}_y(B_1)} > 0,$$

then, after passing to a subsequence, either

(i) *a dyadic exterior velocity-stress occupation state is present:*

$$\limsup_{n \rightarrow \infty} \sum_{k \geq K_n} 2^{-2k} \|U_n \otimes U_n\|_{L^q_{\tau} L^{3/2}_y(\mathcal{A}_k)} > 0$$

for some $K_n \rightarrow \infty$; or

(ii) *a fixed dyadic annulus carries a positive annular stress lower bound, which is then tested by lemma 9.60 and is either a regular annular source or a local CKN carrier.*

In item (i), the dyadic occupation state is then tested by lemma 9.55: it either produces a dyadic CKN carrier or is regular on the far dyadic annuli. If the dyadic tail is regular but its harmonic pressure contribution remains nonzero on the interior cylinder, the branch is assigned to the ordered pressure row rather than to exterior concentration. Consequently a nonzero far-field harmonic pressure tail is never silently discarded; it is a dyadic exterior carrier, a regular dyadic tail with vanishing harmonic effect, or the ordered pressure row.

Proof. Write

$$H_n = \sum_{k \geq 2} H_{n,k} + h_n, \quad H_{n,k} = \mathcal{R}_i \mathcal{R}_j (\theta_k F_{n,ij}) - \oint_{B_2} \mathcal{R}_i \mathcal{R}_j (\theta_k F_{n,ij}) dy,$$

where h_n is the pressure-gauge/reconstruction remainder. By assumption the gauge and reconstruction rows have not failed, so h_n is absent modulo functions of time in the retained pressure atlas.

For $y \in B_2$ and $z \in \text{supp } \theta_k$, the second derivative of the Newtonian kernel is $O(2^{-3k})$. Hence

$$\|H_{n,k}(\tau)\|_{L^{3/2}(B_1)} \leq C 2^{-3k} \|F_n(\tau)\|_{L^1(\mathcal{A}_k)} \leq C 2^{-2k} \|F_n(\tau)\|_{L^{3/2}(\mathcal{A}_k)},$$

because $|\mathcal{A}_k|^{1/3} \simeq 2^k$. Summing in k and taking the L_τ^q norm gives

$$\|H_n\|_{L_\tau^q L^{3/2}(B_1)} \leq C \sum_{k \geq 2} 2^{-2k} \|F_n\|_{L_\tau^q L^{3/2}(\mathcal{A}_k)}.$$

If the left-hand side has positive limsup, the weighted dyadic stress sum has positive limsup. Either a fixed finite set of dyadic annuli carries a positive portion of that limsup, or the positive portion escapes to $k \rightarrow \infty$. The first case gives item (ii) and is routed by [lemma 9.60](#); a fixed annular stress lower bound is not identified with a singular carrier until the local CKN test has selected one. The escaping case is precisely the dyadic occupation state in item (i). Applying [lemma 9.55](#) to that occupation state gives a dyadic CKN carrier or a regular dyadic tail. In the regular-tail case, either the corresponding harmonic contribution vanishes modulo functions of time on B_1 , or the nonzero harmonic contribution is by definition the ordered pressure row. No whole-space pressure or velocity norm is used. $\square$

Lemma 9.62 (Compact local-state factor with failure faces). *Fix a model cylinder Q and choose a countable exhaustion $Q^{(\ell)} \Subset Q$ by compact subcylinders, where $B^{(\ell)}$ denotes the spatial ball of $Q^{(\ell)}$. Choose also countable dense families of smooth test fields $\psi_j, \Psi_j, \varphi_j$ on those subcylinders. Consider represented suitable branches on Q . For each compact subcylinder record the coordinates*

$$\int_{Q^{(\ell)}} V \cdot \psi_j, \quad \int_{Q^{(\ell)}} \nabla V : \Psi_j, \quad \int_{Q^{(\ell)}} (P - (P)_{B^{(\ell)}}(t)) \varphi_j,$$

together with the local-energy, pressure-atlas, gauge, compact-upper, rough-core, modulation, and cost diagnostic coordinates, each compactified by adding the value ∞ or a named failure face.

Then the closure of all represented local states satisfying the retained finite bounds is compact metrizable. Any limit point either:

- (i) *is again a represented local suitable state with the inherited pressure atlas and local energy inequality on smaller compact cylinders; or*
- (ii) *lies on one of the explicitly named failure faces.*

Moreover, the occupation coordinates used in the pressure, shell, and reselection compactifications are continuous coordinate projections in this compactification.

Proof. On each retained chart the compact-window estimates furnish uniform weak- $*$ $L_t^\infty L_x^2$, weak $L_t^2 H_x^1$, and weak $L^{3/2}$ pressure bounds on the model cylinder. Banach–Alaoglu therefore gives compactness of the recorded linear coordinates, and Aubin–Lions gives strong local velocity compactness on every subcylinder for which the time-derivative/modulation row remains finite.

More explicitly, the compactified state space is defined as the closed subset of the product

$$\prod_j K_j$$

where each K_j is one of the following compact factors: a weakly compact closed ball for a retained finite linear coordinate, a one-point compactified interval $[0, \infty]$ for a nonnegative diagnostic size, or a finite discrete set of named failure faces. The occupation coordinates are included as coordinate projections in this product topology. Thus their positive superlevel sets are closed, and weak- $*$ convergence of normalized occupation measures is exactly weak- $*$ convergence against these continuous coordinate projections. The distributional equation and local energy inequality are imposed as closed conditions on the retained finite-coordinate part; if any condition fails, the state lies on the corresponding failure face rather than inside the retained suitable-state component.

Passing to a diagonal subsequence over the countable test families yields a compact metrizable closure. Lower semicontinuity of the local energy inequality and the stability of the local pressure

atlas preserve the represented local suitable structure on smaller compact cylinders. If one of the compactness inputs is not finite, then by definition the corresponding coordinate lands on its named failure face. This is exactly the stated alternative. The pressure, shell, and reselection occupation variables are among the product coordinates just recorded, so their projections are continuous in the compactified topology. $\square$

Lemma 9.63 (Normalized defect mass: atomic or diffuse visible row). *Let μ_n be normalized nonnegative defect measures on a compactified local state space X , and suppose*

$$\mu_n \rightharpoonup \mu, \quad \mu(X) = 1.$$

Assume the compactification includes a countable separating family of nonnegative local occupation coordinates

$$\mathcal{O}_j : X \rightarrow [0, \infty],$$

generated by compactly supported nonnegative test functions ϕ_j which separate nonzero local L_+^1 -mass on compact subcylinders. These coordinates are assumed to be continuous coordinate projections in the product compactification, and for every j and $\eta > 0$ the superlevel set

$$\{x \in X : \mathcal{O}_j(x) \geq \eta\}$$

is closed. Then exactly one of the following occurs:

- (i) *there exist j and $\eta > 0$ such that*

$$\mu(\{x \in X : \mathcal{O}_j(x) \geq \eta\}) > 0.$$

In this case a support point with positive compact local activity is realized by descendant states;

- (ii) *for every j and every $\eta > 0$,*

$$\mu(\{x \in X : \mathcal{O}_j(x) \geq \eta\}) = 0.$$

Then the defect is genuinely diffuse. The nonzero probability mass μ itself is a visible diffuse occupation-state coordinate, and it must be routed as a named diffuse row rather than replaced by a nonzero support profile.

Proof. If alternative (i) holds, the set $\{x \in X : \mathcal{O}_j(x) \geq \eta\}$ is closed and has positive μ -mass, so it meets $\text{supp } \mu$. Any support point in that set carries positive compact local activity, and the relevant support-realization theorem turns it into an actual descendant realization.

If alternative (i) fails, then alternative (ii) holds by definition. Because the coordinates are countable and separate compact local mass, no support point can carry positive compact local activity at a positive level. However $\mu(X) = 1$, so the limiting defect is still nonzero. Since the compactification records the normalized defect mass itself and not only its support points, this nonzero diffuse measure remains a visible state-space coordinate. The two alternatives are manifestly exclusive and exhaustive. $\square$

Lemma 9.64 (Residual selected-time support transfer). *Let*

$$V_n = U_n + S_n$$

be an admissible retained terminal decomposition on a compact terminal cylinder $Q_{2R,I} = B_{2R} \times I$. Fix selected times $\tau_n \in I$ and a compact ball $B_\rho \Subset B_R$. Let

$$\mu_n^{\text{res}} := |S_n(\tau_n)|^3 dy|_{B_\rho}$$

and assume $\mu_n^{\text{res}}(B_\rho) > 0$ for all n . Recenter the selected times by [lemma 10.70](#) so that they are represented on a fixed model interval. For each point carrying residual mass, record its compactified relative center-scale data, the rebased local occupation state on a fixed model cylinder,

and the named failure faces of the compact local-state factor. Then, after normalizing by $\mu_n^{\text{res}}(B_\rho)$ and passing to a subsequence, these pushforwards converge weakly-* to a probability measure Π_{res} on a compact metrizable residual state space

$$\mathfrak{X}_{\text{res}} := \bar{\mathcal{C}}_{\text{res}} \times \mathfrak{D}_{\text{res}} \times \mathfrak{F}_{\text{res}}.$$

The local occupation coordinates of $\mathfrak{X}_{\text{res}}$ are continuous coordinate projections, so the dichotomy of [lemma 9.63](#) applies to Π_{res} . Every point of $\text{supp } \Pi_{\text{res}}$ is realized by an actual sequence of selected-time residual descendants and their corresponding rebased local states from the original represented branch.

Proof. Use [lemma 10.70](#) to place the selected time at the origin of a fixed model interval in the local validation family. For each n , send a point of B_ρ carrying residual mass to its compactified relative center-scale data together with the associated rebased local state on that fixed model cylinder and the named failure faces of [lemma 9.62](#). The compactified relative scale-center chart is compact, and the local-state factor with failure faces is compact metrizable by [lemma 9.62](#). Hence their product $\mathfrak{X}_{\text{res}}$ is compact metrizable.

Set

$$\hat{\mu}_n^{\text{res}} := \frac{\mu_n^{\text{res}}}{\mu_n^{\text{res}}(B_\rho)}$$

and push $\hat{\mu}_n^{\text{res}}$ forward by the residual state map to $\mathfrak{X}_{\text{res}}$. By compactness, after passing to a subsequence, these pushforwards converge weakly-* to a probability measure Π_{res} .

Let $\mathfrak{r} \in \text{supp } \Pi_{\text{res}}$. Choose a nested neighborhood basis (U_m) of $\mathfrak{r}$ with diameters tending to zero. Since $\Pi_{\text{res}}(U_m) > 0$, weak-* convergence implies that the approximating pushforwards assign positive mass to U_m for all large n . For each m , select an actual residual descendant whose rebased state lies in U_m . These genuine rebased local states converge to $\mathfrak{r}$, proving the support-realization statement.

The occupation coordinates of $\mathfrak{X}_{\text{res}}$ come from the product coordinates of the compact local-state factor, hence are continuous coordinate projections with closed positive superlevel sets. Therefore [lemma 9.63](#) applies directly to the probability measure Π_{res} . $\square$

Lemma 9.65 (Pressure-defect support points are realized by descendant rebasing). *Fix once and for all a canonical dyadic family of descendant cylinders inside the unit cylinder. For each normalized pressure defect π_j arising from the nested parent cylinder $Q_{\rho_j}(z_j)$, record for every dyadic descendant: its compactified relative scale-center coordinates inside the parent, the rebased velocity/pressure tuple on a fixed model cylinder, the named failure faces for pressure atlas, pressure gauge, pressure reconstruction, compact-upper, local-energy, and rough-core breakdown, and the harmonic jets of the rebased harmonic remainder on a slightly smaller model cylinder. By [lemma 9.62](#), the local-state factor with these named failure faces is compact metrizable, so the closure of this family in the product topology is a compact metrizable pressure-defect state space*

$$\mathfrak{X}_{\text{pres}} := \bar{\mathcal{C}}_{\text{pres}} \times \mathfrak{D}_{\text{pres}} \times \mathfrak{H}_{\text{pres}},$$

where $\bar{\mathcal{C}}_{\text{pres}}$ is the compactified relative scale-center chart of the descendants, $\mathfrak{D}_{\text{pres}}$ is the local occupation-state space obtained from the compact-cylinder suitable bounds, local energy structure, rough-core data, and the local pressure atlas on the rebased model cylinder, together with the corresponding named failure faces, and $\mathfrak{H}_{\text{pres}}$ is the compact harmonic-jet factor supplied by interior harmonic estimates. The local occupation coordinates in $\mathfrak{D}_{\text{pres}}$ are continuous coordinate projections of this product compactification; hence each positive superlevel set of a pressure-defect occupation coordinate is closed in $\mathfrak{X}_{\text{pres}}$.

After passing to a subsequence, the pushforwards of π_j to $\mathfrak{X}_{\text{pres}}$ converge weakly-* to a probability measure Π_{pres} . The limiting probability measure itself is retained as part of the compactification

data, so later diffuse arguments may use either its support points or the diffuse measure itself via [lemma 9.63](#). Every point $\mathbf{p} \in \text{supp } \Pi_{\text{pres}}$ is realized by an actual sequence of descendant cylinders

$$Q_{\sigma_j \rho_j}(w_j) \subset Q_{\rho_j}(z_j)$$

and the corresponding rebased suitable solutions from the original branch. Moreover, on every realizing sequence the decomposition

$$P_j = P_j^{\text{loc}} + H_j$$

transports to the rebased cylinders: the rebased P_j^{loc} is the local Calderón–Zygmund pressure generated by the rebased truncated stress, and the rebased H_j is harmonic on a slightly smaller model cylinder. Hence every realized support profile inherits the local Calderón–Zygmund/harmonic splitting.

Proof. On each parent cylinder, first normalize to the unit cylinder and then assign to every point its full itinerary of canonical dyadic descendants. For a descendant of relative radius σ , translate and dilate by $\sigma \rho_j$ and subtract the appropriate pressure mean. Translation, dilation, and subtraction of time-dependent means preserve the suitable-solution equation, the local energy inequality, and the local pressure decomposition on a fixed interior model cylinder. The rebased velocity variables therefore lie in the compact local-state factor furnished by [lemma 9.62](#), while the rebased harmonic remainders have compact jet coordinates by interior harmonic estimates. Together with the compactified relative scale-center chart, this gives the compact metrizable state space $\mathfrak{X}_{\text{pres}}$.

Push π_j forward by this measurable state map. Since $\mathfrak{X}_{\text{pres}}$ is compact metrizable, Prokhorov and Banach–Alaoglu give a weak-* convergent subsequence $\Pi_{\text{pres}}^{(j)} \rightarrow \Pi_{\text{pres}}$. Now let $\mathbf{p} \in \text{supp } \Pi_{\text{pres}}$. Choose a nested neighborhood basis $(U_m)_{m \geq 1}$ of $\mathbf{p}$ with diameters tending to zero. Since $\Pi_{\text{pres}}(U_m) > 0$, weak-* convergence implies that for each m and all sufficiently large j one has $\Pi_{\text{pres}}^{(j)}(U_m) > 0$. Pick an index $j_m \uparrow \infty$ and an actual descendant cylinder/state from the original sequence whose rebased state lies in U_m . By construction, these rebased states converge to $\mathbf{p}$, so $\mathbf{p}$ is realized by a subsequential rebasing of the original suitable solutions.

For each realizing descendant cylinder choose a cutoff equal to one on the smaller rebased model cylinder and supported inside the parent cylinder. The usual local pressure atlas gives

$$P_j^{\text{rebased}} = \mathcal{R}_i \mathcal{R}_j (\zeta(V_j^{\text{rebased}} \otimes V_j^{\text{rebased}})_{ij}) + H_j^{\text{rebased}}$$

modulo functions of time, with H_j^{rebased} harmonic on a slightly smaller cylinder. This is exactly the transported Calderón–Zygmund/harmonic splitting claimed above. $\square$

In the pressure-defect compactification we use the same dyadic activity ladder as in the shell compactification: a nonzero compact local stress or velocity source is recorded at the first dyadic level $\eta_m > 0$ and first basis test function for which the corresponding local occupation coordinate is positive. Thus positive compact source is routed at an all-positive local level, not only at the original retained threshold.

Theorem 9.66 (Pressure-defect compactification and support transfer). *Let*

$$Q_{\rho_j}(z_j)$$

be a nested pressure-only selection with normalized pressure defects

$$\pi_j := \frac{\int |P_j - (P_j)_{B_{\rho_j}}(t)|^{3/2} dx dt}{\iint_{Q_{\rho_j}(z_j)} |P_j - (P_j)_{B_{\rho_j}}(t)|^{3/2} dx dt}.$$

After passing to a subsequence, the induced defect states of [lemma 9.65](#) converge to a probability measure Π_{pres} on the compact metrizable pressure-defect state space $\mathfrak{X}_{\text{pres}}$. Every point in $\text{supp } \Pi_{\text{pres}}$ is realized by a subsequential rebasing of the original suitable solutions, and every such realized support profile satisfies exactly one of the following alternatives:

- (i) a nonzero local Calderón–Zygmund source generated by velocity stress;
- (ii) a non-affine harmonic tail;
- (iii) an affine or pure-gauge harmonic pressure profile;
- (iv) a pressure-atlas, pressure-gauge, pressure-reconstruction, compact-upper, local-energy, or rough-core failure face.

Proof. Apply [lemma 9.65](#). This gives the compact metrizable state space, the weak-* limit measure Π_{pres} , and an explicit realizing sequence of descendant cylinders for every support point.

Fix $\mathfrak{p} \in \text{supp } \Pi_{\text{pres}}$ and let $(\tilde{V}_m, \tilde{P}_m)$ be a realizing rebased suitable sequence on a fixed model cylinder. If $\mathfrak{p}$ lies on a compactification face at which admissible pressure reconstruction, gauge normalization, the pressure atlas, compact-cylinder upper control, local energy, or rough-core control breaks down, then by the ordered state-space test we are in alternative (iv). Hence we may assume that the realizing sequence remains in the represented part of the state space. The support-realization lemma then supplies the rebased splitting

$$\tilde{P}_m = \tilde{P}_m^{\text{loc}} + \tilde{H}_m$$

modulo functions of time, with $\tilde{P}_m^{\text{loc}}$ the local Calderón–Zygmund pressure of the rebased stress and $\tilde{H}_m$ harmonic on a slightly smaller model cylinder. Passing to the limit in the occupation-state and harmonic-jet coordinates yields the same splitting for the realized support profile.

If the limiting Calderón–Zygmund component is nonzero, then some compact occupation coordinate of the local velocity-stress source is positive. By the dyadic activity ladder fixed above, this source is recorded at the first positive local level $\eta_m > 0$, so alternative (i) occurs at an all-positive local level rather than only at the original retained threshold. If the Calderón–Zygmund component vanishes, the realized profile is harmonic on the interior cylinder. A nonzero non-affine harmonic profile is exactly alternative (ii), while a pure-gauge harmonic profile or an affine harmonic profile after admissible accelerating rebasement is alternative (iii). Thus one and only one of the represented support-point alternatives (i)–(iv) occurs. $\square$

Lemma 9.67 (Nested pressure-only tests are well-founded). *Let*

$$Q_{\rho_j}(z_j)$$

be a nested sequence of cylinders selected by the exterior pressure-only test. Assume that each cylinder carries a pressure lower bound at the pressure-only threshold and that no velocity carrier has been selected at any previous stage. Then, after passing to a subsequence, one of the following named rows occurs:

- (i) a velocity-stress carrier row;
- (ii) a far-field harmonic pressure-tail row;
- (iii) a pressure-gauge or pressure-reconstruction failure row;
- (iv) a same-point pressure cascade or scale-collapse row;
- (v) a diffuse pressure-only defect row.

Proof. On each selected cylinder, decompose the pressure into the local Calderón–Zygmund part and the harmonic part:

$$P_j = P_j^{\text{loc}} + H_j.$$

If P_j^{loc} carries a fixed fraction of the pressure lower bound, then the local pressure is generated by the local velocity stress $V_j \otimes V_j$. The fixed-annular stress dichotomy gives either a velocity

carrier on a slightly enlarged cylinder or a stress concentration at a smaller nested scale. In the first case we are in alternative (i). In the second case the nested scale test continues.

If H_j carries a fixed fraction of the pressure lower bound, then the harmonic interior estimates and the far-field harmonic-tail lemma give either a far-field harmonic pressure-tail row or a pressure-gauge/reconstruction failure. These are alternatives (ii) and (iii).

It remains to consider the possibility that the pressure-only test continues indefinitely through smaller nested cylinders without producing a velocity carrier, harmonic tail, or gauge failure. Compare the selected centers and scales. If the centers remain at the same physical point and the scales form a strict nested chain, the branch is a same-point pressure cascade. If the relative centers or relative scales fail to have a compact limit, the branch is a scale-collapse row. This gives alternative (iv).

Finally, normalize the pressure-defect measures on the nested cylinders and pass to the compactified pressure-defect state space. Apply [lemma 9.63](#). If the atomic/activity alternative occurs, choose a realized support point with positive compact local activity. Then [theorem 9.66](#) routes that point to one of the previously listed rows (i)–(iv). If the genuinely diffuse alternative occurs, the nonzero diffuse pressure-defect measure itself is the diffuse pressure-only row, giving alternative (v). The alternatives exhaust all infinite nested pressure-only selections. $\square$

Lemma 9.68 (Diffuse pressure-only defects are not terminal). *A diffuse pressure-only defect produced by [lemma 9.67](#) cannot be a terminal retained branch.*

Proof. Let π be the normalized diffuse pressure-defect measure. By [lemma 9.63](#) to the normalized pressure-defect measures.

If the atomic/activity alternative occurs, choose a realized support profile with positive compact local activity furnished by [theorem 9.66](#). If that support profile has nonzero local Calderón–Zygmund source, then some compact test in the pressure-defect occupation chart detects a positive velocity-stress level. By the dyadic activity ladder, this is recorded at a positive local level $\eta_m > 0$. The fixed-annular stress dichotomy and the local CKN test then produce either a regular annular source or a velocity carrier at that positive level. The regular source is absorbed by the local pressure atlas; the carrier case contradicts pressure-only terminality.

If the atomic support profile is harmonic and remains non-affine after quotienting by time-dependent constants and every admissible local accelerating rebasement, then the far-field harmonic-tail row occurs.

If the realized atomic harmonic profile is spatially constant, it is removed by the ordinary pressure gauge. If it is spatially affine with nonzero linear part, then [lemma 9.69](#) either absorbs that affine forcing into the repaired accelerating-center modulation on a smaller local cylinder, reducing the harmonic remainder to a time-gauge term, or else the branch enters the affine-pressure/modulation failure row. Thus no terminal pressure-only branch can be supported by an untreated affine harmonic profile.

If the genuinely diffuse alternative occurs, then the pressure-defect measure itself is a visible diffuse pressure occupation coordinate. Apply [Theorem 10.86](#) to this visible pressure coordinate. The output is either a named pressure/gauge/rough-core/local-energy row or a retained comparable class. The retained comparable class is impossible because pressure-only descendants are not retained velocity classes.

Hence only named rows remain, and no diffuse pressure-only defect is terminal. $\square$

Lemma 9.69 (Affine harmonic pressure absorption by local accelerating rebasement). *Let (u, p) be a suitable weak solution on a local cylinder $Q := B_R(x_0) \times I$, and suppose that on a smaller cylinder $Q' := B_{R'}(x_0) \times I' \Subset Q$ the harmonic pressure component has the form*

$$H(x, t) = a(t) \cdot x + c(t)$$

with $a, c \in L^1_{\text{loc}}(I')$. Let $B \in W^{2,1}_{\text{loc}}(I'; \mathbb{R}^3)$. On every compact subinterval, $\dot{B}$ is bounded and absolutely continuous. Assume

$$\ddot{B}(t) = -a(t) \quad \text{a.e. on } I'.$$

Define on

$$\tilde{Q}' := \{(y, t) : (y + B(t), t) \in Q'\}$$

the rebased fields

$$v(y, t) := u(y + B(t), t) - \dot{B}(t), \quad q(y, t) := p(y + B(t), t) + \ddot{B}(t) \cdot y.$$

Then:

- (i) (v, q) is again a suitable weak solution on every compact $\tilde{Q}'' \Subset \tilde{Q}'$;
- (ii) the transformed harmonic component is spatially constant,

$$\tilde{H}(y, t) = a(t) \cdot B(t) + c(t),$$

hence removable by the ordinary pressure gauge modulo functions of time;

- (iii) all conclusions are local on compact subcylinders.

If the repaired gauge/modulation package on the branch does not admit this accelerating rebasement, then the first failed interface is the affine-pressure/modulation failure row.

Proof. Set $x = y + B(t)$. Since $\dot{B}$ is spatially constant, $\nabla_y \cdot v = \nabla_x \cdot u = 0$. Also

$$\partial_t v = \partial_t u(x, t) + \dot{B}(t) \cdot \nabla_x u(x, t) - \ddot{B}(t),$$

while

$$(v \cdot \nabla_y)v = (u - \dot{B}) \cdot \nabla_x u = u \cdot \nabla_x u - \dot{B} \cdot \nabla_x u.$$

Therefore

$$\partial_t v + (v \cdot \nabla)v - \Delta v = \partial_t u + (u \cdot \nabla)u - \Delta u - \ddot{B},$$

and, because

$$\nabla_y q(y, t) = \nabla_x p(x, t) + \ddot{B}(t),$$

the Navier–Stokes equation for (u, p) implies

$$\partial_t v + (v \cdot \nabla)v - \Delta v + \nabla q = 0$$

distributionally on $\tilde{Q}'$.

Decompose $p = p_{\text{loc}} + H$ on Q' , with $H(x, t) = a(t) \cdot x + c(t)$. Then

$$q(y, t) = p_{\text{loc}}(y + B(t), t) + a(t) \cdot (y + B(t)) + c(t) + \ddot{B}(t) \cdot y.$$

Since $\ddot{B} = -a$, the y -affine part cancels and

$$q(y, t) = p_{\text{loc}}(y + B(t), t) + a(t) \cdot B(t) + c(t).$$

Hence the transformed harmonic remainder is spatially constant. It is therefore removable by subtracting a function of time.

It remains to justify suitability. The transformed test

$$\psi(x, t) := \phi(x - B(t), t)$$

is compactly supported after shrinking the cylinder, but since $B \in W^{2,1}_{\text{loc}}$, it is in general only absolutely continuous in time. This is admissible by the same AC-test approximation used for repaired-gauge charts.

For completeness, choose smooth functions $B^\varepsilon \rightarrow B$ in $W^{2,1}$ on the compact time interval of $\tilde{Q}''$, with $\dot{B}^\varepsilon \rightarrow \dot{B}$ uniformly and $\ddot{B}^\varepsilon \rightarrow \ddot{B}$ in L^1 . Define

$$\psi^\varepsilon(x, t) := \phi(x - B^\varepsilon(t), t).$$

Then $\psi^\varepsilon \in C_c^\infty(Q')$ after shrinking once. Apply the local energy inequality for (u, p) to ψ^ε , perform the change of variables $x = y + B^\varepsilon(t)$, and rewrite the result in terms of

$$v^\varepsilon(y, t) := u(y + B^\varepsilon(t), t) - \dot{B}^\varepsilon(t), \quad q^\varepsilon(y, t) := p(y + B^\varepsilon(t), t) + \ddot{B}^\varepsilon(t) \cdot y.$$

The transport terms generated by the moving center cancel against the constant velocity shift and the affine pressure correction. Passing $\varepsilon \rightarrow 0$ uses the local finite-energy bounds, the uniform convergence of $\dot{B}^\varepsilon$, and the L^1 -convergence of $\ddot{B}^\varepsilon$. The limit is the local energy inequality for (v, q) on $\tilde{Q}''$. $\square$

Lemma 9.70 (Exterior pressure-only concentration is not terminal). *Let an exterior concentration branch on $A_{r,R}$ satisfy the pressure-carrier alternative of [lemma 9.49](#) and suppose that no velocity carrier occurs on any smaller exterior or separated-core cylinder after passing to subsequences. Then the pressure carrier is not an independent terminal exterior row. More precisely, after local pressure decomposition on nested cylinders, one of the following occurs:*

- (i) *the local Calderón–Zygmund pressure part forces a velocity-stress carrier on a slightly larger exterior cylinder;*
- (ii) *the far-field harmonic pressure tail produces a dyadic exterior occupation state, which is regular or gives a dyadic CKN carrier by [lemma 9.55](#);*
- (iii) *the pressure reconstruction or pressure gauge fails, which is the ordered non-exterior pressure row;*
- (iv) *the pressure-only selection continues through nested cylinders and is well-foundedly routed by [lemma 9.67](#), with its diffuse case excluded by [lemma 9.68](#).*

Proof. Work on a nested exterior cylinder and decompose the pressure into the local Calderón–Zygmund part generated by $\zeta u \otimes u$, the harmonic part, and a function of time. If the local Calderón–Zygmund part carries the pressure lower bound, Calderón–Zygmund boundedness and the quantitative converse in [lemma 9.21](#) force a nonzero velocity-stress carrier on a slightly larger exterior cylinder. Apply the local split [lemma 9.49](#) on that slightly larger cylinder. If the resulting carrier is velocity-supported, this is item (i).

Thus any persistent pressure lower bound must lie in the harmonic remainder. Apply [lemma 9.22](#) on nested cylinders. A nonvanishing harmonic remainder is either pressure-reconstruction/gauge failure, which is item (iii), or it is generated by a far-field tail. For the far-field tail, use [lemma 9.61](#). A fixed annular stress level is first tested by [lemma 9.60](#); a selected local CKN carrier is then split by [lemma 9.49](#). Its velocity-supported subcase gives item (i), while its pressure-supported subcase is handled by [lemma 9.67](#). A regular annular source is removable and cannot be an independent pressure-only exterior exit. An escaping dyadic stress occupation state is tested by [lemma 9.55](#); it is either regular on far dyadic annuli or gives a dyadic CKN carrier, which is item (ii). If the local pressure-only selection attempts to continue to smaller nested cylinders, apply [lemma 9.67](#). Alternatives (i)–(iv) there are the already named velocity-stress, harmonic tail, pressure-gauge, and scale-cascade rows. Alternative (v) is excluded by [lemma 9.68](#). Hence an exterior pressure-only concentration cannot be terminal. $\square$

Lemma 9.71 (Locality audit for exterior routing). *The exterior routing lemmas from [lemma 9.49](#) through [lemma 9.70](#) use only local annular CKN tests, finite-overlap covering, local pressure decomposition modulo functions of time, and the ordered local state-space stratification. They do not use a whole-space $L^3(\mathbb{R}^3)$ bound, a global critical profile decomposition, global Kato smallness, or an assumption that localized dissipation vanishes.*

Proof. The CKN split is a local identity on the selected cylinder. Positive-distance decoupling uses kernel smoothness on sets separated by a fixed physical distance, local L^1 bounds for the recorded source, and the ordered state-space rule that any positive compact-core mass must be

retained or routed before the exterior row is declared. The exterior and separated-core lemmas use the local scale-center compactification and the already stated separated-point and same-point reductions. The packing lemma uses only the Vitali finite-overlap property on a fixed annulus and explicitly avoids any cardinality bound from a global critical norm. Noncompact exterior routing is a countable dyadic application of the same local annular test. A positive fixed-annulus stress lower bound is not promoted directly to a singular carrier; it first passes through the local CKN dichotomy of [lemma 9.60](#). Exterior scale collapse is handled by rebasing the active local CKN carrier and then applying the local scale-collapse first-failure ledger. Pressure-only routing uses the local Calderón–Zygmund pressure atlas, the nested-ball harmonic-tail estimate, and the harmonic pressure routing lemma. None of these steps imports a global L^3 estimate, a whole-space profile theorem, global Kato perturbation, or unproved vanishing of localized dissipation. $\square$

Theorem 9.72 (Exterior concentration routes to closed or non-exterior rows). *Fix $0 < r < R < \infty$. If a Type II branch admits an exterior concentration branch on $A_{r,R}$, then that branch is not an independent terminal exterior exit. After passing to subsequences and applying the ordered local tests, one of the following holds:*

- (i) *the exterior contribution is regular or perturbative for the selected compact core, by [theorem 9.47](#) or [lemma 9.51](#);*
- (ii) *the branch enters the separated-core, multicore, or same-point cascade rows, which are closed by [theorem 11.6](#);*
- (iii) *the branch enters an exterior scale-collapse row, which after rebasing is handled by [lemma 9.59](#);*
- (iv) *the branch enters the noncompact exterior row, which routes to the preceding alternatives by [lemma 9.58](#);*
- (v) *the branch enters a pressure-only exterior row, which is converted to an exterior velocity/dyadic carrier, absorbed as a regular annular or dyadic source, or routed to an ordered non-exterior pressure row by [lemma 9.70](#).*

Consequently, once the non-exterior rows in [definition 11.35](#) are handled, exterior concentration and noncompact exterior concentration contribute no additional obstruction.

Proof. Start with the exterior concentration branch and apply the stratification [theorem 9.39](#). In the exterior-profile subcase, [lemma 9.51](#) gives perturbative decoupling from the selected compact core. In the separated-core or multicore subcase, [lemma 9.52](#) places the branch in the multibubble/cascade package, which is excluded by [theorem 11.6](#). If exterior carriers accumulate at one positive-distance point with strict scale separation, then [lemma 9.56](#) gives the same cascade closure.

If the stratification gives noncompact exterior behavior, first apply [lemma 9.57](#). The no-concentration dyadic case is regular and removable by the exterior regularity/removal theorems. The dyadic carrier case is routed by [lemma 9.58](#) to regular removal, separated/multicore, cascade, exterior scale-collapse, or pressure-only routing. The first three are already covered. The exterior scale-collapse subcase is rebased and tested by [lemma 9.59](#): finite-cost and retained compact or scale-rigid scale-collapse subcases are closed by the cited scale-collapse theorems, while any first failure is one of the ordinary non-exterior rows in the final ledger.

It remains only to treat the possibility that the CKN lower bound is pressure-only. The velocity-pressure split [lemma 9.49](#) either supplies a velocity carrier, returning to the preceding cases, or gives a pressure carrier. By [lemma 9.70](#), a pressure-only carrier either produces a velocity/dyadic carrier, is absorbed as a regular annular or dyadic source, exposes the ordered pressure row, or contradicts the pressure lower bound after all local and harmonic tails have been

removed. Thus every exterior concentration subtype has been routed to a closed structural row, regular removal, rebased scale-collapse, or one of the remaining non-exterior first-failure rows.

The proof is local by [lemma 9.71](#). Hence exterior concentration is discharged as an independent named exit without adding a global critical-norm, global profile, global Kato, or hidden dissipation assumption. $\square$

Corollary 9.73 (Exterior concentration exit discharged). *Exterior concentration and noncompact exterior concentration are not remaining independent named exits in the local Type II assembly. They are either regular and removable, closed by the structural multibubble/cascade or scale-collapse packages after rebasing, or assigned to the remaining non-exterior rows of [definition 11.35](#).*

Proof. This is [theorem 9.72](#). $\square$

10. CORRECTED MONOTONICITY AND SCALE-DRIFT COST CRITERIA

This section records the local analytic inputs for the localized cost-divergence exclusion argument on renormalized Type II branches. The statements below are local criteria, and the later windowwise closure theorem verifies the criteria on the canonical retained route. The statements are written entirely in the variables of the repaired gauge:

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(\tau)(V + y \cdot \nabla V) + b(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0.$$

All pressure terms are understood modulo functions of time.

10.1. Data for the Cost Criterion.

Definition 10.1 (Cost criterion data). A renormalized local Type II branch satisfies the cost criterion data if the following local analytic facts hold.

- (i) the branch lies in the local Type II concentration alternative;
- (ii) the physical suitable-solution local energy inequality gives finite localized energy on every compact selected window used by the branch, with the renormalized local energy inequality available after the repaired-gauge pullback; no whole-space energy bound is part of this data;
- (iii) the branch is written in the repaired-gauge variables (V, P, a, b) above, including the scale, center, pressure normalization, and modulation data;
- (iv) the localized cost is identified with

$$\tilde{\mathfrak{D}}_{R_0}(\tau) = \nu \int_{\mathbb{R}^3} |\nabla V(y, \tau)|^2 \phi_{R_0}(y) dy + a_+(\tau) \int_{\mathbb{R}^3} |V(y, \tau)|^2 \phi_{R_0}(y) dy,$$

or with an explicitly stated tail-comparable nonnegative replacement;

- (v) a localized energy identity or suitable-solution local energy inequality is available for the cutoff used in the cost;
- (vi) a corrected localized monotonicity inequality, in the sense of [definition 10.6](#), is available.

Definition 10.2 (Cost-divergence exclusion conclusion). A class of renormalized local Type II branches has the cost-divergence exclusion conclusion if every branch in the class satisfying [definition 10.1](#) is excluded from the local Type II class whenever the corresponding nonnegative localized cost diverges:

$$\int_{\tau_0}^{\infty} \mathcal{D}_{R_0}(\tau) d\tau = \infty.$$

This is a conclusion verified below on the canonical retained route, not an active input.

Lemma 10.3 (Concentration and localized energy). *If a branch lies in the local Type II concentration alternative and satisfies the physical suitable-solution local energy inequality on the compact cylinders selected by the branch, then it satisfies [definition 10.1](#)(i) and (ii).*

Proof. The local Type II concentration alternative supplies the selected concentration branch. Suitability gives the local energy inequality on each compact physical cylinder, and [lemma 10.59](#) transports it to the renormalized compact windows whenever the repaired-gauge chart is available. These are precisely items (i) and (ii); no global energy quantity is used. $\square$

Lemma 10.4 (Cost identification). *Assume the localized cost is the Navier–Stokes cost $\tilde{\mathfrak{D}}_{R_0}$. Then divergence of*

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau$$

is the cost divergence in [definition 10.2](#).

Proof. The active cost is assumed to be equal to $\tilde{\mathfrak{D}}_{R_0}$. Thus the cost appearing in the local Navier–Stokes estimates and the cost appearing in the local exclusion conclusion are the same nonnegative functional. The divergence condition is therefore identical in the two formulations. $\square$

Lemma 10.5 (Parabolic representation class). *If the repaired-gauge representation theorem applies to a Type II branch, then the branch has the parabolic transport-diffusion structure required in [definition 10.1](#).*

Proof. The repaired-gauge theorem supplies the renormalized orbit (V, P, a, b) , the scale and center, the pressure reconstruction modulo functions of time, and the modulation coefficients. The displayed equation is a parabolic Navier–Stokes equation with transport, dilation, and translation drift terms. Consequently the branch has the parabolic transport-diffusion structure required by the cost criterion data. $\square$

Definition 10.6 (Corrected localized monotonicity input). A renormalized branch has corrected localized monotonicity if there are a localized renormalized energy $\mathcal{E}_{R_0}$, a nonnegative cost $\mathcal{D}_{R_0}$, a constant $c_0 > 0$, and a nonnegative remainder $\mathcal{R}_{R_0}$ such that, on the tail of the branch,

$$\frac{d}{d\tau} \mathcal{E}_{R_0}(\tau) + c_0 \mathcal{D}_{R_0}(\tau) + \mathcal{R}_{R_0}(\tau) \leq 0$$

in distributions. If the localized energy identity has an integrable error B_{R_0} , the corrected energy is

$$\mathcal{E}_{R_0}^{\text{corr}}(\tau) = \mathcal{E}_{R_0}(\tau) + \int_{\tau}^{\infty} B_{R_0}(s) ds,$$

and the same inequality is required for $\mathcal{E}_{R_0}^{\text{corr}}$. Multiplication of a nonnegative cost by $c_0 > 0$ does not change whether its tail integral diverges.

Lemma 10.7 (Energy and Caccioppoli bounds alone are not monotonicity). *The global physical energy inequality and compact-cylinder Caccioppoli estimates do not, by themselves, imply the corrected localized monotonicity input of [definition 10.6](#).*

Proof. The global physical energy inequality controls total physical energy and physical dissipation, but it does not identify a localized renormalized energy at the selected Type II scale whose derivative is monotone after pressure flux, cutoff flux, transport, and modulation terms are included. The Caccioppoli estimate gives local spacetime H^1 control from the local energy inequality. It is an a priori estimate, not a corrected monotonicity formula. Therefore an additional absorption or tail-correction argument is needed. $\square$

Proposition 10.8 (Cost-divergence exclusion from the recorded conclusion). *Let a renormalized Type II branch satisfy [definition 10.1](#) and let its ambient class have the cost-divergence exclusion conclusion of [definition 10.2](#). If its nonnegative localized cost has divergent tail integral, then the branch is excluded from the local Type II class.*

Proof. [Lemma 10.3](#) gives the local Type II concentration and bounded-energy inputs. [Lemma 10.4](#) identifies the localized Navier–Stokes cost with the cost used in the exclusion criterion. [Lemma 10.5](#) gives the parabolic transport-diffusion class. The corrected monotonicity input is one item of [definition 10.1](#). Once the cost diverges, the conclusion recorded in [definition 10.2](#) gives the stated exclusion. $\square$

10.2. Fixed-Cutoff Monotonicity. Let $\phi = \phi_{R_0} \in C_c^\infty(\mathbb{R}^3)$, $0 \leq \phi \leq 1$, be the radial nonincreasing cutoff used in the localized cost. Define

$$\mathcal{E}_\phi(\tau) = \frac{1}{2} \int_{\mathbb{R}^3} |V(y, \tau)|^2 \phi(y) dy,$$

$$A_\phi(\tau) = \int_{\mathbb{R}^3} |\nabla V(y, \tau)|^2 \phi(y) dy, \quad M_\phi(\tau) = \int_{\mathbb{R}^3} |V(y, \tau)|^2 \phi(y) dy,$$

and

$$\tilde{\mathfrak{D}}_{R_0}(\tau) = \nu A_\phi(\tau) + a_+(\tau) M_\phi(\tau), \quad a_+(\tau) = \max\{a(\tau), 0\}.$$

The localized energy identity used below is

$$(13) \quad \begin{aligned} \frac{d}{d\tau} \mathcal{E}_\phi = & -\nu A_\phi + \frac{\nu}{2} \int |V|^2 \Delta \phi + \frac{1}{2} \int |V|^2 V \cdot \nabla \phi + \int P V \cdot \nabla \phi \\ & - \frac{a}{2} M_\phi - \frac{a}{2} \int |V|^2 y \cdot \nabla \phi - \frac{1}{2} b \cdot \int |V|^2 \nabla \phi. \end{aligned}$$

Definition 10.9 (Fixed-cutoff error density). Set $a_-(\tau) = \max\{-a(\tau), 0\}$ and define

$$\begin{aligned} B_{R_0}(\tau) := & \frac{\nu}{2} \left| \int |V|^2 \Delta \phi \right| + \frac{1}{2} \left| \int |V|^2 V \cdot \nabla \phi \right| + \left| \int P V \cdot \nabla \phi \right| \\ & + \frac{1}{2} |b(\tau)| \left| \int |V|^2 \nabla \phi \right| + \frac{1}{2} a_-(\tau) M_\phi(\tau) + \frac{1}{2} a_+(\tau) \int |V|^2 (-y \cdot \nabla \phi). \end{aligned}$$

The fixed-cutoff finite-error condition is

$$\int_{\tau_0}^{\infty} B_{R_0}(\tau) d\tau < \infty.$$

Remark 10.10 (Pressure normalization). The pressure term in B_{R_0} is independent of the representative of the pressure modulo functions of time. Replacing P by $P + c(\tau)$ changes the pressure flux by

$$c(\tau) \int V \cdot \nabla \phi = -c(\tau) \int \phi \nabla \cdot V = 0.$$

Theorem 10.11 (Fixed-cutoff corrected monotonicity). Assume (13) holds on (τ_0, ∞) and that $B_{R_0} \in L^1(\tau_0, \infty)$. Define

$$\mathcal{E}_\phi^{\text{corr}}(\tau) = \mathcal{E}_\phi(\tau) + \int_{\tau}^{\infty} B_{R_0}(s) ds.$$

Then, in distributions on (τ_0, ∞) ,

$$\frac{d}{d\tau} \mathcal{E}_\phi^{\text{corr}}(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R_0}(\tau) \leq 0.$$

Proof. Add

$$\frac{1}{2} \tilde{\mathfrak{D}}_{R_0} = \frac{\nu}{2} A_\phi + \frac{1}{2} a_+ M_\phi$$

to both sides of (13). The gradient term becomes

$$-\nu A_\phi + \frac{\nu}{2} A_\phi = -\frac{\nu}{2} A_\phi \leq 0.$$

The core scale term satisfies

$$-\frac{a}{2} M_\phi + \frac{1}{2} a_+ M_\phi = \frac{1}{2} a_- M_\phi \leq B_{R_0}(\tau).$$

Since ϕ is radial nonincreasing, one has

$$y \cdot \nabla \phi(y) \leq 0 \quad \text{for all } y \in \mathbb{R}^3.$$

Hence

$$-\frac{a}{2} \int |V|^2 y \cdot \nabla \phi \leq \frac{1}{2} a_+(\tau) \int |V|^2 (-y \cdot \nabla \phi),$$

because the a_- -part contributes a nonpositive term. Every remaining term in (13) is bounded above by its corresponding absolute-value contribution in B_{R_0} . Therefore

$$\frac{d}{d\tau} \mathcal{E}_\phi(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R_0}(\tau) \leq B_{R_0}(\tau)$$

in distributions. Since $B_{R_0} \in L^1(\tau_0, \infty)$, the correction tail is finite for every $\tau \geq \tau_0$, absolutely continuous on compact subintervals, and

$$\frac{d}{d\tau} \int_{\tau}^{\infty} B_{R_0}(s) ds = -B_{R_0}(\tau).$$

Adding this identity gives the claimed inequality. $\square$

Corollary 10.12 (Fixed-cutoff monotonicity input). *If the repaired-gauge representation, the cost identification, the localized energy identity, and the fixed-cutoff finite-error condition all hold, then [definition 10.6](#) holds with $\mathcal{D}_{R_0} = \tilde{\mathfrak{D}}_{R_0}$ and $c_0 = 1/2$.*

Proof. The repaired-gauge representation supplies (V, P, a, b) . The cost identification supplies $\tilde{\mathfrak{D}}_{R_0}$. The localized energy identity and finite-error condition allow [theorem 10.11](#) to be applied. $\square$

Remark 10.13 (Finite-tail subconditions). The fixed-cutoff finite-error condition is equivalent to the conjunction of tail integrability for the viscous cutoff term, the convective flux, the pressure flux, the center-drift cutoff term, the negative-scale term, and the positive scale-shell term:

$$\begin{aligned} & \left| \int |V|^2 \Delta \phi \right|, \quad \left| \int |V|^2 V \cdot \nabla \phi \right|, \quad \left| \int P V \cdot \nabla \phi \right|, \\ & |b| \left| \int |V|^2 \nabla \phi \right|, \quad a_- M_\phi, \quad a_+ \int |V|^2 (-y \cdot \nabla \phi). \end{aligned}$$

For radial nonincreasing cutoffs, only the positive part a_+ contributes to the scale-shell error. The negative part is favorable and does not belong to the residual finite-error burden. Uniform compact-window bounds do not by themselves give this global tail integrability on $[\tau_0, \infty)$.

Definition 10.14 (Fixed-cutoff cost hypotheses). A renormalized Type II branch satisfies the fixed-cutoff cost hypotheses if it has the local Type II concentration alternative, bounded physical energy, repaired-gauge representation, cost identification with $\tilde{\mathfrak{D}}_{R_0}$, divergent nonnegative localized cost, localized energy identity, and the finite-tail bound $\int_{\tau_0}^{\infty} B_{R_0} < \infty$.

Theorem 10.15 (Fixed-cutoff cost criterion). *Every branch satisfying the fixed-cutoff cost hypotheses is covered by [proposition 10.8](#).*

Proof. By [corollary 10.12](#), the finite-tail condition supplies the corrected monotonicity input. The remaining inputs are precisely the other items in [definition 10.1](#), together with cost divergence. The local exclusion step is later proved on the canonical windowwise schedule in [theorem 10.42](#). $\square$

Definition 10.16 (Compact cost-divergence replacement). A compact cost-divergence replacement is available when the compact Type II hypotheses of the preceding sections imply

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty,$$

and this divergence is the cost divergence used in [proposition 10.8](#). In particular, it may be supplied by the combination of classification of local alternatives, uniform critical tightness, local windowed H^1 control, and the cost identification above whenever those are the active hypotheses in the compact Type II theorem.

Theorem 10.17 (Compact cost-divergence criterion). *If a renormalized branch has all fixed-cutoff cost hypotheses except that cost divergence is supplied through the compact cost-divergence replacement of [definition 10.16](#), then [proposition 10.8](#) applies.*

Proof. The compact cost-divergence replacement supplies the missing divergent cost. Substituting this into the fixed-cutoff hypotheses gives the hypotheses of [theorem 10.15](#). $\square$

10.3. Moving-Cutoff Monotonicity. Let $\phi \in C_c^\infty(\mathbb{R}^3)$ be radial and nonincreasing, $0 \leq \phi \leq 1$, $\phi = 1$ on B_1 , and $\phi = 0$ outside B_2 . For $R(\tau) \geq R_0$, locally absolutely continuous and nondecreasing, set

$$\Phi(\tau, y) = \phi(y/R(\tau)), \quad A_{R(\tau)} = \{R(\tau) \leq |y| \leq 2R(\tau)\},$$

$$\mathcal{E}_R(\tau) = \frac{1}{2} \int |V(y, \tau)|^2 \Phi(\tau, y) dy,$$

and

$$\tilde{\mathfrak{D}}_{R(\tau)}(\tau) = \nu \int |\nabla V|^2 \Phi + a_+(\tau) \int |V|^2 \Phi.$$

Definition 10.18 (Moving-cutoff error density). Define

$$\begin{aligned} B_{\text{mov}}(\tau) := & \frac{1}{2} \left| \int |V|^2 \partial_\tau \Phi \right| + \frac{\nu}{2} \left| \int |V|^2 \Delta \Phi \right| + \frac{1}{2} \left| \int |V|^2 V \cdot \nabla \Phi \right| + \left| \int P V \cdot \nabla \Phi \right| \\ & + \frac{1}{2} |b(\tau)| \left| \int |V|^2 \nabla \Phi \right| + \frac{1}{2} a_-(\tau) \int |V|^2 \Phi + \frac{1}{2} a_+(\tau) \int |V|^2 (-y \cdot \nabla \Phi). \end{aligned}$$

The moving finite-error condition is

$$\int_{\tau_0}^{\infty} B_{\text{mov}}(\tau) d\tau < \infty.$$

Theorem 10.19 (Moving-cutoff corrected monotonicity). *Assume the localized energy identity is valid with the time-dependent cutoff Φ , and assume $B_{\text{mov}} \in L^1(\tau_0, \infty)$. Then*

$$\mathcal{E}_R^{\text{corr}}(\tau) = \mathcal{E}_R(\tau) + \int_{\tau}^{\infty} B_{\text{mov}}(s) ds$$

satisfies

$$\frac{d}{d\tau} \mathcal{E}_R^{\text{corr}}(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \leq 0$$

in distributions.

Proof. Use the localized energy identity with the time-dependent test function $\Phi(\tau, y)$. Compared with the fixed-cutoff identity, the only new term is

$$\frac{1}{2} \int |V|^2 \partial_\tau \Phi,$$

which comes from differentiating the cutoff inside $\mathcal{E}_R$. All other terms are the fixed-cutoff terms with ϕ replaced by Φ . Add $\frac{1}{2} \tilde{\mathfrak{D}}_{R(\tau)}$ to both sides. The gradient term leaves

$$-\frac{\nu}{2} \int |\nabla V|^2 \Phi \leq 0,$$

and the core scale term contributes

$$\frac{1}{2} a_- \int |V|^2 \Phi.$$

Since ϕ is radial nonincreasing, one has

$$y \cdot \nabla \Phi(\tau, y) \leq 0 \quad \text{for all } (\tau, y).$$

Therefore

$$-\frac{a}{2} \int |V|^2 y \cdot \nabla \Phi \leq \frac{1}{2} a_+(\tau) \int |V|^2 (-y \cdot \nabla \Phi),$$

because the a_- -part is nonpositive. Each remaining term is bounded above by the corresponding absolute-value term in B_{mov} . Hence

$$\frac{d}{d\tau} \mathcal{E}_R(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \leq B_{\text{mov}}(\tau).$$

The moving finite-error condition makes the correction tail finite and absolutely continuous on compact subintervals. Differentiating that tail subtracts B_{mov} , giving the stated inequality. $\square$

Definition 10.20 (Summable moving-annulus errors). The moving-annulus summability condition is the existence of a nondecreasing locally absolutely continuous $R(\tau) \rightarrow \infty$ such that each of the seven quantities

$$\begin{aligned} & \left| \int |V|^2 \partial_\tau \Phi \right|, \quad \left| \int |V|^2 \Delta \Phi \right|, \quad \left| \int |V|^2 V \cdot \nabla \Phi \right|, \quad \left| \int P V \cdot \nabla \Phi \right|, \\ & |b| \left| \int |V|^2 \nabla \Phi \right|, \quad a_- \int |V|^2 \Phi, \quad a_+ \int |V|^2 (-y \cdot \nabla \Phi) \end{aligned}$$

is integrable on $[\tau_0, \infty)$.

Lemma 10.21 (Summable annuli imply moving finite error). *If [definition 10.20](#) holds, then $B_{\text{mov}} \in L^1(\tau_0, \infty)$.*

Proof. The density B_{mov} is a finite positive linear combination of the seven quantities listed in [definition 10.20](#). If each is integrable, then their linear combination is integrable. $\square$

Definition 10.22 (Summable tightness schedule). A summable tightness schedule is a nondecreasing sequence $R_n \rightarrow \infty$ such that, on unit windows $I_n = [\tau_0 + n, \tau_0 + n + 1]$,

$$\sum_{n=0}^{\infty} \sup_{\tau \in I_n} \|V(\tau)\|_{L^3(A_{R_n})}^3 < \infty,$$

and the same schedule is compatible with the pressure-tail normalization used for P . This is stronger than uniform critical tightness, which permits large radii on each window but does not by itself give a monotone radius schedule with a summable series.

Definition 10.23 (Moving-cost compatibility). The moving localized cost is compatible with the cost-divergence criterion if

$$\tilde{\mathfrak{D}}_{R(\tau)}(\tau) = \nu \int |\nabla V|^2 \phi_{R(\tau)} + a_+(\tau) \int |V|^2 \phi_{R(\tau)}$$

is the localized cost appearing in [definition 10.2](#), or is compared to that cost by a positive constant on the tail so that cost divergence is unchanged.

Theorem 10.24 (Moving-cutoff replacement). *Assume the repaired-gauge representation, localized energy identity, moving-cost compatibility, and summable moving-annulus errors. Then the moving cutoff supplies the corrected monotonicity input needed in [proposition 10.8](#).*

Proof. [Lemma 10.21](#) gives $B_{\text{mov}} \in L^1(\tau_0, \infty)$. Applying [theorem 10.19](#) gives corrected monotonicity with cost $\tilde{\mathfrak{D}}_{R(\tau)}$. The moving-cost compatibility input identifies this cost, or a tail-comparable version of it, with the localized cost in the recorded cost-divergence conclusion. $\square$

10.4. Negative Scale Drift. Set

$$M_R(\tau) = \int |V(y, \tau)|^2 \Phi(\tau, y) dy.$$

The negative-scale contribution in the moving cutoff argument is

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau.$$

Definition 10.25 (Negative scale-drift alternatives). We say that the branch has finite negative drift if

$$\int_{\tau_0}^{\infty} a_-(\tau) d\tau < \infty.$$

It has bounded moving L^2 core mass if

$$\sup_{\tau \geq \tau_0} M_R(\tau) < \infty.$$

It has a moving L^2 core floor if there are $m_0 > 0$ and $\tau_1 \geq \tau_0$ such that

$$M_R(\tau) \geq m_0 \quad \text{for a.e. } \tau \geq \tau_1.$$

The scale-collapse drift obstruction is the divergent alternative

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau = \infty.$$

Lemma 10.26 (Finite negative drift controls the scale term). *If the branch has bounded moving L^2 core mass and finite negative drift, then*

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau < \infty.$$

Proof. Let $M_* = \sup_{\tau \geq \tau_0} M_R(\tau) < \infty$. Since $a_- \geq 0$,

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau \leq M_* \int_{\tau_0}^{\infty} a_-(\tau) d\tau < \infty.$$

□

Theorem 10.27 (Finite-or-infinite scale-negative alternative). *For every renormalized orbit for which a_- and M_R are measurable and nonnegative, precisely one of the following alternatives holds:*

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau < \infty, \quad \int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau = \infty.$$

Proof. The integral

$$I_R = \int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau$$

is an extended nonnegative number. Hence either $I_R < \infty$ or $I_R = \infty$, and the two cases are mutually exclusive. □

Lemma 10.28 (Logarithmic scale identity). *For the repaired-gauge convention*

$$\tau_t = \lambda(t)^{-2}, \quad a(\tau) = \lambda(t) \lambda_t(t),$$

one has

$$\int_{\tau_1}^{\tau_2} a(\tau) d\tau = \log \lambda(\tau_2) - \log \lambda(\tau_1)$$

whenever $\tau_0 \leq \tau_1 < \tau_2 < \infty$.

Proof. The chain rule gives

$$\frac{d}{d\tau} \log \lambda = \frac{\lambda_t}{\lambda} \frac{dt}{d\tau} = \frac{\lambda_t}{\lambda} \lambda^2 = \lambda \lambda_t = a.$$

Integrating over $[\tau_1, \tau_2]$ proves the identity. $\square$

Theorem 10.29 (Collapse with core floor forces scale-drift obstruction). *Assume*

$$\lambda(\tau) \rightarrow 0 \quad \text{as } \tau \rightarrow \infty,$$

and assume the moving L^2 core floor of [definition 10.25](#). Then

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau = \infty.$$

Proof. By [lemma 10.28](#),

$$\int_{\tau_1}^{\tau_2} a(\tau) d\tau = \log \lambda(\tau_2) - \log \lambda(\tau_1).$$

Since $\lambda(\tau) \rightarrow 0$, the right-hand side tends to $-\infty$ as $\tau_2 \rightarrow \infty$. If $\int_{\tau_1}^{\infty} a_-(\tau) d\tau < \infty$, then

$$\int_{\tau_1}^{\tau_2} a(\tau) d\tau = \int_{\tau_1}^{\tau_2} a_+(\tau) d\tau - \int_{\tau_1}^{\tau_2} a_-(\tau) d\tau \geq - \int_{\tau_1}^{\infty} a_-(\tau) d\tau > -\infty,$$

which contradicts the preceding limit. Thus $\int_{\tau_1}^{\infty} a_-(\tau) d\tau = \infty$.

By the moving L^2 core floor, $M_R(\tau) \geq m_0$ for a.e. $\tau \geq \tau_1$. Therefore

$$\int_{\tau_0}^{\infty} a_-(\tau) M_R(\tau) d\tau \geq m_0 \int_{\tau_1}^{\infty} a_-(\tau) d\tau = \infty.$$

$\square$

Corollary 10.30 (Scale-drift alternatives). *In the moving-cutoff setup, the negative-scale term has the following exhaustive alternatives. If bounded moving L^2 core mass and finite negative drift hold, then the negative-scale term is integrable. If the weighted negative-drift integral is infinite, the branch lies in the scale-collapse drift obstruction. If genuine collapse and a moving L^2 core floor also hold, then the obstruction is forced.*

Proof. The first case is [lemma 10.26](#). The exhaustive split is [theorem 10.27](#). The final assertion is [theorem 10.29](#). $\square$

10.5. Scale-Drift Cost Criteria.

Definition 10.31 (Scale-collapse drift cost). For a renormalized branch define

$$\mathfrak{C}_{\text{sc}}(\tau) = a_-(\tau) M_R(\tau), \quad M_R(\tau) = \int |V(y, \tau)|^2 \phi_{R(\tau)}(y) dy,$$

and

$$\mathcal{C}_{\text{sc}}[\tau_0, \infty) = \int_{\tau_0}^{\infty} \mathfrak{C}_{\text{sc}}(\tau) d\tau.$$

Thus the scale-collapse drift obstruction is precisely $\mathcal{C}_{\text{sc}}[\tau_0, \infty) = \infty$. The density is nonnegative and local in the renormalized core.

Definition 10.32 (Scale-collapse cost compatibility). The scale-collapse drift cost is compatible with the cost-divergence criterion when:

- (i) $\mathfrak{C}_{\text{sc}}$ is measurable, nonnegative, and locally integrable on finite τ -intervals for renormalized branches;

- (ii) divergence of $\mathcal{C}_{\text{sc}}[\tau_0, \infty)$ is one of the divergent-cost cases covered by the conclusion recorded in [definition 10.2](#);
- (iii) the resulting conclusion is the same local Type II exclusion conclusion used for the fixed localized cost.

Theorem 10.33 (Scale-collapse drift cost criterion). *If the scale-collapse drift obstruction holds and the cost $\mathfrak{C}_{\text{sc}}$ is compatible with the cost-divergence criterion, then the divergent-cost hypothesis of [proposition 10.8](#) is satisfied for the scale-collapse branch.*

Proof. By [definition 10.31](#), the obstruction is precisely

$$\mathcal{C}_{\text{sc}}[\tau_0, \infty) = \infty.$$

By [definition 10.32](#), this divergence is a nonnegative cost-divergence hypothesis with the same local conclusion. $\square$

Corollary 10.34 (Scale-collapse exclusion by drift cost). *Assume a renormalized local Type II branch has scale-collapse drift obstruction, compatible scale-collapse drift cost, and the remaining cost criterion data. If [definition 10.2](#) holds, then the branch is excluded from the local Type II class under consideration.*

Proof. [Theorem 10.33](#) supplies the divergent localized cost. Applying [proposition 10.8](#) gives the exclusion. $\square$

Definition 10.35 (Absolute scale-drift cost). Define

$$\mathfrak{C}_{\text{abs}}(\tau) = \nu \int |\nabla V|^2 \phi_{R(\tau)} + |a(\tau)| M_R(\tau).$$

The absolute scale-drift cost is compatible with the cost-divergence criterion if divergence of $\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{abs}}(\tau) d\tau$ is included as a Type II cost-divergence condition with the same local conclusion.

Lemma 10.36 (Scale-collapse drift forces absolute-cost divergence). *If the scale-collapse drift obstruction holds, then*

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{abs}}(\tau) d\tau = \infty.$$

Proof. Since

$$\mathfrak{C}_{\text{abs}}(\tau) = \nu \int |\nabla V|^2 \phi_{R(\tau)} + |a(\tau)| M_R(\tau) \geq a_-(\tau) M_R(\tau),$$

divergence of $\int a_- M_R$ forces divergence of $\int \mathfrak{C}_{\text{abs}}$. $\square$

Corollary 10.37 (Absolute scale-cost criterion). *If the scale-collapse drift obstruction holds and the absolute scale-drift cost is compatible with the cost-divergence criterion, then the divergent-cost hypothesis of [proposition 10.8](#) is satisfied.*

Proof. [Lemma 10.36](#) gives divergence of the absolute scale-drift cost, and compatibility makes this divergence a Type II cost-divergence condition. $\square$

Theorem 10.38 (Moving cutoff with scale-collapse alternative). *Assume the moving-cutoff setup and a renormalized local Type II branch.*

- (i) *If the negative-scale term is integrable and the remaining moving-annulus errors are summable, then the moving-cutoff corrected monotonicity estimate is available.*
- (ii) *If the scale-collapse drift obstruction holds and either the scale-collapse drift cost or the absolute scale-drift cost is compatible with the cost-divergence criterion, then the divergent-cost hypothesis is satisfied directly.*

- (iii) If, in addition, the remaining local data of [definition 10.1](#) and [definition 10.2](#) hold, then the branch is excluded.

Proof. The first branch is [theorem 10.24](#). The second branch is [theorem 10.33](#) or [corollary 10.37](#), depending on which compatible cost is used. The third branch is [proposition 10.8](#). $\square$

10.6. Cost-Divergence Exclusion on the Canonical Windowwise Schedule. The exclusion conclusion recorded in [definition 10.2](#) is, given the corrected monotonicity input that is already part of [definition 10.1](#), an arithmetic consequence of the nonnegativity of the localized renormalized energy together with the windowwise alternative argument of [corollaries 8.126](#) and [8.131](#). We make this implication explicit here so that no cost-divergence input remains in the assembly.

Lemma 10.39 (Nonnegativity of the corrected localized energy). *Let $\phi = \phi_{R_0} \in C_c^\infty(\mathbb{R}^3)$ be a radial nonnegative cutoff, and let $R(\tau)$ be either fixed at R_0 or the canonical windowwise interpolated radius of [lemma 8.116](#). Then $\mathcal{E}_\phi(\tau) = \frac{1}{2} \int |V(y, \tau)|^2 \phi(y) dy \geq 0$ and $\mathcal{E}_R(\tau) = \frac{1}{2} \int |V(y, \tau)|^2 \Phi(\tau, y) dy \geq 0$ for every τ , and the correction tails $\int_\tau^\infty B_{R_0}(s) ds$ and $\int_\tau^\infty B_{\text{mov}}(s) ds$ of [definitions 10.9](#) and [10.18](#) are nonnegative. Consequently the corrected energies $\mathcal{E}_\phi^{\text{corr}}$ and $\mathcal{E}_R^{\text{corr}}$ of [theorems 10.11](#) and [10.19](#) are nonnegative.*

Proof. Both $\mathcal{E}_\phi$ and $\mathcal{E}_R$ are integrals of $|V|^2$ against a nonnegative weight, hence nonnegative. Inspecting the seven summands of B_{R_0} ([definition 10.9](#)) and the seven summands of B_{mov} ([definition 10.18](#)): five are absolute values of integrals and are therefore nonnegative; the remaining two are $\frac{1}{2}a_-(\tau) \int |V|^2 \Phi$ and $\frac{1}{2}a_+(\tau) \int |V|^2 (-y \cdot \nabla \Phi)$ (and likewise with Φ replaced by ϕ). In the first, $a_- \geq 0$ and the integrand $|V|^2 \Phi \geq 0$, so the term is nonnegative. In the second, $a_+ \geq 0$, and $-y \cdot \nabla \Phi \geq 0$ because $\Phi(\tau, y) = \phi(y/R(\tau))$ with ϕ radial nonincreasing forces $\nabla \Phi(\tau, y)$ to point opposite to y ; the term is nonnegative. Hence $B_{R_0}, B_{\text{mov}} \geq 0$ pointwise, and their tail integrals are nonnegative. Adding a nonnegative tail to a nonnegative quantity preserves nonnegativity, so $\mathcal{E}_\phi^{\text{corr}}$ and $\mathcal{E}_R^{\text{corr}}$ are nonnegative. $\square$

Lemma 10.40 (Monotonicity versus cost divergence). *Let $\mathcal{E}^{\text{corr}}: [\tau_0, \infty) \rightarrow [0, \infty)$ be locally integrable with finite initial value $\mathcal{E}^{\text{corr}}(\tau_0) < \infty$, let $\mathcal{D}: [\tau_0, \infty) \rightarrow [0, \infty)$ be locally integrable, and let $c_0 > 0$. If*

$$\frac{d}{d\tau} \mathcal{E}^{\text{corr}}(\tau) + c_0 \mathcal{D}(\tau) \leq 0 \quad \text{in distributions on } (\tau_0, \infty),$$

then

$$\int_{\tau_0}^\infty \mathcal{D}(\tau) d\tau \leq \frac{1}{c_0} \mathcal{E}^{\text{corr}}(\tau_0) < \infty.$$

In particular, the simultaneous validity of corrected monotonicity in this sense and of $\int_{\tau_0}^\infty \mathcal{D} = \infty$ is contradictory. The same conclusion holds if the inequality is strengthened to the form $(d/d\tau)\mathcal{E}^{\text{corr}} + c_0 \mathcal{D} + \mathcal{R} \leq 0$ with $\mathcal{R} \geq 0$, as in [definition 10.6](#).

Proof. The distributional inequality $(\mathcal{E}^{\text{corr}})' + c_0 \mathcal{D} \leq 0$ on (τ_0, ∞) , with $\mathcal{D} \geq 0$ and $\mathcal{E}^{\text{corr}} \geq 0$ locally integrable, implies that $\mathcal{E}^{\text{corr}}$ admits a non-increasing right-continuous representative. Testing the inequality against a smooth approximation of $\mathbf{1}_{[\tau_0, \tau]}$ and passing to the limit using $\mathcal{D} \geq 0$ and $\mathcal{E}^{\text{corr}} \geq 0$ (so the boundary terms have a definite sign) yields

$$c_0 \int_{\tau_0}^\tau \mathcal{D}(\sigma) d\sigma \leq \mathcal{E}^{\text{corr}}(\tau_0) - \mathcal{E}^{\text{corr}}(\tau) \leq \mathcal{E}^{\text{corr}}(\tau_0).$$

Letting $\tau \rightarrow \infty$ and using monotone convergence ($\mathcal{D} \geq 0$) yields the displayed bound. If $\int \mathcal{D} = \infty$ the left-hand side is infinite while the right-hand side is finite, a contradiction. $\square$

Theorem 10.41 (Cost-divergence/corrected-monotonicity dichotomy on the canonical schedule). *Let ω be a renormalized local Type II branch in the pre-carrier validated compact-cost regime of [definition 8.87](#). Equip ω with the canonical windowwise interpolated cutoff $\Phi(\tau, y) = \phi(y/R(\tau))$ of [lemma 8.116](#), and let $\tilde{\mathfrak{D}}_{R(\tau)}$ be the moving localized cost of [definition 10.23](#). Assume in addition the localized energy identity for the moving cutoff $\Phi(\tau, \cdot)$ of [definition 10.1\(v\)](#) (supplied, e.g., by [lemma 8.89](#) on the regime). If*

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) d\tau = \infty,$$

then the following sharpened dichotomy holds: at least one of the three components

$$J_1, \quad J_6, \quad J_7$$

fails to be in $L^1(\tau_1, \infty)$ for every $\tau_1 \geq \tau_0$. More precisely, exactly one of the following two mutually exclusive alternatives holds:

- (D1) *at least one of J_1, J_6, J_7 fails to be in $L^1(\tau_1, \infty)$ for every $\tau_1 \geq \tau_0$;*
- (D2) *every $J_k, k \in \{1, \dots, 7\}$, belongs to $L^1(\tau_0, \infty)$.*

Moreover, case (D2) is incompatible with cost divergence, so divergence of the cost forces case (D1), i.e. failure of one of J_1, J_6, J_7 .

Proof. The moving error density of [definition 10.18](#) satisfies the identity

$$B_{\text{mov}}(\tau) = \frac{1}{2}J_1(\tau) + \frac{\nu}{2}J_2(\tau) + \frac{1}{2}J_3(\tau) + J_4(\tau) + \frac{1}{2}J_5(\tau) + \frac{1}{2}J_6(\tau) + \frac{1}{2}J_7(\tau),$$

with J_k as in [definition 8.100](#), because the seven nonnegative terms in B_{mov} coincide with $J_1, \dots, J_7$ up to the listed multiplicative constants, all of which are strictly positive (recall $\nu > 0$ is fixed throughout). In particular $B_{\text{mov}} \in L^1(\tau_0, \infty)$ if and only if every $J_k \in L^1(\tau_0, \infty)$.

By [lemma 8.116](#), on the canonical windowwise schedule one has automatically

$$J_2, J_3, J_4, J_5 \in L^1(\tau_0, \infty).$$

Therefore the only J_k that can fail to be in L^1 on a tail are J_1, J_6, J_7 . We obtain the following dichotomy: either at least one of J_1, J_6, J_7 fails to be in $L^1(\tau_1, \infty)$ for every $\tau_1 \geq \tau_0$ (case (D1)), or each of J_1, J_6, J_7 is integrable on $[\tau_0, \infty)$ and hence (combined with the automatic integrability of J_2, J_3, J_4, J_5) every J_k is integrable on $[\tau_0, \infty)$ (case (D2)).

In case (D2), [lemma 10.21](#) gives $B_{\text{mov}} \in L^1(\tau_0, \infty)$, and the localized energy identity of [definition 10.1\(v\)](#) together with [theorem 10.19](#) produces

$$\frac{d}{d\tau} \mathcal{E}_R^{\text{corr}}(\tau) + \frac{1}{2} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) \leq 0 \quad \text{in distributions on } (\tau_0, \infty),$$

with $c_0 = \frac{1}{2}$ in the sense of [definition 10.6](#). We verify the finite initial value locally. By [lemma 8.96](#), after a harmless terminal restriction the initial time is a compact-cylinder local-energy time for the finite ball supporting $\Phi(\tau_0, \cdot)$. Hence

$$\mathcal{E}_R(\tau_0) = \frac{1}{2} \int |V(\tau_0)|^2 \Phi(\tau_0, \cdot) dy < \infty.$$

Combined with the case-(D2) integrability $\|B_{\text{mov}}\|_{L^1(\tau_0, \infty)} < \infty$,

$$\mathcal{E}_R^{\text{corr}}(\tau_0) = \mathcal{E}_R(\tau_0) + \int_{\tau_0}^{\infty} B_{\text{mov}}(s) ds < \infty.$$

By [lemma 10.39](#), $\mathcal{E}_R^{\text{corr}} \geq 0$, so [lemma 10.40](#) yields

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) d\tau \leq 2 \mathcal{E}_R^{\text{corr}}(\tau_0) < \infty,$$

contradicting the cost divergence. Hence case (D2) cannot occur, and case (D1) must hold; that is, at least one of J_1, J_6, J_7 fails to be in L^1 on every later tail. $\square$

Theorem 10.42 (CostPhase 6: cost-divergence exclusion on the canonical windowwise schedule). *Fix the canonical windowwise moving family of lemma 8.116 on the pre-carrier validated compact-cost regime of definition 8.87, and assume the localized energy identity for the moving cutoff (lemma 8.89). Assume further that (R1) failure of $J_6 \in L^1(\tau_1, \infty)$ on every later tail $[\tau_1, \infty)$ places the branch into the negative scale-drift alternative (vi) of definition 8.110, in the sense of theorem 8.133 (this is the same negative-drift routing statement used in that theorem, realized through definition 10.25 and corollary 10.30); and (R2) the named adjacent strata of theorem 8.92 together with the exterior/translation/positive-scale/negative-drift alternatives (iii)–(vi) of definition 8.110 are closed by their assigned local arguments in the ambient Type II class. Then for every renormalized local Type II branch ω in that regime satisfying definition 10.1, divergence of the localized cost*

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R(\tau)}(\tau) d\tau = \infty$$

excludes ω from the local Type II class. Equivalently, the negative-drift alternative (R1) together with the adjacent-stratum closures (R2) imply definition 10.2 as a theorem on the canonical windowwise schedule.

Proof. Let ω be a candidate branch. By theorem 10.41, case (D2) is impossible, so case (D1) occurs: at least one of J_1, J_6, J_7 fails to be in $L^1(\tau_1, \infty)$ for every $\tau_1 \geq \tau_0$.

Apply the alternative statement assigned to that index:

- if J_1 fails, the windowwise transition/leakage alternative of corollary 8.126 (proved unconditionally on the regime via theorem 8.125) places ω into the exterior/leakage alternative (iii) of definition 8.110, entering the named critical-tightness failure stratum of theorem 8.92;
- if J_6 fails, hypothesis (R1) places ω into the negative scale-drift alternative (vi);
- if J_7 fails, the windowwise positive-scale alternative of corollary 8.131 (proved unconditionally via the argument theorem 8.128 and proposition 5.42) places ω into the positive-scale alternative (v).

By hypothesis (R2), each of these named alternatives is closed by an already proved argument in the ambient Type II class. Therefore ω is excluded. This proves the conclusion of definition 10.2 on the canonical windowwise schedule under premises (R1) and (R2); the next corollary verifies these premises on the physical local Type II branch. $\square$

Corollary 10.43 (CostPhase 6: unconditional physical cost-divergence exclusion). *The hypotheses (R1)–(R2) of theorem 10.42 on the canonical windowwise schedule coincide, after combining with the unconditional windowwise alternatives corollaries 8.126 and 8.131, with the verified negative-drift routing statement for J_6 and the adjacent-stratum closure of corollary 8.134. In particular, on any ambient Type II class on which the conclusion of corollary 8.134 (compact-cost adjacent-stratum closure) is established, definition 10.2 is a theorem and may be removed from the hypothesis lists of definitions 10.138 and 10.140. For the physical local Type II class considered in the final assembly, this adjacent closure is established by corollary 8.139.*

Proof. The phase 5 windowwise transition/leakage alternative corollary 8.126 and windowwise positive-scale alternative corollary 8.131 are unconditional on the regime. Combined with the negative-drift alternative for J_6 verified on physical local Type II branches by lemma 8.137, and closed as an independent terminal row by theorem 10.118, the alternative component (R1) of theorem 10.42 is satisfied. The adjacent-stratum closure component (R2) is exactly the closure

statement of [corollary 8.134](#); in the physical local Type II class it is supplied by [corollary 8.139](#). Therefore [theorem 10.42](#) produces exclusion without any further hypothesis, and [definition 10.2](#) is a theorem in that ambient class. $\square$

Remark 10.44 (Status of the cost-divergence conclusion). The proof has three independent ingredients:

- (i) the trivial arithmetic consequence [lemma 10.40](#) that corrected localized monotonicity together with finite initial corrected energy and $\mathcal{D} \geq 0$ forces $\int \mathcal{D} < \infty$;
- (ii) the moving error-density identity $B_{\text{mov}} = \sum_k c_k J_k$ of [definitions 8.100](#) and [10.18](#) (positive linear combination), combined with the canonical schedule integrability $J_2, J_3, J_4, J_5 \in L^1(\tau_0, \infty)$ of [lemma 8.116](#);
- (iii) the windowwise alternative for J_1 and J_7 ([corollaries 8.126](#) and [8.131](#), proved unconditionally on the regime) plus the negative-drift alternative for J_6 supplied by [theorem 8.133](#).

Thus [definition 10.2](#) is not an independent analytic input. Its content is the conjunction of the trivial monotonicity-vs-divergence arithmetic and the windowwise alternative argument that is already used to close adjacent strata in [corollary 8.134](#).

Lemma 10.45 (CostPhase 7: acyclic order for cost-divergence closure). *In the physical local Type II assembly, the cost-divergence closure is used only after the local carrier and adjacent-routing rows have been verified. The dependency order is:*

- (i) *the suitable local energy inequality and repaired-gauge pullback supply the local carrier subidentities on compact windows;*
- (ii) *the canonical windowwise moving family supplies the unconditional transition and positive-scale routing alternatives [corollaries 8.126](#) and [8.131](#);*
- (iii) *nonintegrability of the negative scale-drift component J_6 is routed by the local negative-drift branch [lemma 8.137](#);*
- (iv) *the preceding rows give ambient adjacent-stratum closure through [corollary 8.139](#);*
- (v) *only then is [theorem 10.42](#) applied to remove a retained branch with divergent validated localized cost;*
- (vi) *terminal retained-state elimination is applied after the adjacent and cost rows have been routed or closed.*

No statement in (i)–(iv) invokes the cost-divergence exclusion step or the final retained Type II closure theorem.

Proof. Item (i) is [lemmas 8.89](#) and [10.59](#). Item (ii) is exactly the pair of windowwise routing corollaries cited there. Item (iii) is [lemma 8.137](#), whose proof routes J_6 failure through the local scale-drift, scale-collapse, compact scale-collapse, and scale-rigid packages, not through cost-divergence closure. Item (iv) is [corollary 8.139](#); its proof cites the windowwise routing and negative-drift branch but not the cost-divergence theorem. The proof of [theorem 10.42](#) then uses those already verified rows as inputs, which is item (v). Item (vi) is an ordering convention for the later final assembly: terminal retained-state elimination is allowed to use the fact that the named adjacent and cost rows have already been routed or closed, but none of the preceding rows uses terminal retained-state elimination as an input. $\square$

10.7. Absolutely Continuous Repaired-Gauge Charts and Renormalized Equations.

The following statements isolate the parts of the repaired-gauge representation that follow from an absolutely continuous concentration chart and an absolutely continuous repaired-gauge path. They do not prove extraction of the chart from arbitrary data and do not prove solvability of the gauge constraints.

Definition 10.46 (Absolutely continuous concentration chart). Let u, p solve the Navier–Stokes equations on $\mathbb{R}^3 \times (t_0, T^*)$,

$$\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u, \quad \nabla \cdot u = 0.$$

An absolutely continuous concentration chart consists of absolutely continuous functions

$$x_c : (t_0, T^*) \rightarrow \mathbb{R}^3, \quad \lambda : (t_0, T^*) \rightarrow (0, \infty),$$

and a strictly increasing absolutely continuous time ρ such that

$$\rho_t = \lambda^{-2} \quad \text{a.e.},$$

$$y = \frac{x - x_c(t)}{\lambda(t)}, \quad u(x, t) = \lambda(t)^{-1} U(y, \rho(t)),$$

and

$$p(x, t) = \lambda(t)^{-2} Q(y, \rho(t)) + c(t)$$

for some scalar function $c(t)$. The chart is assumed to have the local integrability required for the distributional chain rule on compact sets in the y -variables. Define

$$A(\rho(t)) := \lambda(t) \lambda_t(t), \quad B(\rho(t)) := \lambda(t) x'_c(t),$$

or equivalently

$$A = \partial_\rho \log \lambda, \quad B = \lambda^{-1} \partial_\rho x_c,$$

where λ and x_c are regarded as functions of ρ .

Lemma 10.47 (Initial renormalized equation). *Assume Definition 10.46. Then U, Q satisfy*

$$\partial_\rho U + (U \cdot \nabla)U + \nabla Q = \nu \Delta U + A(\rho)(U + y \cdot \nabla U) + B(\rho) \cdot \nabla U, \quad \nabla \cdot U = 0$$

distributionally on compact y -sets.

Proof. For smooth functions the computation is pointwise. The distributional case is obtained by testing against compactly supported functions and using the assumed chart chain rule. The spatial derivatives are

$$\nabla_x u = \lambda^{-2} \nabla_y U, \quad \Delta_x u = \lambda^{-3} \Delta_y U, \quad (u \cdot \nabla_x)u = \lambda^{-3} (U \cdot \nabla_y)U.$$

Since

$$y_t = -\frac{x'_c}{\lambda} - \frac{\lambda_t}{\lambda} y, \quad \rho_t = \lambda^{-2},$$

one has

$$\partial_t u = \lambda^{-3} \partial_\rho U - \lambda^{-2} \lambda_t (U + y \cdot \nabla U) - \lambda^{-2} x'_c \cdot \nabla U.$$

Also

$$\nabla_x p = \lambda^{-3} \nabla_y Q.$$

Substitution into the physical Navier–Stokes equation and multiplication by λ^3 gives

$$\partial_\rho U - \lambda \lambda_t (U + y \cdot \nabla U) - \lambda x'_c \cdot \nabla U + (U \cdot \nabla)U + \nabla Q = \nu \Delta U.$$

Moving the two chart-velocity terms to the right gives the displayed equation. Finally,

$$\nabla_x \cdot u = \lambda^{-2} \nabla_y \cdot U,$$

so $\nabla_y \cdot U = 0$. □

Definition 10.48 (Absolutely continuous repaired-gauge path). An absolutely continuous repaired-gauge path for the initial chart consists, on every compact final time interval, of absolutely continuous functions

$$\mu(\tau) > 0, \quad q(\tau) \in \mathbb{R}^3, \quad \rho = \rho(\tau),$$

such that

$$\begin{aligned} \rho_\tau &= \mu(\tau)^2 \quad \text{a.e.}, \\ V(Y, \tau) &= \mu(\tau)U(\mu(\tau)Y + q(\tau), \rho(\tau)), \end{aligned}$$

and

$$P(Y, \tau) = \mu(\tau)^2 Q(\mu(\tau)Y + q(\tau), \rho(\tau)).$$

The scale and centering constraints defining the repaired gauge are assumed to hold along this path. This hypothesis assumes the path has already been selected; it does not prove solvability of the gauge constraints.

Lemma 10.49 (Physical chart after gauge repair). *Assume Definitions 10.46 and 10.48. Let $t = t(\rho(\tau))$, and define*

$$\Lambda(\tau) := \lambda(t(\rho(\tau)))\mu(\tau), \quad X(\tau) := x_c(t(\rho(\tau))) + \lambda(t(\rho(\tau)))q(\tau).$$

Then, on every compact final time interval, Λ and X are absolutely continuous, $\Lambda > 0$, and

$$\frac{d\tau}{dt} = \Lambda(\tau)^{-2}.$$

Moreover,

$$x = X(\tau) + \Lambda(\tau)Y$$

is the final physical chart and

$$u(x, t) = \Lambda(\tau)^{-1}V(Y, \tau), \quad p(x, t) = \Lambda(\tau)^{-2}P(Y, \tau) + c(t).$$

Proof. The intermediate time ρ is strictly increasing because $\rho_t = \lambda^{-2} > 0$ a.e.; hence it has an absolutely continuous inverse on compact subintervals of its range. The compositions $\lambda(t(\rho(\tau)))$ and $x_c(t(\rho(\tau)))$ are therefore absolutely continuous on compact final windows. Products and sums with the absolutely continuous functions μ, q are absolutely continuous, so Λ and X are absolutely continuous.

The identity $\rho_\tau = \mu^2$ implies

$$\frac{d\tau}{d\rho} = \mu^{-2}.$$

Since $d\rho/dt = \lambda^{-2}$,

$$\frac{d\tau}{dt} = \frac{d\tau}{d\rho} \frac{d\rho}{dt} = \mu^{-2} \lambda^{-2} = (\lambda\mu)^{-2} = \Lambda^{-2}.$$

If $x = X + \Lambda Y$, then

$$x = x_c + \lambda q + \lambda\mu Y,$$

and hence the intermediate coordinate is

$$y = \frac{x - x_c}{\lambda} = q + \mu Y.$$

Therefore

$$u(x, t) = \lambda^{-1}U(q + \mu Y, \rho) = \lambda^{-1}\mu^{-1}V(Y, \tau) = \Lambda^{-1}V(Y, \tau),$$

and similarly

$$p(x, t) = \lambda^{-2}Q(q + \mu Y, \rho) + c(t) = \lambda^{-2}\mu^{-2}P(Y, \tau) + c(t) = \Lambda^{-2}P(Y, \tau) + c(t).$$

□

Lemma 10.50 (Canonical final-chart differential identities). *Let $X, \Lambda \in AC_{\text{loc}}(I)$, with $\Lambda > 0$, and let*

$$\frac{d\tau}{dt} = \Lambda(\tau)^{-2}, \quad Y = \frac{x - X(\tau)}{\Lambda(\tau)}, \quad u = \Lambda^{-1}V, \quad p = \Lambda^{-2}P + c(t).$$

Set $a = \partial_\tau \log \Lambda$ and $b = \Lambda^{-1}X_\tau$. For smooth fields, at fixed (x, t) ,

$$\begin{aligned} \partial_t u &= \Lambda^{-3}(\partial_\tau V - (b + aY) \cdot \nabla_Y V - aV), \\ \nabla_x u &= \Lambda^{-2} \nabla_Y V, \quad \Delta_x u = \Lambda^{-3} \Delta_Y V, \quad \nabla_x p = \Lambda^{-3} \nabla_Y P, \end{aligned}$$

and

$$\nabla_x \cdot u = \Lambda^{-2} \nabla_Y \cdot V.$$

The same identities hold distributionally for absolutely continuous charts after testing on compact subsets of the chart image.

Proof. At fixed x ,

$$Y_\tau = -\Lambda^{-1}X_\tau - \frac{\Lambda_\tau}{\Lambda}Y = -(b + aY), \quad \tau_t = \Lambda^{-2}.$$

Consequently

$$\begin{aligned} \partial_t u &= \tau_t \partial_\tau (\Lambda^{-1}V(Y, \tau)) \\ &= \Lambda^{-3}(\partial_\tau V - (b + aY) \cdot \nabla_Y V - aV). \end{aligned}$$

Since $\nabla_x Y = \Lambda^{-1}I$, the spatial derivative, Laplacian, pressure-gradient, and divergence identities follow directly. The term $c(t)$ disappears under the spatial gradient. For an absolutely continuous chart these calculations are justified in distributions by pulling back a compactly supported smooth test function, approximating the chart parameters by smooth absolutely continuous approximants on the compact time interval, and passing to the limit. The chart derivatives converge in L^1 , while Λ and Λ^{-1} converge uniformly there; hence every term in the displayed identities converges in the corresponding distributional pairing. $\square$

Lemma 10.51 (Modulation coefficients in repaired gauge). *Define*

$$a(\tau) := \partial_\tau \log \Lambda(\tau), \quad b(\tau) := \Lambda(\tau)^{-1}X_\tau(\tau),$$

where derivatives are understood a.e. on compact final windows. Under the hypotheses of Lemma 10.49,

$$a(\tau) = \mu(\tau)^2 A(\rho(\tau)) + \frac{\mu_\tau(\tau)}{\mu(\tau)}$$

and

$$b(\tau) = \mu(\tau)B(\rho(\tau)) + \mu(\tau)A(\rho(\tau))q(\tau) + \frac{q_\tau(\tau)}{\mu(\tau)}.$$

In particular $a, b \in L^1_{\text{loc}}$ on every compact final time interval.

Proof. Since $\Lambda = \lambda(\rho(\tau))\mu(\tau)$,

$$\partial_\tau \log \Lambda = \rho_\tau \partial_\rho \log \lambda + \partial_\tau \log \mu = \mu^2 A + \frac{\mu_\tau}{\mu}.$$

This proves the formula for a . Next

$$X = x_c(\rho(\tau)) + \lambda(\rho(\tau))q(\tau).$$

Using

$$\partial_\rho x_c = \lambda B, \quad \partial_\rho \lambda = \lambda A, \quad \rho_\tau = \mu^2,$$

we get

$$X_\tau = \mu^2 \lambda B + \mu^2 \lambda A q + \lambda q_\tau.$$

Division by $\Lambda = \lambda\mu$ gives

$$\Lambda^{-1}X_\tau = \mu B + \mu Aq + \frac{q_\tau}{\mu}.$$

Local integrability follows from absolute continuity and positivity of μ . On every compact final interval, μ is bounded above and below away from zero, q is bounded, and $\mu_\tau, q_\tau \in L^1$. The A -terms are locally integrable because

$$\int |A(\rho(\tau))|\mu(\tau)^2 d\tau = \int |A(\rho)| d\rho,$$

and the extra factors μ, μ^{-1} , and q are bounded on the same compact window. Likewise, since μ is bounded away from zero,

$$\int |B(\rho(\tau))|\mu(\tau) d\tau \leq C \int |B(\rho(\tau))|\mu(\tau)^2 d\tau = C \int |B(\rho)| d\rho,$$

so the B -term in b is locally integrable. $\square$

Theorem 10.52 (Canonical renormalized equation in the repaired gauge). *Let X, Λ, V, P be a final chart satisfying the hypotheses of [lemma 10.50](#). If u, p solve the Navier–Stokes equations distributionally on the physical image, then V, P satisfy*

$$(14) \quad \partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(\tau)(V + Y \cdot \nabla V) + b(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0$$

distributionally on compact Y -sets. Conversely, every distributional solution of this displayed equation pulls back through the same chart to a distributional Navier–Stokes solution on the physical image. In particular, the conclusion applies to the two-stage absolutely continuous chart of [definitions 10.46](#) and [10.48](#), with a, b given by [lemma 10.51](#).

Proof. By [lemma 10.50](#), every term in the physical equation has the common factor Λ^{-3} . Removing that factor gives

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(V + Y \cdot \nabla V) + b \cdot \nabla V.$$

The divergence identity in the same lemma gives $\nabla \cdot V = 0$. The argument is distributional because the identities there were proved by test function pullback. Running those identities in reverse proves the converse. For the two-stage chart, [lemma 10.49](#) supplies the final physical representation and [lemma 10.51](#) identifies its coefficients with the displayed a, b . $\square$

Lemma 10.53 (Pressure reconstruction under the charts). *Assume the physical pressure satisfies*

$$-\Delta_x p = \partial_i \partial_j (u_i u_j)$$

in distributions. Then the intermediate pressure satisfies

$$-\Delta_y Q = \partial_i \partial_j (U_i U_j)$$

modulo functions of ρ , and the final pressure satisfies

$$-\Delta_Y P = \partial_i \partial_j (V_i V_j)$$

modulo functions of τ .

Proof. The additive function $c(t)$ in

$$p = \lambda^{-2}Q + c(t)$$

has zero spatial derivatives. Since

$$\Delta_x p = \lambda^{-4} \Delta_y Q, \quad \partial_{x_i} \partial_{x_j} (u_i u_j) = \lambda^{-4} \partial_{y_i} \partial_{y_j} (U_i U_j),$$

the intermediate pressure equation follows. For the repaired variables,

$$P(Y) = \mu^2 Q(\mu Y + q), \quad V(Y) = \mu U(\mu Y + q).$$

Therefore

$$-\Delta_Y P = \mu^4(-\Delta_z Q)(z) = \mu^4 \partial_{z_i} \partial_{z_j} (U_i U_j)(z) = \partial_{Y_i} \partial_{Y_j} (V_i V_j)(Y).$$

Adding functions of time to Q or P does not change the identity. $\square$

Lemma 10.54 (Scaling of local L^r norms). *For every $1 \leq r < \infty$ for which the norms are finite,*

$$\|V(\tau)\|_{L_Y^r}^r = \Lambda(\tau)^{r-3} \|u(t(\tau))\|_{L_x^r}^r.$$

In particular,

$$\|V(\tau)\|_{L_Y^3} = \|u(t(\tau))\|_{L_x^3}.$$

Thus membership in $L^3(\mathbb{R}^3)$, finiteness, infiniteness, and positivity are invariant under the final chart.

Proof. Using

$$u(X + \Lambda Y, t) = \Lambda^{-1} V(Y, \tau), \quad dx = \Lambda^3 dY,$$

we obtain

$$\|u(t)\|_{L_x^r}^r = \int |\Lambda^{-1} V(Y, \tau)|^r \Lambda^3 dY = \Lambda^{3-r} \|V(\tau)\|_{L_Y^r}^r.$$

Rearranging gives the stated formula. The case $r = 3$ is scale invariant. $\square$

Lemma 10.55 (Regularity of scale, center, and modulation). *If λ, x_c, μ, q are absolutely continuous and $\mu > 0$, then on compact final time intervals the final scale Λ , center X , and coefficients a, b are absolutely continuous or locally integrable as specified above: $\Lambda, X \in W_{\text{loc}}^{1,1}$ and*

$$a, b \in L_{\text{loc}}^1.$$

Proof. This follows from Lemmas 10.49 and 10.51. Positivity of μ and compactness of the time window give a positive lower bound for μ , so the divisions by μ in the formulas for a, b are legitimate. $\square$

Theorem 10.56 (Consequences of absolutely continuous repaired-gauge charts). *Assume an absolutely continuous concentration chart and an absolutely continuous repaired-gauge path in the sense of Definitions 10.46 and 10.48. Then the following analytic conclusions hold:*

- (i) *the initial and repaired renormalized Navier–Stokes equations hold distributionally on compact renormalized sets;*
- (ii) *the final physical chart satisfies*

$$u(x, t) = \Lambda(\tau)^{-1} V(Y, \tau), \quad p(x, t) = \Lambda(\tau)^{-2} P(Y, \tau) + c(t), \quad \frac{d\tau}{dt} = \Lambda^{-2};$$

- (iii) *pressure reconstruction holds modulo functions of time;*
- (iv) *the final modulation coefficients are the functions a, b in Lemma 10.51;*
- (v) *the critical L^3 norm is invariant under the final chart;*
- (vi) *the final scale, center, and modulation coefficients have the local absolute-continuity and integrability stated in Lemma 10.55.*

Proof. Lemma 10.47 gives the initial renormalized equation. Lemma 10.49 and theorem 10.52 give the final chart and the final repaired-gauge equation with the stated coefficients. Lemma 10.53 gives pressure reconstruction modulo functions of time. Lemma 10.54 gives critical norm invariance, and Lemma 10.55 gives the regularity inheritance. $\square$

Corollary 10.57 (Representation inputs supplied by the repaired gauge). *Once the absolutely continuous concentration chart and absolutely continuous repaired-gauge path are available, pressure reconstruction, modulation coefficient realization, local $W^{1,1}$ regularity of the chart, and critical L^3 -invariance are analytic consequences rather than separate inputs. The representation inputs that must be supplied before applying this corollary are:*

- (i) *extraction of an absolutely continuous concentration chart from the selected Type II branch;*
- (ii) *solvability of the repaired scale and centering equations on the admissible profile class;*
- (iii) *selection of the repaired scale and center as an absolutely continuous path satisfying $\rho_\tau = \mu^2$;*
- (iv) *terminal admissibility of the final scale and time, including $\tau \rightarrow \infty$ and preservation of the Type II core scale.*

Proof. This is Theorem 10.56. The enumerated items are the complementary analytic inputs that must be available before that theorem can be applied: one must produce the initial concentration chart, solve the repaired gauge equations, select the resulting scale and center absolutely continuously, and verify that the resulting final chart is terminally admissible. $\square$

10.8. Local Caccioppoli Estimate for Absolutely Continuous Suitable Branches.

Definition 10.58 (Suitable repaired-gauge Type II branch). A suitable repaired-gauge Type II branch is a repaired-gauge Type II branch for which:

- (i) an absolutely continuous concentration chart and an absolutely continuous repaired-gauge path are available;
- (ii) the conclusions of Theorem 10.56 hold, including pressure reconstruction and final modulation coefficients;
- (iii) the physical branch (u, p) is a local suitable weak solution on the terminal interval:

$$u \in L_t^\infty L_{\text{loc}}^2 \cap L_t^2 H_{\text{loc}}^1, \quad p \in L_{\text{loc}}^{3/2}, \quad \nabla \cdot u = 0,$$

the Navier–Stokes equations hold distributionally, and the physical local energy inequality holds for every nonnegative smooth compactly supported test function.

Lemma 10.59 (Local energy inequality under absolutely continuous Type II charts). *Let (u, p) be physically suitable on a compact physical cylinder, and let (V, P) be obtained from an absolutely continuous final chart*

$$u(x, t) = \Lambda(s)^{-1} V(Y, s), \quad p(x, t) = \Lambda(s)^{-2} P(Y, s) + c(t), \quad Y = \frac{x - X(s)}{\Lambda(s)},$$

with $s = \tau$, $dt = \Lambda(s)^2 ds$, $\Lambda > 0$, and

$$X, \Lambda \in W_{\text{loc}}^{1,1}.$$

Then (V, P) satisfies the renormalized local energy inequality on compact (Y, s) -cylinders:

$$\begin{aligned} & \int_{\mathbb{R}^3} \frac{|V(s_2)|^2}{2} \phi(s_2) dY + \nu \int_{s_1}^{s_2} \int_{\mathbb{R}^3} |\nabla V|^2 \phi dY ds \\ & \leq \int_{\mathbb{R}^3} \frac{|V(s_1)|^2}{2} \phi(s_1) dY + \int_{s_1}^{s_2} \int_{\mathbb{R}^3} \frac{|V|^2}{2} (\partial_s \phi + \nu \Delta \phi) dY ds \\ & \quad + \int_{s_1}^{s_2} \int_{\mathbb{R}^3} \left(\frac{|V|^2}{2} + P \right) V \cdot \nabla \phi dY ds \\ & \quad + \int_{s_1}^{s_2} \int_{\mathbb{R}^3} a(s) \frac{|V|^2}{2} (-\phi - Y \cdot \nabla \phi) dY ds \\ & \quad - \int_{s_1}^{s_2} \int_{\mathbb{R}^3} \frac{|V|^2}{2} b(s) \cdot \nabla \phi dY ds, \end{aligned}$$

for every nonnegative $\phi \in C^\infty([s_1, s_2] \times \mathbb{R}^3)$ that is compactly supported in the spatial variable. Here a, b are the final modulation coefficients.

Proof. First suppose X, Λ are C^1 . Use the physical test function

$$\psi(x, t) = \Lambda(s)^{-1} \phi \left(\frac{x - X(s)}{\Lambda(s)}, s \right), \quad t = t(s),$$

where $dt = \Lambda^2 ds$. The spatial support of ϕ is compact, and Λ is bounded above and below on the selected compact time interval, so ψ is an admissible compactly supported test function in the physical local energy inequality.

Let

$$Y = \frac{x - X(s)}{\Lambda(s)}.$$

At fixed x ,

$$\frac{dY}{ds} = -\frac{X_s}{\Lambda} - \frac{\Lambda_s}{\Lambda} Y.$$

Consequently,

$$\nabla_x \psi = \Lambda^{-2} \nabla_Y \phi, \quad \Delta_x \psi = \Lambda^{-3} \Delta_Y \phi,$$

and, since $ds/dt = \Lambda^{-2}$,

$$\partial_t \psi = \Lambda^{-3} \left[\partial_s \phi - \frac{\Lambda_s}{\Lambda} (\phi + Y \cdot \nabla \phi) - \frac{X_s}{\Lambda} \cdot \nabla \phi \right].$$

Changing variables

$$x = X(s) + \Lambda(s)Y, \quad dt = \Lambda^2 ds,$$

and using

$$u = \Lambda^{-1} V, \quad p = \Lambda^{-2} P + c(t),$$

each term in the physical local energy inequality has the following renormalized form. The endpoint and dissipation terms transform as

$$\int \frac{|u(t_i)|^2}{2} \psi(t_i) dx = \int \frac{|V(s_i)|^2}{2} \phi(s_i) dY,$$

and

$$\nu \int_{t_1}^{t_2} \int |\nabla_x u|^2 \psi dx dt = \nu \int_{s_1}^{s_2} \int |\nabla_Y V|^2 \phi dY ds.$$

The test-function derivative term becomes

$$\begin{aligned} & \int_{t_1}^{t_2} \int \frac{|u|^2}{2} (\partial_t \psi + \nu \Delta_x \psi) dx dt \\ &= \int_{s_1}^{s_2} \int \frac{|V|^2}{2} [\partial_s \phi + \nu \Delta_Y \phi - a(\phi + Y \cdot \nabla \phi) - b \cdot \nabla \phi] dY ds, \end{aligned}$$

where

$$a = \frac{\Lambda_s}{\Lambda}, \quad b = \frac{X_s}{\Lambda}.$$

The transport and pressure term becomes

$$\begin{aligned} & \int_{t_1}^{t_2} \int \left(\frac{|u|^2}{2} + p \right) u \cdot \nabla_x \psi dx dt \\ &= \int_{s_1}^{s_2} \int_{\mathbb{R}^3} \left(\frac{|V|^2}{2} + P \right) V \cdot \nabla_Y \phi dY ds + \int_{s_1}^{s_2} \Lambda(s)^2 c(t(s)) \left(\int_{\mathbb{R}^3} V \cdot \nabla_Y \phi dY \right) ds. \end{aligned}$$

The pressure constant drops out because

$$\int_{\mathbb{R}^3} V \cdot \nabla_Y \phi dY = - \int_{\mathbb{R}^3} \phi \nabla_Y \cdot V dY = 0.$$

Combining the transformed terms gives the displayed renormalized local energy inequality.

For $X, \Lambda \in W_{\text{loc}}^{1,1}$, the same identity follows at the level of test functions. On the selected compact time interval, Λ has positive lower and finite upper bounds. The pulled-back test

$$\psi(x, t) = \Lambda(s)^{-1} \phi((x - X(s))/\Lambda(s), s)$$

is compactly supported, nonnegative, bounded, has bounded spatial derivatives on the compact cylinder, and has time derivative in $L_t^1 L_x^\infty$ because $X', \Lambda' \in L^1$ and ϕ is smooth compactly supported. The physical local energy inequality, initially stated for C_c^∞ tests, extends to this class by standard mollification of the test function: the suitability classes give the local integrability needed for every term in the inequality, and the mollified tests keep support inside a slightly larger compact cylinder. Applying the extended inequality to ψ and changing variables gives the displayed renormalized inequality. The only derivatives of the chart that appear are the L^1 -functions X', Λ' , hence the modulation terms are locally integrable against local $|V|^2$. $\square$

Lemma 10.60 (Pressure and modulation for suitable repaired-gauge branches). *For a suitable repaired-gauge branch, the pressure identity*

$$-\Delta P = \partial_i \partial_j (V_i V_j)$$

holds modulo functions of τ , and the final coefficients a, b in the renormalized equation are measurable and locally integrable on compact renormalized windows.

Proof. This is the pressure and modulation part of Theorem 10.56. Pressure pullback comes from the physical pressure equation and the chart/gauge scaling. The final coefficients are forced by the absolutely continuous scale-translation gauge transform:

$$a = \mu^2 A + \frac{\mu\tau}{\mu}, \quad b = \mu B + \mu A q + \frac{q\tau}{\mu}.$$

Since the chart and gauge parameters are absolutely continuous on compact windows, these coefficients are locally integrable. $\square$

Theorem 10.61 (Canonical local Caccioppoli estimate for repaired gauges). *Every suitable repaired-gauge Type II branch satisfies the local renormalized Caccioppoli estimate on each compact repaired-gauge cylinder. More precisely, let $I' \Subset I$, $0 < r < R$, and choose $\eta \in C_c^\infty(I)$, $\zeta \in C_c^\infty(B_R)$, with*

$$0 \leq \eta, \zeta \leq 1, \quad \eta = 1 \text{ on } I', \quad \zeta = 1 \text{ on } B_r.$$

Then

$$\begin{aligned} (15) \quad & \text{ess sup}_{\tau \in I'} \int_{B_r} \frac{|V(\tau)|^2}{2} dY + \nu \int_{I'} \int_{B_r} |\nabla V|^2 dY d\tau \\ & \leq C \int_I \int_{B_R} \frac{|V|^2}{2} (|\eta'| \zeta^2 + \nu \eta |\Delta(\zeta^2)|) dY d\tau \\ & \quad + C \int_I \int_{B_R} \left(\frac{|V|^2}{2} + |P - (P)_{B_R}(\tau)| \right) |V| \eta |\nabla(\zeta^2)| dY d\tau \\ & \quad + C \int_I |a(\tau)| \int_{B_R} \frac{|V|^2}{2} \eta (\zeta^2 + |Y| |\nabla(\zeta^2)|) dY d\tau \\ & \quad + C \int_I |b(\tau)| \int_{B_R} \frac{|V|^2}{2} \eta |\nabla(\zeta^2)| dY d\tau. \end{aligned}$$

Here C is absolute. In the bounded-modulation regime, if [proposition 6.1](#) holds on $Q_{2R,J}$, $J' \Subset J$, $a, b \in L^\infty(J)$, and $\mathcal{C}_J(2R) + \mathcal{D}_J(2R) < \infty$, then

$$\mathcal{A}_{J'}(R) + \mathcal{H}_{J'}(R) \leq C_*(1 + \mathcal{C}_J(2R) + \mathcal{D}_J(2R)),$$

where

$$C_* = C_*(R, |J|, \nu, \|a\|_{L^\infty(J)}, \|b\|_{L^\infty(J)}, \text{dist}(J', \partial J)).$$

No local Caffarelli–Kohn–Nirenberg smallness, global windowed H^1 , critical L^3 -norm, tightness, finite energy at infinity, or whole-space compactness conclusion is asserted.

Proof. By the suitable repaired-gauge hypotheses and Theorem 10.56, the branch has an absolutely continuous final chart, an admissible absolutely continuous repaired-gauge path, pressure reconstruction, and final modulation coefficients. Lemma 10.59 pulls the physical local energy inequality to the repaired variables. Lemma 10.60 supplies the pressure and coefficient identities needed to interpret the renormalized inequality in the same PDE class used by the local windowed estimates.

Choose a standard compactly supported Caccioppoli cutoff

$$\phi(Y, \tau) = \eta(\tau)\zeta(Y)^2.$$

Let $I = (\tau_1, \tau_2)$, let $I' \Subset I$, and let $0 < r < R$. Take $0 \leq \eta \leq 1$, $\eta \in C_c^\infty(I)$, with $\eta \equiv 1$ on I' , and take $0 \leq \zeta \leq 1$, $\zeta \in C_c^\infty(B_R)$, with $\zeta \equiv 1$ on B_r . Applying Lemma 10.59 on (τ_1, σ) and then taking the essential supremum over $\sigma \in I'$ controls the endpoint term. Applying the same inequality on the full interval I controls the dissipation term. A function of τ may be subtracted from P , since

$$\int_{B_R} (P)_{B_R}(\tau) V \cdot \nabla(\zeta^2) dY = 0$$

by incompressibility and the compact support of ζ . Adding the two estimates therefore gives exactly (15). Since the cutoff derivatives are bounded by constants depending only on I, I', r, R , this is the local Caccioppoli estimate on the compact renormalized cylinder.

No global estimate is used in this step. On the compact cylinder under consideration, the absolutely continuous chart has

$$0 < \Lambda_- \leq \Lambda(\tau) \leq \Lambda_+ < \infty.$$

Pulling back physical suitability therefore gives, for every compact spatial set K ,

$$V \in L_\tau^\infty L_Y^2(K) \cap L_\tau^2 H_Y^1(K), \quad P \in L_{\tau,Y}^{3/2}(K).$$

Local Sobolev interpolation on the finite cylinder gives

$$V \in L_{\tau,Y}^{10/3}(K),$$

hence $V \in L_{\tau,Y}^3(K)$. Thus $PV \cdot \nabla \phi$ is integrable by Hölder's inequality. The representation identities give

$$a, b \in L_{\text{loc}}^1(d\tau),$$

and

$$\tau \mapsto \int_K |V(Y, \tau)|^2 dY$$

is locally essentially bounded; therefore the scale and translation modulation terms in the pulled-back local energy inequality are integrable on the same compact cylinder. This proves the local Caccioppoli estimate, but not a global windowed H^1 or critical L^3 -norm bound.

Under the additional bounded-modulation and finite $\mathcal{C} + \mathcal{D}$ hypotheses in the statement, lemma 6.4 applies to the same canonical local energy inequality and gives the displayed quantitative $\mathcal{A} + \mathcal{H}$ estimate, with precisely the stated dependence of the constant. $\square$

Corollary 10.62 (No separate Caccioppoli obstruction for suitable absolutely continuous branches). *On suitable repaired-gauge branches, failure of local Caccioppoli regularity, failure of pressure reconstruction, and failure of local gauge regularity are not separate obstructions to the local H^1 analysis. The additional inputs used when invoking the uniform windowed H^1 estimates are the bounded critical L^3 and modulation bounds on the retained window; if either input fails, the ordered local state-space test records the corresponding active-core, rough-core, modulation, or scale-drift row before the uniform estimate is invoked.*

Proof. [Theorem 10.61](#) gives local Caccioppoli regularity. [Theorem 10.56](#) gives pressure reconstruction and the final modulation coefficients from the final chart and absolutely continuous gauge path. Thus the listed failures cannot occur separately on suitable repaired-gauge branches. The bounded critical norm and bounded modulation used by the uniform windowed H^1 estimate are therefore retained-window inputs or routed first failures, not independent Caccioppoli obstructions. $\square$

10.9. Caccioppoli Regularity Criterion.

Definition 10.63 (Suitable physical hypotheses). Let u, p be a local suitable weak solution of the three-dimensional Navier–Stokes equations on $\mathbb{R}^3 \times (0, T^*)$. Thus

$$u \in L_t^\infty L_{\text{loc}}^2 \cap L_t^2 H_{\text{loc}}^1, \quad p \in L_{\text{loc}}^{3/2}, \quad \nabla \cdot u = 0,$$

the equations hold distributionally, and for every nonnegative $\psi \in C_c^\infty(\mathbb{R}^3 \times (0, T^*))$ and a.e. $0 < t_1 < t_2 < T^*$,

$$\begin{aligned} & \int_{\mathbb{R}^3} \frac{|u(t_2)|^2}{2} \psi(t_2) dx + \nu \int_{t_1}^{t_2} \int_{\mathbb{R}^3} |\nabla u|^2 \psi dx dt \\ & \leq \int_{\mathbb{R}^3} \frac{|u(t_1)|^2}{2} \psi(t_1) dx + \int_{t_1}^{t_2} \int_{\mathbb{R}^3} \frac{|u|^2}{2} (\partial_t \psi + \nu \Delta \psi) dx dt \\ & \quad + \int_{t_1}^{t_2} \int_{\mathbb{R}^3} \left(\frac{|u|^2}{2} + p \right) u \cdot \nabla \psi dx dt. \end{aligned}$$

Lemma 10.64 (Renormalized local energy inequality). *Assume the suitable physical hypotheses of [definition 10.63](#). Let*

$$u(x, t) = \lambda(t)^{-1} V \left(\frac{x - x_c(t)}{\lambda(t)}, \tau(t) \right), \quad \frac{d\tau}{dt} = \lambda(t)^{-2},$$

with

$$P(y, \tau) = \lambda(t)^2 p(x_c(t) + \lambda(t)y, t),$$

where $\lambda > 0$, x_c , and τ are C^1 on compact windows. Assume that V, P satisfy the repaired-gauge equation

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(\tau)(V + y \cdot \nabla V) + b(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0.$$

Then, for every nonnegative $\phi \in C_c^\infty(\mathbb{R}^3 \times (\tau_0, \infty))$ and a.e. $\tau_1 < \tau_2$ in the corresponding renormalized time interval,

$$\begin{aligned} & \int_{\mathbb{R}^3} \frac{|V(\tau_2)|^2}{2} \phi(\tau_2) dy + \nu \int_{\tau_1}^{\tau_2} \int_{\mathbb{R}^3} |\nabla V|^2 \phi dy d\tau \\ & \leq \int_{\mathbb{R}^3} \frac{|V(\tau_1)|^2}{2} \phi(\tau_1) dy + \int_{\tau_1}^{\tau_2} \int_{\mathbb{R}^3} \frac{|V|^2}{2} (\partial_\tau \phi + \nu \Delta \phi) dy d\tau \\ & \quad + \int_{\tau_1}^{\tau_2} \int_{\mathbb{R}^3} \left(\frac{|V|^2}{2} + P \right) V \cdot \nabla \phi dy d\tau \\ & \quad + \int_{\tau_1}^{\tau_2} \int_{\mathbb{R}^3} a(\tau) \frac{|V|^2}{2} (-\phi - y \cdot \nabla \phi) dy d\tau \\ & \quad - \int_{\tau_1}^{\tau_2} \int_{\mathbb{R}^3} \frac{|V|^2}{2} b(\tau) \cdot \nabla \phi dy d\tau. \end{aligned}$$

Proof. For smooth solutions, multiply the repaired-gauge equation by $V\phi$ and integrate by parts. Put $e = |V|^2/2$. Since $\nabla \cdot V = 0$,

$$V \cdot (V \cdot \nabla V) = V \cdot \nabla e = \nabla \cdot (eV), \quad V \cdot \nabla P = \nabla \cdot (PV),$$

and

$$V \cdot (y \cdot \nabla V) = y \cdot \nabla e, \quad V \cdot (b \cdot \nabla V) = b \cdot \nabla e.$$

The scale contribution is

$$\int a(2e + y \cdot \nabla e) \phi dy = \int ae(-\phi - y \cdot \nabla \phi) dy,$$

because

$$\int y \cdot \nabla e \phi dy = - \int e \nabla \cdot (y\phi) dy = - \int e(3\phi + y \cdot \nabla \phi) dy.$$

The translation contribution is

$$\int b \cdot \nabla e \phi dy = - \int e b \cdot \nabla \phi dy,$$

and the viscous term gives

$$\int \nu \Delta V \cdot V \phi dy = -\nu \int |\nabla V|^2 \phi dy + \nu \int e \Delta \phi dy.$$

Combining these identities gives the displayed equality in the smooth case.

For suitable weak solutions, apply [lemma 10.59](#) with

$$\Lambda = \lambda, \quad X = x_c, \quad s = \tau.$$

Its physical test function is

$$\psi(x, t) = \lambda(t)^{-1} \phi \left(\frac{x - x_c(t)}{\lambda(t)}, \tau(t) \right),$$

which is exactly the critically scaled pullback needed for the endpoint and dissipation terms. Since the present C^1 case is a special case of the absolutely continuous pullback lemma, the same inequality follows directly. $\square$

Lemma 10.65 (Caccioppoli estimate from the renormalized local energy inequality). *Assume [lemma 10.64](#) on a compact renormalized cylinder $B_{2R} \times I$, integrated modulation on I , and local bounds*

$$V \in L^\infty(I; L^2(B_{2R})) \cap L^3(I \times B_{2R}), \quad P \in L^{3/2}(I \times B_{2R}).$$

Then the compact-cylinder renormalized Caccioppoli estimate used in [lemma 10.105](#) is justified.

Proof. Choose a standard nonnegative cutoff

$$\phi(y, \tau) = \eta(\tau)\zeta(y)^2,$$

where $\zeta \equiv 1$ on the inner ball and ζ is supported in the outer ball. Insert this test function in [lemma 10.64](#). The left-hand side contains

$$\nu \int |\nabla V|^2 \eta \zeta^2.$$

Every term on the right-hand side is finite under the stated hypotheses. The term

$$\int |V|^2 |\partial_\tau \phi + \Delta \phi|$$

is controlled by $V \in L^\infty(I; L^2(B_{2R}))$ and finite measure. The convective term

$$\int |V|^3 |\nabla \phi|$$

is controlled by the local spacetime L^3 bound. For the pressure term, choose a time-dependent constant $c(\tau)$. Since

$$\int c(\tau) V \cdot \nabla \phi \, dy = - \int c(\tau) \phi \nabla \cdot V \, dy = 0,$$

Hölder's inequality controls

$$\int |P - c(\tau)| |V| |\nabla \phi|$$

by $P - c(\tau) \in L^{3/2}(I \times B_{2R})$ and $V \in L^3(I \times B_{2R})$. The modulation terms are bounded by

$$\int_I (|a(\tau)| + |b(\tau)|) \, d\tau$$

times the local $L_I^\infty L^2$ -mass. These estimates give the renormalized Caccioppoli estimate on compact cylinders. $\square$

Proposition 10.66 (Caccioppoli regularity criterion). *Assume the suitable physical hypotheses, the repaired-gauge representation, pressure reconstruction modulo functions of time, and C^1 scale and center functions on compact renormalized windows. Then the compact-cylinder Caccioppoli regularity hypothesis used in the local windowed H^1 estimate holds.*

Proof. The repaired-gauge representation writes the suitable weak solution in the variables V, P, a, b . [Lemma 10.64](#) gives the renormalized local energy inequality, and [lemma 10.65](#) shows that this inequality justifies the renormalized Caccioppoli estimate on every compact cylinder. $\square$

Corollary 10.67 (Caccioppoli estimates give windowed control under boundedness hypotheses). *Assume the suitable physical hypotheses, the repaired-gauge representation, pressure reconstruction, a bounded compact spacetime L^3 amount on compact cylinders, the local $L_\tau^\infty L_y^2$ bound from the suitable local energy package, and bounded integrated modulation on unit windows. Then local windowed H^1 control follows from [lemma 10.105](#). Thus failure of the Caccioppoli regularity hypothesis is not an independent rough-core alternative under these hypotheses.*

Proof. [Proposition 10.66](#) supplies the Caccioppoli regularity hypothesis. The bounded compact spacetime L^3 amount, the local $L_\tau^\infty L_y^2$ control, pressure reconstruction, and the windowwise integrated modulation bound are the remaining local inputs used in [lemma 10.105](#). Therefore the local windowed H^1 estimate holds. $\square$

10.10. Terminal Local State-Space Compactness.

Definition 10.68 (Admissible retained terminal decomposition and completeness conclusion). Let $Q_{2R,I} := B_{2R} \times I$ be a compact terminal-frame cylinder and let

$$V_n = U_n + S_n$$

be the local decomposition into the retained terminal component U_n and the nonretained residual S_n . The critical-profile interface is expressed through the following local notions.

- (i) The decomposition is an admissible retained terminal decomposition if, on every $Q_{2R,I}$, the represented fields satisfy the repaired-gauge equation, local energy inequality, pressure reconstruction modulo functions of time, bounded modulation, and the compact-cylinder estimates

$$\sup_n \left(\|V_n\|_{L^\infty(I; L^3(B_{2R}))} + \|V_n\|_{L^\infty(I; L^2(B_{2R}))} + \|V_n\|_{L^2(I; H^1(B_{2R}))} + \|P_n - c_{n,2R}\|_{L^{3/2}(Q_{2R,I})} \right) < \infty,$$

and the same local velocity bounds hold for U_n , hence also for $S_n = V_n - U_n$ after enlarging the constant. The represented time derivatives $\partial_s V_n$ are bounded in $L^1(I; H^{-m}(B_{2R}))$ for some m large enough for the Aubin–Lions compactness argument. No time compactness of U_n or S_n is assumed unless a componentwise time-derivative row is explicitly present; otherwise those components are recorded by weak local velocity coordinates, selected-time residual measures, and mixed-interaction stress occupation coordinates.

- (ii) A positive residual local label is assigned as follows. If, after passing to a subsequence,

$$\limsup_{n \rightarrow \infty} \|S_n\|_{L^\infty(I; L^3(B_R))} > 0,$$

then a selected subwindow and subcylinder carry an augmented local occupation state for the tuple $(V_n, P_n, U_n, S_n, a_n, b_n)$, together with the local residual concentration measures $|S_n|^3 dy$ on selected time slices, whose residual coordinate or residual concentration measure is nonzero. The first closed label of this augmented state is read in the ordered local validation list of [lemma 8.73](#), with the compact local Type II row refined by [theorem 11.18](#).

- (iii) The local terminal completeness conclusion is the statement that, if all named adjacent alternatives are closed by their local arguments and comparable retained classes have been grouped before selecting the terminal frame, then no nonretained residual with a positive residual local label remains on the retained terminal branch. Equivalently,

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_R))$$

on every compact terminal cylinder of the retained branch.

No global L^3 profile decomposition, global heat-flow norm, or global Kato estimate is part of this statement.

Remark 10.69 (Terminal $L^\infty(I; L^3)$ coordinates are interface inputs). The $L^\infty(I; L^3)$ bounds in [definition 10.68](#)(i) are part of the local terminal decomposition interface inherited from the admissible perturbative family of [definition 5.16](#). They are not produced by the spacetime compactness package of [definition 7.27](#). When a physical branch fails this pointwise critical control on a compact terminal cylinder, the branch is routed first by [lemma 10.78](#) or by an earlier active-family failure, rather than being treated inside the perturbative terminal decomposition layer.

Local proof of terminal state-space completeness.

Lemma 10.70 (Interior-time selection on a compact terminal window). *Let $I = [\tau_-, \tau_+]$ be a compact time-interval in the terminal-frame coordinates, and let $\tau_n \in I$ be a sequence of selected times for which $\|S_n(\tau_n)\|_{L^3(B_R)} \geq \delta > 0$. Then, after passing to a subsequence, exactly one of the following two admissible selections holds:*

- (i) *there is a compact subinterval $J \Subset I$ such that $\tau_n \in J$ for all large n ;*
- (ii) *the selected times approach an endpoint of I , and the terminal window is replaced by the admissible recentered windows*

$$I_n^* := \tau_n + J_0, \quad J_0 = [-\sigma, \sigma],$$

with fixed $\sigma > 0$, written in coordinates centered at τ_n . In these recentered coordinates the selected time is $0 \in J_0^\circ$ and

$$\|S_n(0)\|_{L^3(B_R)} \geq \delta.$$

No lower bound is transferred from one time to another.

Proof. After passing to a subsequence, $\tau_n \rightarrow \tau^* \in I$. If $\tau^* \in (\tau_-, \tau_+)$, choose $J \Subset I$ with $\tau^* \in J^\circ$. Then $\tau_n \in J$ for all large n , giving (i).

If $\tau^* \in \partial I$, we do not move the lower bound to a nearby time. Instead we use the admissibility of the terminal validation family under finite covariant time recentering: selected compact windows are local objects, and composing the terminal-frame chart with the finite translation $\tau \mapsto \tau - \tau_n$ produces another member of the same local validation family on every fixed model interval $J_0 = [-\sigma, \sigma]$. The compact-cylinder estimates are then the estimates of [definition 10.68\(i\)](#) on these recentered selected windows, with constants depending on the model cylinder and not on n . In the recentered coordinates the original selected time τ_n is exactly 0, so the lower bound remains $\|S_n(0)\|_{L^3(B_R)} \geq \delta$. This proves (ii) and uses no continuity in time of the L^3 norm. $\square$

Lemma 10.71 (Local compactness of terminal occupation states). *Assume the estimates in [Definition 10.68\(i\)](#) on $Q_{2R,I}$. Then every sequence of selected subwindows and subcylinders has a subsequential local occupation state.*

More precisely, after passing to a subsequence, the full represented branch V_n converges weakly in $L^2(I; H^1(B_R))$ and strongly in $L^2(I; L^q(B_R))$ for every $q < 6$. Its pressure gradients converge distributionally modulo functions of time. In the canonical pressure split, the source-generated parts converge with the velocity products, while the mean-normalized harmonic components converge weakly on inner balls after subsequence extraction; failure of that harmonic compactness is recorded as the pressure/exterior row.

The components U_n and S_n have the same strong compactness only on subbranches for which their componentwise time-derivative rows are available. In the general terminal decomposition, U_n and S_n are retained through their weak $L^2 H^1$ bounds, selected-time measures

$$|U_n(\tau_n)|^3 dy, \quad |S_n(\tau_n)|^3 dy,$$

and mixed-stress occupation coordinates.

Proof. The Aubin–Lions argument applies directly to V_n , because the represented Navier–Stokes equation gives the required time-derivative bound. The local velocity bounds and the pressure bound give compactness of V_n on smaller cylinders and pressure convergence modulo functions of time.

For U_n and S_n , the same strong spacetime compactness is available only when their componentwise time-derivative rows are part of the retained terminal decomposition. Otherwise the augmented state space records their weak velocity bounds, their selected-time concentration

measures, and their mixed-stress/pressure diagnostics. These are precisely the coordinates used in the local residual exhaustion theorem. $\square$

Lemma 10.72 (Closed local labels for terminal residuals). *In the augmented local topology for*

$$(V_n, P_n, U_n, S_n, a_n, b_n)$$

supplied by lemma 10.71, the ordered validation rows in lemma 8.73, with the compact local Type II row refined by theorem 11.18, are closed under subsequential limits of selected compact terminal windows, after shrinking the tested cylinders inside the fixed compact cylinder when necessary.

Proof. The topology records the represented velocities V_n weakly in $L^2 H^1$ and strongly in $L^2 L_{\text{loc}}^q$ for every $q < 6$. The component fields U_n and S_n are recorded weakly in their local velocity topologies; they are recorded strongly only on subbranches where the corresponding componentwise time-derivative rows are selected. In all cases the augmented topology also records the selected-time residual measures $|S_n(\tau_n)|^3 dy$ weakly-* as finite Radon measures on compact balls and the mixed stress coordinates $U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n$ as occupation coordinates. It records the pressures weakly in $L_{\text{loc}}^{3/2}$ modulo functions of time, and the modulation, scale, center, cutoff, and cost variables in the compactified repaired-gauge coordinates used in the ordered validation procedure. The compactification adds boundary faces exactly for loss of representation, scale admissibility, bounded modulation, pressure reconstruction, local H^1 control, critical tightness, componentwise time-compactness when needed, and cost compatibility.

Positive local velocity-pressure concentration is closed on smaller cylinders. For spacetime rows, the velocity part follows from strong local L^3 convergence after interpolating the strong $L^2 L^q$ convergence with the uniform $L^\infty L^3$ bound. For the pressure part, the weak $L^{3/2}$ bound alone gives only lower semicontinuity and would allow positive pressure mass to disappear in the weak limit. We therefore upgrade pressure convergence to strong local convergence on the retained branch as follows. After subtracting the reconstructed pressure of the limit branch, write the pressure difference on B_{R_1} , modulo functions of time, as

$$\tilde{P}_n = P_n - P - c_n = \tilde{P}_n^{\text{loc}} + \tilde{H}_n + c_{n,R_1}(\tau), \quad \tilde{P}_n^{\text{loc}} = \mathcal{R}_i \mathcal{R}_j (\zeta((V_n \otimes V_n) - (V \otimes V))_{ij}),$$

with $\zeta \equiv 1$ on B_{R_1} and $\text{supp } \zeta \Subset B_{R_2}$. Since $V_n \rightarrow V$ strongly in $L_\tau^2 L_y^3(B_{R_2})$, while V_n and V are uniformly bounded in $L_\tau^\infty L_y^3(B_{R_2})$, interpolation gives strong convergence in $L_\tau^3 L_y^3(B_{R_2})$. Hence

$$V_n \otimes V_n \rightarrow V \otimes V \quad \text{strongly in } L_\tau^{3/2} L_y^{3/2}(B_{R_2}).$$

Calderón–Zygmund continuity therefore gives strong convergence of $\tilde{P}_n^{\text{loc}}$ to zero in $L_\tau^{3/2} L_y^{3/2}(B_R)$. The harmonic remainder $\tilde{H}_n$ is not controlled by a weak lower-semicontinuity argument: it is routed with $(q, p) = (3/2, 3/2)$ by the pressure/exterior alternative in lemma 9.22. Thus on the retained branch the harmonic part vanishes locally modulo functions of time, while off that branch the nonzero harmonic part is a named pressure or exterior label. Consequently positive pressure mass cannot be hidden by weak convergence on the retained branch; pressure rows are strongly closed in the local $L_\tau^{3/2} L_y^{3/2}$ topology after shrinking cylinders. For selected-time rows, the velocity part is recorded by the weak-* limits of the local measures $|S_n(\tau_n)|^3 dy$, so a positive lower bound persists as mass of the limiting measure. Hence a lower CKN threshold cannot disappear after passing to a subsequence and shrinking the cylinder.

Multibubble and multicore splitting are finite collections of such positive CKN lower bounds with separated centers or separated scales; the separation relations are closed in the scale-center compactification. Same-point cascade labels are closed because they are the boundary faces $\lambda_i^n / \lambda_j^n \rightarrow 0$ or ∞ at a common physical center. Comparable same-point labels are closed by convergence of the relative scale-center parameters and are grouped by lemma 9.3.

Exterior leakage and critical-tightness failure are closed because they are defined by a positive lower bound for local L^3 or CKN mass outside every retained compact set. If that mass returns to a fixed compact set along a subsequence, the state is relabelled by the corresponding compact concentration row; otherwise it remains on the exterior/tightness boundary face. Loss of local windowed H^1 control is the face at infinity for the nonnegative L^2H^1 seminorm on a compact cylinder.

Finite-cost scale collapse, negative-drift, and compact cost-divergence labels are determined by nonnegative localized cost functionals. Their sublevel or positive-carrier conditions are lower semicontinuous in the recorded occupation-state topology. The scale-rigid residual label is closed by the relative scale-center convergence and by the compact realization of the scale-rigid branch in [definition 5.19](#). Representation, modulation, pressure, carrier, and cost-compatibility failures are not hidden analytic assumptions: they are the explicitly added boundary faces of the ordered validation compactification. Thus every first row of the ordered local validation is a closed diagnostic set. $\square$

Lemma 10.73 (Residual measure routing for retained decompositions). *Let*

$$V_n = U_n + S_n$$

be an admissible retained terminal decomposition on a compact terminal cylinder $Q_{2R,I} = B_{2R} \times I$. Assume the compact-window bounds, pressure reconstruction, and augmented occupation-state coordinates of [Definition 10.68](#). Suppose that for selected times $\tau_n \in I$ and a compact ball $B_\rho \Subset B_R$,

$$|S_n(\tau_n)|^3 dy|_{B_\rho} \xrightarrow{*} \mu_S, \quad \mu_S(B_\rho) > 0.$$

Then, after passing to a subsequence and shrinking B_ρ if necessary, one of the following occurs:

- (i) *the residual measure is supported on a separated-center, separated-scale, noncompact, or cascade face; this is a named adjacent row;*
- (ii) *the residual measure is supported on a same-point comparable retained class; then the class is grouped into U_n by [Lemma 10.77](#);*
- (iii) *the mixed-stress occupation coordinate*

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n$$

has positive mass; this is a named pressure, rough-core, compactness, multibubble/cascade, exterior, or cost/carrier row;

- (iv) *the residual measure is genuinely diffuse in the augmented local occupation-state coordinates; then [Theorem 10.86](#) routes it to a named adjacent row or to a retained comparable class, and the retained comparable class is grouped into U_n ;*
- (v) *none of the above occurs, in which case*

$$S_n \rightarrow 0 \quad \text{in } L^3_{\text{loc}},$$

the residual is perturbative in the complete terminal PDE topology, and the ancestry terminalization absorbs it before the nonretained residual is fixed.

Proof. The augmented terminal state records the residual measures

$$|S_n(\tau_n)|^3 dy$$

and the mixed-stress occupation coordinates

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n.$$

If the mixed-stress coordinate has positive mass, then the pressure/stress diagnostics in the augmented state record one of the named pressure, rough-core, compactness, multibubble/cascade, exterior, or cost/carrier rows. This is (iii).

Assume therefore that the mixed-stress coordinate is not positive. Normalize the positive residual measure on B_ρ and apply [Lemma 9.64](#). This furnishes a probability measure Π_{res} on a compact residual local-state space to which [Lemma 9.63](#) applies, and every point of $\text{supp } \Pi_{\text{res}}$ is realized by an actual sequence of selected-time residual descendants from the original represented branch.

If the atomic/activity alternative occurs, the support-realization principle of [Lemma 9.64](#) realizes a descendant state with positive compact local activity. The relative scale-center partition places that descendant in a separated-center, separated-scale, noncompact, cascade, or same-point comparable class. The separated, noncompact, and cascade cases are named adjacent rows, giving (i). The same-point comparable case is grouped into U_n by [Lemma 10.77](#), giving (ii).

If the normalized residual measure is genuinely diffuse, then the diffuse probability measure Π_{res} itself is a visible residual coordinate of the augmented state space. Applying [Theorem 10.86](#) to that visible diffuse residual coordinate yields either a named adjacent local alternative or a retained comparable class. The retained comparable class is grouped into U_n by [Lemma 10.77](#). Thus the diffuse case gives (iv).

If none of (i)–(iv) occurs, then all compact spacetime occupation coordinates and all mixed-stress coordinates vanish in the augmented topology. The diagonal compact exhaustion in [proposition 10.80](#)(iii) gives $S_n \rightarrow 0$ in L^3_{loc} , and [lemma 9.30](#) gives perturbativity in every terminal PDE topology. The absorption step of [proposition 10.80](#) then removes this coordinate before the nonretained residual is finalized. This is (v). $\square$

Lemma 10.74 (Residual concentration is visible in the augmented state). *Let $V_n = U_n + S_n$ on a compact terminal cylinder, with the local $L^\infty L^3$ bounds of [definition 10.68](#)(i). Suppose that, for selected times τ_n and a compact ball B_ρ ,*

$$|S_n(\tau_n)|^3 dy|_{B_\rho} \xrightarrow{*} \mu_S, \quad \mu_S(B_\rho) > 0.$$

Then, after shrinking B_ρ if necessary, the residual is assigned to a named adjacent row, grouped into the retained comparable class, recorded as a visible diffuse residual coordinate routed by [theorem 10.86](#) to a named adjacent local alternative or the retained comparable class, or is perturbative in the complete terminal PDE topology and absorbed before the nonretained residual is finalized. In particular, residual L^3 concentration cannot be invisible in the augmented local state space by cancellation between V_n and U_n .

Proof. Apply [Lemma 10.73](#). The lemma routes the positive residual measure either to a named adjacent row, to a same-point comparable retained class grouped into U_n , to a positive mixed-stress row, to a visible diffuse residual coordinate routed by [theorem 10.86](#) to a named adjacent local alternative or a retained comparable class, or to perturbativity in the complete terminal PDE topology followed by ancestry terminal absorption. Thus residual mass is visible in the augmented local state and cannot be hidden by cancellation between V_n and U_n . $\square$

Lemma 10.75 (Active-witness compact-chart/tail dichotomy). *Let (x_n^0, λ_n^0) be a retained frame and let (x_n, λ_n) be a candidate active witness based at the same terminal time origin. Set*

$$\rho_n := \frac{\lambda_n}{\lambda_n^0}, \quad \xi_n := \frac{x_n - x_n^0}{\lambda_n^0}, \quad \sigma_n := \log \rho_n.$$

After passing to a subsequence, exactly one of the following alternatives holds:

- (i) Compact relative chart. *There is $L < \infty$ such that*

$$|\xi_n| \leq L, \quad |\sigma_n| \leq L$$

for all large n .

- (ii) Comparable-scale center escape. *The sequence (ρ_n) is bounded above and away from zero, while $|\xi_n| \rightarrow \infty$. The two frames are separated-point frames.*
- (iii) Strict inner scale. *One has $\rho_n \rightarrow 0$. If (ξ_n) is bounded, the candidate and retained frames form a same-point cascade with the candidate as the inner frame; if $|\xi_n| \rightarrow \infty$, they are separated-point frames.*
- (iv) Strict outer scale. *One has $\rho_n \rightarrow \infty$. If (ξ_n/ρ_n) is bounded, the candidate and retained frames form a same-point cascade with the candidate as the outer frame; if $|\xi_n|/\rho_n \rightarrow \infty$, they are separated-point frames.*

The separated-point alternatives are assigned to the separated-center or exterior row. The two same-point alternatives are assigned respectively to the strict inner-cascade row and to the outer-scale/cascade row. If the rebased witness fails represented compactness, the local pressure atlas, the local energy inequality, or the rough-core bound on its model cylinder, its first failed coordinate is instead the corresponding named failure face of [lemma 9.62](#); no compact-chart packing is applied to that witness.

If a tail witness is recentered and the bounded-window selection is repeated, the iteration is governed by [lemma 11.15](#): it terminates in a finite named row or produces one of the named infinite same-point cascade, scale-collapse, noncompact exterior, or diffuse reselection-tree rows.

Proof. Pass first to a subsequence on which ρ_n converges in the compactified half-line $[0, \infty]$. If its limit belongs to $(0, \infty)$, pass to a further subsequence on which (ξ_n) is bounded or $|\xi_n| \rightarrow \infty$. The bounded case is item (i), after enlarging L , and the escaping case is item (ii). In the latter case,

$$\frac{|x_n - x_n^0|}{\lambda_n^0} = |\xi_n| \rightarrow \infty, \quad \frac{|x_n - x_n^0|}{\lambda_n} = \frac{|\xi_n|}{\rho_n} \rightarrow \infty,$$

so the frames are separated-point in the sense of [definition 5.4](#).

If $\rho_n \rightarrow 0$, pass to a further subsequence on which (ξ_n) is bounded or $|\xi_n| \rightarrow \infty$. In the bounded case,

$$\frac{\lambda_n}{\lambda_n^0} \rightarrow 0, \quad \frac{|x_n - x_n^0|}{\lambda_n^0} = O(1),$$

which is the same-point cascade relation with the candidate as the inner frame. In the escaping case, both normalized center distances diverge, because $|\xi_n|/\rho_n \geq |\xi_n|$ for all large n ; hence the frames are separated-point. This proves item (iii).

If $\rho_n \rightarrow \infty$, pass to a further subsequence on which (ξ_n/ρ_n) is bounded or $|\xi_n|/\rho_n \rightarrow \infty$. The bounded case satisfies

$$\frac{\lambda_n^0}{\lambda_n} \rightarrow 0, \quad \frac{|x_n - x_n^0|}{\lambda_n} = O(1),$$

so it is the same-point cascade relation after interchanging the two frames. In the escaping case,

$$\frac{|x_n - x_n^0|}{\lambda_n} = \frac{|\xi_n|}{\rho_n} \rightarrow \infty, \quad \frac{|x_n - x_n^0|}{\lambda_n^0} = |\xi_n| \rightarrow \infty,$$

and the frames are separated-point. This proves item (iv) and exhausts all limits of ρ_n .

The stated assignments are the corresponding cases of the exhaustive parameter partition [lemma 5.10](#). Strict outer states in the retained inner frame are treated by [theorem 9.6](#); escaping centers are routed by the exterior scale-center package, in particular [lemma 9.58](#). Failure of a rebased analytic input lands on the named coordinate added in [lemma 9.62](#). Finally, [lemma 11.15](#) gives the asserted finite or infinite named route for repeated bounded-window reselection. $\square$

Lemma 10.76 (Finite active packing in a compact relative chart). *Let*

$$Q_{\text{work}} := B_{R_{\text{work}}} \times I_{\text{work}}$$

be a fixed represented working cylinder and suppose

$$\sup_n \|V_n\|_{L^\infty(I_{\text{work}}; L^3(B_{R_{\text{work}}}))} \leq M < \infty.$$

Let Γ be a finite or countable family of candidate active witnesses relative to a retained frame (x_n^0, λ_n^0) , all based at the same terminal time origin. For $\gamma \in \Gamma$, write

$$\rho_n^\gamma := \frac{\lambda_n^\gamma}{\lambda_n^0}, \quad \xi_n^\gamma := \frac{x_n^\gamma - x_n^0}{\lambda_n^0}, \quad \sigma_n^\gamma := \log \rho_n^\gamma.$$

Assume that, for some fixed $L, r_*, \eta_* > 0$ and compact interval J_* , every $\gamma \in \Gamma$ satisfies on a tail depending on γ

$$|\xi_n^\gamma| \leq L, \quad |\sigma_n^\gamma| \leq L,$$

and has a normalized representative on $Q_* := B_{r_*} \times J_*$ given by

$$W_n^\gamma(z, s) = \rho_n^\gamma V_n(\xi_n^\gamma + \rho_n^\gamma z, (\rho_n^\gamma)^2 s)$$

with the fixed velocity floor

$$\int_{J_*} \int_{B_{r_*}} |W_n^\gamma|^3 dz ds \geq \eta_*^3.$$

Assume also that the image cylinders

$$\hat{Q}_n^\gamma := (\xi_n^\gamma + \rho_n^\gamma B_{r_*}) \times (\rho_n^\gamma)^2 J_*$$

are contained in Q_{work} . For $\delta > 0$, call a subfamily $\Gamma_\delta \subset \Gamma$ δ -separated if, for distinct $\gamma, \beta \in \Gamma_\delta$,

$$\liminf_{n \rightarrow \infty} (|\xi_n^\gamma - \xi_n^\beta| + |\sigma_n^\gamma - \sigma_n^\beta|) \geq \delta.$$

Choose Γ_δ maximal before assembling same-point comparable representatives into compound classes. Then $N := \#\Gamma_\delta$ is finite and satisfies

$$N \leq \left\lfloor \frac{C_{L, \delta, r_*, J_*} e^{2L} |I_{\text{work}}| M^3}{\eta_*^3} \right\rfloor,$$

where $C_{L, \delta, r_*, J_*} < \infty$ is the overlap constant of the image cylinders. Every unselected compact-chart witness lies in one of these parameter cells and is grouped with its same-point comparable compound class.

Proof. The bounds on σ_n^γ give

$$e^{-L} \leq \rho_n^\gamma \leq e^L.$$

Let $F \subset \Gamma_\delta$ be finite. Partition $[-L, L]$ into finitely many intervals of length less than $\delta/4$, and in each scale interval apply the Euclidean packing lemma to the centers $\xi_n^\gamma \in \overline{B_L}$. For a finite δ -separated subfamily in the metric

$$d((\xi, \sigma), (\xi', \sigma')) := |\xi - \xi'| + |\sigma - \sigma'|,$$

this gives, for all sufficiently large n depending on that finite subfamily, an overlap bound

$$\sum_{\gamma \in F} \mathbf{1}_{\hat{Q}_n^\gamma} \leq C_{L, \delta, r_*, J_*}$$

on Q_{work} . Equivalently, this is the finite-band parabolic Vitali estimate; its constant depends only on the displayed compact-chart data.

The change of variables

$$y = \xi_n^\gamma + \rho_n^\gamma z, \quad \tau = (\rho_n^\gamma)^2 s$$

has $dy d\tau = (\rho_n^\gamma)^5 dz ds$. Since $W_n^\gamma = \rho_n^\gamma V_n$ in these coordinates,

$$\begin{aligned} \int_{\widehat{Q}_n^\gamma} |V_n(y, \tau)|^3 dy d\tau &= (\rho_n^\gamma)^2 \int_{Q_*} |W_n^\gamma(z, s)|^3 dz ds \\ &\geq e^{-2L} \eta_*^3. \end{aligned}$$

Take n large enough that the parameter separation and mass floor hold for every $\gamma \in F$. Summing the last inequality and using the overlap bound gives

$$\begin{aligned} \#F e^{-2L} \eta_*^3 &\leq \sum_{\gamma \in F} \int_{\widehat{Q}_n^\gamma} |V_n|^3 \\ &\leq C_{L, \delta, r_*, J_*} \int_{Q_{\text{work}}} |V_n|^3 \\ &\leq C_{L, \delta, r_*, J_*} |I_{\text{work}}| M^3. \end{aligned}$$

Thus every finite subfamily obeys the displayed cardinality bound, which proves the same bound for Γ_δ . Enumerate the finite or countable family Γ and select greedily whenever the next parameter is δ -separated from those already selected. If this procedure did not terminate within the displayed bound, its first one-more-than-the-bound members would contradict the estimate just proved. It therefore produces a finite maximal δ -separated family. Every remaining parameter lies within δ of one of its members; the corresponding frames have comparable scales and bounded relative centers, and hence are assembled into the same-point comparable compound classes of [lemma 9.3](#). The estimate uses only the fixed working-cylinder norm stated in the lemma. $\square$

Lemma 10.77 (Maximality of the selected retained frame). *The terminal selection rule for the retained component U_n on a compact terminal cylinder $Q_{2R, I}$ first assigns every candidate outside all compact relative scale-center charts to the parameter-tail alternatives of [lemma 10.75](#). On the remaining compact chart it produces a finite collection of compound retained classes $\{U_n^\alpha\}_{\alpha \in A}$ such that:*

- (i) Pairwise scale-center separation. *For $\alpha \neq \beta \in A$, the relative scale-center parameters $(\lambda_n^\alpha / \lambda_n^\beta, (x_n^\alpha - x_n^\beta) / \lambda_n^\beta)$ either tend to 0 or to ∞ (separated scales) or stay bounded with the centers escaping each other in scale-comparable units (separated centers).*
- (ii) Maximality. *If a candidate retained class $(\tilde{\lambda}_n, \tilde{x}_n)$ at the same physical point has relative scale-center parameters in a compact chart with respect to some $(\lambda_n^\alpha, x_n^\alpha)$ and survives the selection rule (i.e. carries a positive active local mass floor in the sense of [theorem 9.33](#)), then it belongs to the same compound comparable class and is absorbed into U_n^α .*

Proof. The terminal selection rule is local on the fixed compact cylinder $Q_{2R, I}$. It extracts only classes which carry the active local mass floor of [theorem 9.33](#) on a selected compact subcylinder. First normalize those detecting subcylinders using the rescaling in [theorem 9.33](#), decreasing the threshold once if necessary so that one fixed $\eta_* > 0$ holds on a common model cylinder. The relative parameters below are the parameters of these normalized witnesses. If no compact relative scale-center chart contained all remaining candidates, a diagonal choice of candidates would leave every such chart. Applying [lemma 10.75](#) to that diagonal sequence would produce a separated-center/exterior, strict inner cascade, outer-scale/cascade, or named analytic failure row before the retained frame was fixed. Repeated bounded-window rebasing is also routed by the well-foundedness clause of that lemma. Therefore all candidates that reach retained-frame packing lie in one fixed compact relative chart.

The working-cylinder buffer contains the images of the normalized cores. Now apply [lemma 10.76](#) with the local bound

$$\sup_n \|V_n\|_{L^\infty(I; L^3(B_{2R}))} < \infty.$$

It gives a finite maximal separated family by an estimate on $Q_{2R,I}$. Every remaining compact-chart candidate is comparable with a selected representative and is grouped into its compound class. Hence the retained list on the compact terminal branch is finite. No whole-space critical norm is used.

For two retained classes in this finite list, either their relative scale-center parameters leave every compact subset of the relative scale-center chart, giving separated scales or separated centers, or those parameters remain in a compact subset. In the latter case the two concentrations are comparable in the same local observer coordinates and are grouped before the terminal frame is fixed. This proves the separation alternative in (i) for distinct retained classes.

For maximality (ii), suppose toward contradiction that $(\tilde{\lambda}_n, \tilde{x}_n)$ remains in a compact relative chart with respect to $(\lambda_n^\alpha, x_n^\alpha)$ but survives outside U_n^α . Then the rescaled profile $\tilde{\lambda}_n V_n(x_n^\alpha + \tilde{\lambda}_n y, \cdot)$ carries an active local mass floor in a relative scale-center chart compactly comparable to the chart of U_n^α . Thus the two candidates are not separated by any ordered local boundary face. By [lemma 9.3](#), all same-point comparable concentrations in such a compact relative chart are assembled into one compound retained class before terminal selection is closed. The candidate $(\tilde{\lambda}_n, \tilde{x}_n)$ must therefore already be absorbed into U_n^α , contradicting the assumption that it is a surviving unabsorbed comparable class. $\square$

Lemma 10.78 (Terminal local critical-mass failures are adjacent labels). *Let a physical local Type II branch have passed covered entry and let a terminal active frame be selected from its local state-space decomposition. On each fixed compact terminal cylinder, exactly one of the following occurs after passing to a subsequence:*

- (i) *the retained terminal component has positive finite local critical spacetime L^3 mass on the cylinder;*
- (ii) *the retained spacetime activity floor fails, so the candidate is not a retained terminal component;*
- (iii) *the local critical-mass upper control fails on a compact spacetime subcylinder, in which case the ordered local state space records an additional compact active label: a critical-mass adjacent label at the same retained frame, a comparable finite-mass retained class which must be grouped into the retained component, a multibubble/multicore label, a cascade label, a rough-core label, or an exterior/separated label.*

In the retained terminal branch only (i) remains.

Proof. The lower bound in (i) is the active local spacetime mass floor of [theorem 9.33](#). If this floor is absent, the candidate frame is not retained by the terminal selection rule, giving (ii).

Assume the lower floor is present and the finite local upper bound fails on a compact spacetime subcylinder $B_R \times J \Subset B_{2R} \times I$. Then

$$\int_J \int_{B_R} |V_n(y, \tau)|^3 dy d\tau \rightarrow \infty.$$

Cover $B_R \times J$ by finitely many smaller parabolic cylinders. On one member of this finite cover, after passing to a subsequence and renormalizing at the corresponding local scale if necessary, the branch carries a further positive compact spacetime L^3 diagnostic. The ordered local state-space compactification records the first label of that diagnostic. If the diagnostic stays at the same comparable retained frame and the compact upper control still fails, the first label is the critical-mass adjacent face itself; by the retained-state definition this is not a retained terminal branch.

If the diagnostic is comparable and finite-mass, then [lemma 10.77](#) says it was already grouped into the retained component. If it is not comparable, it is by definition one of the adjacent compact active labels: multibubble/multicore splitting, cascade, rough core, or exterior/separated concentration. These are precisely the labels in (iii).

On a retained terminal branch all comparable labels have already been grouped, and all non-comparable labels are adjacent local alternatives routed by the state-space closure package. Hence only the positive finite local critical mass alternative remains. This is a compact-cylinder statement; it does not require $V(\tau) \in L^3(\mathbb{R}^3)$ or a whole-space critical norm bound. $\square$

Corollary 10.79 (No unclassified compact-limit obstruction). *Let an ancestry-generated retained terminal branch produce full represented windows (V_n, P_n, a_n, b_n) satisfying the selected compactness package on a compact terminal cylinder. Then every nonzero compact-window limit has one of the following already recorded routes:*

- (i) *a first compactness-interface failure gives its named local row;*
- (ii) *a regular limit gives the bounded selected-window row;*
- (iii) *a singular limit gives an actual CKN-active descendant of the full represented branch. Under the ordinary local selection, that descendant either leaves every compact relative scale-center chart and enters a named separated-center, exterior, inner-cascade, outer-cascade, or analytic-failure row, or it remains in a compact relative chart. In the latter case a descendant carrying the retained activity floor is already part of the compound retained class, while failure of the floor or of the finite compact upper row is an adjacent critical-mass label.*

Consequently there is no additional retained terminal obstruction arising from a nonzero compact limit. The comparable retained class in (iii) is not eliminated by this corollary; it is identified with the class already selected before terminalization.

Proof. Apply [lemma 7.30](#). Its first alternative is item (i). Its nonsingular alternative supplies a nonzero bounded selected-window limit, which is the bounded-window row named in [theorem 5.27](#). Its physical closure is deferred to the bounded-window entry of the final ledger. This is item (ii).

In the singular alternative, the diagonal rebase constructed in that lemma is performed on the full represented tuple and therefore enters the ordinary local descendant selection. That selection writes each rebased window in coordinates with its own selected time at the common model origin; it does not move the CKN lower bound to a different time. Relative to the retained frame in those coordinates, apply [lemma 10.75](#). Its three noncompact alternatives are precisely the named separated/exterior and inner- or outer-cascade rows, with any failed analytic input recorded before scale-center packing. In the compact-chart alternative, a descendant with the retained activity floor is absorbed by [lemma 10.77](#). If the activity floor or the compact upper control needed for retention fails, the first label is one of the alternatives of [lemma 10.78](#). These cases are exhaustive, which proves item (iii) and the conclusion. $\square$

Proposition 10.80 (Residual-spike routing without pointwise persistence). *Let*

$$V_n = U_n + S_n$$

be an ancestry-generated retained terminal decomposition, and suppose that on some compact terminal cylinder

$$\limsup_{n \rightarrow \infty} \|S_n\|_{L^\infty(I; L^3(B_R))} > 0.$$

Choose times τ_n carrying a fixed lower bound and apply the admissible finite recentering of [lemma 10.70](#). After passing to a diagonal subsequence, the following ordered alternatives are exhaustive.

- (i) *The ordered representation, pressure, modulation, local-energy, rough-core, compact-upper, critical-mass, exterior, scale-center, or represented-equation validation first fails on the recentered full branch. The corresponding named local row occurs.*
- (ii) *All preceding validation rows pass, and either a compact subcylinder Q' has positive residual spacetime occupation,*

$$\liminf_{n \rightarrow \infty} \int_{Q'} |S_n|^3 dy d\tau > 0,$$

or a parabolic rebase of the selected-time spike reaches the fixed local CKN activity threshold. At the first occupation coordinate and first positive level detecting this amount, rerun the local state selection on the full represented branch V_n . A candidate leaving every compact relative scale-center chart is routed by [lemma 10.75](#). A candidate remaining in a compact relative chart is grouped into the existing compound retained class by [lemma 10.77](#), or it produces the corresponding multicore, same-point-cascade, gauge-degenerate, bounded-window, pressure/exterior, rough-core, or other named local row.

- (iii) *All preceding validation rows pass, no compact subcylinder has positive limiting residual spacetime occupation, and no parabolic rebase reaches the local CKN activity threshold. Then*

$$S_n \rightarrow 0 \quad \text{in } L^3_{\text{loc}},$$

and [lemma 9.30](#) makes S_n perturbative in every nonlinear-stress, pressure, modulation, linear-residual, and cutoff topology used downstream. The ancestry terminalization therefore absorbs this perturbative component before fixing the nonretained terminal residual.

Thus a selected-time L^3 spike never creates an additional terminal Type II branch. It either has an active descendant of the full represented branch, is assigned to an earlier named row, is grouped into the retained comparable class, or is removed by perturbative terminal absorption.

Proof. Choose $\delta > 0$, a subsequence, and times $\tau_n \in I$ such that

$$\|S_n(\tau_n)\|_{L^3(B_R)} \geq \delta.$$

[Lemma 10.70](#) places these same times in the interior of a fixed model window; it does not transfer the lower bound to another time. Run the ordinary ordered validation on the full represented branch on the recentered windows. A first failure gives (i). In particular, the represented equation for U_n used below is either available with its vanishing local error, or its failure has already closed the corresponding row.

Assume that every preceding local input passes. Fix a countable exhaustion of the selected model window by compact subcylinders and the countable family of occupation tests used in the augmented local state space. Pass diagonally so that every nonnegative residual occupation coordinate has a limit. If one such coordinate is positive, choose the first coordinate in the fixed enumeration and the first dyadic level below its limiting value. This gives the positive compact spacetime occupation in (ii) without selecting a new time.

Run the local state selection at that window and level on (V_n, P_n, a_n, b_n) , retaining U_n and S_n only as augmented coordinates. The relative scale-center alternatives are exhaustive by [lemma 10.75](#). Outside all compact charts that lemma gives a separated-center, cascade, noncompact, or named analytic row. Inside a compact chart, a candidate carrying the retained activity floor is grouped by [lemma 10.77](#); failure of that floor or of any compact local input is routed by [lemma 10.78](#) and the ordered validation rows. If the full selected compactness package passes on this subbranch, [corollary 10.79](#) additionally shows that its strong compact limit can introduce only the same active-descendant, bounded-window, named-interface, or already retained comparable routes. This proves the spacetime-occupation part of (ii).

The same conclusion holds when the fixed-window occupations vanish but a smaller parabolic rebase reaches the CKN activity threshold. Indeed, [lemma 9.64](#) realizes the rebased state as an actual descendant of the original represented branch. Applying the same relative scale-center partition and maximality argument to that descendant gives (i) or (ii). No Navier–Stokes equation is imposed on S_n as an independent solution. In particular, [lemma 7.30](#) is applied only to the full represented sequence, never to the residual coordinate.

It remains to consider the case in which all compact residual spacetime occupation coordinates vanish and there is no active rebase. The diagonal choice and the compact exhaustion imply

$$\|S_n\|_{L^3(Q')} \rightarrow 0$$

for every compact Q' in the selected terminal window. Since the represented-equation row has passed, all hypotheses of [lemma 9.30](#) hold on nested compact cylinders. That lemma proves perturbativity in the complete terminal PDE topology.

For completeness, the terminal absorption in this case preserves the exact decomposition and every retained limiting coordinate. On a fixed compact window replace

$$(U_n, S_n) \quad \text{by} \quad (\tilde{U}_n, \tilde{S}_n) := (U_n + S_n, 0) = (V_n, 0).$$

The full represented equation and its pressure belong to $\tilde{U}_n$. The difference between the old and new retained equations consists precisely of the cross stress, cross pressure, modulation, linear-residual, and cutoff terms shown in [lemma 9.30](#) to vanish. Moreover, $S_n \rightarrow 0$ in spacetime L^3 , so the positive finite spacetime mass of the retained class and every compact occupation coordinate are unchanged in the limit. Performing this replacement on the countable compact exhaustion by the same diagonal procedure is the perturbative-pruning step of the ancestry construction. Hence such a spike is not left in the finalized nonretained residual. This is (iii) and completes the ordered routing. $\square$

Theorem 10.81 (Adjacent terminal labels are locally routed). *For every admissible retained terminal decomposition, every first label in the ordered augmented local state space is either the retained comparable class or one of the named local alternatives. The adjacent labels are closed and routed out of the retained terminal component as follows: representation, modulation, pressure, critical-mass, carrier, cost-compatibility, exterior tightness, local H^1 -loss, multibubble, multicore, cascade, finite-cost scale-collapse, and scale-rigid residual labels are assigned to their named local alternatives and do not remain on the retained terminal branch. This theorem is a routing theorem only: it does not invoke the final state-elimination proposition and does not claim, inside the terminal-completeness argument, that the adjacent alternatives have already been globally excluded.*

Proof. [Lemma 10.72](#) proves closedness of the rows in the augmented local topology. Representation and modulation failures are the boundary faces of the repaired-gauge compactification and are exactly the representation alternatives. Pressure rows are routed by [lemma 9.22](#) and by the pressure reconstruction stability [lemma 3.12](#). Critical-mass rows are routed by [lemma 10.78](#): failure of the retained spacetime activity floor is not retained, and an upper local L^3 -failure produces a comparable or adjacent active label. Exterior tightness and exterior leakage are routed by [theorem 9.39](#). Local H^1 -loss is exactly the rough-core row supplied by [theorem 11.7](#). Multibubble, multicore, and cascade are primitive structural labels in the ordered active selection record: they arise from failure of unique retained-core selection, strict nested same-point descendants, or separated comparable descendants before any terminal retained class is fixed. Gauge-degenerate labels are already part of the repaired-gauge representation failure record. Finite-cost scale-collapse is the finite-cost scale-collapse row of [theorem 11.9](#), and scale-rigid residual labels are recorded as the scale-rigid row governed by [proposition 5.42](#). Carrier and

compact-cost compatibility rows are the carrier/cost adjacent rows of the ordered compact-cost validation.

The point needed here is only separation from the retained terminal component. Each listed row is a boundary face of the ordered local state space, not the retained comparable class. Its later exclusion is supplied in the assembly layer by the local closure theorems and, for carrier/cost rows, by the ambient adjacent wrapper. The terminal-completeness proof therefore uses this theorem only to say that a positive residual label cannot stay in the nonretained remainder of a retained terminal branch. The only first label not routed away is the retained comparable class, which is grouped into the selected terminal component before the terminal frame is fixed. $\square$

Theorem 10.82 (Residual local mass gives an admissible augmented state). *Let $V_n = U_n + S_n$ be an ancestry-generated retained terminal decomposition, and assume the local compact-window estimates of [definition 10.68\(i\)](#). If*

$$\limsup_{n \rightarrow \infty} \|S_n\|_{L^\infty(I; L^3(B_R))} > 0,$$

then, after passing to a subsequence, there are selected times, still denoted τ_n in the selected coordinates, a compact selected or recentered subcylinder $Q_{\rho, J}$, and an augmented local occupation state

$$\mathfrak{s} = \lim_{n \rightarrow \infty} (V_n, P_n, U_n, S_n, a_n, b_n)|_{Q_{\rho, J}}$$

whose residual coordinate or selected-time residual L^3 concentration measure is nonzero. The Navier–Stokes admissibility conditions in the ordered local validation are imposed on the full represented branch (V_n, P_n, a_n, b_n) , not on S_n as an independent solution. Consequently the first nonretained residual row is either a named adjacent local alternative or the retained comparable local class.

Proof. Choose $\delta > 0$, times $\tau_n \in I$, and a subsequence such that

$$\|S_n(\tau_n)\|_{L^3(B_R)} \geq \delta.$$

By [lemma 10.70](#), after passing to a further subsequence, either the original selected times lie in a fixed interior subinterval $J \Subset I$, or we replace the terminal window by the admissible recentered windows supplied by that lemma and write the selected time as $0 \in J_0^\circ$. In both cases we keep the same selected time rather than transferring the lower bound to a different time; below J denotes the fixed model interval in the selected coordinates and τ_n denotes the selected interior time in those coordinates. Cover B_R by finitely many smaller balls compactly contained in B_R . One ball, after passing to a subsequence, carries residual L^3 mass bounded below by a fixed multiple of δ . The compact subinterval J supplied by [lemma 10.70](#) contains the selected times with positive distance from the boundary of the model window; we use this J without further time movement. The estimates in [definition 10.68\(i\)](#) and [lemma 10.71](#) give a subsequential limit of the augmented tuple on $Q_{\rho, J}$. In addition, the measures

$$|S_n(\tau_n)|^3 dy|_{B_\rho}$$

have uniformly bounded total mass on compact balls and therefore converge weakly-*, after passing to a subsequence, to a finite nonzero Radon measure. The lower mass bound on the chosen ball passes to this measure. Thus the augmented state has a nonzero residual coordinate or, in the concentration case, a nonzero residual L^3 concentration measure. [Lemma 10.74](#) shows that this residual diagnostic is visible either as a named adjacent local label of the full represented branch, as a positive mixed-stress row, as a visible diffuse residual coordinate in the augmented state space, or as a positive comparable retained label. Hence it cannot be hidden by cancellation with U_n .

The PDE structure used by the state-space decomposition is carried by (V_n, P_n, a_n, b_n) : it satisfies the repaired-gauge Navier–Stokes equation, the local energy inequality, pressure reconstruction modulo functions of time, bounded modulation, and the compact-cylinder bounds. The retained component U_n and residual coordinate S_n are bookkeeping variables in the augmented state space. Thus no step applies the Navier–Stokes decomposition theorem to S_n alone.

Apply [proposition 10.80](#) to the selected-time lower bound. Its first alternative is already a named adjacent local row. In its second alternative, the positive spacetime occupation or CKN-active rebase is an actual descendant of the full represented branch and therefore enters the ordered first-row extraction below. Its third alternative is perturbative in the full terminal PDE topology and is absorbed before the ancestry-generated nonretained residual is finalized; hence it cannot be the S_n appearing in the conclusion of this theorem.

Evaluate the ordered local validation list of [lemma 8.73](#) on the augmented state. If a representation, pressure, modulation, exterior-tightness, local H^1 , finite-cost, scale-rigid, or carrier/cost row fails first, that row is the named adjacent alternative. If the compact local Type II row is reached, apply [theorem 10.81](#) to the first nonretained row of the full represented branch (V_n, P_n) on the selected compact window. If that first row is not the retained comparable class, then it is a named adjacent local alternative. If the visible residual diagnostic is diffuse rather than atomic, we apply the same first-row extraction to that diffuse residual coordinate in the augmented state space; closedness of the ordered rows again yields a named adjacent local alternative or the retained comparable class. The remaining possibility is therefore that the residual diagnostic is carried by the retained compact class. That class is comparable to the selected terminal frame; otherwise the separated-center or separated-scale decoupling lemmas would put it into one of the preceding adjacent rows. Hence the first nonretained residual row is either a named adjacent local alternative or the retained comparable class. $\square$

Lemma 10.83 (Positive residual mass produces a local state-space label). *Assume the local compact-window estimates of [definition 10.68\(i\)](#). If*

$$\limsup_{n \rightarrow \infty} \|S_n\|_{L^\infty(I; L^3(B_R))} > 0,$$

then, after passing to a subsequence and selecting times and subcylinders, the residual produces one of the first closed labels listed in [definition 10.68\(ii\)](#).

Proof. [Theorem 10.82](#) constructs the augmented occupation state and verifies that the Navier–Stokes admissibility belongs to the full represented branch (V_n, P_n) . The ordered first row of that state is closed by [lemma 10.72](#). Therefore, after passing to a subsequence, the positive residual mass has a first closed local label in the sense of [definition 10.68\(ii\)](#). $\square$

Lemma 10.84 (Residual mass cannot remain in the retained terminal branch). *Assume the local terminal frame has been chosen after grouping comparable retained classes. Then a nonretained residual satisfying the hypotheses of [lemma 10.83](#) cannot remain on the retained terminal branch.*

Proof. If the residual has positive local mass, then [lemma 10.83](#) assigns it to a first closed local label. By [theorem 10.81](#), an adjacent label is a named local alternative and is not part of the retained branch. If the label is the retained compact class, then the corresponding component is comparable to the selected retained terminal frame; by the terminal selection rule together with [lemma 10.77](#), comparable retained classes based at the same physical point were already grouped into U_n . It therefore cannot remain in S_n . Both cases contradict the assumption that positive residual mass persists in the nonretained component of the retained terminal branch. $\square$

Lemma 10.85 (Subsequential exhaustion of local residual mass). *Under the hypotheses of [lemma 10.84](#), every compact terminal cylinder satisfies*

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_R)).$$

Proof. If the convergence failed, then a subsequence and selected times would satisfy the positive residual lower bound in [proposition 10.80](#). In alternative (i) of that proposition, the branch has entered a named adjacent row and is not the retained terminal branch. In alternative (ii), the active descendant is either assigned to a named adjacent row or grouped into the retained comparable class by [lemma 10.77](#); in neither case can it persist in the finalized nonretained residual. In alternative (iii), [lemma 9.30](#) proves that the spike is perturbative in the complete terminal PDE topology, and the ancestry-terminalization step absorbs it before fixing S_n . All ordered alternatives therefore contradict persistence of the selected-time lower bound in the finalized nonretained residual. Hence

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_R)),$$

without transferring the lower bound between times. $\square$

Theorem 10.86 (Visible local residual exhaustion theorem). *Assume the local compact-window estimates, local pressure atlas, repaired-gauge chart on the retained branch, and local energy inequality on a fixed compact terminal cylinder. Assume moreover that the positive residual diagnostic is visible in the augmented local occupation-state coordinates used in [lemma 10.72](#); this includes, in particular, positive local L^3 residual mass and any pressure/cost/compactness defect already encoded there as a closed diagnostic row. Then after passing to a subsequence and, if necessary, shrinking to a selected compact subcylinder, one of the explicitly named local rows occurs. More precisely, the first closed row is either a named adjacent local alternative, or the retained comparable class, in which case the corresponding component must already have been grouped into the retained terminal class before the frame was fixed.*

Proof. For positive residual L^3 mass, [proposition 10.80](#) first separates a fixed-window occupation or CKN-active rebase from a spacetime-small selected-time spike. The latter is perturbatively absorbed and therefore is absent from the finalized nonretained residual. In the former case, [theorem 10.82](#) produces an augmented local occupation state for the tuple

$$(V_n, P_n, U_n, S_n, a_n, b_n)$$

on a selected compact subcylinder. The no-cancellation step is supplied before the first-row extraction by [lemma 10.73](#): a positive selected-time residual measure produces a named adjacent row, a same-point comparable retained class, a positive mixed-stress row, or a visible diffuse residual coordinate; if none of these occurs, the residual is perturbative in the terminal residual topology and no positive residual measure remains. For any other residual diagnostic covered by the statement, visibility in the same augmented coordinates is part of the hypothesis. By [lemma 10.72](#), each row of the ordered local validation is a closed diagnostic set in that augmented topology. Therefore a positive residual defect has a first closed row after subsequence extraction.

If the first row is representation, pressure, modulation, rough-core, exterior, multibubble, cascade, carrier, or cost incompatibility, then [theorem 10.81](#) identifies it as a named adjacent local alternative. If the first row is positive residual L^3 mass, then [lemma 10.74](#) routes it through the no-cancellation alternatives: a named adjacent row, a positive mixed-stress row, a visible diffuse residual coordinate, or a retained comparable class. By [lemma 10.72](#), the diffuse residual coordinate also has a first closed row in the ordered validation. By [lemma 10.77](#), every such retained comparable class was already grouped into the selected retained component before the terminal frame was fixed.

Consequently no positive nonretained residual defect can remain unlabelled on the retained terminal branch: it is exhausted by a named local row or absorbed into the retained comparable class. This is precisely the required local exhaustion statement. $\square$

Theorem 10.87 (Local terminal completeness from state-space stratification). *Every admissible retained terminal decomposition in the sense of [definition 10.68](#)(i) satisfies the local terminal completeness conclusion of [definition 10.68](#)(iii), after the terminal frame is selected by grouping comparable retained classes.*

Proof. [Lemma 10.71](#) gives the local compactness needed to pass to occupation states. The actual exhaustion step for positive local residual L^3 mass, and more generally for every visible residual diagnostic recorded in the augmented state space, is [theorem 10.86](#): such a defect is assigned, after subsequence extraction, to a named adjacent row or to the retained comparable class. The latter possibility is removed by [lemma 10.77](#), and the former is incompatible with persistence of nonretained residual mass on the retained branch by [lemma 10.84](#). Hence [lemma 10.85](#) gives $S_n \rightarrow 0$ on compact terminal cylinders. This is the local terminal completeness conclusion named in [definition 10.68](#)(iii). $\square$

Remark 10.88 (Status of the terminal completeness statement). [Theorem 10.87](#) is local. It uses only compact-cylinder estimates, Aubin–Lions compactness, pressure reconstruction modulo functions of time, and the local state-space stratification. It does not use a global critical profile decomposition, a global heat-flow smallness norm, or a global Kato perturbation theorem.

Definition 10.89 (Local compact terminal windows). The local compact terminal-window hypothesis is the assertion that every terminal active sequence used in the terminal analysis is represented, on each compact terminal-frame cylinder $Q_{2R,I}$, by fields satisfying the local compact-cylinder bounds in [definition 10.68](#)(i). Equivalently, the analysis is carried out on the local state variables

$$(V_n, U_n, S_n, P_n, a_n, b_n)|_{Q_{2R,I}},$$

not on an arbitrary global sequence of initial data. Here the time interval I is always taken from the fixed local validation family of model intervals, up to the admissible finite covariant recenterings of [lemma 10.70](#). In particular the terminal compactness package is local on fixed-size model windows rather than on arbitrary long global intervals.

Theorem 10.90 (Local perturbative residual stability). *Let $Q_{2R,I}$ be a compact terminal cylinder. Assume*

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_{2R})),$$

and assume U_n and S_n are uniformly bounded in

$$L^2(I; H^1(B_{2R})) \cap L^\infty(I; L^3(B_{2R})).$$

Then all velocity interactions involving at least one factor S_n vanish on the smaller cylinder:

$$U_n \otimes S_n + S_n \otimes U_n + S_n \otimes S_n \rightarrow 0 \quad \text{in } L^1(I; L^2(B_R)).$$

If the corresponding cross pressures are reconstructed by the local Calderón–Zygmund decomposition, then their gradients vanish in

$$L^1(I; H^{-1}(B_R)).$$

Proof. By Sobolev on B_{2R} , the $L^2 H^1$ bounds give uniform $L^2(I; L^6(B_{2R}))$ bounds. Hölder’s inequality gives

$$\|U_n \otimes S_n\|_{L^1(I; L^2(B_R))} \leq \|S_n\|_{L^\infty(I; L^3(B_R))} \|U_n\|_{L^1(I; L^6(B_R))} \rightarrow 0,$$

and the same estimate treats $S_n \otimes U_n$. The term $S_n \otimes S_n$ is identical, using the uniform $L^1 L^6$ bound for S_n . The pressure statement follows from the local pressure-kernel estimate applied to the cross stress just shown to vanish in $L^1 L^2$, with the pressure normalized modulo functions of time. Concretely, the local Calderón–Zygmund piece of the cross pressure vanishes in $L^1_\tau L^2_y(B_R)$ by

Calderón–Zygmund continuity applied to the vanishing cross stress, and the harmonic remainder is routed by [lemma 9.22](#) with $(q, p) = (1, 2)$ (using the quantitative nested-ball harmonic-tail estimate of [lemma 9.21](#)): on the retained terminal branch any non-vanishing harmonic remainder triggers a named pressure/exterior-leakage alternative, which is excluded. The two contributions sum to zero in $L^1_\tau H^{-1}_y(B_R)$. $\square$

Definition 10.91 (Local perturbative stability). Local perturbative stability means that every residual component which is small in $L^\infty(I; L^3(B_R))$ on compact terminal cylinders contributes only vanishing velocity-stress, pressure, and linear-residual terms in the local topologies used in [definition 9.13](#).

Corollary 10.92 (Local residual stability gives perturbative stability). *Theorem 10.90 implies the local perturbative stability condition of [definition 10.91](#).*

Proof. This is exactly the conclusion of [theorem 10.90](#) on each compact terminal cylinder, combined with the local pressure and residual estimates already included in terminal nonlinear decoupling. $\square$

Discharge of the retained compact-test closure.

Lemma 10.93 (Logarithmic shell support points are realized after shell rebasing). *Fix compact model intervals $J_{\text{sh}} \in J_* \in \mathbb{R}$ with nonempty interior and set*

$$Q_{\text{sh}} := B_1 \times J_{\text{sh}}.$$

Let J_n be selected terminal windows modeled on J_ . For each occupied spacetime shell block*

$$A_{n,k} \times J_n = \{2^k R_n < |y| \leq 2^{k+1} R_n\} \times J_n$$

cover the block by finitely overlapping translated shell cylinders

$$Q_{\text{sh}}(y_{n,k,\alpha}, \tau_{n,k,\alpha}) := B_1(y_{n,k,\alpha}) \times (\tau_{n,k,\alpha} + J_{\text{sh}}) \subset \mathbb{R}^3 \times J_n,$$

with $y_{n,k,\alpha} \in A_{n,k}$. Record the logarithmic shell radius

$$\rho_{n,k} := \log(2^k R_n),$$

the compactified shell-location data of the chosen cylinder inside the shell block, the rebased local velocity/pressure state on the fixed model cylinder Q_{sh} obtained by spatial translation and admissible finite time recentering only, and the boundary faces where representation, pressure-atlas, gauge, local-energy, or rough-core admissibility fails. The local-state factor is compact by [lemma 9.62](#), so the closure of this family is a compact metrizable shell-state space

$$\mathfrak{X}_{\text{shell}} := (\mathbb{N} \cup \{\infty\}) \times \overline{\mathcal{C}}_{\text{shell}} \times \mathfrak{D}_{\text{shell}} \times \mathfrak{F}_{\text{shell}}.$$

The shell occupation coordinates in $\mathfrak{D}_{\text{shell}}$ are continuous coordinate projections of the product compactification. Consequently every positive superlevel set of a shell occupation coordinate is closed in $\mathfrak{X}_{\text{shell}}$. After passing to a subsequence, the normalized spacetime shell measures induce a weak- probability limit on $\mathfrak{X}_{\text{shell}}$. The limiting shell measure itself is retained as part of the shell compactification, so later diffuse arguments may apply [lemma 9.63](#) directly to it. Every support point is realized by an actual sequence of occupied shell blocks and corresponding translated shell-cylinder extractions from the original represented branch.*

If the support point avoids the failure faces $\mathfrak{F}_{\text{shell}}$, then the realizing shell sequence satisfies the represented local equation, pressure atlas modulo functions of time, repaired-gauge structure, local energy inequality, and compact-cylinder bounds on the fixed model cylinder Q_{sh} ; hence the ordered local state-space validation applies there. If the support point lies on a face of $\mathfrak{F}_{\text{shell}}$, that face itself is the corresponding named local row.

Proof. For each occupied shell block, choose a finitely overlapping family of translated shell cylinders $Q_{\text{sh}}(y_{n,k,\alpha}, \tau_{n,k,\alpha}) \subset \mathbb{R}^3 \times J_n$. Translate each chosen cylinder by $-y_{n,k,\alpha}$ and use the admissible finite covariant recentering of lemma 10.70 to identify its time window with the fixed model interval J_{sh} . No dilation by the shell radius is performed in the local Navier–Stokes state; the shell radius enters only through the logarithmic coordinate $\rho_{n,k}$ and the compactified shell-location data. Spatial translation, finite time recentering, and subtraction of the appropriate pressure mean preserve the represented local equation, divergence-free condition, local energy inequality, and local pressure reconstruction on the fixed model cylinder Q_{sh} . Thus every rebased shell cylinder either remains in a uniform local compactness class on Q_{sh} or hits one of the explicit failure faces recorded in $\mathfrak{F}_{\text{shell}}$. The logarithmic shell index belongs to the one-point compactification of $\mathbb{N}$, the shell-location chart is compactified exactly as in the local state-space package, and the rebased local-state factor is compact by lemma 9.62. This gives the compact metrizable space $\mathfrak{X}_{\text{shell}}$.

Push forward the normalized spacetime shell measures to $\mathfrak{X}_{\text{shell}}$ and extract a weak-* limit. If $\mathfrak{z}$ lies in the support of that limit, choose a nested neighborhood basis (U_m) of $\mathfrak{z}$ with diameters tending to zero. Positive limit mass of U_m implies positive mass of U_m for the approximating shell measures along a subsequence, so for each m we may select an actual occupied shell block and a rebased model-cylinder state lying in U_m . These genuine rebased shell states converge to $\mathfrak{z}$, proving realization. The last paragraph is built into the definition of the failure faces and the represented part of the shell-state space. $\square$

Theorem 10.94 (Logarithmic tail support-profile transfer). *Let*

$$M_n := \int_{J_n} \int_{|y| > R_n} |V(y, \tau)|^3 dy d\tau,$$

and

$$\nu_n := M_n^{-1} |V(y, \tau)|^3 \mathbf{1}_{\{\tau \in J_n, |y| > R_n\}} dy d\tau$$

be the normalized spacetime tail measures associated to a failure sequence of windowed critical tightness, where the selected terminal windows J_n are all modeled on one fixed compact interval $J_* \in \mathbb{R}$, $M_n \geq \varepsilon_0 > 0$, and $R_n \rightarrow \infty$. After passing to a subsequence, the logarithmic shell compactification of the spacetime tail carries a probability limit state. Every point in the support of that limit is realized by a translated shell-cylinder extraction of the original represented branch, with no shell-radius Navier–Stokes dilation, and each realized support profile has a first closed local row in the ordered local state-space validation.

Proof. Apply lemma 10.93. This gives the compact metrizable shell-state space, a probability limit state, and an actual translated shell-cylinder realizing sequence for every support point.

Let Π_{shell} denote the weak-* limit shell state furnished by that lemma.

Let $\mathfrak{z} \in \text{supp } \Pi_{\text{shell}}$ and let $(W_m, \tilde{P}_m)$ be the corresponding realizing rebased shell sequence on a fixed model cylinder Q_{sh} , obtained only by spatial translation and finite time recentering. If $\mathfrak{z}$ lies on one of the failure faces of $\mathfrak{F}_{\text{shell}}$, then that face is already the first closed local row. Otherwise $(W_m, \tilde{P}_m)$ is a genuine represented local sequence with the pressure atlas, repaired gauge, local energy inequality, and compact-cylinder bounds on the fixed model cylinder. The ordered local rows are closed in this topology by lemma 10.72 and theorem 10.81, so the realized support profile has a first closed local row. This proves the support-transfer statement. $\square$

From this point onward, “rebased unit-shell” means a local state on the fixed model cylinder Q_{sh} obtained from an occupied shell block only by spatial translation and admissible finite time recentering. The shell radius is recorded as a logarithmic occupation-state coordinate and is not used as a Navier–Stokes scaling parameter.

Theorem 10.95 (Minimal mesoscopic extraction for diffuse critical tails). *Assume the support profiles furnished by [theorem 10.94](#) carry no positive compact retained local velocity activity at any predeclared dyadic ladder level η_m . Then, after passing to a subsequence, there exists either a realized support profile or a genuinely diffuse shell measure whose first closed local row is minimal among the atomic/activity and diffuse critical-tail alternatives. This first row is one of the named local rows below the compact retained single-core branch.*

Proof. The ordered local validation has finitely many rows. Apply [lemma 9.63](#) to the shell measures furnished by [theorem 10.94](#).

If the atomic/activity alternative occurs, choose among the realized support profiles with positive compact local activity one whose first closed local row is minimal in that finite ordered list. By the assumption that no support profile carries positive compact retained local velocity activity, the carrier rows of [lemma 10.103](#) do not occur for this minimal profile. Therefore its first row is one of the remaining named local rows: pressure/gauge/reconstruction failure, rough-core or local-energy failure, affine/pure-gauge behavior, hidden-scale/scale-collapse behavior, or the diffuse coherent critical-tail row.

If the genuinely diffuse alternative occurs, then the nonzero shell measure itself is the required diffuse critical-tail state, and its first closed row is the diffuse critical-tail row. This exhausts the minimal mesoscopic alternatives. $\square$

From this point onward, fix once and for all a dyadic activity ladder

$$\eta_m := 2^{-m}\eta_{\text{ref}}, \quad m \in \mathbb{N},$$

for some $\eta_{\text{ref}} > 0$. Fix also a countable family of nonnegative functions $\{\phi_j\}_{j \geq 1} \subset C_c^\infty(Q_{\text{sh}})$ with the following separation property: for every compact subcylinder $Q' \Subset Q_{\text{sh}}$ and every nonzero $f \in L_+^1(Q')$, there exists j with $\text{supp } \phi_j \Subset Q'$ and

$$\int_{Q_{\text{sh}}} \phi_j f > 0.$$

The local occupation-state coordinates are generated by the integrals

$$f \mapsto \int \phi_j f.$$

Positive local L^3 mass is recorded at the first basis pair (j, m) for which $\int \phi_j |W|^3 \geq \eta_m$.

Theorem 10.96 (Young-variance compactness dichotomy for realized diffuse tails). *Let W_n be a realized diffuse critical-tail extraction on a rebased unit-shell cylinder whose earlier representation, pressure, gauge, and rough-core rows have passed. Then exactly one of the following occurs:*

- (i) *one of the remaining hypotheses needed for the selected compactness package—namely the local modulation bound, the retained inner spacetime L^3 floor, or the selected upper local bound, or the selected upper local L^3 rows—fails on the rebased unit-shell window, and the first such failure is already a named lower local row;*
- (ii) *all hypotheses of the selected compactness package hold on the rebased unit-shell window, and [lemma 7.29](#) gives strong local L^3 -compactness on the rebased unit-shell windows.*

Proof. Passing the earlier representation, pressure, gauge, and rough-core rows gives the represented equation, local pressure reconstruction modulo functions of time, and local compact-cylinder H^1 control on the rebased unit-shell window. What remains to verify for the selected compactness package of [definition 7.27](#) are precisely the local modulation bound, the retained inner spacetime L^3 floor, and the upper local spacetime L^3 rows.

If the modulation bound fails, the ordered local validation records the autonomous-modulation or modulation-driven recentring row, which is one of the named lower rows in [lemma 7.32](#). If

the retained active spacetime L^3 floor is lost, then the branch enters the failure-of-retained-spacetime-activity / pressure-only critical-mass row, again named in [lemma 7.32](#). If the upper local L^3 rows fail, then [lemma 10.78](#) supplies the named lower local label. These are exactly the missing hypotheses noted in alternative (i).

If none of those failures occurs, then all clauses of the selected compactness package hold on the rebased unit-shell window. We may therefore apply [lemma 7.29](#), which yields strong local L^3 -compactness on the rebased unit-shell windows. This is alternative (ii). $\square$

Lemma 10.97 (Diffuse shell states enter the augmented residual coordinates). *Let $(W_n, \tilde{P}_n, a_n, b_n)$ be a realized shell-state sequence from [theorem 10.94](#) on a fixed unit-shell model cylinder Q_{sh} . Assume that:*

- (i) *the branch lies away from the pressure/gauge/local-energy failure faces, so the shell realization package supplies the local pressure splitting, repaired gauge coordinates, and local energy inequality on compact subcylinders of Q_{sh} ;*
- (ii) *case (ii) of [theorem 10.96](#) holds on compact subcylinders of Q_{sh} , so that after passing to a subsequence*

$$W_n \rightarrow W_\infty \quad \text{strongly in } L^3_{\text{loc}}(Q_{\text{sh}});$$

- (iii) *W_∞ carries positive local L^3 mass on some compact subcylinder $Q' \Subset Q_{\text{sh}}$:*

$$\int_{Q'} |W_\infty|^3 dy ds > 0.$$

Then there exists a basis test function ϕ_j from the fixed countable local occupation chart on Q_{sh} such that

$$\liminf_{n \rightarrow \infty} \int_{Q_{\text{sh}}} \phi_j |W_n|^3 dy ds > 0.$$

Consequently, after setting

$$U_n := W_\infty, \quad S_n := W_n - W_\infty,$$

the augmented tuple

$$(W_n, \tilde{P}_n, U_n, S_n, a_n, b_n)$$

belongs to the local occupation-state framework of [lemmas 10.71](#) and [10.72](#). Hence every positive nonretained local mass/pressure defect that survives after subtracting the strong local limit is visible either as a residual coordinate in that augmented state space or as an already named pressure/gauge/rough-core failure row. In particular, the shell branch legitimately enters the augmented residual coordinates of [theorem 10.86](#).

Proof. Choose a nonnegative $\phi \in C_c^\infty(Q_{\text{sh}})$ supported in Q' with

$$\int_{Q_{\text{sh}}} \phi |W_\infty|^3 dy ds > 0.$$

Strong local L^3 -convergence gives

$$\int \phi |W_n|^3 dy ds \rightarrow \int \phi |W_\infty|^3 dy ds > 0.$$

Because the augmented local occupation-state topology is built from a fixed countable family of nonnegative compact test functions on Q_{sh} , choose ϕ_j in that family so that ϕ_j is supported in Q' and $\int \phi_j |W_\infty|^3 > 0$. Then

$$\liminf_{n \rightarrow \infty} \int \phi_j |W_n|^3 dy ds > 0,$$

which is the required visible occupation coordinate.

By [theorem 10.94](#), the realized shell sequence is a represented local sequence on the fixed model cylinder and hence satisfies the repaired-gauge equation, the local energy inequality, pressure reconstruction modulo functions of time, and the compact-cylinder bounds on smaller compact windows. Since $U_n \equiv W_\infty$ and $S_n = W_n - W_\infty \rightarrow 0$ strongly in local L^3 , the tuple $(W_n, \tilde{P}_n, U_n, S_n, a_n, b_n)$ belongs to the augmented local occupation-state framework on every smaller compact subcylinder.

For the pressure, let $\tilde{P}_\infty^{\text{loc}}$ be the local Calderón–Zygmund pressure generated by W_∞ on a slightly larger compact cylinder. By [lemma 7.29](#), either

$$\tilde{P}_n - c_{n,R_1}(s) - \tilde{P}_\infty^{\text{loc}} \rightarrow 0 \quad \text{strongly in } L_{\text{loc}}^{3/2}$$

modulo functions of time, or the branch has already entered the named pressure/exterior row. Thus every surviving local pressure defect is visible in the same augmented local state space. Consequently, if a positive nonretained local diagnostic remains after subtracting the strong local limit, then on a smaller compact subcylinder it is either a nonzero residual coordinate for S_n in that augmented state space, or one of the already named pressure/gauge/rough-core failure faces. This is exactly the interface needed to apply [theorem 10.86](#) on the diffuse shell branch. $\square$

Lemma 10.98 (Strong diffuse-tail limits create a named row). *Let $(W_n, \tilde{P}_n)$ be a realized diffuse critical-tail extraction on a rebased unit-shell model cylinder Q_{sh} , and suppose case (ii) of [theorem 10.96](#) holds on compact subcylinders of Q_{sh} . Let*

$$W_n \rightarrow W_\infty \quad \text{strongly in } L_{\text{loc}}^3(Q_{\text{sh}}).$$

If $W_\infty \not\equiv 0$, then after shrinking to a smaller compact shell subwindow exactly one of the following ordered alternatives occurs:

- (i) *the pressure/exterior, repaired-chart, modulation, represented-equation, or local-energy limit passage first fails; its named local row occurs;*
- (ii) *all limit interfaces pass and W_∞ has a singular point. Then an actual CKN-active descendant of the full represented shell sequence is produced. The ordinary scale-center and critical-mass routing assigns that descendant to a named adjacent row or to the already selected comparable retained class;*
- (iii) *all limit interfaces pass and W_∞ has no singular point. Then there is a compact shell subwindow $Q' \Subset Q_{\text{sh}}$ on which $W_\infty \in L^\infty(Q')$, $W_\infty \not\equiv 0$, and $W_n \rightarrow W_\infty$ strongly in $L^3(Q')$. Hence the bounded-window routing applies. If this positive limit lies below the shell's previous activity level, the fixed occupation chart also records it, without creating another alternative: for some basis function ϕ_j and some level η_m of the predeclared dyadic activity ladder,*

$$\int_{Q_{\text{sh}}} \phi_j |W_\infty|^3 dy ds \geq \eta_m > 0,$$

$$\liminf_{n \rightarrow \infty} \int_{Q_{\text{sh}}} \phi_j |W_n|^3 dy ds \geq \eta_m/2.$$

Thus the same bounded carrier is visible at a predeclared lower level.

In particular, a strong diffuse-tail limit creates no additional regularity type and no unclassified terminal row.

Proof. Case (ii) of [theorem 10.96](#) says precisely that the selected compactness package holds on the rebased shell windows. Apply [lemma 7.30](#). A first failed limit interface gives alternative (i).

If the resulting suitable limit has a singular point, that lemma produces an actual CKN-active diagonal rebase of the full represented shell sequence. The routing corollary [corollary 10.79](#) then sends it through the compact-chart/tail partition: a noncompact relative parameter is a

named separated, exterior, or cascade row; a compact comparable carrier is already part of the compound retained class; and failure of a retained compact input is its named adjacent row. This is alternative (ii).

If the suitable limit has no singular point, the regular alternative of [lemma 7.30](#) supplies a compact subwindow Q' on which it is nonzero and bounded. Therefore [theorems 5.27](#) and [11.16](#) give the bounded-window route in alternative (iii). If the original shell bookkeeping has not registered this mass at its previous activity level, choose a nonnegative basis function ϕ_j supported in Q' with $\int \phi_j |W_\infty|^3 > 0$. The fixed dyadic ladder contains η_m with

$$0 < \eta_m \leq \int \phi_j |W_\infty|^3 dy ds.$$

Strong local L^3 -convergence then gives

$$\int \phi_j |W_n|^3 dy ds \longrightarrow \int \phi_j |W_\infty|^3 dy ds,$$

so the same bounded carrier is visible at the predeclared lower level. This does not split the regular case into a second mathematical alternative. The three outcomes are exclusive by first-failure order and by the singular versus nonsingular dichotomy. $\square$

Theorem 10.99 (Coherent critical-tail closure). *A nonzero realized diffuse critical-tail state cannot remain terminal on the retained compact single-core branch. More precisely, every such state either produces a lower named local row immediately, or else after strong local compactness its local limit enters the bounded-window/local-routing package and any surviving shell remainder is exhausted by the augmented residual state space. In particular no realized diffuse critical-tail state is terminal.*

Proof. Let W_n be a realized diffuse critical-tail state. By [theorem 10.96](#), either the selected compactness package fails and we immediately obtain a lower named local row, or else W_n is strongly compact in local L^3 on rebased unit-shell windows. Choose a compact unit-shell subwindow on which the strong local limit is nonzero. By [lemma 10.98](#), after shrinking that subwindow if necessary, either a compactness interface fails and gives its named row, the limit has a singular point and creates an actual active descendant routed by the ordinary compact-chart/tail partition, or the limit is regular and nonzero on a smaller cylinder and the bounded-window row closes that compact core. When the regular limit lies below the previous shell activity level, the predeclared dyadic ladder records the same bounded carrier at its first positive lower level; this is bookkeeping inside the regular alternative, not a fourth terminal state.

In the singular alternative, a compact-chart descendant assigned to the comparable retained class cannot remain a diffuse-tail state: comparable same-point carriers were grouped before the diffuse shell extraction. It is therefore reidentified with the existing compound retained class, while every noncomparable descendant is a named lower row.

If, after subtracting the strong local limit on such a compact subwindow, a positive nonretained shell remainder still survives, then [lemma 10.97](#) places that remainder inside the augmented residual coordinates of [theorem 10.86](#). That theorem yields either a named adjacent local row or a retained comparable class. A retained comparable class cannot occur for a diffuse tail, because comparable retained compact carriers were already grouped before the shell extraction and the diffuse shell selection was made outside every retained compact carrier. Hence only named lower rows remain.

Therefore no realized diffuse critical-tail state is terminal on the retained compact single-core branch. $\square$

Definition 10.100 (Windowed critical tightness). A selected terminal window modeled on a compact interval $J_* \in \mathbb{R}$ means a window in the local terminal validation family obtained from J_* by the admissible finite covariant recenterings of [lemma 10.70](#). A retained represented branch is windowed critically tight on the terminal tail if for every $\varepsilon > 0$ and every compact model interval $J_* \in \mathbb{R}$ there exists $R_{\varepsilon, J_*} < \infty$ such that for every selected terminal window J modeled on J_* ,

$$\int_J \int_{|y| > R_{\varepsilon, J_*}} |V(y, \tau)|^3 dy d\tau < \varepsilon.$$

Failure of windowed critical tightness means that there exist $\varepsilon_0 > 0$, a compact model interval $J_* \in \mathbb{R}$, selected terminal windows J_n modeled on J_* , and radii $R_n \rightarrow \infty$ such that

$$\int_{J_n} \int_{|y| > R_n} |V(y, \tau)|^3 dy d\tau \geq \varepsilon_0.$$

Remark 10.101 (Pointwise time-slice L^3 persistence is not used in the unconditional routing). The unconditional terminal routing does not convert a single-time L^3 tail lower bound into a spacetime row. Instead it works directly with the windowed spacetime tail defect of [definition 10.100](#), whose shell occupation measures are available without an additional time-persistence argument.

Lemma 10.102 (Windowed diffuse critical tail becomes a spacetime row). *Let V be a retained compact single-core branch, and suppose there exist $\varepsilon_0 > 0$, a compact model interval $J_* \in \mathbb{R}$, selected terminal windows J_n modeled on J_* , and radii $R_n \rightarrow \infty$ such that*

$$\int_{J_n} \int_{|y| > R_n} |V(y, \tau)|^3 dy d\tau \geq \varepsilon_0.$$

Assume no rebased unit shell cylinder carries a positive compact spacetime L^3 activity level that is already known to realize a retained carrier. Then exactly one of the following occurs:

- (i) *during shell extraction the represented local structure, pressure atlas, gauge chart, local energy, rough-core, or compact-upper package fails on the rebased shell cylinders, producing the corresponding named failure row;*
- (ii) *after passing to a subsequence, the normalized spacetime tail measures define a nonzero diffuse shell occupation coordinate in the logarithmic shell compactification;*
- (iii) *a parabolic shell support profile with positive compact spacetime activity is realized, reducing to the carrier case of [Lemma 10.103](#).*

Proof. Decompose the spacetime tail into dyadic shell blocks

$$A_{n,k} \times J_n = \{2^k R_n < |y| \leq 2^{k+1} R_n\} \times J_n.$$

Cover each occupied block by finitely overlapping unit parabolic cylinders in rebased shell coordinates. If the represented local structure or one of the pressure-atlas, gauge, local-energy, rough-core, or compact-upper inputs fails on those rebased shell cylinders, then the first failed test is exactly the corresponding named row. This is alternative (i).

Assume these failure faces do not occur. If some rebased unit shell cylinder carries a positive compact spacetime L^3 level, then [lemma 10.93](#) and [theorem 10.94](#) produce a realized shell support profile with positive compact activity on a fixed model cylinder. This is alternative (iii).

If no unit shell cylinder carries such a positive compact spacetime level, normalize the spacetime tail measures by their total mass and pass to the logarithmic shell compactification. Apply [lemma 9.63](#). If the atomic/activity alternative occurs, a realized support point with positive compact local activity appears on a fixed shell cylinder, returning us to alternative (iii). If the genuinely diffuse alternative occurs, then the nonzero limiting shell measure itself is a diffuse shell occupation coordinate. This is alternative (ii). $\square$

Lemma 10.103 (Failure of windowed critical tightness is a named local row). *Let V be a retained compact single-core branch in the renormalized variables, with scale function $\lambda(\tau)$. Suppose windowed critical tightness fails in the sense of Definition 10.100: there exist $\varepsilon_0 > 0$, a compact model interval $J_* \in \mathbb{R}$, selected terminal windows J_n modeled on J_* , and radii $R_n \rightarrow \infty$ such that*

$$\int_{J_n} \int_{|y| > R_n} |V(y, \tau)|^3 dy d\tau \geq \varepsilon_0.$$

Then, after passing to a subsequence, one of the following named rows occurs:

- (i) *a same-point larger-scale cascade row;*
- (ii) *a separated physical exterior concentration row;*
- (iii) *a noncompact escaping exterior row;*
- (iv) *a pressure/gauge/reconstruction failure row;*
- (v) *a rough-core, local-energy, or compact-upper failure row;*
- (vi) *a diffuse critical-tail row.*

In particular, failure of critical tightness cannot remain inside the retained compact single-core terminal stratum.

Proof. Decompose the spacetime tail

$$\{(y, \tau) : \tau \in J_n, |y| > R_n\}$$

into dyadic spacetime shell blocks

$$A_{n,k} \times J_n = \{2^k R_n < |y| \leq 2^{k+1} R_n\} \times J_n.$$

Cover each occupied shell by finitely overlapping translated shell cylinders modeled on Q_{sh} . If the represented local structure, pressure atlas, or repaired-gauge package fails during this extraction, we are in item (iv). If the local-energy, rough-core, or compact-upper package fails, we are in item (v).

Assume these failure faces do not occur. Suppose first that some rebased unit shell cylinder carries a positive compact spacetime L^3 mass level. Then lemma 10.93 and theorem 10.94 furnish a realized shell support profile coming from translated shell cylinders with centers (y_n, τ_n) in the original renormalized variables and $|y_n| \simeq R_n$. Set

$$d_n := \lambda(\tau_n) |y_n|.$$

Passing to a subsequence, exactly one of the following occurs. If $d_n \rightarrow 0$, the support profile stays at the same physical point while moving to a larger renormalized radius, giving item (i). If $d_n \rightarrow d \in (0, \infty)$, the profile lies at positive physical distance from the selected core, giving item (ii). If $d_n \rightarrow \infty$, the profile escapes every compact physical neighborhood in the selected chart, giving item (iii).

The only remaining possibility is that no positive local carrier exists at any fixed level. Normalize the spacetime tail measures by their total mass and pass to the logarithmic shell compactification. Apply lemma 9.63. If the atomic/activity alternative occurs, then theorem 10.94 realizes a shell support point with positive compact spacetime activity on a translated shell cylinder modeled on Q_{sh} , returning us to the carrier cases already handled above and therefore to one of the items (i)–(iii). If the genuinely diffuse alternative occurs, then the nonzero limiting shell measure itself is a genuinely diffuse shell occupation coordinate, yielding item (vi). This exhausts all possibilities. $\square$

Lemma 10.104 (Diffuse critical-tail row is closed). *A diffuse critical-tail defect produced by lemma 10.103 cannot remain as a terminal retained compact single-core branch. It either produces a lower named row already present in the state-space ledger, or it contradicts the minimal mesoscopic/critical-tail closure theorem.*

Proof. Let ν be the normalized limiting tail-defect measure produced in [lemma 10.103](#). By construction, ν is nonzero. Apply [lemma 10.93](#) to the logarithmic shell compactification. If the shell compactification hits a pressure, gauge, local-energy, or rough-core failure face first, that face is already the corresponding named local row.

Apply [lemma 9.63](#) to the logarithmic shell defect measures. If the atomic/activity alternative occurs, then [theorem 10.94](#) realizes a support point with positive compact local activity on a fixed shell model cylinder. This is a retained local carrier in the original branch, contrary to the diffuse case.

If the genuinely diffuse alternative occurs, then the nonzero diffuse shell measure itself is a visible diffuse occupation-state coordinate. If the selected compactness package fails before the strong-shell stage, the first failed test is already a named lower row. Otherwise the same augmented-shell mechanism used in [lemma 10.97](#) places this diffuse coordinate inside the residual coordinates of [theorem 10.86](#). That theorem yields either a named local row or a retained comparable class. The retained comparable class is impossible because comparable retained compact carriers were grouped before the shell extraction.

Thus a diffuse critical-tail defect cannot be terminal in the retained compact single-core stratum. $\square$

Lemma 10.105 (Windowed Caccioppoli with integrated modulation). *Let (V, P, a, b) be a suitable repaired-gauge branch on a terminal tail. Fix $R > 0$ and a unit window $I_n = [\tau_0 + n, \tau_0 + n + 1]$, and set*

$$I_n^* := [\tau_0 + n - 1, \tau_0 + n + 2], \quad A_{n,R} := 1 + \int_{I_n^*} (|a(\tau)| + \|b(\tau)\|_\infty) d\tau.$$

Assume on $I_n^ \times B_{4R}$ the local pressure reconstruction, the compact-cylinder critical spacetime bound*

$$\int_{I_n^*} \int_{B_{4R}} |V|^3 dy d\tau \leq M_{3,4R},$$

the local energy bound

$$\sup_{\tau \in I_n^*} \int_{B_{4R}} |V(y, \tau)|^2 dy \leq E_{2,4R},$$

and the local Caccioppoli estimate of [theorem 10.61](#). Then

$$\int_{I_n} \int_{B_R} |\nabla V|^2 dy d\tau \leq C_R (M_{3,4R} + E_{2,4R} (1 + A_{n,R})),$$

where C_R depends only on the fixed compact-cylinder cutoffs and ν .

Proof. Apply the local Caccioppoli estimate on $I_n^* \times B_{2R}$, with a cutoff equal to one on $I_n \times B_R$. The endpoint and cutoff-error terms are bounded by the local energy control:

$$\int_{I_n^*} \int_{B_{2R}} |V|^2 dy d\tau \leq C_R E_{2,4R}.$$

The cubic convection term is controlled directly by the spacetime L^3 bound:

$$\int_{I_n^*} \int_{B_{2R}} |V|^3 dy d\tau \leq M_{3,4R}.$$

The pressure term is controlled, after subtracting spatial means, by the local pressure reconstruction and Calderón–Zygmund on the larger ball:

$$\|P - c_{4R}\|_{L^{3/2}(I_n^* \times B_{2R})} \leq C_R \|V \otimes V\|_{L^{3/2}(I_n^* \times B_{4R})} \leq C_R M_{3,4R}^{2/3}.$$

Hence Hölder yields

$$\int_{I_n^*} \int_{B_{2R}} |P - c_{4R}| |V| dy d\tau \leq C_R M_{3,4R}.$$

The modulation terms in the local energy inequality have the form

$$\int_{I_n^*} (|a(\tau)| + \|b(\tau)\|_\infty) \int_{B_{2R}} |V(\tau)|^2 dy d\tau,$$

up to cutoff constants; the local $L_\tau^\infty L_y^2$ bound controls the inner integral by $E_{2,4R}$. Therefore

$$\int_{I_n^*} (|a(\tau)| + \|b(\tau)\|_\infty) \int_{B_{2R}} |V(\tau)|^2 dy d\tau \leq E_{2,4R} A_{n,R}.$$

Combining these estimates gives the displayed bound. No pointwise-in-time critical L^3 control or pointwise-in-time modulation bound is used. $\square$

Lemma 10.106 (Unbounded modulation windows are adjacent exits). *Let a physical branch be in the retained compact single-core stratum after the ordered local validation tests. Fix $R > 0$. Then either*

$$\sup_n \int_{\tau_0+n-1}^{\tau_0+n+2} (|a(\tau)| + \|b(\tau)\|_\infty) d\tau < \infty,$$

or the branch leaves the retained compact single-core stratum through one of the named modulation, scale-drift, modulation-driven recentering, exterior, or gauge-degenerate local alternatives.

Proof. If the displayed window integrals are unbounded, choose windows $I_{n_j}^*$ on which the integral tends to infinity. The coefficients a, b are the absolutely continuous scale and translation velocities of the repaired-gauge chart, as identified in [lemma 10.51](#). Large accumulated negative scale motion is the negative scale-drift branch of [definition 10.25](#); large positive scale-shell work is the positive-scale carrier branch routed by [corollary 8.131](#); and large translation work either is absorbed by the admissible recentering rule or produces the modulation-driven recentering/exterior label in [definition 8.110](#). If the accumulated modulation cannot be represented by the absolutely continuous chart variables, the first failed row is the repaired-gauge or gauge-degenerate alternative. These are all named adjacent local alternatives, not retained compact single-core states. Hence on the retained compact single-core stratum the window integrals are uniformly bounded. $\square$

Lemma 10.107 (Local windowed H^1 -loss routes to the rough-core alternative). *Assume the retained compact single-core stratum, with a suitable repaired-gauge representation, local pressure reconstruction, compact-cylinder critical spacetime bounds, local $L_\tau^\infty L_y^2$ energy control, and the local Caccioppoli estimate on the compact cylinders used in [lemma 10.105](#). These are the retained-stratum inputs supplied by the ordered tests in [lemma 8.73](#)(i),(ii),(iv),(vii) and by the suitable local energy package. Assume also that the unbounded-modulation window exit of [lemma 10.106](#) has not occurred. Then, for every $R > 0$, the local windowed H^1 bound*

$$\sup_{T \geq \tau_0+1} \int_T^{T+1} \int_{B_R} |\nabla V|^2 dy d\tau \leq C \left(R, \nu, M_{3,4R}, E_{2,4R}, \sup_n A_{n,R} \right)$$

holds, so the local windowed H^1 -control condition of [definition A.26](#) is automatic on the retained stratum. Equivalently, failure of local windowed H^1 control is incompatible with the retained-stratum input package and is routed to the rough-core alternative (iii) of [theorem 11.18](#), which by [theorem 11.7](#) forces failure of compact CKN control on enclosing cylinders and therefore returns to the multibubble/cascade alternative (ii), already excluded from the retained branch.

Proof. By lemma 10.106, absence of the modulation-window exit gives, for every fixed R ,

$$\sup_n A_{n,R} < \infty.$$

The compact-cylinder critical spacetime bound and local pressure reconstruction are retained-stratum inputs, together with the local $L_\tau^\infty L_y^2$ energy control. Applying lemma 10.105 on each unit window therefore gives a bound independent of n . This proves the displayed windowed H^1 estimate. The standard mollification and lower semicontinuity argument passes the estimate from smooth approximants to suitable repaired-gauge branches.

If, contrapositively, the windowed H^1 bound fails for some $R > 0$ along a terminal sequence, then by lemma 10.105 and lemma 10.106 at least one retained input must fail along that sequence: local pressure reconstruction, compact-cylinder critical spacetime control, the local $L_\tau^\infty L_y^2$ energy control, local Caccioppoli regularity, or the uniform integrated modulation-window bound. By the ordered tests lemma 8.73 this earlier-test failure routes the branch out of the retained stratum into the corresponding representation, repaired-gauge, critical-mass, local-regularity, or modulation/scale-drift alternative. Independently, in the local state-space terminology of theorem 11.18 the same failure is recorded as the rough-core alternative (iii), and by theorem 11.7 that alternative is equivalent to failure of compact CKN control on enclosing cylinders, hence to active multibubble or cascade behaviour, alternative (ii) of theorem 11.18. Both routings exit the retained single-core stratum. Therefore on the retained stratum H^1 -loss is impossible. $\square$

Theorem 10.108 (Retained compact-test closure). *On the retained compact single-core stratum, both compact tests of definition A.26 pass: critical tightness holds and the local windowed H^1 control holds. Equivalently, the entries (v),(vi) of lemma 8.73 are not genuine new assumptions on the retained stratum; their failure is routed to one of the named non-retained alternatives.*

Proof. Windowed critical tightness on selected terminal windows holds by lemma 10.103: failure routes the branch to a named local row. The carrier rows are handled by the same-point cascade, separated exterior, or noncompact exterior closures. The pressure/gauge and the rough-core, local-energy, and compact-upper rows are closed by the corresponding first-failure theorems. The diffuse tail case is closed by lemma 10.104. Local windowed H^1 control holds by lemma 10.107: failure routes the branch to the rough-core alternative, which by theorem 11.7 is absorbed into the multibubble/cascade alternative, also non-retained. Hence both compact tests are theorems, not assumptions, on the retained stratum. $\square$

Corollary 10.109 (Discharge of the retained compact-test closure row). *The retained compact-test closure obligation listed as item 3 of the local closure ledger is discharged on every retained compact single-core branch produced by the ordered local validation procedure of lemma 8.73. No additional global tightness or global windowed H^1 hypothesis is required.*

Proof. Theorem 10.108 reduces both compact tests to local consequences of inputs already present on the retained stratum or to named non-retained alternatives. No global tightness or H^1 bound is invoked. $\square$

10.11. Closure of the thin negative-drift exit. The remaining point in the scale-collapse drift branch is the case in which the weighted negative drift

$$g_{\text{sc}}(\tau) := a_-(\tau)M_R(\tau), \quad M_R(\tau) = \int |V(y, \tau)|^2 \Phi(\tau, y) dy,$$

has infinite tail integral but does not produce fixed-length windows with a uniform positive average. The following argument closes this row as an independent terminal alternative. It does not turn an infinite integral into a pointwise lower bound; instead it passes to normalized occupation blocks and uses the local state-space labels.

Definition 10.110 (Thin negative-drift occupation blocks). Let a represented branch satisfy the scale-collapse drift obstruction

$$\int_{\tau_0}^{\infty} g_{\text{sc}}(\tau) d\tau = \infty.$$

A sequence of intervals $[s_m, t_m]$, $s_m \rightarrow \infty$, is a normalized negative-drift occupation block sequence if

$$\int_{s_m}^{t_m} g_{\text{sc}}(\tau) d\tau = 1 \quad \text{for every } m.$$

It is diffuse if

$$\sup_{T \in [s_m, t_m]} \int_{[T, T+1] \cap [s_m, t_m]} g_{\text{sc}}(\tau) d\tau \rightarrow 0.$$

The associated occupation measures are

$$d\mu_m(\tau) = g_{\text{sc}}(\tau) \mathbf{1}_{[s_m, t_m]}(\tau) d\tau.$$

Definition 10.111 (Canonical thin-drift terminal schedule). In the canonical windowwise pre-carrier regime, after a harmless terminal restriction, the terminal unit-window schedule is

$$K_k = [\tau_{\sharp} + k, \tau_{\sharp} + k + 1], \quad k = 0, 1, 2, \dots$$

These are the canonical cells on which the windowwise moving family and the ordered local validation rows are tested.

Lemma 10.112 (Canonical terminal windows are admissible or already labelled). *Let a physical local Type II branch be in the canonical windowwise pre-carrier regime, and let K_k be the canonical cells of [definition 10.111](#). For every sequence K_{k_m} with $k_m \rightarrow \infty$, after passing to a subsequence, exactly one of the following alternatives holds:*

- (A1) *a non-thin first-failure row occurs on these windows: autonomous modulation failure, modulation-driven recentering, compactness-package first failure routed by [theorem 7.34](#), exterior or separated-core concentration, rough-core loss, multicore or same-point cascade, carrier/cost validation failure outside the canonical retained channel, or loss of the retained active core;*
- (A2) *after the finite covariant recentering $s = \tau - (\tau_{\sharp} + k_m)$, the canonical windows K_{k_m} are admissible selected terminal windows for the ordered local validation: the repaired-gauge representation, local pressure reconstruction, local suitable-energy bounds, moving L^2 core floor, retained local critical-mass row, compact tests, and modulation bookkeeping hold on every fixed compact cylinder, with constants independent of m .*

Proof. The canonical windowwise moving family is the schedule on which the carrier, compactness, repaired-gauge, pressure, and modulation rows are validated. The finite translation $\tau \mapsto \tau - (\tau_{\sharp} + k_m)$ does not transfer any lower bound from one physical time to another; it only writes the canonical cell K_{k_m} as the model interval $[0, 1]$ in its terminal chart, as in the finite recentering step of [lemma 10.70](#). Thus every row in the ordered validation list is tested on the same canonical cells.

Run the finite ordered validation list on K_{k_m} . If any required local row fails before retention, the first failed row is recorded by [lemmas 7.26](#) and [8.73](#) and belongs to the list in (A1) after excluding the thin-drift row itself. If no such first failure occurs, all retained local inputs listed in (A2) hold on the recentered windows. The moving L^2 floor is part of this retained local state data on the same canonical windows; if that floor fails, the ordered validation records the first such failure before any thin-drift conclusion. Thus only the stated alternatives remain. $\square$

Lemma 10.113 (Thin drift gives diffuse occupation blocks). *Assume $g_{\text{sc}} \geq 0$, $g_{\text{sc}} \in L^1_{\text{loc}}([\tau_0, \infty))$, and*

$$\int_T^\infty g_{\text{sc}}(\tau) d\tau = \infty \quad \text{for every } T \geq \tau_0.$$

Then every terminal tail contains a normalized negative-drift occupation block sequence. After passing to a subsequence, exactly one of the following holds:

(D1) *there are $\gamma > 0$ and unit windows $U_m = [\sigma_m, \sigma_m + 1]$, $\sigma_m \rightarrow \infty$, such that*

$$\int_{U_m} g_{\text{sc}}(\tau) d\tau \geq \gamma;$$

(D2) *the retained block sequence is diffuse in the sense of [definition 10.110](#).*

If the branch is in the thin-drift alternative of [theorem 7.24](#), then case (D1) either produces one of the non-thin first-failure rows of [lemma 10.112](#) (A1), or is impossible on the retained selected terminal tail. Hence, after non-thin first failures are removed, only (D2) remains.

Proof. Apply [lemma 8.123](#) with $g = g_{\text{sc}}$. Its block construction gives disjoint terminal intervals $[s_m, t_m]$ with $\int_{s_m}^{t_m} g_{\text{sc}} = 1$. The same lemma gives the two alternatives: persistent unit-window carrier mass or diffuse carrier mass. These are exactly (D1) and (D2) above.

If (D1) holds, compare U_m with the canonical schedule $K_k = [\tau_{\sharp} + k, \tau_{\sharp} + k + 1]$. Each unit interval U_m meets at most three canonical cells. Hence, after passing to a subsequence, there are canonical cells K_{k_m} with $k_m \rightarrow \infty$ and

$$\int_{K_{k_m}} g_{\text{sc}}(\tau) d\tau \geq \gamma/3.$$

Apply [lemma 10.112](#) to these canonical cells. If alternative (A1) occurs, the branch has already left the thin terminal row through a non-thin first failure. Otherwise the canonical cells are admissible selected terminal windows. On those selected windows,

$$\int_{K_{k_m}} a_-(\tau) M_R(\tau) d\tau \geq \gamma/3.$$

This is the negation of the thin-drift alternative in [theorem 7.24](#), which was precisely the absence of a uniform positive lower bound for fixed-length averaged weighted negative drift on the selected terminal tail. Thus a retained thin-drift branch with no non-thin first failure can only produce diffuse occupation blocks. $\square$

Lemma 10.114 (Small weighted negative drift is perturbative). *Let $I \subset [\tau_0, \infty)$ be an interval of length at most one. Assume that the branch has a moving L^2 core floor*

$$M_R(\tau) \geq m_0 > 0 \quad \text{for a.e. } \tau \in I,$$

and that on a compact set $K^+ \Subset \mathbb{R}^3$, containing a unit neighborhood of K , the retained local energy bound

$$\text{ess sup}_{\tau \in I} \|V(\tau)\|_{L^2(K^+)} \leq L$$

holds. If

$$\int_I g_{\text{sc}}(\tau) d\tau \leq \varepsilon,$$

then the negative scale correction

$$F_-(\tau) := a_-(\tau)(V(\tau) + y \cdot \nabla V(\tau))$$

satisfies

$$\|F_-\|_{L^1(I; H^{-1}(K))} \leq C_K L m_0^{-1} \varepsilon,$$

where C_K depends only on K and the fixed enlargement K^+ . Moreover, for every $\psi \in C_c^\infty(K)$,

$$\int_I |\langle F_-(\tau), V(\tau)\psi \rangle| d\tau \leq C_{K,\psi} L^2 m_0^{-1} \varepsilon.$$

Proof. The core floor gives

$$\int_I a_-(\tau) d\tau \leq m_0^{-1} \int_I a_-(\tau) M_R(\tau) d\tau \leq m_0^{-1} \varepsilon.$$

Let $\varphi \in C_c^\infty(K; \mathbb{R}^3)$ with $\|\varphi\|_{H_0^1(K)} \leq 1$. By integration by parts in y ,

$$\int_K (y \cdot \nabla V) \cdot \varphi dy = - \int_K V \cdot (y \cdot \nabla \varphi + 3\varphi) dy.$$

Hence, for a.e. $\tau \in I$,

$$|\langle V + y \cdot \nabla V, \varphi \rangle| \leq C_K \|V(\tau)\|_{L^2(K^+)} \|\varphi\|_{H_0^1(K)} \leq C_K L.$$

Taking the supremum over φ and integrating against $a_-(\tau)$ gives

$$\|F_-\|_{L^1(I; H^{-1}(K))} \leq C_K L \int_I a_-(\tau) d\tau \leq C_K L m_0^{-1} \varepsilon.$$

For the local energy pairing, integrate the scale-generator term by parts:

$$\int_K (y \cdot \nabla V) \cdot V \psi dy = \frac{1}{2} \int_K y \cdot \nabla (|V|^2) \psi dy = -\frac{1}{2} \int_K |V|^2 (3\psi + y \cdot \nabla \psi) dy.$$

Therefore

$$|\langle V + y \cdot \nabla V, V\psi \rangle| \leq C_{K,\psi} \|V(\tau)\|_{L^2(K^+)}^2 \leq C_{K,\psi} L^2.$$

Integrating against a_- and using $\int_I a_- \leq m_0^{-1} \varepsilon$ gives the second estimate. No bound on ∇V and no global norm is used. $\square$

Lemma 10.115 (Active core persistence on admissible drift windows). *Let $I_m = [\sigma_m, \sigma_m + 1]$ be admissible selected terminal windows in the sense of lemma 10.112 (A2). Then either the retained spacetime activity floor fails (equivalently, the loss-of-retained-active-core row occurs), or, after passing to a subsequence, there are a compact cylinder $Q_{r_*} = B_{r_*} \times I_{r_*} \subseteq B_R \times (0, 1)$ and $\eta_* > 0$ such that the represented sequence on I_m carries the retained active local mass floor*

$$\liminf_{m \rightarrow \infty} \int_{I_{r_*}} \int_{B_{r_*}} |V_m|^3 dy ds \geq \eta_*^3,$$

where $V_m(y, s) = V(y, \sigma_m + s)$ and $P_m(y, s) = P(y, \sigma_m + s)$. Consequently the minimal active-family extraction and the same-point/separated reductions are applicable on these windows; if pressure-only CKN detection is needed instead of velocity detection, the branch is routed to the pressure/exterior row rather than retained here.

Proof. The retained local critical-mass row in lemma 10.112 (A2) is interpreted by lemma 10.78. If the retained spacetime activity floor fails, the selected candidate is not a retained active component; this is exactly the loss-of-retained-active-core row.

If the retained spacetime activity floor does not fail, then theorem 9.33 supplies a positive active local spacetime velocity floor on a compact subcylinder of the selected terminal window. A finite cover of the retained compact core and the terminal maximality rule lemma 10.77 allow one fixed compact subcylinder Q_{r_*} and one threshold $\eta_* > 0$ after passing to a subsequence. If the witnesses escaped every such fixed subcylinder, the first label would be exterior/separated concentration, cascade, or compactness failure, all of which are non-thin rows already listed in lemma 10.112 (A1). If the active CKN diagnostic were carried only by pressure while the velocity floor above failed, then lemmas 9.21 and 9.22 routes the branch to the pressure/exterior row, not

to the retained thin-drift row. Thus the displayed active spacetime velocity floor holds on the retained branch.

The hypotheses of the minimal active-family extraction and of the same-point/separated reductions are exactly the repaired-gauge representation, local pressure reconstruction, compact-cylinder suitable bounds, and retained active local mass floor now available on Q_{r*} . Hence those reductions may be applied by the terminal active-frame machinery. $\square$

Lemma 10.116 (Vanishing negative scale forcing has no terminal label). *Let $I_m = [\sigma_m, \sigma_m + 1]$ be admissible selected terminal windows and assume that on every compact $K \Subset \mathbb{R}^3$*

$$a_-(\sigma_m + \cdot)(V_m + y \cdot \nabla V_m) \longrightarrow 0 \quad \text{in } L^1(0, 1; H^{-1}(K)).$$

Assume also the corresponding local energy pairings with compact cutoffs tend to zero. Then the negative scale-drift term contributes no nonzero local defect measure and no independent terminal state-space label. Any retained limit or occupation state obtained from the windows is the same local Navier–Stokes state that would be obtained after moving this term to a vanishing $L^1 H_{\text{loc}}^{-1}$ forcing; the ordered labels are therefore the ordinary active-core, compactness, exterior, multicore/cascade, scale-rigid, or bounded-window labels.

Proof. Let $\varphi \in C_c^\infty((0, 1) \times K; \mathbb{R}^3)$. The assumed $L^1 H^{-1}$ convergence gives

$$\int_0^1 \langle a_-(\sigma_m + s)(V_m + y \cdot \nabla V_m), \varphi(s) \rangle ds \longrightarrow 0.$$

Hence the negative scale term disappears in the distributional limit of the represented equation. The compact-cutoff energy pairings tend to zero by assumption, so the local energy inequality also passes to the same limit with no additional defect measure. The pressure reconstruction is unchanged: since $\nabla \cdot V_m = 0$, the field $V_m + y \cdot \nabla V_m$ is divergence-free in the spatial variables, so this lower-order velocity forcing contributes no new pressure source.

Thus, in the augmented local compactness topology, the extra coordinate recording the negative scale forcing converges to the zero face. Closedness of the ordered local diagnostics [lemmas 8.121](#) and [10.72](#) then implies that this vanishing coordinate cannot be the first nonzero label of a terminal state. The only possible first labels are the ordinary local state-space labels listed in the statement.

This is a zero-face statement for the augmented local state space, not a small-data statement for a velocity remainder. The proof uses only distributional closedness, compact-cutoff local-energy closedness, and the closedness of the already defined local diagnostics. No global Kato theory and no small $L^\infty L^3$ velocity remainder is invoked. $\square$

Lemma 10.117 (Diffuse thin drift routes to another local row). *Let a physical local Type II branch be in the canonical windowwise pre-carrier regime and suppose that, after the finite-cost and finite absolute-cost alternatives have been removed, it satisfies the scale-collapse drift obstruction and the thin-drift alternative. Let $[s_m, t_m]$ be the diffuse blocks supplied by [lemma 10.113](#). Then exactly one of the following occurs after passing to a subsequence:*

- (T1) *one of the non-thin local first-failure rows occurs: autonomous modulation failure, modulation-driven recentering, compactness-package first failure routed by [theorem 7.34](#), exterior or separated-core concentration, rough-core loss, multicore or same-point cascade, carrier/cost validation failure outside the canonical retained channel, or loss of the retained active core;*
- (T2) *a closed retained row occurs: compact single-core, scale-rigid residual behavior, retained compact scale-collapse, or a nonzero bounded selected-window limit;*
- (T3) *the negative scale-drift term is perturbative on the canonical selected compact unit windows and no additional terminal thin-drift state remains.*

Proof. Because the blocks are diffuse,

$$\varepsilon_m := \sup_{T \in [s_m, t_m]} \int_{[T, T+1] \cap [s_m, t_m]} g_{\text{sc}}(\tau) d\tau \longrightarrow 0.$$

The block lengths tend to infinity. Indeed, if a subsequence of blocks had length bounded by L , each such block would be covered by at most $\lceil L \rceil + 2$ unit intervals, and its total g_{sc} -mass would be bounded by $(\lceil L \rceil + 2)\varepsilon_m \rightarrow 0$, contradicting the normalization $\int_{s_m}^{t_m} g_{\text{sc}} = 1$. Hence, after discarding a finite prefix, each block contains a canonical unit cell

$$K_{k_m} = [\tau_{\sharp} + k_m, \tau_{\sharp} + k_m + 1] \subset [s_m, t_m], \quad k_m \rightarrow \infty.$$

The same estimate holds for every canonical cell contained in the block. In particular, since $K_{k_m} \subset [s_m, t_m]$, diffuseness gives

$$\int_{K_{k_m}} g_{\text{sc}} \leq \varepsilon_m.$$

Apply [lemma 10.112](#) to the canonical cells K_{k_m} . If alternative (A1) occurs, the branch is in (T1). Hence assume that alternative (A2) holds: after recentring, K_{k_m} are admissible selected terminal windows with the retained local representation, pressure, energy, compactness, modulation, and moving core-floor data. Write $I_m := K_{k_m}$ for these canonical selected windows.

On these admissible windows, the moving L^2 core floor and the retained local $L_t^\infty L_x^2$ bounds are part of [lemma 10.112](#) (A2). Hence [lemma 10.114](#) implies

$$\|a_-(V + y \cdot \nabla V)\|_{L^1(I_m; H^{-1}(K))} \longrightarrow 0$$

for every compact $K \Subset \mathbb{R}^3$, and its compact-cutoff local energy pairings also vanish. By [lemma 10.116](#), the diffuse negative drift is a vanishing local forcing on the selected unit windows; it is not a new compact state-space object and cannot be a terminal first label.

It remains to classify the retained compact sequence on these same admissible windows. Apply [lemma 10.115](#). If the retained spacetime activity floor is lost, the branch is in (T1). Otherwise a fixed compact subcylinder carries a positive retained local spacetime mass floor, so the structural reductions are applicable. Apply the minimal active-family extraction and the same-point/separated-point reductions given by [theorems 5.11, 5.20](#) and [5.21](#). If those reductions produce another active core, a strict same-point cascade, a separated physical point, exterior pressure persistence, rough-core loss, or loss of compact estimates, then the branch is in (T1). If they produce the compact single-core row, the retained compact scale-collapse row, or the scale-rigid residual row, those rows are the closed alternatives in (T2), closed respectively by [theorems 7.37](#) and [11.5](#) and [proposition 5.42](#). In the scale-rigid case, if the compact limit has no further CKN subconcentration, [lemma 5.37](#) gives a nonzero bounded selected-window limit; if it has a further subconcentration, the same lemma routes it to the active structural rows already listed in (T1). Finally, [theorem 5.27](#) excludes the bounded-limit case from the local Type II class.

Thus, after all ordered first failures and closed retained rows have been accounted for, the only remaining effect of thin negative drift on the selected unit windows is the vanishing H_{loc}^{-1} perturbative term proved above. By [lemma 10.116](#), such a vanishing forcing carries no independent terminal state-space label. This proves (T3) and the asserted trichotomy. $\square$

Theorem 10.118 (Thin negative drift is not a terminal exit). *On every physical local Type II branch in the canonical windowwise pre-carrier regime, the thin negative-drift alternative is not an independent remaining terminal exit. More precisely, if the branch enters the scale-collapse drift obstruction and does not have persistent fixed-length negative-drift windows, then the branch either enters one of the other named local exit rows, or it enters an already closed retained row. Consequently the row “thin drift or negative scale-drift without thick autonomous windows” is not an unresolved entry in the later remaining-exit ledger.*

Proof. The scale-collapse drift obstruction gives the infinite tail integral of g_{sc} . By lemma 10.113, every terminal tail has either persistent unit-window weighted negative drift or diffuse normalized drift blocks. In the first alternative, finite overlap with the canonical schedule of definition 10.111 gives canonical cells carrying a uniform positive weighted negative-drift lower bound, unless one of the non-thin first-failure rows in lemma 10.112(A1) has already occurred. If such a first failure occurs, the branch has left through another named row. Otherwise those canonical cells are admissible selected terminal windows, and the branch is in the thick-window ledger of theorem 7.24: selected-window compactness can fail, autonomous modulation convergence can fail, or the retained thick compact scale-collapse state is reached. The first two are other named rows, and the retained thick state is closed by theorem 7.37. No part of this persistent-window branch validates arbitrary intervals or uses the diffuse thin-drift closure itself.

On a genuine thin-drift tail, only the diffuse block alternative remains. Lemma 10.117 then gives three possibilities. The first is one of the other named local exit rows. The second is a closed retained row, including the bounded selected-window exit excluded by theorem 5.27. The third says that the negative drift is a vanishing local H^{-1} perturbation on the selected compact windows and therefore does not define a terminal state-space row after the retained perturbative remainder has been formed.

Thus no local Type II branch can remain whose first and only unresolved label is thin negative drift. The argument used only finite-window local estimates, occupation blocks, the retained spacetime activity floor, and the ordered local state-space partition; it did not use a global L^3 bound, a whole-space profile decomposition, or a hidden conversion of diffuse drift into thick windows. $\square$

10.12. Closure of the Type II analytic-data exhaustiveness row. The exhaustion row behind theorem A.4 is the entry condition for every admissible local Type II branch. It is closed by an ordered local extraction package: each component of definition A.2 is either supplied by an existing local lemma on the admissible class $\mathcal{U}_{\Pi}^{NS}$, or its first failure is routed to a named alternative that is class-incompatible or already covered by the local state-space stratification. No global critical profile theorem and no global L^3 bound are invoked.

Lemma 10.119 (Concentration extraction routes). *Let $\omega = (u, p, T^*) \in \mathcal{U}_{\Pi}^{NS}$. Exactly one of the following holds:*

- (a) *the local concentration dichotomy theorem 3.6 produces $x_0 \in \Sigma(T^*)$, $\eta > 0$, and $r_n \downarrow 0$ with $C((x_0, T^*), r_n) + D((x_0, T^*), r_n) \geq \eta$ for all n , and the concentration point lies in the local Type II alternative;*
- (b) *the singular set $\Sigma(T^*)$ is empty;*
- (c) *$\Sigma(T^*) \neq \emptyset$, and every concentration point of $\Sigma(T^*)$ lies in the local Type I alternative.*

Cases (b) and (c) are class-incompatible with definition A.1. Hence on $\mathcal{U}_{\Pi}^{NS}$ only case (a) occurs, and the first ordered failure of definition A.3 (“concentration extraction fails”) does not arise on the admissible class.

Proof. If $\Sigma(T^*) = \emptyset$, case (b) holds. If $\Sigma(T^*) \neq \emptyset$, apply the local concentration dichotomy at singular points. If some singular point is not in the local Type I alternative, choose that point as x_0 ; the Caffarelli–Kohn–Nirenberg failure at x_0 supplies $\eta > 0$ and $r_n \downarrow 0$ with

$$C((x_0, T^*), r_n) + D((x_0, T^*), r_n) \geq \eta,$$

and case (a) holds. If no such point exists, every concentration point lies in the local Type I alternative and case (c) holds. The three cases are therefore disjoint and exhaustive.

In case (b) there is no terminal singular concentration point, so the branch is the local continuation/no-terminal-concentration alternative rather than a local Type II branch. This is

excluded by clause (iv) of [definition A.1](#); no global continuation estimate is used. In case (c) every concentration point lies in the local Type I alternative of [theorem 3.6](#), which is exactly the local Type I scale-invariant criterion. This contradicts clause (iii) of [definition A.1](#). In case (a) the dichotomy supplies the positive CKN concentration scales and the exclusion of the Type I alternative at (x_0, T^*) ; this is the content of the first analytic datum of [definition A.2](#). $\square$

Lemma 10.120 (Supercritical scale alternative routes). *Let $\omega \in \mathcal{U}_{\text{II}}^{NS}$ and let (x_0, T^*) be a concentration point produced by [lemma 10.119](#)(a). Then the supercritical scale alternative of [definition A.2](#)(ii) holds at (x_0, T^*) .*

Proof. The dichotomy of [theorem 3.6](#) at (x_0, T^*) is exclusive: either a local Type I scale-invariant bound holds, or no such local Type I bound holds. The first option contradicts clause (iii) of [definition A.1](#), which excludes branches in the local Type I class. Hence the second option holds at every concentration point produced by the previous lemma; this is exactly the supercritical Type II scale alternative. The second ordered failure of [definition A.3](#) therefore does not arise on $\mathcal{U}_{\text{II}}^{NS}$. $\square$

Lemma 10.121 (Local profile data routes). *Let $\omega \in \mathcal{U}_{\text{II}}^{NS}$ and let (x_0, T^*) be a concentration point in the local Type II alternative as produced by [lemmas 10.119](#) and [10.120](#). Run the local scale-translation selection of [proposition 3.7](#) at the positive concentration scales and test whether it produces a nondegenerate compact single core. Then exactly one of the following holds:*

- (a) *the localized scale and center are nondegenerate and no represented selected-window subsequence has a nonzero bounded selected-window limit on any compact cylinder; [proposition 3.8](#) supplies represented variables (V_n, P_n, a_n, b_n) on compact cylinders satisfying the represented Navier–Stokes equation [\(14\)](#); the represented selected windows form a local Type II sequence in the sense of [definition 3.2](#), providing the local profile data of [definition A.2](#)(iii);*
- (b) *the localized nondegeneracy fails, and the branch falls into the multibubble, cascade, or gauge-degenerate alternative of [definition 5.6](#)(ii)–(v);*
- (c) *nondegeneracy holds but the represented selected windows admit a nonzero bounded selected-window limit on some compact cylinder; the branch then enters the bounded-selected-window alternative of [definition 3.1](#) and leaves the strict selected-window branch by [theorem 5.27](#), after which the physical bounded-window event is routed by [theorem 11.16](#).*

In cases (b) and (c) the branch is non-retained: case (b) is treated by the multibubble/cascade machinery of [Section 5](#), and case (c) is the bounded selected-window exit of the scale-rigid stratum. Neither case is a compact single-core branch.

Proof. The local scale-translation selection at (x_0, T^*) attempts to extract a single compact core at the positive CKN concentration scales $r_n \downarrow 0$. Either the localized scale and center are nondegenerate (cases (a) and (c)), or they are degenerate (case (b)); under nondegeneracy, either no represented selected-window subsequence has a nonzero bounded selected-window limit (case (a)), or some such subsequence does (case (c)). The three cases are exhaustive.

In case (a), [proposition 3.7](#) applies and rewrites the shrinking cylinders $Q_{r_n}(x_0, T^*)$ as fixed selected windows. The represented sequence has positive local CKN concentration on a fixed compact cylinder; combined with the defining case-(a) absence of nonzero bounded selected-window limits, it is a local Type II sequence in the sense of [definition 3.2](#). [Proposition 3.8](#) then supplies the represented variables (V_n, P_n, a_n, b_n) on compact cylinders satisfying [\(14\)](#); this is the local profile datum of [definition A.2](#)(iii).

In case (b), failure of localized nondegeneracy is, by the second sentence of [proposition 3.8](#), the multibubble, cascade, or gauge-degenerate alternative. The corresponding stratum is named

in [definition 5.6](#) (ii)–(v) and is non-retained: it is handled by the multibubble/cascade machinery and never enters the retained compact single-core stratum produced by [lemma 8.73](#).

In case (c), the represented selected-window sequence has a nonzero bounded selected-window limit on some compact cylinder, placing the branch in [definition 3.1](#); [theorem 5.27](#) then routes the branch out of Type II via the scale-rigid bounded-window exit, which is also non-retained.

Hence the third ordered failure of [definition A.3](#) (“local profile data are insufficient”) is routed to the multibubble/cascade/gauge-degenerate non-retained alternative or to the scale-rigid bounded-window exit; on the retained compact single-core stratum it does not arise. $\square$

Theorem 10.122 (Type II analytic-data routing on the admissible class). *The Type II analytic data of [definition A.2](#) are exhaustive by ordered local routing on $\mathcal{U}_{\text{II}}^{NS}$: every $\omega \in \mathcal{U}_{\text{II}}^{NS}$ either has the concentration profile, the supercritical scale alternative, and the local profile data, or its first applicable failure in [definition A.3](#) (i)–(iv) is matched to a named alternative that is class-incompatible with [definition A.1](#) or to a non-retained state-space branch already covered by the local stratification. In particular, after those named exits are removed, every retained admissible branch has the three Type II analytic data.*

Proof. Process the three analytic data of [definition A.2](#) in the order specified by [definition A.3](#).

Datum (i): concentration extraction. [Lemma 10.119](#) shows that on $\mathcal{U}_{\text{II}}^{NS}$ the only realised case is the dichotomy case (a) producing (x_0, T^*) with positive CKN concentration in the local Type II alternative; the other two cases are class-incompatible.

Datum (ii): supercritical scale alternative. [Lemma 10.120](#) then forces the supercritical Type II scale alternative at (x_0, T^*) ; the Type I alternative is class-incompatible by clause (iii) of [definition A.1](#).

Datum (iii): local profile data. [Lemma 10.121](#) either supplies the local profile data via the local scale-translation selection and the representation input [proposition 3.8](#), or routes the branch to the multibubble/cascade/gauge-degenerate non-retained alternative [definition 5.6](#) (ii)–(v), or to the bounded selected-window exit [theorem 5.27](#).

Residual datum (iv). The fourth ordered failure of [definition A.3](#), “the admissible Type II branch lies outside the concentration, scale, and profile alternatives under consideration”, is the catch-all for a branch that is not one of the local Navier–Stokes Type II alternatives after the first three tests have been applied. Clause (iv) of [definition A.1](#) excludes continuation, Type I, boundary, and non-Navier–Stokes-domain alternatives. Any remaining local branch outside the retained profile alternatives is one of the named non-retained state-space exits produced in the preceding steps.

In every case the branch either has all three analytic data or its first ordered failure is recorded as a class-incompatible alternative or as a non-retained state-space branch already covered by the local stratification. This is the ordered local-routing exhaustiveness statement associated with [corollary A.5](#). Removing the named exits leaves only the first alternative of that corollary, namely the full Type II analytic data. $\square$

Corollary 10.123 (Type II analytic-data exhaustiveness discharged). *The ordered-routing exhaustion input used by [theorem A.4](#) is a theorem on $\mathcal{U}_{\text{II}}^{NS}$: the analytic-data entry row is discharged by ordered local extraction. On retained admissible branches the three Type II analytic data hold; non-retained failures are assigned to named local alternatives before the downstream compact analysis is invoked. No global critical profile theorem and no global L^3 bound are invoked.*

Proof. [Theorem 10.122](#) establishes ordered local routing of all failures of the Type II analytic data on $\mathcal{U}_{\text{II}}^{NS}$. By [corollary A.5](#), after the named routed exits are removed, the remaining retained branch is in the first alternative and therefore has the Type II analytic data required by

[theorem A.4](#). Each step uses only the local concentration dichotomy [theorem 3.6](#), the local scale-translation selection [proposition 3.7](#), the local representation input [proposition 3.8](#), the local Type II definition [definition 3.2](#), the admissibility clauses of [definition A.1](#), and the multibubble/cascade stratum [definition 5.6](#), together with the bounded-window exit [theorem 5.27](#). $\square$

Definition 10.124 (Terminal local compactness input hypotheses). The terminal local compactness *input* hypotheses consist of the local compact terminal-window hypothesis [definition 10.89](#) and local perturbative stability [definition 10.91](#), together with item (i) of [definition 10.68](#) (the admissibility estimates on the represented branch). These hypotheses do *not* include the local terminal completeness conclusion (iii) of [definition 10.68](#).

Definition 10.125 (Terminal local compactness package). The terminal local compactness package is the input hypotheses [definition 10.124](#) together with the local terminal completeness theorem [theorem 10.87](#). This is a record of what has been proved and is not an additional assumption.

Theorem 10.126 (Terminal critical local compactness). *Assume the terminal local compactness input hypotheses of [definition 10.124](#). Assume also that the ordered adjacent local labels are routed as in [theorem 10.81](#) and that comparable retained classes are grouped before the terminal frame is fixed. Then, on every retained terminal branch and every compact terminal cylinder, all nonretained residual components vanish in the local perturbative topology of [definition 9.13](#).*

Proof. The local compact terminal-window hypothesis gives the compact-cylinder estimates. Applying [theorem 10.87](#) under the stated routing and grouping gives

$$S_n \rightarrow 0 \quad \text{in } L^\infty(I; L^3(B_R))$$

on each compact terminal cylinder. Local perturbative stability then converts this velocity convergence into the nonlinear-stress, pressure, and residual convergences required in terminal nonlinear decoupling. $\square$

Corollary 10.127 (Terminal completeness from local compactness hypotheses). *The terminal local completeness statement in the final argument may be replaced by the terminal local compactness package of [definition 10.125](#).*

Proof. [Theorem 10.126](#) gives the asserted terminal local completeness. $\square$

Remark 10.128 (Boundary of the terminal compactness hypothesis). The preceding compactness conclusion applies only to terminal windows on which the repaired-gauge representation, pressure reconstruction, local Caccioppoli control, bounded modulation, and local state-space labels are available. If one of these inputs is absent, the branch is assigned to the corresponding representation, pressure, rough-core, modulation, or state-space alternative; it is not treated by importing a global profile theorem.

Definition 10.129 (Terminal sequences sampled from the renormalized orbit). A terminal active sequence is sampled from the renormalized orbit if it is of the form

$$u_{0,n} = \Lambda_{\lambda_n, x_n} V(\tau_n), \quad \tau_n \rightarrow \infty, \quad \lambda_n > 0, \quad x_n \in \mathbb{R}^3,$$

or by the equivalent inverse critical change of variables, where

$$(\Lambda_{\lambda, x_0} f)(x) = \lambda^{-1} f\left(\frac{x - x_0}{\lambda}\right).$$

Lemma 10.130 (Critical symmetry preserves the L^3 norm). *For every $f \in L^3(\mathbb{R}^3)$, $\lambda > 0$, and $x_0 \in \mathbb{R}^3$,*

$$\|\Lambda_{\lambda, x_0} f\|_{L^3} = \|f\|_{L^3}.$$

Proof. By the change of variables $x = x_0 + \lambda y$,

$$\|\Lambda_{\lambda, x_0} f\|_{L^3}^3 = \int_{\mathbb{R}^3} \lambda^{-3} \left| f\left(\frac{x - x_0}{\lambda}\right) \right|^3 dx = \int_{\mathbb{R}^3} |f(y)|^3 dy.$$

Taking cube roots gives the claim. $\square$

Theorem 10.131 (Bounded terminal sequences from local critical bounds). *Assume the renormalized orbit satisfies the local compact-window bounds*

$$0 < \inf_{\tau \in I} \|V(\tau)\|_{L^3(B_R)} \leq \sup_{\tau \in I} \|V(\tau)\|_{L^3(B_{2R})} < \infty,$$

on the terminal interval under consideration, together with the local Caccioppoli and pressure bounds on $Q_{2R, I}$. If every terminal active window is sampled from the renormalized orbit, then every terminal active window satisfies the local compact terminal-window hypothesis.

Proof. The local L^3 upper bound is one component of the compact-window hypothesis. Local $L^\infty L^2$, $L^2 H^1$, and pressure bounds are supplied by the local Caccioppoli estimate and pressure reconstruction on $Q_{2R, I}$. The repaired-gauge equation gives the required time-derivative bound in H^{-m} . The lower bound records that the retained window is nontrivial. $\square$

Theorem 10.132 (Terminal compactness alternatives). *For a candidate terminal active sequence in the Type II analysis, the following alternatives exhaust the possible local compactness outcomes, after passing to a subsequence when necessary:*

- (i) *the sequence is not sampled from the renormalized orbit;*
- (ii) *the repaired-gauge representation, pressure reconstruction, or local Caccioppoli estimate is unavailable on the selected compact cylinder;*
- (iii) *the local L^3 upper bound fails along a terminal sampling subsequence;*
- (iv) *the retained local L^3 lower bound vanishes along a terminal sampling subsequence;*
- (v) *after excluding the preceding alternatives, the terminal active sequence satisfies the local compact terminal-window hypothesis.*

Proof. If the terminal sequence is not sampled from the renormalized orbit, then (i) holds. If one of the local PDE inputs is missing, then (ii) holds and the branch is assigned to the corresponding local alternative. If the local critical upper bound fails, then (iii) holds after passing to a subsequence. If the retained lower bound vanishes, then (iv) holds. Otherwise the local hypotheses of [theorem 10.131](#) hold, and (v) follows. $\square$

Definition 10.133 (Terminal coordinate-frame construction). The terminal coordinate-frame construction means that every terminal active frame used in the analysis is obtained from the repaired-gauge orbit as follows. Choose sampling times $\tau_n \rightarrow \infty$ and critical symmetry parameters (λ_n, x_n) ; form the local terminal windows by applying the critical coordinate-frame map to V on compact time intervals around τ_n ; group comparable retained local classes based at the same physical point before selecting the terminal active frame. No external data are introduced.

Theorem 10.134 (Terminal coordinate-frame construction from local state selection). *Assume terminal active frames are selected by first partitioning local active states by physical concentration point, then grouping comparable retained states at the same point into compound local states, and finally choosing a minimal-scale retained class at that point. Then the terminal coordinate-frame construction of [definition 10.133](#) holds for the selected terminal frames.*

Proof. By the stated selection rule, each terminal active frame is attached to a compound active class after all comparable retained states at the same physical point have been grouped. The scale and center of the terminal frame are the scale and center of this compound class. The

terminal data are then obtained by applying the Navier–Stokes critical change of variables to the local renormalized branch along the chosen sampling sequence $\tau_n \rightarrow \infty$. Thus the construction uses only the repaired-gauge orbit and the Navier–Stokes critical symmetries, and it satisfies all clauses of [definition 10.133](#). $\square$

Theorem 10.135 (Terminal coordinate frames are sampled from the orbit). *Assume the terminal coordinate-frame construction. Then every terminal active sequence belongs to the class described in [definition 10.129](#).*

Proof. By definition of the terminal coordinate-frame construction, each terminal active frame is obtained by choosing times τ_n , selecting scale and center parameters (λ_n, x_n) , and applying the Navier–Stokes critical change of variables to the renormalized solution $V(\tau_n)$. Hence the associated terminal data have the form

$$u_{0,n} = \Lambda_{\lambda_n, x_n} V(\tau_n),$$

or the equivalent inverse formulation, which is the orbit-sampling condition. $\square$

Theorem 10.136 (Repaired-gauge representation gives terminal orbit sampling). *Assume a repaired-gauge representation of the Type II branch and the terminal coordinate-frame construction. Then every terminal active sequence is sampled from the renormalized orbit.*

Proof. The repaired-gauge representation supplies the renormalized orbit $V(\tau)$ and the critical coordinate maps between physical variables and renormalized variables. The terminal coordinate-frame construction supplies the terminal sampling times and scale-center parameters. Combining these hypotheses gives

$$u_{0,n} = \Lambda_{\lambda_n, x_n} V(\tau_n),$$

which is the assertion of [definition 10.129](#). $\square$

Corollary 10.137 (Failure of terminal orbit sampling). *If a sequence reaches the terminal local-state stage but is not sampled from the renormalized orbit, then either the repaired-gauge representation is not available for that sequence or the terminal coordinate-frame construction has not been applied as specified.*

Proof. This is the contrapositive of [theorem 10.136](#). $\square$

10.13. Terminal Retained-State Assembly.

Definition 10.138 (Abstract terminal elimination data). The abstract terminal elimination data for this section are:

- (i) every renormalized local Type II branch enters one of the local alternatives used in [theorem 11.18](#);
- (ii) compact single-core vanishing-dissipation branches are excluded by [theorem 11.5](#);
- (iii) multibubble, cascade, and gauge-degenerate branches are excluded by [theorem 11.6](#), and scale-rigid terminal branches are excluded by [proposition 5.42](#);
- (iv) finite-cost scale-collapse branches are excluded by [theorem 11.8](#), while compact retained single-core branches with divergent localized cost are covered by [theorem A.21](#) and the cost-divergence theorem [theorem 10.42](#) and [corollary 10.43](#); scale-collapse branches with divergent negative scale-drift cost are covered by [theorem 10.33](#) or [corollary 10.37](#) and then excluded by the same cost-divergence theorem;
- (v) failure of the windowed critical-tightness test on selected terminal windows is excluded, equivalently that test holds for every renormalized branch in the class;
- (vi) failure of local windowed H^1 control is excluded, equivalently the required local windowed H^1 control holds for every renormalized branch in the class;

Theorem 10.139 (Abstract terminal Type II exclusion). *Under the abstract terminal elimination data of [definition 10.138](#), no renormalized local Type II branch in the stated class remains.*

Proof. Let a renormalized local Type II branch be given. By item (i), it lies in one of the alternatives in [theorem 11.18](#). Item (ii) excludes the vanishing-dissipation compact single-core subcase, while item (iv) excludes the retained compact single-core subcase with divergent localized cost. Item (iii) excludes multibubble, cascade, gauge-degenerate, and scale-rigid alternatives. Item (iv) also excludes finite-cost scale collapse and the divergent negative-drift scale-collapse branch whenever the corresponding cost is one of the costs covered by the cost-divergence theorem [theorem 10.42](#) and its status corollary [corollary 10.43](#). Items (v) and (vi) exclude failure of the windowed critical-tightness test and failure of local windowed H^1 control. These cases exhaust the alternatives, so no branch remains. $\square$

Definition 10.140 (Expanded terminal data). The following analytic data imply the abstract terminal elimination data:

- (i) the local alternative decomposition for renormalized Type II branches;
- (ii) the compact single-core vanishing-dissipation, multibubble/cascade, finite-cost scale-collapse, compact-cost-divergence, and scale-rigidity exclusions from the preceding sections;
- (iii) a critical-tightness estimate giving the windowed critical-tightness condition on selected terminal windows;
- (iv) an exhaustive terminal compactness partition and a local finite critical-mass statement for retained active terminal strata;
- (v) local perturbative residual stability and terminal state-space completeness on compact windows;
- (vi) removal of physically exterior regular components;
- (vii) the repaired-gauge representation theorem;
- (viii) the local Caccioppoli/windowed regularity input;
- (ix) the local compact scale-collapse routing theorem [theorem 7.37](#).

Lemma 10.141 (Terminal data imply critical tightness). *The expanded terminal data imply the windowed critical-tightness condition of [definition A.26](#).*

Proof. The repaired-gauge representation turns each renormalized orbit into the terminal windows used in the local state-space analysis. The terminal compactness partition, finite critical-mass statement, local perturbative stability, and terminal state-space completeness give the required terminal completeness. Exterior regular components are removed by the exterior-removal datum in [definition 10.140](#). Together with the multibubble and scale-collapse exclusions, this leaves only possible loss of windowed critical tightness on selected terminal windows; the critical-tightness estimate excludes that loss. $\square$

Lemma 10.142 (Terminal data imply windowed H^1 control). *The expanded terminal data imply local windowed H^1 control.*

Proof. The windowed H^1 estimate requires a repaired-gauge branch, local pressure and modulation control, compact-cylinder critical spacetime bounds, the local $L^\infty_\tau L^2_y$ energy control, and the Caccioppoli/windowed regularity input. These are among the expanded terminal data. Hence the required local windowed H^1 control holds, and the failure of local windowed H^1 control is unavailable. $\square$

Lemma 10.143 (Expanded terminal data imply abstract terminal data). *The expanded terminal data imply the abstract terminal elimination data of [definition 10.138](#).*

Proof. Items (i) and (ii) of [definition 10.140](#) give the local alternative decomposition and the compact vanishing-dissipation, multibubble, scale-collapse, compact-cost-divergence, and scale-rigidity exclusions. The cost-divergence criterion used in the divergent retained-core subcase is supplied by the cost-divergence theorem [theorem 10.42](#) and its status corollary [corollary 10.43](#). By [lemma 10.141](#) they also give critical tightness, and by [lemma 10.142](#) they give the local windowed H^1 control. These are the data listed in [definition 10.138](#). $\square$

Theorem 10.144 (Expanded terminal Type II exclusion). *If the expanded terminal data hold, then no renormalized local Type II branch in the stated class remains.*

Proof. [Lemma 10.143](#) gives the abstract terminal elimination data. Applying [theorem 10.139](#) gives the exclusion. $\square$

Remark 10.145 (Branchwise reading). Compact single-core branches are split by [lemma 11.10](#): vanishing-dissipation cores are removed by the compact criterion, and divergent retained-local-cost cores are removed by [theorem A.21](#) and [corollary 10.43](#). Retained scale-collapse branches are removed by the local compact scale-collapse routing theorem [theorem 7.37](#), or by [theorem 10.33](#) when the scale-collapse drift cost is compatible with the cost-divergence theorem; multibubble and non-tight cases are removed by the terminal profile and critical-tightness inputs; failures of local windowed H^1 control are removed by the local Caccioppoli estimate; and every divergent compatible localized cost on the canonical windowwise schedule is handled by [theorem 10.42](#).

Theorem 10.146 (Expanded terminal data hold branchwise). *Let ω be a physical local Type II branch after [corollary 11.26](#). Then either ω exits through a named local or ambient alternative assigned to its specified closure row in the acyclic local state-space order, or the nine items of [definition 10.140](#) hold on a retained terminal tail.*

Proof. The proof is an ordered verification of the nine terminal inputs.

Local decomposition and structural exclusions. Covered entry supplies a positive local Type II concentration sequence or an earlier named exit. Such an exit is not silently discarded; it is assigned to the closure row specified by [corollary 8.139](#) and [lemma 10.45](#). In the retained case [theorem 11.18](#) gives the local alternative decomposition. The compact single-core alternative is split by [lemma 11.10](#): the vanishing-dissipation subcase is closed by [theorem 11.5](#), and the divergent retained-local-cost subcase is closed, after the carrier and adjacent rows have been verified, by [theorem A.21](#) and [corollary 10.43](#). The multibubble/cascade, finite-cost scale-collapse, retained compact scale-collapse, and scale-rigid alternatives are closed by [theorems 7.37, 11.6 and 11.8](#) and [proposition 5.42](#).

Critical tightness and compact terminal completeness. On the retained compact stratum, critical tightness and local windowed H^1 -control are theorems by [theorem 10.108](#). If either test fails before the retained stratum is reached, the failure is an exterior, rough-core, or multibubble/cascade state-space exit, routed by [theorems 9.39, 11.6 and 11.7](#). The finite critical-mass input used here is local on compact terminal windows: [lemma 10.78](#) proves that failure of positive finite local L^3 mass is either non-retention or an adjacent active label. The proof does not require the whole-space critical norm $\|V(\tau)\|_{L^3(\mathbb{R}^3)}$. With this local finite-mass input, the terminal compactness partition is supplied by [theorem 10.132](#), and terminal completeness on the retained branch is supplied by [theorem 10.126](#). In the retained compact-cost path, the compactness actually used by the cost-divergence argument is supplied from the local compact tests by [lemma A.32](#). The adjacent labels in the partition are routed by [theorem 10.81](#). That theorem only removes adjacent labels from the retained terminal component; the later exclusion of carrier and cost adjacent rows is supplied in the assembly layer by the ambient adjacent wrapper [corollary 8.139](#).

Residual stability, representation, exterior removal, and local regularity. Terminal state-space completeness is [theorem 9.16](#), proved locally from the discharged critical profile decomposition [theorem 10.87](#). Local perturbative residual stability is [corollary 10.92](#). Physically exterior regular components are removed by [theorem 9.39](#). The repaired-gauge representation is the ordered consequence of [corollaries A.12](#) and [A.19](#) together with the absolutely continuous chart consequences [theorem 10.56](#); if the chart or gauge row fails, that failure is the representation alternative and not a retained terminal branch. The local Caccioppoli/windowed regularity input is supplied by [theorem 10.61](#) and [corollary 10.62](#) on suitable repaired-gauge branches, with retained windowed control supplied by [theorem 10.108](#). Finally, compact scale-collapse routing is exactly [theorem 7.37](#).

Thus each item in [definition 10.140](#) is either verified on the retained terminal tail or its first failure is one of the named local or ambient alternatives assigned to a closure row before terminal retained-state elimination is invoked. The dependency order is the one recorded in [lemma 10.45](#): carrier and adjacent rows are routed before cost-divergence closure, and terminal retained-state elimination is last. The verification uses only local compact windows, ordered state-space labels, suitable local energy, and compact terminal estimates; it does not import a global profile theorem, global Kato smallness, or a whole-space L^3 bound. $\square$

Corollary 10.147 (Expanded terminal wrapper discharged). *The expanded terminal data are verified on the retained terminal path: every physical local Type II branch either satisfies those data on a retained terminal tail or exits through a named local or ambient alternative assigned to the closure row specified by the acyclic local state-space order. This corollary verifies the retained terminal inputs; it does not by itself exclude every named exit row.*

Proof. This is [theorem 10.146](#). $\square$

11. ASSEMBLY OF THE LOCAL TYPE II EXCLUSION

11.1. Local Entry. Let (u, p) be a suitable weak solution of the three-dimensional Navier–Stokes equations

$$\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u, \quad \nabla \cdot u = 0,$$

on $\mathbb{R}^3 \times (T - \delta, T)$. For $z_0 = (x_0, T)$ and $r > 0$, set

$$Q_r(z_0) = B_r(x_0) \times (T - r^2, T),$$

$$C(z_0, r) = r^{-2} \int_{Q_r(z_0)} |u|^3 dx dt,$$

and

$$D(z_0, r) = r^{-2} \int_{T-r^2}^T \int_{B_r(x_0)} |p(x, t) - (p)_{B_r(x_0)}(t)|^{3/2} dx dt.$$

The pressure is always taken modulo functions of time, and the spatial mean normalization is the one used in the Caffarelli–Kohn–Nirenberg criterion.

Definition 11.1 (Singular set and local Type II point). Let $\Sigma(T)$ be the set of points x_0 for which u is not locally bounded in any backward cylinder $Q_r(x_0, T)$. A point $x_0 \in \Sigma(T)$ is in the local Type II alternative if it carries a positive local Caffarelli–Kohn–Nirenberg concentration sequence and no local Type I scale-invariant bound holds at (x_0, T) .

Definition 11.2 (Positive local Type II concentration sequence). A sequence $r_n \downarrow 0$ is a positive local Type II concentration sequence at (x_0, T) if there exists $\eta > 0$ such that

$$C((x_0, T), r_n) + D((x_0, T), r_n) \geq \eta \quad \text{for all } n,$$

and (x_0, T) is in the local Type II alternative. The associated parabolic rescalings are

$$u_n(y, s) = r_n u(x_0 + r_n y, T + r_n^2 s), \quad p_n(y, s) = r_n^2 p(x_0 + r_n y, T + r_n^2 s).$$

Theorem 11.3 (Local concentration and Type I/Type II dichotomy). *If $\Sigma(T) \neq \emptyset$, then for every $x_0 \in \Sigma(T)$ there are $\eta > 0$ and $r_n \downarrow 0$ such that*

$$C((x_0, T), r_n) + D((x_0, T), r_n) \geq \eta \quad \text{for all } n.$$

At the same point exactly one of the following alternatives holds: there is a local Type I scale-invariant bound, or (x_0, T) is in the local Type II alternative.

Proof. This is the pointwise CKN concentration dichotomy of [theorem 3.6](#), restated in the notation used in the final assembly. $\square$

11.2. Local Theorems Used in the Assembly. The local results needed for the final assembly are restated in the following form.

Theorem 11.4 (Repaired-gauge representation). *Let a positive local Type II concentration sequence contain a nondegenerate selected core. Then, after passing to selected windows, there are final scale-center variables $\Lambda(\tau) > 0$, $X(\tau)$, renormalized fields*

$$V(Y, \tau) = \Lambda(\tau)u(X(\tau) + \Lambda(\tau)Y, t(\tau)), \quad P(Y, \tau) = \Lambda(\tau)^2 p(X(\tau) + \Lambda(\tau)Y, t(\tau)) + \pi(\tau),$$

and modulation coefficients a, b satisfying the single canonical represented equation (14). On each compact renormalized window one has the local suitable energy inequality, local pressure reconstruction modulo functions of time, stability under subsequences, and the compact-window bounds proved in [Section 4](#). Failure of the nondegenerate scale-center selection is part of the multibubble, cascade, or scale-rigidity alternatives.

Proof. [Theorem 4.37](#) constructs the compact-window chart and its suitable structure. Its absolutely continuous final variables enter [theorem 10.52](#), and [theorem 4.25](#) supplies the canonical local pressure decomposition. These are exactly the conclusions stated here. $\square$

Theorem 11.5 (Compact single-core exclusion). *No compact single-core vanishing-dissipation branch can be a local Type II sequence. More explicitly, let a represented selected-window sequence have:*

- (i) *retained positive local velocity concentration on a fixed compact cylinder;*
- (ii) *uniformly bounded compact-cylinder suitable estimates $A + E + C + D$;*
- (iii) *nondegenerate local scale-center representation;*
- (iv) *bounded modulation coefficients;*
- (v) *local pressure reconstruction modulo spatial means;*
- (vi) *vanishing localized dissipation on some retained compact cylinder.*

Then the sequence is incompatible with the local Type II condition.

Proof. [Theorem 3.18](#) gives the stated exclusion. $\square$

Theorem 11.6 (Multibubble and cascade exclusion). *In the local terminal class where terminal admissibility is supplied by [proposition 5.43](#) and scale-rigidity is discharged by [proposition 5.42](#), no local multibubble, cascade, or gauge-degenerate Type II branch can occur in the class produced by the local active-core extraction. Equivalently, every failure of the single nondegenerate compact-core alternative is excluded by the local multibubble/cascade analysis.*

Proof. [Theorem 5.44](#), together with its removal corollary [corollary 5.45](#), gives the stated exclusion. $\square$

Theorem 11.7 (Rough-core reduction). *On any compact renormalized window satisfying the local energy inequality and bounded modulation hypotheses, finite outer velocity-pressure CKN control $\mathcal{C} + \mathcal{D}$ implies finite inner windowed H^1 control. Hence failure of local windowed H^1 control on a retained inner cylinder forces failure of the compact-cylinder $\mathcal{C} + \mathcal{D}$ bounds on every relevant enclosing cylinder.*

Proof. [Corollary 6.7](#) gives the stated reduction. $\square$

Theorem 11.8 (Finite-cost scale-collapse exclusion). *No renormalized branch with finite scale-collapsing cost, and hence no branch with finite absolute scale cost, can have genuine scale collapse together with a nonvanishing selected core.*

Proof. [Theorem 7.7](#) and [corollary 7.9](#) give the stated exclusion. $\square$

Theorem 11.9 (Scale-collapse decomposition of alternatives). *Let a represented Type II branch have genuine scale collapse and a nonvanishing selected core. After passing to selected terminal windows, exactly one of the alternatives isolated in [Section 7](#) occurs:*

- (i) *finite scale-collapsing cost;*
- (ii) *finite absolute scale cost;*
- (iii) *thin-drift alternative, routed by [theorem 10.118](#);*
- (iv) *autonomous-modulation alternative;*
- (v) *compactness-package first failure, routed by [theorem 7.34](#);*
- (vi) *a nonzero autonomous reduced limit in the scale-rigid state.*

In the local Type II class of alternatives used by the multibubble/cascade and scale-collapse analyses, alternatives (iii)–(vi) do not constitute additional compact single-core mechanisms: thin drift is routed to another local row or a closed retained row, the modulation alternative lies outside the retained compact scale-collapse branch, the compactness-package failure is routed to

a more specific named local row, and the remaining retained scale-collapse state is removed by [theorem 7.37](#).

Proof. The assertion follows by combining the scale-collapse decomposition of alternatives with the local compact scale-collapse routing theorem [theorem 7.37](#). $\square$

Lemma 11.10 (Compact retained core dissipation dichotomy). *Let a represented retained compact core carry positive local Type II concentration on a fixed compact cylinder and satisfy the retained compact-cylinder bounds. Then, after passing to selected terminal windows, one of the following ordered alternatives holds:*

- (i) *there is a retained compact cylinder on which the selected windows have vanishing localized dissipation, hence the branch is a compact single-core vanishing-dissipation branch in the sense of [definition 3.4](#);*
- (ii) *an exterior, rough-core, modulation, scale, carrier, or cost adjacent label occurs before the branch is a retained compact single-core state;*
- (iii) *on a retained core cylinder carrying the positive local velocity concentration, the localized dissipation has divergent tail integral, hence the canonical localized cost diverges for every cutoff which is identically one on that core.*

Proof. Fix a retained cylinder Q_ρ on which the positive velocity core is recorded. If

$$\int_{\tau_0}^{\infty} \int_{B_\rho} |\nabla V|^2 dy d\tau < \infty,$$

then choose retained selected windows I_{n_j} , with disjoint enlarged windows after passing to a subsequence, such that

$$\int_{I_{n_j}} \int_{B_\rho} |\nabla V|^2 dy d\tau \rightarrow 0.$$

The retained compact-cylinder bounds, pressure reconstruction, admissible modulation control on the selected windows, and positive velocity core are exactly the remaining clauses of [definition 3.4](#); if any retained compact-cylinder, pressure, exterior, rough-core, modulation, scale, carrier, or cost input needed for that definition fails, the ordered first-failure rule assigns item (ii). In particular, failure of the integrated modulation control is the modulation-window exit of [lemma 10.106](#). If no such exit occurs, the selected subsequence gives item (i).

If the displayed dissipation integral is infinite, then for any cutoff $\phi_{R_0} \equiv 1$ on B_ρ the canonical localized cost satisfies

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau \geq \nu \int_{\tau_0}^{\infty} \int_{B_\rho} |\nabla V|^2 dy d\tau = \infty,$$

unless one of the ordered local tests needed to define the retained compact core has failed first. Such first failures are precisely the exterior, rough-core, modulation, scale, or carrier/cost adjacent labels of the local state-space validation, giving item (ii). If no first failure occurs, the displayed divergence is item (iii). $\square$

11.3. Assembly of Alternatives.

Lemma 11.11 (Bounded selected-window limits are locally regular modulo named remaining activity). *Let (V_n, P_n) be represented suitable weak solutions on Q_{2R} , with the local pressure reconstruction and compact-cylinder suitable bounds supplied by the repaired-gauge state-space construction. Assume*

$$V_n \rightarrow V_* \quad \text{strongly in } L^3(Q_{2R}), \quad V_* \in L^\infty(Q_{2R}).$$

Then for every compact $K \Subset Q_R$, exactly one of the following occurs after passing to a subsequence:

- (i) V_n is uniformly bounded on a finite subcover of K , hence the corresponding physical branch is regular on the physical image of K ;
- (ii) a positive CKN amount survives outside that regular subcover and is captured by one of the named ordered rows: a smaller active core, an exterior core, a pressure-only carrier, a multibubble/cascade branch, a gauge-degenerate branch, or a rough-core/local-energy failure row.

Proof. Fix $z_0 = (y_0, s_0) \in K$. Since $V_* \in L^\infty(Q_{2R})$, for sufficiently small $\rho = \rho(z_0) > 0$ with $Q_{2\rho}(z_0) \subseteq Q_{2R}$,

$$\rho^{-2} \int_{Q_{2\rho}(z_0)} |V_*|^3$$

is below a fixed fraction of the represented-gauge CKN threshold. Strong L^3 -convergence gives the same smallness for V_n on $Q_{2\rho}(z_0)$ for all large n .

By [lemma 5.25](#), the bounded selected-window limit has uniformly vanishing velocity–pressure CKN quantity at sufficiently small scales on compact subcylinders. The strong convergence of V_n , the strong convergence of the local Calderón–Zygmund pressure pieces, and harmonic pressure decay therefore transfer this smallness to (V_n, P_n) for all large n , unless the pressure-atlas or harmonic-pressure bounds fail. Such a failure is one of the named pressure rows in alternative (ii). In the non-failure case, the represented-gauge CKN regularity lemma gives uniform boundedness of V_n on a smaller cylinder around z_0 .

A finite subcover of K gives alternative (i) unless, during the finite covering process, a positive amount of CKN mass is found outside the regularized subcover. That remaining positive CKN mass is, by the ordered state-space construction, a smaller active core, exterior core, pressure-only carrier, multibubble/cascade, gauge-degenerate, or local-energy/rough-core failure. These are exactly the named alternatives in (ii). $\square$

Definition 11.12 (Model carrier versus original-window carrier). A model carrier is a positive CKN lower bound on a fixed cylinder after rebasing a selected descendant to the model cylinder. An original-window carrier is a positive CKN lower bound on a fixed subcylinder of the original bounded selected window, independent of the generation of the reselection tree.

The diffuse reselection-tree case means that model carriers exist at every generation, but no subsequence yields an original-window carrier.

Lemma 11.13 (Reselection-tree support points are realized by descendant states). *Consider an infinite bounded-window reselection tree whose selected descendant cylinders all satisfy the fixed lower CKN threshold. Normalize the activity mass on each generation and record for every selected descendant its compactified relative scale-center coordinates inside the original bounded window, together with the rebased local occupation state on a fixed model cylinder and the usual failure faces for representation, pressure atlas, gauge chart, compact-upper, local-energy, and rough-core breakdown. After passing to a subsequence in the generation index, these normalized tree measures converge weakly-* on a compact metrizable reselection state space. The limiting tree measure itself is retained as part of the compactification, so later diffuse arguments may apply [lemma 9.63](#) directly to it. The tree occupation coordinates are continuous coordinate projections in this product compactification, and their positive superlevel sets are closed in the reselection state space.*

Every support point of the limit state is realized by an actual sequence of descendant cylinders and corresponding rebased local states from the original reselection tree. If such a realized support profile carries positive compact local activity on some fixed model cylinder, then it gives a model carrier. It gives an original-window carrier only if the relative scale-center coordinates of the realizing descendants remain in a compact chart of the original bounded window.

Proof. The compactified relative scale-center chart is compact by construction, and the rebased local-state factor is compact by [lemma 9.62](#) once one records failure faces for loss of represented structure, compact-upper control, local energy, pressure atlas, gauge chart, and rough-core breakdown. Hence the normalized tree measures are tight and possess weak-* limits.

Let Θ be one such limit and let $\mathfrak{s} \in \text{supp } \Theta$. Choose nested neighborhoods (U_m) of $\mathfrak{s}$ with diameters tending to zero. Since $\Theta(U_m) > 0$, for every m and all sufficiently large generations there is an actual selected descendant cylinder whose rebased state lies in U_m . Selecting one such descendant for each m produces a sequence of genuine rebased local states converging to $\mathfrak{s}$, proving realization.

If the realized support profile has positive compact local activity on a fixed model cylinder Q_ρ , then some neighborhood of $\mathfrak{s}$ is defined by a quantitative lower bound for the CKN quantity on Q_ρ . The support property then forces actual descendants from the original tree to satisfy the same lower bound for all large m , which is exactly a model carrier. If the relative scale-center coordinates of those descendants remain in a compact chart of the original bounded window, undoing the rebasing sends them to physical subcylinders of the original bounded window and produces an original-window carrier. $\square$

Lemma 11.14 (Diffuse reselection-tree row is closed). *Let an infinite bounded-window reselection tree be produced after a bounded selected-window region has been regularized. Assume that the tree carries the fixed CKN lower threshold at each selected generation, but that no subsequence of descendants produces an original-window carrier in the original reselection branch.*

Then the diffuse reselection-tree alternative is not terminal. More precisely, after passing to the compactified reselection-tree state space of [Lemma 11.13](#), one of the following named rows occurs:

- (i) *a representation, pressure atlas, gauge chart, compact-upper, local-energy, or rough-core failure face;*
- (ii) *a same-point cascade, scale-collapse, or noncompact exterior row;*
- (iii) *an original-window carrier, contradicting the diffuse/no-original-window-carrier case;*
- (iv) *a visible positive residual defect in the augmented local occupation-state coordinates of [Theorem 10.86](#), hence a named adjacent local row or a retained comparable class already grouped before the bounded-window reselection was declared.*

Consequently a diffuse bounded-window reselection tree cannot remain as a terminal physical Type II branch.

Proof. Apply [Lemma 11.13](#). The normalized tree measures have a weak-* limit on a compact metrizable reselection state space, and every support point is realized by actual descendant cylinders from the original reselection tree.

If a support point lies on a failure face for representation, pressure atlas, gauge chart, compact upper estimates, local energy, or rough-core control, then the corresponding named row occurs. This is alternative (i).

Apply [Lemma 9.63](#) to the normalized reselection-tree measures. If the atomic/activity alternative occurs, then the support-realization lemma gives a model carrier. If the relative scale-center coordinates of the corresponding descendants stay in a compact chart, that model carrier yields an original-window carrier, contradicting the diffuse case. If those relative coordinates do not stay in a compact chart, then the ordered relative scale-center partition gives the same-point cascade, scale-collapse, or noncompact exterior row. This is alternatives (ii) and (iii).

If the genuinely diffuse alternative occurs, then the normalized tree measure itself is a visible diffuse occupation coordinate. Applying [Theorem 10.86](#) gives a named adjacent row or a retained comparable class. The retained comparable class is impossible because comparable retained

classes are grouped before bounded-window reselection is declared. Hence only named adjacent rows remain.

Thus the diffuse reselection-tree alternative is not terminal. $\square$

Lemma 11.15 (Bounded-window reselection is well-founded). *Let a physical local Type II germ produce a bounded selected-window limit on a selected compact cylinder. Remove the regular portion of that bounded window using the local CKN criterion, and rerun the ordered local selection on any remaining positive CKN concentration in the physical image of a smaller terminal neighborhood.*

Then exactly one of the following alternatives occurs:

- (i) *there is a physical neighborhood $Q_{\rho_0}(x_0, T)$ and a scale $r_0 > 0$ such that every admissible subcylinder $Q_r(z) \subset Q_{\rho_0}(x_0, T)$, $0 < r < r_0$, satisfies*

$$C(z, r) + D(z, r) < \varepsilon_{\text{CKN}};$$

hence the germ is locally regular there;

- (ii) *after finitely many reselections, the branch enters one of the named state-space rows: active core, pressure-only row, separated exterior row, multibubble row, same-point cascade row, gauge-degenerate row, rough-core row, or retained terminal compact branch;*
- (iii) *the reselection process does not terminate finitely. Then the selected centers and scales generate one of the infinite rows: same-point scale cascade, scale-collapse row, noncompact exterior row, or diffuse reselection-tree row.*

Proof. Each reselection produces a selected cylinder $Q_{r_j}(z_j)$ with CKN quantity bounded from below by the fixed activity threshold. If no such cylinder exists below some physical scale, the finite-cover form of the CKN criterion gives item (i).

Assume therefore that selected cylinders continue to exist. Compare each new selected cylinder with the previous selected core through the local complexity triple

$$\Xi_j := (\mathfrak{c}_j, \mathfrak{s}_j, \eta_*),$$

where $\mathfrak{c}_j$ records whether the new center is same-point, physically separated at comparable scale, or escapes every compact relative chart, and $\mathfrak{s}_j$ records whether the new scale is strictly smaller, comparable, strictly larger in the renormalized variables, or has no compact relative limit. This is a finite ordered partition of the relative scale-center chart.

If infinitely many reselections fall in the comparable separated-center class, then a finite-cover selection on the corresponding comparable cylinders gives a separated exterior or multibubble row. If infinitely many reselections stay at the same physical point with strictly decreasing physical scales, then the selected cylinders form an infinite nested chain and we are in the same-point scale-cascade row. If infinitely many reselections stay at the same physical point while the renormalized radius strictly increases, then the branch is the same-point larger-scale cascade row. If the relative centers or relative scales leave every compact subset of the relative scale-center chart, then by definition the branch is in the scale-collapse or noncompact exterior row.

Finally, if the CKN mass persists but no subsequence produces an original-window carrier, normalize the remaining CKN mass on the reselection tree at the first scale on which each descendant cylinder is selected. Apply [lemma 11.13](#) and then [lemma 9.63](#). If the atomic/activity alternative occurs, then one obtains a model carrier. If its relative scale-center coordinates stay in a compact chart, this produces an original-window carrier, contradiction. If not, the branch is in the same-point cascade, scale-collapse, or noncompact exterior row. If the genuinely diffuse alternative occurs, then by definition the branch is in the diffuse reselection-tree row. Thus infinite reselection is routed to one of the named infinite rows. $\square$

Theorem 11.16 (Bounded selected-window branch is physically routed). *Let (u, p, x_0, T) be a physical local Type II singular germ, and suppose that one selected-window extraction has a nonzero bounded selected-window limit in the sense of [definition 3.1](#). Then the bounded-window branch is not a terminal unrecorded physical Type II branch. More precisely, after applying [lemma 11.11](#), exactly one of the following occurs:*

- (i) *the selected window is regular/removable and the ordered local concentration algorithm continues on any surviving smaller, exterior, pressure, multicore, cascade, gauge-degenerate, or rough-core row;*
- (ii) *no surviving CKN concentration remains in the physical neighborhood under consideration, in which case the local Type I/regular alternative holds there, contradicting the assumption that the germ is a physical local Type II singular germ;*
- (iii) *the surviving concentration is one of the named rows in the local state-space ledger and is closed by the corresponding row theorem.*

Consequently bounded selected-window limits are closed physical routing rows, not terminal residual Type II branches.

Proof. Apply [lemma 11.11](#) on an exhausting finite family of compact subcylinders inside the selected window. If the bounded-window region regularizes and no CKN concentration remains, the physical image has a local boundedness/Type I neighborhood, contrary to the physical Type II assumption. If CKN concentration remains, the ordered state-space tests are not restarted ad infinitum without classification: apply [lemma 11.15](#) to any remaining positive CKN concentration. Alternative (i) gives physical local regularity. Alternative (ii) enters one of the finite named rows. Alternative (iii) enters one of the infinite named rows. All rows are closed by the terminal ledger. Hence the bounded selected-window branch cannot remain as a physical Type II germ. $\square$

Theorem 11.17 (Ordered local state-space exhaustion). *Let $r_n \downarrow 0$ be a positive local Type II concentration sequence in the sense of [definition 11.2](#). Starting from the associated rescalings, run the following ordered local tests:*

*C1 admissibility $\rightarrow$ representation/gauge $\rightarrow$ single-core versus multibubble,
pressure-only versus velocity carrier $\rightarrow$ compact upper estimates $\rightarrow$ rough-core/ H^1 control,
scale-collapse/scale-rigid $\rightarrow$ canonical cost validation $\rightarrow$ terminal retained branch.*

After passing to a diagonal subsequence, exactly one first successful retained state or one first failed local row occurs. Every first failed row is one of the named local rows later closed in the ledger, and every retained state is one of the following: a represented compact single-core branch, a multibubble/cascade/gauge-degenerate branch, loss of local windowed H^1 control on a retained compact cylinder, a finite-cost scale-collapse branch with nonvanishing selected core, or a retained compact scale-collapse or scale-rigid terminal state.

Proof. The C1 admissibility test is the pointwise local Type I/Type II dichotomy of [theorem 11.3](#); on a physical local Type II germ its class-side failures are recorded by [lemma 11.24](#) and [corollary 11.26](#). Given admissibility, the ordered repaired-gauge representation tests of [theorem 11.4](#) and [corollaries A.12](#) and [A.19](#) either produce a represented local branch on selected windows or record the first representation/gauge failure.

On a represented branch, the active-core extraction is tested first. Failure of uniqueness, same-point splitting, separated cores, or unrepaired scale degeneracy is the multibubble/cascade side of [theorem 11.6](#). If a single retained core is selected, then the ordered pressure atlas and carrier tests decide whether one is on a retained velocity core, a pressure-only row, or an exterior/separated carrier row; this is exactly the local routing package of [lemma 11.32](#) and [corollary 9.73](#). The compact-window upper estimates are then supplied by [proposition 7.14](#), while the retained

spacetime mass floor is supplied by [theorem 7.22](#) whenever the loss-of-active-core row has not already been recorded.

Before applying any repaired gauge chart, the strong gauge-chart entry test [lemma 4.9](#) is run. Failure of this test is the repaired-gauge degeneracy row. If a nonzero bounded selected-window limit appears, then [lemma 5.25](#) gives the local velocity–pressure CKN smallness of the bounded limit, and [lemma 11.15](#) shows that infinite reselection is not left open: it is routed to the same-point cascade, scale-collapse, noncompact exterior, or diffuse reselection-tree rows.

If a nonzero bounded selected-window limit appears at this stage, the branch is not retained in the strict selected-window Type II row by [theorem 5.27](#); in the physical assembly that event is closed by [theorem 11.16](#), and in the local branchwise assembly it is returned to the ordered active-window decomposition rather than declared terminal. If no bounded selected-window limit appears, the rough-core test [theorem 11.7](#) decides whether local H^1 control is retained or the branch returns to the multibubble or compact-cylinder-failure side.

When critical tightness fails, apply [lemma 10.103](#). This produces one of the named rows: same-point larger-scale cascade, separated exterior concentration, noncompact exterior, pressure/gauge/reconstruction failure, rough-core/local-energy/compact-upper failure, or diffuse critical-tail. When pressure-only exterior selection nests indefinitely, apply [lemma 9.67](#); this produces one of the named rows: same-point pressure cascade, separated exterior concentration, scale-collapse, or diffuse pressure-only. The later closure of the diffuse critical-tail and diffuse pressure-only rows is deferred to the remaining-exit closure layer and is not used in the routing step here.

The remaining represented retained branch is tested by the scale package of [theorem 11.9](#): either one reaches finite-cost scale-collapse, thin drift, autonomous modulation, compactness-package first failure, or the retained compact scale-collapse/scale-rigid states. On a pass branch where the selected compactness package holds on a nonzero retained window, [lemma 7.30](#) gives the exhaustive regular–singular split. The regular case is the bounded-window row. The singular case is an actual active descendant, and [corollary 10.79](#) routes it to a named parameter-tail or analytic row, or identifies it with the already assembled comparable retained class. Hence this compact-limit step adds no sixth retained state.

Finally, the canonical cost-validation layer is available on every retained compact branch by [lemma 11.30](#), and its first-failure map is the ordered list in [lemma 8.73](#). Because the list of tests is finite and each stage either records its first failure or passes to the next named layer after subsequence extraction, the diagonal subsequence construction terminates in exactly one first failed row or one retained terminal state. The retained terminal states are exactly the five retained alternatives listed in the statement. $\square$

Theorem 11.18 (Local Type II decomposition into local alternatives). *Every positive local Type II concentration sequence in [definition 11.2](#), after passing to subsequences and selected windows, enters one of the following alternatives:*

- (i) *a represented compact-core branch carrying a single retained active core, which by [lemma 11.10](#) either has vanishing-dissipation selected windows, exits through a named adjacent local row before retention, or enters the retained local cost-divergence branch;*
- (ii) *a multibubble, cascade, or gauge-degenerate branch;*
- (iii) *loss of local windowed H^1 control on a retained compact cylinder;*
- (iv) *a finite-cost scale-collapse branch with nonvanishing selected core;*
- (v) *a retained compact scale-collapse or scale-rigid terminal state.*

Proof. This is the list of retained terminal outcomes produced by the ordered exhaustion theorem [theorem 11.17](#). The single-core represented branch is refined by [lemma 11.10](#); all failures of

unique retained-core selection belong to the multibubble/cascade side of [theorem 11.6](#); rough-core loss is returned by [theorem 11.7](#) to the compact-cylinder failure side; and genuine scale collapse is divided by [theorem 11.9](#) into finite-cost scale collapse or one of the retained compact scale-collapse/scale-rigid terminal states. No other retained terminal state survives the ordered first-failure algorithm; in particular, [corollary 10.79](#) rules out an extra state arising only from the regularity type of a nonzero compact limit. $\square$

Proposition 11.19 (Elimination of local Type II states). *No positive local Type II concentration sequence can remain in any retained state among alternatives (i)–(v) of [theorem 11.18](#). If the ordered validation inside alternative (i) exits through a named adjacent local row before retention, that branch is not a retained local state and is closed in the ambient adjacent-routing layer.*

Proof. Alternative (i) is excluded by the compact single-core criterion, [theorem 11.5](#), in the vanishing-dissipation subcase of [lemma 11.10](#). The adjacent-label subcase of that lemma is not a retained compact single-core state: the ordered validation has stopped earlier at a named local row. In the divergent localized-dissipation subcase, the canonical localized cost diverges on the retained core and the cost-divergence closure [theorem 10.42](#) and [corollary 10.43](#) excludes the branch after the windowwise adjacent alternatives have been routed.

Alternative (ii) is excluded by the local multibubble and cascade theorem, [theorem 11.6](#).

Alternative (iii) is not an independent terminal state. By [theorem 11.7](#), rough-core loss of local H^1 control forces failure of the compact-cylinder $\mathcal{C} + \mathcal{D}$ bounds on enclosing compact cylinders. Such failure yields an active compact-cylinder concentration branch and therefore belongs to the multibubble/cascade alternative, already excluded by [theorem 11.6](#).

Alternative (iv) is excluded by the finite-cost scale-collapse theorem, [theorem 11.8](#).

Alternative (v) is the retained compact scale-collapse or scale-rigid terminal state. Retained compact scale-collapse states are excluded by [theorem 7.37](#). Scale-rigid terminal states are excluded by [proposition 5.42](#): the state either has a nonzero bounded selected-window limit, in which case the strict selected-window branch stops and the bounded-window event is routed by [theorem 11.16](#) in the physical assembly, or it reduces to the compact single-core branch, which is excluded by [theorem 11.5](#). Hence it cannot remain as a retained local Type II branch. $\square$

11.4. Physical Entry into the Covered Local Type II Path.

Definition 11.20 (Physical local Type II singular germ). Let (u, p) be a suitable weak solution on $\mathbb{R}^3 \times (T - \delta, T)$. A tuple (u, p, x_0, T) is a physical local Type II singular germ if:

- (i) $x_0 \in \Sigma(T)$;
- (ii) the local Type I scale-invariant bound of [theorem 3.6](#) fails at (x_0, T) ;
- (iii) the germ is an interior Navier–Stokes germ, meaning that all sufficiently small backward cylinders $Q_r(x_0, T)$ are contained in the Navier–Stokes domain under consideration.

Definition 11.21 (Covered local Type II entry). A physical local Type II singular germ has covered local Type II entry if one of the following two alternatives occurs:

- (i) one of its Caffarelli–Kohn–Nirenberg concentration sequences, after passing to a subsequence and applying the selected-window reductions used in this paper, enters one of the alternatives in [theorem 11.18](#);
- (ii) before the selected-window state decomposition is reached, the ordered local tests assign the germ to a class-incompatible alternative or to a named non-retained local state-space exit.

Lemma 11.22 (Physical Type II germs are admissible local branches). *Let (u, p, x_0, T) be a physical local Type II singular germ. After translating space and time and restricting to a smaller*

terminal interval if necessary, the terminal germ defines an element of $\mathcal{U}_{\text{II}}^{NS}$ in the sense of [definition A.1](#).

Proof. Choose $0 < \delta' < \delta$ so that the cylinders under consideration remain interior. Set

$$\tilde{u}(y, s) = u(y + x_0, T - \delta' + s), \quad \tilde{p}(y, s) = p(y + x_0, T - \delta' + s), \quad 0 < s < \delta',$$

and $T^* = \delta'$. This space-time translation and restriction preserve the local Navier–Stokes equations, suitability, and the pressure normalization modulo functions of time. The terminal time T^* is finite by construction. Since $x_0 \in \Sigma(T)$, the terminal germ is not a local continuation branch. By the definition of physical local Type II germ, the local Type I scale-invariant bound fails at (x_0, T) , so the germ is not in the local Type I class. The interior-domain clause excludes boundary and non-Navier–Stokes-domain alternatives. These are exactly the admissibility clauses in [definition A.1](#). $\square$

Lemma 11.23 (Physical Type II germs carry local Type II concentration sequences). *Let (u, p, x_0, T) be a physical local Type II singular germ. Then there are $\eta > 0$ and $r_n \downarrow 0$ such that*

$$C((x_0, T), r_n) + D((x_0, T), r_n) \geq \eta \quad \text{for all } n,$$

and r_n is a positive local Type II concentration sequence in the sense of [definition 11.2](#).

Proof. Since $x_0 \in \Sigma(T)$, [theorem 3.6](#) gives $\eta > 0$ and $r_n \downarrow 0$ such that the displayed CKN lower bound holds. The same dichotomy has exactly two pointwise alternatives: the local Type I scale-invariant bound holds, or the point lies in the local Type II alternative. The first alternative is excluded by [definition 11.20](#). Hence (x_0, T) lies in the local Type II alternative, and the displayed concentration sequence is exactly a positive local Type II concentration sequence as defined in [definition 11.2](#). $\square$

Lemma 11.24 (No extra outside-class row for physical Type II germs). *Let (u, p, x_0, T) be a physical local Type II singular germ and let ω be the admissible terminal branch supplied by [lemma 11.22](#). Any failure of the C1 analytic data for ω is either class-incompatible or one of the named non-retained local exits in [theorem 10.122](#). In particular, there is no additional outside-covered-class alternative for such a physical Type II germ.*

Proof. By [lemma 11.22](#), $\omega \in \mathcal{U}_{\text{II}}^{NS}$. The ordered C1 discharge [theorem 10.122](#) applies on this class. Its conclusion is precisely that the first failure of [definition A.3](#) is assigned either to a class-incompatible alternative excluded by [definition A.1](#), or to a named non-retained state-space branch already covered by the local stratification. Thus a physical Type II germ has no residual “outside the covered class” row after the ordered local tests are run. $\square$

Theorem 11.25 (Covered-class entry for physical Type II germs). *Every physical local Type II singular germ has covered local Type II entry in the sense of [definition 11.21](#). More explicitly, after passing to a CKN concentration sequence and to the selected windows used in this paper, the germ enters one of the alternatives in [theorem 11.18](#) through the ordered exhaustion route of [theorem 11.17](#), unless it has already exited through one of the class-incompatible or named non-retained local alternatives of [theorem 10.122](#).*

Proof. Let (u, p, x_0, T) be a physical local Type II singular germ. By [lemma 11.22](#), the germ defines an admissible terminal branch. Run the ordered C1 analytic-data tests on that branch. If a first failure occurs, then [lemma 11.24](#) routes it to a class-incompatible alternative or to a named non-retained local state-space exit. This is covered entry in the sense of [definition 11.21](#) (ii).

It remains to consider the retained case, when no C1 failure occurs. By [lemma 11.23](#), the germ carries a positive local Type II concentration sequence. The ordered exhaustion theorem [theorem 11.17](#) is formulated exactly for such sequences, and its retained outcomes are listed in

[theorem 11.18](#). After passing to subsequences and selected windows, the sequence therefore enters one of the five alternatives listed there. This is covered entry in the sense of [definition 11.21](#) (i). Thus every case is a covered local state-space outcome. The proof uses only pointwise CKN concentration, local first-failure routing, and selected-window state-space decomposition; it does not use a global $L^3(\mathbb{R}^3)$ bound, a global profile decomposition, global Kato smallness, or vanishing localized dissipation beyond what is produced inside the selected local branch. $\square$

Corollary 11.26 (Discharge of covered-class entry). *The covered-class entry obligation for the local Type II state-space analysis is discharged: every physical local Type II singular germ either enters the local state-decomposition path of [theorem 11.18](#), or exits earlier through a class-incompatible or named non-retained local alternative.*

Proof. This is [theorem 11.25](#). $\square$

Remark 11.27 (Scope of the covered-entry discharge). [Corollary 11.26](#) proves only the entry statement: a physical local Type II germ is not lost before the local state-space analysis begins. It does not by itself supply the ambient adjacent-stratum closure data or the expanded terminal data used in the final assembly. Nor does it assert that every named non-retained exit has already been eliminated; it asserts that any such exit is recorded in the local state-space taxonomy rather than being an untracked outside-class branch. The closure of those named exits is supplied by the local closure theorems and by the ambient/terminal wrappers [corollaries 8.139](#) and [10.147](#). The entry argument is also independent of [theorem 11.39](#); it uses the pointwise CKN dichotomy, the local Type I/Type II split, and the ordered C1 routing theorem, so the discharge is not circular.

11.5. Closure of the remaining named exit rows. The three closure theorems in this subsection convert the rows previously listed in [definition 11.35](#) into closed adjacent rows, i.e. rows that are routed by an existing local closure or wrapper into a retained-state alternative that has already been excluded. The arguments are local on compact terminal cylinders. They use only the ordered local state-space tests, the modulation-window dichotomy, the local pressure reconstruction with harmonic remainder, and the existing wrappers [corollaries 8.139](#) and [10.147](#). No global L^3 bound, no global profile decomposition, no global Kato smallness, no hidden vanishing localized dissipation, and no noncanonical schedule are introduced.

Lemma 11.28 (Bounded modulation defects are local carrier labels). *Let (a_n, b_n) be the repaired-gauge modulation coefficients on a fixed compact terminal window $[0, T]$ for a represented branch with a nonvanishing selected core V_n , and suppose*

$$\sup_n \int_0^T (|a_n(\tau)| + \|b_n(\tau)\|) d\tau < \infty.$$

After passing to a subsequence, the ordered local test gives one of the following outcomes:

(i) *writing*

$$\bar{a}_n = T^{-1} \int_0^T a_n(\tau) d\tau, \quad \bar{b}_n = T^{-1} \int_0^T b_n(\tau) d\tau,$$

there are constants $a_\infty \in \mathbb{R}$ and $b_\infty \in \mathbb{R}^3$ such that, after passing to the same subsequence, $(\bar{a}_n, \bar{b}_n) \rightarrow (a_\infty, b_\infty)$ and either

$$a_n \rightarrow a_\infty \quad \text{in } L^1(0, T), \quad b_n \rightarrow b_\infty \quad \text{in } L^1(0, T; \mathbb{R}^3);$$

or the nonconstant part of the modulation is invisible in the terminal local residual topology:

$$(a_n - \bar{a}_n)(V_n + y \cdot \nabla V_n) + (b_n - \bar{b}_n) \cdot \nabla V_n \rightarrow 0$$

in $L^1(0, T; H^{-1}(K))$ for every compact K ;

- (ii) a signed component of the normalized modulation defect is visible in the local residual topology and defines a local occupation state. The ordered state-space test records one of the adjacent labels: modulation-driven recentering, positive scale-shell work, negative scale-drift, loss of the selected active core, or repaired-gauge degeneracy.

The conclusion is local on the fixed terminal window and uses no compactness stronger than finite-measure weak-* compactness of the modulation defect measures.

Proof. Write $m_n = T^{-1} \int_0^T (a_n, b_n) d\tau \in \mathbb{R}^4$. The L^1 bound implies that, after passing to a subsequence, $m_n \rightarrow m_\infty = (a_\infty, b_\infty)$. Set

$$\alpha_n(\tau) = a_n(\tau) - \bar{a}_n, \quad \beta_n(\tau) = b_n(\tau) - \bar{b}_n, \quad m_n = (\bar{a}_n, \bar{b}_n).$$

If

$$\|\alpha_n\|_{L^1(0,T)} + \|\beta_n\|_{L^1(0,T;\mathbb{R}^3)} \rightarrow 0,$$

then item (i) holds.

Assume instead that the displayed defect norm has a positive liminf. Define the local modulation-defect residual

$$\mathcal{F}_n^{\text{mod}} := \alpha_n(V_n + y \cdot \nabla V_n) + \beta_n \cdot \nabla V_n.$$

If $\mathcal{F}_n^{\text{mod}} \rightarrow 0$ in $L^1(0,T;H^{-1}(K))$ for every compact K , then the represented local equation has the same terminal residual limit as the sequence of constant-coefficient local equations with coefficients m_n . Since $m_n \rightarrow m_\infty$ in the finite-dimensional coefficient space, this is the constant-coefficient autonomous channel after the retained local compactness test is run. Item (i) holds in its perturbative form.

It remains to consider the case in which the defect residual is not perturbative on some compact K . Passing to a further subsequence and applying the ordered component convention, one of

$$\int_0^T (\alpha_n)_+ d\tau, \quad \int_0^T (\alpha_n)_- d\tau, \quad \int_0^T |\beta_n| d\tau$$

has a subsequence whose normalized signed residual contribution is nonzero in the local terminal residual topology. Normalize that visible component by its mass. The resulting probability measures are bounded in $\mathcal{M}([0,T])$, and hence have a weak-* subsequential limit. Together with the nonzero residual pairing, this is a visible local modulation-defect occupation state on the selected window.

The ordered local state-space test evaluates visible signed defect residuals before declaring a retained autonomous window. A visible β_n -component is the modulation-driven recentering/moving-translation label of [definition 8.110](#). A visible $(\alpha_n)_+$ -component is the positive scale-shell label, and a visible $(\alpha_n)_-$ -component is the negative scale-drift label. If the selected core on which the modulation acts has lost its active mass floor, the first failed row is the retained active-core row. If the defect cannot be represented in the repaired chart, the first failed row is repaired-gauge degeneracy. These alternatives are ordered before autonomous retention, and they exhaust the nonperturbative signed components of the modulation defect. Hence item (ii) holds. $\square$

Theorem 11.29 (Autonomous-modulation failure is a closed adjacent row). *Let (u, p, x_0, T) be a physical local Type II singular germ in the sense of [definition 11.20](#), and assume the germ has admissible covered entry into the local state-space path. If a represented selected scale-collapse branch with nonvanishing selected core enters the autonomous-modulation alternative [theorem 7.24\(v\)](#) on a compact terminal cylinder, then the branch enters one of the following closed adjacent rows: the modulation-window-unbounded row routed by [lemma 10.106](#); the modulation-driven recentering or moving-carrier row routed by [definition 8.110](#) and discharged in the ambient adjacent wrapper [corollary 8.139](#); the negative scale-drift row closed by [theorem 10.118](#); the*

positive scale-shell row routed by [corollary 8.131](#); the loss-of-selected-active-core row closed by [theorem 11.33](#); or the gauge-degenerate / repaired-gauge row already excluded from the retained branch by the ordered representation tests. Equivalently, the autonomous-modulation alternative is not an independent remaining terminal exit.

Proof. Fix a represented selected scale-collapse branch on a compact terminal cylinder $B_R \times [\tau_n, \tau_n + T]$ in the autonomous regime, and let $a_n(\tau) = a(\tau_n + \tau)$ and $b_n(\tau) = b(\tau_n + \tau)$ on $[0, T]$ be the absolutely continuous repaired-gauge modulation coefficients identified in [lemma 10.51](#). Cover the selected window by finitely many unit windows of the type used in [lemma 10.106](#), with the same repaired-gauge chart variables as that lemma; the construction of the chart variables and their absolute continuity are supplied for any represented branch by [lemma 10.51](#) and the ordered representation tests [corollaries A.12](#) and [A.19](#). Exactly one of the following two cases occurs.

Case A (window-integrals unbounded). On at least one unit window, the supremum $\sup_n \int (|a_n| + \|b_n\|_\infty) d\tau$ is infinite. By the same stratification of accumulated chart velocities used in the proof of [lemma 10.106](#), namely the splitting into large accumulated negative scale motion, large positive scale-shell work, large translation work, and unrepresentable accumulated modulation, the branch is recorded as one of the named modulation, scale-drift, modulation-driven recentering, exterior, or gauge-degenerate local alternatives. The same stratification applies on the scale-collapse stratum because it depends only on the absolutely continuous chart variables of [lemma 10.51](#) and on the carrier-routing definition [definition 8.110](#), neither of which uses the single-core hypothesis. All five rows are closed adjacent rows: the negative scale-drift row by [theorem 10.118](#); the positive scale-shell component by [corollary 8.131](#); modulation-driven recentering by [definition 8.110](#) together with the ambient adjacent wrapper [corollary 8.139](#); the exterior row by [corollary 9.73](#); and the gauge-degenerate / representation row by the ordered representation tests [corollaries A.12](#) and [A.19](#). This is the conclusion claimed.

Case B (window-integrals bounded). On every unit window in the cover, the opposite alternative of [lemma 10.106](#) holds, so $\sup_n \int_0^T (|a_n| + \|b_n\|_\infty) d\tau < \infty$. By the local modulation-defect routing lemma [lemma 11.28](#), after passing to a subsequence either the modulation coefficients have a constant L^1 limit on the window, the nonconstant defect is perturbative and the represented local equation has the same terminal residual limit as the sequence of constant-coefficient branches with the same window-average coefficients, or a nonzero visible local modulation-defect occupation state is recorded. In the first two alternatives the finite-dimensional window averages converge, and the retained local compactness test gives the corresponding constant-coefficient autonomous local limit, which is excluded by the autonomous-modulation alternative [theorem 7.24\(v\)](#). In the remaining nonperturbative alternative the defect is recorded, before autonomous retention, as modulation-driven recentering, positive scale-shell work, negative scale-drift, loss of selected active core, or repaired-gauge degeneracy. These rows are closed respectively by [corollary 8.139](#), [corollary 8.131](#), [theorem 10.118](#), [theorem 11.33](#), and the ordered representation tests [corollaries A.12](#) and [A.19](#). Thus the bounded-modulation case also leaves the autonomous-modulation alternative through a closed adjacent row.

The two cases exhaust the possibilities. No part of the argument used a whole-space critical norm $\|V(\tau)\|_{L^3(\mathbb{R}^3)}$, a global profile decomposition, a global Kato smallness condition, a noncanonical schedule, or a hidden vanishing localized dissipation: only finite-window L^1 boundedness, weak-* compactness of normalized defect measures, and the existing closure theorems for the named adjacent rows are used. $\square$

Lemma 11.30 (Canonical identity-cost channel is always available on retained compact branches). *Let ω be a retained compact branch that has passed the representation, pressure-atlas, compactness,*

and retained-core tests leading to the local cost-validation layer. Then the canonical identity-cost validation channel coming from the suitable local energy inequality is available on that branch. If this canonical validation fails, the first failed canonical test is one of the already named local rows in [lemma 8.73](#).

Proof. On every retained compact branch, the represented equation, repaired-gauge chart, and local suitable energy inequality are available by [theorem 11.4](#) and [lemma 10.64](#). Hence the canonical identity-cost channel is defined using the identity cutoff and the localized dissipation/carrier package attached to that suitable-energy inequality. The ordered validation tests are exactly those of [lemma 8.73](#). If one fails before retained canonical divergence is validated, that failure is by definition the first failed canonical test and is one of the named local rows listed there. Thus the canonical channel is always available once the retained compact branch reaches the cost-validation layer, and any failure is already part of the existing first-failure ledger. $\square$

Theorem 11.31 (Noncanonical carrier and cost validation route to the canonical channel). *Let ω be a renormalized Type II solution with active supercritical scaling on a terminal tail. Let a candidate continuation lie in a compact cost-divergence regime through a noncanonical carrier or cost-validation channel different from the canonical retained identity-cost channel. Then exactly one of the following occurs after running the ordered local validation tests of [lemma 8.73](#):*

- (a) *the candidate is, after the comparable-cost compatibility supplied by [lemma 8.67](#), the canonical identity-cost retained channel up to the terminal one-sided comparison of [definition A.23](#), in which case the cost-divergence exclusion on the retained stratum is supplied by [corollary 8.82](#), hence closed;*
- (b) *the first failed test lies in the cost-side list [lemma 8.73\(b\)](#). These are the cost-compatibility, carrier-subidentity, finite fixed-cutoff error, summable moving-annulus error, carrier-comparison for the validated cost, modulation-integrability, or compact cost-divergence failures; each such first failure is a closed adjacent row, routed in the ambient adjacent wrapper [corollary 8.139](#).*

Consequently the row “carrier or cost-validation failure outside the canonical retained identity-cost channel” is not an independent remaining terminal exit.

Proof. By [lemma 8.73](#), after running the ordered local validation tests (i)–(ix) of that lemma, exactly one of (a) all tests pass with ω on the retained cost-divergence regime, or (b) the first failed test is one of the named local cost-side first failures listed there.

In case (a), the comparable-cost compatibility lemma [lemma 8.67](#) identifies the validated noncanonical compact Navier–Stokes cost on the retained stratum with the canonical identity-cost up to the terminal one-sided comparison of [definition A.23](#). By [lemma 11.30](#), the canonical identity-cost validation channel is available on the retained compact branch. Therefore the retained cost-divergence exclusion supplied by [corollary 8.82](#) applies, and the candidate is excluded.

In case (b), the first failed test is by construction one of the named cost-side adjacent rows enumerated in [lemma 8.73\(b\)](#). We show each such row is closed.

Class-side first failures. Ordered exhaustion failure, representation failure, and repaired-gauge or scale-degeneracy failure are the C1/C2 first failures already routed before the state-decomposition path is reached. On a physical Type II germ they are not left as analytic gaps: the covered-entry theorem records them either as class-incompatible alternatives or as named local state-space rows. Thus they are discharged by [corollary 11.26](#) together with the ordered representation tests [corollaries A.12](#) and [A.19](#).

State-decomposition first failures. Multibubble/multicore splitting is closed by [theorem 11.6](#) and [theorem 11.5](#); cascade is closed by the finite-cost scale-collapse theorem [theorem 11.8](#) together

with [corollary 5.45](#); exterior-tightness failure is closed by [corollary 9.73](#); loss of local windowed H^1 control is closed by [theorem 11.7](#), which returns to the multibubble/cascade closed row; finite-cost scale-collapse is closed by [theorem 11.8](#); the scale-rigid residual is closed by [proposition 5.42](#); critical-mass adjacent labels are closed by [theorem 11.33](#); and modulation-integrability failure is closed by [theorem 11.29](#) (and by [lemma 10.106](#) in the ambient single-core stratum).

Cost-side first failures. These are the irreducible noncanonical-cost rows. The remaining items in [lemma 8.73\(b\)](#) are the cost-side labels: cost-compatibility failure, carrier-subidentity failure, finite fixed-cutoff error failure, summable moving-annulus error failure, carrier-comparison failure for the validated cost, and failure of compact cost divergence for the validated cost. These rows are not closed by declaring the Navier–Stokes solution impossible; they are closed by the ordered canonical-channel rule: the final retained proof uses only the canonical identity cost, or an explicitly comparable compact cost, and every failure of a noncanonical diagnostic is routed back to the canonical local validation layer.

Cost-compatibility and carrier-comparison failure. By definition of [lemmas 8.66](#) and [8.67](#), cost-compatibility failure means that the validated cost used to record divergence is neither the canonical identity cost nor a comparable cost in the sense of [definition A.23](#). Such a diagnostic is not used to retain a Type II branch. By [lemma 11.30](#), the canonical identity-cost channel from the suitable local energy inequality is available on every retained compact branch. Hence the ordered local validation is rerun on that canonical channel. If the canonical channel passes, the retained stratum is closed by [corollary 8.82](#). If it fails, the first failed canonical test is one of the ordered local rows already treated above. Thus an incomparable noncanonical cost is only a bookkeeping failure of that diagnostic and does not create a new terminal exit.

Carrier-subidentity, finite fixed-cutoff error, and summable moving-annulus error failure. By [lemma 8.89](#), suitable local energy on the repaired-gauge branch supplies the moving carrier subidentities and the corresponding fixed and moving carrier closure packages of [definitions 8.90](#) and [8.91](#), with finite-error selection [definition 8.98](#) and the carrier completion of [corollaries 8.107](#) and [8.108](#). If one of these carrier rows fails before the retained carrier is validated, the failure is not repaired by a global estimate. The ordered componentwise carrier decomposition records the first nonsummable local term. Transition, pressure, convection, and cutoff terms route to exterior/leakage, critical tightness failure, or local H^1 -loss; translation terms route to modulation-driven recentering; positive scale terms route to the positive scale-shell/scale-rigid row; and negative scale terms route to the thin-drift row. These are precisely the alternatives in [definition 8.110](#), closed by [corollary 8.139](#), [corollary 8.131](#), and [theorem 10.118](#). If the suitable-energy carrier identity itself is unavailable, the first failed row is the already routed representation/repaired-gauge or suitable-energy row, not a new cost-validation exit.

Failure of compact cost divergence for the validated cost. This is the negation of the entry hypothesis: if the validated cost does not diverge on the compact terminal cylinder, the candidate is not in the compact cost-divergence regime under consideration. The final assembly then continues through the ordinary retained-state decomposition: finite cost is handled by [theorem 11.8](#), while failure before finite-cost classification is one of the ordered local rows already listed in [lemma 8.73](#). Thus nondivergence of a validated noncanonical cost cannot survive as an independent terminal exit.

The argument uses only the ordered finite-step local validation tests, the comparable-cost compatibility on compact windows, the suitable local energy package on the retained stratum, and the canonical-channel fallback for diagnostics that are not explicitly comparable. It introduces no whole-space L^3 bound, no global profile decomposition, and no noncanonical schedule outside the comparison framework already established by [lemmas 8.67](#) and [8.80](#). $\square$

Lemma 11.32 (Pressure-only critical mass routes through local pressure atlas). *Let $B_R \Subset B_{R_1} \Subset B_{R_2}$ be nested balls in a retained terminal frame on a compact time interval I , and fix $1 \leq q < \infty$. Suppose a pressure oscillation modulo functions of time persists in $L_\tau^q L_y^{3/2}(B_R)$, while the ordered terminal selection has not retained a velocity active core for that candidate. With the local pressure decomposition*

$$P_n = P_n^{\text{loc}} + H_{n,R_1} + c_{n,R_1}(\tau) \quad \text{on } B_{R_1},$$

where P_n^{loc} is generated by $\zeta(V_n \otimes V_n)$ on B_{R_2} and H_{n,R_1} is mean-zero harmonic on B_{R_1} , one of the following ordered alternatives occurs:

- (i) *the localized velocity source $\zeta(V_n \otimes V_n)$ has a nonzero compact-window $L_\tau^q L_y^{3/2}$ limsup; the local CKN test on the selected compact subcylinders then either absorbs the source as regular, or records a comparable retained class, a multibubble/multicore label, a cascade label, a rough-core label, or an exterior/separated label;*
- (ii) *the localized velocity source tends to zero in $L_\tau^q L_y^{3/2}(B_{R_2})$ for the fixed exponent q , and any remaining pressure oscillation is carried by the harmonic part H_{n,R_1} , which is routed by [lemma 9.22](#) to the ordered pressure/exterior row;*
- (iii) *the pressure reconstruction or gauge needed for the displayed decomposition fails, which is itself the ordered pressure/repaired-gauge row.*

Thus a pressure-only critical-mass signal is never closed by assuming vanishing velocity mass; it is routed by the local source/harmonic dichotomy.

Proof. The pressure decomposition is one of the ordered local pressure-atlas tests. If that test fails, item (iii) is recorded. Assume it holds.

If

$$\limsup_{n \rightarrow \infty} \|\zeta(V_n \otimes V_n)\|_{L_\tau^q L_y^{3/2}(B_{R_2})} > 0,$$

then the compact-window velocity stress is nonzero. Since $\|V_n \otimes V_n\|_{L^{3/2}(B)} = \|V_n\|_{L^3(B)}^2$ on each spatial ball, a finite cover of B_{R_2} and a subsequence select a compact subcylinder with positive local velocity critical mass in the corresponding spacetime norm. Run the local CKN test on that subcylinder. If all smaller admissible cylinders are below the CKN threshold, the source is regular on the compact window and is absorbed into the regular pressure atlas. Otherwise a local CKN carrier is selected, and the ordered terminal critical-mass routing [lemma 10.78](#) records either a comparable retained class or one of the adjacent active labels listed in item (i). This uses only the compact-window stress norm, local CKN testing, and finite covering.

If the localized source tends to zero in $L_\tau^q L_y^{3/2}(B_{R_2})$, local Calderón–Zygmund estimates give

$$P_n^{\text{loc}} \rightarrow 0 \quad \text{in } L_\tau^q L_y^{3/2}(B_{R_1})$$

modulo functions of time. Hence any persistent pressure oscillation on B_R must lie in the harmonic remainder. Applying [lemma 9.22](#) with $p = 3/2$ and the fixed time exponent q routes that nonvanishing harmonic remainder to the pressure/exterior row. No global pressure compactness or lower semicontinuity argument is used. $\square$

Theorem 11.33 (Loss of retained active core is a closed adjacent row). *Let (u, p, x_0, T) be a physical local Type II singular germ that has covered entry into the local state-space path, and let a terminal active frame be selected from its local state-space decomposition. If on a fixed compact terminal cylinder the candidate fails to retain an active core in the sense of [lemma 10.78](#), then the candidate enters one of the following closed adjacent rows:*

- (i) *the retained spacetime activity-floor failure, in which case the candidate is not a retained terminal frame by the terminal selection rule and the state-space partition records the*

next available retained frame; if no velocity retained frame remains, the original CKN signal is pressure/exterior, rough-core, or regular, and is routed by the pressure-only and exterior closures;

- (ii) the local L^3 -upper-bound failure, which by [lemma 10.78\(iii\)](#) is one of: a comparable finite-mass class absorbed into the retained frame by [lemma 10.77](#), a multibubble or multicore splitting closed by [theorem 11.6](#) and [theorem 11.5](#), a cascade closed by the finite-cost scale-collapse theorem [theorem 11.8](#) together with the cascade absorption used in [corollary 5.45](#), a rough-core label closed by [theorem 11.7](#), or an exterior/separated label closed by [corollary 9.73](#);
- (iii) the pressure-only critical-mass row, in which case the local pressure atlas routes the signal by [lemma 11.32](#): the local Calderón–Zygmund source is either regular on the compact window, produces a velocity-supported terminal label, or the harmonic remainder is routed into the pressure/exterior-leakage alternative, closed by [corollary 9.73](#).

Consequently the row “loss of a retained active core, including pressure-only critical-mass detection” is not an independent remaining terminal exit.

Proof. Apply [lemma 10.78](#) on the fixed compact terminal cylinder $B_R \times I$. After passing to a subsequence, exactly one of its three alternatives (i)–(iii) holds.

If alternative (i) of that lemma holds, the candidate has positive finite local critical L^3 mass on the cylinder; this is the retained case and not the loss-of-core case under consideration. If alternative (ii) holds, the lower active-mass floor of [theorem 9.33](#) fails, so the candidate frame is not retained by the terminal selection rule. The selection rule then moves to the next available velocity frame. If there is no such frame, the positive CKN concentration of the physical germ cannot be a retained velocity core; it is either pressure-only, exterior/separated, rough-core, or below the CKN threshold on all smaller cylinders. The first three are routed by [lemma 11.32](#), [corollary 9.73](#), and [theorem 11.7](#); the last contradicts singular CKN concentration by [theorem 9.35](#). This is item (i) of the present conclusion. If alternative (iii) holds, the local L^3 upper bound fails on the compact cylinder. By [lemma 10.78\(iii\)](#), the ordered local state space then records a comparable finite-mass class absorbed into the retained component, a multibubble or multicore label, a cascade label, a rough-core label, or an exterior/separated label. Each such label is closed by the theorems cited in item (ii) of the present conclusion: the rough-core label is closed by [theorem 11.7](#), which forces failure of compact-cylinder CKN control on every relevant enclosing cylinder and therefore returns to the multibubble/cascade closed row; the exterior/separated label is closed by [corollary 9.73](#).

It remains to handle the pressure-only subcase, in which the active core fails because the critical mass detected on the cylinder is carried by the pressure rather than by a velocity profile, i.e. the local velocity L^3 -mass remains below the active floor while a pressure oscillation modulo a function of time persists on B_R along a subsequence. Apply the local pressure-only routing lemma [lemma 11.32](#) with the time exponent in which the pressure oscillation persists. If the localized Calderón–Zygmund source has a nonzero compact-window stress limsup, the finite-cover and CKN test in that lemma either absorbs the source as regular, or records a velocity-supported terminal label; the latter is either absorbed into a comparable retained class by [lemma 10.77](#) or routed to one of the closed multibubble/cascade, rough-core, or exterior/separated rows. If the localized source tends to zero, the local pressure part vanishes by Calderón–Zygmund estimates and any remaining pressure oscillation is harmonic; [lemma 9.22](#) routes it to the ordered pressure/exterior-leakage alternative, closed by [corollary 9.73](#). If pressure reconstruction fails, that is the ordered pressure/repaired-gauge row. This is item (iii) of the present conclusion.

Each step uses only compact terminal windows, the local active-mass floor, the ordered local state-space partition, and the local pressure reconstruction with harmonic remainder. No global

L^3 bound, no global profile decomposition, no global Kato smallness, and no pressure compactness by weak lower semicontinuity is used. $\square$

Lemma 11.34 (Acyclic terminal ledger). *The routing layer used before the remaining-exit closure subsection is acyclic. More precisely, [lemmas 7.30, 9.67, 10.103 and 11.15](#), [corollary 10.79](#), and [theorems 10.81, 10.86 and 11.17](#) produce only named local rows or the retained comparable class. They do not use the later closure theorems [lemmas 9.68, 10.104 and 11.14](#) to create those rows.*

Consequently the terminal routing package first records a named row, and only later, in the remaining-exit closure layer, invokes the theorem that closes that row.

Proof. By construction, [lemmas 9.67, 10.103 and 11.15](#) stop at the diffuse pressure-only, diffuse critical-tail, and diffuse reselection-tree rows by name. They do not appeal to the later theorems that eliminate those rows. The regular–singular routing lemma [lemma 7.30](#) produces only a named limit interface, an active descendant, or a bounded-window row, and [corollary 10.79](#) routes the active descendant to a named parameter-tail/analytic row or to the already retained comparable class. It records the bounded-window row by name and defers its physical closure to the final ledger. Theorem [theorem 10.81](#) is a routing theorem only, and [theorem 10.86](#) routes a visible residual defect to a named adjacent row or to the retained comparable class. The proof of [theorem 11.17](#) uses these routing outputs only as named ledger rows and explicitly defers their closure to the remaining-exit closure layer. This is exactly the asserted acyclic ordering. $\square$

Definition 11.35 (Canonical final named-exit ledger). The present retained-state assembly does not count a first-failure row as closed merely because it has been named. Rank 1 denotes an unresolved named exit at the terminal-routing stage, while rank 0 denotes a row already discharged to a regular/removable or previously closed structural state. The following table is the single final ledger; each displayed closure edge strictly decreases rank.

Named exit	Rank	Closure source	Target rank
Bounded selected-window limit	1	lemmas 5.25 and 11.15 and theorem 11.16	0
Autonomous-modulation failure or modulation-driven recentering	1	theorem 11.29	0
Noncanonical carrier or cost-validation failure	1	theorem 11.31	0
Loss of a retained active core, including pressure-only, critical-tail, and nested exterior-pressure behavior	1	theorem 11.33 and lemmas 9.67, 9.68, 10.103 and 10.104	0

The noncanonical-cost row is tested before the identity-cost channel of [lemma 11.30](#) is invoked. All local rows not displayed in the table have already been discharged by [proposition 11.19](#) and [corollary 8.139](#). By [lemma 11.34](#), the four rank-1 rows are recorded in the routing layer only by name; their listed closure sources are invoked only in [Section 11.5](#). This ledger controls [theorem 11.37](#), and its unconditional upgrade is [theorem 11.38](#).

Corollary 11.36 (All four remaining named exit rows are closed). *Every rank-1 row in the canonical ledger [definition 11.35](#) is carried by its displayed closure source to a rank-0 row. Hence all four remaining named exits are closed adjacent rows.*

Proof. Apply the closure source in each row of [definition 11.35](#). Each target has rank 0, and [lemma 11.34](#) prevents any of these closure edges from returning to a rank-1 routing state. $\square$

Theorem 11.37 (Retained Type II closure with named-exit ledger). *Let (u, p) be a suitable weak solution on $\mathbb{R}^3 \times (T - \delta, T)$, and let (u, p, x_0, T) be a physical local Type II singular germ in the sense of [definition 11.20](#). Then either the germ enters one of the remaining named exit rows in [definition 11.35](#), or its retained local branch is impossible. Consequently, together with [corollary 11.36](#), no physical local Type II singular germ remains.*

Proof. Let (u, p, x_0, T) be a physical local Type II singular germ. By [corollary 11.26](#), the germ either enters the local state-decomposition path or exits earlier through a class-incompatible or named non-retained local alternative.

A class-incompatible exit contradicts the definition of physical local Type II germ and the admissibility theorem [lemma 11.22](#). If a named non-retained exit occurs, then either it is one of the closed local rows cited in [proposition 11.19](#) and [corollary 8.139](#), or it belongs to the remaining named-exit ledger [definition 11.35](#). In the latter case the first alternative in the statement holds.

It remains to consider the case in which the branch reaches the retained terminal path without entering any remaining named exit. The expanded terminal wrapper [corollary 10.147](#) supplies the expanded terminal data of [definition 10.140](#), and [theorem 10.144](#) excludes the branch.

By [lemma 11.34](#), all uses of critical-tightness, pressure-only, and bounded-window reselection inside the terminal routing layer occur only at the row-naming stage. Their closure is invoked only now through [corollary 11.36](#).

Equivalently, starting from a physical local Type II germ, the ordered state-space exhaustion produces one first retained row or one first failed row. The strong gauge-chart entry test prevents the gauge implicit-function theorem from being applied outside its C^1 chart. Bounded selected-window limits are locally regular or well-foundedly reselected by [lemmas 5.25](#) and [11.15](#). Failure of critical tightness is routed by [lemma 10.103](#), and the diffuse tail case is closed by [lemma 10.104](#). Exterior pressure-only nesting is well-founded by [lemma 9.67](#), and diffuse pressure-only defects are excluded by [lemma 9.68](#). All remaining rows are the previously closed active, exterior, multibubble, cascade, rough-core, pressure-gauge, cost, and terminal compact rows. Hence no row can contain a physical local Type II germ. $\square$

Theorem 11.38 (No physical local Type II germ). *Let (u, p) be a suitable weak solution on $\mathbb{R}^3 \times (T - \delta, T)$. No physical local Type II singular germ (u, p, x_0, T) in the sense of [definition 11.20](#) exists.*

Proof. By [theorem 11.37](#), every physical local Type II singular germ either enters one of the rows in [definition 11.35](#) or is excluded on the retained branch. By [corollary 11.36](#), every row in [definition 11.35](#) is a closed adjacent row: it is routed by its local closure theorem either to a closed retained-state alternative, to a structural row already excluded, or to a regular/removable row. Hence no physical local Type II singular germ remains. $\square$

11.6. Assembly Theorem.

Theorem 11.39 (Local Type II exclusion). *Let (u, p) be a suitable weak solution on $\mathbb{R}^3 \times (T - \delta, T)$. No point $x_0 \in \Sigma(T)$ lies in the local Type II alternative.*

Proof. [Theorem 11.38](#) excludes every physical local Type II singular germ. If a point $x_0 \in \Sigma(T)$ lay in the local Type II alternative, then it would satisfy the three clauses of [definition 11.20](#): singularity gives $x_0 \in \Sigma(T)$, the local Type I bound fails by definition of the Type II alternative, and the theorem is stated for interior cylinders of $\mathbb{R}^3 \times (T - \delta, T)$. This is a physical local Type II singular germ, contradicting [theorem 11.38](#). $\square$

Corollary 11.40 (Final local dichotomy after Type I separation). *Every singular point at time T belongs to the local Type I alternative. Equivalently, once the Type I case is covered by the Type I part of the theory, there is no remaining local Type II singular branch.*

Proof. Let $x_0 \in \Sigma(T)$. By [theorem 11.3](#), x_0 carries a positive CKN concentration sequence and lies either in the local Type I alternative or in the local Type II alternative. [Theorem 11.39](#) excludes the latter. $\square$

12. CONCLUSION

The paper gives a local branchwise exclusion mechanism for Type II Navier–Stokes concentration. Beginning with a positive Caffarelli–Kohn–Nirenberg concentration sequence, the argument passes to repaired-gauge variables around a retained core and keeps the analysis inside local PDE estimates: suitable energy inequalities, pressure reconstruction, compactness on fixed cylinders, Caccioppoli estimates, profile decoupling, and localized scale-cost identities.

The compact single-core mechanism is the central obstruction. A retained core with positive local velocity density and vanishing localized dissipation has only two compact limits after passing to selected windows. The zero limit contradicts the persistence of the retained velocity core by strong L^3 convergence, while a nonzero spatially constant limit gives a bounded selected-window profile and hence exits the Type II regime. Pressure-only harmonic modes are routed through the pressure/exterior rows and are not used to close the compact single-core velocity row. The rest of the proof shows that the remaining alternatives to this compact mechanism are also accounted for: rough cores are reduced to compact-cylinder control, active multibubbles and cascades are separated by the profile analysis, and finite-cost scale collapse is incompatible with a nonvanishing retained core.

The scale-collapse analysis is intentionally explicit about its routing. Finite scale-collapsing cost, and hence finite absolute scale cost in the covered setting, are excluded within the renormalized setting. Thick autonomous windows produce nontrivial autonomous limits only after the relevant local compactness and modulation rows have passed. The retained compact autonomous branch is not closed by silently invoking a global Liouville theorem: it is routed through [theorem 7.37](#) into the scale-rigid discharge. If a retained input fails, the failure is one of the named modulation, compactness, exterior, rough-core, or multicore local rows; the compactness, exterior, and thin-drift boundary faces are routed by their local closure theorems, including [theorem 7.34](#), [corollary 9.73](#), and [theorem 10.118](#).

Thus every physical local Type II concentration branch is either eliminated on the retained path or routed through one of the named local rows. The remaining named rows are closed in [corollary 11.36](#), and the final assembly [theorem 11.38](#) excludes physical local Type II singular germs. The conclusion remains local and branchwise: it does not import a whole-space critical estimate, a global profile decomposition, global Kato smallness, or hidden vanishing localized dissipation.

APPENDIX A. PDE INTERFACE DEFINITIONS USED IN THE LOCAL TYPE II PROOF

This appendix records, in PDE notation, the local Type II classification interfaces used earlier in the proof. The statements are interface definitions and local consequences for admissible Type II branches; they do not provide a separate post-assembly proof path. In the final assembly, their inputs are verified on the retained local path or routed to closed named rows.

A.1. Admissible Type II branches and exhaustion.

Definition A.1 (Admissible local Type II class). Let $\mathcal{U}_{\Pi}^{NS}$ be the class of admissible terminal branches (u, p, T^*) such that:

- (i) (u, p) is a three-dimensional incompressible Navier–Stokes solution on a terminal interval $[0, T^*)$ in the analytic class under consideration;
- (ii) $T^* < \infty$ is the terminal time of the branch;
- (iii) the branch is not in the local Type I class associated with the scale-invariant Type I criterion;
- (iv) the branch is a Type II branch rather than a continuation, Type I, boundary, or non-Navier–Stokes-domain alternative.

Membership in $\mathcal{U}_{\Pi}^{NS}$ specifies the class under analysis; no nonemptiness assertion is made.

Definition A.2 (Type II analytic data). A branch $\omega = (u, p, T^*) \in \mathcal{U}_{\Pi}^{NS}$ has the Type II analytic data if the following three analytic assertions hold:

- (i) a concentration profile or blow-up germ modulo the Navier–Stokes symmetry group;
- (ii) a supercritical scale alternative, meaning that the branch enters the Type II scale alternative rather than the Type I scale alternative;
- (iii) local profile data in the Navier–Stokes profile class used for the subsequent compactness and representation arguments.

The analytic data are *strictly exhaustive* if every admissible Type II branch has these three properties. They are *exhaustive by ordered local routing* if every branch which fails one of the three properties is assigned, at its first failed C1 test, to one of the ordered failures in [definition A.3](#) and that failure is then routed to a class-incompatible alternative or to a named non-retained local state-space exit. On a retained subbranch, ordered local-routing exhaustion reduces to strict exhaustion.

Definition A.3 (Ordered exhaustion failures). For an admissible Type II branch, the ordered failures of the Type II analytic data are:

- (i) concentration extraction fails;
- (ii) concentration extraction succeeds but the supercritical scale alternative does not hold;
- (iii) concentration extraction and the scale alternative hold but the local profile data are insufficient for the subsequent compactness and representation arguments;
- (iv) the admissible Type II branch lies outside the concentration, scale, and profile alternatives under consideration.

The first applicable failure in this order is the exhaustion alternative.

Theorem A.4 (Type II branch exhaustion). *Assume the Type II analytic data of [definition A.2](#) are exhaustive by ordered local routing on $\mathcal{U}_{\Pi}^{NS}$, and that each ordered failure is assigned to a class-incompatible alternative or to a named non-retained local state-space exit. Then every $\omega \in \mathcal{U}_{\Pi}^{NS}$ either has the concentration profile, the supercritical scale alternative, and the local profile data, or is already in one of those named exits. In particular, every retained admissible Type II branch has the three analytic data.*

Proof. Fix $\omega = (u, p, T^*) \in \mathcal{U}_{\text{II}}^{NS}$. By ordered local routing, exactly one of two possibilities occurs. If no ordered failure occurs, then the concentration-extraction, scale, and profile components all hold, so ω has the three Type II analytic data. If an ordered failure occurs, the first such failure is, by hypothesis, assigned to a class-incompatible alternative or to a named non-retained local state-space exit. Therefore a retained admissible branch cannot be in the failure case, and must have the three analytic data. $\square$

Corollary A.5 (Ordered exhaustion classification). *For each $\omega \in \mathcal{U}_{\text{II}}^{NS}$, exactly one ordered alternative occurs: the Type II analytic data hold, or one of the four failures in [definition A.3](#) occurs. The data are strictly exhaustive on $\mathcal{U}_{\text{II}}^{NS}$ if and only if the first alternative holds for every $\omega \in \mathcal{U}_{\text{II}}^{NS}$. They are exhaustive by ordered local routing if every failure alternative is assigned to a class-incompatible alternative or to a named non-retained local state-space exit; on the retained subbranch this is equivalent to the first alternative.*

Proof. Test the four analytic requirements in the stated order. If concentration extraction, the scale alternative, and the profile condition all hold, the branch has the Type II analytic data. If a test fails, the first failed test gives the corresponding ordered failure. The alternatives are disjoint by the first failure convention and exhaustive by construction. The final equivalence is the definition of strict exhaustion. The ordered-routing statement follows by recording the target assigned to each first failure. Once class-incompatible and named non-retained exits have been removed, a retained branch cannot lie in any failure alternative, so it lies in the data alternative. $\square$

A.2. Repaired-gauge representation from Type II analytic data.

Definition A.6 (Raw Type II chart). A raw Type II chart consists of

$$(u, p, T^*, x_c, \lambda, \tau, V, P, \mathcal{I}), \quad \mathcal{I} = (t_0, T^*),$$

where u, p solve

$$\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u, \quad \nabla \cdot u = 0,$$

$\lambda(t) > 0$, $x_c(t) \in \mathbb{R}^3$, $x_c, \lambda \in AC_{\text{loc}}$, $\lambda(t) \rightarrow 0$ as $t \uparrow T^*$, and

$$\frac{d\tau}{dt} = \lambda(t)^{-2}, \quad \tau(t) \rightarrow \infty \quad \text{as } t \uparrow T^*.$$

The renormalized variables are

$$y = \frac{x - x_c(t)}{\lambda(t)}, \quad u(x, t) = \lambda(t)^{-1}V(y, \tau(t)), \quad p(x, t) = \lambda(t)^{-2}P(y, \tau(t)),$$

with enough regularity to justify the distributional chain rule on compact y -sets.

Lemma A.7 (Raw renormalized equation). *For a raw Type II chart, define*

$$a_{\text{raw}}(\tau(t)) = \lambda(t)\lambda_t(t), \quad b_{\text{raw}}(\tau(t)) = \lambda(t)x'_c(t).$$

Then V, P satisfy, distributionally on compact y -sets,

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a_{\text{raw}}(\tau)(V + y \cdot \nabla V) + b_{\text{raw}}(\tau) \cdot \nabla V, \quad \nabla \cdot V = 0.$$

Proof. Write

$$u(x, t) = \lambda(t)^{-1}V(y, \tau(t)), \quad y = \frac{x - x_c(t)}{\lambda(t)}, \quad \tau_t = \lambda^{-2}.$$

Then

$$\nabla_x u = \lambda^{-2} \nabla_y V, \quad \Delta_x u = \lambda^{-3} \Delta_y V, \quad (u \cdot \nabla_x)u = \lambda^{-3} (V \cdot \nabla_y)V.$$

Since

$$y_t = -\frac{x'_c(t)}{\lambda(t)} - \frac{\lambda_t(t)}{\lambda(t)}y,$$

one has

$$\partial_t u = \lambda^{-3} \partial_\tau V - \lambda_t \lambda^{-2} (V + y \cdot \nabla V) - \lambda^{-2} x'_c(t) \cdot \nabla V.$$

Moreover $\nabla_x p = \lambda^{-3} \nabla_y P$. Substitution in Navier–Stokes and multiplication by λ^3 give

$$\partial_\tau V - \lambda \lambda_t (V + y \cdot \nabla V) - \lambda x'_c(t) \cdot \nabla V + (V \cdot \nabla) V + \nabla P = \nu \Delta V.$$

Moving the modulation terms to the right-hand side gives the equation with $a_{\text{raw}} = \lambda \lambda_t$ and $b_{\text{raw}} = \lambda x'_c$. Finally, $\nabla_x \cdot u = \lambda^{-2} \nabla_y \cdot V$, so incompressibility of u implies $\nabla_y \cdot V = 0$. $\square$

Definition A.8 (Repaired-gauge realization). Let $0 < p < 3$, $\Theta_0 > 0$, and let ψ_R be a smooth compactly supported centering cutoff. A raw Type II chart admits a repaired-gauge realization if, after an admissible time-dependent scale and translation correction, the profile satisfies

$$G_{\text{sc}}(V) = \int_{\mathbb{R}^3} |y|^{-p} |V(y)|^3 dy - \Theta_0 = 0,$$

and

$$G_j(V) = \int_{\mathbb{R}^3} y_j |V(y)|^2 \psi_R(y) dy = 0, \quad j = 1, 2, 3,$$

for a.e. renormalized time. The corrected profile belongs to the admissible gauge class $\mathcal{A}_p$, has the regularity needed to differentiate these functionals, and all time-dependent symmetry corrections have been absorbed into final modulation coefficients a, b .

Definition A.9 (Repaired-gauge representation). A Type II branch has a repaired-gauge representation if it has:

- (i) a raw Type II chart;
- (ii) a repaired-gauge realization;
- (iii) pressure reconstruction modulo functions of time;
- (iv) final modulation coefficients a, b appearing in the repaired-gauge equation.

The representation consists of the tuple

$$(V, P, a, b, G_{\text{sc}}, G_1, G_2, G_3, \tau_0).$$

Theorem A.10 (Representation implication). *Assume a Type II profile branch has a raw Type II chart, a repaired-gauge realization, pressure reconstruction modulo functions of time, and final modulation coefficients. Then the branch has a repaired-gauge representation in the sense of [definition A.9](#). In particular, (V, P, a, b) satisfies the PDE class required by the compact Type II cost-divergence criterion.*

Proof. The raw chart gives V, P and, by [lemma A.7](#), the renormalized Navier–Stokes equation with raw modulation coefficients. The repaired-gauge realization places the profile on the scale and centering gauge surface and absorbs the time-dependent symmetry correction into updated coefficients. The pressure reconstruction identifies the pressure associated with V , modulo the usual functions of time. The modulation data identify the final coefficients a, b . These are the components of [definition A.9](#). $\square$

Definition A.11 (Ordered representation failures). If a repaired-gauge representation is not obtained, the ordered failures are:

- (i) no raw orbit is supplied by the concentration and profile data;
- (ii) the raw orbit exists but the repaired gauge cannot be realized;
- (iii) pressure reconstruction modulo functions of time is unavailable;

(iv) the final modulation coefficients are unavailable.

The first applicable failure is the representation alternative.

Corollary A.12 (Ordered representation classification). *Every Type II profile branch in the admissible class has exactly one ordered alternative: a repaired-gauge representation, or one of the four failures in [definition A.11](#).*

Proof. Verify the four representation inputs in order. If all are present, [theorem A.10](#) gives the repaired-gauge representation. If not, the first missing input gives the corresponding ordered failure. This is exhaustive and disjoint by first-failure ordering. $\square$

A.3. Minimal chart and gauge realization.

Definition A.13 (Minimal representation data). The minimal data for a Type II branch satisfying the analytic data consist of:

(i) an absolutely continuous concentration chart

$$(u, p, T^*, x_c, \lambda, \rho, U, Q, \mathcal{I}),$$

where $\mathcal{I} = (t_0, T^*)$, $\lambda(t) > 0$, $\lambda(t) \rightarrow 0$, $\rho_t = \lambda^{-2}$, and

$$u(x, t) = \lambda(t)^{-1}U(y, \rho(t)), \quad p(x, t) = \lambda(t)^{-2}Q(y, \rho(t)) + c(t), \quad y = \frac{x - x_c(t)}{\lambda(t)};$$

(ii) an absolutely continuous repaired-gauge correction consisting of $\rho = \rho(\tau)$, $\mu(\tau) > 0$, and $q(\tau) \in \mathbb{R}^3$, with

$$\rho_\tau = \mu(\tau)^2,$$

$$V(Y, \tau) = \mu(\tau)U(\mu(\tau)Y + q(\tau), \rho(\tau)), \quad P(Y, \tau) = \mu(\tau)^2Q(\mu(\tau)Y + q(\tau), \rho(\tau)),$$

and

$$G_{\text{sc}}(V(\tau)) = 0, \quad G_j(V(\tau)) = 0, \quad j = 1, 2, 3.$$

The condition $\rho_\tau = \mu^2$ preserves the viscosity coefficient after the time-dependent scale correction.

Lemma A.14 (Chart data imply the raw orbit). *Under the chart data in [definition A.13](#), U, Q satisfy*

$$\partial_\rho U + (U \cdot \nabla)U + \nabla Q = \nu \Delta U + A(\rho)(U + y \cdot \nabla U) + B(\rho) \cdot \nabla U, \quad \nabla \cdot U = 0,$$

where

$$A(\rho(t)) = \lambda(t)\lambda_t(t), \quad B(\rho(t)) = \lambda(t)x'_c(t).$$

Proof. The chart data are the raw orbit data with raw time ρ . The additive pressure function $c(t)$ has zero spatial gradient. The chain-rule computation is the one in [lemma A.7](#), with τ replaced by ρ . Thus spatial derivatives scale as $\nabla_x u = \lambda^{-2}\nabla_y U$, $\Delta_x u = \lambda^{-3}\Delta_y U$, the convection term scales as $\lambda^{-3}(U \cdot \nabla)U$, and $\rho_t = \lambda^{-2}$. The terms generated by λ_t and x'_c give $A = \lambda\lambda_t$ and $B = \lambda x'_c$. Incompressibility pulls back as $\nabla_x \cdot u = \lambda^{-2}\nabla_y \cdot U$. $\square$

Lemma A.15 (Pressure pullback). *Under the chart data,*

$$-\Delta Q = \partial_i \partial_j (U_i U_j)$$

in distributions, modulo functions of ρ . After any admissible repaired-gauge correction, the transformed pressure satisfies

$$-\Delta P = \partial_i \partial_j (V_i V_j)$$

in distributions, modulo functions of τ .

Proof. Taking divergence of the physical Navier–Stokes equation gives

$$-\Delta p = \partial_i \partial_j (u_i u_j),$$

with pressure determined up to a function of time. Pulling this identity back through $p = \lambda^{-2}Q + c(t)$, $u = \lambda^{-1}U$, and $y = (x - x_c)/\lambda$ gives the pressure equation for Q . If

$$V(Y) = \mu U(\mu Y + q), \quad P(Y) = \mu^2 Q(\mu Y + q),$$

then

$$-\Delta_Y P = \mu^4 (-\Delta Q)(\mu Y + q) = \mu^4 \partial_i \partial_j (U_i U_j)(\mu Y + q) = \partial_i \partial_j (V_i V_j).$$

Adding a function of ρ or τ to the pressure does not affect the identity. $\square$

Lemma A.16 (Forced modulation coefficients). *Assume the minimal data of [definition A.13](#). Then (V, P) satisfies*

$$\partial_\tau V + (V \cdot \nabla)V + \nabla P = \nu \Delta V + a(\tau)(V + Y \cdot \nabla V) + b(\tau) \cdot \nabla V,$$

where

$$a(\tau) = \mu(\tau)^2 A(\rho(\tau)) + \frac{\mu_\tau(\tau)}{\mu(\tau)}$$

and

$$b(\tau) = \mu(\tau) B(\rho(\tau)) + \mu(\tau) A(\rho(\tau)) q(\tau) + \frac{q_\tau(\tau)}{\mu(\tau)}.$$

Proof. Let

$$z = \mu(\tau)Y + q(\tau), \quad V(Y, \tau) = \mu(\tau)U(z, \rho(\tau)), \quad \rho_\tau = \mu^2.$$

The raw equation in (z, ρ) is

$$U_\rho + (U \cdot \nabla_z)U + \nabla_z Q = \nu \Delta_z U + A(U + z \cdot \nabla_z U) + B \cdot \nabla_z U.$$

The identities

$$\nabla_Y V = \mu^2 \nabla_z U, \quad \Delta_Y V = \mu^3 \Delta_z U, \quad (V \cdot \nabla_Y)V = \mu^3 (U \cdot \nabla_z)U, \quad \nabla_Y P = \mu^3 \nabla_z Q$$

together with differentiation in τ give

$$V_\tau + (V \cdot \nabla_Y)V + \nabla_Y P - \nu \Delta_Y V = \left(\frac{\mu_\tau}{\mu} + \mu^2 A \right) (V + Y \cdot \nabla_Y V) + \left(\mu B + \mu A q + \frac{q_\tau}{\mu} \right) \cdot \nabla_Y V.$$

This is the displayed equation. Since μ, q, ρ are absolutely continuous and A, B are the raw chart coefficients, a, b are measurable and locally integrable on every finite renormalized window. $\square$

Theorem A.17 (Minimal representation theorem). *If a Type II branch satisfying the analytic data has the minimal data of [definition A.13](#), then it has a repaired-gauge representation in the sense of [definition A.9](#).*

Proof. [Lemma A.14](#) gives the raw orbit. The gauge correction places the profile on the repaired scale and centering surface, hence gives the repaired-gauge realization. [Lemma A.15](#) gives pressure reconstruction modulo functions of time. [Lemma A.16](#) gives the final modulation coefficients. These are the four inputs in [definition A.9](#). $\square$

Definition A.18 (Minimal representation failures). Outside the admissible repaired-gauge representation class, the ordered failures are:

- (i) no absolutely continuous concentration chart is available;
- (ii) the chart exists but the repaired scale and centering gauges cannot be solved admissibly and absolutely continuously.

Pressure-pullback and modulation-coefficient failures are not independent alternatives when the chart and the absolutely continuous gauge solve exist.

Corollary A.19 (Ordered minimal representation classification). *Every Type II branch satisfying the analytic data has exactly one ordered alternative: a repaired-gauge representation, failure of the raw chart, or failure of the repaired-gauge solve.*

Proof. If the chart is absent, the first failure occurs. If the chart exists but the admissible absolutely continuous gauge solve is absent, the second failure occurs. If both are present, [theorem A.17](#) gives the repaired-gauge representation. The alternatives are disjoint by their order and exhaustive for the representation row. $\square$

A.4. Local Retained Mass Cost Divergence.

Lemma A.20 (Spatially constant retained limits are bounded windows). *Let $0 < r < R_*$, let $J \Subset (-\theta, \theta)$, and let*

$$V_j(y, s) = V(y, \sigma_j + s)$$

be selected retained windows on $B_{R_} \times (-\theta, \theta)$. Assume*

$$V_j \rightarrow c(s) \quad \text{strongly in } L^3(B_{R_*} \times (-\theta, \theta)), \quad \nabla_y c = 0,$$

and assume the retained time-derivative estimate

$$\sup_j \|\partial_s V_j\|_{L^1(J; H^{-m}(B_r))} < \infty$$

for some m with $C_c^\infty(B_r) \subset H_0^m(B_r)$. Then c has a representative in $BV(J; \mathbb{R}^3)$, hence $c \in L^\infty(J)$. If $c \not\equiv 0$ on $(-\theta, \theta)$, then after restricting to a smaller selected parabolic subwindow and recentering its time coordinate, the sequence has a nonzero bounded selected-window limit in the sense of [definition 3.1](#). Consequently the branch exits the local Type II class by [theorem 5.27](#).

Proof. Choose $\zeta \in C_c^\infty(B_r)$ with $\int_{B_r} \zeta dy = 1$. For each component k , set

$$C_{j,k}(s) = \int_{B_r} V_{j,k}(y, s) \zeta(y) dy.$$

The strong L^3 convergence implies $C_{j,k} \rightarrow c_k$ in $L^1(J)$. In distributions on J ,

$$\frac{d}{ds} C_{j,k}(s) = \langle \partial_s V_{j,k}(s), \zeta \rangle,$$

and therefore

$$\text{Var}_J(C_{j,k}) \leq \|\zeta\|_{H_0^m(B_r)} \|\partial_s V_j\|_{L^1(J; H^{-m}(B_r))} \leq C_{J,r}.$$

Lower semicontinuity of total variation under L^1 convergence gives $\text{Var}_J(c_k) \leq C_{J,r}$. Since $c_k \in L^1(J)$, this proves $c_k \in BV(J)$. In one dimension $BV(J) \subset L^\infty(J)$, with the bound depending only on $\|c_k\|_{L^1(J)}$, $\text{Var}_J(c_k)$, and $|J|$. Thus $c \in L^\infty(J)$.

If $c \not\equiv 0$, choose a density point $s_0 \in (-\theta, \theta)$ of the set where $|c| > 0$. Since the preceding argument applies to every compact subinterval $J \Subset (-\theta, \theta)$, choose $J_0 \Subset (-\theta, \theta)$ containing s_0 and use the $BV(J_0) \subset L^\infty(J_0)$ representative. Then choose a radius $\rho > 0$ such that

$$B_\rho \times (s_0 - \rho^2, s_0) \Subset B_r \times J_0, \quad \int_{s_0 - \rho^2}^{s_0} |c(s)|^3 ds > 0.$$

Replacing the selected centers by $\sigma_j + s_0$ identifies this subwindow with Q_ρ . This finite time recentering is admissible by the same local selected-window covariance used in [lemma 10.70](#): no lower bound is transferred to a different physical time; the original subwindow is merely rewritten with its own center as the origin. The strong convergence remains valid there and the limit is bounded and nonzero. This is exactly a nonzero bounded selected-window limit, so [theorem 5.27](#) applies. $\square$

Theorem A.21 (Local retained-mass cost divergence). *Let (V, P, a, b) be a repaired-gauge local Type II branch on a terminal tail. Assume there are $R_* > 0$, $\theta > 0$, $\eta_* > 0$, and an unbounded family $\mathcal{S} \subset [\tau_0, \infty)$ of retained selected active-window centers such that every $\sigma \in \mathcal{S}$ carries positive local velocity mass,*

$$\int_{\sigma-\theta}^{\sigma+\theta} \int_{B_{R_*}} |V(y, \tau)|^3 dy d\tau \geq \eta_*.$$

Assume also the retained compact-window bounds, the retained $L^1 H^{-m}$ time-derivative bounds on compact subwindows, the local Aubin–Lions compactness on $B_{R_} \times (-\theta, \theta)$, local windowed H^1 control, and the scale-rigid routing of nonzero bounded selected-window limits. These are all local compact-window hypotheses; no whole-space critical bound, global profile theorem, or a priori vanishing dissipation is assumed. Then finite total localized dissipation on the retained core is impossible. In particular, for the canonical identity cost on a cutoff $\phi_{R_0} \equiv 1$ on B_{R_*} ,*

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty$$

unless the branch has already exited through a bounded selected-window, rough-core, exterior, or adjacent state-space alternative.

Proof. Assume the localized cost is finite on the terminal tail and no listed exit has occurred. Since $\phi_{R_0} \equiv 1$ on B_{R_*} , finite cost gives finite localized dissipation on B_{R_*} . Choose retained selected active-window centers $\sigma_j \in \mathcal{S}$, $\sigma_j \rightarrow \infty$, with disjoint enlarged windows and such that

$$\int_{\sigma_j-\theta}^{\sigma_j+\theta} \int_{B_{R_*}} |\nabla V|^2 dy d\tau \rightarrow 0.$$

This selection is possible because the total dissipation on the tail is finite and $\mathcal{S}$ is unbounded; if no unbounded retained active-window family with the stated mass floor remains, the lower active-mass row has failed and the branch has already exited by the terminal critical-mass routing.

Set $V_j(y, s) = V(y, \sigma_j + s)$. The retained compact-window bounds give, after passing to a subsequence,

$$V_j \rightarrow V_\infty \quad \text{strongly in } L^3(B_{R_*} \times (-\theta, \theta)),$$

by the local Aubin–Lions compactness in the retained state-space topology. The vanishing dissipation gives $\nabla V_\infty = 0$ distributionally on $B_{R_*} \times (-\theta, \theta)$. Hence $V_\infty(y, s) = c(s)$ on the retained core cylinder.

If $c \not\equiv 0$, [lemma A.20](#) turns the spatially constant compact limit into a nonzero bounded selected-window limit on a smaller selected parabolic subwindow. The scale-rigid exit theorem [theorem 5.27](#) therefore removes the branch from the retained local Type II path. If $c \equiv 0$, strong convergence on the core cylinder contradicts the retained local mass floor:

$$\int_{-\theta}^{\theta} \int_{B_{R_*}} |V_j(y, s)|^3 dy ds \rightarrow 0 \quad \text{while} \quad \int_{-\theta}^{\theta} \int_{B_{R_*}} |V_j(y, s)|^3 dy ds \geq \eta_*.$$

Both alternatives are impossible on the retained branch. Therefore the localized cost cannot be finite. $\square$

A.5. Localized cost and the Type II scale-exclusion condition.

Definition A.22 (Localized Type II cost). For a repaired-gauge Type II branch, let

$$\tilde{\mathfrak{D}}_{R_0}(\tau) = \nu \int_{\mathbb{R}^3} |\nabla V(y, \tau)|^2 \phi_{R_0}(y) dy + a_+(\tau) \int_{\mathbb{R}^3} |V(y, \tau)|^2 \phi_{R_0}(y) dy,$$

where $a_+ = \max(a, 0)$, ϕ_{R_0} is the compact Type II cutoff, and $\nu > 0$. A Type II scale-exclusion cost is a measurable nonnegative function $\mathfrak{C}_{\text{II}}$ that is locally integrable on finite τ -intervals and whose divergence on $[\tau_0, \infty)$ gives the Type II scale-exclusion condition.

Definition A.23 (Cost comparison). A cost comparison for a repaired-gauge Type II branch consists of $R_0, \phi_{R_0}, \mathfrak{C}_{\text{II}}, c_0, \tau_1$, with $c_0 > 0$ and $\tau_1 \geq \tau_0$, such that

$$\mathfrak{C}_{\text{II}}(\tau) \geq c_0 \tilde{\mathfrak{D}}_{R_0}(\tau) \quad \text{for a.e. } \tau \geq \tau_1.$$

The canonical Navier–Stokes choice is

$$\mathfrak{C}_{\text{II}}(\tau) = \tilde{\mathfrak{D}}_{R_0}(\tau), \quad c_0 = 1, \quad \tau_1 = \tau_0.$$

Lemma A.24 (Identity-cost comparison). *If $\tilde{\mathfrak{D}}_{R_0}$ is measurable, nonnegative, and locally integrable on finite τ -intervals, then the canonical choice $\mathfrak{C}_{\text{II}} = \tilde{\mathfrak{D}}_{R_0}$ defines a valid Type II scale-exclusion cost and satisfies the cost comparison with $c_0 = 1$ and $\tau_1 = \tau_0$.*

Proof. The assumed measurability, nonnegativity, and local finite-interval integrability give the cost properties in [definition A.22](#). The identity

$$\mathfrak{C}_{\text{II}} = \tilde{\mathfrak{D}}_{R_0}$$

gives the comparison in [definition A.23](#) with $c_0 = 1$ and $\tau_1 = \tau_0$. $\square$

Theorem A.25 (Cost comparison theorem). *Let (V, P, a, b) be a repaired-gauge Type II branch. Assume:*

- (i) $\tilde{\mathfrak{D}}_{R_0}$ is measurable, nonnegative, and locally integrable on finite τ -intervals;
- (ii)

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty;$$

- (iii) either the canonical cost $\mathfrak{C}_{\text{II}} = \tilde{\mathfrak{D}}_{R_0}$ is used, or a cost comparison as in [definition A.23](#) is available.

Then the Type II scale-exclusion condition associated with $\mathfrak{C}_{\text{II}}$ holds.

Proof. In the canonical case, [lemma A.24](#) supplies the cost comparison. In the general case, the comparison is assumed. Hence there are $c_0 > 0$ and $\tau_1 \geq \tau_0$ such that

$$\mathfrak{C}_{\text{II}}(\tau) \geq c_0 \tilde{\mathfrak{D}}_{R_0}(\tau)$$

for a.e. $\tau \geq \tau_1$. Since $\tilde{\mathfrak{D}}_{R_0} \geq 0$ and is locally integrable on finite intervals, divergence on $[\tau_0, \infty)$ implies

$$\int_{\tau_1}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty.$$

Therefore

$$\int_{\tau_1}^{\infty} \mathfrak{C}_{\text{II}}(\tau) d\tau \geq c_0 \int_{\tau_1}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty.$$

Since $\mathfrak{C}_{\text{II}}$ is nonnegative and locally integrable on finite intervals,

$$\int_{\tau_0}^{\infty} \mathfrak{C}_{\text{II}}(\tau) d\tau = \int_{\tau_0}^{\tau_1} \mathfrak{C}_{\text{II}}(\tau) d\tau + \int_{\tau_1}^{\infty} \mathfrak{C}_{\text{II}}(\tau) d\tau = \infty.$$

By definition of the Type II scale-exclusion cost, this divergence is the Type II scale-exclusion condition. $\square$

A.6. Local compact-cost classification.

Definition A.26 (Remaining compact cost-divergence tests). For a repaired-gauge Type II orbit, define the critical-tightness condition by

$$\forall \varepsilon > 0 \forall J_* \in \mathbb{R} \exists R_{\varepsilon, J_*} : \int_J \int_{|y| > R_{\varepsilon, J_*}} |V(y, \tau)|^3 dy d\tau < \varepsilon$$

for every selected terminal window $J \in [\tau_0, \infty)$ modeled on J_* in the sense of [definition 10.100](#). Failure of critical tightness is the condition

$$\exists \varepsilon_0 > 0 \exists J_* \in \mathbb{R} \exists J_n \in [\tau_0, \infty), R_n \rightarrow \infty : \int_{J_n} \int_{|y| > R_n} |V(y, \tau)|^3 dy d\tau \geq \varepsilon_0.$$

Here each J_n is modeled on the same fixed compact interval J_* . Define local windowed H^1 control by

$$\forall m \geq 1 : \sup_{n \in \mathbb{N}} \int_{\tau_0+n}^{\tau_0+n+1} \|V(\tau)\|_{H^1(B_m)}^2 d\tau < \infty.$$

Failure of local windowed H^1 control is the condition

$$\exists m \geq 1 : \sup_{n \in \mathbb{N}} \int_{\tau_0+n}^{\tau_0+n+1} \|V(\tau)\|_{H^1(B_m)}^2 d\tau = \infty.$$

Definition A.27 (Compact Type II data). A Type II branch has compact Type II data if it has:

- (i) the Type II analytic data of [definition A.2](#);
- (ii) the repaired-gauge representation of [definition A.9](#);
- (iii) retained positive finite local critical mass on every compact terminal window, in the sense of [lemma 10.78](#);
- (iv) the locally integrable nonnegative cost and cost comparison of [definitions A.22](#) and [A.23](#).

This condition does not include critical tightness or windowed H^1 control; those are the two remaining compact cost-divergence tests.

Definition A.28 (Routed pre-compact-data exits). A routed pre-compact-data exit is the first failed ordered test before the compact cost-divergence tests are reached. Thus it is one of the routed C1 exhaustion failures in [definition A.3](#), one of the ordered representation failures in [definition A.11](#) or [definition A.18](#), one of the local terminal critical-mass failures routed by [lemma 10.78](#), or a failure of local cost well-posedness or cost comparison before [definitions A.22](#) and [A.23](#) is available. The first-failure convention makes these exits disjoint from the compact Type II data of [definition A.27](#).

Lemma A.29 (Pre-compact-data routing). *Every admissible Type II branch has exactly one of the following ordered outcomes:*

- (i) *it has compact Type II data in the sense of [definition A.27](#);*
- (ii) *it has a routed pre-compact-data exit in the sense of [definition A.28](#).*

Proof. Fix an admissible branch. The ordered C1 routing hypothesis of [theorem A.4](#) is supplied by [theorem 10.122](#). Applying [theorem A.4](#), if the concentration, scale, and local profile data are not all present, the first failed C1 test is a routed pre-compact-data exit. Otherwise the first component of [definition A.27](#) holds.

Next run the ordered representation classification [corollary A.12](#), equivalently the minimal representation classification [corollary A.19](#) when the minimal chart package is used. A first representation, repaired-gauge, pressure, or modulation failure is a routed pre-compact-data exit. If no such failure occurs, the branch has the repaired-gauge representation of [definition A.9](#).

On that repaired-gauge branch, apply the local terminal critical-mass routing [lemma 10.78](#). Failure of the retained spacetime activity floor means the candidate was not a retained active core; an upper local L^3 -failure produces a comparable retained class or an adjacent active label. These are routed pre-compact-data exits unless they are grouped into the retained terminal component. In the retained case the positive finite local critical-mass row holds on every compact terminal window.

Finally test whether the localized cost is well posed and compatible with the Navier–Stokes identity cost. Availability of the local cost and comparison is the fourth component of [definition A.27](#). If it is absent, the first absence is the local cost-well-posedness or cost-comparison failure listed in [definition A.28](#); if it is present, all four components of compact Type II data hold. The first-failure convention makes the two outcomes disjoint and exhaustive. $\square$

Lemma A.30 (Uniform retained active-window mass floor). *Assume an admissible Type II branch has compact Type II data in the sense of [definition A.27](#) and remains in the retained compact cost-divergence stratum after the ordered validation tests of [lemma 8.73](#). Then there exist*

$$R_* > 0, \quad \theta > 0, \quad \eta_* > 0,$$

and an unbounded family $\mathcal{S} \subset [\tau_0, \infty)$ of retained selected active-window centers such that

$$\int_{\sigma-\theta}^{\sigma+\theta} \int_{B_{R_*}} |V(y, \tau)|^3 dy d\tau \geq \eta_* \quad \text{for every } \sigma \in \mathcal{S}.$$

If such a fixed retained active-window family does not exist, the branch has already left the retained compact stratum through a lower-mass, pressure-only, exterior/separated, cascade, or comparable-class grouping row.

Proof. The retained finite critical-mass component of [definition A.27](#) is read through [lemma 10.78](#). Thus the lower active mass row does not fail on the retained branch; if it did, the candidate would not be a retained terminal component.

Apply the CKN detection argument in [theorem 9.33](#) to the retained active component on the selected compact terminal cylinders. Its proof gives a fixed local threshold $\eta_{\text{det}} > 0$. If the velocity part of the detected CKN quantity were not bounded below on a compact subcylinder, then either the component would be perturbative by [corollary 9.32](#), or the detected mass would be pressure-only and would be routed by [lemma 9.22](#) to the pressure/exterior row. Neither row remains in the retained compact cost-divergence stratum. Therefore each retained active witness carries a positive spacetime velocity mass on one member of a finite compact covering of the selected terminal core.

The covering is finite and the terminal selection rule groups all comparable same-point retained classes by [lemma 10.77](#). If the active witnesses escaped every fixed member of this finite compact chart along the tail, the first failure would be an exterior/separated or cascade row in the ordered local state space. Since the branch remains retained, one fixed compact member of the cover occurs along an unbounded family of selected windows. Writing that member as $B_{R_*} \times (-\theta, \theta)$, and absorbing the scale-invariant covering constants into $\eta_* := c\eta_{\text{det}}R_*^2 > 0$, gives the displayed lower bound for an unbounded family $\mathcal{S}$ of retained active-window centers. $\square$

Lemma A.31 (Local upper estimates give retained Aubin–Lions compactness). *Let (V, P, a, b) be a retained repaired-gauge branch on a terminal tail. Fix a compact model interval $I \subseteq \mathbb{R}$, a radius $R > 0$, and selected centers $\sigma_j \rightarrow \infty$. Assume that on the shifted windows*

$$V_j(y, s) = V(y, \sigma_j + s), \quad P_j(y, s) = P(y, \sigma_j + s), \quad s \in I,$$

the ordered retained rows supply:

(i) a uniform local critical upper bound

$$\sup_j \|V_j\|_{L^\infty(I; L^3(B_{2R}))} < \infty;$$

(ii) local windowed H^1 control on B_{2R} ;

(iii) local pressure reconstruction modulo functions of time;

(iv) the repaired-gauge equation with uniformly bounded selected-window modulation integrals

$$\sup_j \int_I (|a(\sigma_j + s)| + |b(\sigma_j + s)|) ds < \infty.$$

Then, after subtracting spatial pressure means, the sequence satisfies

$$\sup_j \left(\|V_j\|_{L^\infty(I; L^2(B_{2R}))} + \|V_j\|_{L^2(I; H^1(B_{2R}))} + \|P_j - c_{j,2R}\|_{L^{3/2}(B_{2R} \times I)} \right) < \infty,$$

and for some large m ,

$$\sup_j \|\partial_s V_j\|_{L^1(I; H^{-m}(B_R))} < \infty.$$

Consequently, after passing to a subsequence,

$$V_j \rightarrow V_\infty \quad \text{strongly in } L^3(B_R \times I).$$

Proof. The $L_s^\infty L_y^2(B_{2R})$ bound follows from the local $L_s^\infty L_y^3(B_{2R})$ bound and finite measure:

$$\|V_j(s)\|_{L^2(B_{2R})} \leq |B_{2R}|^{1/6} \|V_j(s)\|_{L^3(B_{2R})}.$$

The $L_s^2 H_y^1(B_{2R})$ bound is the local windowed H^1 control; a fixed compact model interval I meets only finitely many unit windows, so the uniform unit-window bound gives a uniform bound on I . The pressure bound is the local Calderón–Zygmund pressure reconstruction, modulo functions of time, applied to $V_j \otimes V_j$; the finite length of I and the uniform $L_s^\infty L_y^3$ bound give the displayed $L_{s,y}^{3/2}$ estimate.

It remains to bound the time derivative. On B_R the repaired-gauge equation with

$$a_j(s) := a(\sigma_j + s), \quad b_j(s) := b(\sigma_j + s)$$

gives

$$\partial_s V_j = \nu \Delta V_j - \nabla \cdot (V_j \otimes V_j) - \nabla P_j + a_j(V_j + y \cdot \nabla V_j) + b_j \cdot \nabla V_j.$$

Choose m so large that

$$L^1(B_{2R}) + L^{3/2}(B_{2R}) + H^{-1}(B_{2R}) \hookrightarrow H^{-m}(B_R).$$

The diffusion term is bounded in $L^1(I; H^{-m})$ by the $L^2 H^1$ bound. The nonlinear term is bounded because

$$\|V_j \otimes V_j\|_{L^1(I; L^{3/2}(B_{2R}))} \leq |I| \|V_j\|_{L^\infty(I; L^3(B_{2R}))}^2.$$

The pressure-gradient term is bounded by the pressure estimate. For the modulation terms, integrate by parts in the distributional pairing:

$$y \cdot \nabla V_j = \nabla \cdot (y \otimes V_j) - 3V_j, \quad b_j \cdot \nabla V_j = \nabla \cdot (b_j \otimes V_j).$$

Thus their $H^{-m}(B_R)$ norms are bounded by

$$C_R (|a_j(s)| + |b_j(s)|) \|V_j(s)\|_{L^2(B_{2R})}.$$

The retained modulation row gives $a_j, b_j \in L^1(I)$ with uniform selected window bounds; if this row failed, the ordered validation would route the branch to the modulation/scale-drift exit before the retained compact stratum was declared. Hence $\partial_s V_j$ is bounded in $L^1(I; H^{-m}(B_R))$.

The Aubin–Lions–Simon theorem applied to

$$H^1(B_R) \Subset L^2(B_R) \hookrightarrow H^{-m}(B_R)$$

gives strong convergence in $L^2(B_R \times I)$, after passing to a subsequence. The same bounds give a uniform $L^{10/3}(B_R \times I)$ bound by the standard parabolic interpolation between $L^\infty L^2$ and $L^2 H^1$. Interpolating the strong L^2 convergence with this uniform $L^{10/3}$ -bound gives strong convergence in L^p for every $2 \leq p < 10/3$, in particular in $L^3(B_R \times I)$. $\square$

Lemma A.32 (Retained compact data supply the retained mass and compactness package). *Assume an admissible Type II branch has compact Type II data in the sense of [definition A.27](#), satisfies the two compact tests of [definition A.26](#), and remains in the retained compact cost-divergence stratum after the ordered validation tests of [lemma 8.73](#). Then, on every compact retained terminal window after a harmless terminal restriction, the requirements of [theorem A.21](#) are satisfied. More precisely:*

- (i) *there are $R_* > 0$, $\theta > 0$, $\eta_* > 0$, and an unbounded retained selected active-window family $\mathcal{S}$ satisfying the positive spacetime local mass floor of [theorem A.21](#);*
- (ii) *the local compact-cylinder bounds*

$$V \in L^\infty(I; L^3(B_{2R})) \cap L^\infty(I; L^2(B_{2R})) \cap L^2(I; H^1(B_{2R})), \quad P - c_{2R} \in L^{3/2}(B_{2R} \times I)$$
hold for each fixed R and compact terminal interval I ;
- (iii) *on smaller balls, $\partial_\tau V$ is bounded in $L^1(I; H^{-m}(B_R))$ for some large m , hence the local Aubin–Lions compactness used in [theorem A.21](#) is available;*
- (iv) *any nonzero bounded selected-window limit is routed out of the retained Type II branch by [theorem 5.27](#).*

If any item in this list fails before the retained stratum is reached, the first failure is one of the named representation, pressure, modulation, critical-mass, rough-core, exterior, or adjacent compact-cost exits of [lemma 8.73](#); it is not repaired by inserting a global estimate.

Proof. The positive spacetime retained-window mass floor is [lemma A.30](#). Its proof uses the active CKN detection threshold, pressure-only routing, and the terminal maximality rule; if the fixed unbounded active-window family does not exist, the branch has already left the retained compact stratum through a named first-failure row.

The repaired-gauge representation, pressure reconstruction, and modulation data are item (ii) of [definition A.27](#) together with the passed representation rows of [lemma 8.73](#). The two compact tests give critical tightness and local windowed H^1 control. Equivalently, by [theorem 10.108](#) and [corollary 10.109](#), their failure is a named non-retained alternative. Combining the retained local L^3 upper bound, local windowed H^1 control, local pressure reconstruction, and the retained modulation-window bound of [lemma 10.106](#), [lemma A.31](#) gives the displayed local compact-cylinder bounds, the $L^1 H^{-m}$ time-derivative bound, and the strong L^3 Aubin–Lions compactness on selected windows. This compactness lemma uses only local upper bounds and local PDE estimates; it does not require the pointwise-in-time lower bound in [theorem 10.131](#).

Finally, a nonzero bounded selected-window limit leaves the strict selected-window retained branch by [theorem 5.27](#); on a physical germ it is then closed by [theorem 11.16](#). Hence all hypotheses used in [theorem A.21](#) are local consequences of the retained compact data and validation rows. $\square$

Lemma A.33 (Compact positive case is excluded). *Assume an admissible Type II branch has compact Type II data and satisfies both compact cost-divergence tests in [definition A.26](#). If the branch remains in the retained compact cost-divergence stratum after the ordered validation tests of [lemma 8.73](#), then the branch satisfies the Type II exclusion conclusion.*

Proof. The compact Type II data give the supercritical scale alternative, a repaired-gauge orbit (V, P, a, b) , retained positive local critical mass, local compact-cylinder critical upper bounds, and the cost comparison. By [lemma A.32](#), the retained branch also has the local compact-window, pressure, H^1 , time-derivative, and bounded-limit routing package required by [theorem A.21](#). Thus

$$\int_{\tau_0}^{\infty} \tilde{\mathfrak{D}}_{R_0}(\tau) d\tau = \infty.$$

[Theorem A.25](#) converts this divergence into the Type II scale-exclusion condition, hence the validated compact cost-divergence conclusion. Since the branch remains in the retained compact cost-divergence stratum, [corollary 8.82](#) gives the Type II exclusion conclusion. $\square$

REFERENCES

- [Aub63] Jean-Pierre Aubin. Un théorème de compacité. *Comptes Rendus de l'Académie des Sciences de Paris*, 256:5042–5044, 1963.
- [BL83] Haïm Brezis and Elliott H. Lieb. A relation between pointwise convergence of functions and convergence of functionals. *Proceedings of the American Mathematical Society*, 88(3):486–490, 1983.
- [CF88] Peter Constantin and Ciprian Foias. *Navier–Stokes Equations*. University of Chicago Press, Chicago, 1988.
- [CKN82] Luis Caffarelli, Robert Kohn, and Louis Nirenberg. Partial regularity of suitable weak solutions of the navier–stokes equations. *Communications on Pure and Applied Mathematics*, 35(6):771–831, 1982.
- [CZ52] Alberto P. Calderon and Antoni Zygmund. On the existence of certain singular integrals. *Acta Mathematica*, 88:85–139, 1952.
- [ESS03] Luis Escauriaza, Gregory Seregin, and Vladimir Sverak. $L_{3,\infty}$ -solutions of navier–stokes equations and backward uniqueness. *Russian Mathematical Surveys*, 58(2):211–250, 2003.
- [Gal01] Isabelle Gallagher. Profile decomposition for solutions of the navier–stokes equations. *Bulletin de la Société Mathématique de France*, 129(2):285–316, 2001.
- [Kat84] Tosio Kato. Strong L^p -solutions of the navier–stokes equation in $\mathbb{R}^m$, with applications to weak solutions. *Mathematische Zeitschrift*, 187:471–480, 1984.
- [Lad69] Olga A. Ladyzhenskaya. *The Mathematical Theory of Viscous Incompressible Flow*. Gordon and Breach, New York, 1969.
- [Ler34] Jean Leray. Sur le mouvement d’un liquide visqueux emplissant l’espace. *Acta Mathematica*, 63:193–248, 1934.
- [Lin98] Fanghua Lin. A new proof of the caffarelli-kohn-nirenberg theorem. *Communications on Pure and Applied Mathematics*, 51(3):241–257, 1998.
- [LM72] Jacques-Louis Lions and Enrico Magenes. *Non-Homogeneous Boundary Value Problems and Applications. Vol. I*. Springer, Berlin, 1972.
- [LR02] Pierre Gilles Lemarié-Rieusset. *Recent Developments in the Navier–Stokes Problem*. Chapman and Hall/CRC, Boca Raton, FL, 2002.
- [NRŠ96] Jindřich Nečas, Michael Růžička, and Vladimír Šverák. On leray’s self-similar solutions of the navier–stokes equations. *Acta Mathematica*, 176(2):283–294, 1996.
- [Sch76] Vladimir Scheffer. Partial regularity of solutions to the navier–stokes equations. *Pacific Journal of Mathematics*, 66(2):535–552, 1976.
- [Sim87] Jacques Simon. Compact sets in the space $L^p(0, T; B)$. *Annali di Matematica Pura ed Applicata*, 146:65–96, 1987.
- [Ste70] Elias M. Stein. *Singular Integrals and Differentiability Properties of Functions*. Princeton University Press, Princeton, NJ, 1970.
- [Tem77] Roger Temam. *Navier–Stokes Equations: Theory and Numerical Analysis*. North-Holland, Amsterdam, 1977.

Email address: guillem@fragile.tech